

Teaching an old DAG new tricks

...

Migrating a decade old pipeline to Airflow

Outline

Cloud native deployment

- Cloud native deployment
- Multi-repo DAG management
- Manage Airflow Variables with code through Terraform
- Airflow monitoring best practices with Datadog and Pagerduty

Airflow Migration

- Simulate production run to surface issues early
- Plan and execute with incremental deliverables

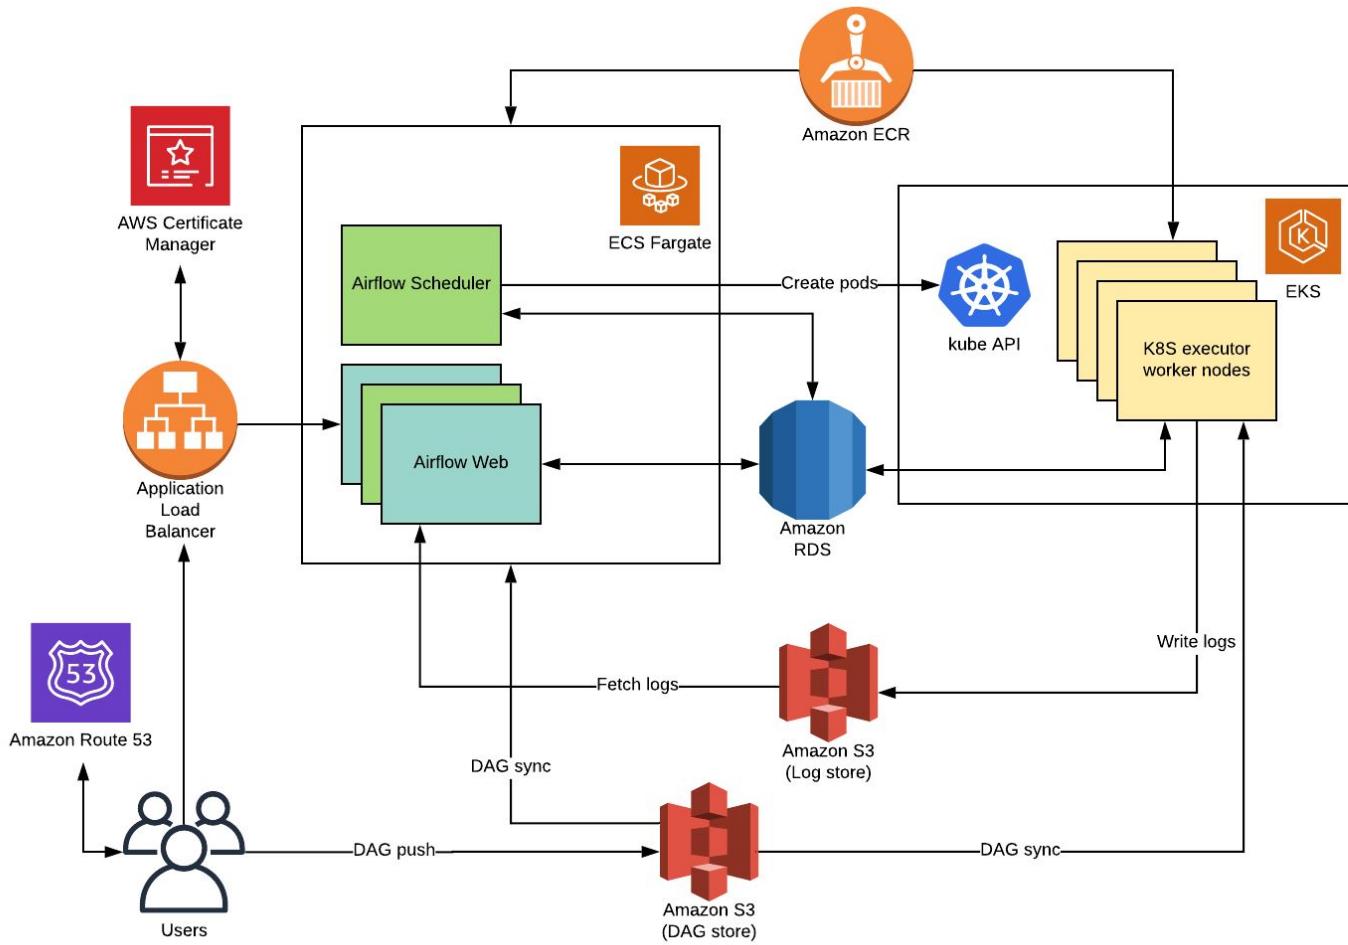


Scribd is moving to the cloud

<https://tech.scribd.com/blog/2019/building-the-library.html>

Cloud native Airflow

- Use managed service whenever possible
- Separation of stateless compute and stateful data store
- Separation of infrastructure (Airflow cluster) and application (DAG)
- Separation of environments
- Automate Infrastructure provisioning with code
- Running on development branch of Airflow for latest improvements and bug fixes

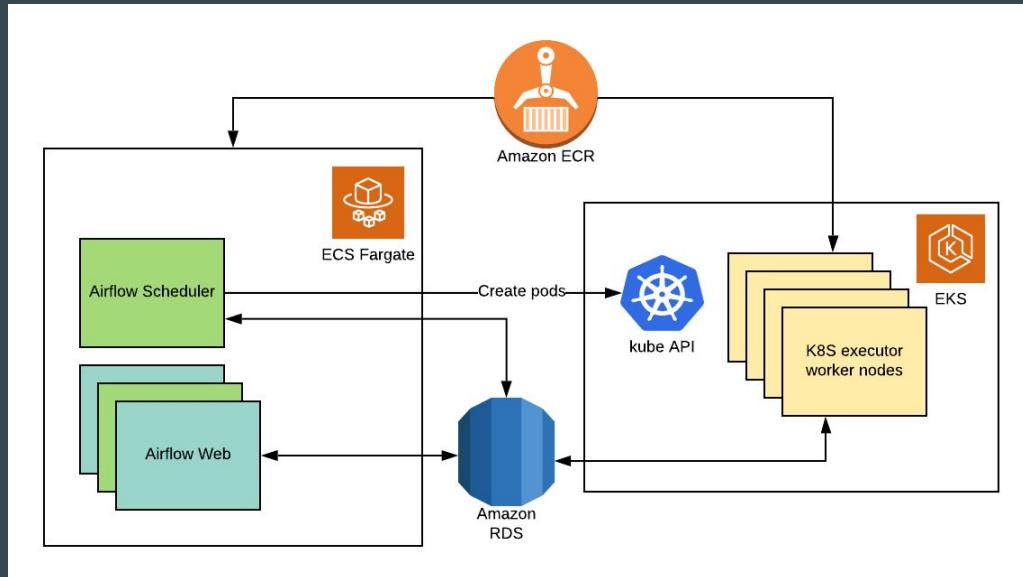


PagerDuty



ECS and EKS?!

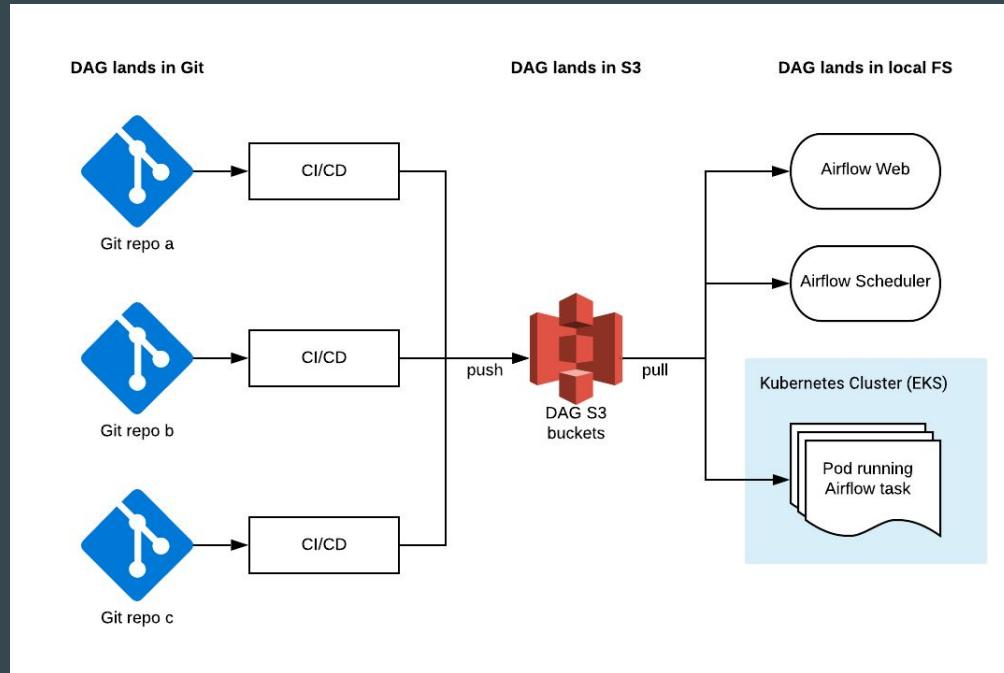
- Different crash zones
- Reduce maintenance burden with ECS fargate



Out of cluster Kubernetes executor support for EKS

- Kubernetes Python client doesn't work well with EKS
- API token generated by aws-iam-authenticator expires about every 14 minutes
- Python client fix backported to Airflow:
<https://github.com/apache/airflow/pull/5731>

Develop DAGs across multiple repos



<https://tech.scribd.com/blog/2020/breaking-up-the-dag-repo.html>

DAG sync daemon

- Background daemon written in Golang with small CPU and memory footprint
- Single binary ready to run in any environment
- File list and checksums are cached in memory to minimize network and disk IO
- DAG release gets picked up within seconds
 - Future plan to use S3 event notification to make it near realtime
- Expose operational metrics as prometheus format through HTTP
 - DAG Update/Delete/Create statistics
 - Time spent on DAG sync
 - Daemon uptime

Project Github: <https://github.com/scribd/objinsync>

Manage Variables with Terraform

We use variables to templatize a lot of things



- IAM roles for Databricks clusters
- Glue catalog id
- EC2 Instance profile ARN
- Application Jar release version
- ...

```
{"assume_role_arn":"arn:aws:iam::1234567:role/automated-job-role","glue_catalogid":"2234567","instance_profile_arn":"arn:aws:iam::3234567:instance-profile/foo","instance_profile_arn":"arn:aws:iam::4234567:instance-profile/databricks-jobs-dev-profile"}
```

Airflow Terraform Provider

```
locals {
  team_a_remote_state = data.terraform_remote_state.team_a.outputs
  dev_vars = {
    "team_a_cluster" = {
      "assume_role_arn"      = local.team_a_remote_state.databricks_role_arn,
      "glue_catalogid"       = local.team_a_remote_state.glue_catalogid,
      "instance_profile_arn" = local.team_a_remote_state.databricks_instance_profile_arn,
    }
  }
}

provider "airflow" {
  variables_output_path = "development.variables.json"
}

resource "airflow_variable" "development" {
  for_each = local.dev_vars
  key      = each.key
  value    = jsonencode(each.value)
}
```

Airflow Terraform Provider

- Project Github: <https://github.com/houqp/terraform-provider-airflow>
- Experimental branch using Airflow Go client:
 - <https://github.com/houqp/terraform-provider-airflow/tree/openapi>
 - <https://github.com/apache/airflow-client-go/pull/1>

Monitor Airflow with Datadog

Datadog agent as sidecar container within ECS

```
datadog_container = {
    name      = "datadog-agent",
    image     = "datadog/agent:latest",
    essential = true,
    environment = [
        { name = "DD_API_KEY",
          value = var.datadog_api_key
        },
        { name = "DD_TAGS",
          value = "env:${local.env} application:airflow"
        },
        { name = "DD_DOGSTATSD_TAGS",
          value = "[\"env:${local.env}\", \"application:airflow\"]"
        },
        { name = "ECS_FARGATE",
          value = "true"
        },
    ],
    portMappings = [
        { protocol = "tcp"
          containerPort = 8125
        }
    ]
}
```

Statsd config for scheduler

```
# monitoring
{ name   = "AIRFLOW__SCHEDULER__STATSD_ON"
  value = "True"
},
{ name   = "AIRFLOW__SCHEDULER__STATSD_HOST"
  value = "127.0.0.1"
},
{ name   = "AIRFLOW__SCHEDULER__STATSD_PORT"
  value = "8125"
},
{ name   = "AIRFLOW__SCHEDULER__STATSD_PREFIX"
  value = "airflow"
},
```



Monitor Airflow with Datadog

- Synchronize ALB, RDS, S3, ECS and EKS Cloudwatch metrics to Datadog using Terraform (<https://github.com/scribd/terraform-aws-datadog>)

The screenshot shows the HashiCorp Terraform Registry interface. At the top, there's a navigation bar with the Terraform logo, a search bar labeled "Search for modules", and links for "Browse", "Publish", and "Sign-in". Below the header, the main content area displays a module card for "datadog" by AWS. The card features the AWS logo, the module name "datadog", and a brief description: "Terraform module for setting up AWS Datadog integration". A dropdown menu indicates the version is "Version 1.2.0". To the right, a box titled "Provision Instructions" contains sample Terraform code for integrating with Datadog. The code includes a module declaration, a source URL pointing to the GitHub repository, a specific version number, and a note about inserting a required variable.

datadog
AWS

Terraform module for setting up AWS Datadog integration

Version 1.2.0 ▾

Provision Instructions

```
module "datadog" {  
  source = "scribd/datadog/aws"  
  version = "1.2.0"  
  # insert the 1 required variable here  
}
```

Published July 1, 2020 by scribd
Module managed by jim80net
Total provisions: 1,247
Source Code: github.com/scribd/terraform-aws-datadog (report an issue)

Examples ▾

Search... | dev

\$env dev



> Overview 18

> Scheduler 19

> Web 11

> RDS 28

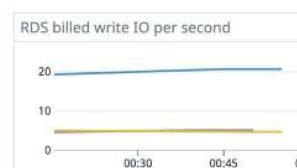
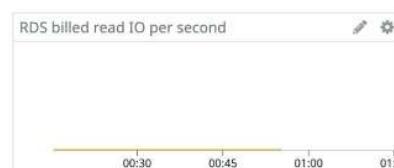
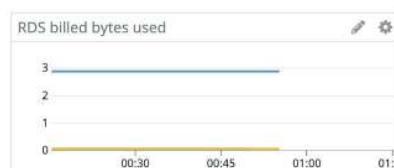
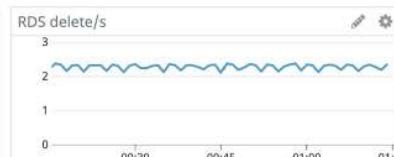
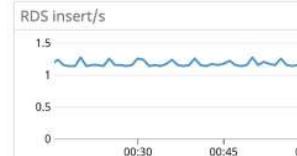
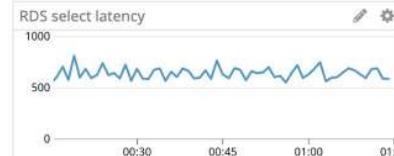
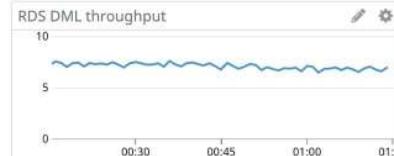
> DAG sync 6

> ECS 4

> EKS 13

> S3 2

🔍 Search... dev \$env dev ⚙️



> DAG sync 6

> ECS 4

> EKS 13

> S3 2

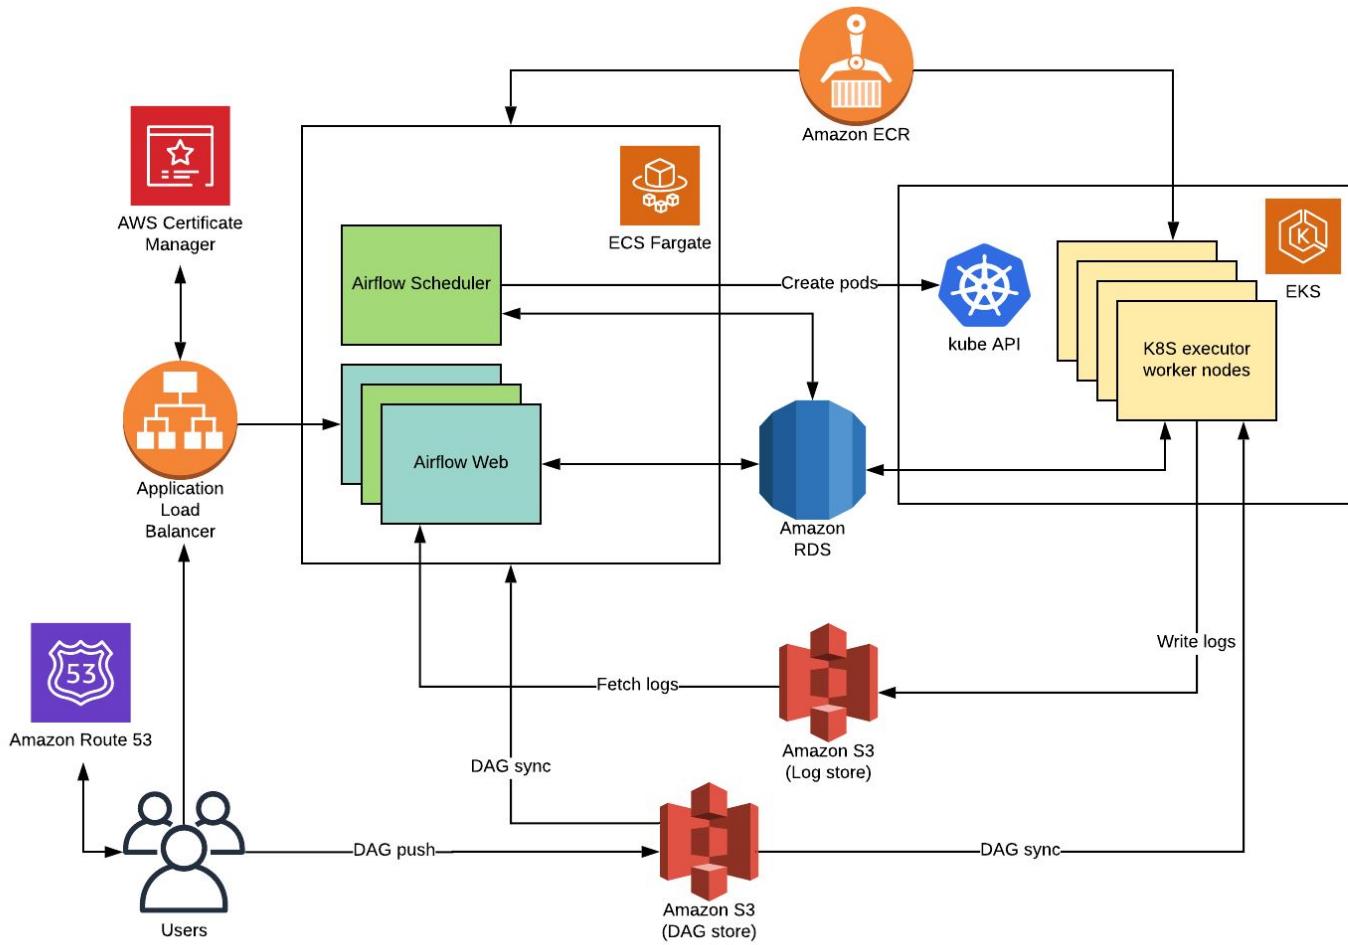
Incident response with Pagerduty



- Paging for infrastructure incidents
 - Through Datadog monitors
- Paging for application incidents
 - Pagerduty event emitted from Airflow for
 - Task failures
 - SLA misses
 - Adhoc events

Integration with Pagerduty

	<p>When all conditions are met</p> <ul style="list-style-type: none">• <code>payload.group</code> contains <code>pdservice:airflow</code>• <code>payload.group</code> contains <code>env:prod</code>	<ul style="list-style-type: none">• Route to Airflow• Then stop processing
	<p>When all conditions are met</p> <ul style="list-style-type: none">• <code>payload.group</code> contains <code>env:preprod</code>• <code>payload.group</code> contains <code>pdservice:airflow</code>	<ul style="list-style-type: none">• Route to Airflow• Set severity to warning• Then stop processing
	<p>When all conditions are met</p> <ul style="list-style-type: none">• <code>payload.group</code> contains <code>env:dev</code>• <code>payload.group</code> contains <code>pdservice:airflow</code>	<ul style="list-style-type: none">• Route to Airflow• Set severity to warning• Then stop processing
	<p>When all conditions are met</p> <ul style="list-style-type: none">• <code>payload.source</code> equals <code>dag_nightly</code>	<ul style="list-style-type: none">• Route to Nightly• Then stop processing



PagerDuty



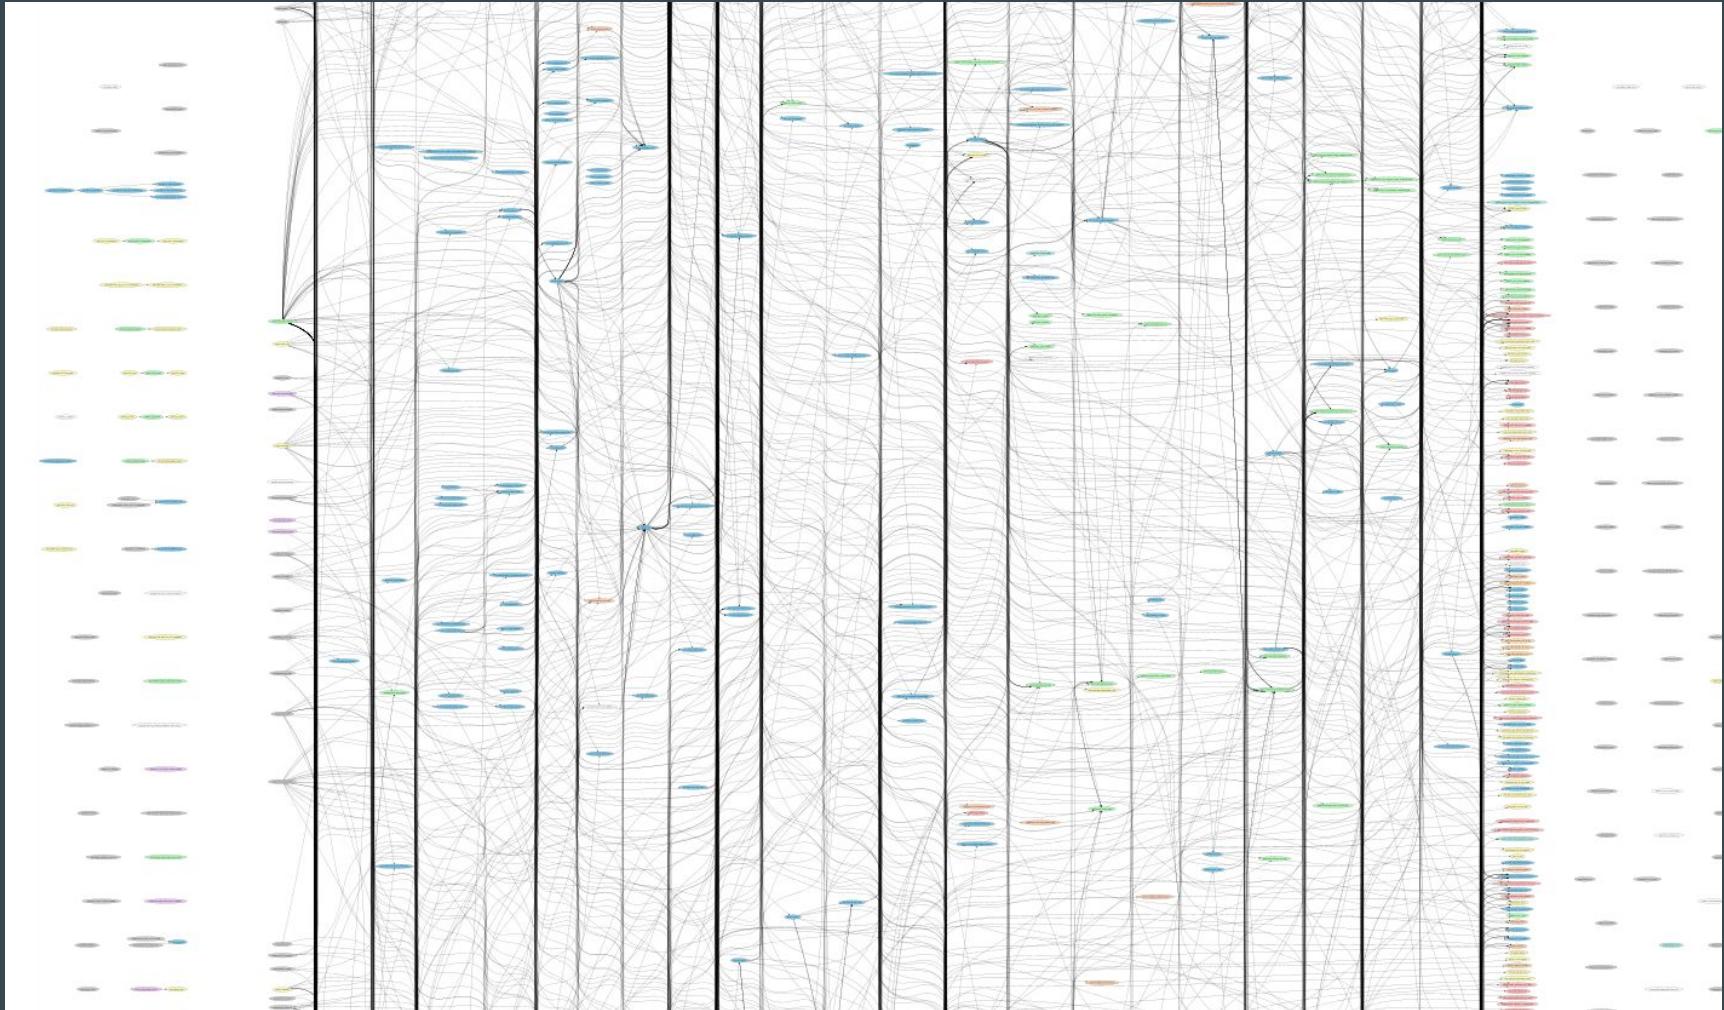
Migration

A decade old data pipeline

- In house workflow orchestration system called Datapipe
- First commit dates back to 2010
- 1500+ tasks with 1200+ of them in a single DAG
- Depend on features not supported by Airflow out of the box
- Data storage: HDFS, S3, Kafka, MySQL, Redis, ES
- Compute: Hive, Implala, Spark 1, Spark 2, Ruby

A brave new world

- Orchestrated through Airflow
- Data storage: S3 with Delta lake, Kafka, RDS, ElasticCache
- Compute: Spark 3 (Databricks)



Simulate production run early

- Automation to transpile Ruby DSL to Airflow DAG
 - Each task is a dummy operator that sleeps to simulate a run
 - Task sleep time calculated based off Avg runtime recorded by in-house system
- Scheduler was able to handle this DAG out of the box



503

Service Unavailable

How to render a 1500+ tasks DAG in Airflow

- It takes a long time to generate and render a 100MB page (tree view)
- Optimizations:
 - Avoid serialize the whole ORM object
 - Remove unnecessary if statements
 - Serialize JSON as string to be parsed with JSON.parse in the frontend
 - ...
 - <https://github.com/apache/airflow/pull/7492>
- Reduced page size by more than 10X
- Improved page load time by 5X



To the cloud, with incremental deliverables

- Incremental daily sync for new data lake in S3
 - Wrote a mini Python parser in Ruby
- Move ad-hoc read-only interactive queries
- Trim the dependency graph
- Move output phase of the pipeline to unblock external services
- Move remaining of the pipeline

About me (QP Hou)

Engineer at Scribd's Core Platform team

New Airflow committer

Maintainer and contributor of many other open-source projects

You can find me at:

- Airflow slack and mailing list
- <https://about.houqp.me>

Closing

- Truly a team effort within different engineering teams at Scribd
 - Driven by Platform Engineering
 - Core platform team
 - Data engineering team
- Embrace the open-source community
 - 41 PRs merged into upstream Airflow, many more to come
- Openings: <https://www.scribd.com/about/engineering>