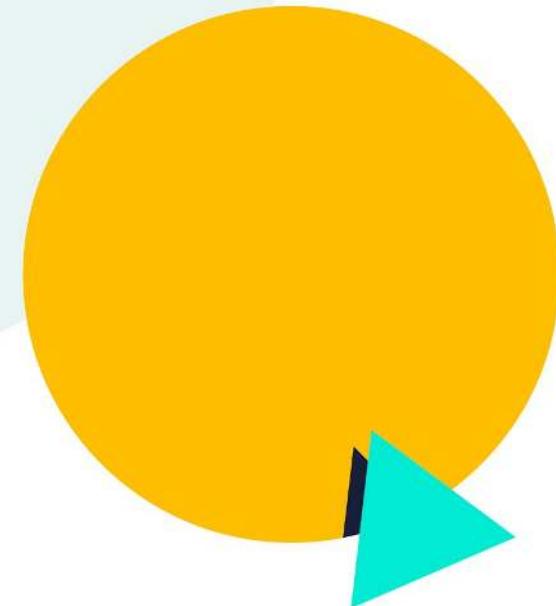




# Reasoning Reliability in Wrike's Data Pipeline



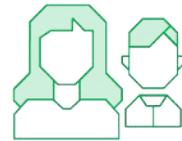
# Wrike - A Collaborative Work Management Platform



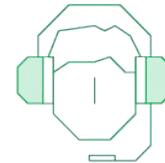
Founded in  
2006



10 Offices  
Globally



20,000+  
Customers  
Globally



1000+  
Employees



5 years in  
the Fast 500

# 20,000+

Organizations choose Wrike to  
orchestrate their digital work

With an additional 35,000 starting  
trials each month

- 2M users
- 130+ countries
- 10 languages
- 100M+ completed tasks



iOS and Android app icon for Beautyfilter

Beautyfilter Design Current Sprint

New Coordinator Anna by Anna on Nov 14

Set a date 0:00 Add subtask 1 field Attach files 18

**What needs to be designed? Icon**

**For what purpose it needs to be designed?** Icon for mobile application – Beautyfilter

**Do you have a reference idea on how it should look?** Crown/Magic Stick/Mirror

Coordinator Anna @Designer Matt Can you please assist with the design of the icon please?

Designer Matt @Coordinator Anna Got the first draft here, let me know what you think.

Backend updates

Product Current Week

Active

Default Workflow

Active

Cartification and Compliance  
Community Feedback  
Content Workflow  
CSO Operations  
Customer Education Content  
Customer Interview Workflow  
Data Engineering Workflow  
Data Subject Request  
Deployment  
Doc + Comm Workflow  
Event Workflow  
Finance Workflow

In Progress  
Completed  
Deferred  
Cancelled  
Duplicate

Change Task Workflow

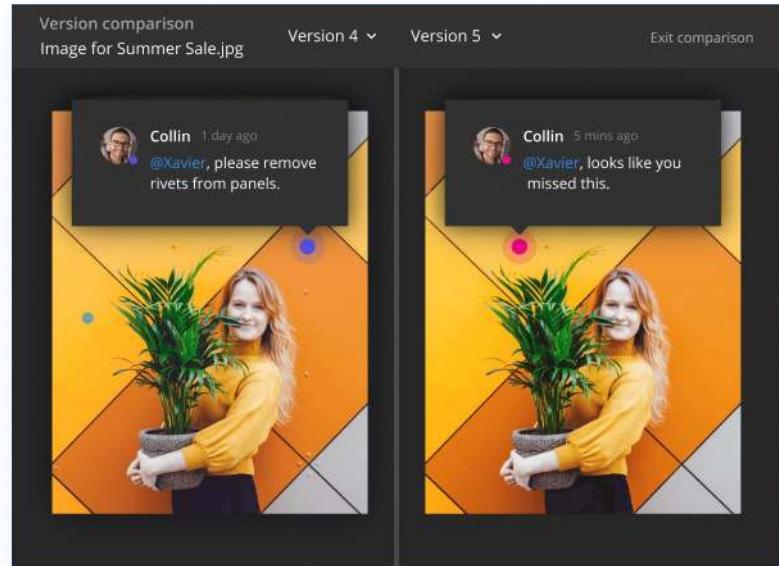
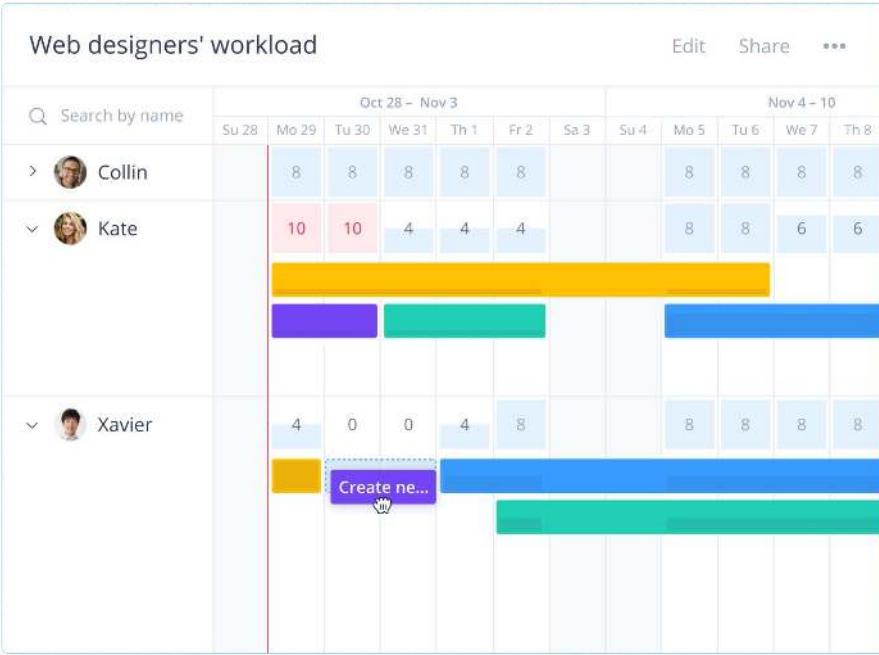
New  
Backlog  
In Queue  
Blocked  
In Progress  
Testing  
Technical review  
Ready for Deploy  
Completed  
Deferred  
Cancelled

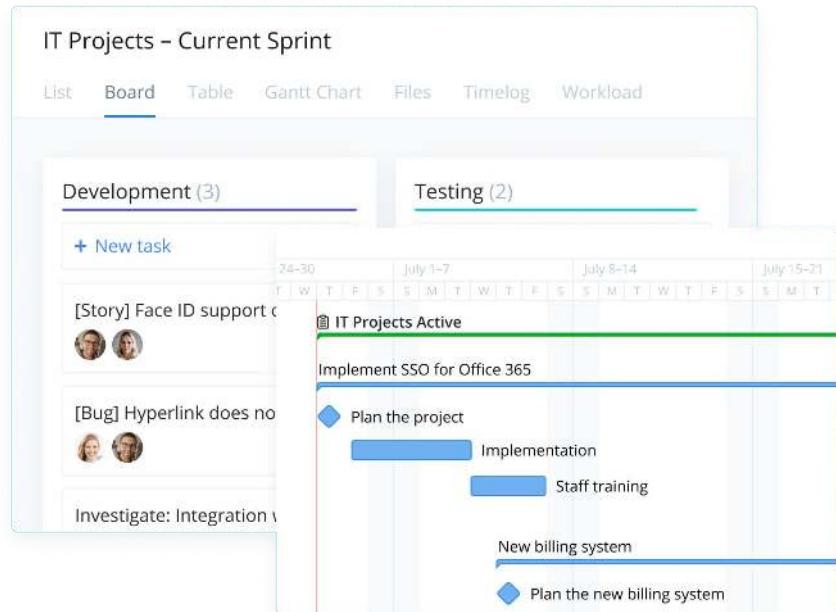
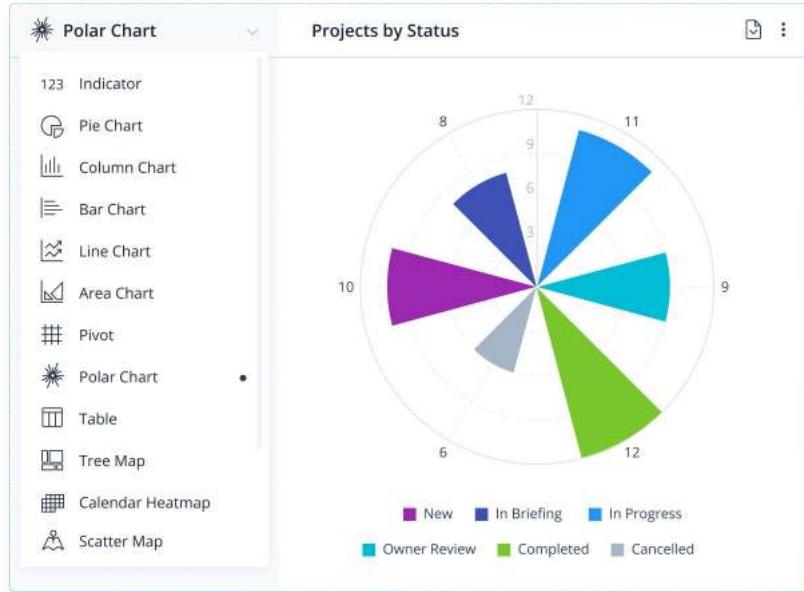
by Ashley S. on Nov 14

Attach files Add dependency 18

Add comment

Amanda





★ Featured

New/updated project in Wrike  
will create/update opportunity  
in Salesforce



Wrike  
Integrate

★ Featured

New Project created in Wrike gets  
assigned Custom Job Number



Wrike  
Integrate

★ Featured

New/Updated Task in Wrike  
creates an Event in Outlook  
Calendar



Wrike  
Integrate

★ Featured

Completed project in Wrike  
creates invoice in Quickbooks



Wrike  
Integrate

## Clients

LIST BOARD TABLE MORE ▾

+ New task

Emily Moore - Wrike

⋮ ⌂ ...

Active



Emily Moore - Wrike Active 1 field Attach files Shared with 1 group

Jessica Brown - Brig... Active Client's website: [wrike.com](http://wrike.com)

Paul Jones - Unity Inc. Active Account notes: Senior Corporate Communications Manager

Margaret Jenninston... Paused Contact information: [emily.team.wrike.com](mailto:emily.team.wrike.com)

Jason Reed - Clever.ly Active

Add comment

⋯ @ ⌂ Aa



Intro (0 out of 3)

# Data Engineering in Wrike

- SaaS means that we
  - Create
  - Support
  - Sell our product, and
  - Attract leads
- Help these teams speak the language of data
- We've got big space for data democratization



# Data Engineering Team in Wrike

- 16 data engineers in 4 teams
- We're supporting 250+ DAGs on production
- Up to 1200 tasks
- With median of 13 tasks
- ~10 updates of production or acceptance each day
- Helped 5 other teams to start using Airflow
- ~10-15% of our colleagues are using data engineering infrastructure and sources every month directly (>50% are using analytical reports or through integrations)



# We've Started With

- First analysts using new Data Warehouse based on Google BigQuery
- Data provided by a single instance of Airflow
  - A lot of bugs found on production data
  - A lot of changes during review
  - A lot of delays in data
  - Partially available data
  - Lack of the full picture during code review and architecture problems
- And we wanted to start democratization
  - Reliable production
  - No changes on production, at least unexpected ones
    - No changes in Data Structure
    - No changes in Data Freshness

# Acceptance Could Help

## DEV

- Quickly run your pipeline on a very small subset of your data
- In our case 0.0025% of all data
- Nothing will make sense, but it's a nice integration test

## TST

- Select a subset of your data for data that you know
- Immediately see if something is off
- Still quick to run

## ACC

- Carbon copy of production
- You can check if you feel comfortable pushing to PRD
- Give access to a Product Owner for them to check

## PRD

- Greenlight procedure for merging from ACC to PRD
- Manual operation
- Great for git blame

[Via Data's Inferno](#) by Wholesale Banking Advanced Analytics

# Acceptance Environment

- Acceptance is an environment where changes are welcome
- To make sure that we aren't going to need them on production

# No Changes on Production, at Least Unexpected Ones

- No Changes in Data Structure
- No Changes in Data Freshness
- No Changes during release from Acceptance to Production



# No Changes in Data Structure



# Implementation of Acceptance

## DEV

- Quickly run your pipeline on a very small subset of your data
- In our case 0.0025% of all data
- Nothing will make sense, but it's a nice integration test

## TST

- Select a subset of your data for data that you know
- Immediately see if something is off
- Still quick to run

## ACC

- Carbon copy of production
- You can check if you feel comfortable pushing to PRD
- Give access to a Product Owner for them to check

## PRD

- Greenlight procedure for merging from ACC to PRD
- Manual operation
- Great for git blame

[Via Data's Inferno](#) by Wholesale Banking Advanced Analytics



No Changes in Data Structure (1 out of 3)

# Acceptance on DB Side. BigQuery

- Acceptance and production are **different projects** in the notation of BigQuery
- Isolated quotas and limits (resources)
- BigQuery allows for cross-project queries
  - So we store on acceptance only changed data
  - And take source data from production.

Resources + ADD DATA

product\_aggregations

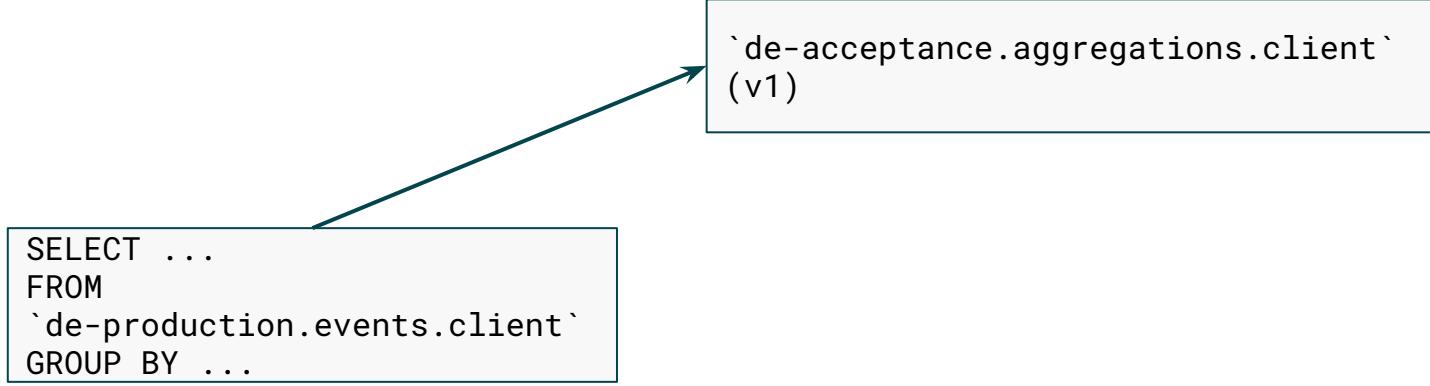
de-acceptance

- product\_aggregations
- stage\_product\_aggregations
  - accs\_adoption\_wau
  - client\_events\_schema\_st...

de-production

- product\_aggregations
- stage\_product\_aggregations
  - accs\_adoption\_wau

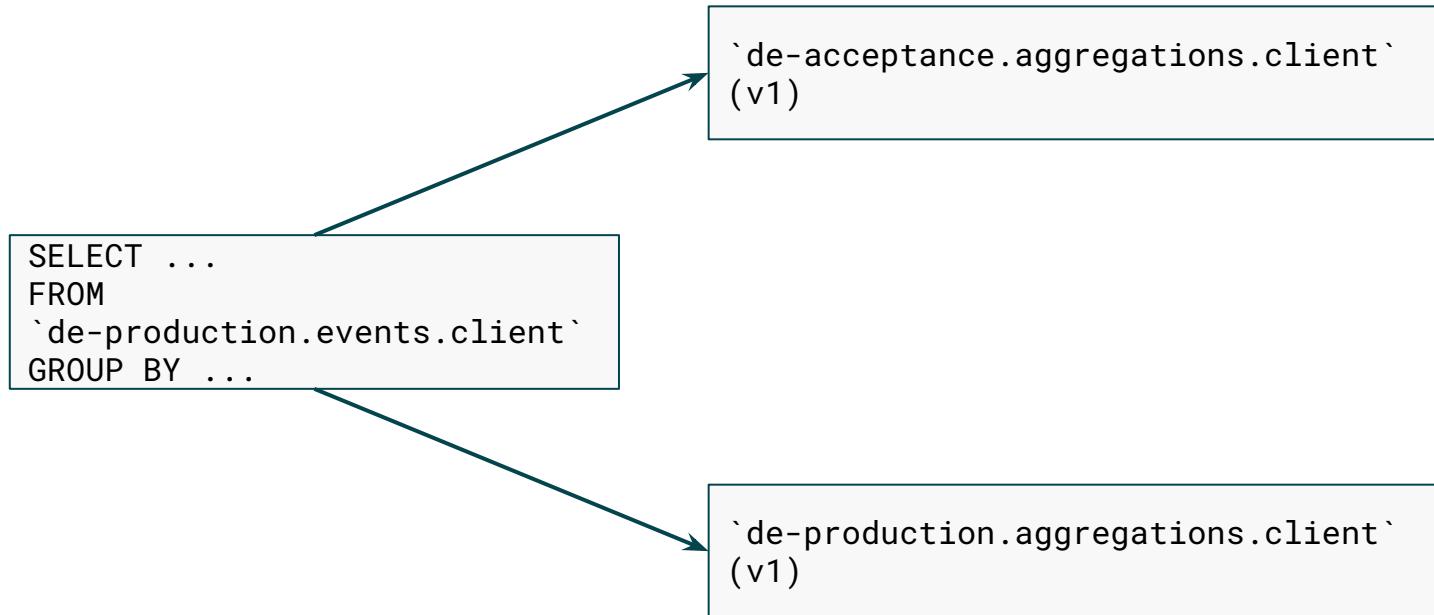
# Dataflow Example



No Changes in Data Structure (1 out of 3)

17

# Dataflow Example



No Changes in Data Structure (1 out of 3)

# Dataflow Example

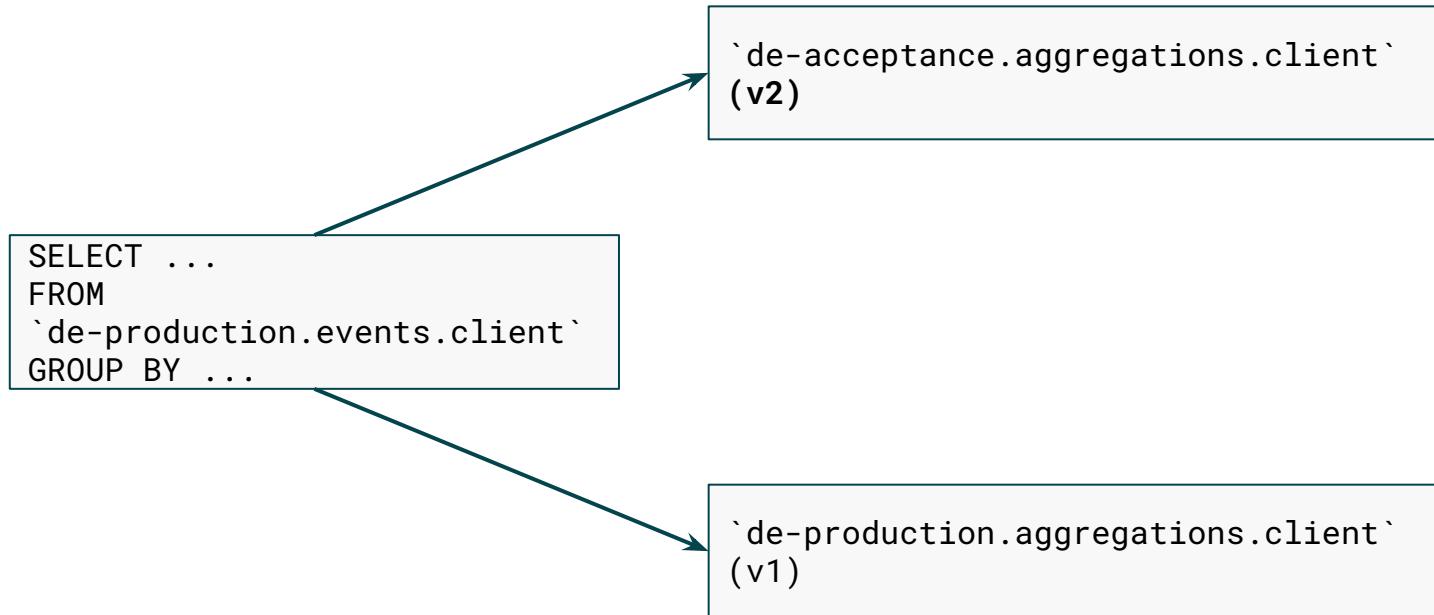
```
SELECT ...  
FROM  
`de-production.events.client`  
GROUP BY ...
```

```
`de-production.aggregations.client`  
(v1)
```



No Changes in Data Structure (1 out of 3)

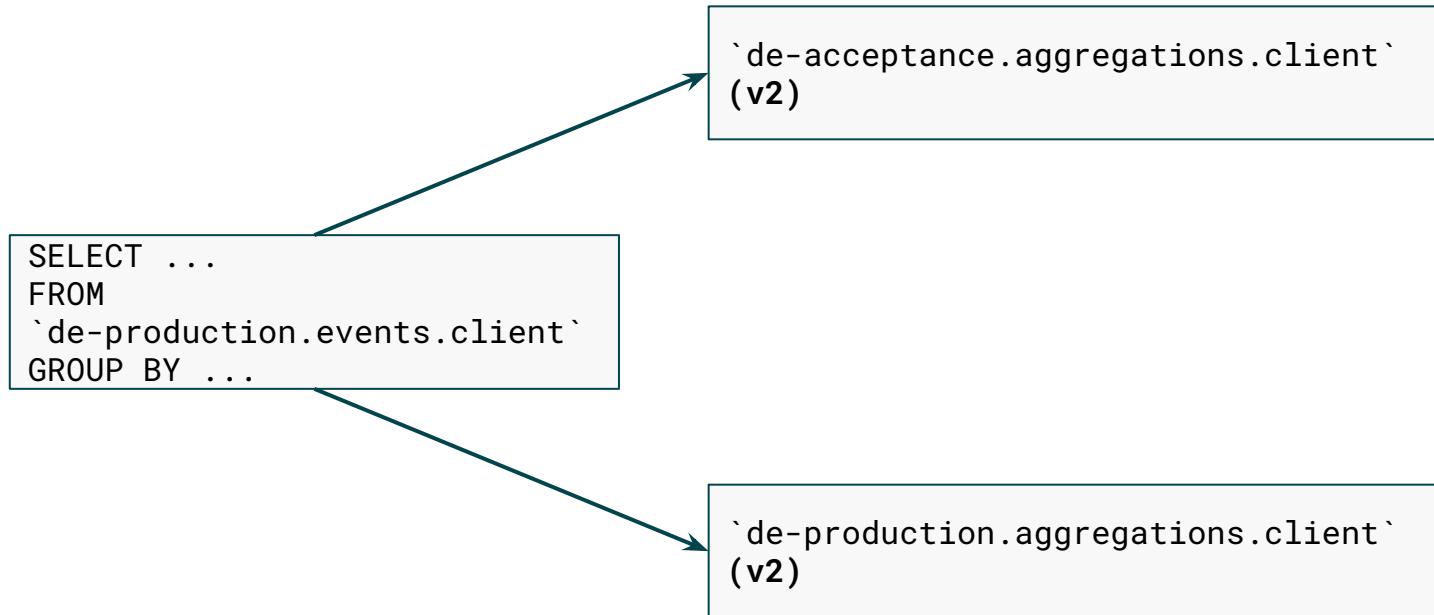
# Dataflow Example



No Changes in Data Structure (1 out of 3)

20

# Dataflow Example



No Changes in Data Structure (1 out of 3)

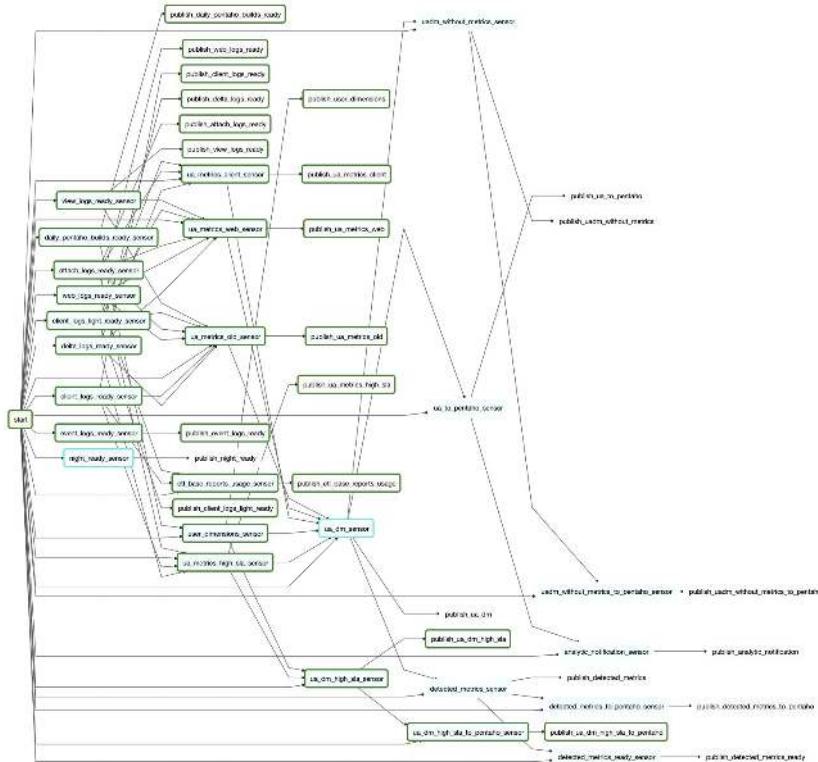
# Interface Separation on Other DBs

- Look for interface separation and resource isolation
  - And think about cost tradeoffs
- Approaches for **interface separation**
  - Schemas
  - Base directory name
  - Naming (bucket names for example)
  - Separate DBs
- Approaches for **resource isolation** (several trade offs with cost)
  - On service layer (separate DBs)
  - On DB side (e.g. roles, connection pools, quotas)
  - Airflow side (e.g. pools, priority, parallelism limit)
  - On monitoring side (e.g. query killer)

# No Changes in Data Freshness

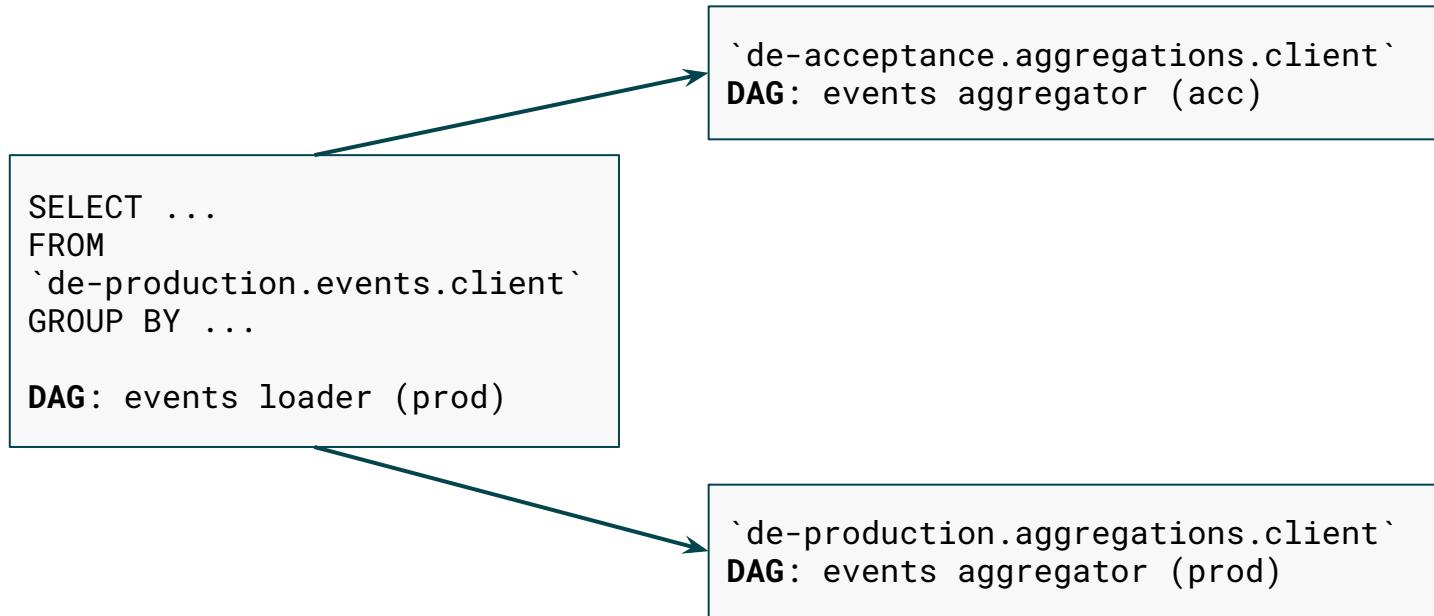


# Beautiful DAG with 150 Tasks



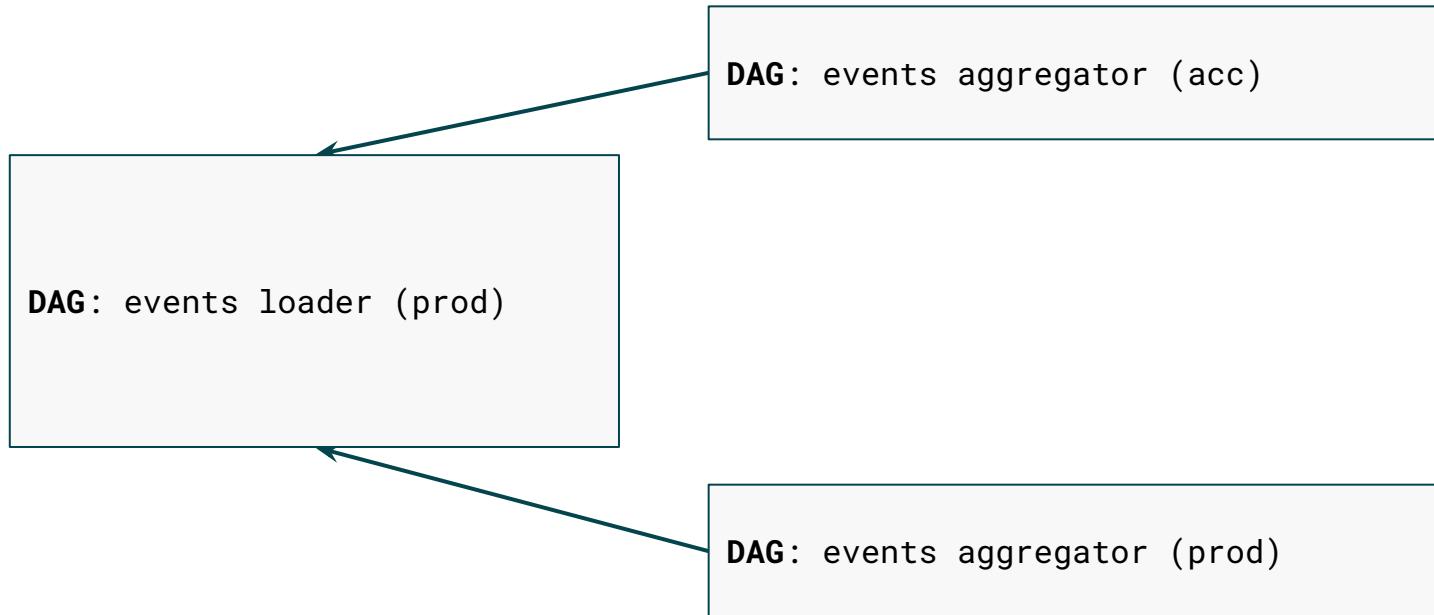
No Changes in Data Freshness (2 out of 3)

# Dataflow Example



No Changes in Data Freshness (2 out of 3)

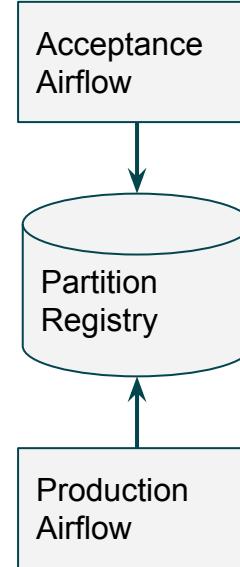
# Execution Example



No Changes in Data Freshness (2 out of 3)

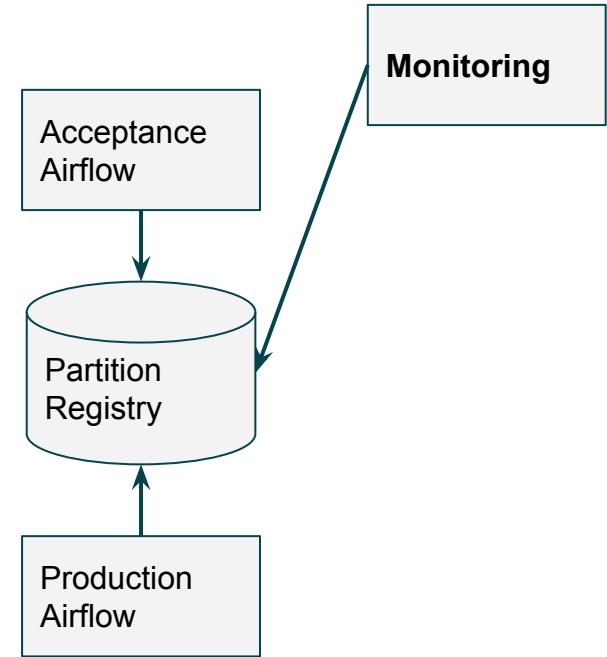
# Separate Airflows

- Coordinated via Postgres database named Partition Registry
  - Inspired by [Functional Data Engineering](#) by Maxime Beauchemin
  - Partition — unit of work for DAG, typically hour/day/week in a table
- State of partition published using operator
  - Explicitly publish sources
  - After all data validations have passed
- Wait for dependent sources using sensor
  - Automatically identify the strategy for interval
    - Week-on-hour, Month-on-day, custom catch-ups, etc.



# Partition Registry Now

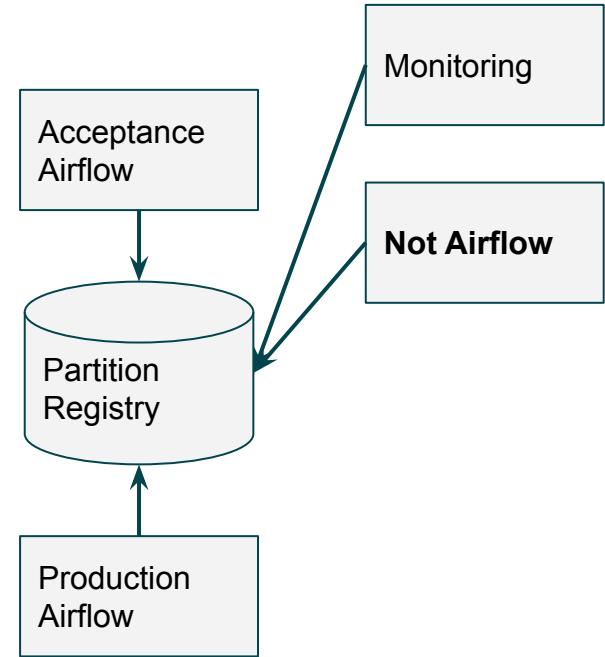
- Custom monitoring and alerts:
  - Severity of delays for partitions (DAG SLAs)
  - Base for data lineage



No Changes in Data Freshness (2 out of 3)

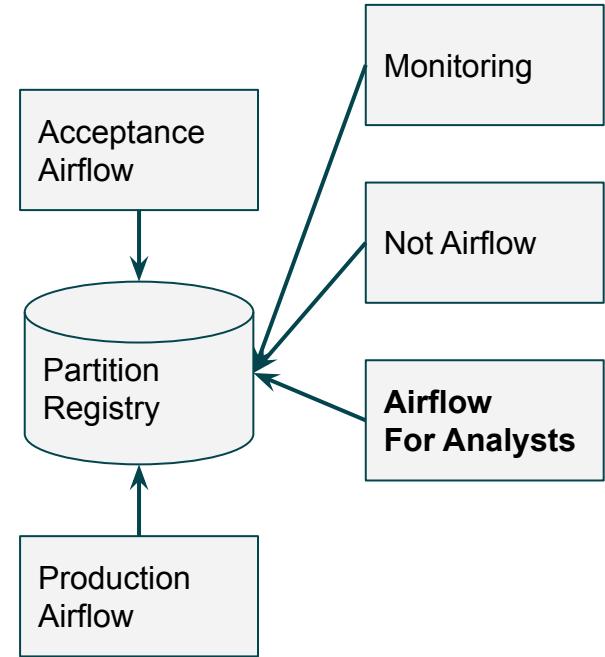
# Partition Registry Now

- Custom monitoring and alerts:
  - Severity of delays for partitions (DAG SLAs)
  - Base for data lineage
- Not Airflow: Pentaho DI and Old Jenkins Pipelines



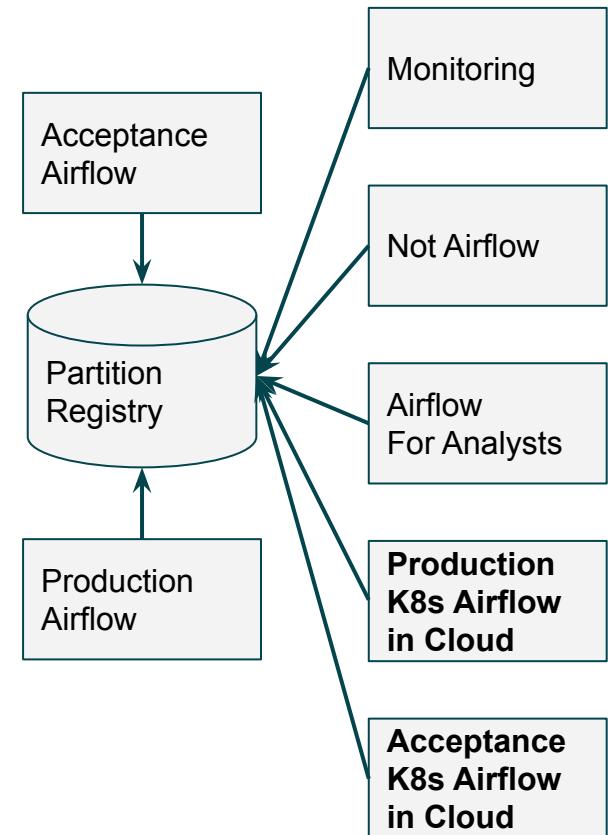
# Partition Registry Now

- Custom monitoring and alerts:
  - Severity of delays for partitions (DAG SLAs)
  - Base for data lineage
- Not Airflow: Pentaho DI and Old Jenkins Pipelines
- Airflow for Analysts: isolated resources and credentials



# Partition Registry Now

- Custom monitoring and alerts:
  - Severity of delays for partitions (DAG SLAs)
  - Base for data lineage
- Not Airflow: Pentaho DI and Old Jenkins Pipelines
- Airflow for Analysts: isolated resources and credentials
- K8s Airflow in Cloud
  - Easy switch with on-prem
  - Zero downtime migration
  - Data locality





A decorative graphic in the background features three overlapping circles: a large white circle at the top, a smaller yellow circle at the bottom left, and a green semi-circle on the right. A small teal triangle points towards the top center of the white circle.

# No Changes During Release from Acc to Prod

# Acceptance Told Us Where We Went Wrong

696 Commits 47 Branches



acceptance protected

- ed4b4018 · remove test\_source from alerts · 2 days ago

42 804

Merge request

production default protected

- ff1c9375 · new non\_pii\_backend views · 2 days ago



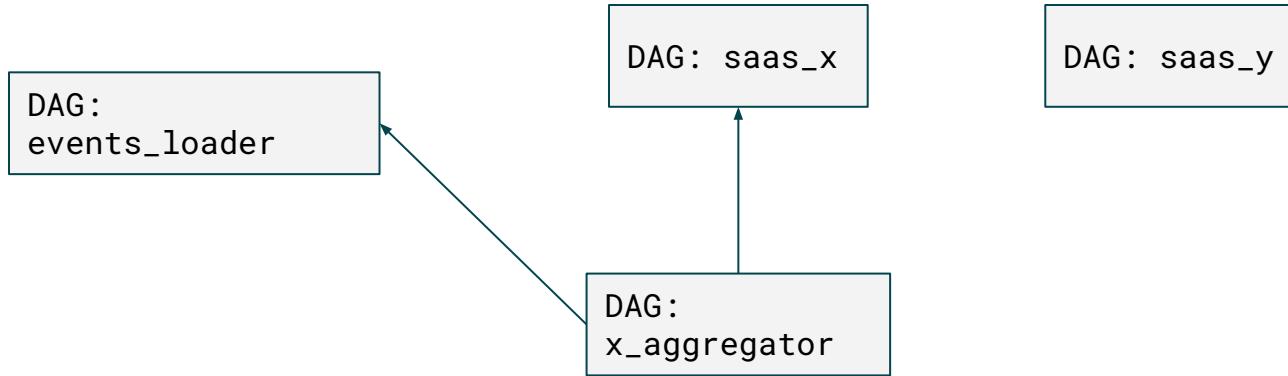
No Changes During Release Process (3 out of 3)

# Fast and Reliable Release

- We need code freeze to test dependent parts
- But we need 10 releases per day
  - So, we need to freeze as little as possible
    - But still review and test every change made

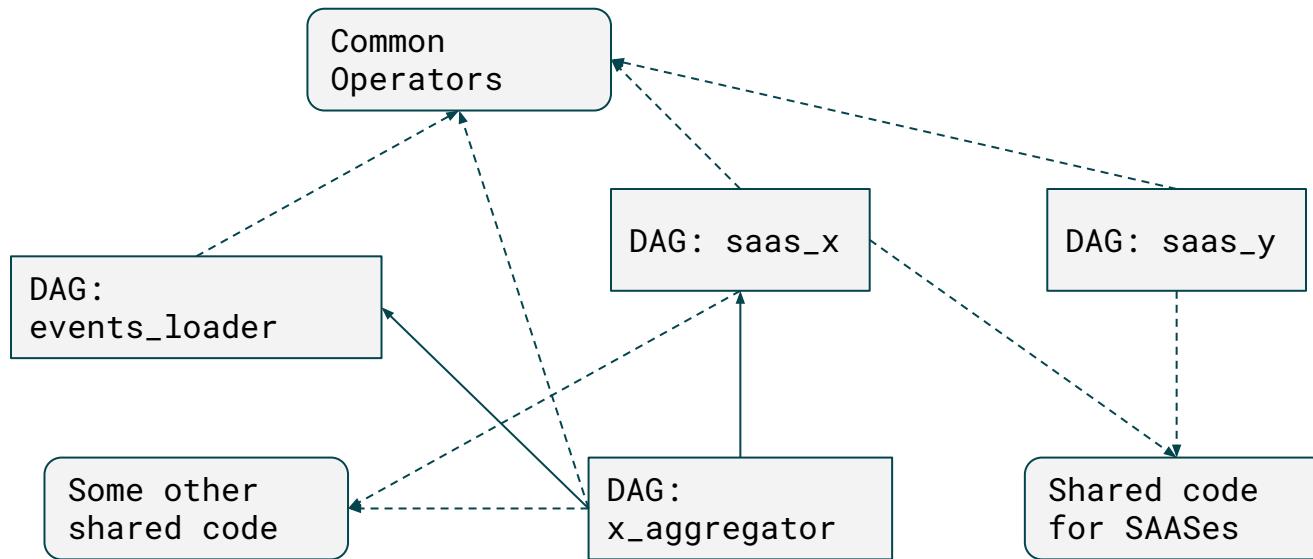


# Dependency Scheme



No Changes During Release Process (3 out of 3)

# Dependency Scheme with Code



No Changes During Release Process (3 out of 3)

# No Changes During Release Process Means

- Good data isolation during release
- Good code isolation during release



# Bad Data Isolation Is When

- You recalculate your data and get different results
- Data distribution changes
- Data distribution does not change when it should
- Analytical dashboard starts to focus on the wrong things
- You achieve your results a lot faster :)
- Something else is wrong and you don't know about it.



# So if Data Changes

- It's safe to assume
  - Review is no longer valid
  - Manual testing is no longer valid
  - Data sources may be corrupted
- So before the release of data change
  - Notifying all stakeholders of all changed dependent sources
  - Checking that everything works correctly on acceptance
  - Making atomic release
- We're helping to implement recalculation strategies
  - Recalculating everything and keeping it up-to-date
  - Preserving history for metrics in prestaging
  - Supporting and gradual deprecation of old version of metrics



# Keeping Track of Data Isolation

- Knowing when dependencies are updated after release to production
  - Notifications from other teams
  - Dependency on exact version of partition
    - Makes it easier to switch between acc and prod in code
  - Validation of data on your side
    - [Great Expectations](#) to explicitly specify your assumptions on data nature
    - Anomaly detection
- Finding all dependent sources before release to the production
  - Manual
    - BigQuery history
    - Search in git repository
  - Data Lineage + release process
  - Autotests

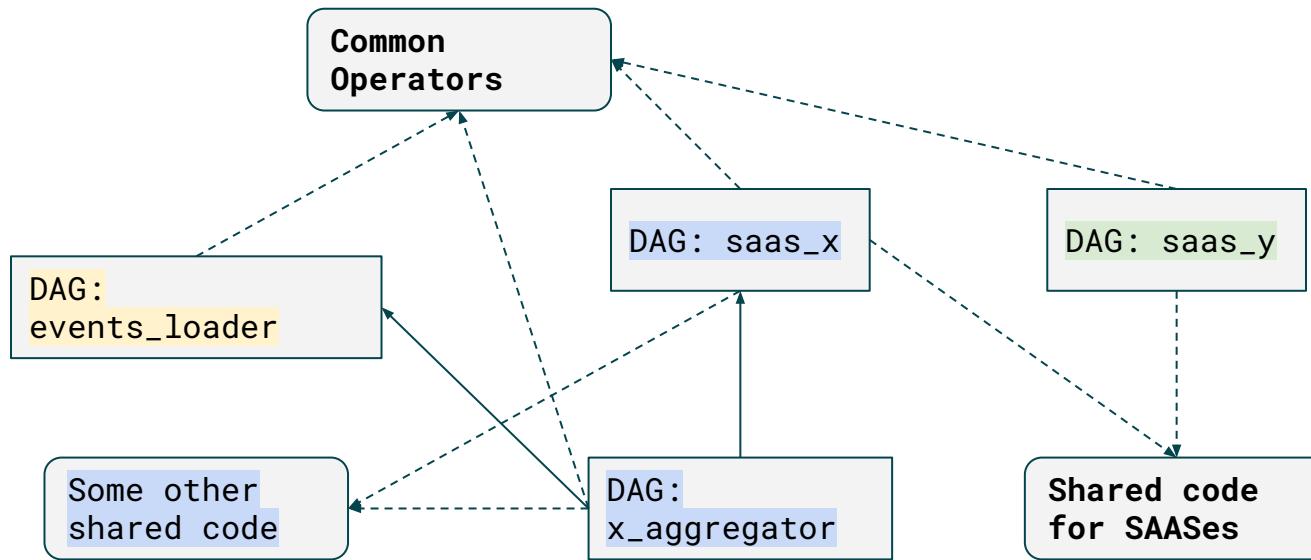


# Good Code Isolation

- Bad code isolation means you have a bug and your pipeline is not working
- This happens when 2+ DAGs use the same code
  - You update code or library and other DAG fails
- Two types of failure
  - Scheduler/Web Server — appears immediately, hard isolation (fat-zip, boilerplate)
  - Worker — visible during execution, easy isolation (k8s, venv)
    - Can be at the end of a 4 hour-long task at the start of the next month :(
- How do we avoid this?
  - There is 20% of code used in 80% of cases
    - We're moving it to the library, test and track backward compatibility
  - We have a shared code that is changed rarely
    - This code should be as private as possible to make sure that we're not reusing it
      - The main reason for DAGs to be included in the single repo or merge request



# Dependency Scheme with Code



No Changes During Release Process (3 out of 3)

# How Do We Reason About Reliability?

- Our production is very predictable
- All interface changes reviewed on separate environment
  - We keep track of all data dependencies and communicate the change to all stakeholders throughout the pipeline
  - Every source on production is reviewed, supported by several data engineers, have a clear time of readiness and all errors are communicated to all stakeholders
- We're using partition registry
  - To isolate resources of acceptance
    - As little recalculation as possible
  - To integrate Airflow with separate creds and resources to other teams
- Acceptance could be made cheaper





**Thank You!  
Any Questions?**

Alexander Eliseev at Airflow Slack  
[alexander.eliseev@team.wrike.com](mailto:alexander.eliseev@team.wrike.com)  
<https://github.com/eliseealex>

