

# Ensuring your DAGs work before going to production

*Airflow Summit 2020*

Bas Harenslak



GO   
DATA  
DRIVEN


# About me

- Background in Software Engineering (Haagse Hogeschool) and Computer Science (University of Leiden)
- Data Engineer at GoDataDriven last 4,5 years



# About me


- Apache Airflow committer
- (Co-)author of Manning's  
“Data Pipelines with Apache Airflow”
- First 10 chapters available in preview, new  
chapters added on a regular basis
- <https://www.manning.com/books/data-pipelines-with-apache-airflow>



# Ensuring your DAGs work before going to production

1. Unit testing
2. Testing in Python with pytest
3. Mocking
4. Using fake external systems
5. Debugging

# Typical very first deployment and test



On DAG: hello\_airflow

Graph View Tree View Task Duration Task Tries Landing Times Gantt Details Code Trigger DAG Refresh Delete

Base date: 2020-07-10T02:00:00+1 Number of runs: 25 Go


BashOperator PythonOperator

[DAG] hello airflow

A DAG visualization showing a single task named "hello" which depends on the DAG itself and the "airflow" operator. The task is currently running, indicated by a green circle. A red arrow points upwards from the DAG visualization towards the production box in the flow diagram.

GO DATA DRIVEN

# Second deployment



Development:

DAG: hello\_airflow

Graph View Tree View Task Duration Task Tries Landing Times Gantt Details Refresh Delete


Base date: 2020-07-11T19:51:57+ Number of runs: 25 Go

BashOperator PythonOperator

[DAG] hello airflow

A red arrow points upwards from the DAG view towards the 'Trigger DAG' button.


Production:




GO DATA DRIVEN

“Testing Airflow is hard”


# Testing Airflow is hard




# Testing Airflow is hard



# Testing Airflow is hard



# Testing Airflow is hard




# Unit testing

GO   
DATA  
DRIVEN

# What is a unit test?


```
def count_lines_in_postgres_table(table_name):  
    # SELECT COUNT(*) FROM table_name...  
  
    ↓  
result = count_lines_in_postgres_table("foobar")  
  
    ↓  
assert result == 5
```

# What is a unit test?



# Unit testing and Airflow

- A DAG consists of one or more tasks, represented as “Operators”
- We can test Operators as a whole
- Or callable functions (e.g. when using PythonOperator)
- Or underlying code (e.g. custom hooks/operators)
- Does *not* validate “integration” of multiple tasks



# So how do I unit test an operator?

- “It depends”
- Example test for PythonOperator:

```
def test_python_operator():
    test = PythonOperator(task_id="test", python_callable=lambda: "testme")
    result = test.execute(context={})
    assert result == "testme"
```

# So how do I unit test an operator? (2)

- “It depends”
- Example test for BashOperator:

```
def test_bash_operator():
    test = BashOperator(task_id="test", bash_command="echo testme", xcom_push=True)
    result = test.execute(context={})
    assert result == "testme"
```

“it depends”



# So how do I unit test an operator? (3)

- It depends:

```
class BaseOperator():
    # ...

    def execute(self, context):
        raise NotImplementedError()
```


- Note context is a required argument, hence the “`context={}`”.

# How do I test with context?

- If you need one single, fixed, thing from the context:

```
def next_week(**context):  
    return context["execution_date"] + datetime.timedelta(days=7)
```

```
def test_python_operator():  
    test = PythonOperator(task_id="test", python_callable=next_week, provide_context=True)  
    testdate = datetime.datetime(2020, 1, 1)  
    result = test.execute(context={"execution_date": testdate})  
    assert result == testdate + datetime.timedelta(days=7)
```



provide the context yourself

# Testing with pytest

GO   
DATA  
DRIVEN

# pytest

- Default testing package in Python in unittest
- Why 3rd party library pytest?
  - Because nice features (e.g. fixtures)
- <https://pytest.org>
- pip install pytest



pytest

# pytest vs unittest example

pytest

```
@pytest.fixture
def a():
    return 1

@pytest.fixture
def b():
    return 2

def test_sum(a, b):
    assert sum([a, b]) == 3
```

unittest

```
class TestWithUnittest(unittest.TestCase):
    def setUp(self):
        self.a = 1
        self.b = 2

    def test_sum(self):
        self.assertEqual(sum([self.a, self.b]), 3)
```

# pytest vs unittest example (2)

```
@pytest.fixture
def a():
    return 1

@pytest.fixture
def b():
    return 2

@pytest.fixture
def c():
    return 3

def test_sum_ab(a, b):
    assert sum([a, b]) == 3

def test_sum_ac(a, c):
    assert sum([a, c]) == 4
```

```
class TestWithUnittest(unittest.TestCase):
    def setUp(self):
        self.a = 1
        self.b = 2
        self.c = 3

    def test_sum_ab(self):
        self.assertEqual(sum([self.a, self.b]), 3)

    def test_sum_ac(self):
        self.assertEqual(sum([self.a, self.c]), 4)
```

# pytest fixture scopes

- Scope of fixtures can be:
  - Function (default)
  - Class
  - Module
  - Package
  - Session
- Especially useful to define when (not) to re-use variables

# pytest fixture scopes example

```
@pytest.fixture(scope="module")
def a():
    return [1]
```

Initialized *once*  
per module!

```
@pytest.fixture
def c():
    return [2]
```

```
class TestBla:
    def test_something(self, a):
        a.append(1)
        assert sum(a) == 2
```

```
def test_something_ac(self, a, c):
```

failure

# pytest fixture scopes

- Module scope is especially useful for “expensive” things
  - e.g. database clients

# pytest built-in fixtures

- capfd
- capfdbinary
- caplog
- capsys
- capsysbinary
- cache
- doctest\_namespace
- monkeypatch
- pytestconfig
- record\_property
- record\_testsuite\_property
- recwarn
- request
- testdir
- tmp\_path
- tmp\_path\_factory
- tmpdir ←————
- tmpdir\_factory

<https://docs.pytest.org/en/stable/fixture.html#fixtures>

# pytest tmpdir fixture

```
def test_writing_to_disk(tmpdir):
    tmpfile = tmpdir.join("hello.txt")  
  
    task = BashOperator(task_id="test", bash_command=f"echo 'hello' > {tmpfile}")
    task.execute(context={})  
  
    assert len(tmpdir.listdir()) == 1
    assert tmpfile.read().replace("\n", "") == "hello"  
  
    ↑  
    Assertions
```

Define tmp file path

Run BashOperator


Write "hello" to tmpfile

# Enough pytest – show me Airflow!

- One useful pytest plugin: `pytest-helpers-namespace`
  - `pip install pytest-helpers-namespace`
- If you need the “full” context, you will need a DAG:

```
import datetime
import pytest
from airflow.models import DAG

pytest_plugins = ["helpers_namespace"]
```




```
@pytest.fixture
def test_dag():
    """Airflow DAG for testing."""
    return DAG(
        "test_dag",
        start_date=datetime.datetime(2020, 1, 1),
        schedule_interval=datetime.timedelta(days=1),
    )
```

```
@pytest.helpers.register
def run_task(task, dag):
    """Run an Airflow task."""
    dag.clear()
    task.run(start_date=dag.start_date, end_date=dag.start_date)
```

# Place this script in a conftest.py

- Pytest will auto-discover anything in /tests prefixed with `test_`/`Test`
- And automagically read everything defined in `conftest.py`
- “per-directory configuration”



# Running a task with DAG

```
def test_bash_operator_tmpdir(test_dag, tmpdir):
    tmpfile = tmpdir.join("hello.txt")

    task = BashOperator(task_id="test", bash_command=f"echo 'hello' > {tmpfile}", dag=test_dag)
    pytest.helpers.run_task(task=task, dag=test_dag)

    assert len(tmpdir.listdir()) == 1
    assert tmpfile.read().replace("\n", "") == "hello"
```

# Running a task with full context


```
def test_full_context(test_dag, tmpdir):
    def do_magic(**context): ← context
        with open(tmpdir / "test.txt", "w") as f:
            f.write(context["ds"])

    task = PythonOperator(task_id="test", python_callable=do_magic, provide_context=True, dag=test_dag)
    pytest.helpers.run_task(task=task, dag=test_dag)

    with open(tmpdir / "test.txt", "r") as f:
        assert f.readlines()[0] == test_dag.start_date.strftime("%Y-%m-%d")
```

# The “DAG integrity test”

- Asserts the validity of DAG objects,  
i.e. “can Python instantiate this DAG”
- Plus, test for cycles
- <https://medium.com/wbaa/datas-inferno-7-circles-of-data-testing-hell-with-airflow-cef4adff58d8>



# “Hello world” of checking DAG correctness

- `python your_dag.py`
- Does it take long?
- Does it produce many logs?
- Does it make connections to external systems?
- Etc...
- *Anything in the global scope must be avoided!*

# The “DAG integrity test”

```
import glob
from os import path

import pytest
from airflow import models as airflow_models

DAG_PATHS = glob.glob(path.join(path.dirname(__file__), "..", "..", "dags", "*.py"))
```


```
@pytest.mark.parametrize("dag_path", DAG_PATHS) ← For every DAG file
def test_dag_integrity(dag_path):
    """Import DAG files and check for a valid DAG instance."""
    dag_name = path.basename(dag_path)
    module = _import_file(dag_name, dag_path)

    # Validate if there is at least 1 DAG object in the file
    dag_objects = [var for var in vars(module).values() if isinstance(var, airflow_models.DAG)]
    assert dag_objects ← Assert if valid DAG object

    # For every DAG object, test for cycles
    for dag in dag_objects:
        dag.test_cycle() ← Test for cycles

def _import_file(module_name, module_path):
    import importlib.util

    spec = importlib.util.spec_from_file_location(module_name, str(module_path))
    module = importlib.util.module_from_spec(spec)
    spec.loader.exec_module(module)
    return module
```



# The “DAG integrity test”

```
task1 = DummyOperator(task_id="task1", dag=dag)
task2 = DummyOperator(task_id="task2", dag=dag)

task1 >> task2 >> task1
```

```
airflow.exceptions.AirflowDagCycleException: Cycle detected in DAG. Faulty task: task2 to task1
```

# Mocking

GO   
DATA  
DRIVEN

# Mocking – what and why?

- Provide “canned” responses, typically to fake calls to external systems
- For example – to test a function processing an API call, provide it a fake result to validate the processing part.
- Within Airflow – we can use mocking to e.g. provide a Connection object, so we don’t need a metastore for testing.
- Sometimes requires digging in internal code.

# Mocking example

- Let's consider an external API (something not maintained by ourselves):
- <https://api.sunrise-sunset.org/json?lat=52.370216&lng=4.895168&formatted=0>
- Returns sunrise and sunset times of the given location (Amsterdam city centre):

```
{  
  "results": {  
    "sunrise": "2020-07-12T03:34:14+00:00",  
    "sunset": "2020-07-12T19:58:09+00:00",  
    "solar_noon": "2020-07-12T11:46:12+00:00",  
    "day_length": 59035,  
    "civil_twilight_begin": "2020-07-12T02:47:34+00:00",  
    "civil_twilight_end": "2020-07-12T20:44:49+00:00",  
    "nautical_twilight_begin": "2020-07-12T01:37:06+00:00",  
    "nautical_twilight_end": "2020-07-12T21:55:17+00:00",  
    "astronomical_twilight_begin": "1970-01-01T00:00:01+00:00",  
    "astronomical_twilight_end": "1970-01-01T00:00:01+00:00"  
  },  
  "status": "OK"  
}
```

*Sunrise/sunset  
times on the date  
of making this call*

# Mocking example

```
from datetime import datetime  
  
import pytest  
from airflow.hooks.base_hook import BaseHook  
from airflow.models import Connection  
from airflow.operators.http_operator import SimpleHttpOperator  
  
Test definition ──────────> def test_simple_http_operator(test_dag, mocker):  
    mocker.patch.object(  
        BaseHook,  
        "get_connection",  
        return_value=Connection(schema="https", host="api.sunrise-sunset.org")  
    )  
  
Mock BaseHook.get_connection() ──────────>  
  
Fake Connection response ──────────> def _check_light(sunset_sunrise_response):  
    results = sunset_sunrise_response.json()["results"]  
    sunrise = datetime.strptime(results["sunrise"][:-6], "%Y-%m-%dT%H:%M:%S")  
    sunset = datetime.strptime(results["sunset"][:-6], "%Y-%m-%dT%H:%M:%S")  
  
    if sunrise < datetime.utcnow() < sunset:  
        print("It is light!")  
    else:  
        print("It is dark!")  
  
    return True  
  
is_it_light = SimpleHttpOperator(  
    task_id="is_it_light",  
    http_conn_id="my_http_conn",  
    endpoint="json",  
    method="GET",  
    data={"lat": "52.370216", "lng": "4.895168", "formatted": "0"},  
    response_check=_check_light,  
    dag=test_dag,  
)  
  
pytest.helpers.run_task(task=is_it_light, dag=test_dag)
```

# Mocking and pytest

- Unittest comes with a built-in mocking package (`from unittest import mock`)
- While you can use this in pytest
- A convenient plugin is `pytest-mock`, which makes it available as a fixture
- `pip install pytest-mock`

# How to mock the “unmockable”?

- Some systems cannot be mocked
- Cloud services for example don't always provide a mocking library
- Three options:
  1. Mock and assert calls to the external systems
  2. Run a Dockerized version of the external system (e.g. PostgreSQL)
  3. Run a development version of the external system (e.g. Google Cloud Storage)

# Mock and assert calls to external systems

```
def test_simple_http_operator_no_external_call(test_dag, mocker):
    mocker.patch.object(
        BaseHook, "get_connection", return_value=Connection(schema="https", host="api.sunrise-sunset.org")
    )
    mock_run = mocker.patch.object(HttpHook, "run")  
    ←-----  
  
    is_it_light = SimpleHttpOperator(
        task_id="is_it_light",
        http_conn_id="my_http_conn",
        endpoint="json",
        method="GET",
        data={"lat": "52.370216", "lng": "4.895168", "date": "{{ ds }}", "formatted": "0"},
        dag=test_dag,
    )  
  
    pytest.helpers.run_task(task=is_it_light, dag=test_dag)
    mock_run.assert_called_once()
    assert mock_run.call_args_list[0][0][1] == {
        "lat": "52.370216",
        "lng": "4.895168",
        "date": test_dag.start_date.strftime("%Y-%m-%d"),  
    ←-----  
        "formatted": "0",
    }  
}
```

# Faking external systems

GO   
DATA  
DRIVEN

# Testing with Docker



- We can run external systems for the duration of our tests in Docker
- Note: a Docker image of your desired system must exist (which is not the case for e.g. Google Cloud Storage)
- One option: create a Docker Compose for your tests
- Alternative option: use `pytest_docker_tools` plugin

# pytest\_docker\_tools

- `pip install pytest_docker_tools`
- Small wrapper around your Docker client (so Docker must be installed!)
- Provide access to Docker SDK with a set of pytest fixtures

# pytest\_docker\_tools

- Say we implemented a “PostgresToLocalOperator”, which runs a PostgreSQL query and writes the result to a local file.
- We need:
  - PostgreSQL database
  - Some location to write to locally
  - Data in the database to download

# pytest\_docker\_tools – create database

```
from pytest_docker_tools import container, fetch

postgres_image = fetch(repository="postgres:11.1-alpine")  
    ← Defines Docker image

postgres = container(  
    image="{postgres_image.id}",  
    environment={  
        "POSTGRES_USER": "secretuser",  
        "POSTGRES_PASSWORD": "secretpassword",  
    },  
    ports={"5432/tcp": None},  
    volumes={  
        path.join(path.dirname(__file__), "postgres-init.sql"): {  
            "bind": "/docker-entrypoint-initdb.d/postgres-init.sql"  
        }  
    },  
)
```

← Defines Docker container  
(very similar to docker run ...)

# pytest\_docker\_tools – initialize database

**postgres-init.sql:**

```
SET search_path TO public;
CREATE TABLE dummy (
    id integer,
    name character varying(255)
);
INSERT INTO dummy (id,name) VALUES (1, 'dummy1');
INSERT INTO dummy (id,name) VALUES (2, 'dummy2');
INSERT INTO dummy (id,name) VALUES (3, 'dummy3');
```

# pytest\_docker\_tools – the actual test

Mock Postgres connection

```
def test_postgres_to_local_operator(test_dag, mocker, tmpdir, postgres):  
    mocker.patch.object(  
        PostgresHook,  
        "get_connection",  
        return_value=Connection(  
            host="localhost",  
            conn_type="postgres",  
            login=postgres_credentials.username,  
            password=postgres_credentials.password,  
            port=postgres.ports["5432/tcp"][0],  
        ),  
    )
```

Reference to Postgres container

```
    output_path = str(tmpdir / "pg_dump")  
    task = PostgresToLocalOperator(  
        task_id="test",  
        postgres_conn_id="postgres",  
        pg_query="SELECT * FROM dummy",  
        local_path=output_path,  
        dag=test_dag,  
    )  
    pytest.helpers.run_task(task=task, dag=test_dag)
```

Call the operator

Check results

```
# Assert if output file exists  
output_file = Path(output_path)  
assert output_file.is_file()  
  
# Assert file contents, should be the same as in postgres-init.sql  
expected = [  
    {"id": 1, "name": "dummy1"},  
    {"id": 2, "name": "dummy2"},  
    {"id": 3, "name": "dummy3"},  
]  
with open(output_file, "r") as f:  
    assert json.load(f) == expected
```

# Debugging

GO   
DATA  
DRIVEN


# airflow test

- airflow test [dag id] [task id] [execution date]
- E.g. “airflow test mydag mytask 2020-01-01”
- Running a single task from the command line, for a given execution date
- No state recorded in metastore

*“airflow task test” in Airflow 2.0*

# Debugging Airflow code

- Set breakpoints in your IDE locally:



```
task = PostgresToLocalOperator(  
    task_id="test",  
    postgres_conn_id="postgres",  
    pg_query="SELECT * FROM dummy",  
    local_path=output_path,  
    dag=test_dag,  
)  
pytest.helpers.run_task(task=task, dag=test_dag)  
  
# Assert if output file exists  
output_file = Path(output_path)  
assert output_file.is_file()
```

# Debugging in production

- Try to avoid
- But if you must;
- (i)Python DeBugger: (i)pdb

```
import ipdb  
ipdb.set_trace()
```

# Debugging in production (2)

 Python Debugger Cheatsheet 

### Getting started

start pdb from within a script:  
`import pdb;pdb.set_trace()`  
start pdb from the commandline:  
`python -m pdb <file.py>`

### Basics

`h(elp)` print available commands  
`h(elp) command` print help about *command*  
`q(uit)` quit debugger

### Examine

`print(expr)` print the value of *expr*  
`pp expr` pretty-print the value of *expr*  
`w(here)` print current position (including stack trace)  
`l(ist)` list 11 lines of code around the current line  
`l(ist) first, last` list from *first* to *last* line number  
`a(rgs)` print the args of the current function

### Movement

<ENTER> repeat the last command  
`n(ext)` execute the current statement (step over)  
`s(tep)` execute and step into function  
`return` continue execution until the current function returns  
`c(ontinue)` continue execution until a breakpoint is encountered  
`u(p)` move one level up in the stack trace  
`d(own)` move one level down in the stack trace

### Breakpoints

`b(reak)` show all breakpoints  
`b(reak) lineno` set a breakpoint at *lineno*  
`b(reak) func` set a breakpoint at the first line of a *func*

### Manipulation

`!stmt` treat *stmt* as a Python statement instead of a pdb command

---

by Florian Preinstorfer (nblock@archlinux.us) — version 1.0 — license cc-by-nc-sa 3.0  
see <https://github.com/nblock/pdb-cheatsheet> for more information.

GO DATA DRIVEN

# airflow test

- airflow test [dag\_id] [task\_id] [execution\_date]

*“airflow tasks test” in Airflow 2.0*



# DebugExecutor

- <https://airflow.apache.org/docs/stable/executor/debug.html>
- Add to your DAG script:

```
if __name__ == "__main__":
    dag.clear(reset_dag_runs=True)
    dag.run()
```

- export AIRFLOW\_\_CORE\_\_EXECUTOR=DebugExecutor
- And run with python your\_dag.py
- ***Warning:*** will run ALL runs from configured start\_date!

# Coming in Airflow 2.0: airflow dags test

- Complete CLI was rewritten in Airflow 2.0
- Can run a DAG using the DebugExecutor *without* editing the DAG

```
airflow dags test dag_id execution_date
```

- Executes all tasks in the order they're defined in

# Final words

- Airflow operators can be tested (run) by simply calling them
- If you want the context, you need a DAG (or explicitly provide it yourself)
- Mocking is useful for e.g. avoiding metastore calls
- Docker can be useful for running systems temporarily during a test
- “airflow test” can be helpful for debugging

# Final words (2)

- Talk is recorded
- Code is available at <https://github.com/godatadriven/airflow-testing-examples>
- @basph on Airflow Slack