



Modernize a decade old pipeline with Airflow 2.0

Dima Suvorov, Kuntal Basu, Stas Bytsko and QP Hou

Talk Overview


- Migration overview - Stas
- Custom Trigger rules
- Migration to Airflow 2.0
- Running Backfill at scale - QP
- Self-Service Backfill UI plugin - Kuntal
- Fixing bugs in Backfill code - Dima
- Databricks clusters cost optimization

Migration overview

- Compute + Storage => 

 - AWS & Databricks

- Improve security and compliance
- Custom scheduler -> Airflow
- Mono-DAG
 - 1.4K tasks
 - Nestedness: up to 22 layers deep



<https://tech.scribd.com/blog/2020/modernizing-an-old-data-pipeline.html>

Custom Trigger Rules

Example


Migrated DAG gradually. DAG served 2 purposes:

- Run Production tasks
- Validate not-yet migrated tasks


Components:

1. DataSync - All tasks need input data from legacy env
2. Production Operators - generates business value
3. Migration Operators - unreliable, under test
 - Output written to separate database and validated against synced data produced by task in legacy env

Dependencies on Migration Operators - not real, only to simplify migration for teams




Naïve approach



Pretend our migration tasks are the same as production tasks

Naïve approach problem



Migration tasks can fail, bringing down production tasks

Custom TriggerRuleDep

```
class BaseOperator( ... ):  
    class ScribdBaseOperator(BaseOperator):  
        deps: Iterable[BaseTIDep] = frozenset({  
            NotInRetryPeriodDep(),  
            PrevDagrunDep(),  
            TriggerRuleDep(),  
            TriggerRuleDepMigration(),  
            NotPreviouslySkippedDep(),  
        })
```

Custom TriggerRuleDep


```
class TriggerRuleDep(BaseTIDep):
    class TriggerRuleDepMigration(TriggerRuleDep):

        @staticmethod
        def _get_states_count_upstream_ti(ti, finished_tasks):
            counter = Counter(task.state
                counter = Counter(task.state if task.operator not in {"MigrationOperator"} else State.SUCCESS
                                for task in finished_tasks if task.task_id in ti.task.upstream_task_ids)
            return (
                counter.get(State.SUCCESS, 0),
                counter.get(State.SKIPPED, 0),
                counter.get(State.FAILED, 0),
                counter.get(State.UPSTREAM_FAILED, 0),
                sum(counter.values()),
            )
    
```


+ 'trigger_rule'

Always treat Migration tasks as successful

Naïve approach solution - DAG



Perf. Problem



Unoptimized task can hold up all of the production downstreams

Perf. Optimization

```
def _get_dep_statuses(self, ti, session, dep_context):
    ...
    # see if the task name is in the task upstream for our task
    successes, skipped, failed, upstream_failed, done = self._get_states_count_upstream(ti,
        ... ti=ti, finished_tasks=dep_context.ensure_finished_tasks(ti.task.dag, ti.execution_date, session)
    )

    upstream_tis = session.query(TaskInstance).filter(
        TaskInstance.dag_id == ti.task.dag.dag_id,
        TaskInstance.execution_date == ti.execution_date,
        TaskInstance.task_id.in_(ti.task.upstream_task_ids),
    ).all()


    finished_or_migration_tasks = {
        ti for ti in upstream_tis
        if ti.state in State.finished or
            ti.operator in {"MigrationOperator"}
    }

    counter = Counter(ti.state if ti.operator not in {"MigrationOperator"} else State.SUCCESS
                      for ti in finished_or_migration_tasks)
```

1. Get all upstream tasks


2. Take finished + Migration

3. Consider Migration tasks as SUCCESS




inline and rewrite

Perf. Optimization - DAG



Last problem: UPSTREAM_FAILED not propagated




```
Dag finishes finished_or_migration_tasks = {  
    ti for ti in upstream_tis  
    if ti.state in State.finished or  
    ti.operator in {"MigrationOperator"}  
}
```

}

```
counter = Counter(  
    State.SUCCESS  
    if ti.operator in {"MigrationOperator"}  
    else ti.state  
    for ti in finished_or_migration_tasks)
```

If DataSync fails, all tasks have to take this into account and stop -
State.UPSTREAM_FAILED has to be propagated

Task lifecycle refresher



<https://airflow.apache.org/docs/apache-airflow/stable/concepts/tasks.html#task-instances>


Propagate UPSTREAM_FAILED

```
finished_or_migration_tasks = {  
    ti for ti in upstream_tis  
    if ti.state in State.finished or  
        ti.operator in {"MigrationOperator"}  
        (ti.operator in {"MigrationOperator"}) and ti.state ≠ State.NONE)  
}  
  
1. Exclude Migration tasks that  
didn't meet deps criteria
```


```
counter = Counter(State.SUCCESS if ti.operator in {"MigrationOperator"} else ti.state  
counter = Counter(State.SUCCESS if ti.operator in {"MigrationOperator"}  
                    and ti.state ≠ State.UPSTREAM_FAILED else ti.state  
                    for ti in finished_or_migration_tasks)
```

2. Propagate UPSTREAM_FAILED
Status

Propagate UPSTREAM_FAILED - DAG



Airflow 2.0 Upgrade



Spoiler alert: It's a one way trip

Airflow upgrade check is your friend

```
pip install apache-airflow-upgrade-check  
airflow upgrade_check
```

Airflow 2.0 upgrade - MySQL (Aurora RDS)

- MySQL 5.6 not supported by Airflow 2.0
 - Missing JSON column types
- MySQL 5.7 kind of works
- MySQL 8 not supported by Aurora RDS
 - Required for scheduler HA

Airflow 2.0 upgrade - Trigger rules

```
▼ 13 airflow/serialization/serialized_objects.py □ ...
↑ @@-390,12,+390,6 @@ def serialize_operator(cls, op: BaseOperator) -> Dict[str, Any]:
390     .....for dep in op.deps:
391         .....klass = type(dep)
392         .....module_name = klass.__module__
393     - .....if not module_name.startswith("airflow.ti_deps.deps."):
394     - .....raise SerializationError(
395     - .....f"Cannot serialize {{op.dag.dag_id}::{op.task_id}!r} with `deps` from non-core "
396     - .....f"module {module_name!r}"
397     - .....)
398     -
399     .....deps.append(f'{module_name}.{klass.__name__}')
400     .....serialize_op['deps'] = deps
401
↓ @@-504,14,+498,7 @@ def _is_excluded(cls, var: Any, attrname: str, op: BaseOperator):
504     .....def _deserialize_deps(cls, deps: List[str]) -> Set["BaseTIDep"]:
505         .....instances = set()
506         .....for qualname in set(deps):
507     - .....if not qualname.startswith("airflow.ti_deps.deps."):
508     - .....log.error("Dep class %r not registered", qualname)
509     - .....continue
510     -
511     - .....try:
512         .....instances.add(import_string(qualname()))
513     - .....except ImportError:
514     - .....log.warning("Error importing dep %r", qualname,
515         .....exc_info=True)
515     .....return instances
390     .....for dep in op.deps:
391         .....klass = type(dep)
392         .....module_name = klass.__module__
393     - .....deps.append(f'{module_name}.{klass.__name__}')
394     .....serialize_op['deps'] = deps
395
498     .....def _deserialize_deps(cls, deps: List[str]) -> Set["BaseTIDep"]:
499         .....instances = set()
500         .....for qualname in set(deps):
501             .....instances.add(import_string(qualname()))
502
...
...
```

Airflow 2.0 upgrade - Performance improvement

- Faster Web UI
- Faster scheduler
- Scheduler sharding

Scheduler CPU usage after 2.0 upgrade



Running Backfill at Scale

Running Backfill at scale

Our goal

1. Backfill data for 14 years
2. Our intended DAG concurrency (i.e. how many version of single DAG we can run concurrently) was 150, we settled later to 100

Limitless Limits

People say “Sky is the limit”, but to reach the sky there is a small matter of gravity that we have to overcome. Exactly that happened to us. Let us talk about our gravitational boundaries.

Airflow limits

AIRFLOW_CORE_PARALLELISM

The amount of parallelism as a setting to the executor. This defines the max number of task instances that should run simultaneously. Default value is 32. We override that in our backfill execution commands to 100.

AIRFLOW__CORE__MAX_ACTIVE_RUNS_PER_DAG

The maximum number of active DAG runs per DAG. It maps to `max_active_runs` attribute in the DAG definition. Default value is 16. We override it to 100.

AIRFLOW__CORE__DAG_CONCURRENCY

The number of task instances allowed to run concurrently by the scheduler in one DAG. It maps to `concurrency` attribute in the DAG definition. We override it to 100.

Databricks and AWS limits

Apart from Airflow limits, we got restricted by Databricks and AWS account limits while working on the backfill. Here are some examples:

1. AWS account limit of 1000 TB of total GP2 EBS volume size. We increased it to 1500 TB while at the same time reduced our EBS volume size per machine by almost 60%.
2. Databricks API limit. We were getting “429 Too Many Requests” errors from Databricks.
3. Databricks Node creation limit at 200 nodes per minute. We worked with Databricks to get these limits lifted for our account.

How much is too much

Tree View

Graph View

Task Duration


Task Tries

Landing Times

Gantt

Details

Code



Self-Service Backfill UI plugin

Self-Service Backfill UI plugin

Why?

- Switch from Legacy in-house system to Airflow
- Increased load on Airflow Admins
- Give back the ability to run self-serviced backfills to our engineers
- Web UI based backfill trigger is still being discussed by the community

Self-Service Backfill UI plugin


Considered approaches:

- Feed all tasks to the scheduler
- New type of Job in the Web Server
- Use built-in Backfill functionality

Self-Service Backfill UI plugin

How?


- AWS Elastic Container Service



Self-Service Backfill UI plugin

How?


- AWS Elastic Container Service
- AWS Elastic Container Registry



Self-Service Backfill UI plugin

How?


- AWS Elastic Container Service
- AWS Elastic Container Registry
- ECS container to ECS calls



Self-Service Backfill UI plugin

How?


- AWS Elastic Container Service
- AWS Elastic Container Registry
- ECS container to ECS calls
- New ECS container for each Backfill



Self-Service Backfill UI plugin

How?


- AWS Elastic Container Service
- AWS Elastic Container Registry
- ECS container to ECS calls
- New ECS container for each Backfill
- Aurora Relational Database Service (MySQL)



Self-Service Backfill UI plugin

How?

- AWS Elastic Container Service
- AWS Elastic Container Registry
- ECS container to ECS calls
- New ECS container for each Backfill
- Aurora Relational Database Service (MySQL)





DAGs

Security

Browse

Admin

Docs

Backfill

13:07 UTC



DAG: dummy-dag

schedule: 21 2 * * *

Tree View

Graph View

Task Duration

Task Tries

Landing Times

Gantt

Details



< > Code



2021-06-07T02:21:00Z

Runs

25

Update

BashOperator

queued running success failed up_for_retry up_for_reschedule upstream_failed skipped scheduled no_status05:21
|
●
■
■



Backfill

DAG Name

dummy-dag

Start Datetime

01.06.2021, 16:03

 Ignore Dependencies**End Datetime**

02.06.2021, 16:03

**Tasks To Run (Could be Task ID or regex)****Config (JSON string that gets pickled into the DagRun's conf attribute)****Start****Clear**



Wait a minute.

Here's the backfill params and the list of task instances you are about to backfill

Backfill Params:

```
dag_id=dummy-dag
start_date=2021-06-01T16:03
end_date=2021-06-02T16:03
ignore_dependencies=True
```

Task Instances List (2 task instances):

```
<TaskInstance: dummy-dag.task1 2021-06-02 02:21:00+00:00 [None]>
<TaskInstance: dummy-dag.task2 2021-06-02 02:21:00+00:00 [None]>
```

OK!Cancel



Airflow

DAGs

Security ▾

Browse ▾

Admin ▾

Docs ▾

Backfill

Backfill job has been successfully triggered.

DAGs

All 3

Active 1

Paused 2

Filter DAGs by tag

i DAG	Owner	Runs i	Schedule
<input checked="" type="checkbox"/> dummy-dag	data-eng	1	21 2 * * *



New ECS Experience

[Tell us what you think](#)

Amazon ECS

Clusters

Task Definitions

Account Settings

Amazon EKS

Clusters

Amazon ECR

Repositories

AWS Marketplace

Discover software

Subscriptions ↗

Clusters > backfill-cluster > Task: 6f0c0862eb0b40d485d6618f50ae60c9

Task : 6f0c0862eb0b40d485d6618f50ae60c9

Details

Tags

Logs

Cluster backfill-cluster

Launch type FARGATE

Platform version 1.4.0

Task definition backfill-task-definition:17

Group family:backfill-task-definition

Task role ecsTaskExecutionRole

Last status PENDING

Desired status RUNNING

Created at 2021-06-08 16:10:58 +0300

Name	Container Runtime I...	Status	Image	Image Digest	CPU Units	Hard/Soft memory ...	Essential	Resource ID
airflow_ba...	9f426f47ea044e2dbf7...	STOPPED	341828981035.dkr.ecr.us-east-2.a...	sha256:a0c8ced0e98d0...	0	--/--	true	20a90990-3284-403...

Details

Exit Code0

Command["airflow", "dags", "backfill", "--ignore-dependencies", "--start-date", "2021-06-01T16:03", "--end-date", "2021-06-02T16:03", "dummy-dag"]

Network bindings - not configured

Environment Variables

Key	Value
AIRFLOW__CORE__SQLALCHEMY_CONN	mysql+mysqldb://*****:*****@*****.rds.amazonaws.com:*****/airflow
AIRFLOW__LOGGING__COLORED_CONSOLE_LOG	False

Environment Files - not configured

Docker labels - not configured

Extra hosts - not configured

Mount Points - not configured

Volumes from - not configured

Ulimits - not configured


Elastic Inference - not configured

Log Configuration

Log driver: awslogs [View logs in CloudWatch](#)

Key	Value
awslogs-group	/ecs/backfill-task-definition
awslogs-region	us-east-2
awslogs-stream-prefix	ecs

Bug Busters



Backfill vs Scheduler

Backfill vs Scheduler

3 types of jobs in Airflow:

- SchedulerJob
- BackfillJob
- LocalTaskJob

Backfill vs Scheduler

Example:

1. SchedulerJob creates a DagRun and starts Task instances
2. SchedulerJob starts “task1”



3. BackfillJob started for a single task - “task1”
4. BackfillJob overwrites scheduler’s DagRun
5. SchedulerJob forgets about “task2” and it never gets triggered

Backfill vs Scheduler

Upstream fix PR under review: <https://github.com/apache/airflow/pull/16089>

Typos in task regex


Typo in task regex

From the Backfill CLI command help output:

```
-t TASK_REGEX, --task-regex TASK_REGEX The regex to filter specific task_ids to backfill (optional)
```

If you made a typo and typed `--task-regex task3`

You will get:



Databricks clusters cost optimization

Databricks clusters cost optimization

Why?

- AWS has limited number of instances for each AZ
“We currently do not have sufficient capacity in the Availability Zone you requested”

Databricks clusters cost optimization

EC2 spot prices across availability zones:



Databricks clusters cost optimization

How?

- Custom Airflow Databricks operator
- AWS “Describe Spot Price History” API
- Take the cheapest AZ in AWS region
- Fallback to the next cheapest AZ

Databricks clusters cost optimization


Gain:

- 10-20% cost saving
- Reduce chances of running into AWS instance limit

Learn More:

<https://tech.scribd.com/blog/2020/optimize-databricks-cluster-configuration.html>

Q&A



tech.scribd.com
(we're hiring)