



Airflow 2.0 on Amazon MWAA

Airflow Summit 2021

Sam Dengler and John Jackson
July 15th, 2021

Table of contents


- How Amazon MWAA does Apache Airflow
- Staying current 2.x and beyond
 - Upgrades and new Apache Airflow version support
- Staying Open and contributing
 - AWS Contributions to Date
 - Support Plan for the Open Source Community

How Amazon MWAA does Apache Airflow



Amazon Managed Workflows for Apache Airflow (MWAA)

- A managed service for Apache Airflow that makes it easy for data engineers and data scientists to execute data processing workflows on AWS
- Released November 24, 2020, added Airflow 2.0 support May 26, 2021




Environments (5)					
	Name	Status	Created date	Airflow version	Airflow UI
<input type="radio"/>	MWAA-Demo-5	Available	May 25, 2021 17:26:01 (UTC-07:00)	2.0.2	Open Airflow UI
<input type="radio"/>	MWAA-Demo-4	Available	May 25, 2021 14:03:12 (UTC-07:00)	2.0.2	Open Airflow UI
<input type="radio"/>	MWAA-Demo-3	Available	May 24, 2021 18:03:09 (UTC-07:00)	2.0.2	Open Airflow UI
<input type="radio"/>	MWAA-Demo-2	Updating	Dec 14, 2020 08:28:22 (UTC-08:00)	1.10.12	Open Airflow UI
<input type="radio"/>	MWAA-Demo-1	Available	Dec 10, 2020 14:37:40 (UTC-08:00)	1.10.12	Open Airflow UI

Amazon MWAA is Open Source Apache Airflow: No Forks



Amazon MWAA Architecture



How Amazon MWAA Helps


- Deployments and Operations
 - Easy to setup up and maintain
- Availability and Sizing
 - Multi AZ/HA with Airflow 2.0 on ECS Fargate
- Scaling
 - Auto scaling with Celery executor
- Security
 - IAM and VPC



Setup



Upgrades



Scaling




Security



Maintenance

Deployments and Operations

- Easy to set up multiple small environments
- Prod/Dev environments
- CloudWatch
 - Logging
 - Metrics/Alarms
 - Dashboards
- CloudFormation/Terraform
 - Configure User Access, Environment (Execution Role) access, and Airflow Configurations without accessing the UI



Scaling/Environment Sizes

- Using Airflow metrics
- $(\text{Tasks running} + \text{tasks queued}) / \text{tasks per container} = \text{number of containers required}$
- Downscaling occurs when $(\text{Tasks running} + \text{tasks queued}) = 0$

Environment class Info	
Class	mw1.small
Scheduler count	2
Maximum worker count	10
Minimum worker count	1

Security with MWAA

- Auth (IAM/Federation)
- VPC/SG
- Secrets Manager
- Execution Role



CI/CD

- Versioned S3 Bucket deployment target
- Flexible integration with existing CI/CD pipelines
- DAGs automatically synchronized every 30-60s
- Custom Plugin and Python dependency changes require Environment update



Create an MWAA
Environment



Copy your DAGs
and Plugins to S3



Access the Airflow
UI

Amazon MWAA Best Practices

Configuration and Migration

Best Practices – Creating the Environment

- VPC Network
 - Connectivity to AWS Services
 - Private Web Server Option
- IAM Execution Role
 - Access to AWS Services
 - AOK vs CMK
- Python Requirements
 - Compatible Versions
 - Debugging
- Configuration Overrides
 - Performance Tuning
 - Secrets Backend Support

Environment details [Info](#)

Name
MyAirflowEnvironment
Use only letters, numbers, dashes, or underscores. Max 80 characters.

Airflow version
2.0.2 (Latest)

DAG code in Amazon S3 [Info](#)

Amazon MWAA uses your Amazon S3 bucket to load your DAGs and supporting files. Specify your S3 bucket, and the paths of your DAG folder, plugins.zip, and requirements.txt.

S3 bucket 

ⓘ Create or specify an S3 bucket to store your DAG code. The bucket name must have versioning enabled. You can create a new bucket in the [Amazon S3 console](#)

S3 Bucket
The S3 bucket where your source code is stored. Enter an S3 URI or browse and select a bucket.
 [View](#) [Browse S3](#)
Format: s3://mybucketname

Best Practices – Migration

- Moving tasks off the Worker that require
 - Local Docker
 - Custom runtimes
 - Large CPU/Memory
 - Sudo
- Assume Ephemeral Workers
 - Limited Bash Value
 - Local storage between tasks
- Web Server Restrictions
 - Not currently installing requirements.txt or plugins



Contributions and Staying Current



Staying Current

- Target is to provide minor versions on Amazon MWAA within 30 days of Apache Airflow community release
 - We're not there yet, but we're working on it
- Major releases will take a bit longer
- We'll keep older versions (i.e. 1.10.x) available for as long as we can
 - Depends on security, stability, and customer demand
- Version status will be posted at
<https://docs.aws.amazon.com/mwaa/latest/userguide/airflow-versions.html>

Open Source Contributions

- EKS Native operator
 - Improved Apache Airflow AWS integrations
 - Performance and security improvements
 - Native Serverless executor

Resources

- Amazon MWAA docs <https://docs.aws.amazon.com/mwaa>
- Amazon MWAA product page <https://aws.amazon.com/mwaa>
- #airflow-aws Slack Channel: <https://apache-airflow.slack.com>
- GitHub samples: <https://github.com/aws-samples/amazon-mwaa-examples>
- Local Runner <https://github.com/aws/aws-mwaa-local-runner>

Q&A

Sam Dengler @*samdengler*

John Jackson @*JohnJacksonPM*