

Airflow in an On-premise Data Mesh Setup

adyen

Speaker

Jorrick Sleijster

Data Engineer

Data Platform Development

Tennis & Padel tennis



Adyen

We provide a single payment platform
globally to accept payments and grow
revenue online, on mobile, and at the point of
sale

adyen

 Apache
Airflow

FACEBOOK

LinkedIn



Microsoft

Bla Bla Car

Uber

Evernote

Spotify

AliExpress™

Dropbox



eBay

The Workflow of the Big Data Platform



Fully on-premise

We have all servers in-house. We manage everything from hardware to software. Different clusters per env: beta, test and live.



Weekly releases

Each week we start with a new release on beta. Two days later deployed to test, and finally a week later to live.

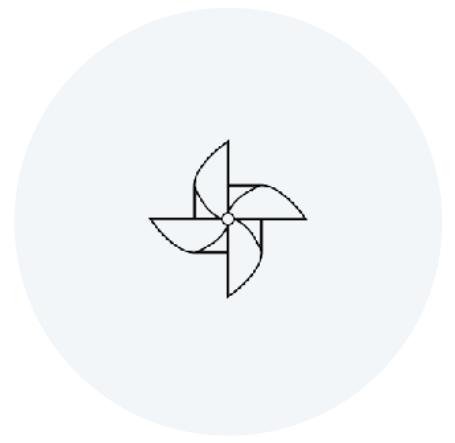


Teams / Streams

We have 20+ streams and 100+ data scientists on the big data platform. Streams are responsible for their own DAGs.

Airflow Cluster Setup at Adyen

Airflow 2.2



Celery workers



Yarn Queue



Postgres DB

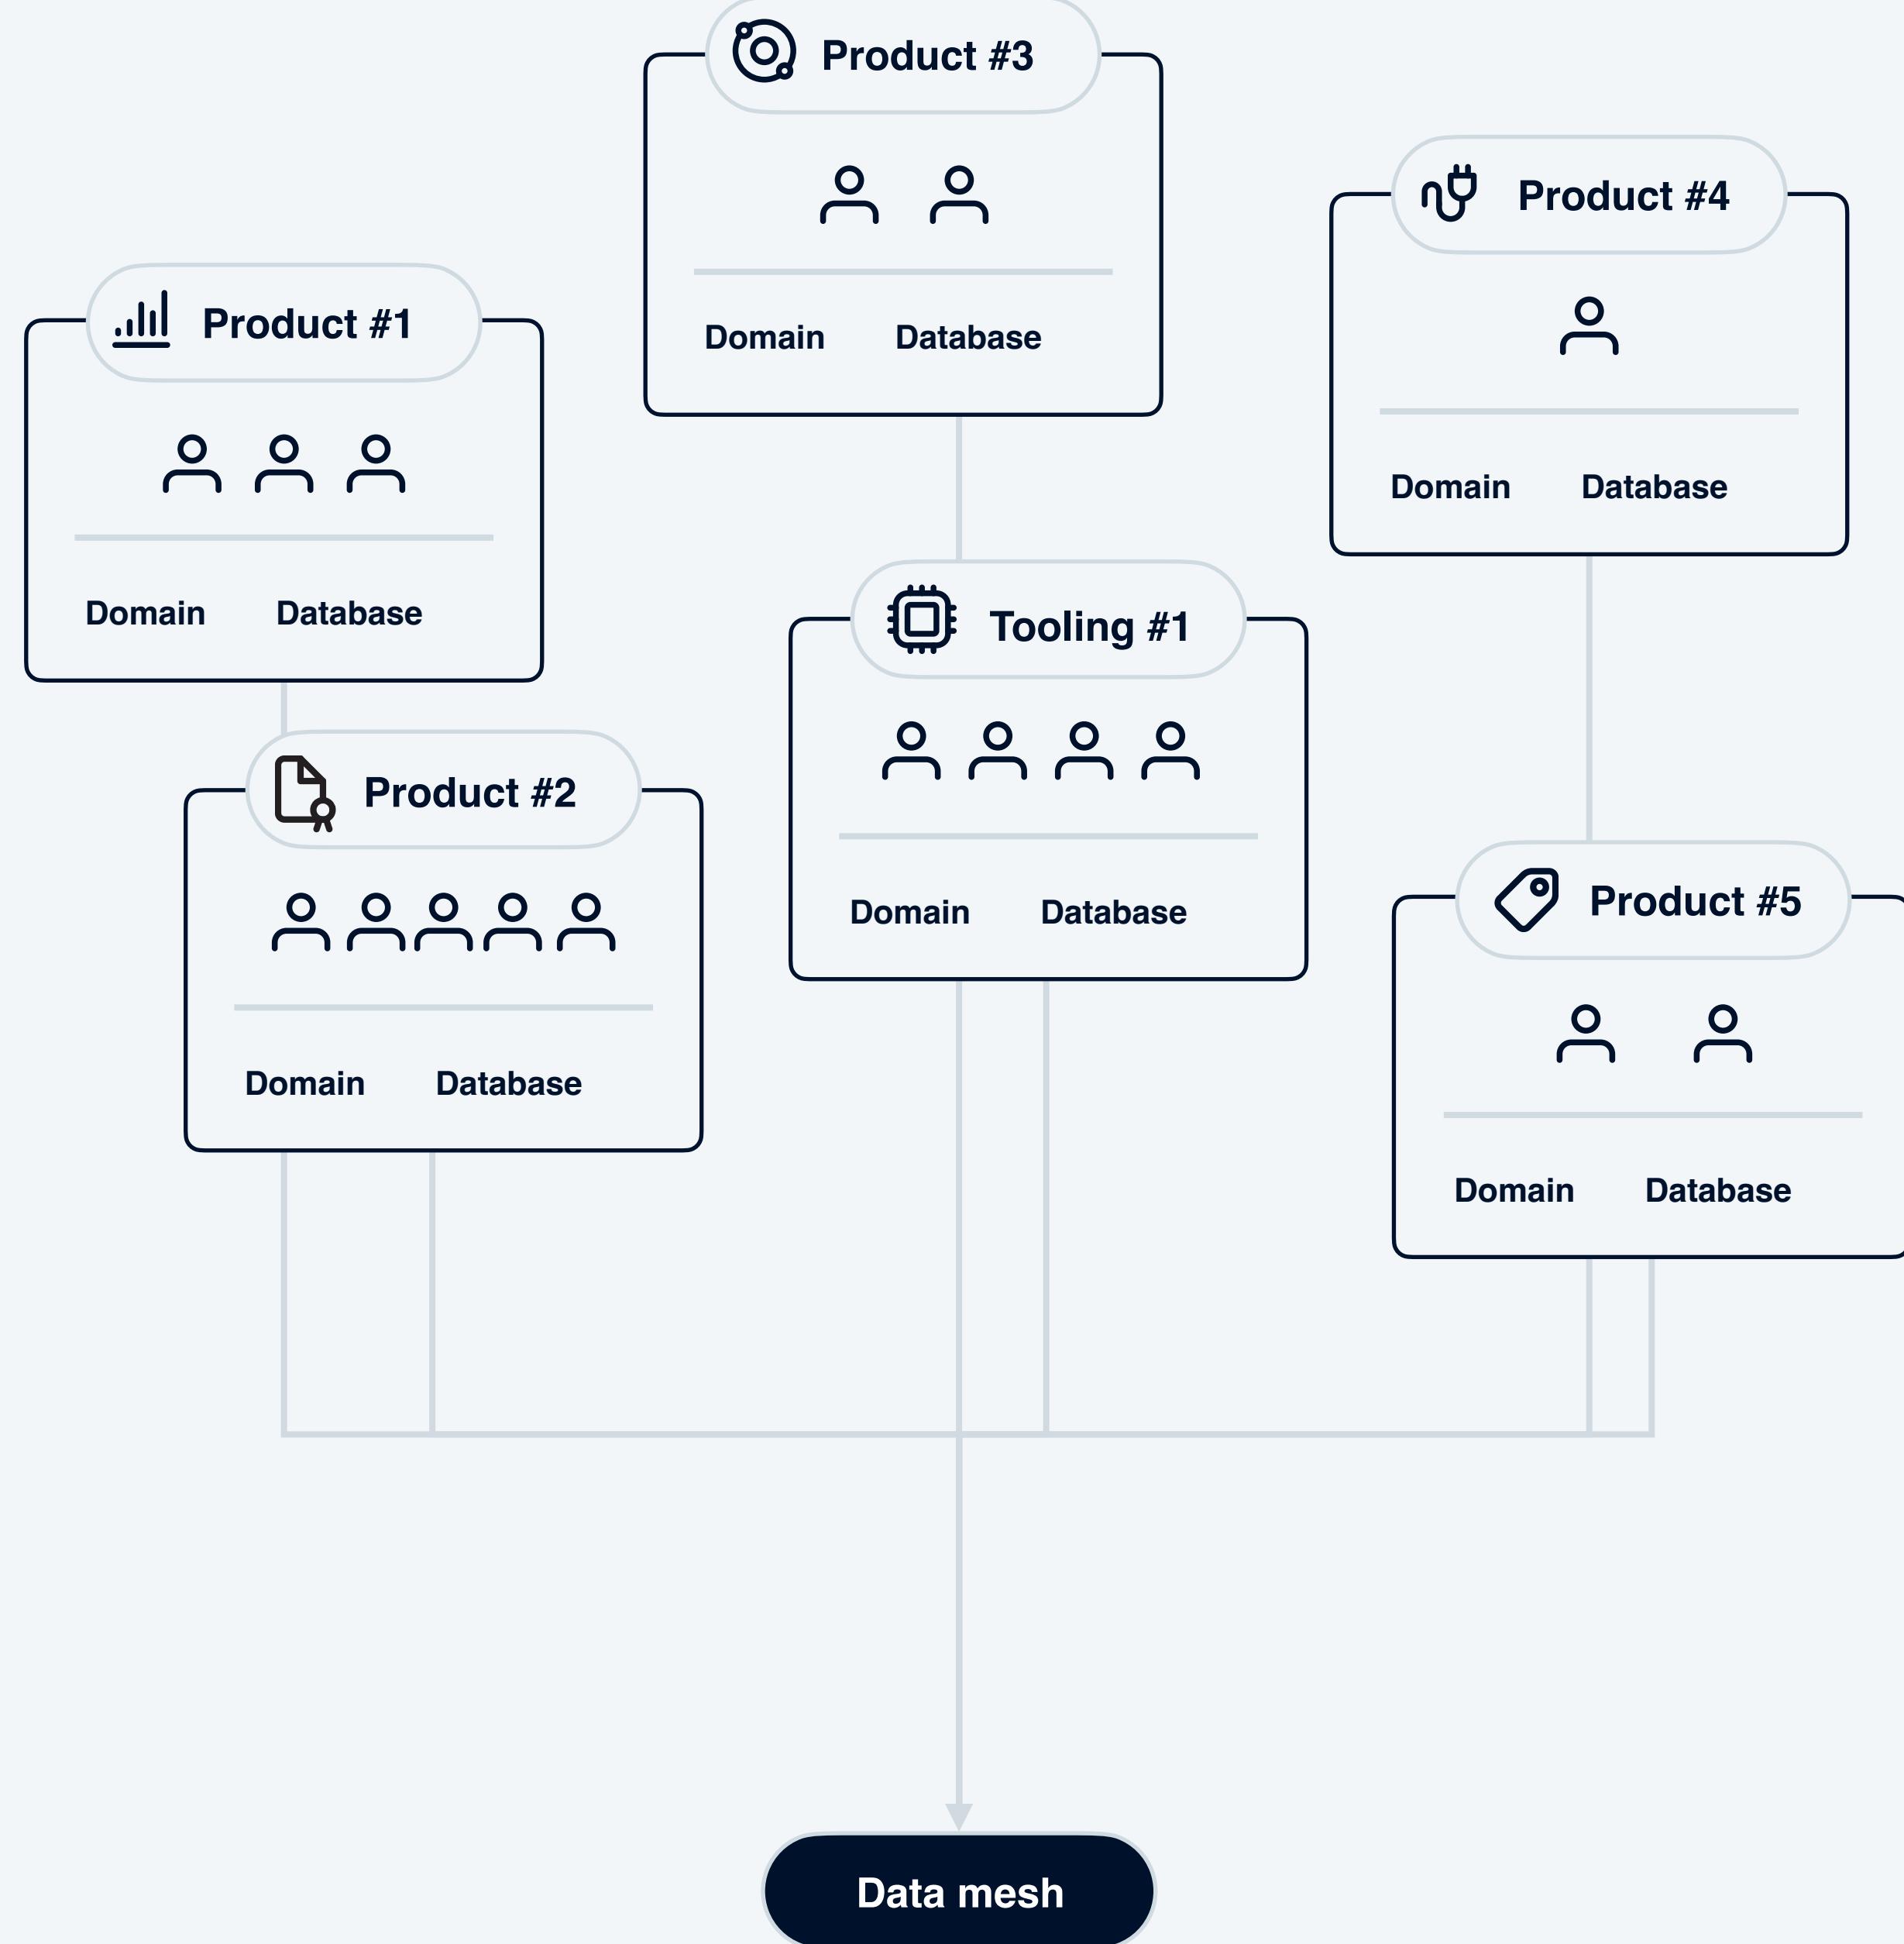


adyen extensions

Models, views, database listeners

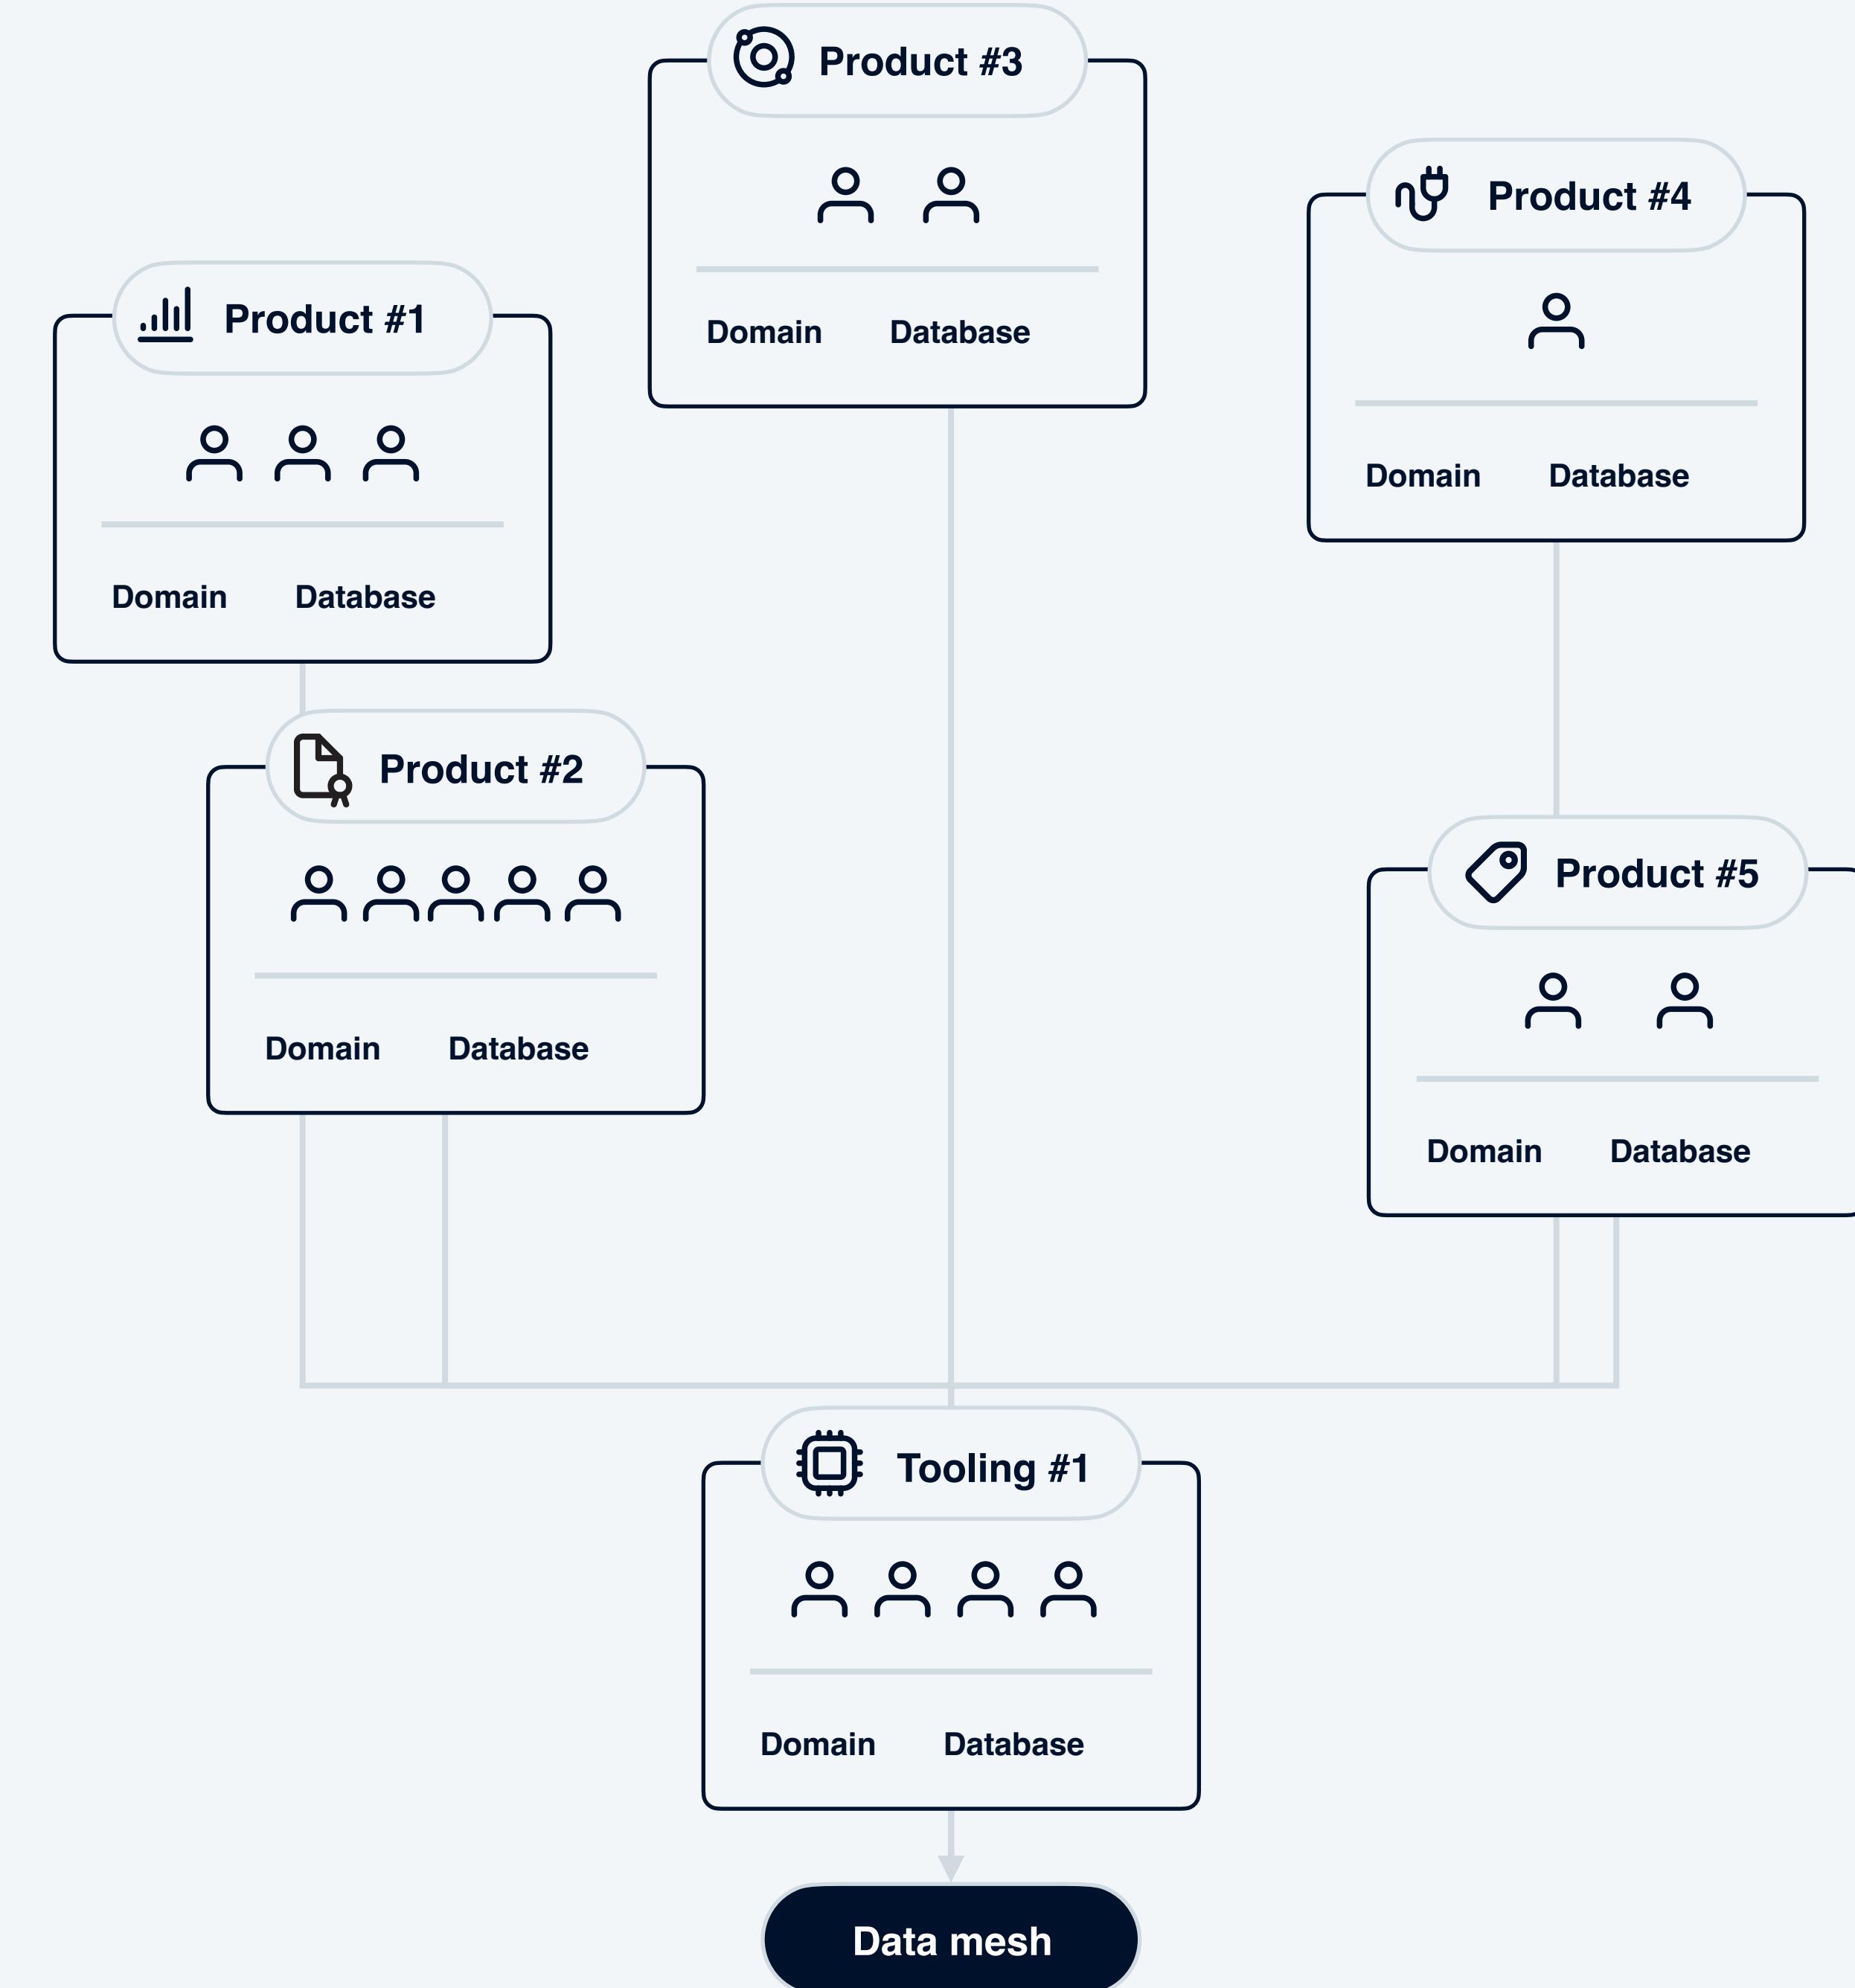
Why Data Mesh

- What is a Data Mesh?
 - Monolithic data infrastructure
 - Distributed domains
 - Domain responsible for their own ETLs & data
- Product teams(domains) with a clear focus
 - Front-enders, back-enders, data scientists
- No central ETL team
- Self service



Data Mesh Teams

- Product teams
- Tooling teams
- Infrastructure teams

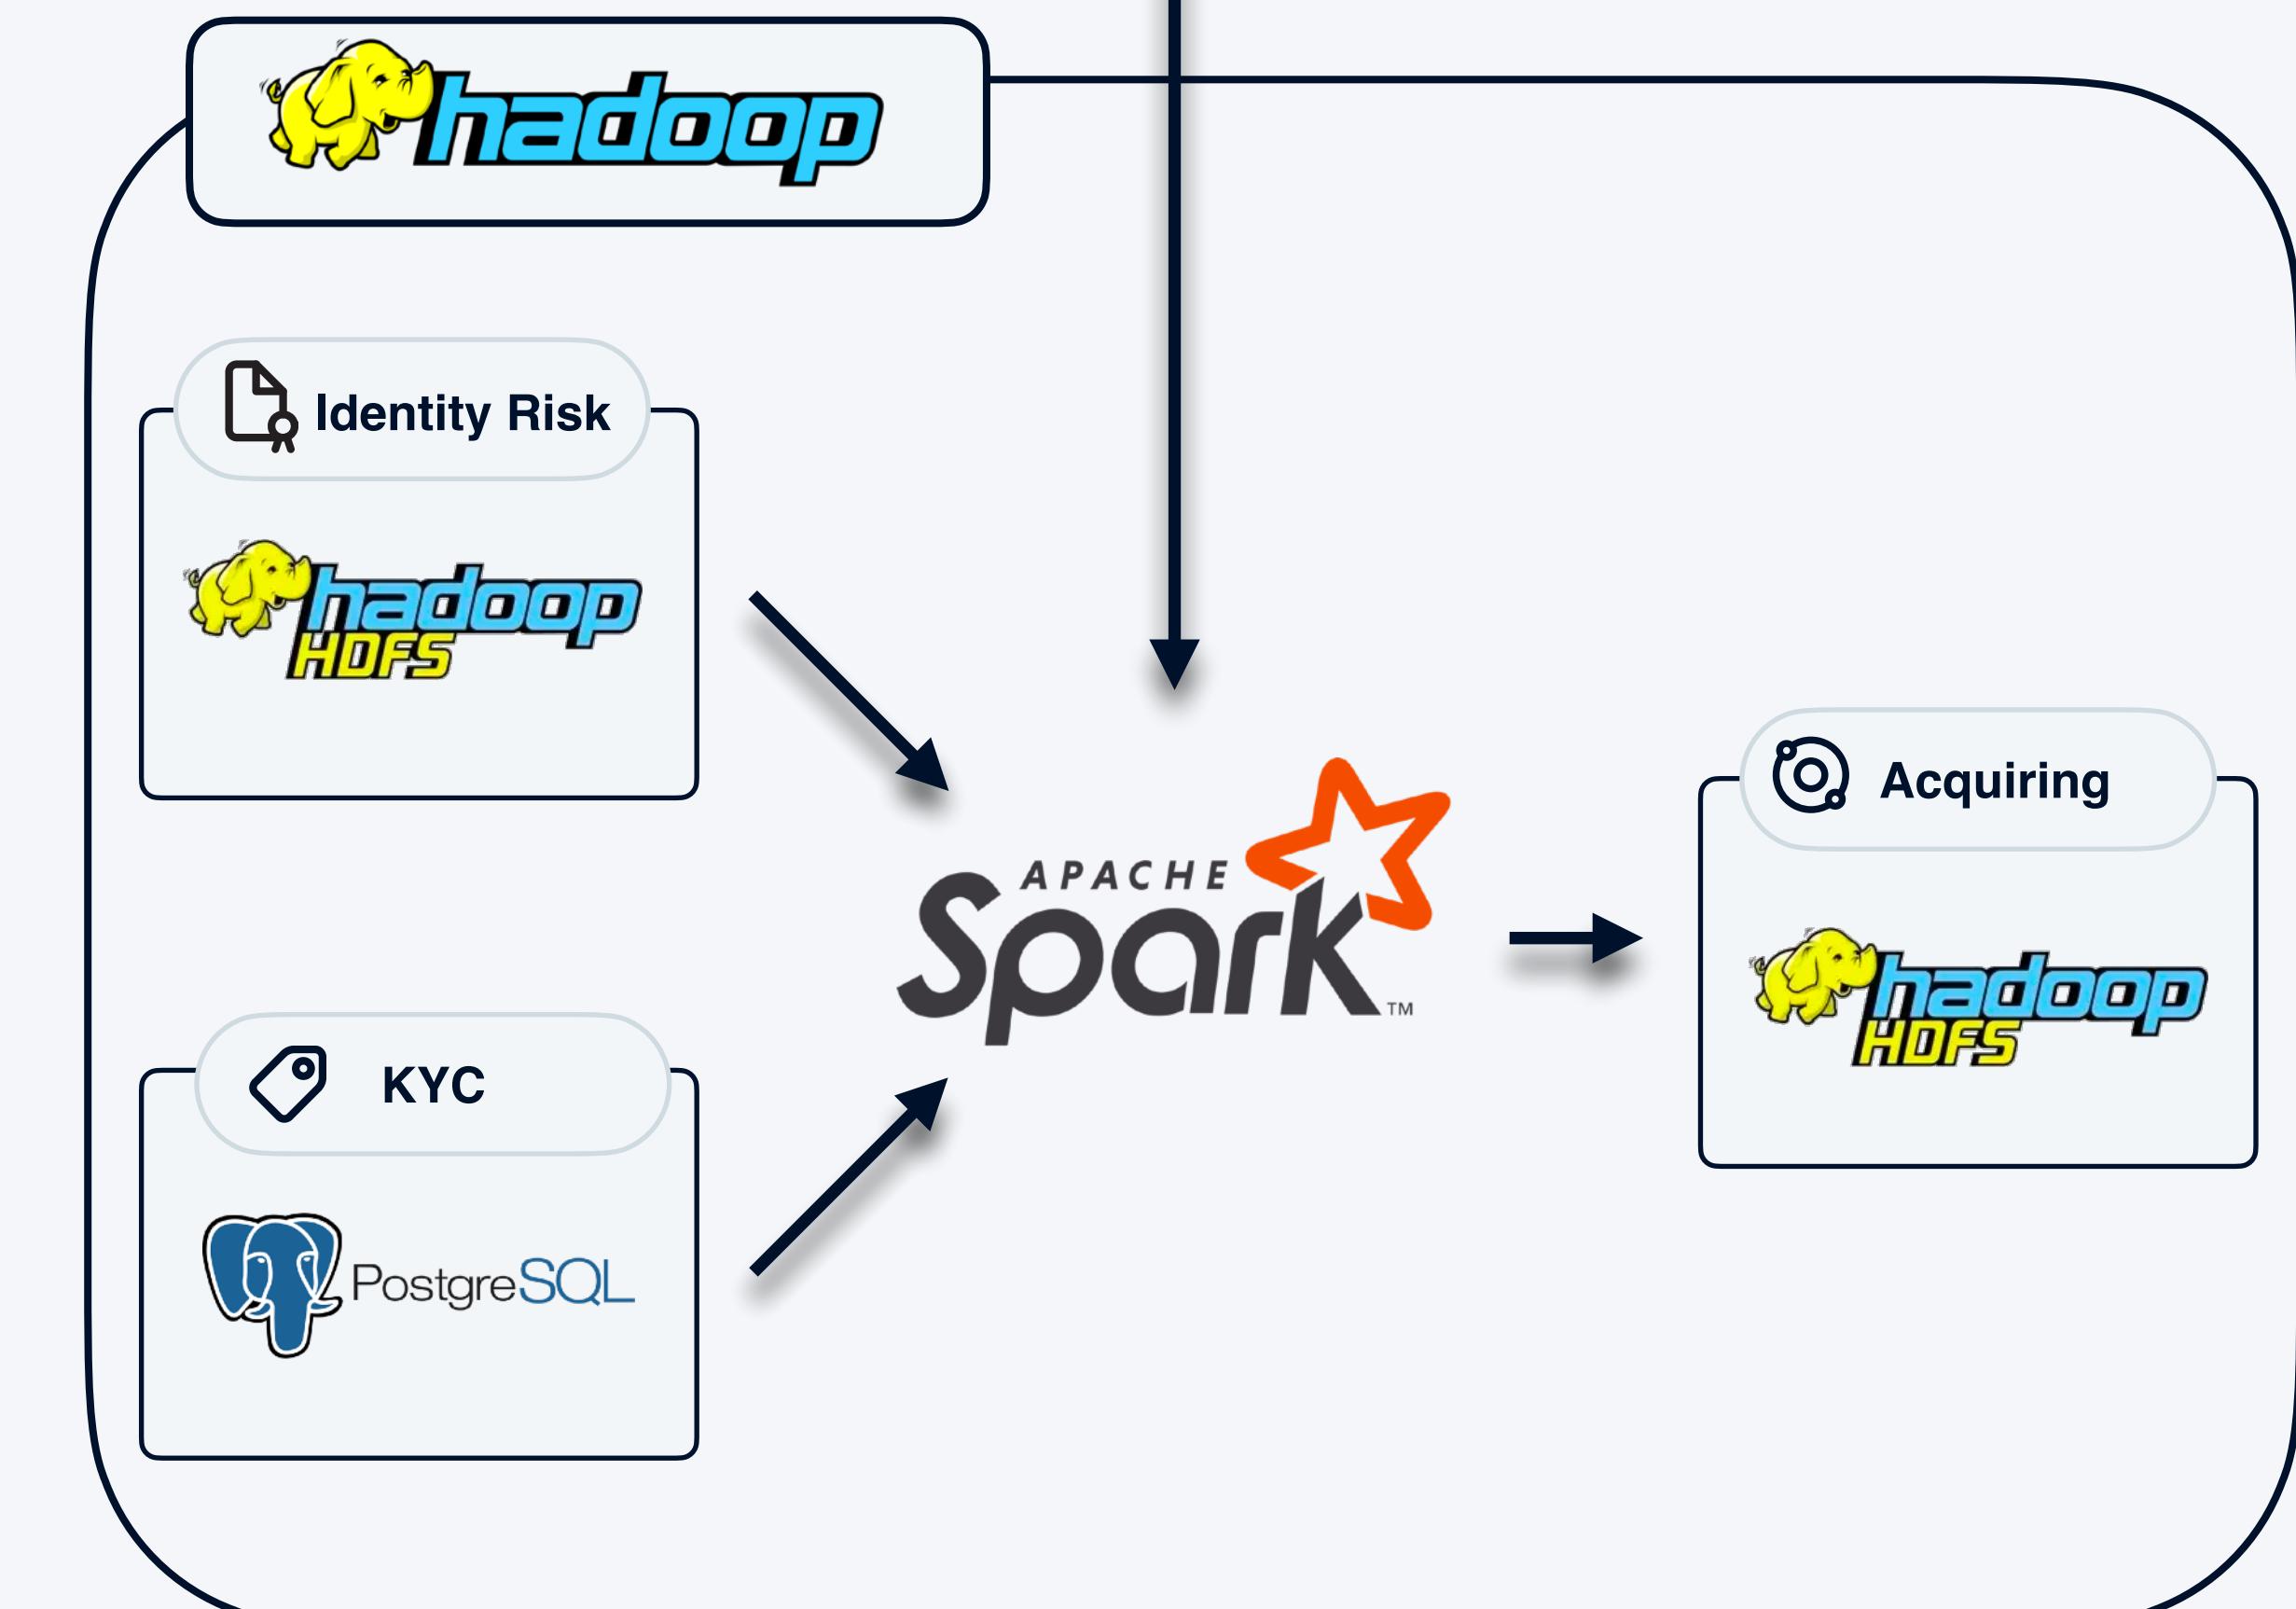


Creating a (new) Data Pipeline

- Create an Airflow DAG
- Create the ETL code



starts



Creating a (new) Data Pipeline

- Create an Airflow DAG
- Create the ETL code
- We need to:
 - Define the schema
 - Have enough resources for the ETL
 - Possibly handle updates of the schema
 - Retention, we don't have infinite storage
 - Remove the table
 - Don't want to (accidentally?) modify other teams resources.

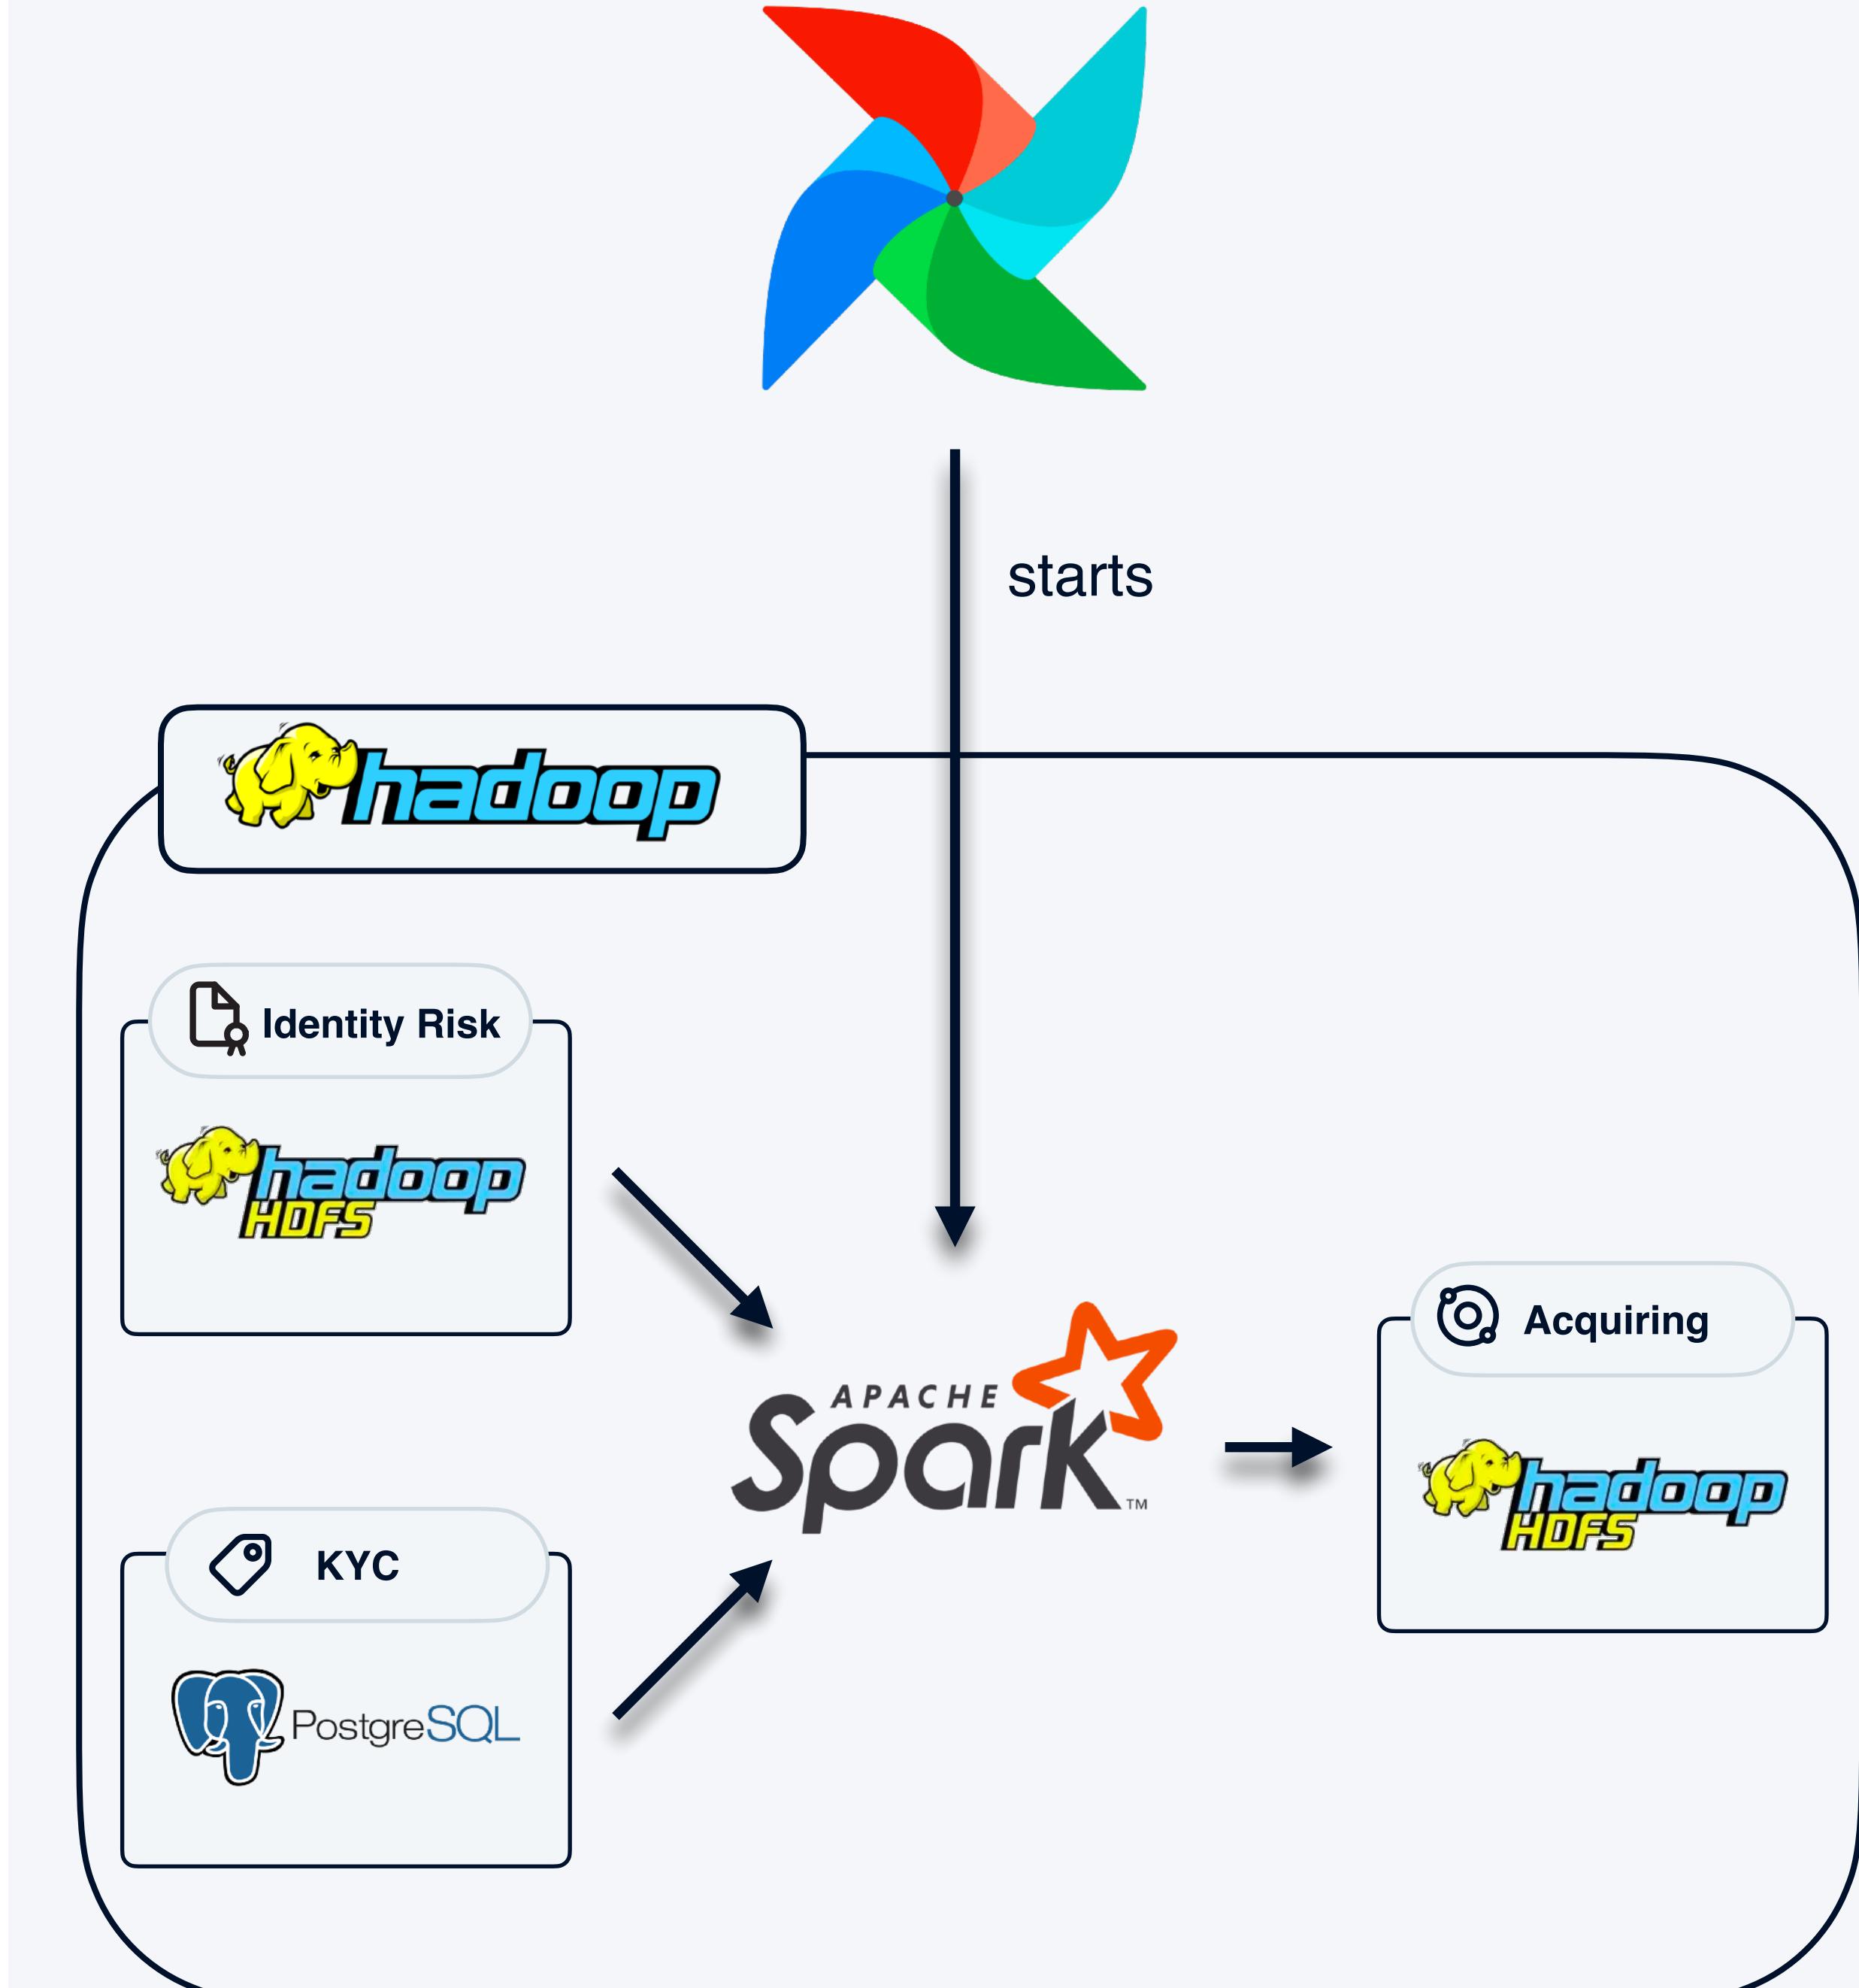


Table Schemas

How can we enable schema evolution?

Tuning Resources

How can we handle data outgrowing the predefined ETL resources?

Retention Period

How can we save HDFS from being filled with data we don't need anymore?

User Permissions

How can we give users access to only the parts that they need?



Table Schemas

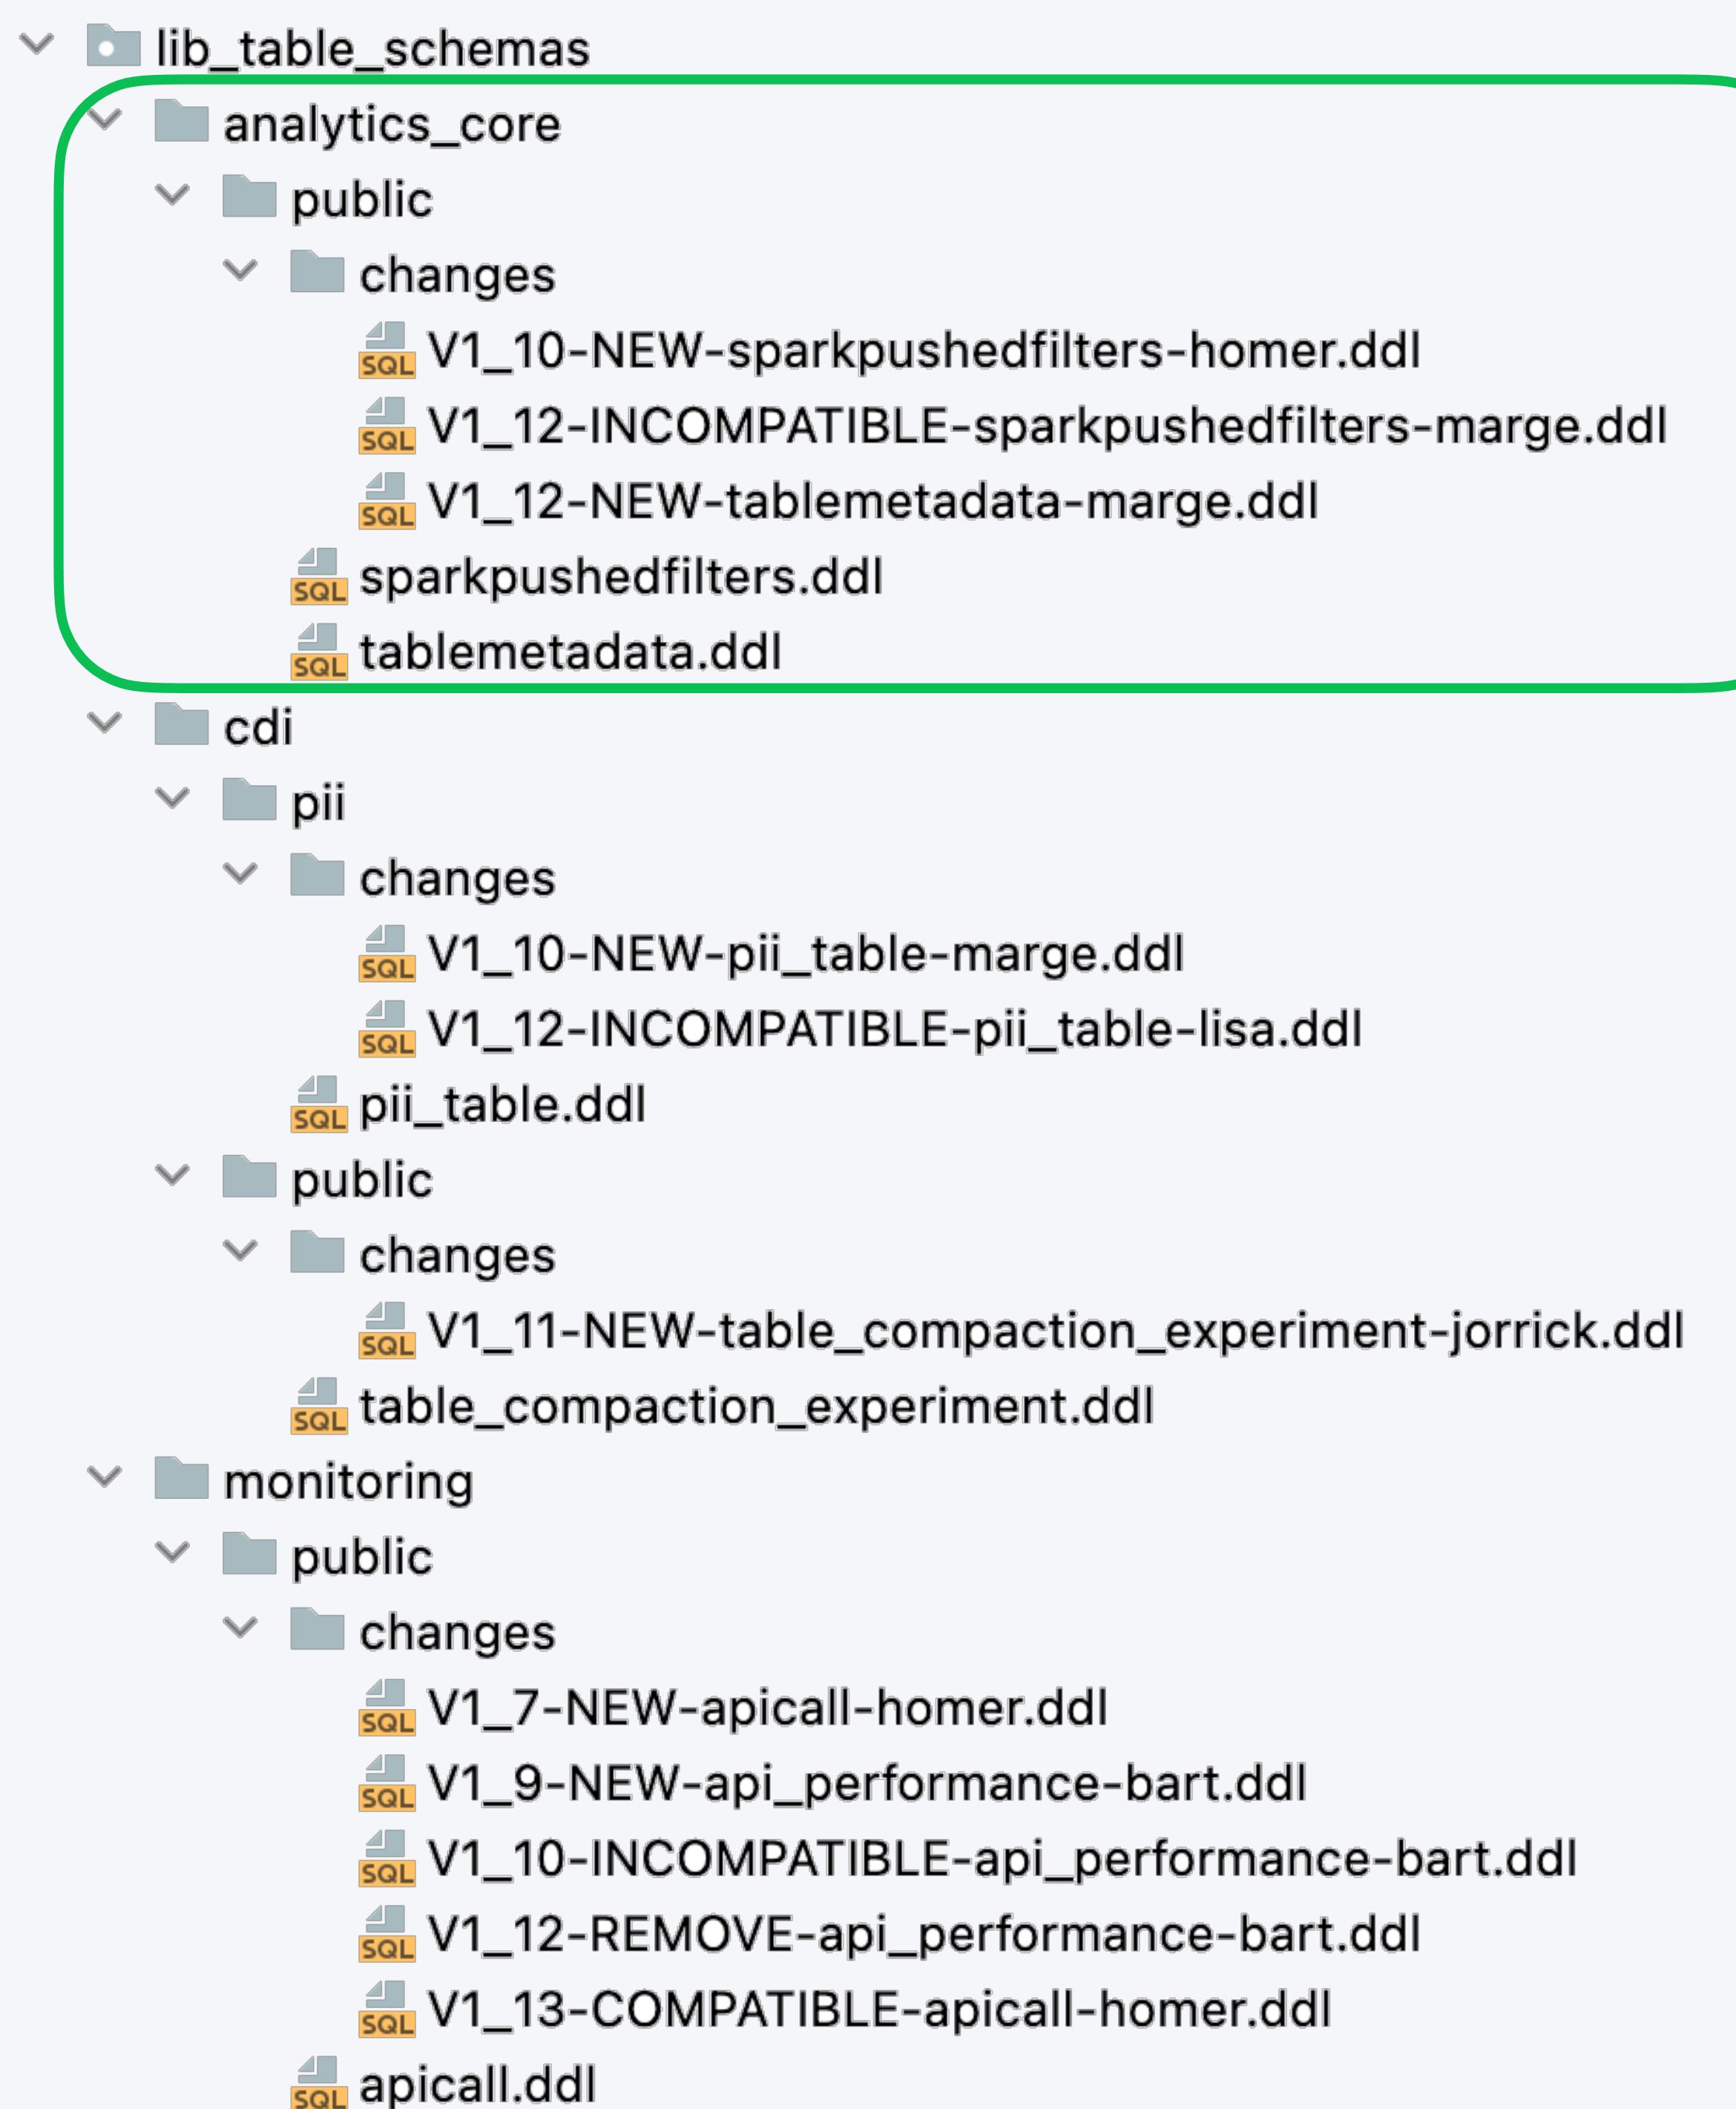
- We are managing an ETL pipeline
 - How & where do we define the schema?
 - How do we update the schema?
 - How do we delete tables we no longer need?

```
1 CREATE TABLE `parquetfilesize` (  
2   `file_size` BIGINT,  
3   `database` STRING,  
4   `table_name` STRING,  
5   `file_path` STRING,  
6   `human_readable_size` STRING,  
7   `write_timestamp` BIGINT,  
8   `partition_date` STRING,  
9   `snapshot_timestamp` STRING,  
10 )  
11 USING parquet  
12 OPTIONS (  
13   `serialization.format` '1',  
14   path '{path}'  
15 )  
16 PARTITIONED BY (partition_date)  
17 TBLPROPERTIES (  
18   sort_keys = 'partition_date, database, table_name',  
19   repartition_factor = 1  
20 )
```

Table Schemas Library

- Each domain has their own database.
- Each database has multiple scopes:
 - pii - Requires white-listed access.
 - private - Access by domain members only.
 - public - Access by everyone.
- Each database scope can have multiple tables.
- Each table has at least one change file.
- Change types: new, (in)compatible, remove

Tool to create schemas from a Spark DataFrame.



Process Flow

1

Update the schema file (.ddl)

2

Introduce the schema change file (.ddl)

3

Create an MR (merge request)

4

Teammate reviews and approves MR

5

MR is merged

```
1 CREATE TABLE `parquetfilesize` (  
2   `file_size` BIGINT,  
3   `database` STRING,  
4   `table_name` STRING,  
5   `file_path` STRING,  
6   `human_readable_size` STRING,  
7   `write_timestamp` BIGINT,  
8   `partition_date` STRING,  
9   `snapshot_timestamp` STRING,  
10 )  
11 USING parquet  
12 OPTIONS (  
13   `serialization.format` '1',  
14   path '{path}'  
15 )  
16 PARTITIONED BY (partition_date)  
17 TBLPROPERTIES (  
18   sort_keys = 'partition_date, database, table_name',  
19   repartition_factor = 1  
20 )
```

Process Flow

- 1 Update the schema file (.ddl)
- 2 Introduce the schema change file (.ddl)
- 3 Create an MR (merge request)
- 4 Teammate reviews and approves MR
- 5 MR is merged
- 6 Release a new version on the cluster
- 7 Duty rolls out all schema change files for release
(Optional) User can also perform ad-hoc table changes like redefining the table with the latest schema.

not-so-live demo



localhost Airflow

Airflow DAGs Security Browse Admin Docs Duty Cluster: dev-local 10:30 UTC AA

Table management View

Logs

roll_out_new_table_changes_and_avro_tables 2022-05-13T05:12:59

Ad-hoc Table Changes

Hive Database ⓘ

Table Name ⓘ

Table Privacy: ⓘ Public Private PII

Delete hive table ⓘ
 Delete table data ⓘ
 Define hive table & Do MSCK repair ⓘ

Roll Out Change

Roll out Table Changes

Please select the changes which you want to roll out from the table below.

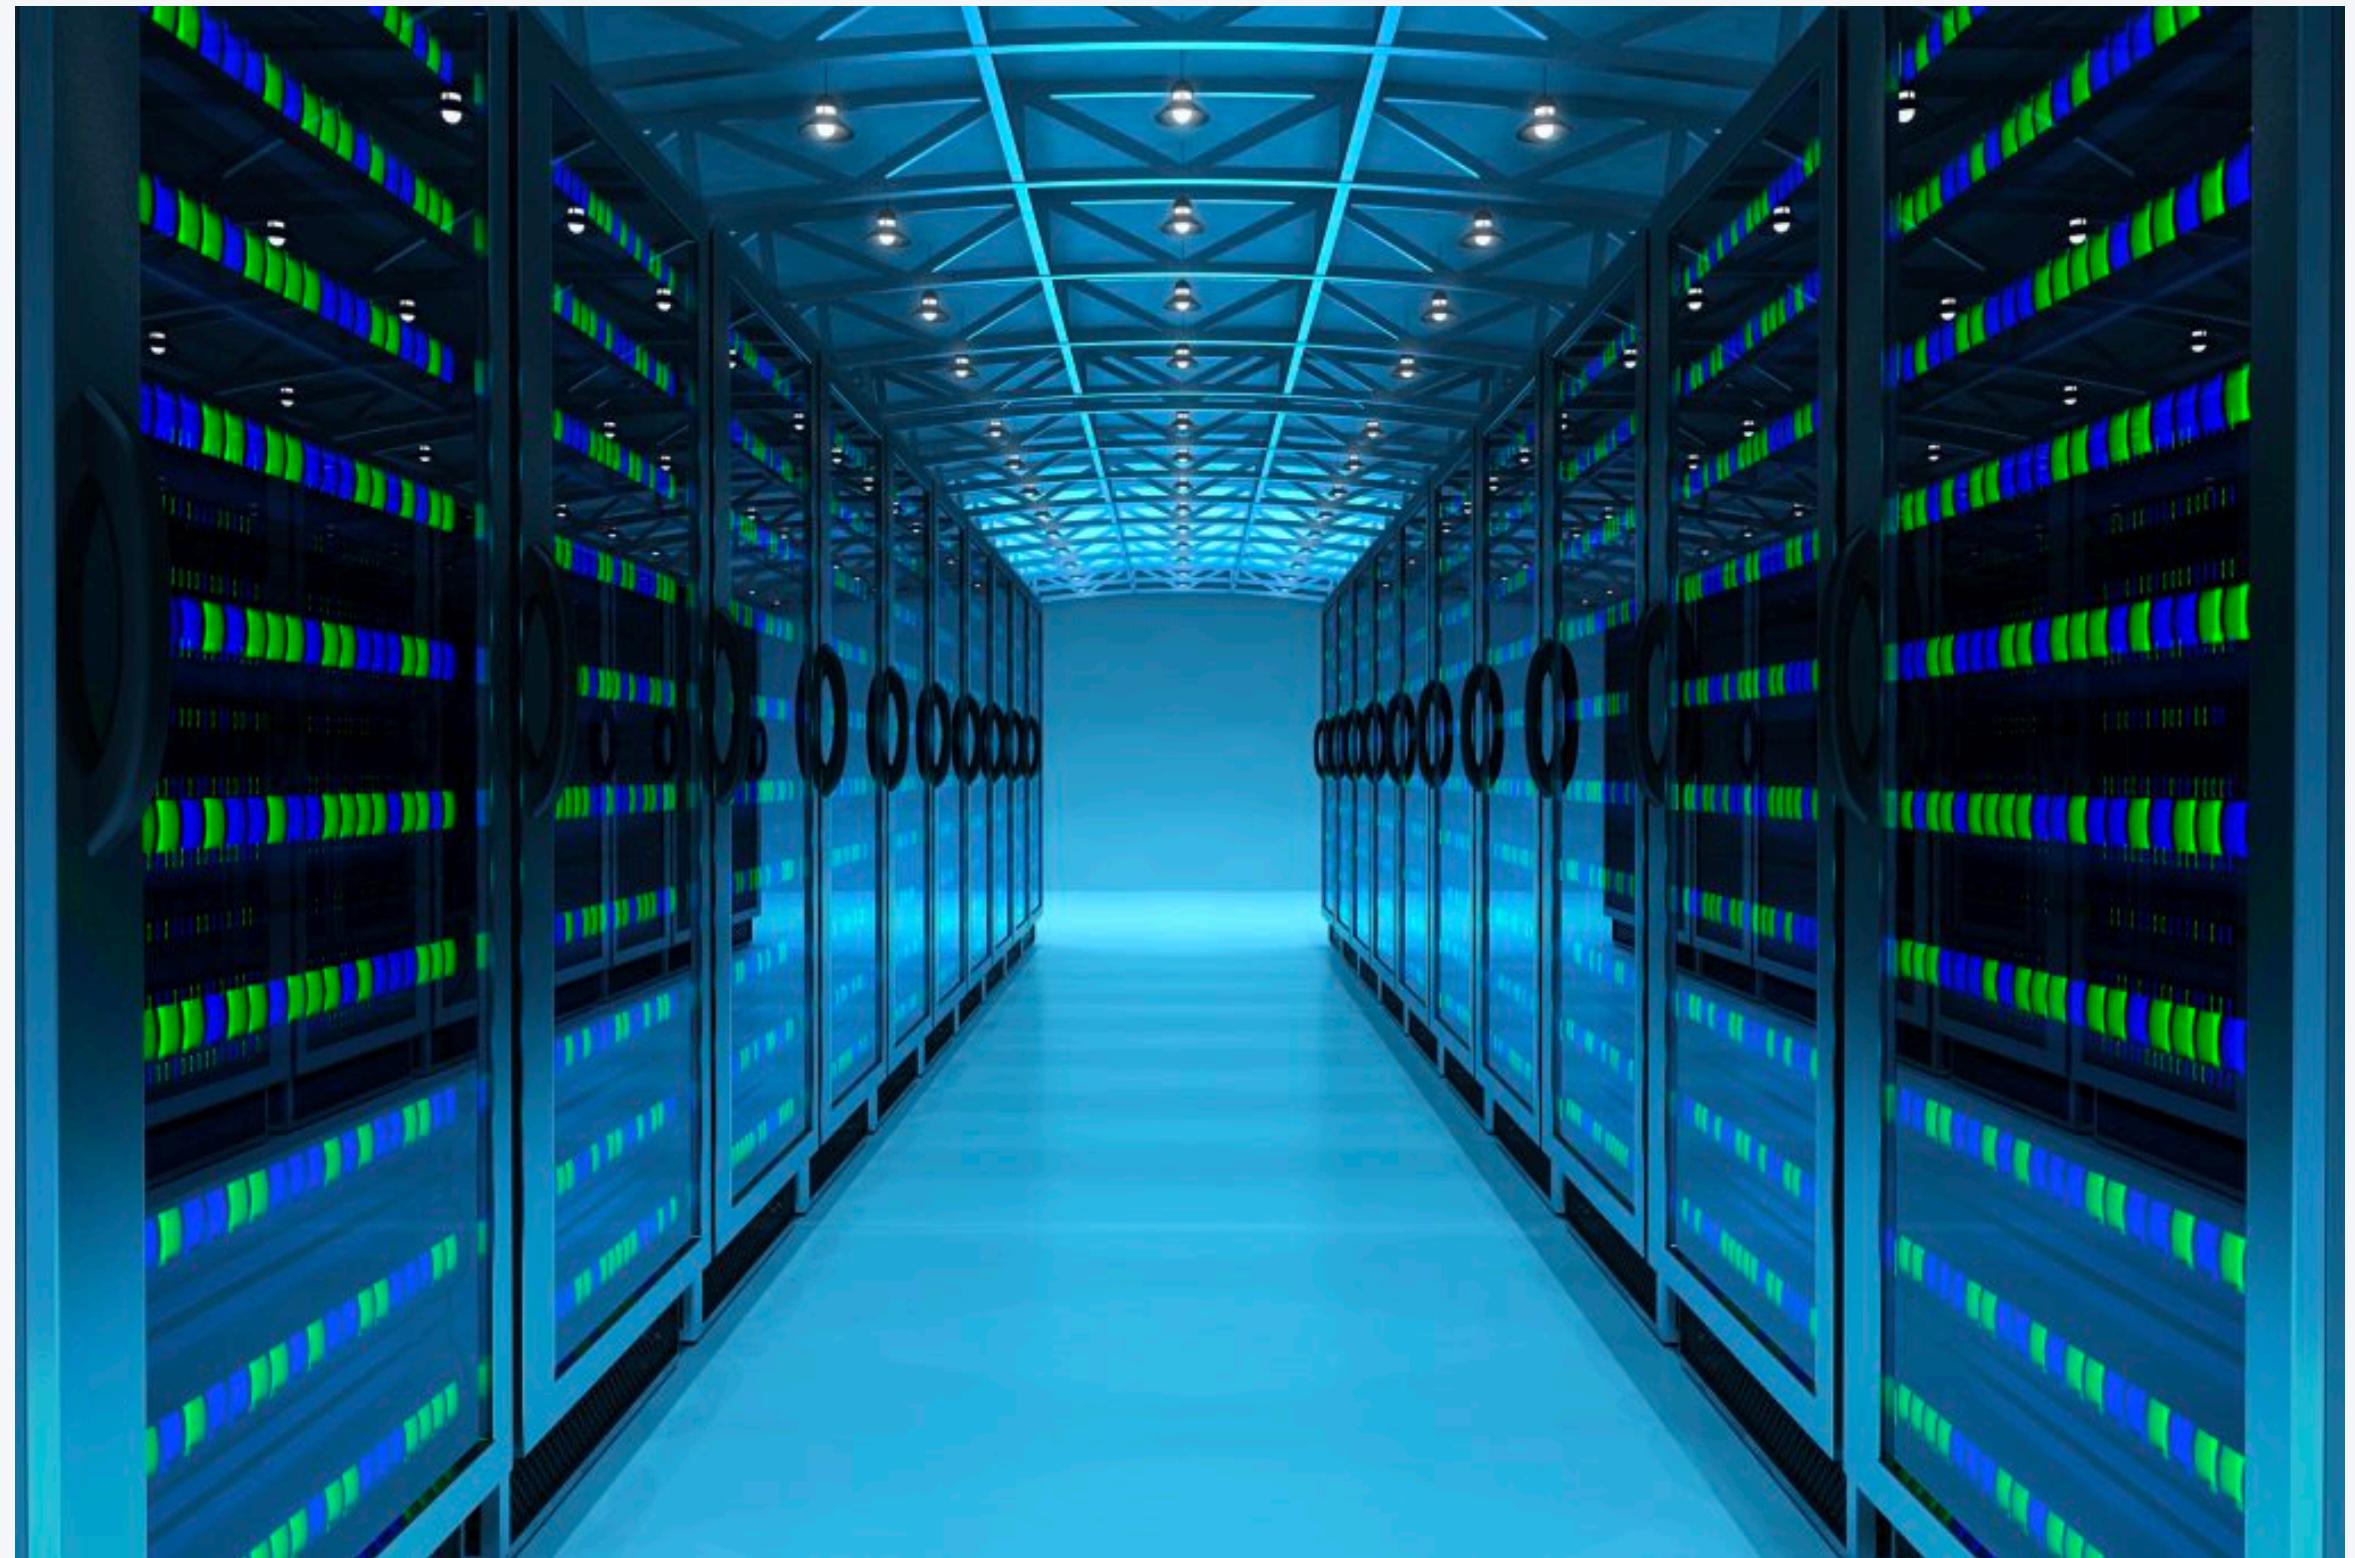
Filter on release version ⓘ Select a release ▾ Filter on other properties ⓘ Filter changes Reset selection Roll Out Changes

Showing 1 to 10 of 10 entries

RC	Change Type	Database	Table Privacy	Table Name	Requested by	Status
V1_12	REMOVE	monitoring	public	api_performance	bart	-
V1_12	NEW	analytics_core	public	tablemetadata	marge	-
V1_12	INCOMPATIBLE	cdi	pii	pii_table	llsa	-
V1_12	COMPATIBLE	analytics_core	public	sparkpushedfilters	marge	-

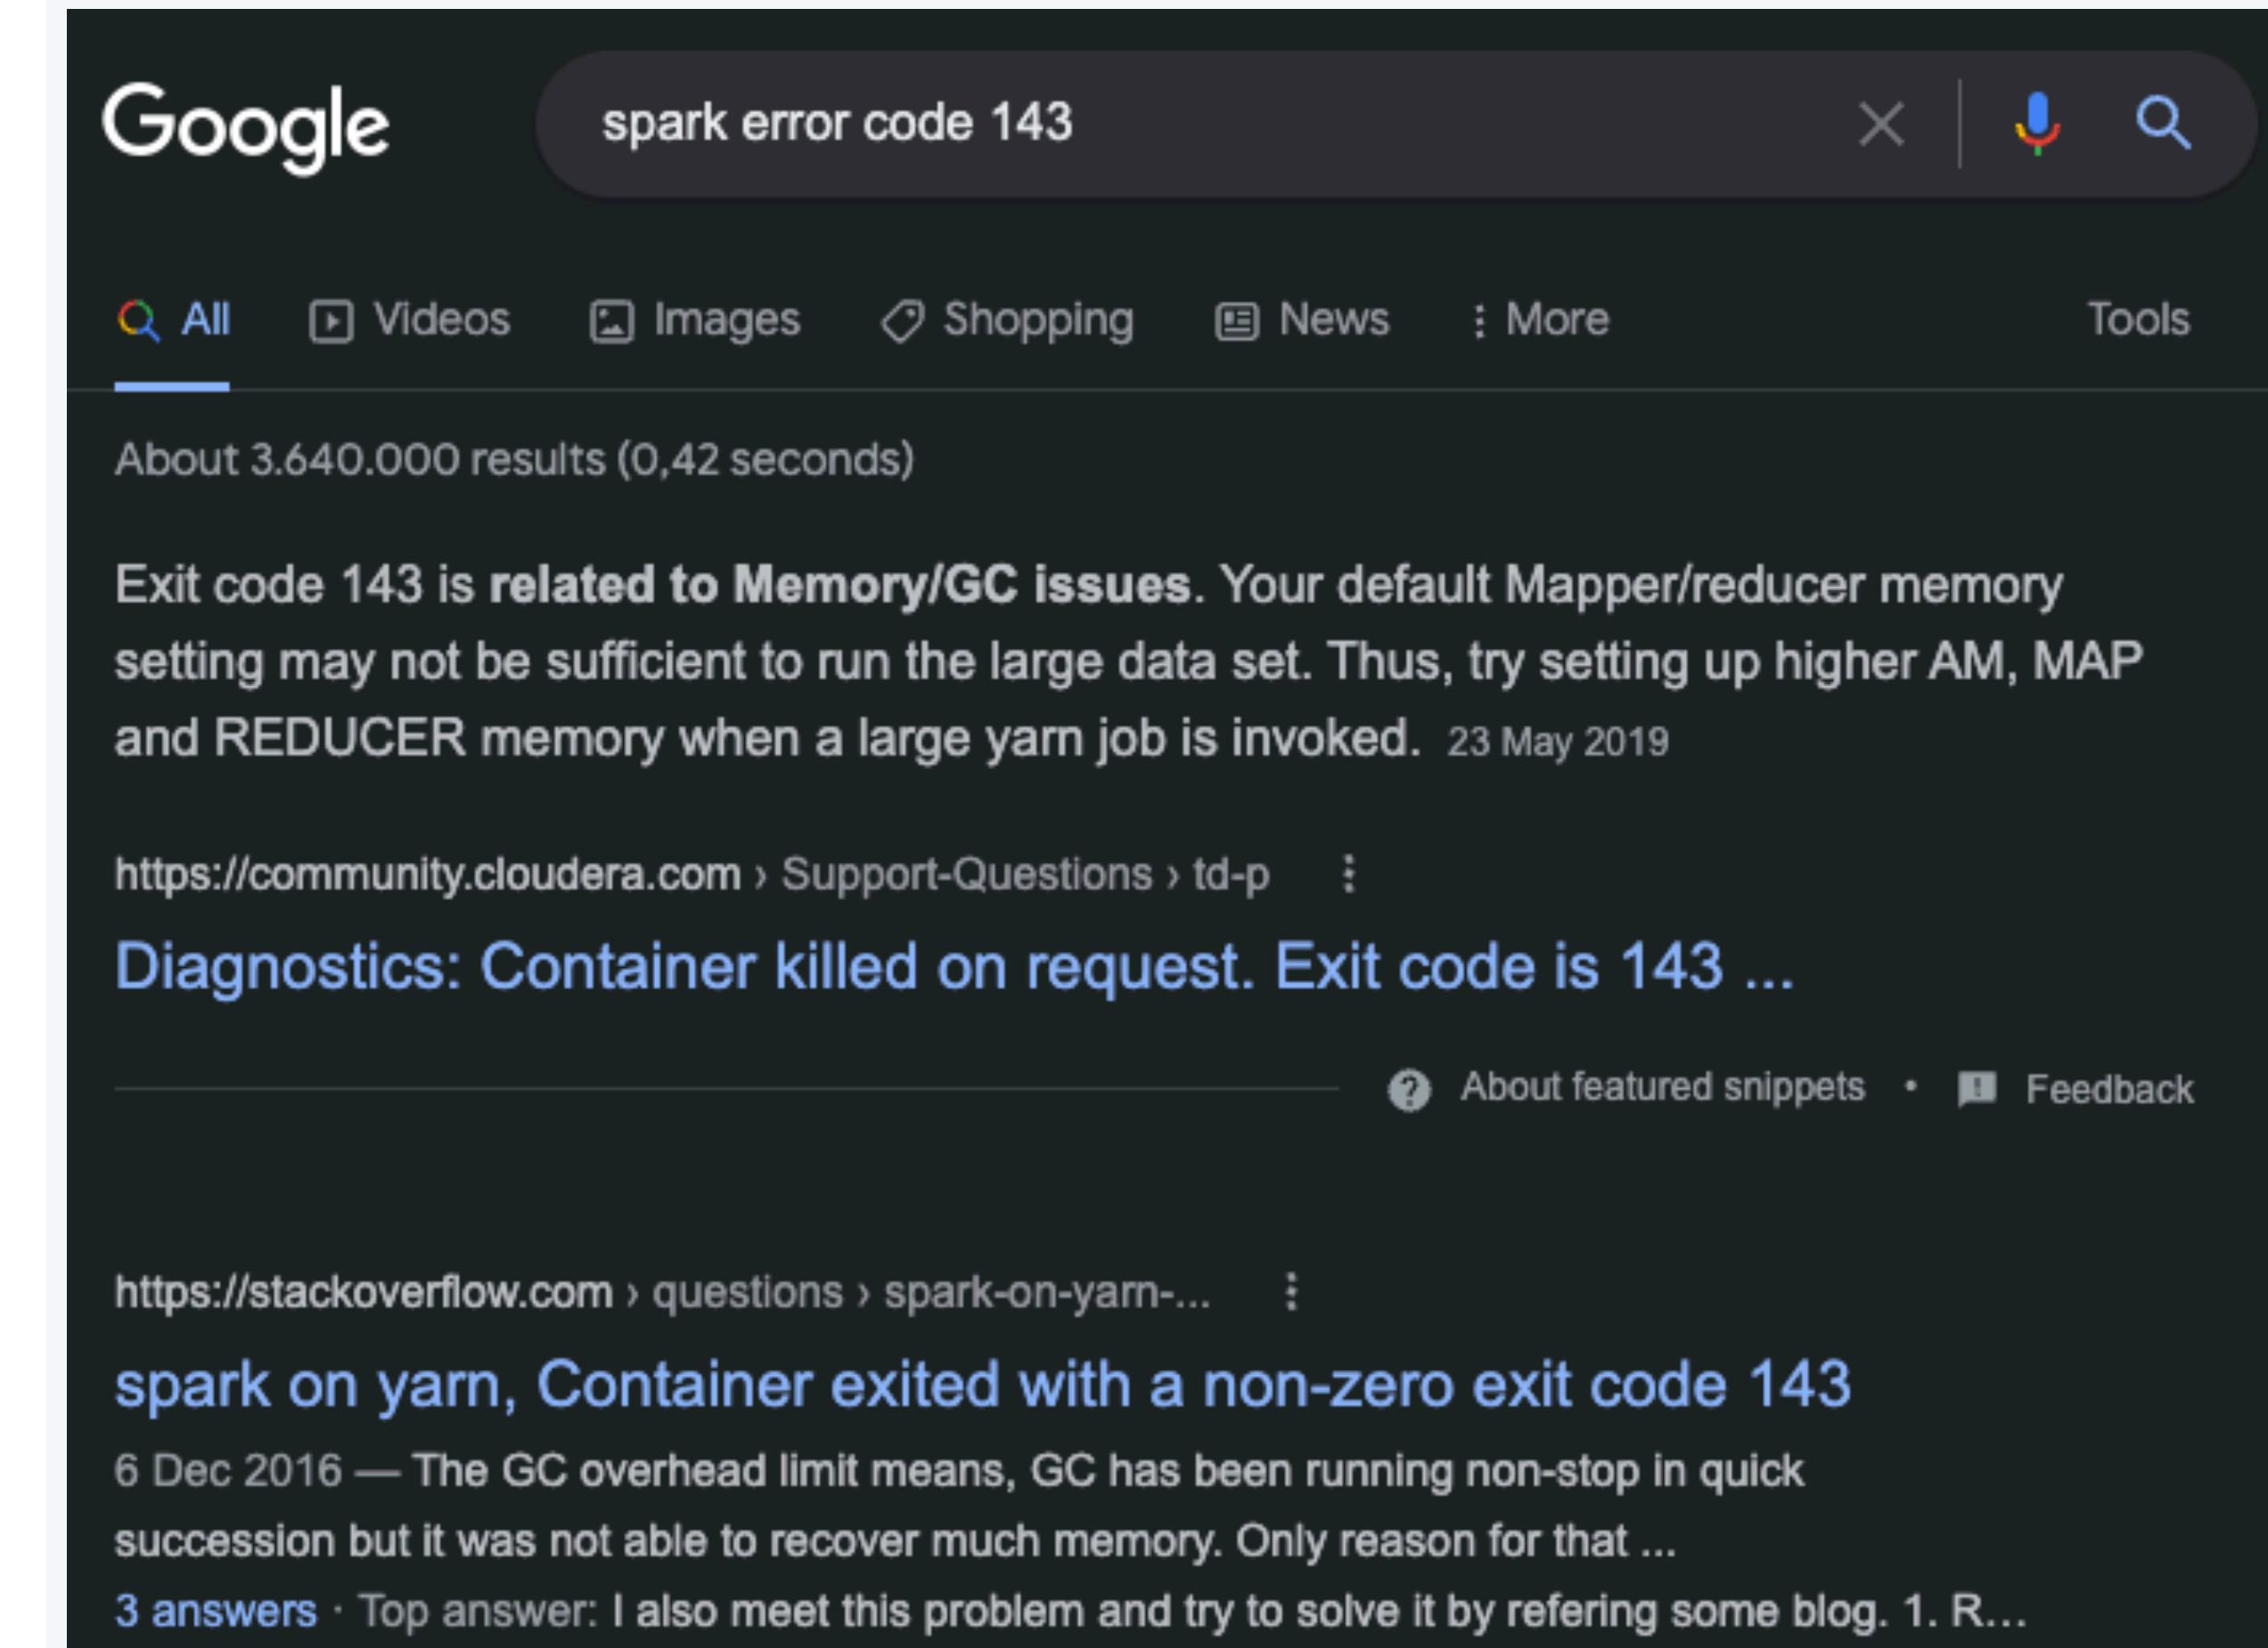
Tuning Resources

- As we are on-premise, we have finite resources
- We hardcode the resources in our code base
- Now imagine:
 - A DAG that has been running for 100 days
 - On day 101 you get ***error code 143***
 - What do you do?



The Usual Workflow

- You google the issue
- Solution: Set a higher Spark driver and/or worker memory limit



A screenshot of a Google search results page. The search query "spark error code 143" is entered in the search bar. The results show approximately 3,640,000 results found in 0.42 seconds. The top result is a snippet from a Cloudera support forum post, which states: "Exit code 143 is related to Memory/GC issues. Your default Mapper/reducer memory setting may not be sufficient to run the large data set. Thus, try setting up higher AM, MAP and REDUCER memory when a large yarn job is invoked." Below this, there are two more snippets. The first is from Stack Overflow, titled "spark on yarn, Container exited with a non-zero exit code 143", dated Dec 6, 2016. It discusses GC overhead limits. The second snippet is a link to a blog post with 3 answers, mentioning that the problem is related to memory management.

Google spark error code 143

All Videos Images Shopping News More Tools

About 3.640.000 results (0,42 seconds)

Exit code 143 is related to Memory/GC issues. Your default Mapper/reducer memory setting may not be sufficient to run the large data set. Thus, try setting up higher AM, MAP and REDUCER memory when a large yarn job is invoked. 23 May 2019

<https://community.cloudera.com> › Support-Questions › td-p ...

Diagnostics: Container killed on request. Exit code is 143 ...

About featured snippets • Feedback

<https://stackoverflow.com> › questions › spark-on-yarn-... ...

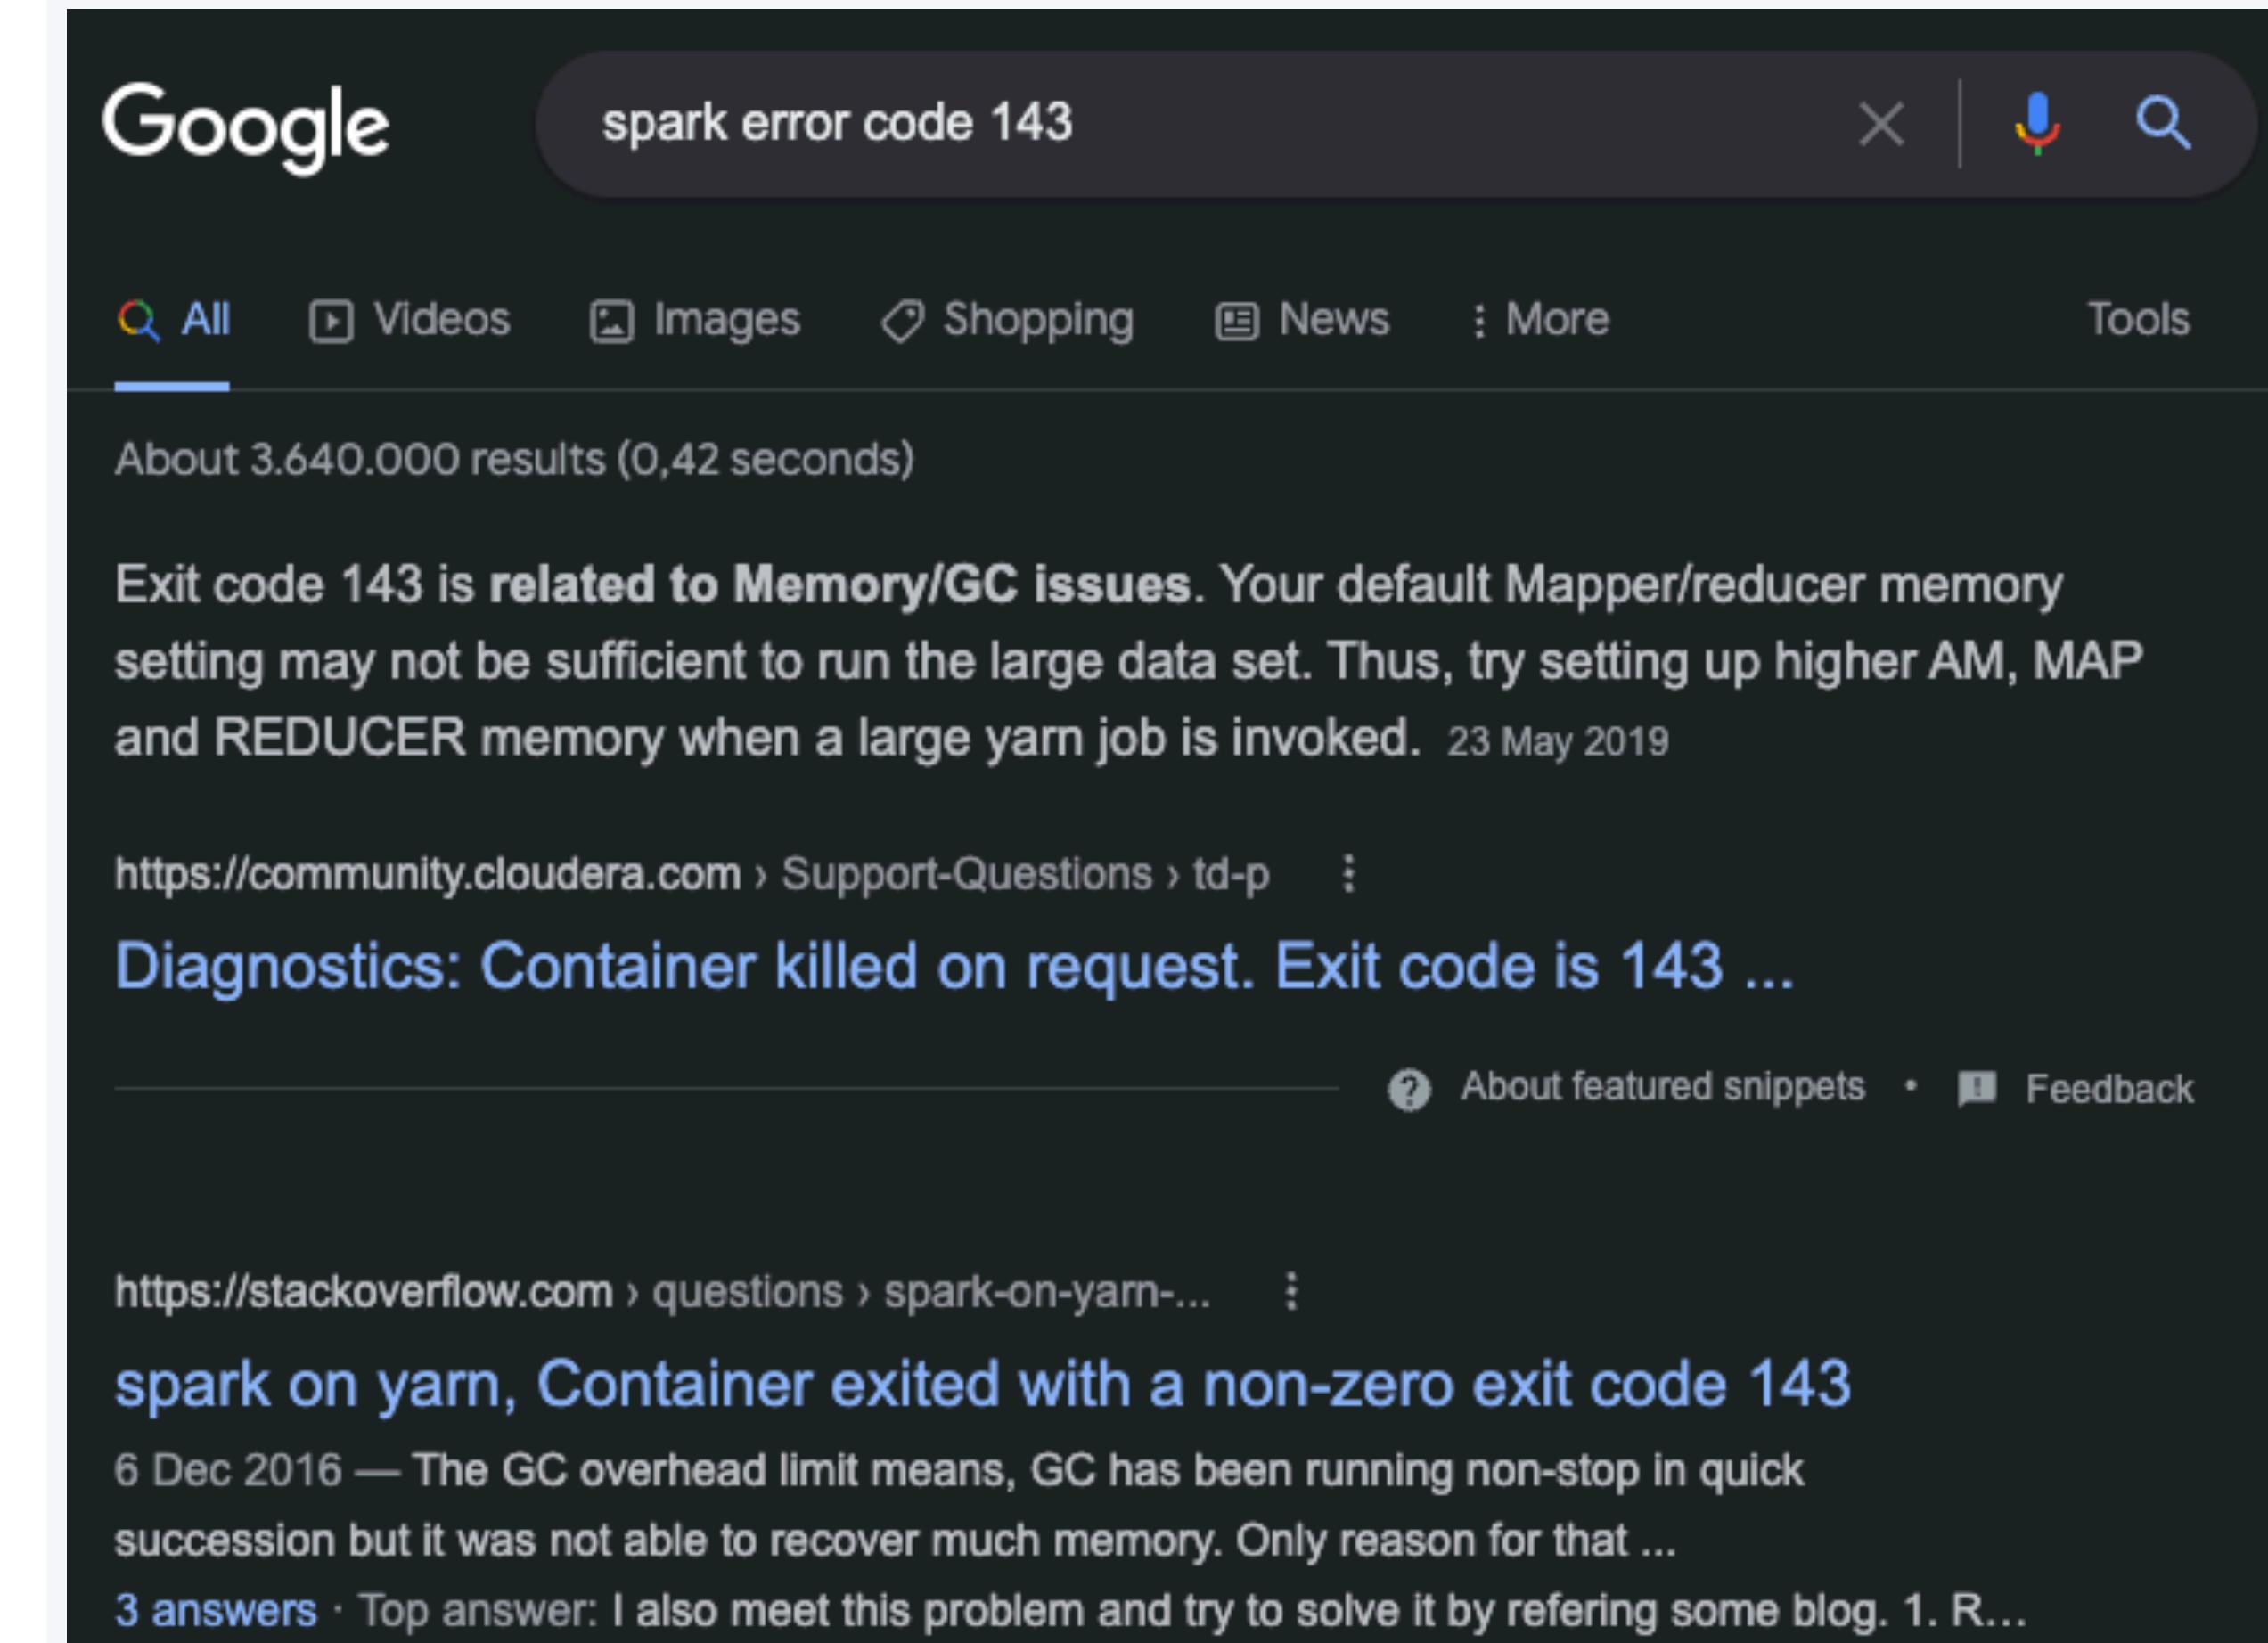
spark on yarn, Container exited with a non-zero exit code 143

6 Dec 2016 — The GC overhead limit means, GC has been running non-stop in quick succession but it was not able to recover much memory. Only reason for that ...

3 answers · Top answer: I also meet this problem and try to solve it by refering some blog. 1. R...

The Usual Workflow

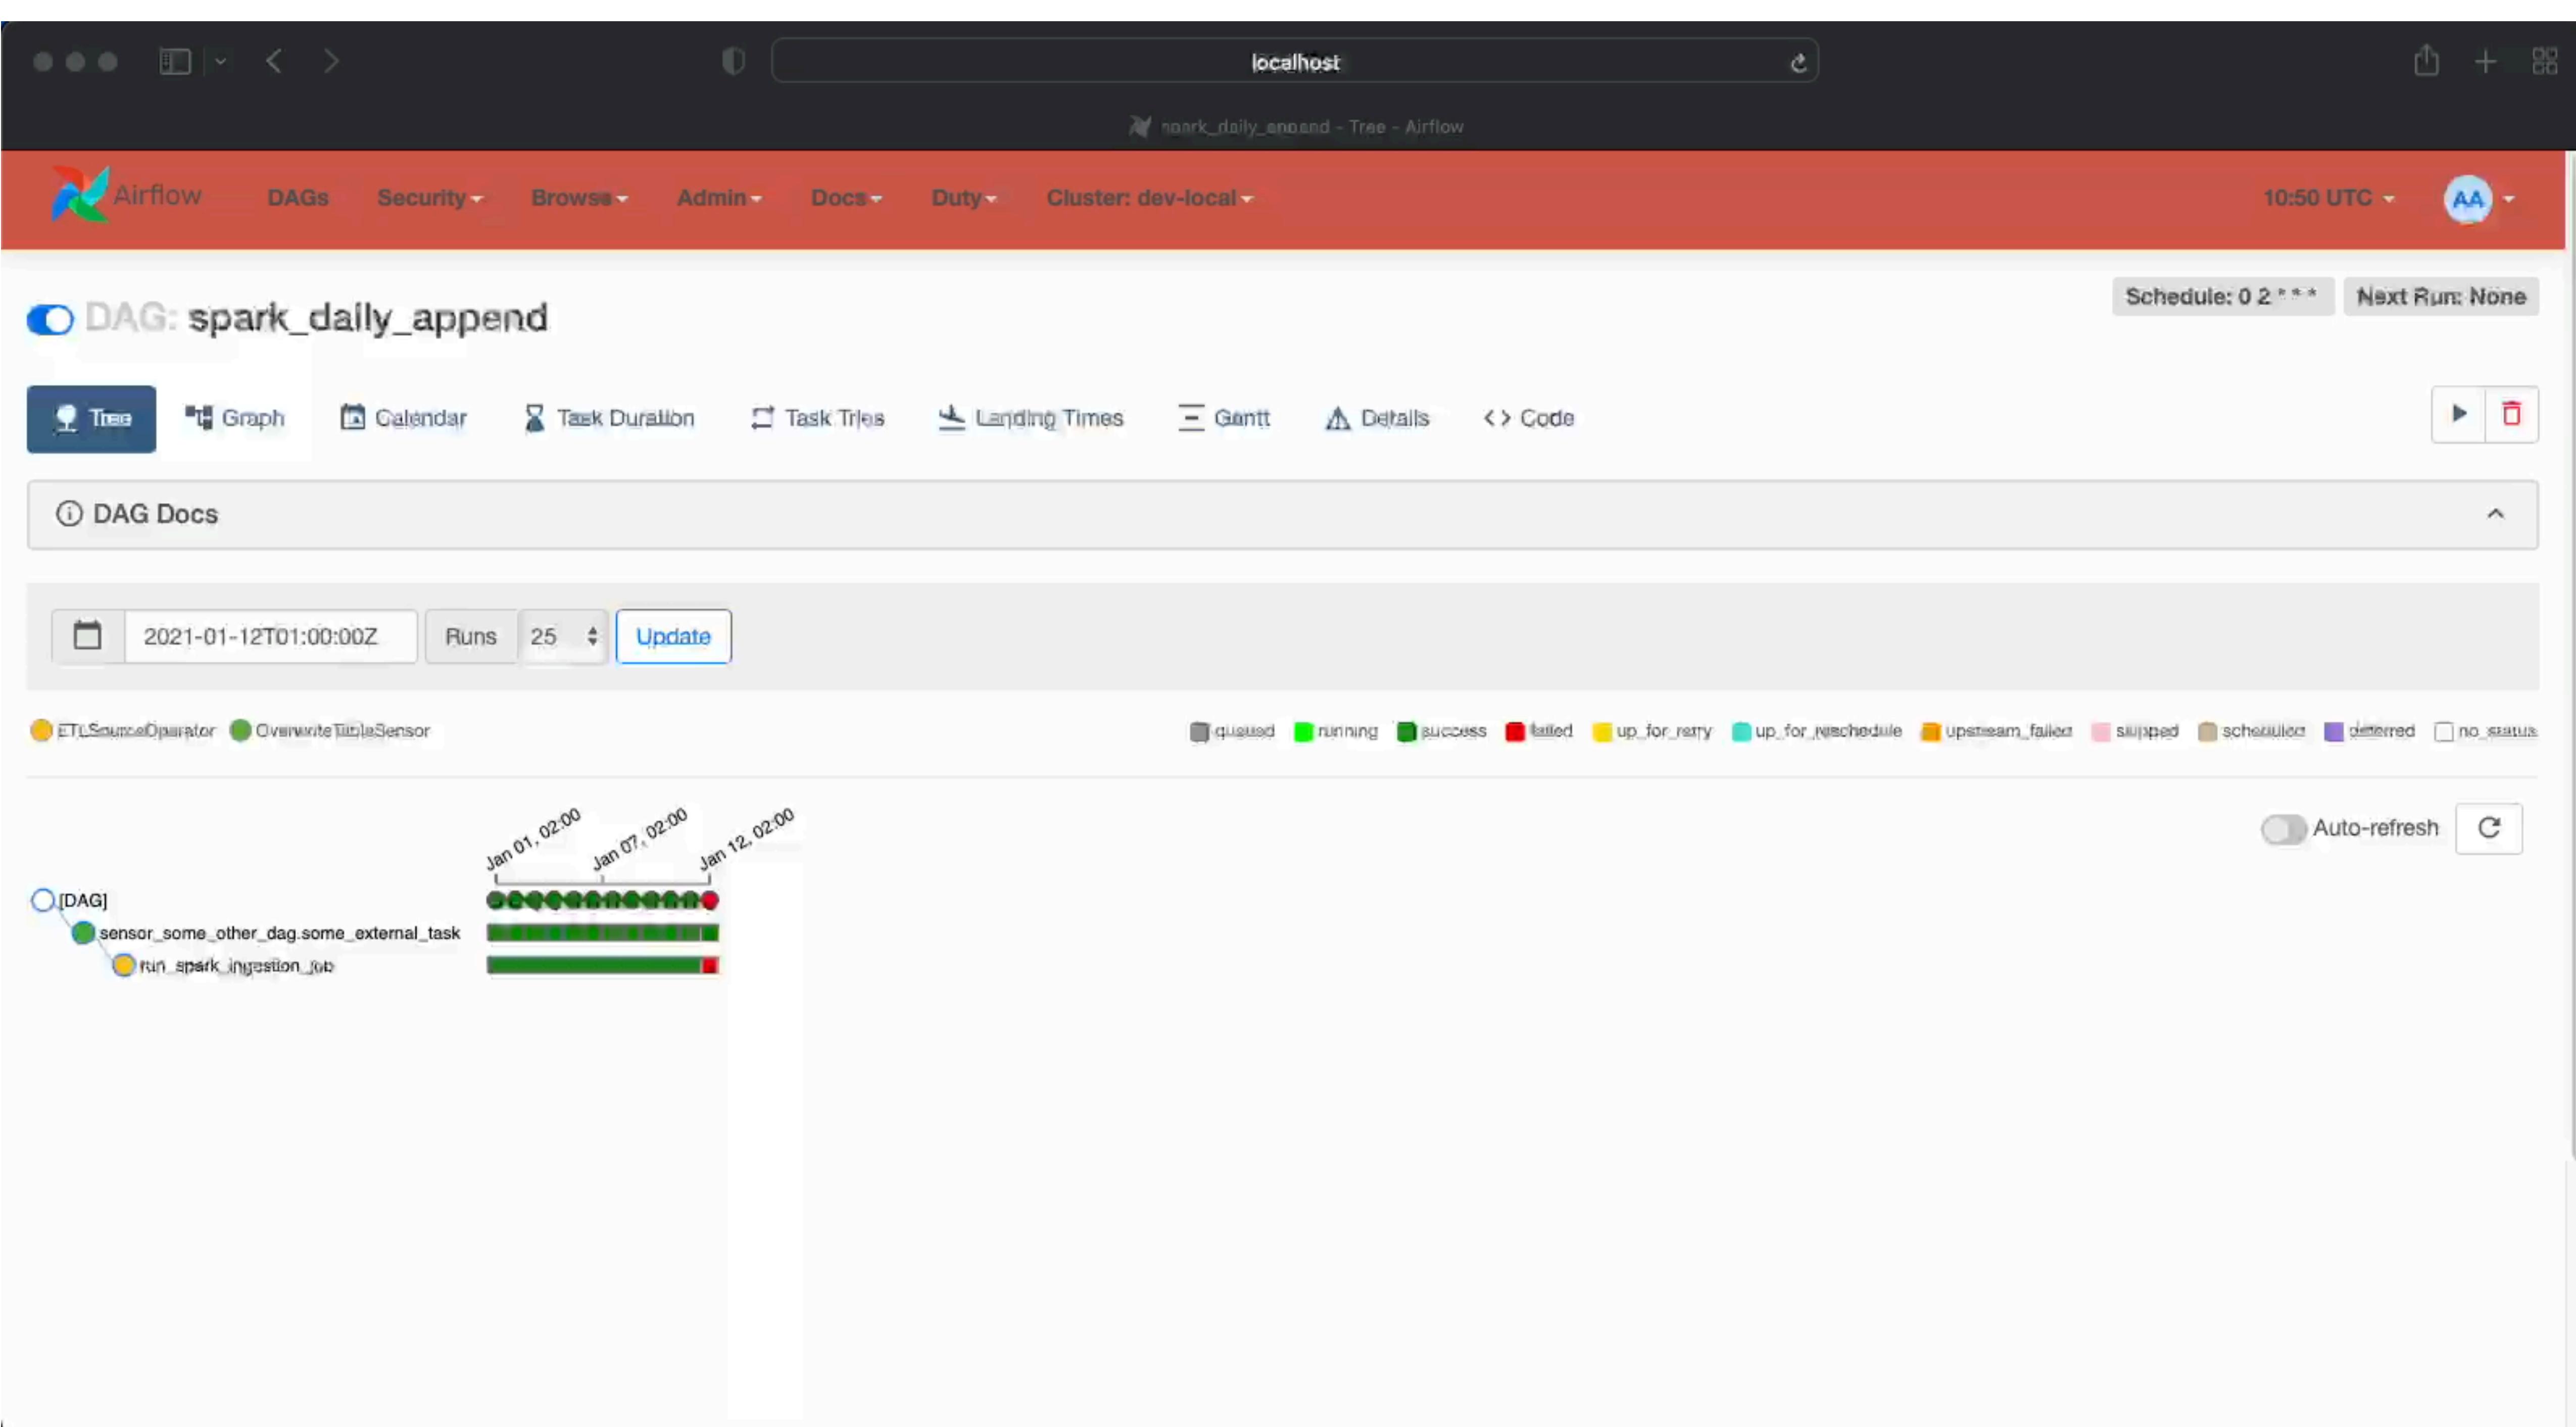
- You google the issue
- Solution: Set a higher Spark driver and/or worker memory limit
- But we hardcoded our resources in the code base, so that would require:
 - Create an MR
 - Request for an approval
 - Creating an official patch request
 - Request for an approval
 - Patching
- How can we prevent patching?



A screenshot of a Google search results page for the query "spark error code 143". The search bar at the top shows the query. Below it, a navigation bar includes "All", "Videos", "Images", "Shopping", "News", and "More" buttons. A "Tools" link is on the far right. The main content area displays search results. The first result is a snippet from Cloudera's community forum: "Exit code 143 is related to Memory/GC issues. Your default Mapper/reducer memory setting may not be sufficient to run the large data set. Thus, try setting up higher AM, MAP and REDUCER memory when a large yarn job is invoked." It includes a timestamp of "23 May 2019". The second result is a snippet from Stack Overflow: "Diagnostics: Container killed on request. Exit code is 143 ...". The third result is another snippet from Stack Overflow: "spark on yarn, Container exited with a non-zero exit code 143". It includes a timestamp of "6 Dec 2016" and a note about GC overhead. At the bottom of the snippet, it says "3 answers · Top answer: I also meet this problem and try to solve it by refering some blog. 1. R...".

not-so-live demo







adyen



Retention Period

- As we are on-premise, we have finite storage.
- If we are close to storage limits, there are two options:
 - Buy more servers
 - Get rid of some of our data
- We should prevent getting close to the limit



Governance library

- Each stream has a Stream file in this library.
- They define all their tables there and their corresponding retention period.
- Disclaimer: This is much more complex when you work with tables that are not partitioned by date

```
 1 class AnalyticsCore(Stream):
 2     def get_all_tables(self) -> List[HDFSTable]:
 3         return [
 4             HDFSTable(
 5                 HiveDatabase.ANALYTICS_CORE, "parquetfilesize", RetentionPeriod.THREE_MONTHS
 6             ),
 7             HDFSTable(
 8                 HiveDatabase.ANALYTICS_CORE, "parquetmetadataRowCount", RetentionPeriod.ONE_MONTH
 9             ),
10             HDFSTable(
11                 HiveDatabase.ANALYTICS_CORE, "sparkpushedfilters", RetentionPeriod.ONE_YEAR
12             ),
13             HDFSTable(
14                 HiveDatabase.ANALYTICS_CORE, "tablemetadata", RetentionPeriod.THREE_YEARS
15             ),
16         ]
17 
```

DAG: retention_based_removal

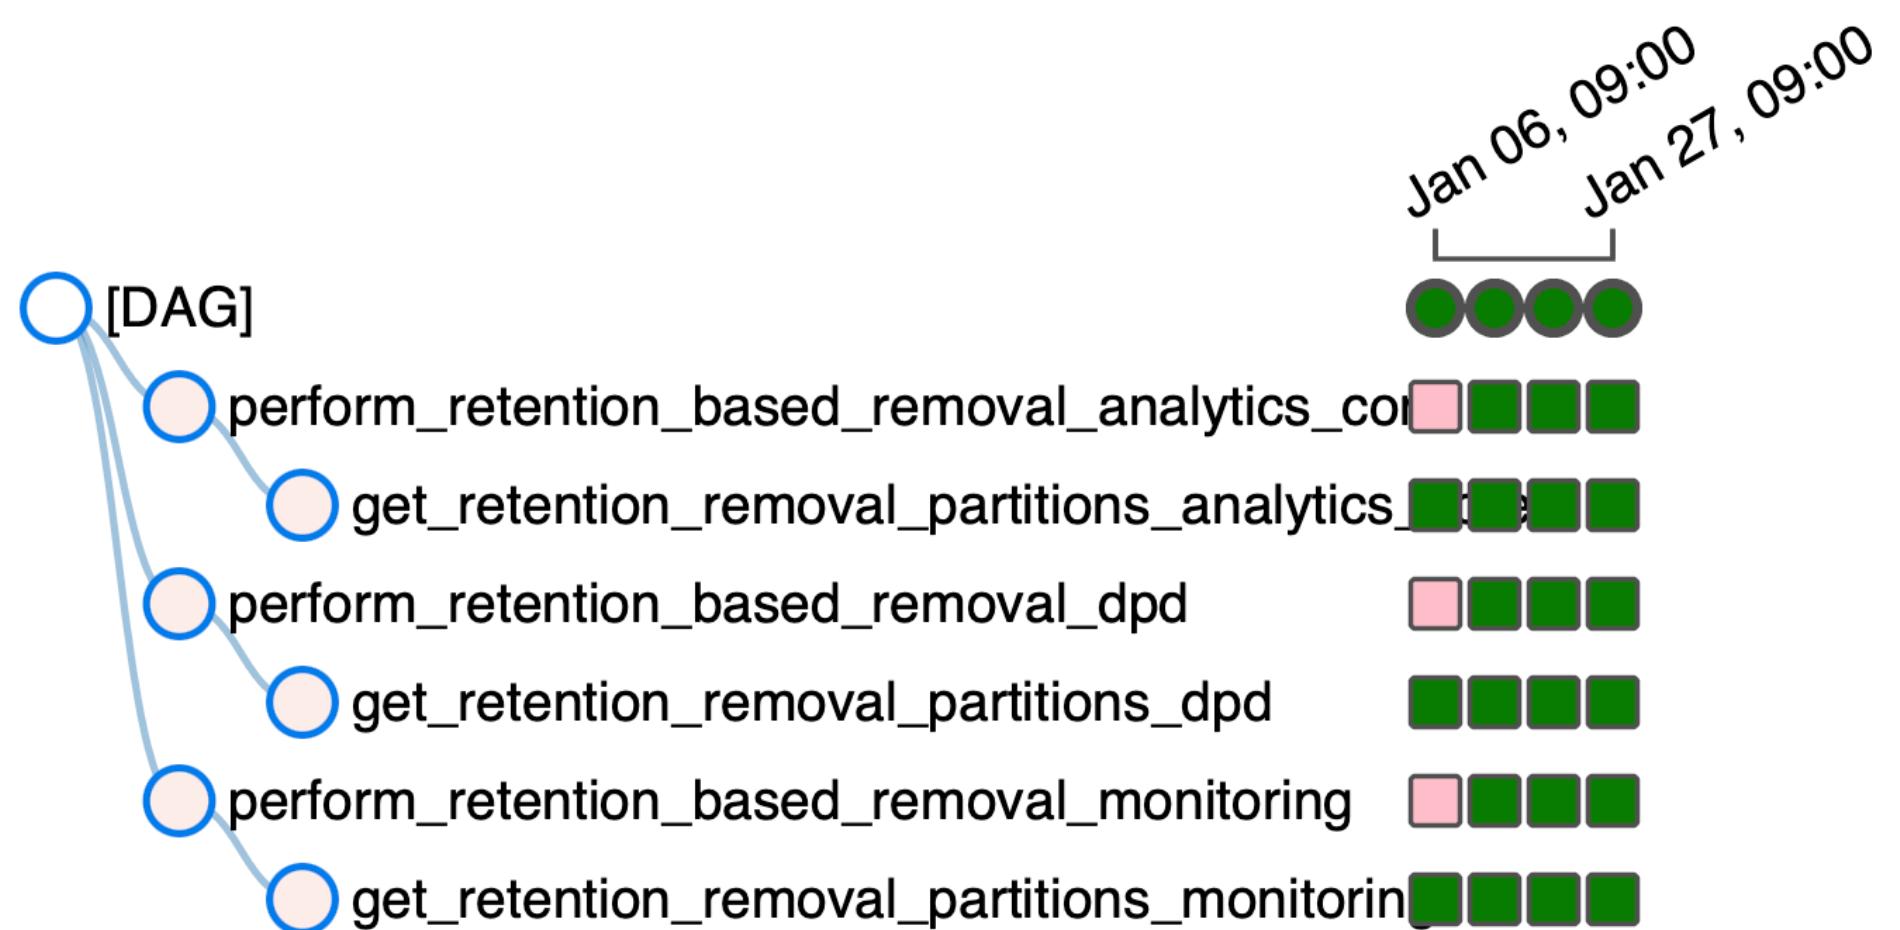
Tree Graph Calendar Task Duration Task Tries Landing Times Gantt Details Code

DAG Docs

2021-01-27T08:00:00Z Runs 25 Update

PythonOperator

queued running success failed up_for_retry up_for_



List Retention Removal Plan Model

Search▼



Actions▼



Record Count: 4

<input type="checkbox"/>	Stream	Hive Database	Table Name	Planned Removal Date	Status	Partitions Start Date	Partitions End Date	Forbid Undo Dag
<input type="checkbox"/>	analytics_core	analytics_core	sparkpushedfilters	2022-05-18	SUCCEEDED	2021-05-11	2021-05-17	False
<input type="checkbox"/>	analytics_core	analytics_core	sparkpushedfilters	2022-05-25	STAGED	2021-05-18	2021-05-24	False
<input type="checkbox"/>	analytics_core	analytics_core	tablemetadata	2022-05-18	SUCCEEDED	2022-05-05	2022-05-11	False
<input type="checkbox"/>	analytics_core	analytics_core	tablemetadata	2022-05-25	STAGED	2022-05-12	2022-05-18	False

List Retention Removal Plan Model

Search▼



Actions▼



Record Count: 4

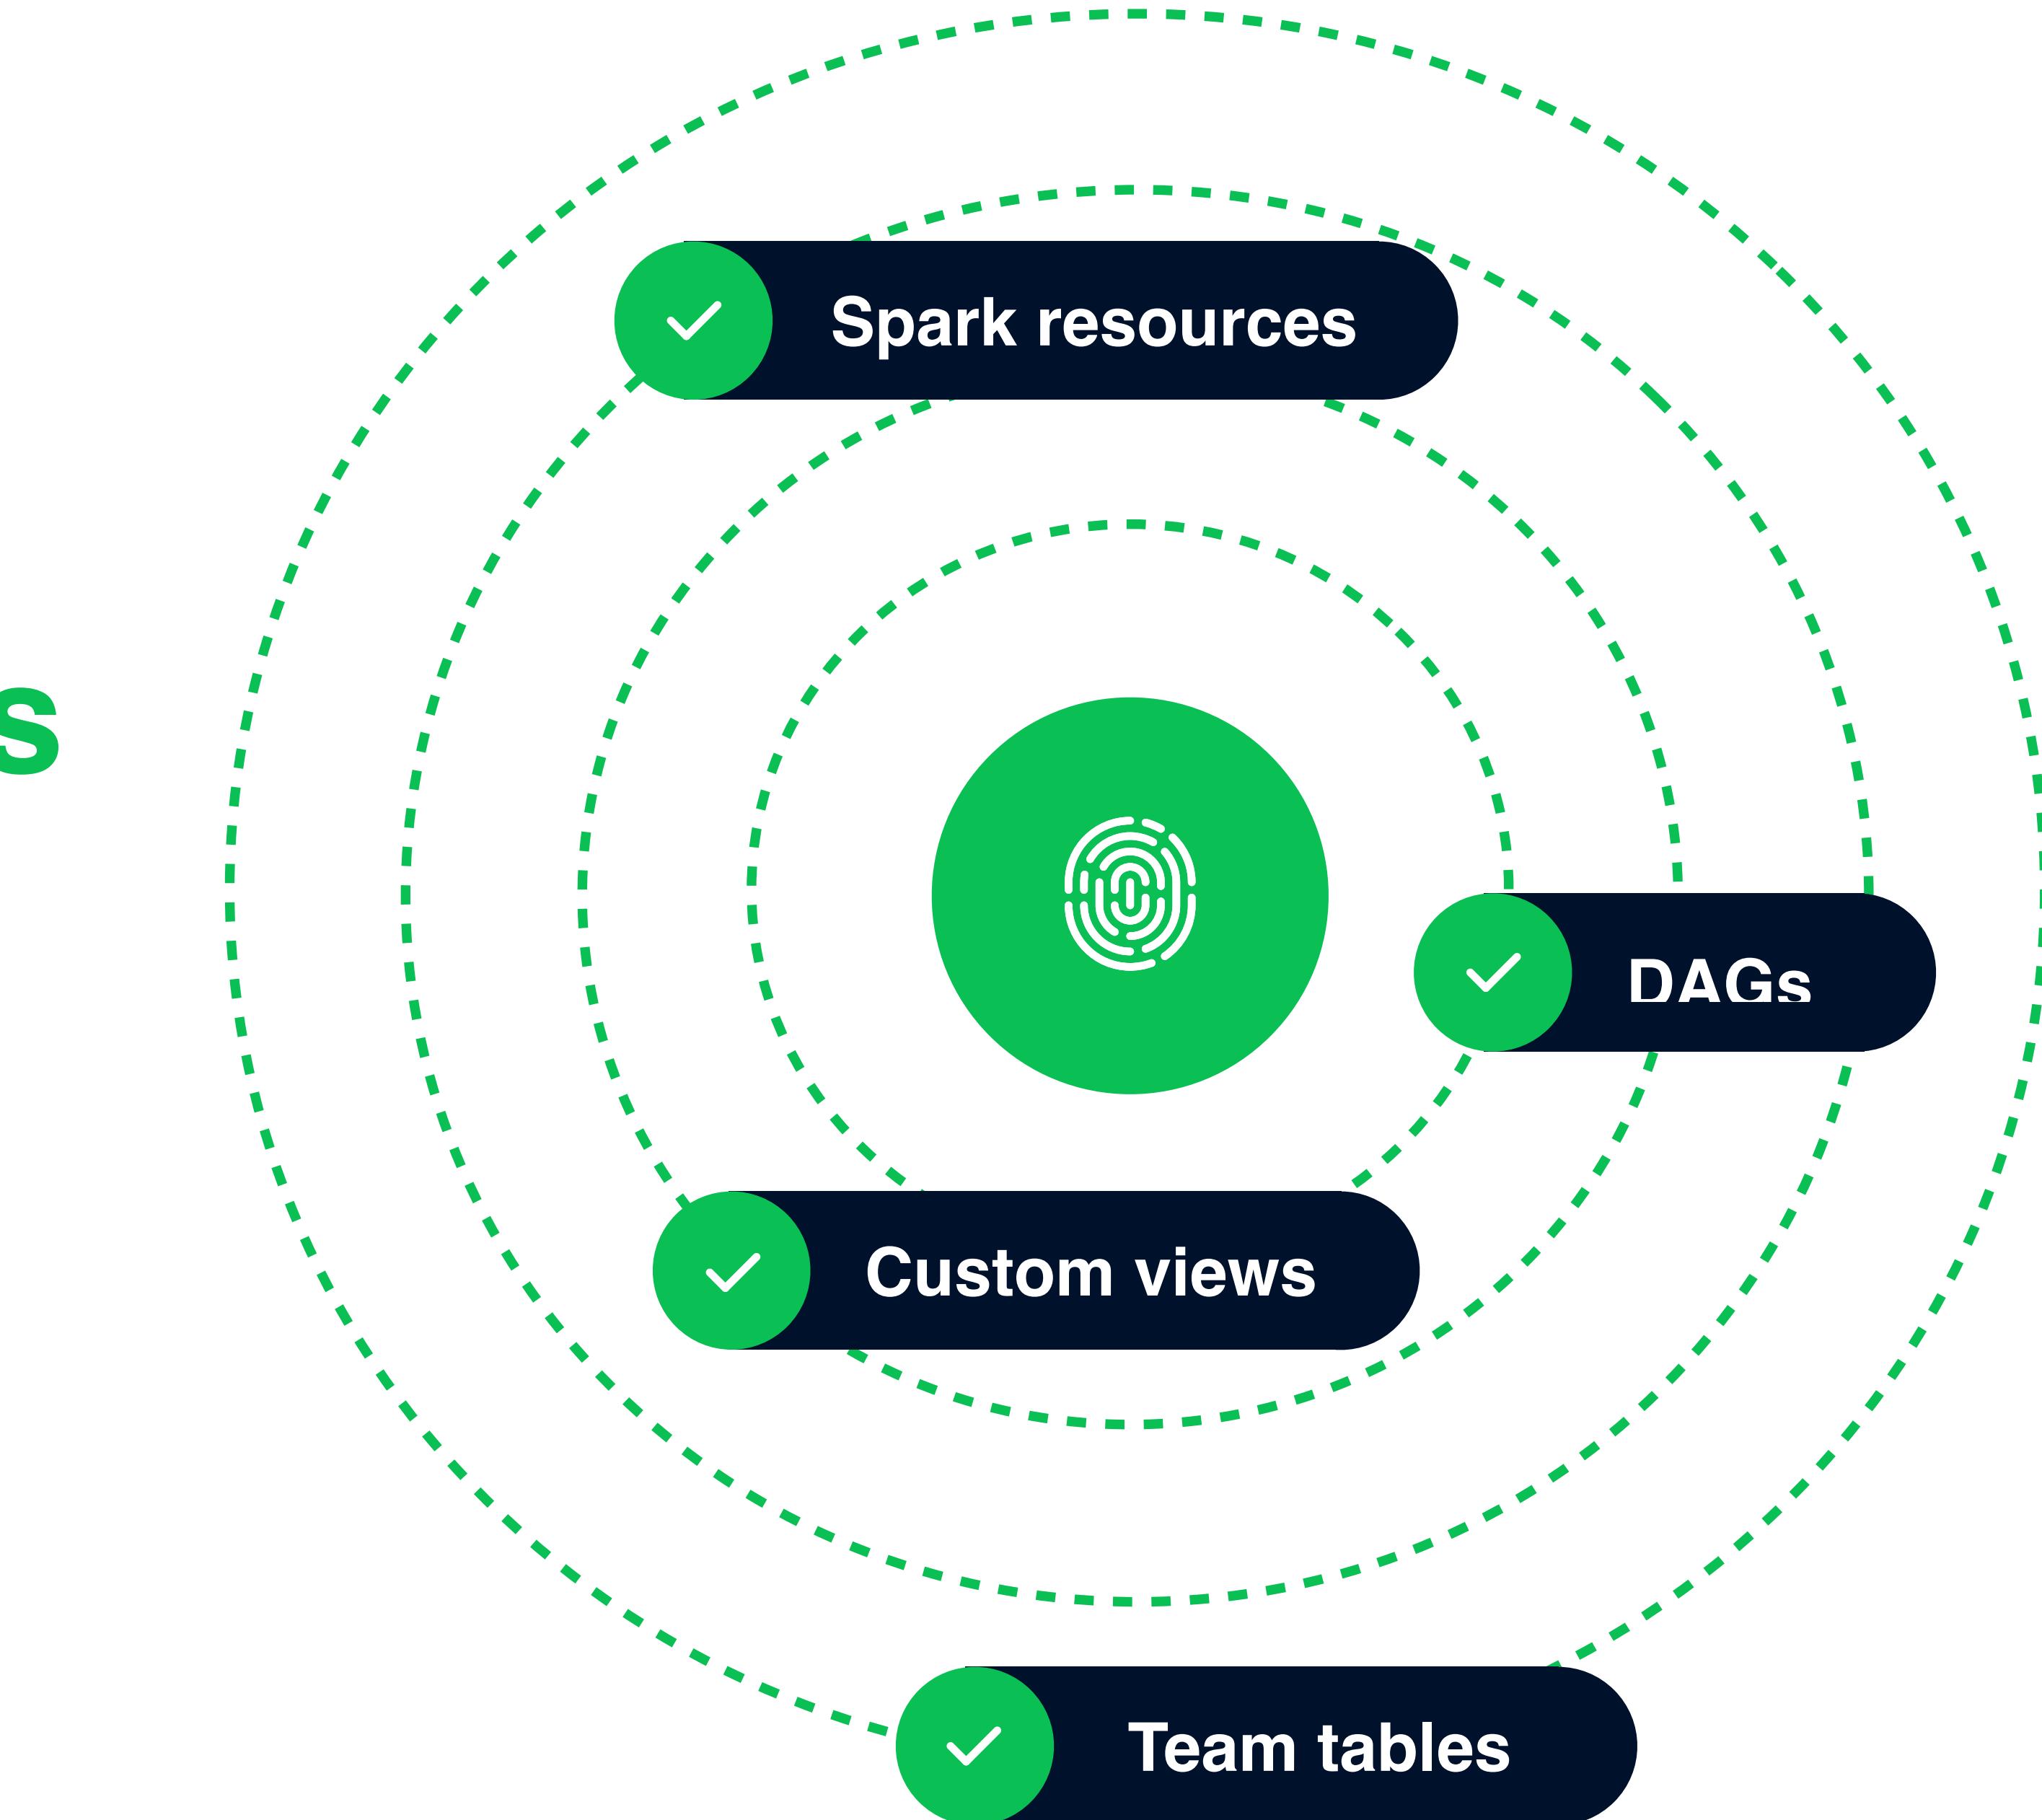
Drop the retention removal plan completely

<input type="checkbox"/>	Table Name	Planned Removal Date	Status	Partitions Start Date	Partitions End Date	Forbid Undo Dag
<input type="checkbox"/>	sparkpushedfilters	2022-05-18	SUCCEEDED	2021-05-11	2021-05-17	False
<input type="checkbox"/>	analytics_core	sparkpushedfilters	STAGED	2021-05-18	2021-05-24	False
<input type="checkbox"/>	analytics_core	tablemetadata	SUCCEEDED	2022-05-05	2022-05-11	False
<input type="checkbox"/>	analytics_core	tablemetadata	STAGED	2022-05-12	2022-05-18	False



User Permissions

- We want to empower the streams so it truly becomes self-service
- Yet, they should only be able to modify their own resources.



User Groups

- We always have one on-duty Admin
- For teams we have two access groups:
 - Stream admin
 - Standard user
- Team admins are able to:
 - Manage all the DAGs of the teams
 - Manage all the tables of their database
 - Manage the Spark resources of their tasks



Admin



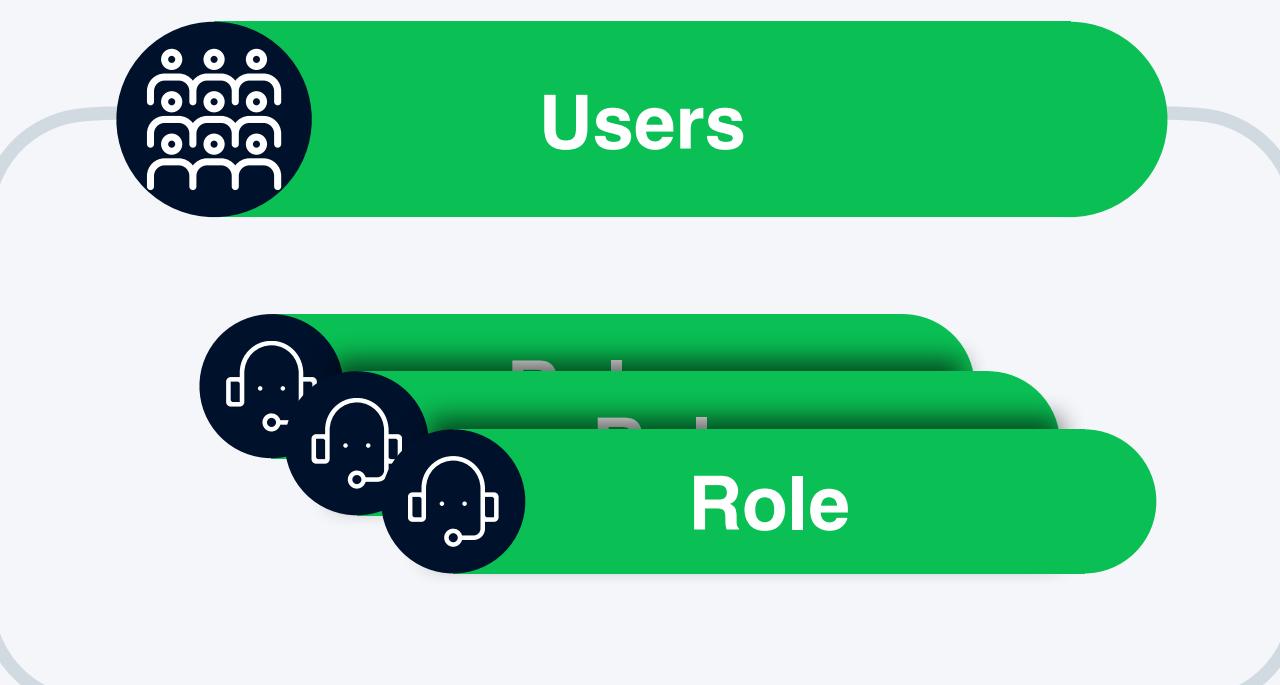
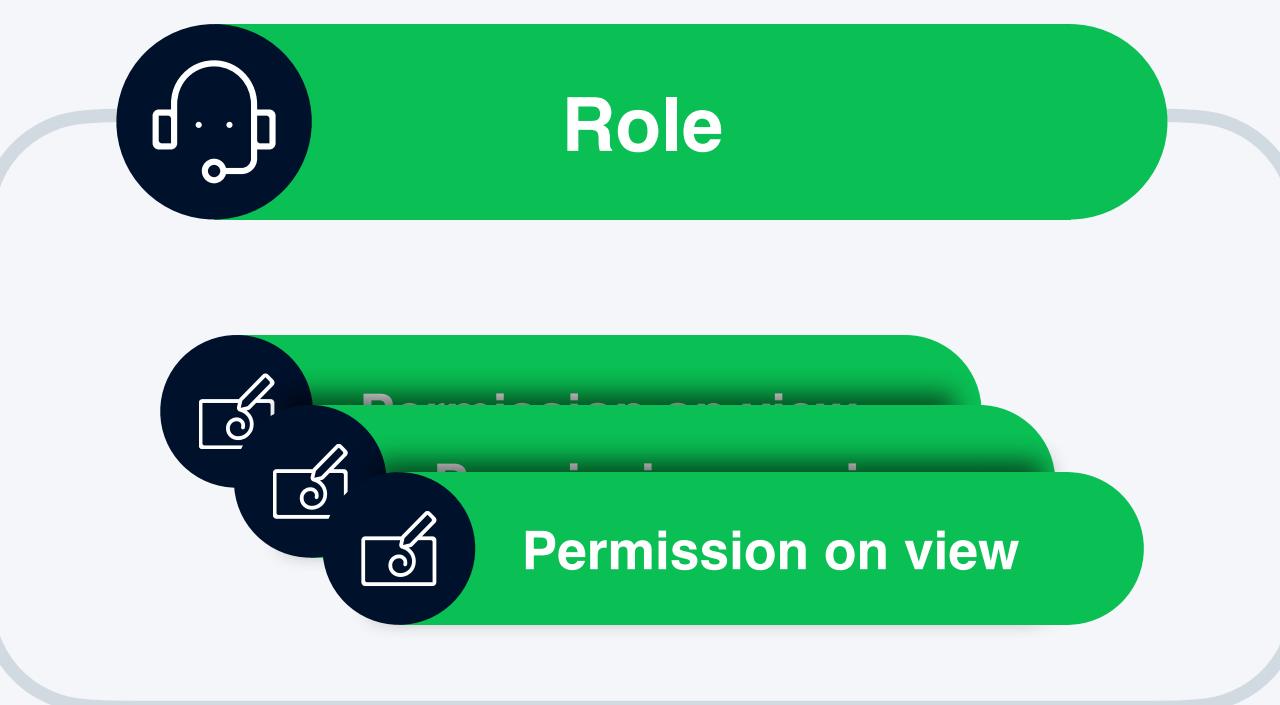
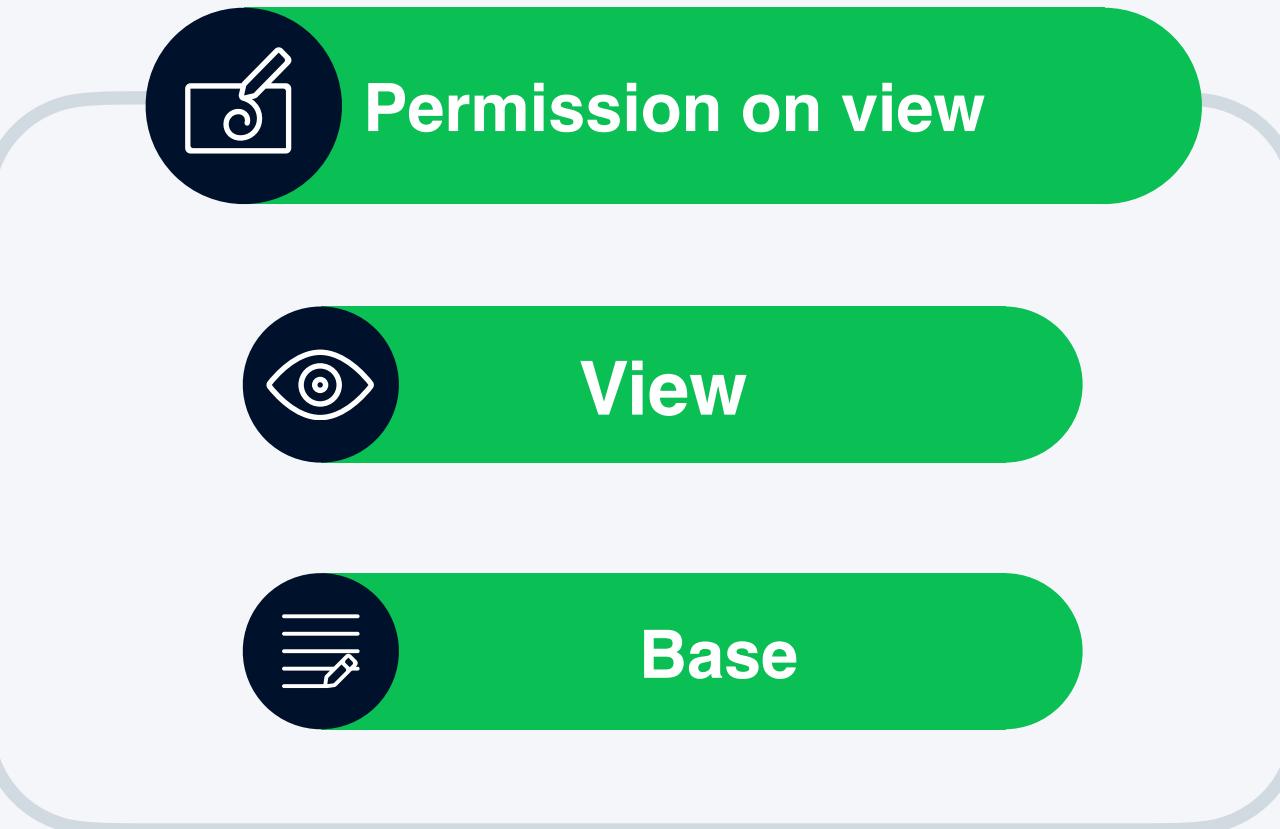
Stream admin



Users

Flask Appbuilder Permissions

- Why are you talking about Flask Appbuilder?
- In terms of permissions on these views:
 - Each DAG has two POVs:
 - `my_dag.can_read`
 - `my_dag.can_edit`
 - Each function in a view has a POV:
 - `spark_configuration.can_add`
 - `spark_configuration.can_edit`
- For roles we create:
 - One basic role for all users
 - One role for each stream admin group



We Implemented

- On each release, all permissions per role are updated.
- Created decorators to indicate who has access.
- Validators that raise exceptions when someone tries to modify a task that it has no access on.
- Modified our views to only show runnable parts.



Admin



Stream admin



Users

not-so-live demo





DAGs

Security

Browse

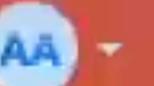
Admin

Docs

Duty

Cluster: dev-local

13:28 UTC



List Users

Search

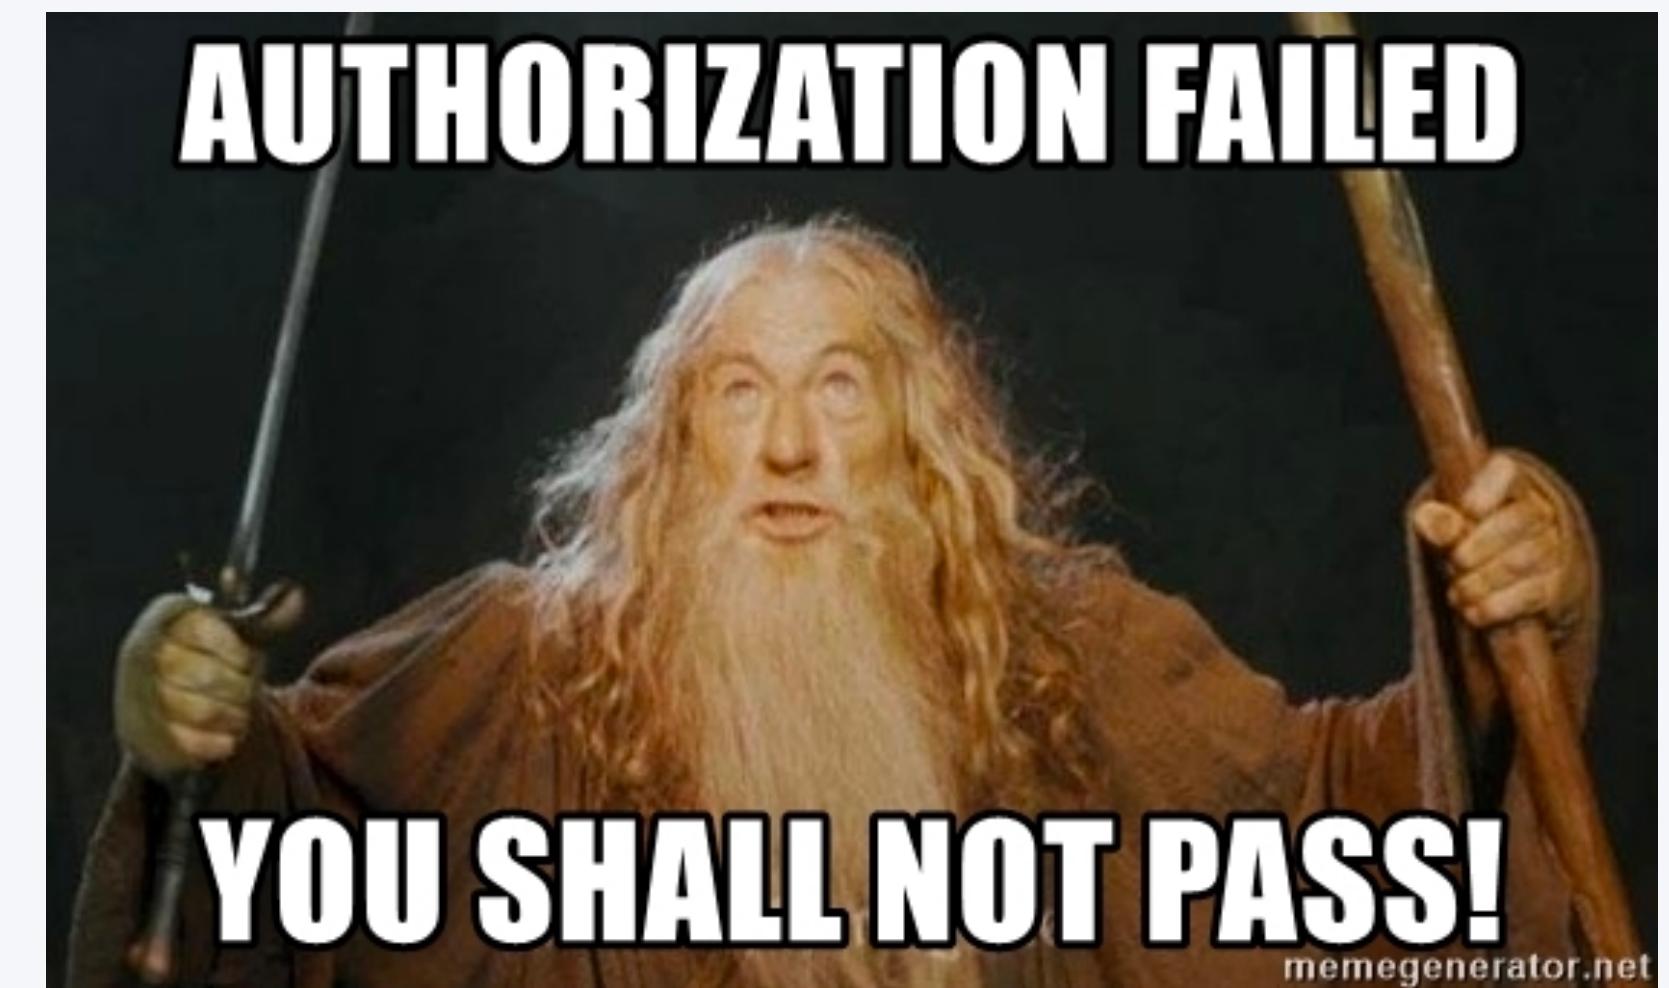


Record Count: 2

First Name	Last Name	User Name	Email	Is Active?	Role
ac	ac	ac	ac@adyen.com	True	[User, stream_analytics_core]
adyen	adyen	adyen	adyen@adyen	True	[Admin]

Mapping Users to Roles

- Now that we have the roles defined, this is the final thing left to do.
- Security team requires us to use LDAP for managing the roles.
- Service that runs each hour.



Wrap up

Table Schemas

How one can enable schema evolution over time and abstract away complexities

Tuning Resources

Preventing patches for data overflowing its ETL resource limits

Retention Period

Work with limited storage

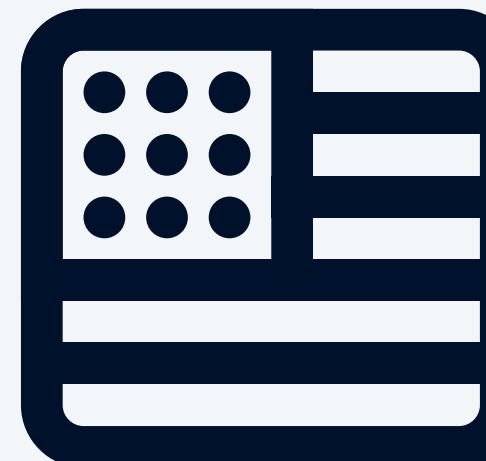
User Permissions

Each domain has their own database.

Before I forget...



**NL - Amsterdam
Head office**



**US - Chicago
Tech hub**



**Spain - Madrid
Tech hub**



Thank you



Slides for questions

Code Example: How we Define Permissions



```
1 class ConfigurationView(AdyenProtectedModelView):
2     datamodel = CustomSQLAInterface(ConfigModel)
3
4     extra_allowed_user_functions = ["list"]
5     extra_allowed_stream_admin_functions = ["add", "delete", "edit"]
6
7     @action("export", "Export the Spark Configurations", "", single=False)
8     @visible_to_users(is_fab_action=True, action_name="export")
9     def export_configs(self, configs: List[ConfigModel]) -> Response:
10         pass
11
12     @action("drop", "Drop Configurations completely", "You sure?", single=False)
13     @visible_to_stream_admins(is_fab_action=True, action_name="drop")
14     def drop_completely(self, configs: List[ConfigModel]) -> Response:
15         pass
```

Code example: Implementation of Decorators



```
1 def visible_to_stream_admins(is_fab_action: bool, action_name: Optional[str] = None):
2     def inner_function(func: Callable):
3         @wraps(func)
4         def wrapper(*args: Any, **kwargs: Any) -> Any:
5             return func(*args, **kwargs)
6
7         if is_fab_action and (action_name is None or not any(action_name)):
8             raise ValueError(f"If {is_fab_action=} you must also set {action_name=}.")
9
10        wrapper.visible_to_stream_admins = True
11        wrapper.is_fab_action = is_fab_action
12        wrapper.fab_action_name = action_name
13        return wrapper
14
15    return inner_function
```

Operations on Tables

For any operation you wish to perform on a table, you can categorise them in one (or multiple) of these 5 operations:

	NEW	COMPATIBLE	INCOMPATIBLE	REMOVE	MOVE
Defines hive table definition	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>		<input checked="" type="checkbox"/>
Delete hive table definition		<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>
Delete table data	<input checked="" type="checkbox"/>		<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	
MSCK Repair		<input checked="" type="checkbox"/>			<input checked="" type="checkbox"/>