



(What's New With)

Amazon Managed Workflows for Apache Airflow

John Jackson (he/him)

Principal Product Manager

Agenda

About Amazon MWAA

Architecture Review

Private vs Public Web Servers

Local Runner

Best Practices

Open Source

What's Next

About Amazon MWAA

What is Amazon MWAA?

Open Source

Same Apache Airflow, no forks

Secure

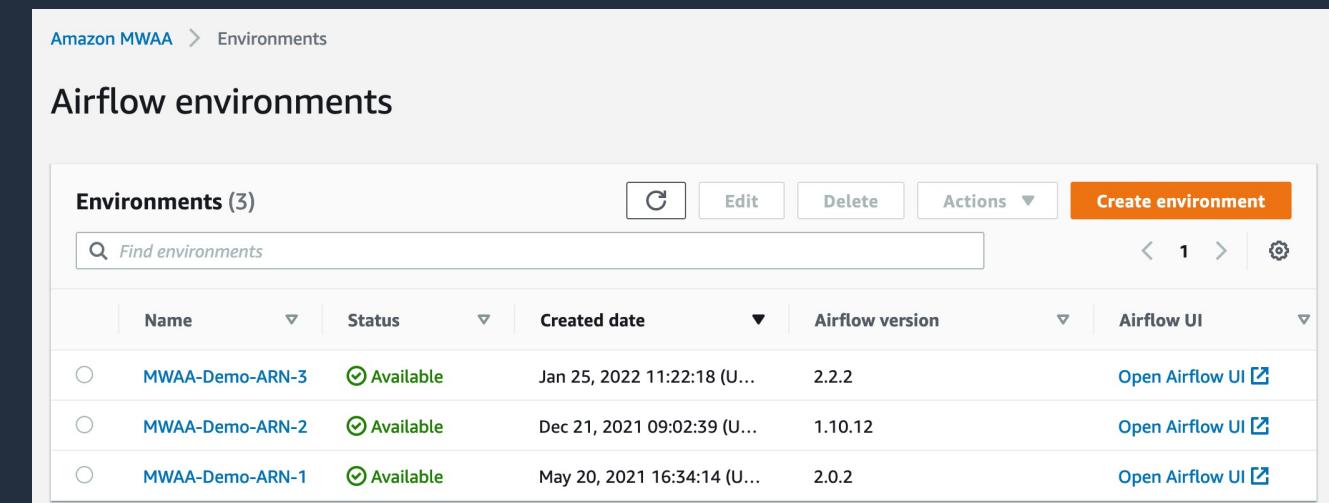
AWS Identity and Access Management (IAM)

Integrated

AWS Fargate, Amazon CloudWatch, S3

Easy to Deploy

Console, CloudFormation, API, CLI



The screenshot shows the AWS Lambda console interface. At the top, there's a navigation bar with 'Amazon MWAA' and 'Environments'. Below it is a section titled 'Airflow environments'. A table lists three environments:

Name	Status	Created date	Airflow version	Airflow UI
MWAA-Demo-ARN-3	Available	Jan 25, 2022 11:22:18 (U...)	2.2.2	Open Airflow UI
MWAA-Demo-ARN-2	Available	Dec 21, 2021 09:02:39 (U...)	1.10.12	Open Airflow UI
MWAA-Demo-ARN-1	Available	May 20, 2021 16:34:14 (U...)	2.0.2	Open Airflow UI

Key Milestones

November 2020

Amazon Managed Workflows for Apache Airflow is generally available

May 2021

Amazon MWAA adds Airflow 2.0 support

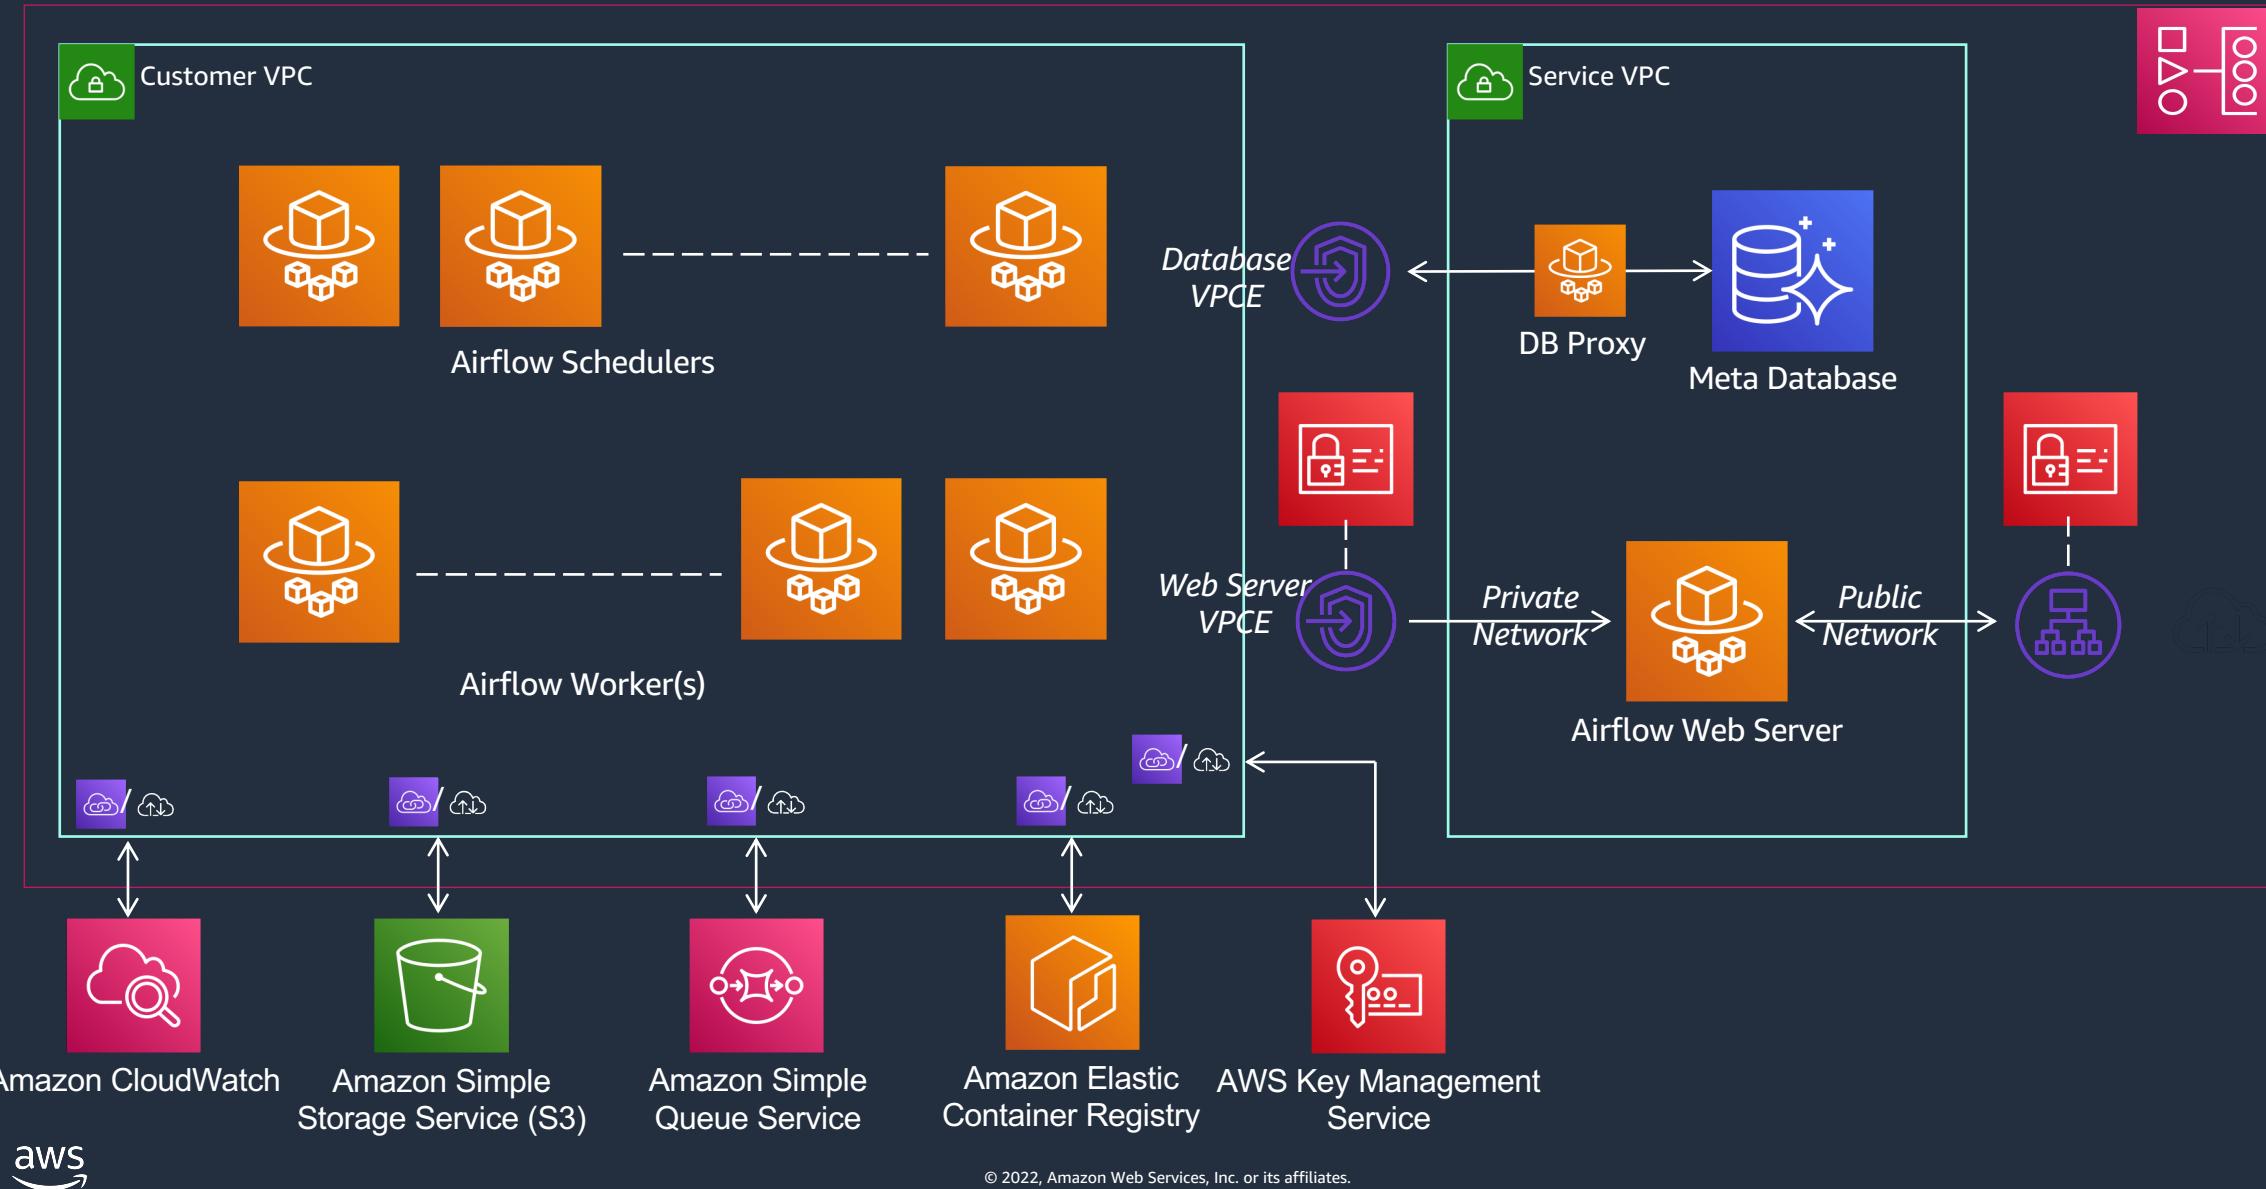
August 2021

Amazon MWAA expands to a total of 15 commercial regions

January 2022

Amazon MWAA adds Airflow 2.2 support with ability to install requirements and plugins on the web server

Amazon MWAA Architecture



Web Server Options

Web Server Options

Private Network

For customers that require their Airflow web server to be completely disconnected from the Internet

Public Network

Uses the same IAM authentication as AWS console pages to allow access to the Airflow web server

Configure advanced settings

Networking [Info](#)

Virtual private cloud (VPC)
Defines the networking infrastructure setup of your Airflow environment. An environment needs 2 private subnets in different availability zones. To create a new VPC with private subnets, choose Create MWAA VPC. [Learn more](#) 

[Choose VPC](#)  [Create MWAA VPC](#) 

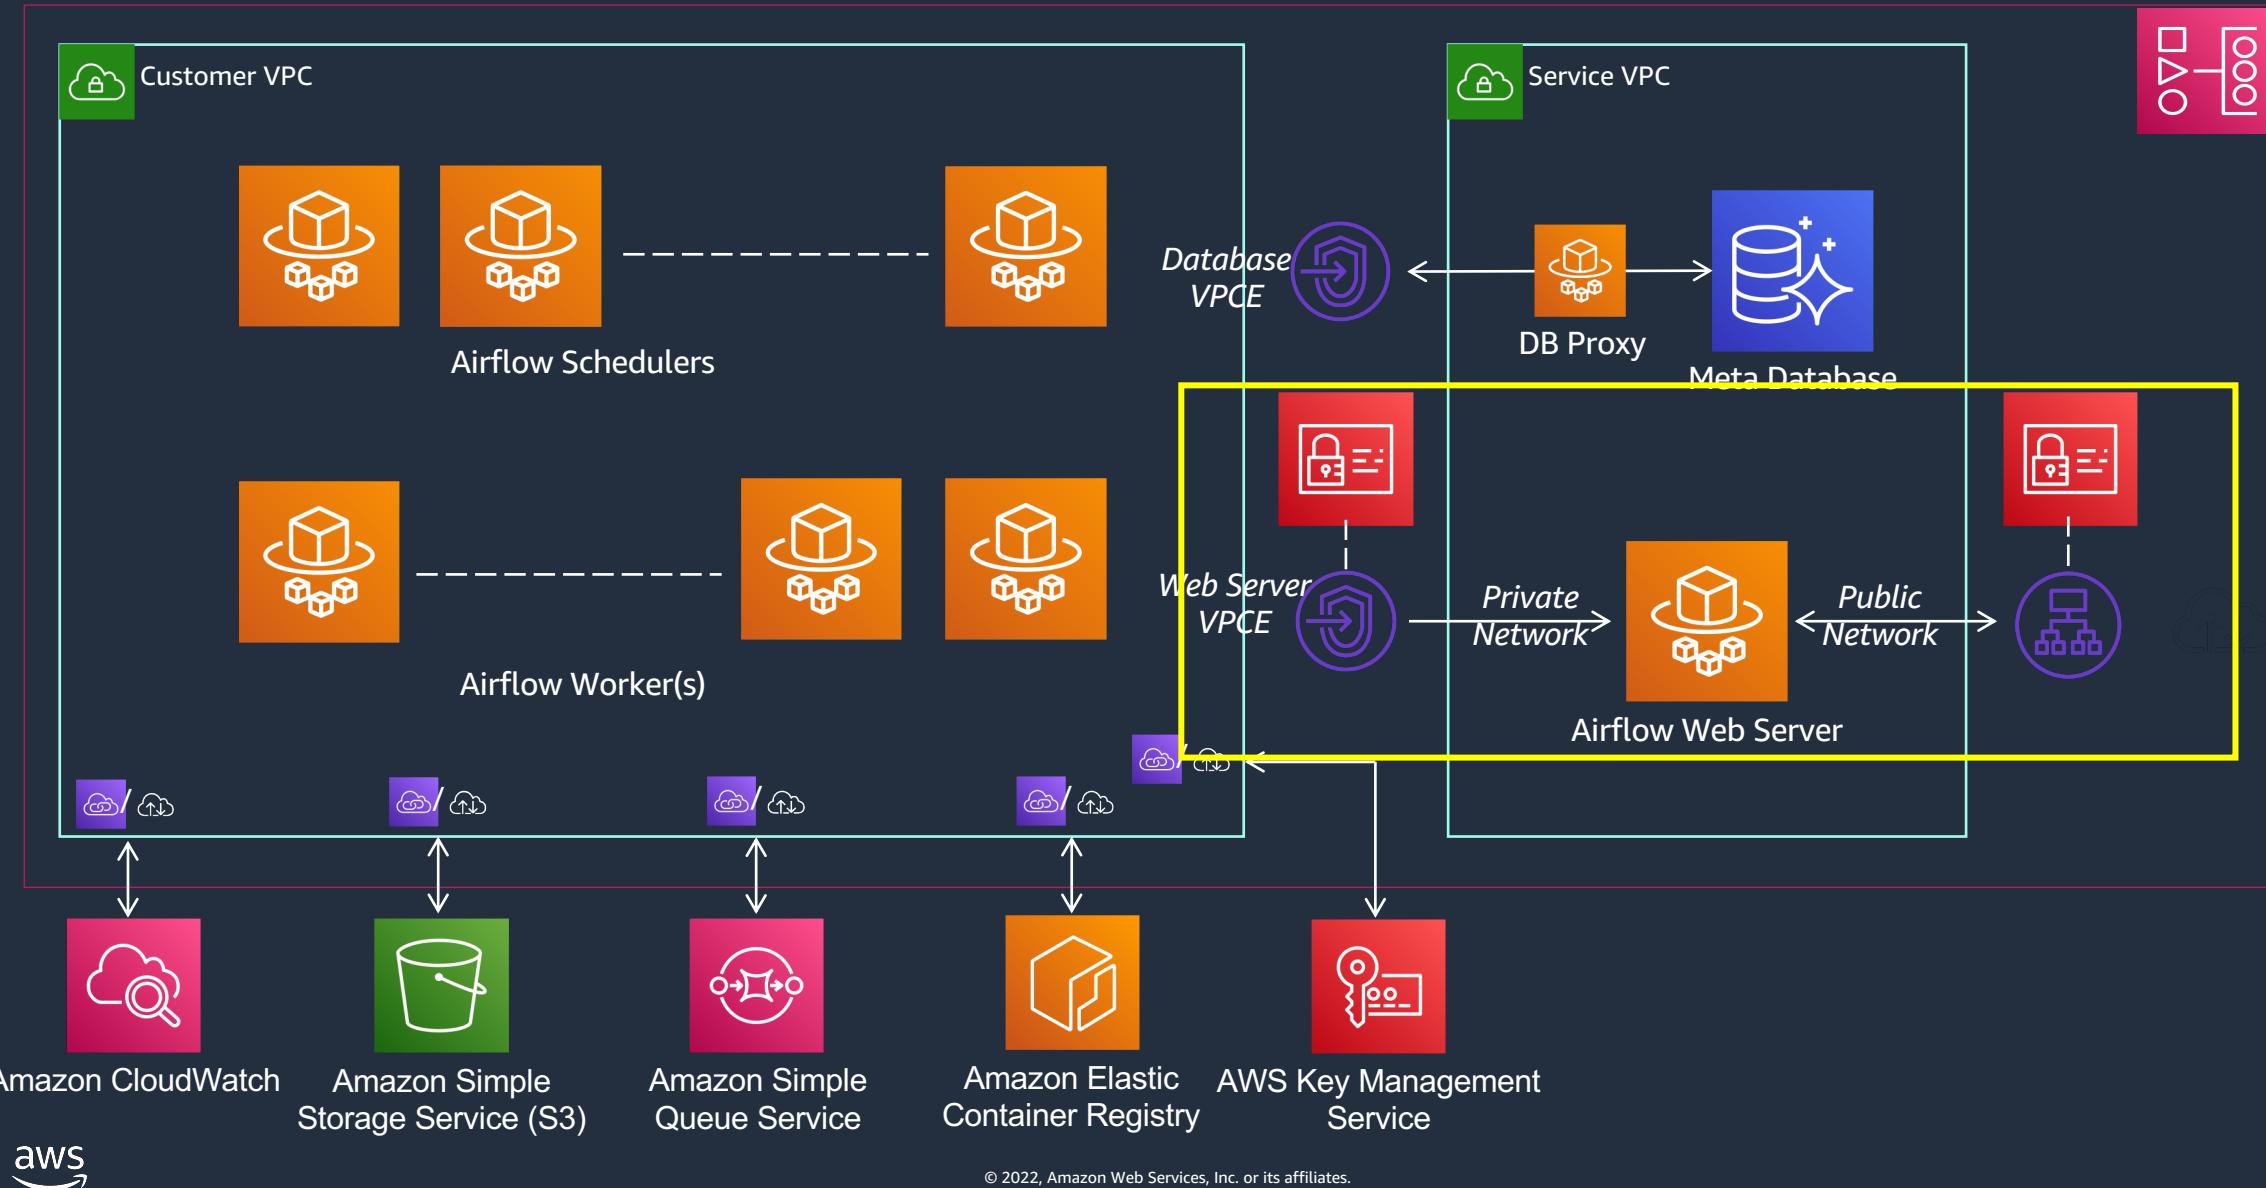
 VPC and subnet selections can't be changed after an environment is created.

Web server access

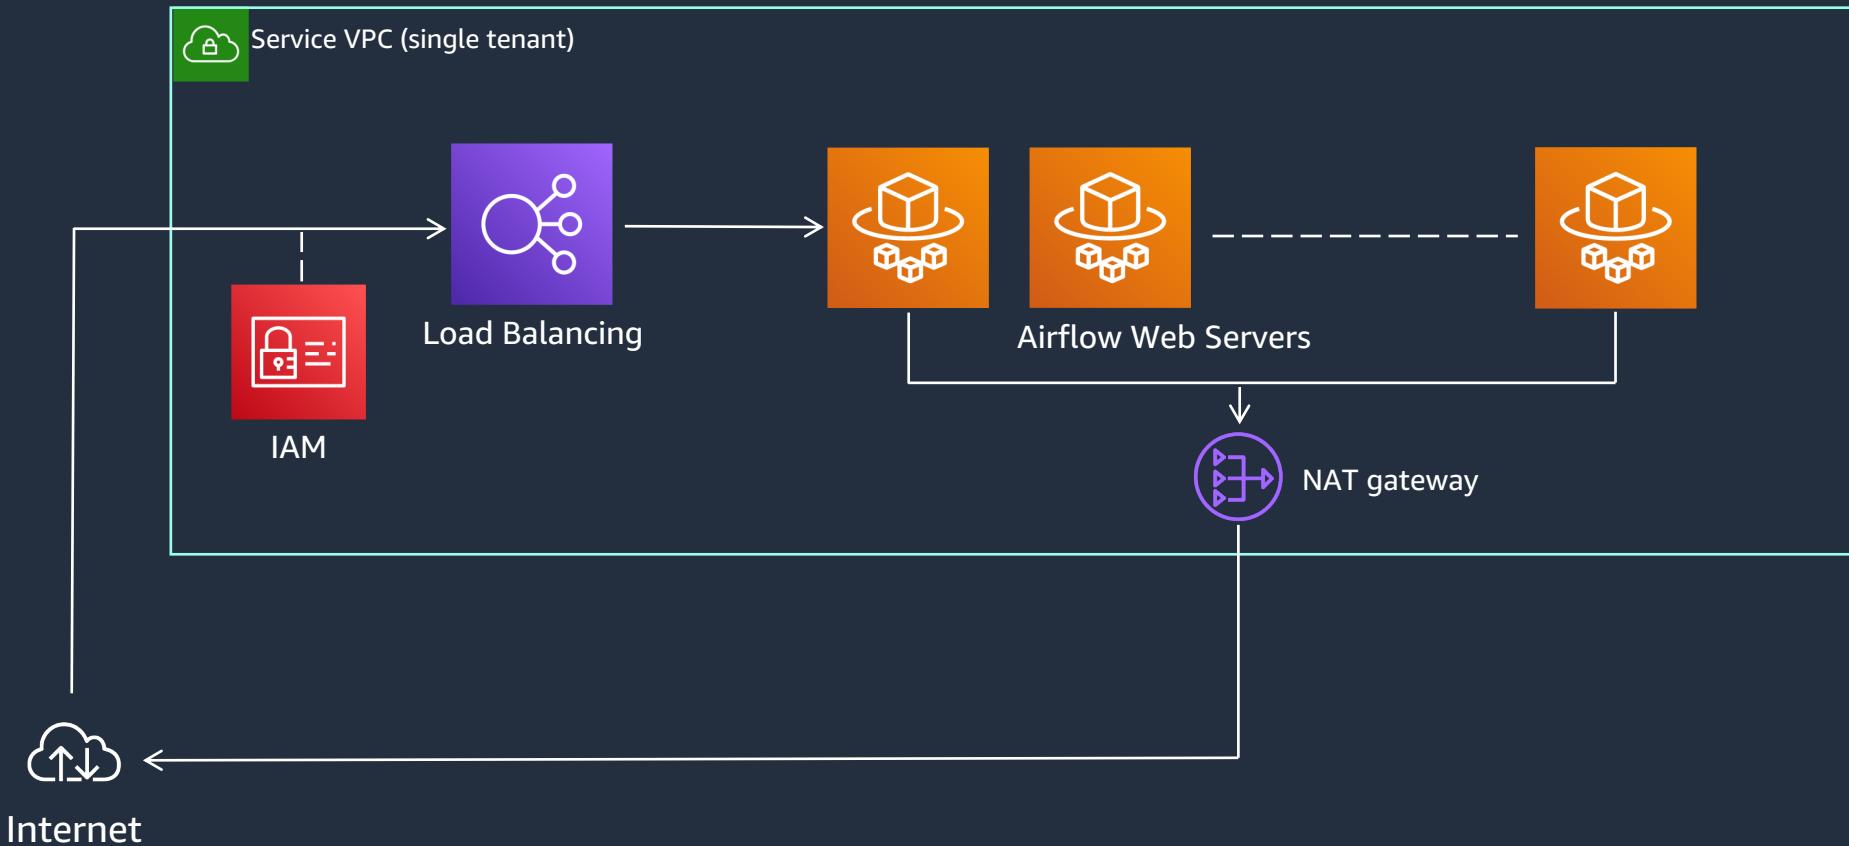
Private network (No internet access)
Additional setup required. Your Airflow UI can only be accessed by secure login behind your VPC. Choose this option if your Airflow UI is only accessed within a corporate network and you do not require a public repository for webserver requirements installation. IAM must be used to handle user authentication.

Public network (Internet accessible)
Your Airflow UI can be accessed by secure login over the Internet. Choose this option if your Airflow UI is accessed outside of a corporate network. IAM must be used to handle user authentication.

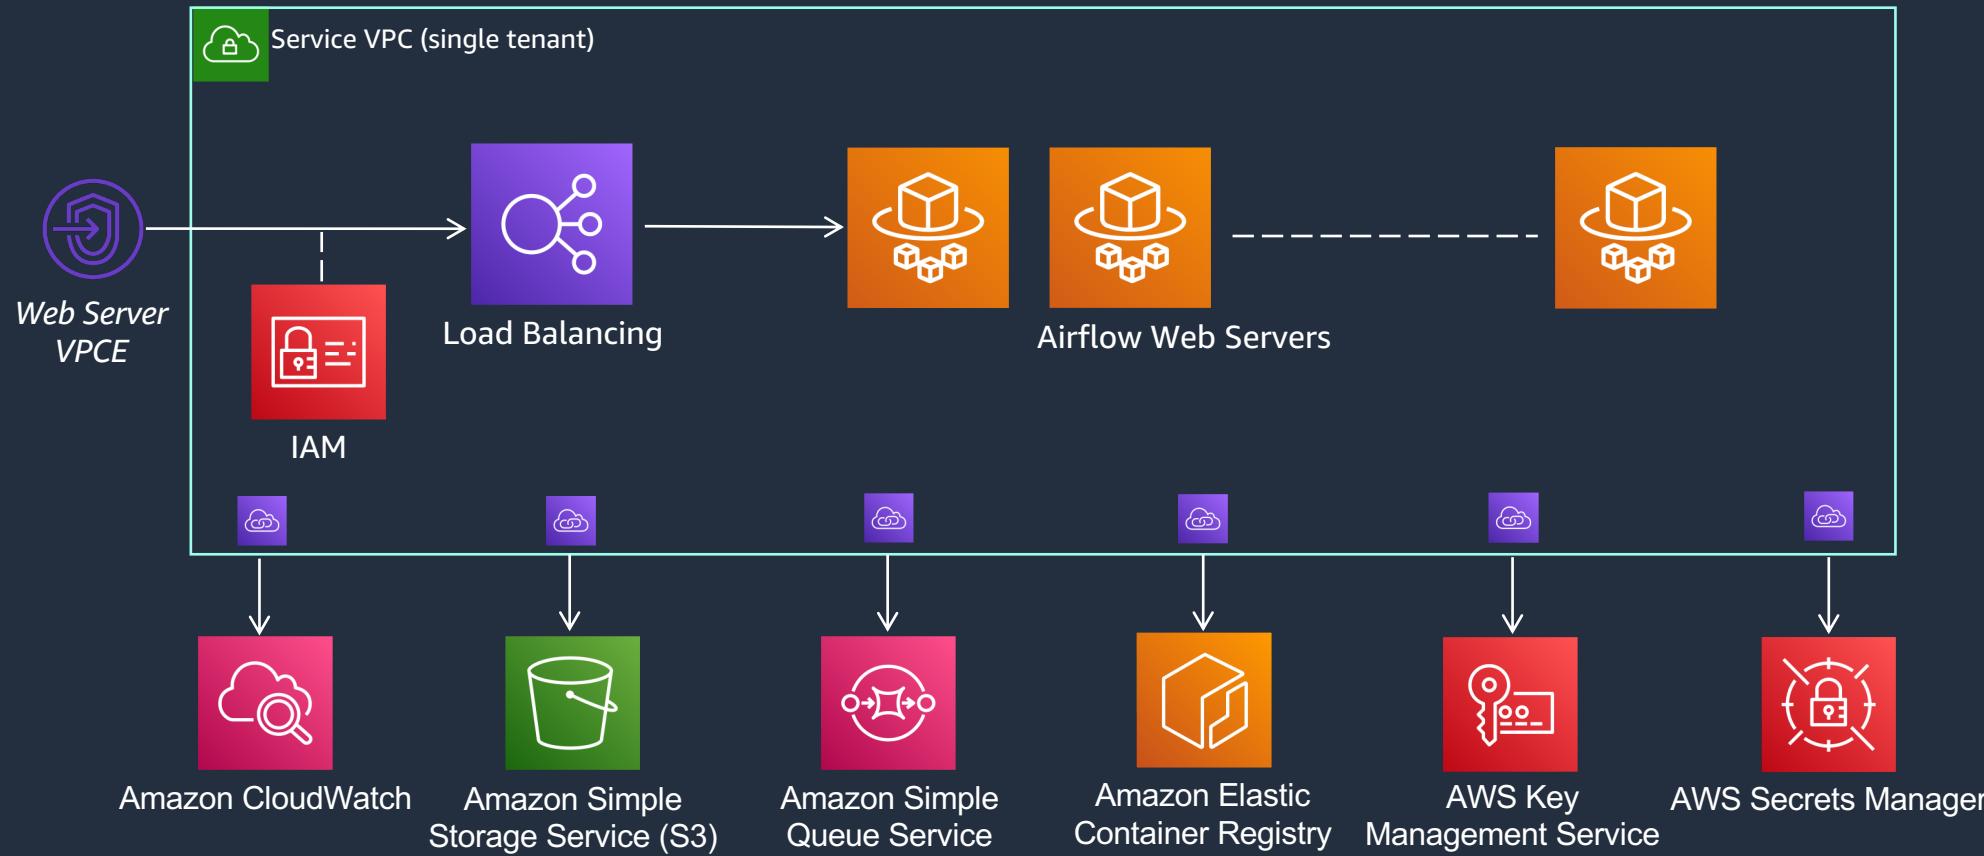
 For private network access, the Airflow web server is reached via a VPC endpoint inside your VPC. Connecting to the endpoint requires additional setup. [Learn more about VPC endpoints](#) 



Public Web Server Option



Private Web Server Option



MWAA Local Runner

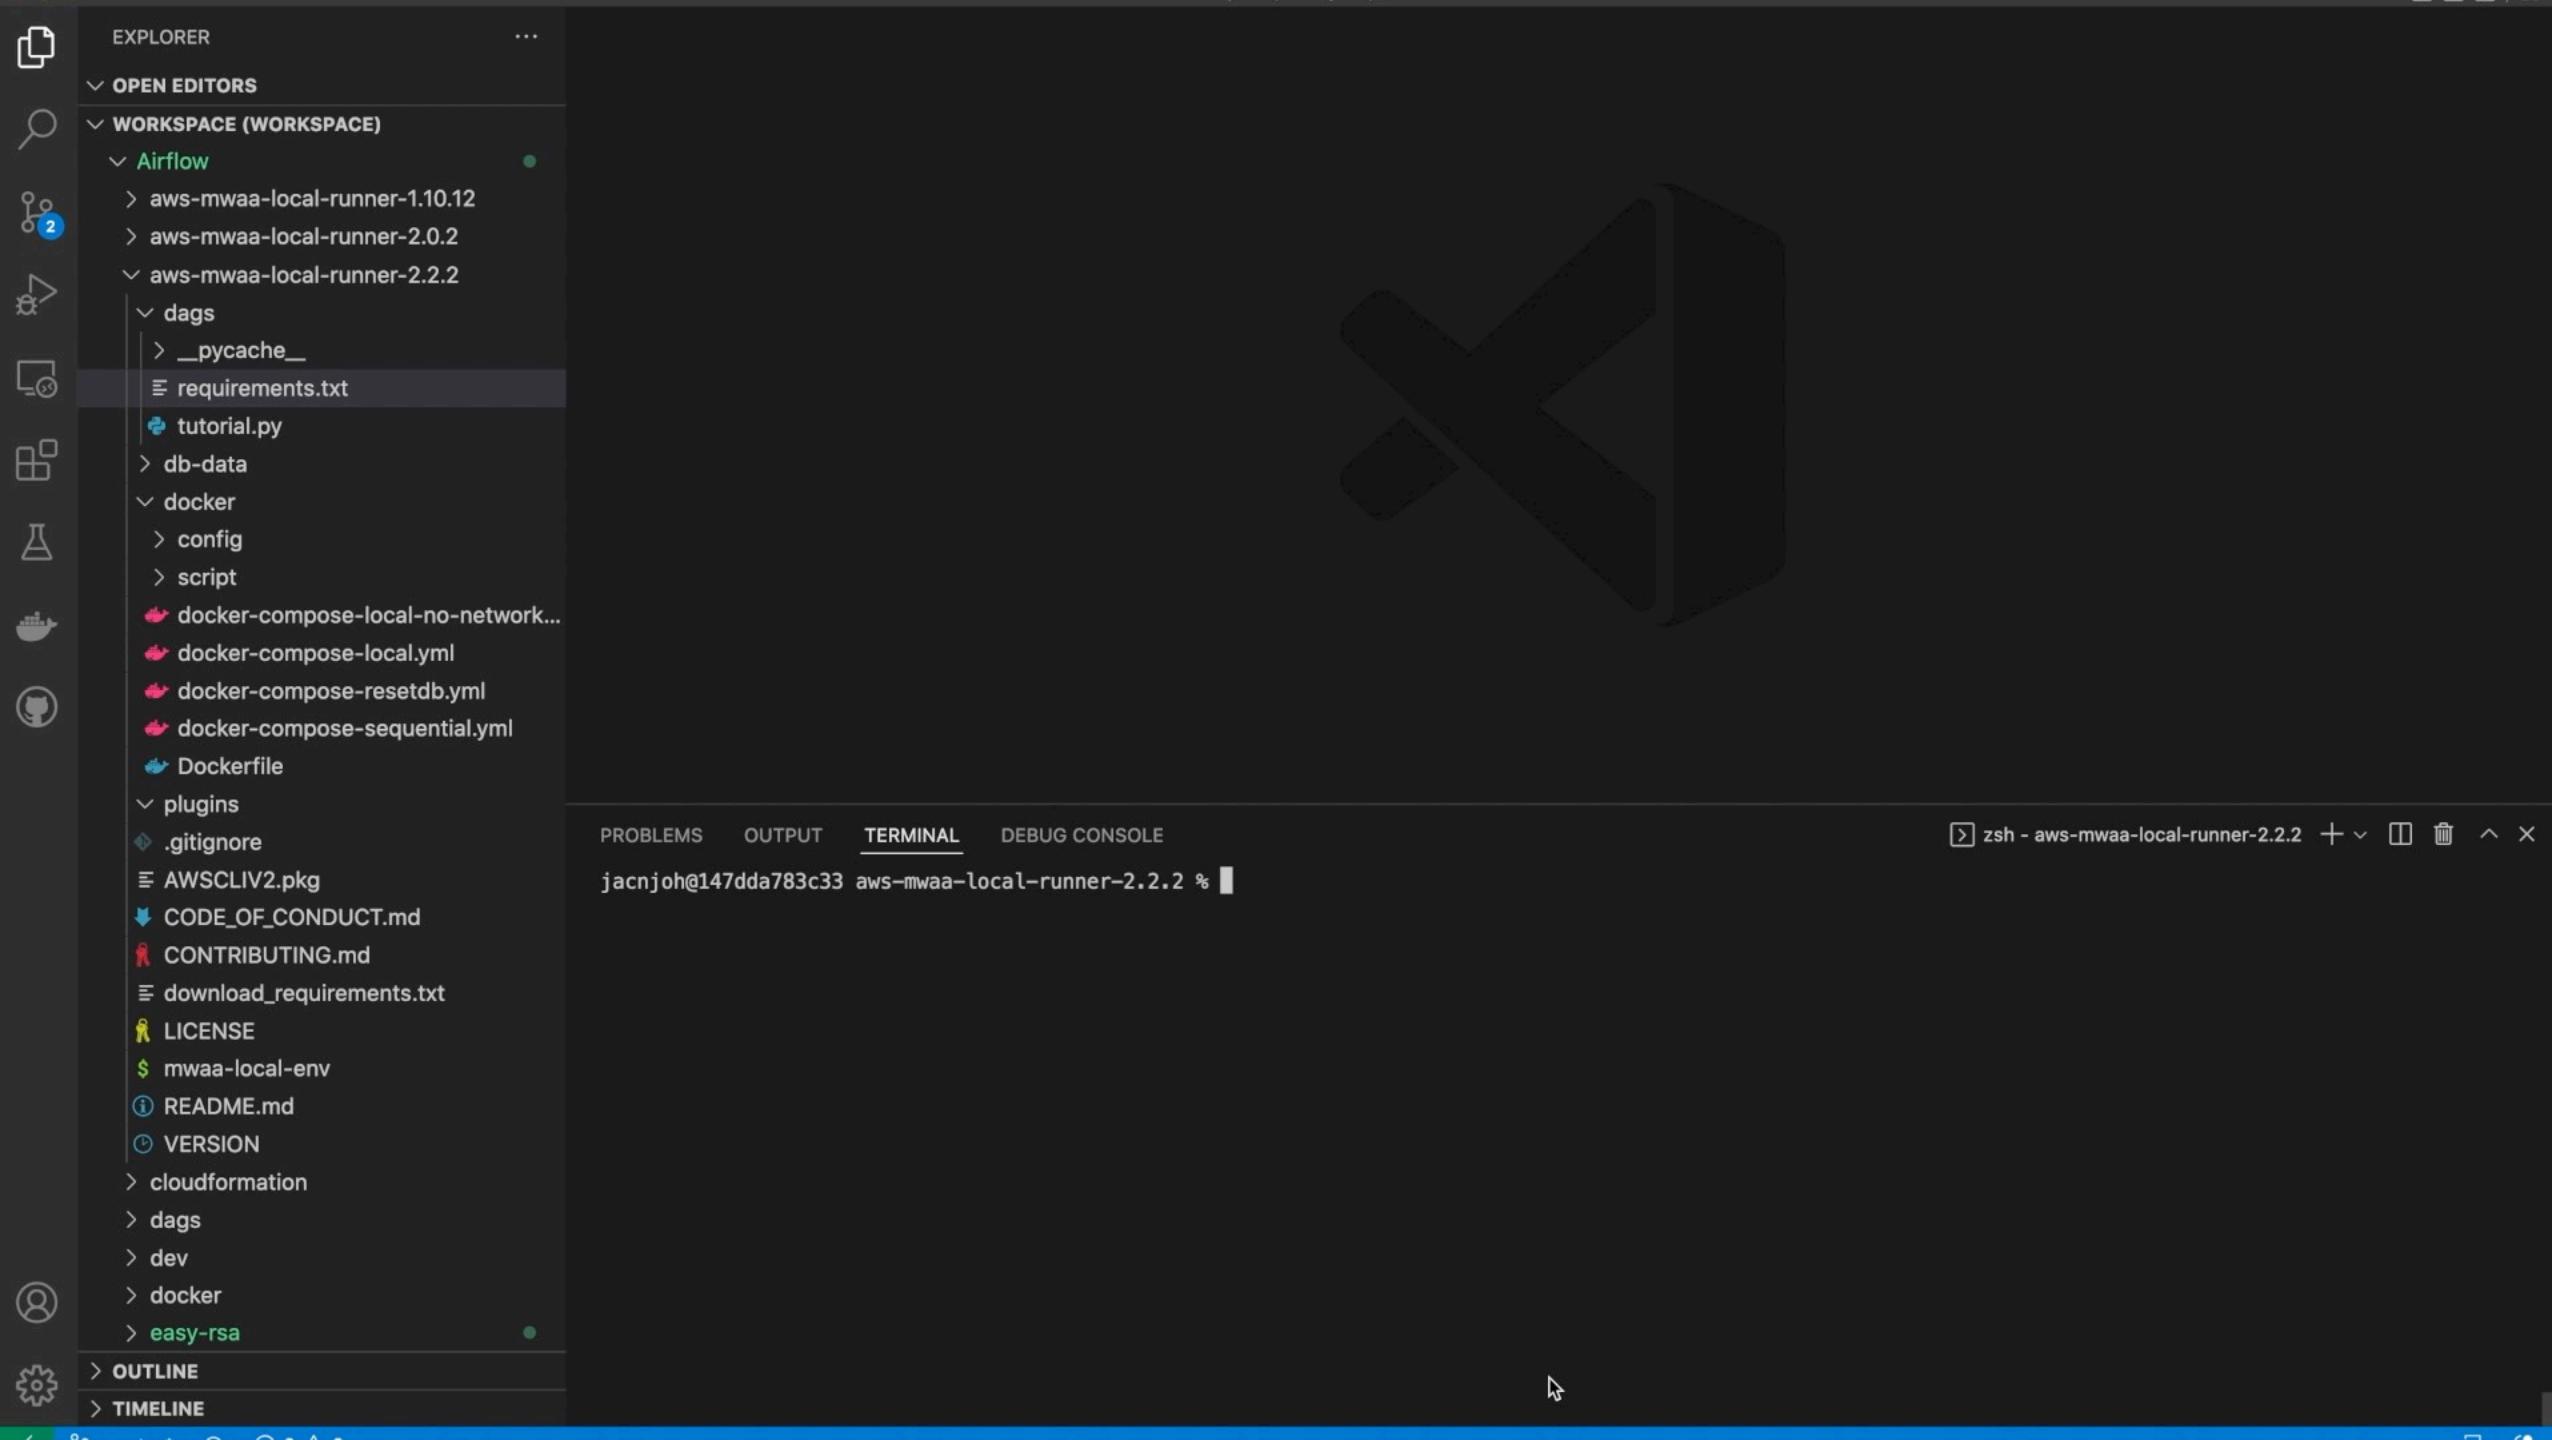
What is MWAA Local Runner?

- Docker Container
- All supported MWAA versions
- Same libraries and options as MWAA
- Requirements management utilities

```
=====
MWAA Local Runner CLI
=====
Syntax: mwaa-local-runner [command]

---commands---
help          Print CLI help
build-image   Build Image Locally
reset-db      Reset local PostgresDB container.
start         Start Airflow local environment. (LocalExecutor, Using postgres DB)
test-requirements Install requirements on an ephemeral instance of the container.
package-requirements Download requirements WHL files into plugins folder.
validate-prereqs Validate pre-reqs installed (docker, docker-compose, python3, pip3)
```

<https://github.com/aws/aws-mwaa-local-runner>



Best Practices

Things to Watch Out For

Logging

MWAA updated to Watchtower 2.0.1 to improve performance.

Can still run into account-wide PutLogEvent throttle

Access

Connectivity to AWS Services (NAT/
PrivateLink/
Security Groups)

IAM Execution Role access to AWS Services, AOK vs CMK

Migration

Everything on the Worker, Custom runtimes, Large CPU/Memory

Ephemeral Workers / Limited Bash Value

Running Airflow at Scale

Shameless self-promotion:

<https://airflowsummit.org/sessions/2022/managing-apache-airflow-at-scale/>

Open Source



AWS Airflow Open Source Team



Niko Oliveira



Shubham Mehta



Dennis Ferruzzi



Syed Hussain



Vincent Beck

What's Next?

What the MWAA Team is Working On

Airflow REST API Access

AWS has a very high bar for security and scalability of APIs

Working with our MWAA service team and Open Source Team

New Airflow Versions Faster

Looking at ways to speed up security reviews and deployments

Improved Performance and Monitoring

Incrementally improving performance and stability, and adding tools that will monitor and when possible recover on your behalf

Resources

Stay “In the Know”...

MWAA docs <https://docs.aws.amazon.com/mwaa>

Amazon MWAA landing page <https://aws.amazon.com/mwaa>

Amazon MWAA resources page <https://aws.amazon.com/managed-workflows-for-apache-airflow/resources/>

#airflow-aws Slack Channel: <https://apache-airflow.slack.com>



Thank you!

John Jackson

jacnjoh@amazon.com

 @JohnJacksonPM