



May 23–27, 2022

AIRFLOW SUMMIT



Hey Maintainer (and user)

Exercise your empathy



Airflow Survey!

<https://bit.ly/AirflowSurvey22>

About me



Jarek Potiuk

Independent Open-Source Contributor and Advisor

Airflow Committer & PMC member

Twitter: [@jarekpotiuk](https://twitter.com/jarekpotiuk)

What's empathy?

Empathy by Merriam-Webster

the action of understanding, being aware of, being sensitive to, and vicariously experiencing the feelings, thoughts, and experience of another of either the past or present without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner

Empathy by Merriam-Webster

the action of **understanding**, being aware of, being sensitive to, and vicariously experiencing the feelings, thoughts, and experience of another of either the past or present without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner

Empathy by Merriam-Webster

the action of understanding, **being aware of, being sensitive to**, and vicariously experiencing the feelings, thoughts, and experience of another of either the past or present without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner

Empathy by Merriam-Webster

the action of understanding, being aware of, being sensitive to, and **vicariously experiencing** the feelings, thoughts, and experience of another of either the past or present without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner

Empathy by Merriam-Webster

the action of understanding, being aware of, being sensitive to, and vicariously experiencing the **feelings, thoughts, and experience** of another of either the past or present without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner

Empathy by Merriam-Webster

the action of understanding, being aware of, being sensitive to, and vicariously experiencing the feelings, thoughts, and experience of **another** of either the past or present without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner

Empathy by Merriam-Webster

the action of understanding, being aware of, being sensitive to, and vicariously experiencing the feelings, thoughts, and experience of another of either the past or present **without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner**

Empathy by Merriam-Webster

the action of understanding, being aware of, being sensitive to, and vicariously experiencing the feelings, thoughts, and experience of another of either the past or present without having the feelings, thoughts, and experience **fully communicated in an objectively explicit manner**

Empathy by Merriam-Webster

the action of understanding, being aware of, being sensitive to, and vicariously experiencing the feelings, thoughts, and experience of another of either the past or present without having the feelings, thoughts, and experience fully communicated in an objectively explicit manner

also : the capacity for this

Why we have to talk about it?

- Some people are more naturally empathetic than others
- Empathy can be taught
- You studied a lot about programming, right?
- When is the last time you took empathy course?

Apache Airflow community

Who “we” are?



Committer
(Maintainer)



Contributor

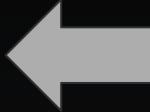


User

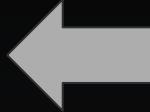
What's the journey ?



Committer
(Maintainer)



Contributor



User

This talk is not about



Contributor

Session

Adding an executor to Airflow: A contributor overflow exception



Vanessa Sochat

Research Software Engineer
@ Stanford University
Research Computing Center

Airflow
Summit 2020



@AirflowSummit
/Airflow-Summit

Join from anywhere.
<https://airflowsummit.org>

<https://youtu.be/RKEmAshcreE>

Adding An Executor to Airflow: A Contributor Overflow Exception



Vanessa Sochat



Adding An Executor to Airflow: A Contributor Overflow Exception



Vanessa Sochat

Be like Elmo!



User and Maintainer



Committer
(Maintainer)



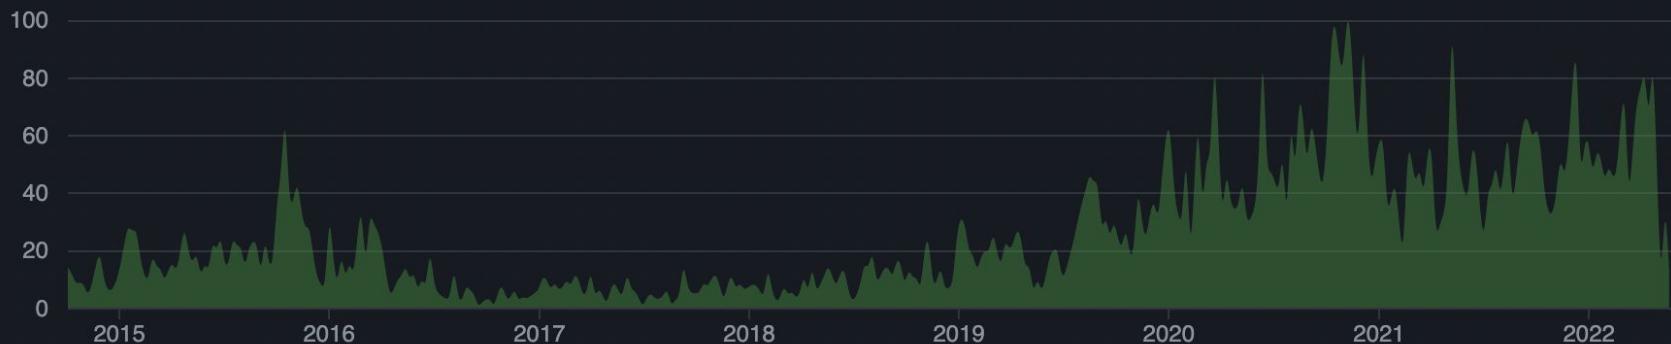
User

Hey maintainer

- Do you remember time when you were user?
- Did you know where to look in sources ?
- Did you know by heart all the configuration values?
- Did you know that someone already discussed it in Issue/devlist/discussion?
- Is it REALLY described in the documentation?
- Is this OK that the user might not know how airflow components are run?
- Is this OK that the user does not know Python?

Apache Airflow Pulse

Contributions to main, excluding merge commits and bot accounts



Hey user - do you know that:

- our maintainer to user ratio is less than 2 : 1000 ?
- we have more than 10 merged PRs a day ?
- there are PR with more than 100 comments ?
- there are about 80 - 100 messages generated a day from airflow repo activity?
- maintainers are not paid to help with troubleshooting ?
- there are SEARCHABLE(!) docs/issues/discussions/stack overflow?

What does it mean to be an empathetic maintainer?

You know nothing - John Snow

And now you can remove all the specific copy operators like S3ToSFTP , AzureBlobbToGCS ... Or, if they don't exist you don't need to implement them anymore!

What do you think @turbaszer? I don't know how the Core development is organized (I know there's a Jira and a mailing list, but know nothing of the organization processes) This will be a core change. But if Airflow aims to thrive, it's necessary (IMHO)



Direct the users

Apache Airflow version
2.1.1

What happened

When I am API calling for users with limit 1000, it returns only 500 requests. How or where can I increase limit more than 500 requests? Looks like by default airflow API returns max 500 requests.

I am running following command but always return 500 requests: https://airflow_url/api/v1/users?limit=1000

mik-laj commented on 19 Apr Member ...

Page size limit
To protect against requests that may lead to application instability, the stable API has a limit of items in response. The default is 100 items, but you can change it using `maximum_page_limit` option in `[api]` section in the `airflow.cfg` file.

<https://airflow.apache.org/docs/apache-airflow/stable/security/api.html>

Direct the users

josh-fell on 28 Jan Collaborator

Have you taken a look at `user_defined_macros` or `user_defined_filters` in the [Jinja templating docs](#)? They are pretty handy when you want to expose custom functions within .Jinja templating.

 Unmark as answer  1 

1 reply

a246530 on 28 Jan Author

Thank you for pointing this out! I was not fully aware of this. I just wrote a user defined filter to decode the xcom base64 value the way I initially suggested in my question and it worked just as expected.
Thanks again!

 1  1

 Write a reply

Answer selected by a246530

Check your assumptions



potiuk commented 13 days ago • edited

Member



...

The `all` is a development set of extras and it only works well in `main` branch.

Please use specific set of extras and providers you want to install rather than installing `all` - this should work fine.

You should not use it to install all dependencies of airflow in released airflow version. It is impossible to make it really without heavy complication of setup. I will make it clear in the docs because I think it is not really described well.

potiuk mentioned this issue 13 days ago

Clarify that bundle extras should not be used for PyPi installs #23697

Merged

potiuk added a commit to potiuk/airflow that referenced this issue 13 days ago

Clarify that bundle extras should not be used for PyPi installs ...

affefa0



potiuk commented 13 days ago

Member



...

Documentation clarification in #23697

Communicate via code

```
252 +     >&2 echo "WARNING! You should run the image with GID (Group ID) set to 0"
253 +     >&2 echo "           even if you use 'airflow' user (UID=50000)"
254 +
255 +     >&2 echo " You started the image with UID=$(id -u) and GID=$(id -g)"
256 +
257 +     >&2 echo " This is to make sure you can run the image with an arbitrary UID in the future."
258 +
259 +     >&2 echo " See more about it in the Airflow's docker image documentation"
260 +     >&2 echo "           http://airflow.apache.org/docs/docker-stack/entrypoint"
261 +
262 +     >&2 echo "
263 # We still allow the image to run with `airflow` user.
264
265 return
```

Invite the users to contribute

Apache Airflow version

2.3.0 (latest released)

What happened

We used to configure the Database using the AIRFLOW__CORE__SQLALCHEMY_CONN_CMD Environment variable.
Now the config option moved from CORE to DATABASE. However, we intended to keep backward compatibility as stated in the [Release Notes](#).

Upon 2.3.0 update however, the _CMD suffixed variables are no longer recognized for database configuration in Core - I think due to a missing entry here:

```
airflow/airflow/configuration.py  
Line 135 in 8622808  
135     ('database', 'sqlalchemy_conn'),
```

What you think should happen instead

We should only get a deprecation warning but the

Are you willing to submit PR?

Yes I am willing to submit a PR!

Indeed seems like a bug. Would you like to submit a PR fixing it @c-thiel ? Seems like an easy fix - and you could become one of the > 2000 contributors.

Add backward compatibility for core_sqlalchemy_conn_cmd #23441

Merged

mik-laj closed this as completed 23 days ago

Invite the users to contribute

I think one could add that to the docs.

Surely. Feel free. <https://airflow.apache.org/docs/apache-airflow/stable/start/local.html> - at the bottom right you will find "suggest a change on this page" and you can open a PR directly clicking on it. I think as a user who struggled with finding it, you are the best to describe it in the way that other users like you will find useful. Airflow is created by more than 1800 contributors - most of them voluntarily - and small contributions like that are the best way to start "giving back" to the community for the free software you get.)

 Unmark as answer

↑ 1



3 replies



boredland on 20 Dec 2021 Author



...

would you accept a PR removing that hard-coded port? "8080" is the default port. I removed that bit and it still started on 8080, as expected. So removing it would not break anything and it would make the configuration more flexible.



potiuk on 20 Dec 2021 Collaborator

sure. But some "docs" around that would be nice too (i.e. I



1



boredland on 20 Dec 2021 Author

ok, lets further discuss that [here](#).



1



Write a reply

fix(standalone): remove hardcoded webserver-port #20429

Merged

kaxil merged 3 commits into apache:main from boredland:patch-1

on 21 Dec 2021



Conversation

4

Commits

3

Checks

36

Files changed

2



boredland commented on 20 Dec 2021 · edited

Contributor

Empathy and assertiveness



In order to be able to help others,
you must help yourself first.

Set the boundaries

The screenshot shows a messaging interface with a dark background. At the top, there's a red circle highlighting a message from a user with a blue profile picture. The message reads "Hi, can I ask you an hand?". Below it, another message from a user with a grey profile picture is highlighted by a red circle, reading "nope". To the right of these messages is a toolbar with icons for smiley face, reply, forward, bookmark, and more. Below the messages, a timestamp "11:45" is followed by a note: "If it is troubleshooting: It is not OK to ask for troubleshooting in private messages. First of all you do not let others learn from help you might get and you do not give a chance for others to help you and the last thing is you demand someone's exclusive attention when this is a place where people help each other for free. If you need private help there are companies and people WHO are offering paid support for airflow. Ask in public and <https://dontasktoask.com/>".

- You cut yourself from help of others
- You demand exclusive attention for free
- You don't let others learn from help you might get

What does it mean to be an empathetic user?

Read The Friendly Manual

How to achieve dependencies between tasks #23032

Unanswered Lenchok asked this question in Q&A

Lenchok on 15 Apr

Apache Airflow version

2.2.2

What happened

My requirement is that the post-task will run only if the pre-dependency runs successfully. If the pre-dependency does not run or fails, the post-task will not run even when the time is up. How to implement this function?

potiuk on 15 Apr Collaborator edited •

Look at our docs and search for "task dependencies". You will find answer on your own without unnecessarily engaging other people. The documentation we have is pretty comprehensive: <https://airflow.apache.org/docs/>

Mark as answer ↑ 1

0 replies

Write a reply

Don't ask to ask

<https://dontasktoask.com/>

Don't ask to ask, just ask

Every now and then, in online chat rooms I hang around in, someone pops in and says something in the lines of,

Foobar123: Any Java experts around?

This is bad form, for several reasons. What the person is *actually* asking here is,

- You ask for too much responsibility
- Laziness
- So, to summarize, don't ask "*Any Java experts around?*", but rather ask "*How do I do [problem] with Java and [other relevant info]?*"

Use Github Discussions more

How to reproduce

What should we do to reproduce the problem? If you are not able to provide a reproducible case, please open a discussion instead.

Write

Preview



Please make sure you provide a reproducible step-by-step case of how to reproduce the problem as minimally and precisely as possible. Keep in mind we do not have access to your cluster or DAGs. Remember that non-reproducible issues will be

Ask a question or get support

Ask a question or request support for using Apache Airflow

Open

Using issues for questions

This issue was moved to a discussion.
You can continue the conversation there. [Go to discussion](#)

How do I install airflow using helm on arm architecture #23777

Closed Upvotes Comments Created 9 days ago · 1 comment

 * commented 9 days ago

Official Helm Chart version
1.5.0 (latest released)

Apache Airflow version
2.3.0 (latest released)

Kubernetes Version
1.18.2

Helm Chart configuration
No response

Docker Image customisations
No response

What happened
No response

What you think should happen instead
No response

How to reproduce
No response

Assignees
No one—assign yourself

Labels
`area:helm-chart` `kind:bug`

Projects
None yet

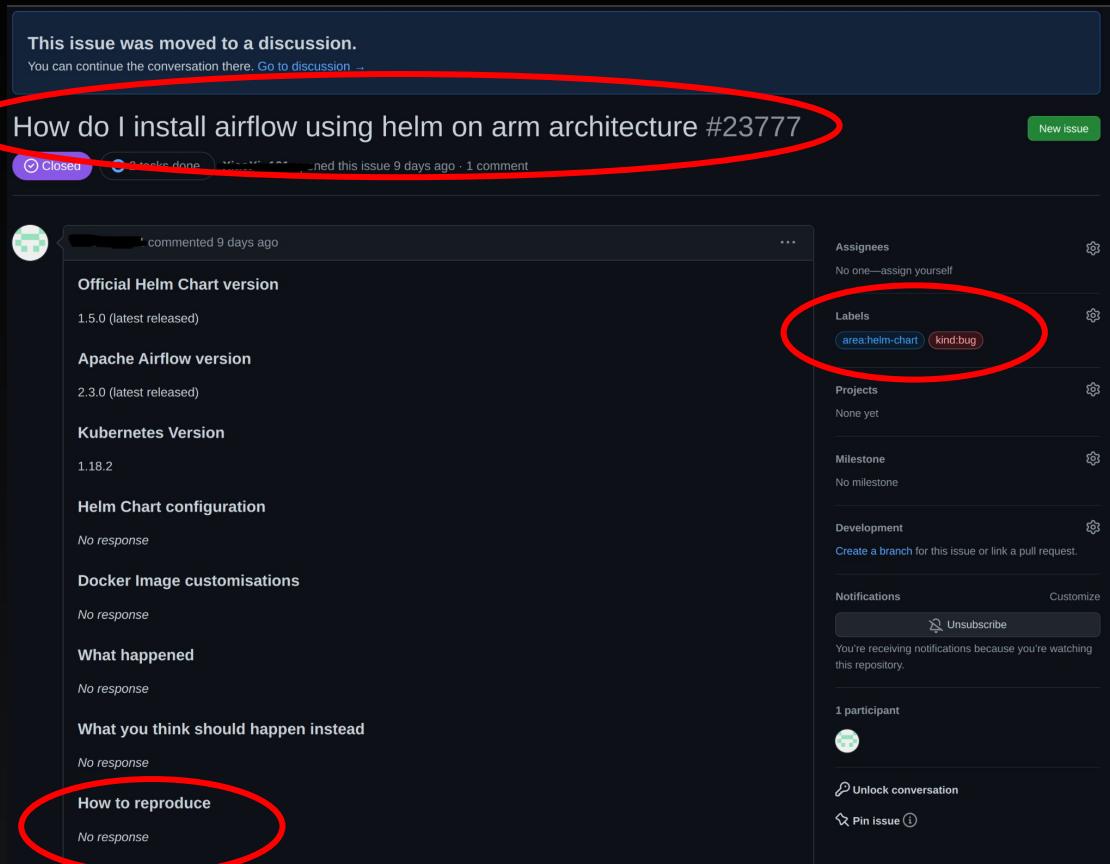
Milestone
No milestone

Development
Create a branch for this issue or link a pull request.

Notifications Customize
Unsubscribe
You're receiving notifications because you're watching this repository.

1 participant


Unlock conversation Pin issue



Help to help yourself

I am deploying airflow in Kubernetes with this chart: <https://github.com/apache/airflow/tree/master/chart> and I am using:

- AKS Kubernetes version: 1.16.13
- PostgreSQL hosted in Azure Database for PostgreSQL version: 10

If I use the PostgreSQL database sub chart it works perfectly.

I modify the values.yaml to connect to the external PostgreSQL database:

```
# Airflow database config
data:
  # If secret names are provided, use those secrets
  metadataSecretName: ~
  resultBackendSecretName: ~

  # Otherwise pass connection values in
  metadataConnection:
    user: <Myuser>
    pass: <Mypass>
    host: <server>
    port: 5432
    db: airflow
    sslmode: disable
  resultBackendConnection:
    user: <Myuser>
    pass: <Mypass>
    host: <server>
    port: 5432
    db: airflow
    sslmode: disable
```

When I deploy it the scheduler and webserver pods fail. I have checked the connection string and using it manually works fine

Add value

The screenshot shows a GitHub issue thread with four visible comments. The first comment from **ashb** on Nov 9, 2021, states: "Just confirmed this works on Mysql with #19425 when going from 2.1.2 to v2-2-test branch, so this will be fixed when 2.2.2 is out in a few days." The second comment from **ashb** on Nov 9, 2021, indicates the issue is closed as completed. The third comment from a user with a checkered profile picture on Jan 24, 2022, asks if anyone has resolved the issue. The fourth comment from **potiuk** on Jan 24, 2022, advises closing the issue and opening a new one if similar details are provided. Red ovals highlight the first two comments from **ashb**, the third comment from the user, and the last comment from **potiuk**.

ashb commented on 9 Nov 2021

Just confirmed this works on Mysql with #19425 when going from 2.1.2 to v2-2-test branch, so this will be fixed when 2.2.2 is out in a few days

ashb closed this as completed on 9 Nov 2021

commented on 24 Jan

hi, i'm facing this issue, anyone resolved it?

potiuk commented on 24 Jan

This issue is closed and fixed in 2.2.2. If you face similiar issue ↳ for 2.2.2 and 2.2.3 - please open a new one where you specify all details - this is likely different issue similar to this one.

It's really not helpful to say "I have the same issue" without providing no details when issue is closed already.

Knowing better

A screenshot of a GitHub pull request interface showing a conversation between two users. The top user, whose profile picture is a blue square with a white 'H', has commented: "Hey guys, what's going on with that task? I need this solution. Please, do it :) It's not complex I know it! Just do it!" A large red oval highlights the entire message. Below this, another user, "potiuk", has commented 9 hours ago: "Hey guys, what's going on with that task? I need this solution. Please, do it :) It's not complex I know it! Just do it!" If you know it's easy, feel free to provide PR. As simple as that. You can become one of the more than 2000 contributors. It's super easy way to repay the debt you have for using free software. A second red oval highlights the last sentence of potiuk's reply.

Hey guys, what's going on with that task? I need this solution. Please, do it :) It's not complex I know it! Just do it!

potiuk commented 9 hours ago

Hey guys, what's going on with that task? I need this solution. Please, do it :) It's not complex I know it! Just do it!

If you know it's easy, feel free to provide PR. As simple as that. You can become one of the more than 2000 contributors. It's super easy way to repay the debt you have for using free software.

Community is not a helpdesk

The screenshot shows a GitHub issue comment thread. The first comment, by a user with a green and white checkered profile picture, was posted on April 7. It reads: "I'm still having this issue on 2.2.4 any updates on this ticket or when it may get fixed. In the mean time i'm using the code snippet provided by @C...". Two words in this sentence are circled in red: "any updates" and "when it may get fixed". The second comment, by a user named potiuk with a profile picture of a man, was posted on April 12 and edited. It reads: "i'm still having this issue on 2.2.5 any updates on this ticket or when it may get fixed. In the mean time i'm using the code snippet provided by C...". Below this, potiuk adds: "Some parts of this issue is marked as fixed. You likely experience other issues. But we will never know unless you open a new issue and provide all the details, your logs and circumstances that you experience the problem. Unless you do that - there is no chance your issue will be even looked at (and even less likely fixed, maybe accidentally) - because no-one has any idea about Your issue details." The word "provide" is circled in red. The third comment, also by potiuk on April 12, reads: "Just be aware C... that you got the software for free - and a lot of the people contribute to it for free. And the best value you can bring is to spend some time on detailing your problems, providing as much as you can - including the circumstances, details and reproducibility. This way me smallest help you can provide for those who try to help people like you - on a free forum, for - again - software that you paid precisely 0 USD for." The words "you got the software for free", "a lot of the people contribute to it for free", "spend some time on detailing your problems", "providing as much as you can", "including the circumstances, details and reproducibility", "This way me smallest help you can provide for those who try to help people like you", "on a free forum", "again", and "software that you paid precisely 0 USD for" are all circled in red.

... commented on 7 Apr

i'm still having this issue on 2.2.4 any updates on this ticket or when it may get fixed. In the mean time i'm using the code snippet provided by @C...

potiuk commented on 12 Apr • edited

i'm still having this issue on 2.2.5 any updates on this ticket or when it may get fixed. In the mean time i'm using the code snippet provided by C...

Some parts of this issue is marked as fixed. You likely experience other issues. But we will never know unless you open a new issue and provide all the details, your logs and circumstances that you experience the problem. Unless you do that - there is no chance your issue will be even looked at (and even less likely fixed, maybe accidentally) - because no-one has any idea about Your issue details.

potiuk commented on 12 Apr

Just be aware C... that you got the software for free - and a lot of the people contribute to it for free. And the best value you can bring is to spend some time on detailing your problems, providing as much as you can - including the circumstances, details and reproducibility. This way me smallest help you can provide for those who try to help people like you - on a free forum, for - again - software that you paid precisely 0 USD for.

Thank you !

Q&A

Survey! <https://bit.ly/AirflowSurvey22>

