

Airflow as an AI Agent's toolkit

Going Beyond MCPs & Unlocking 1000+ Integrations



Kaxil Naik & Pavan Kumar Gopidesu

Introduction



Kaxil Naik

Sr Eng Director @ Astronomer
Airflow Committer & PMC Member



Pavan Kumar Gopidesu

Technical Lead @ Tavant
Airflow Committer & ASF Member






24+ hours of work!

Every. Schema. Change.

2-3 times per week



Customer use case - Relies on a lot of incoming data



Data Sources

- Hundreds of clients → thousands of S3, GCS feeds
- Formats vary (Parquet/CSV/JSON); schemas evolve
- Frequent schema drift
- Data across clouds




Consumers:


- Per-consumer tables (Postgres, Iceberg, Glue)
- Upstream drift breaks ingestion




When things break!



- Debug!
- Schema comparison across files & DB
- Change order processes
- Manual fixes
- Backfills
- Notifying 100+ consumers



Demo: Handling Schema Drift



Demo: Data Quality Checks

Airflow

localhost:28080/dags/lm_dq_check_s3_feed

Incognito

Dag lm_dq_check_s3_feed

Home

Dags

CustomerDataS3Feed

ApproveCustomerDataDQResults

Options

78s

39s

0s

Schedule Latest Run 2025-10-06, 12:47:15 ✓

Next Run Owner airflow Tags

Latest Dag Version v1

Search Dags ⌂ Trigger

Favorite Dag

Overview Runs Tasks Calendar Required Actions Audit Log Code Details

Last 24 Hours 2025-10-05, 13:24:37 - 2025-10-06, 13:24:37

0 Failed Tasks 0 Failed Runs

Last 2 Dag Runs

Duration

Queued Duration Run Duration

2025-10-06, 12:41:45 2025-10-06, 12:47:15

59.87 0.94

Run After

What currently takes 24+ hours:

- ✗ Manual schema comparison
- ✗ Change order approval
- ✗ Code changes and testing
- ✗ Debugging and backfills

What this PROTOTYPE demonstrates:

- ✓ Automated detection (using Apache DataFusion & AI)
- ✓ Human oversight at critical points
- ✓ Cross-cloud validation
- ✓ Business-friendly explanations



Does this problem resonate?

Would this solution work for you?

How does this work?

Automatic Context Injection

For SQL operations:

- Database type and version (PostgreSQL 15.2)
- Full schema from DbApiHook or Asset metadata
- Sample data (first few rows)
- Built-in safety rules

For File operations:

- File format (Parquet, JSON, CSV)
- Storage type (S3, GCS, Azure)
- File size, row count estimates
- Schema information
- Partitioning structure

```
{  
    "database": "PostgreSQL 15.2",  
    "schema": {  
        "customers": {  
            "customer_id": "int64",  
            "email": "string",  
            "created_at": "timestamp"  
        }  
    },  
    "sample_data": [...],  
    "safety_rules": [  
        "No DROP statements",  
        "No DELETE without WHERE"  
    ]  
}
```

Safety Mechanisms

We are not just sending prompts to LLMs.

Safety layers we're exploring:

- ✓ SQL Safety: Blocks DROP, DELETE without WHERE, TRUNCATE
- ✓ Human-in-the-Loop: Required for sensitive operations
- 🚧 Query validation: Parse and analyze before execution
- 🚧 Asset sensitivity: Mark Assets as auto-requiring approval for accessing it (PII)
- 🚧 Audit logging: All AI decisions tracked separately
- 🚧 Read-only by default: Write operations need explicit approval

Why Apache DataFusion?

- Unified query engine across object stores and DB. (S3, Postgres)
- Multiple formats (Parquet, JSON, CSV, Iceberg, Delta Lake)
- Single-node performance (no Spark overhead)
- Performance (in our test): 50M records in 14 seconds (with joins, groupby, min, max etc)



DataFusion is for **READING** only. Write uses DBApiHook

Current Approach - Specialized Operators

Current Implementation:

- `LLMSchemaCompareOperator` - for schema drift
- `LLMDataQualityOperator` - for validation
- `LLMFileAnalysisOperator` - for file analysis
- ... more to come for interacting with API(s) apart from Files & DB

Why specialized: Clear intent, better context for LLM, type safety, focused documentation

Alternative being explored: Unified `LLMOperator` with resource adapters

We're still figuring out the right abstraction. Your feedback will help.

Integration with Assets

Mark Asset as sensitive

Define how to access the Asset

- URI
- Connection

Define Asset type

- Data format
- Schema

Define metadata (for better AI context)

- Description
- Example queries

Future:

- Validations
- Statistics

```
from airflow.sdk import Asset

customer_asset = Asset(
    name="customer_data",
    uri="s3://bucket/customers/",
    conn_id="aws_default",
    schema={
        "customer_id": "int64",
        "email": "string",
        "phone_verified": "boolean"
    },
    sensitivity="pii",
    format_="parquet",
    statistics={"estimated_rows": 50000000}
)
```

Airflow PMC perspective

What Airflow Principles Must Stay

Whatever we build must preserve Airflow's core strengths:

- Deterministic DAG structure - static, reviewable, testable
- Observable - lineage, logging, monitoring
- Reliable - existing retry logic, error handling
- Safe - no breaking changes to existing workflows

Leveraging an LLM is just one task in a predictable pipeline.

We're NOT building AI that changes DAG structure.



What We're NOT Building

✗ NOT:

- AI that changes DAG structure
- Dynamic pipeline generation
- AI that makes architecture decisions
- Replacement for your data engineers

✓ YES:

- AI for repetitive, context-dependent tasks
- Deterministic DAGs with intelligent tasks
- Human oversight at critical points
- Audit trails and observability

Implementation Reality Check

If we proceed, the path would be:

- Phase 1: Experimental provider (`apache-airflow-providers-ai`)
- Phase 2: Community feedback and iteration
- Phase 3: Production-ready provider (if it proves valuable)
- Phase 4: Core integration (only if community demands it)


This could take multiple months to get right. No shortcuts.

Should Airflow have production-ready AI operators?

What We Need From You

Before we go further, we need community input:

- Is this solving real problems you face?
- What safety mechanisms are non-negotiable?
- How should we handle AI errors and edge cases?
- Right balance between intelligence & predictability?
- Should this be a provider or core feature?



Future Possibilities

- Expose 1000s of Hooks as AI Agent's "tools"
- Could AI-detected issues be...
- Should operators propose...
- Multi-agent validation (one...
- All AI calls can be logged i...
- ... any other wild ideas (?)

```
class HookToAIToolsMixin(ABC):
    """Mixin that Hooks in providers implement to expose AI capabilities."""

    @abstractmethod
    def describe_capabilities(self) -> AICapabilities:
        """Describe what this hook can do for AI systems"""

    @abstractmethod
    def get_schema_info(self, path: Optional[str] = None) -> SchemaInfo:
        """Get structural information about the resource"""

    def get_usage_examples(self) -> List[UsageExample]:
        """Provide examples of common operations"""
        return []

    def validate_ai_operation(self, operation: AIOperation) -> tuple[bool, None]:
        return True, None
```

How to get involved?

💬 Mailing list: dev@airflow.apache.org (AIP coming after Summit)

🗣 Slack: #airflow-3-dev channel

✉️ Pavan: gopidesupavan@gmail.com
[He wants your feedback directly]

✉️ Kaxil: kaxil@astronomer.io



Questions?
Concerns?
Ideas?



The 2025 Apache Airflow® Survey is here!

Fill it out to for a free Airflow 3 Fundamentals or DAG Authoring in Airflow 3 certification code

