



Airflow & Bigtop: Modernize and integrate time-proven OSS stack with Apache Airflow

Kengo Seki, Masatake Iwasaki
NTT DATA

3.0

Who we are



- Open source software developers / data engineers @NTT DATA
- Kengo Seki
 - Committer & PMC member of Apache Airflow and Apache Bigtop
 - ASF member
- Masatake Iwasaki
 - Committer & PMC member of Apache Hadoop and Apache Bigtop
 - ASF member

Agenda



- What is Apache Bigtop
- Introducing Bigtop features using the Airflow component
- Why we chose Airflow as our workflow orchestrator
- BigPetStore example DAG

What is Apache Bigtop

What is Apache Bigtop



- "Bigtop is an ASF project for Infrastructure Engineers and Data Scientists looking for comprehensive packaging, testing, and configuration of the leading open source big data components." (from <https://bigtop.apache.org/>)
- Bigtop is a Hadoop distribution developed by an open-source community and makes it easier for users to build and test their data platform.

Bigtop Features



1. Building components and packaging on widely-used Linux distributions
2. Deployment automation
3. Smoke tests
4. Provisioning components onto Docker containers

Each detail will be explained later.

History of Bigtop



2011 2012 2015 2016 2017 2018 2019 2020

	1.0.0	1.1.0	1.2.1	1.3.0	1.4.0	1.5.0
Ambari			2.5.0	2.6.1	2.6.1	2.6.1
Flink			1.1.3	1.4.2	1.6.4	1.6.4
Hadoop	2.6.0	2.7.1	2.7.3	2.8.4	2.8.5	2.10.1
HBase	0.98.12	0.98.12	1.1.9	1.3.2	1.3.2	1.5.0
Hive	1.0.0	1.2.1	1.2.1	2.3.3	2.3.3	2.3.6
Kafka	0.8.1.1	0.8.1.1	0.10.1.1	0.10.2.2	0.10.2.2	2.4.0
Oozie	4.0.1	4.2.0	4.3.0	4.3.0	4.3.0	4.3.0
Spark	1.3.1	1.5.1	2.1.1	2.2.1	2.2.3	2.4.5
Zeppelin		0.5.6	0.7.2	0.7.3	0.7.3	0.8.2
Zookeeper	3.4.6	3.4.6	3.4.6	3.4.6	3.4.6	3.4.13

History of Bigtop



2021

2022

2023

2024

	3.0.0	3.0.1	3.1.0	3.1.1	3.2.0	3.2.1
Ambari	2.7.5	2.7.5	2.7.5	2.7.5	2.7.5	2.8.0
Flink	1.11.3	1.11.6	1.11.6	1.11.6	1.15.0	1.15.3
Hadoop	3.2.2	3.2.2	3.2.3	3.2.4	3.3.4	3.3.5
HBase	2.2.6	2.2.6	2.4.8	2.4.8	2.4.13	2.4.13
Hive	3.1.2	3.1.2	3.1.2	3.1.2	3.1.3	3.1.3
Kafka	2.4.1	2.4.1	2.8.1	2.8.1	2.8.1	2.8.1
Oozie	5.2.1	5.2.1	5.2.1	5.2.1	5.2.1	5.2.1
Spark	3.0.1	3.0.1	3.1.2	3.1.2	3.2.1	3.2.1
Zeppelin	0.9.0	0.10.0	0.10.0	0.10.0	0.10.1	0.10.1
Zookeeper	3.4.14	3.4.14	3.5.9	3.5.9	3.5.9	3.5.9

History of Bigtop



2024 2025

	3.3.0	3.4.0	3.5.0
Airflow			2.10.5
Ambari			
Flink	1.16.2	1.20.0	1.20.1
Hadoop	3.3.6	3.3.6	3.3.6
HBase	2.4.17	2.6.1	2.6.2
Hive	3.1.3	4.0.1	4.0.1
Kafka	2.8.2	3.4.1	3.4.1
Oozie			
Ranger	2.4.0	2.6.0	2.7.0
Spark	3.3.4	3.5.3	3.5.6
Zeppelin	0.11.0	0.11.2	0.11.2
Zookeeper	3.7.2	3.8.4	3.8.4

Introducing Bigtop features using the Airflow component

Bigtop Features



1. Build components as native packages of widely-used Linux distributions
2. Deployment automation
3. Smoke tests
4. Provisioning components onto Docker containers

Building packages locally

```
$ curl -sLO https://downloads.apache.org/bigtop/bigtop-3.5.0/bigtop-3.5.0-project.tar.gz  
$ tar zxf bigtop-3.5.0-project.tar.gz  
$ cd bigtop-3.5.0  
  
$ sudo bigtop_toolchain/bin/puppetize.sh  
$ sudo puppet apply ¥  
--modulepath=".:/etc/puppet/modules:/usr/share/puppet/modules:/etc/puppet/code/modules" ¥  
-e "include bigtop_toolchain::installer"  
(or `./gradlew toolchain` if supported JDK is already available)  
  
$ ./gradlew bigtop-utils-pkg airflow-pkg repo  
$ ls output
```

1. Download and extract Bigtop tarball

2. Install Puppet then tools including compilers and libraries

3. build RPM/DEB packages then create YUM/APT repository

Building packages in containers

(installing Docker)

```
$ curl -sLO https://downloads.apache.org/bigtop/bigtop-3.5.0/bigtop-3.5.0-project.tar.gz  
$ tar zxf bigtop-3.5.0-project.tar.gz  
$ cd bigtop-3.5.0  
  
$ ./gradlew bigtop-utils-pkg-ind airflow-pkg-ind repo-ind  
  -POS=rockylinux-9  
  -Pprefix=3.5.0  
  -Pdocker-run-option="--privileged"  
  -Pmvn-cache-volume=true  
  
$ ls output
```

1. Download and extract Bigtop tarball

2. build RPM/DEB packages then create
YUM/APT repository in Docker container

Installing Bigtop pre-built RPM packages (Rocky Linux 9)

1. Download and install repository definition

```
$ sudo curl -L -o /etc/yum.repos.d/bigtop.repo ¥  
https://downloads.apache.org/bigtop/bigtop-3.5.0/repos/rockylinux-9/bigtop.repo
```

2. Import GPG key

```
$ sudo rpm --import ¥  
https://downloads.apache.org/bigtop/bigtop-3.5.0/repos/GPG-KEY-bigtop
```

```
$ sudo dnf install airflow bigtop-utils  
$ /usr/lib/airflow/bin/airflow version  
2.10.5
```

3. Update repository information and Install components via package manager

Installing Bigtop pre-built DEB packages (Ubuntu 24.04)

1. Download and install repository definition

```
$ sudo curl -o /etc/apt/sources.list.d/bigtop.list ¥  
https://downloads.apache.org/bigtop/bigtop-3.5.0/repos/ubuntu-24.04/bigtop.list
```

2. Import GPG key

```
$ curl -s https://downloads.apache.org/bigtop/bigtop-3.5.0/repos/GPG-KEY-bigtop ¥  
| sudo apt-key add -
```

```
$ sudo apt-get update  
$ sudo apt-get install airflow bigtop-utils  
$ /usr/lib/airflow/bin/airflow version  
2.10.5
```

3. Update repository information and Install components via package manager

Provisioning by Puppet manifest on Rocky Linux 9

```
$ curl -sLO https://downloads.apache.org/bigtop/bigtop-3.5.0/bigtop-3.5.0-project.tar.gz
$ tar zxf bigtop-3.5.0-project.tar.gz
$ cd bigtop-3.5.0
$ sudo bigtop_toolchain/bin/puppetize.sh
$ cat > bigtop-deploy/puppet/hieradata/site.yaml <<'EOF'
bigtop::bigtop_repo_uri: http://repos.bigtop.apache.org/releases/3.5.0/rockylinux/9/$basearch
bigtop::hadoop_head_node: "host1.example.com"
hadoop_cluster_node::cluster_nodes: ["host1.example.com ", "host2.example.com ", host3.example.com "]
hadoop_cluster_node::cluster_components: ["bigtop-utils", "airflow"]
EOF

$ cp -r bigtop-deploy/puppet/hiera* /etc/puppet/
$ puppet apply ¥
  --hiera_config=/etc/puppet/hiera.yaml ¥
  --modulepath=../bigtop-deploy/puppet/modules:/usr/share/puppet/modules:etc/puppet/code/modules ¥
  ./bigtop-deploy/puppet/manifests
```

Provisioning by Puppet manifest on Ubuntu 24.04

```
$ curl -sLO https://downloads.apache.org/bigtop/bigtop-3.5.0/bigtop-3.5.0-project.tar.gz
$ tar zxf bigtop-3.5.0-project.tar.gz
$ cd bigtop-3.5.0
$ sudo bigtop_toolchain/bin/puppetize.sh
$ cat > bigtop-deploy/puppet/hieradata/site.yaml <<'EOF'
bigtop::bigtop_repo_uri: http://repos.bigtop.apache.org/releases/3.5.0/ubuntu/24.04/$(ARCH)
bigtop::hadoop_head_node: "host1.example.com"
hadoop_cluster_node::cluster_nodes: ["host1.example.com ", "host2.example.com ", host3.example.com "]
hadoop_cluster_node::cluster_components: ["bigtop-utils", "airflow"]
EOF

$ cp -r bigtop-deploy/puppet/hiera* /etc/puppet/
$ puppet apply ¥
  --hiera_config=/etc/puppet/hiera.yaml ¥
  --modulepath=./bigtop-deploy/puppet/modules:/usr/share/puppet/modules ¥
  ./bigtop-deploy/puppet/manifests
```

Running smoke tests

```
$ curl -sLO https://downloads.apache.org/bigtop/bigtop-3.5.0/bigtop-3.5.0-project.tar.gz  
$ tar zxf bigtop-3.5.0-project.tar.gz  
$ cd bigtop-3.5.0  
  
$ . /usr/lib/bigtop-utils/bigtop-detect-javahome  
$ export AIRFLOW_HOME=/usr/lib/airflow  
$ ./gradlew bigtop-tests:smoke-tests:airflow:test -Psmoke.tests --info
```

provisioning and smoke-tests on containers

```
(installing Docker and Docker Compose; building packages and creating local repo )
$ cd provisioner/docker
$ ./docker-hadoop.sh --create 1 --memory 16g \
    --image bigtop/puppet:3.5.0-rockylinux-9 \
    --repo file:///bigtop-home/output \
    --enable-local-repo --disable-gpg-check \
    --stack bigtop-utils,airflow \
    --smoke-tests airflow

$ ./docker-hadoop.sh --list
NAME                           IMAGE                               COMMAND      SERVICE     CREATED
STATUS          PORTS
20251002_144838_r3229_bigtop_1  bigtop/puppet:3.5.0-rockylinux-9  "/sbin/init"  bigtop      3 minutes
ago      Up 3 minutes

$ ./docker-hadoop.sh --exec 1 /bin/bash
[root@17a680a4f14a /]#
```

Why we chose Airflow as our new workflow orchestrator



Why we chose Airflow as our new workflow orchestrator

- De-facto standard of workflow orchestrator
- Widespread userbase and vibrant community
- Affinity with the existing Bigtop components



Affinity with the existing Bigtop software stack

- Bigtop is composed of various ASF products, e.g., Hadoop, Spark, Flink, Kafka
- Airflow's diverse providers enable integration between them
- Enterprise-ready features are also supported, e.g., Kerberos support, proven scalability, fine-grained security

Airflow
integration
example:
BigPetStore DAG

BigPetStore: Simple ETL, data analytics and ML application

1. Generate synthetic sales transaction data

`generate_task`

■ success

SparkSubmitOperator

2. Read transaction data, normalize and load them into separated tables

`transform_task`

SparkSubmitOperator

3. Analyze transactions by some axes

`analyze_task`

SparkSubmitOperator

3. Recommend products for customers

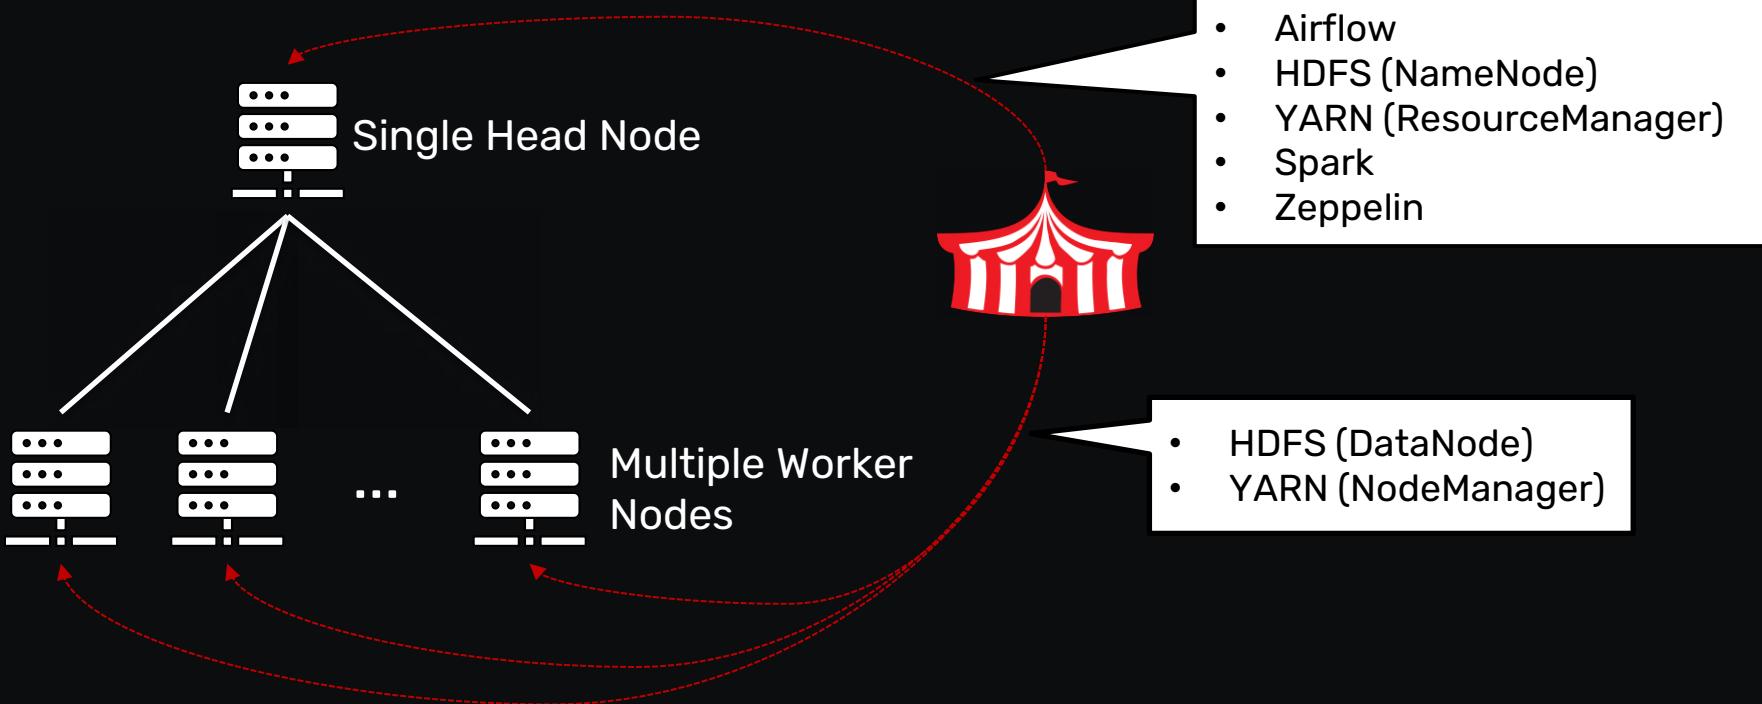
`recommend_task`

SparkSubmitOperator

For details, see

<https://github.com/apache/bigtop/blob/rel/3.5.0/bigtop-bigpetstore/bigpetstore-spark/README.md>

System Composition



Deployment process

- Download Bigtop 3.5.0 and install Puppet on each node

```
$ curl -sLO https://dlcdn.apache.org/bigtop/bigtop-3.5.0/bigtop-3.5.0-project.tar.gz  
$ tar xf bigtop-3.5.0-project.tar.gz  
$ sudo bigtop-3.5.0/bigtop_toolchain/bin/puppetize.sh
```

- Enable "install BigPetStore DAG" option (and some required settings) by modifying the configuration file

```
$ tail -5 bigtop-3.5.0/bigtop-deploy/puppet/hieradata/bigtop/cluster.yaml  
# Airflow  
airflow::server::executor: "LocalExecutor"      Default parameters regarding Airflow  
airflow::server::load_examples: false  
airflow::server::sqlalchemy_conn: "postgresql+psycopg2://..."  
airflow::server::install_bigpetstore_example: false  
$ cat <<EOT > bigtop-3.5.0/bigtop-deploy/puppet/hieradata/site.yaml  
bigtop::hadoop_head_node: fqdn.of.head.node  
hadoop::hadoop_storage_dirs: [/data]          Override parameters in accordance with our requirements  
hadoop_cluster_node::cluster_components: [airflow, hdfs, yarn, spark, zeppelin]  
airflow::server::install_bigpetstore_example: true  
EOT
```

Deployment process (continued)

- Distribute configuration file from head node to all of worker nodes

```
$ for node in $NODES; do scp bigtop-3.5.0/bigtop-deploy/puppet/hieradata/site.yaml ¥  
> ${node}:~/bigtop-3.5.0/bigtop-deploy/puppet/hieradata/site.yaml; done
```

- Copy configuration files to Puppet's default location and apply them on each node

```
$ sudo cp -r bigtop-3.5.0/bigtop-deploy/puppet/hiera* /etc/puppet/  
$ sudo puppet apply --modulepath bigtop-3.5.0/bigtop-deploy/puppet/modules:/usr/share/puppet/modules ¥  
> bigtop-3.5.0/bigtop-deploy/puppet/manifests  
(snip)  
Notice: /Stage[main]/Airflow::Server/Package[airflow]/ensure: created  
Notice: /Stage[main]/Airflow::Server/File[/var/lib/airflow/airflow.cfg]/ensure: defined content as  
'{sha256}acb4d3f8a5e9b1abe56b4b6d0aba5ae611e8fbbe10a17bbaeafbd71be345d7d'  
Notice: /Stage[main]/Airflow::Server/Service[airflow-webserver]/ensure: ensure changed 'stopped' to 'running'  
Notice: /Stage[main]/Airflow::Server/Exec[install-spark-provider]/returns: executed successfully  
Notice: /Stage[main]/Airflow::Server/File[/var/lib/airflow/dags]/ensure: created  
Notice: /Stage[main]/Airflow::Server/File[/var/lib/airflow/dags/example_bigpetstore.py]/ensure: defined  
content as '{sha256}ad9a3c7ef4b76fa1cfbd047207ff48143405b6a687d901690a3d693c7fc5b711'  
(snip)  
Notice: Applied catalog in 570.91 seconds
```



12:20 UTC

A log in icon with the text "Log In".

Sign In

Enter your login and password below:

Username:

A text input field containing the username "admin".

Password:

A password input field containing the masked password ".....".

A large blue rectangular button with the text "Sign In" in white.

Hadoop

Overview

Datanodes

Datanode Volume Failures

Snapshot

Startup Progress

Utilities ▾

Overview

'node0.fmkb5agmaolunimw4xslea2kkh.lx.internal.cloudapp.net:8020'

(✓active)

Started:	Thu Sep 25 20:58:13 +0900 2025
Version:	3.3.6, r79fb2857c7c52d680cbaba0e5cb1e9a0f0d33132
Compiled:	Thu Aug 14 22:06:00 +0900 2025 by jenkins from (HEAD detached at 79fb2857)
Cluster ID:	CID-4449d4e0-e2da-4713-bf61-936c8e079071
Block Pool ID:	BP-914695455-10.0.1.7-1758801490963



Welcome to Zeppelin!

Zeppelin is web-based notebook that enables interactive data analytics.

You can make beautiful data-driven, interactive, collaborative document with SQL, code and even more!

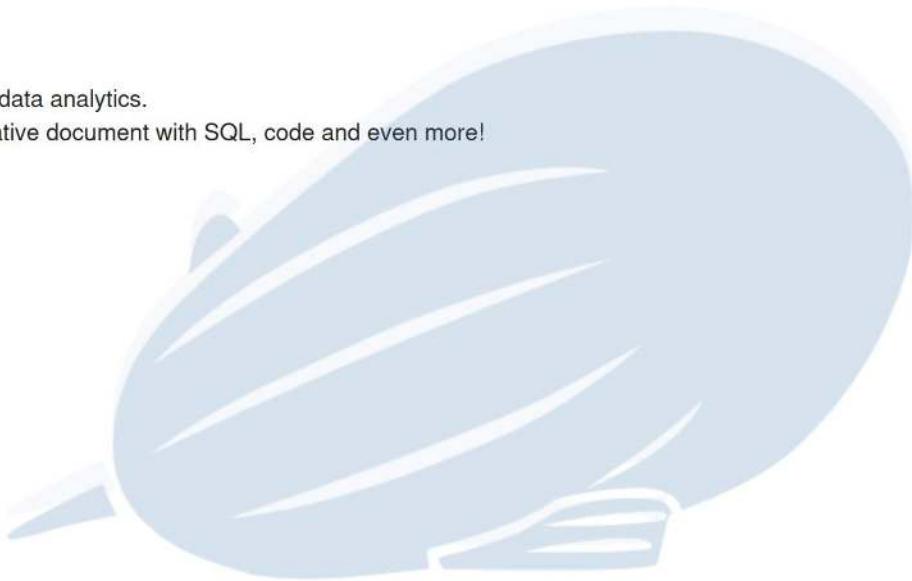
Notebook

Import note

Create new note

Filter

- Flink Tutorial
- Miscellaneous Tutorial
- Python Tutorial
- R Tutorial
- Spark Tutorial



Help

Get started with [Zeppelin documentation](#)

Community

Please feel free to help us to improve Zeppelin,

Any contribution are welcome!

Some additional steps

- Build BigPetStore's jar file ourselves on head node

```
$ cd bigtop-3.5.0/bigtop-data-generators  
$ ../gradlew clean publishToMavenLocal  
$ cd ../bigtop-bigpetstore/bigpetstore-spark  
$ ../../gradlew shadowJar  
$ mv build/libs/bigpetstore-spark-3.5.0-all.jar /tmp/
```

- Create Airflow user's home directory on HDFS

```
$ sudo -u hdfs hdfs dfs -mkdir /user/airflow  
$ sudo -u hdfs hdfs dfs -chown airflow:hadoop /user/airflow
```

- These steps will be automatically done in the future release



DAGs

Cluster Activity

Datasets

Security ▾

Browse ▾

Admin ▾

Docs ▾

DAGs

All 1

Active 0

Paused 1

Running 0

Failed 0

Filter DAGs by tag

DAG	Owner	Runs	Schedule
bigpetstore_dag	airflow	0 0 0 0	None



Airflow

DAGs

Cluster Activity

Datasets

Security ▾

Browse ▾

Admin ▾

Docs ▾

Trigger DAG: bigpetstore_dag

DAG conf Parameters

bigpetstore_jar_path:

/tmp/bigpetstore-spark-3.5.0-all.jar

Generated Configuration JSON and Dagrun Options ▾ Unpause DAG when triggered**Trigger****Cancel**



Airflow

DAGs

Cluster Activity

Datasets

Security

Browse

Admin

Docs

DAG: bigpetstore_dag

2025 / 09 / 12:22:00

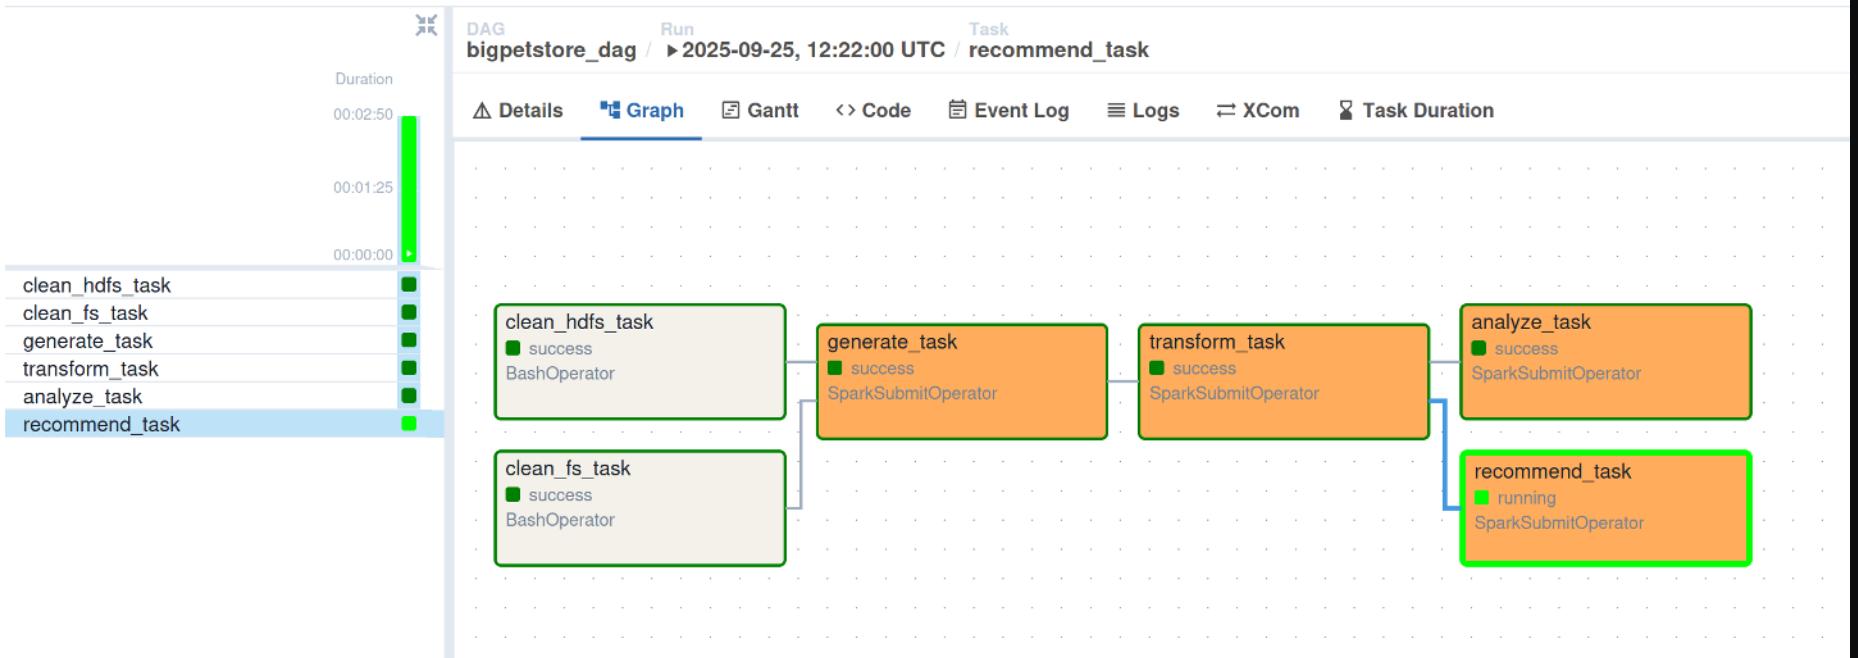
All Run Types

All Run States

Clear Filters

Press shift + / for Shortcuts

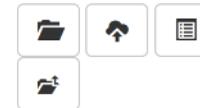
deferred failed queued removed restarting running



Browse Directory

/user/airflow

Go!



Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	airflow	hadoop	0 B	Sep 25 21:25	0	0 B	.sparkStaging	
<input type="checkbox"/>	drwxr-xr-x	airflow	hadoop	0 B	Sep 25 21:23	0	0 B	generated_data	
<input type="checkbox"/>	drwxr-xr-x	airflow	hadoop	0 B	Sep 25 21:24	0	0 B	transformed_data	

Showing 1 to 3 of 3 entries

Previous 1 Next



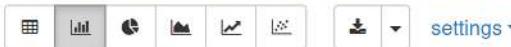
Untitled Note 1

```
202,  
193,  
289,  
...
```

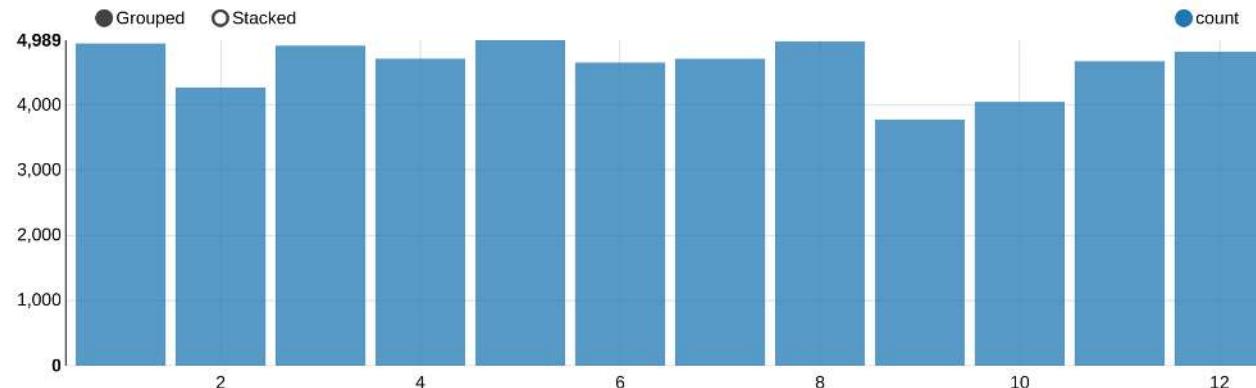
Took 3 sec. Last updated by anonymous at September 25 2025, 9:37:18 PM.

```
%sh  
echo %table  
echo -e month\\tcount  
jq -r '.transactionsByMonth[] | [.month, .count] | @tsv' /tmp/PetStoreStats.json
```

FINISHED ▾



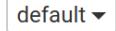
settings ▾



Took 0 sec. Last updated by anonymous at September 25 2025, 9:37:39 PM. (outdated)



Untitled Note 1



```
%sh  
cat /tmp/recommendations.json | jq .recommendations[:10]
```

```
[  
  {  
    "customerId": 0,  
    "productIds": [  
      202,  
      193,  
      141,  
      289,  
      40  
    ]  
  },  
  {  
    "customerId": 1,  
    "productIds": [  
      202,  
      193,  
      289,  
      ...  
    ]  
  }]
```

FINISHED ▶ ✎ 🔍 ⚙

Took 3 sec. Last updated by anonymous at September 25 2025, 9:37:18 PM.

Conclusion



- Bigtop is an ASF project that offers software distribution for big data processing
- Bigtop adopted Airflow as its new workflow orchestrator since v3.5.0
- By using Bigtop, users can easily build their own data platform including Airflow
- Airflow is a keystone of the integration between diverse Bigtop components
- BigPetStore DAG is an example of that integration and we'll add more diverse usecases and workloads in the future, e.g., streaming, GenAI, etc.

Questions?

Kengo.Seki sekikn@apache.org
Masatake Iwasaki iwasakims@apache.org