



AWS Lambda Executor: The Speed of Local Execution with the Advantages of Remote

Nikolas (Niko) Oliveira

3.0

Who Am I?



- Apache Airflow committer
- Sr. software engineer at Amazon
 - Amazon Managed Workflows for Apache Airflow (MWAA)
 - Founding member of the Amazon Apache Airflow Open Source Team

Executors in Airflow



- Executors facilitate the running of Airflow tasks (Task Instances)
- The Scheduler decides *when* a task runs; the executor decides *where* and *how*
- Executors run within the Scheduler process
- Pluggable and extensible, you can write your very own!

Executors in Airflow



- There are many types of Airflow executors, but some major ones include:
 - **Local Executors:** Airflow tasks are executed on the same host that the executor (i.e. scheduler) is running on. E.g.: **LocalExecutor**

Executors in Airflow



- There are many types of Airflow executors, but some major ones include:
 - **Local Executors:** Airflow tasks are executed on the same host that the executor (i.e. scheduler) is running on. E.g.: **LocalExecutor**
 - **Remote Queued/Batched Executors:** Airflow tasks are sent to a central queue where remote workers pull tasks to execute. Often workers are persistent and run multiple tasks at once: E.g.: **CeleryExecutor, AwsBatchExecutor**

Executors in Airflow




- There are many types of Airflow executors, but some major ones include:
 - **Local Executors:** Airflow tasks are executed on the same host that the executor (i.e. scheduler) is running on. E.g.: **LocalExecutor**
 - **Remote Queued/Batched Executors:** Airflow tasks are sent to a central queue where remote workers pull tasks to execute. Often workers are persistent and run multiple tasks at once: E.g.: **CeleryExecutor, AwsBatchExecutor**
 - **Remote Containerized Executors:** Airflow tasks are executed ad hoc inside containers/pods. Each task is isolated in its own environment. E.g.: **KubernetesExecutor, AwsEcsExecutor**



Speed vs Isolation


Speed & convenience vs isolation & scalability can you have it all?





Speed vs Isolation

Yes, with the Lambda Executor!



Introducing: Lambda Executor




- Ephemeral containerized compute environment
- Massively scalable architecture
- Low latency (warm pools, Python runtime)
- Resilient: Highly available, retries, and dead letter queue
- Ideal for short to medium running tasks:
 - 15 min max execution time
 - Restrictive memory and storage (~10GB)



Lambda Executor Architecture

1

Lambda Executor receives task from the Scheduler. Queueing them and ultimately calling an asynchronous invoke for each Airflow Task to be run.






Lambda Executor Architecture

2

Lambda starts a Firecracker VM based off of the provided image.

Users must build the image to their specification and create the Lambda Function.

Example app/handler code and image building tips found [here](#).



Lambda Executor Architecture


3

On Airflow Task completion, a message is sent to a results queue indicating pass/fail (SQS only as of now).

Async Lambdas cannot be described!

It is the responsibility of the Lambda app/handler code to send the results.

See provided example app.py in docs [here](#).




Lambda Executor Architecture

4

The Lambda function should be configured with a Dead Letter Queue (DLQ).

If an invocation fails catastrophically (OOM, uncaught exception, etc) the Lambda function will send a message to the DLQ.




Lambda Executor Architecture

5

The LambdaExecutor periodically reads from both SQS queues for results.

Updating Task state and communicating events back to the Scheduler.



Local vs Lambda



Triggered hello_world with new Run ID manual__2025-10-02T00:43:23.984706+00:00, it should start any moment now.

DAG: hello_world

Schedule: 1 day, 0:00:00

Next Run ID: 2025-10-02, 00:00:00 UTC

10 / 02 / 2025 12:43:25 AM

All Run Types


All Run States

[Clear Filters](#)

Auto-refresh 25

Press **shift** + **!** for Shortcuts

deferred failed queued removed restarting running scheduled shutdown skipped success up_for_reschedule up_for_retry upstream_failed no_status



Local vs Lambda



Triggered hello_world with new Run ID manual__2025-10-02T00:43:23.984706+00:00, it should start any moment now.

DAG: hello_world

Schedule: 1 day, 0:00:00

Next Run ID: 2025-10-02, 00:00:00 UTC

10 / 02 / 2025 12:43:25 AM

All Run Types


All Run States

[Clear Filters](#)

Auto-refresh 25

Press **shift** + **!** for Shortcuts

deferred failed queued removed restarting running scheduled shutdown skipped success up_for_reschedule up_for_retry upstream_failed no_status



Local vs Lambda



Triggered hello_world with new Run ID manual__2025-10-02T00:43:23.984706+00:00, it should start any moment now.

DAG: hello_world

Schedule: 1 day, 0:00:00

Next Run ID: 2025-10-02, 00:00:00 UTC

10 / 02 / 2025 12:43:25 AM

All Run Types


All Run States

[Clear Filters](#)

Auto-refresh 25

Press **shift** + **!** for Shortcuts

deferred failed queued removed restarting running scheduled shutdown skipped success up_for_reschedule up_for_retry upstream_failed no_status



System tests: ECS vs Lambda



ECS Executor

Apache Airflow - Amazon Provider Package Health

View the health of AWS service integrations for Apache Airflow

This live dashboard displays the current health of AWS service integrations available in the [Amazon Provider package](#) of Apache Airflow.

The following table shows data for all runs from the past 7 days of the [AWS System Tests](#) using the latest [Apache Airflow codebase](#).

The data currently being displayed reflects the tests run using the ecs executor.

System name	▲	Successes	▼	Failures	▼	Duration	▼
Local							
ECS executor							
Batch executor							
Lambda executor							
example_appflow_run		28		0		24 minutes	
example_athena		27		1		33 minutes	
example_batch		28		0		an hour	
example_bedrock		28		0		17 minutes	
example_bedrock_batch_inference		28		0		28 minutes	
example_cloudformation		27		1		17 minutes	
example_comprehend		28		0		24 minutes	
example_comprehend_document_classifier		28		0		33 minutes	
example_datasync		27		1		39 minutes	
example_dynamodb		28		0		17 minutes	

Lambda Executor

Apache Airflow - Amazon Provider Package Health

View the health of AWS service integrations for Apache Airflow

This live dashboard displays the current health of AWS service integrations available in the [Amazon Provider package](#) of Apache Airflow.

The following table shows data for all runs from the past 7 days of the [AWS System Tests](#) using the latest [Apache Airflow codebase](#).

The data currently being displayed reflects the tests run using the lambda executor.

System name	▲	Successes	▼	Failures	▼	Duration	▼
Local							
ECS executor							
Batch executor							
Lambda executor							
example_appflow_run		20		7		5 minutes	
example_athena		20		7		6 minutes	
example_batch		20		7		8 minutes	
example_bedrock		20		7		4 minutes	
example_bedrock_batch_inference		20		7		10 minutes	
example_cloudformation		20		7		4 minutes	
example_comprehend		20		7		9 minutes	
example_comprehend_document_classifier		20		7		11 minutes	
example_datasync		20		7		6 minutes	
example_dynamodb		20		7		4 minutes	

System tests: ECS vs Lambda



ECS Executor

Apache Airflow - Amazon Provider Package Health

View the health of AWS service integrations for Apache Airflow

This live dashboard displays the current health of AWS service integrations available in the [Amazon Provider package](#) of Apache Airflow.

The following table shows data for all runs from the past 7 days of the [AWS System Tests](#) using the latest [Apache Airflow codebase](#).

The data currently being displayed reflects the tests run using the ecs executor.

[Local](#) | [ECS executor](#) **Batch executor** | [Lambda executor](#)

System name	▲	Successes	▼	Failures	▼
example_appflow_run		28		0	
example_athena		27		1	
example_batch		28		0	
example_bedrock		28		0	
example_bedrock_batch_inference		28		0	
example_cloudformation		27		1	
example_comprehend		28		0	
example_comprehend_document_classifier		28		0	
example_datasync		27		1	
example_dynamodb		28		0	



Lambda Executor

Apache Airflow - Amazon Provider Package Health

View the health of AWS service integrations for Apache Airflow

This live dashboard displays the current health of AWS service integrations available in the [Amazon Provider package](#) of Apache Airflow.

The following table shows data for all runs from the past 7 days of the [AWS System Tests](#) using the latest [Apache Airflow codebase](#).

The data currently being displayed reflects the tests run using the lambda executor.

[Local](#) | [ECS executor](#) | [Batch executor](#) | **Lambda executor**

System name	▲	Successes	▼	Failures	▼
example_appflow_run		20		7	
example_athena		20		7	
example_batch		20		7	
example_bedrock		20		7	
example_bedrock_batch_inference		20		7	
example_cloudformation		20		7	
example_comprehend		20		7	
example_comprehend_document_classifier		20		7	
example_datasync		20		7	
example_dynamodb		20		7	

System tests: Batch vs Lambda



Batch Executor

Apache Airflow - Amazon Provider Package Health

View the health of AWS service integrations for Apache Airflow

This live dashboard displays the current health of AWS service integrations available in the [Amazon Provider package](#) of Apache Airflow.

The following table shows data for all runs from the past 7 days of the [AWS System Tests](#) using the latest [Apache Airflow codebase](#).

The data currently being displayed reflects the tests run using the batch executor.

System name	▲	Successes	▼	Failures	▼	Duration	▼
Local							
ECS executor							
Batch executor							
System name	▲	Successes	▼	Failures	▼	Duration	▼
example_appflow_run		28		0		14 minutes	
example_athena		28		0		20 minutes	
example_batch		28		0		33 minutes	
example_bedrock		28		0		10 minutes	
example_bedrock_batch_inference		28		0		21 minutes	
example_cloudformation		28		0		10 minutes	
example_comprehend		28		0		18 minutes	
example_comprehend_document_classifier		28		0		25 minutes	
example_datasync		28		0		25 minutes	
example_dynamodb		28		0		10 minutes	

Lambda Executor

Apache Airflow - Amazon Provider Package Health

View the health of AWS service integrations for Apache Airflow

This live dashboard displays the current health of AWS service integrations available in the [Amazon Provider package](#) of Apache Airflow.

The following table shows data for all runs from the past 7 days of the [AWS System Tests](#) using the latest [Apache Airflow codebase](#).

The data currently being displayed reflects the tests run using the lambda executor.

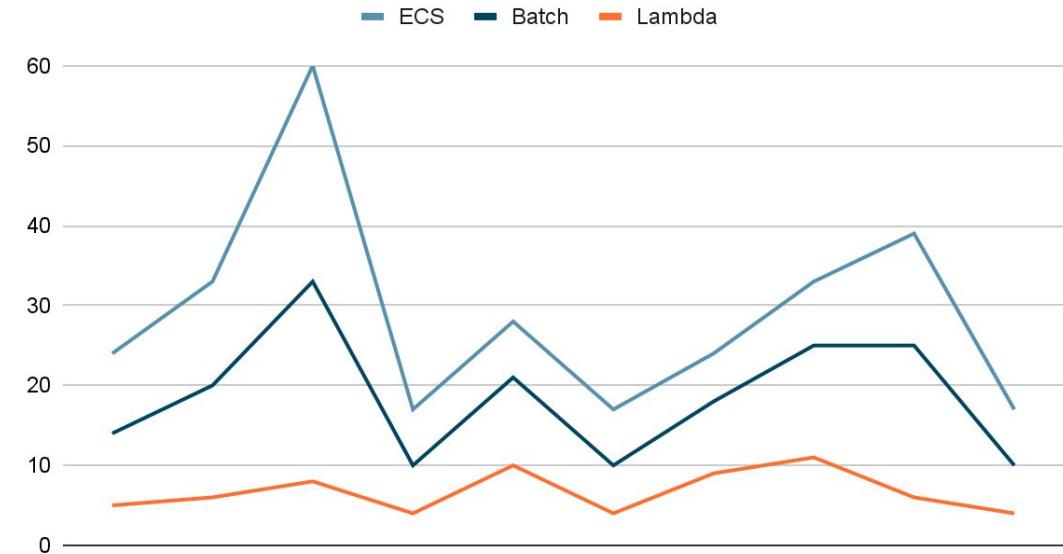
System name	▲	Successes	▼	Failures	▼	Duration	▼
Local							
ECS executor							
Batch executor							
System name	▲	Successes	▼	Failures	▼	Duration	▼
example_appflow_run		20		7		5 minutes	
example_athena		20		7		6 minutes	
example_batch		20		7		8 minutes	
example_bedrock		20		7		4 minutes	
example_bedrock_batch_inference		20		7		10 minutes	
example_cloudformation		20		7		4 minutes	
example_comprehend		20		7		9 minutes	
example_comprehend_document_classifier		20		7		11 minutes	
example_datasync		20		7		6 minutes	
example_dynamodb		20		7		4 minutes	



ECS vs Batch vs Lambda

The same data from the previous slides plotted alongside each other

Runtime per Executor (minutes)

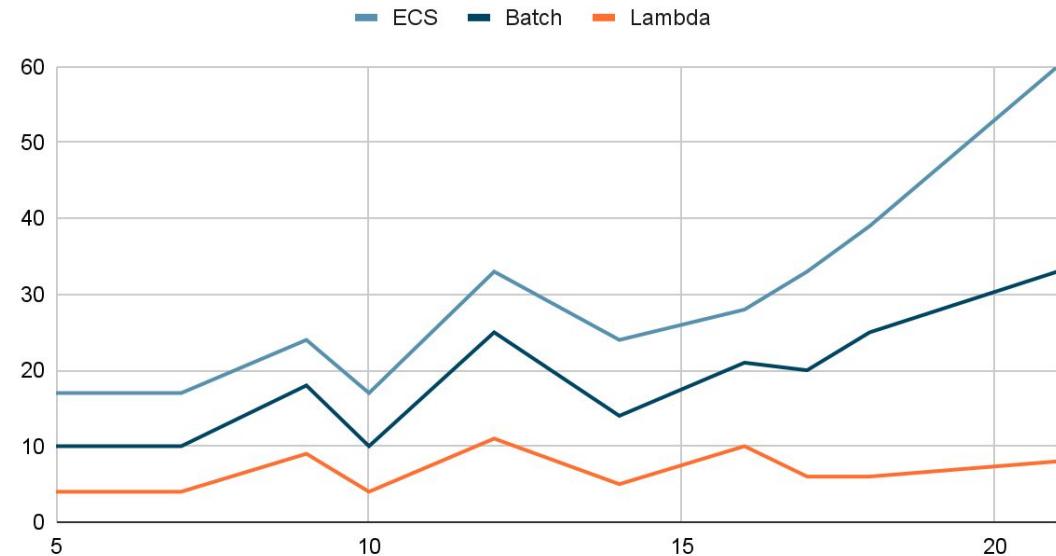




ECS vs Batch vs Lambda

If you order the data by the number of tasks you start to see the defining factor. Startup latency!

Runtime per Executor (minutes) - Ordered by Number of Tasks





ECS vs Batch vs Lambda

Plotting a trend, you can see Lambda runtime as tasks are added grows at a slower rate.

Runtime per Executor (minutes) - Ordered by Number of Tasks





What about the 15min Limit?

Of all the AWS Operators we test regularly. Only 1 was not usable with the Lambda Executor, in its basic use case.

Other Operators or usage patterns may differ.

Deferrable operators are your friend!



150+

A large green circle with a thick green border. Inside the circle, the text "150+" is displayed in a large, bold, green font.

1

A large red circle with a thick red border. Inside the circle, the number "1" is displayed in a large, bold, red font.

Not just Lambda: FaaS(t) Executors



- A new type of executor: Function as a Service (FaaS) --> FaaSt
- Distinctly different:
 - Queued/async execution pattern
 - FAST/resilient/scalable
 - Restrictive memory/storage
 - Execution time limits

Questions?

[Link to the docs](#)

Try it out now! Supports Airflow 2 and 3:

Currently Experimental



github.com/o-nikolas



linkedin.com/in/niko-oliveira-aws

