



Beyond Logs: Unlocking Airflow 3.0 Observability with OpenTelemetry Traces

Christos Bisias, G-Research

A large, bold, blue "3.0" graphic. The "3" is on the left, and the ".0" is on the right, with a small blue square at the bottom right corner of the ".0". The numbers are set against a black background.



Christos Bisias

Introduction

- I'm a software engineer from Greece
- I'm part of the G-Research Open Source team

GR-OSS

- G-Research is a leading quantitative research and technology firm
- G-Research is heavily investing in and using open source software
- The GR-OSS team is trying to leverage OSS to solve business problems

My background

- I have been contributing to Open Source projects on behalf of G-Research for the last 3.5 years
- In the past 1 year I've been doing a lot of work around OpenTelemetry and Airflow
- I adjusted Airflow's OTel implementation to make it support context propagation (more on that later)

Agenda

- What OTel is? Why use it with Airflow?
- OpenTelemetry basics
 - Traces and spans
 - Span data and context
 - Context propagation and mechanisms
- Demo
 - Execute an Airflow dag that exports sub spans
 - One of the tasks, makes a call to a running spring boot app that will also export a span

Why is my Dag slow?



Still running?

What is OpenTelemetry?

- OpenTelemetry is a collection of OS tools that are used for collecting traces, logs and metrics from applications
 - The collected data is later exported to visualization backends
 - This presentation will focus on **traces**

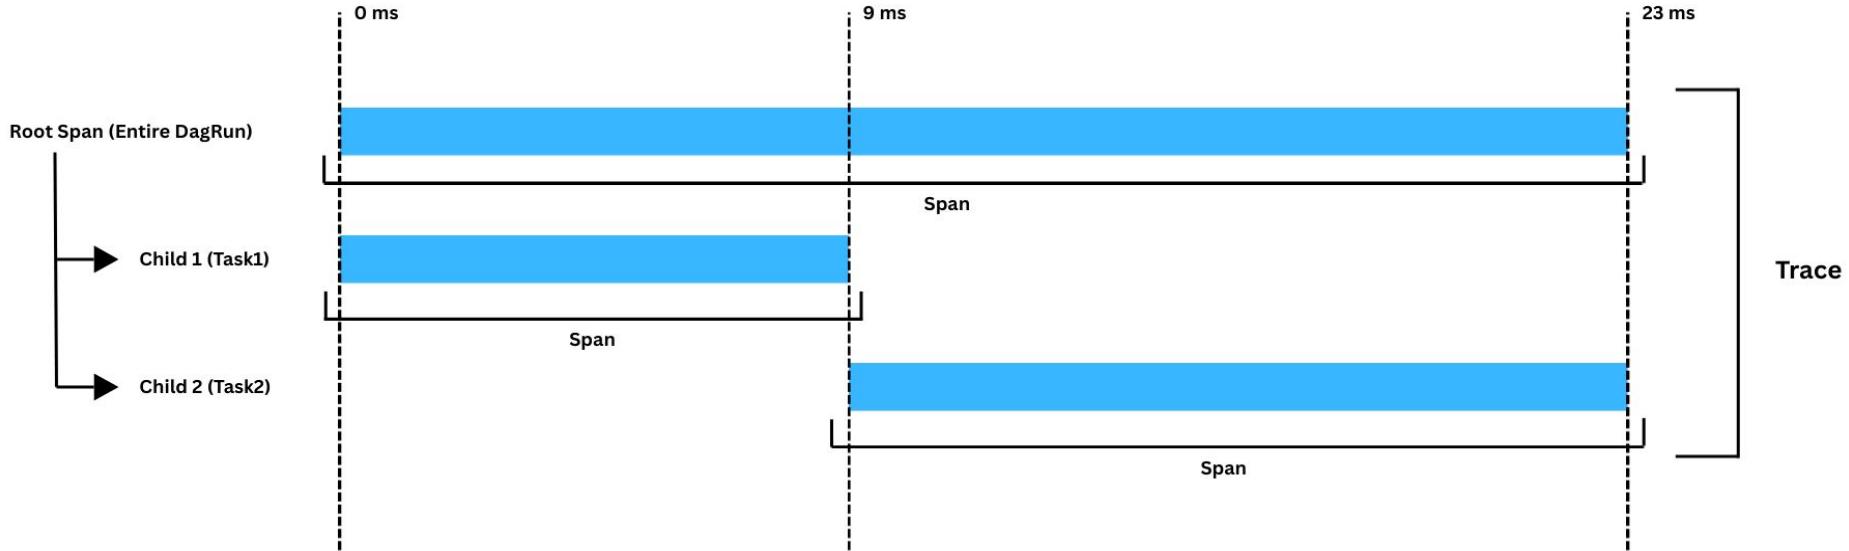
Why OpenTelemetry Traces?

- OTel is **technology agnostic**
- It has a different client implementation for each programming language
 - Python, Java, Golang, etc.
- The OTel SDK can be added as a dependency to any framework or application
- Especially useful in distributed systems
 - We can track an operation spreading across multiple services

Traces & Spans

- Let's assume that we have a large operation which consists of multiple small steps like API calls, DB calls, calculations etc.
- The collection of **all** the steps, will be a **trace**
 - An entire DagRun
- Each **individual** step that we are observing, will be a **span**
 - A single task

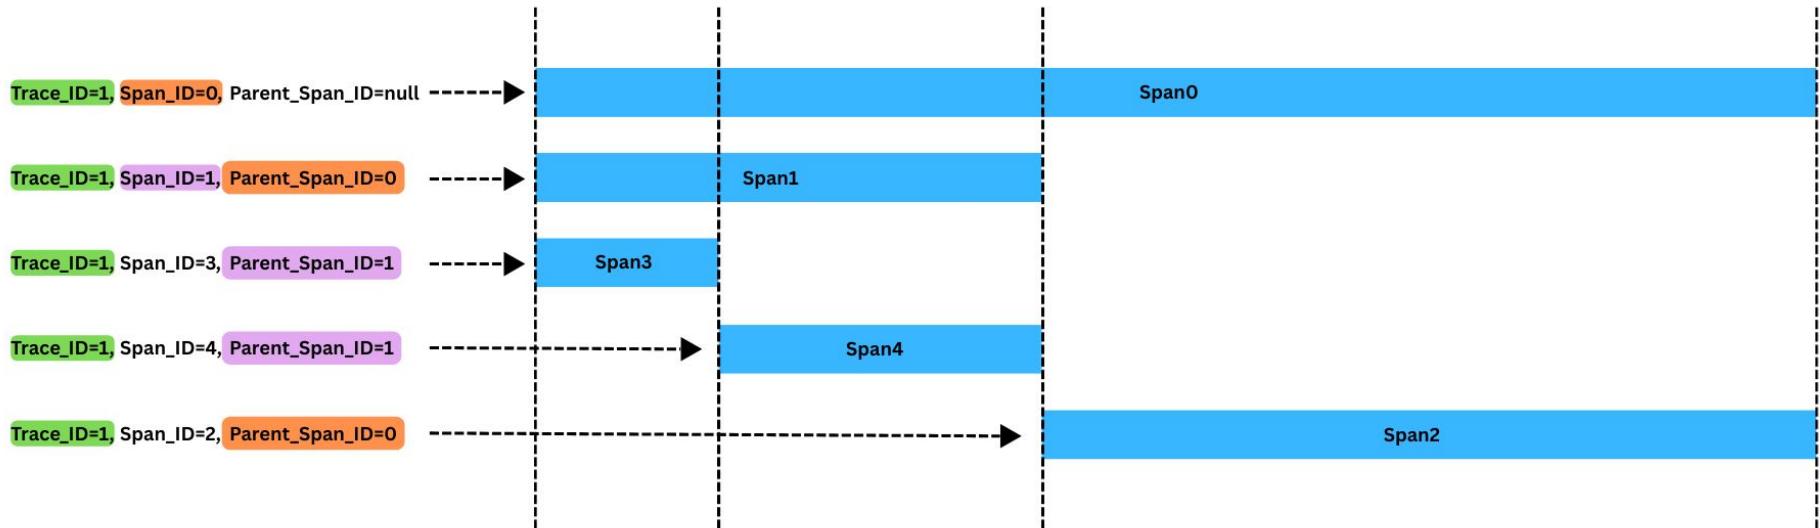
Trace & Spans Visualization



Why use OTel with Airflow?

- We can create our own sub spans under tasks to observe individual operations
- We can also monitor external calls to public APIs or other services running in our network

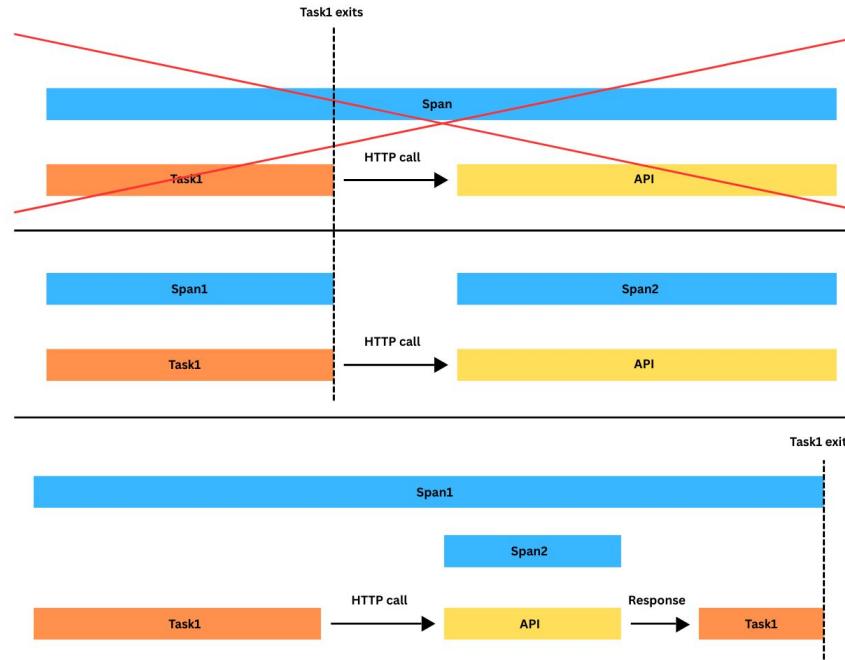
Span data & association



OTel Specifications

- Spans are thread local and cannot be shared outside of the current thread
- This is done to ensure that each process is solely responsible for handling its own spans

Cross-service trace

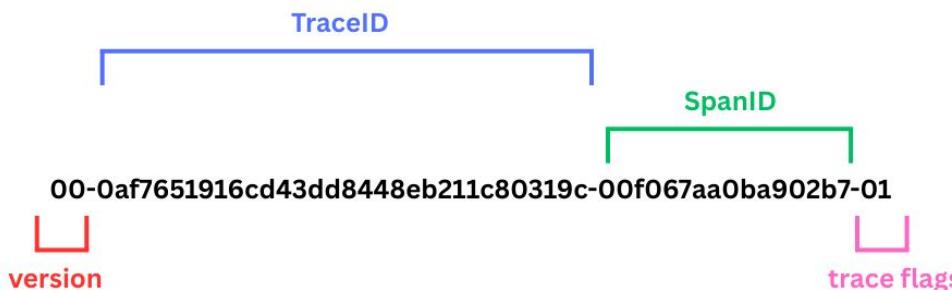


Context Propagation

- Context propagation is the method of sharing a span's context across services in a distributed system
- The context is used to create children spans
- The format follows the W3C specification

W3C Context Example

traceparent: 00-0af7651916cd43dd8448eb211c80319c-00f067aa0ba902b7-01



Propagation Mechanisms

- Common mechanisms of context propagation
 - HTTP calls
 - In the HTTP headers
 - gRPC calls
 - In the metadata
 - Message queues
 - Within the message
- Custom mechanism depending on the system
 - In Airflow we are propagating the context by storing it in the DB and later retrieving it

Capturing & Exporting Spans

- **Auto-instrumentation**
 - the instrumentation library automatically injects telemetry collection code into the application at runtime without requiring manual code changes
 - data are collected for certain frameworks, databases, HTTP clients and other common components
 - this is **strict**, it can't monitor anything that the library doesn't recognize
- **Manual** code injection into the app
 - the user has to write the code
 - it provides **flexibility** to monitor almost anything

Demo

Demo - Setup Explained

- Airflow with OTel traces enabled in the config
- An otel-collector service
- Jaeger as a visualization backend
- A Spring boot application running
 - With OTel SDK for Java
- Both Airflow and the spring boot app will create spans and export them to the common otel-collector
- The otel-collector will forward them to Jaeger

Demo - Dag Code

- The tasks will
 - Create sub spans using context propagation
 - Hook the auto-instrumentation library to monitor a GET request to a public API such as GitHub's
 - Make a call to the spring boot app and pass the current context in the HTTP headers
 - The app will create a sub span with the context

Demo - New Span

```
@task

def task1(ti):
    context_carrier = ti.context_carrier
    parent_context = otel_task_tracer.extract(context_carrier)

    with otel_task_tracer.start_child_span(
        span_name="part1_with_parent_ctx",
        parent_context=parent_context,
        component="dag",
    ) as p1_with_ctx_s:
        # Some work.

        logger.info("From part1_with_parent_ctx.")
```

Demo - Inject context into the headers



```
@task

def task2(ti):
    context_carrier = ti.context_carrier

    res = requests.get(
        "http://java-tester:7777/api/work",
        headers=context_carrier,
        timeout=25
    )
```

Webserver UI

Dag
otel_test_dag

task1
task2

Options ▾ 14s

7s
0s

Schedule Latest Run Next Run

2025-09-07, 14:58:24 ✓

Overview Runs Tasks Events Code Details

Last 24 hours 2025-09-07, 17:38:12 - 2025-09-08, 17:38:12

0 Failed Tasks 0 Failed Runs

Last 1 Dag Run

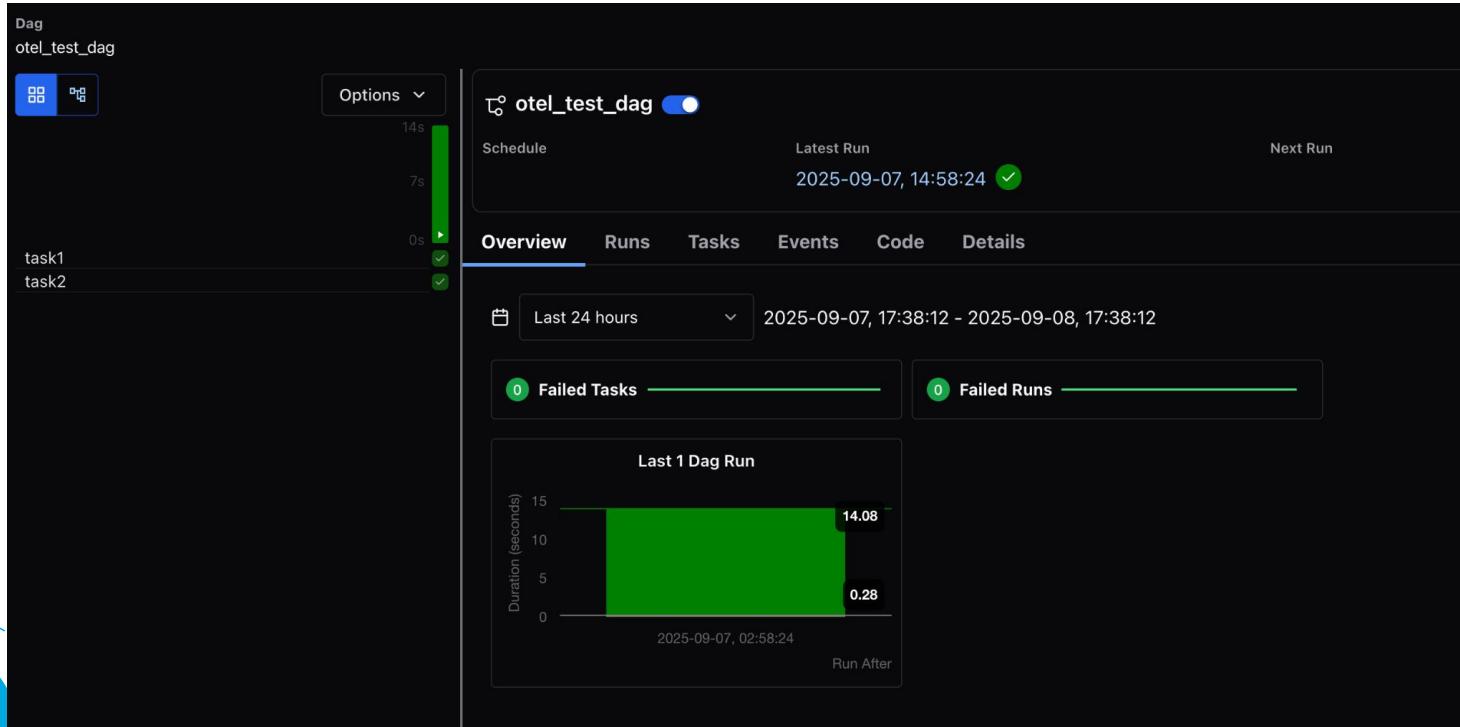
Duration (seconds)

15
10
5
0

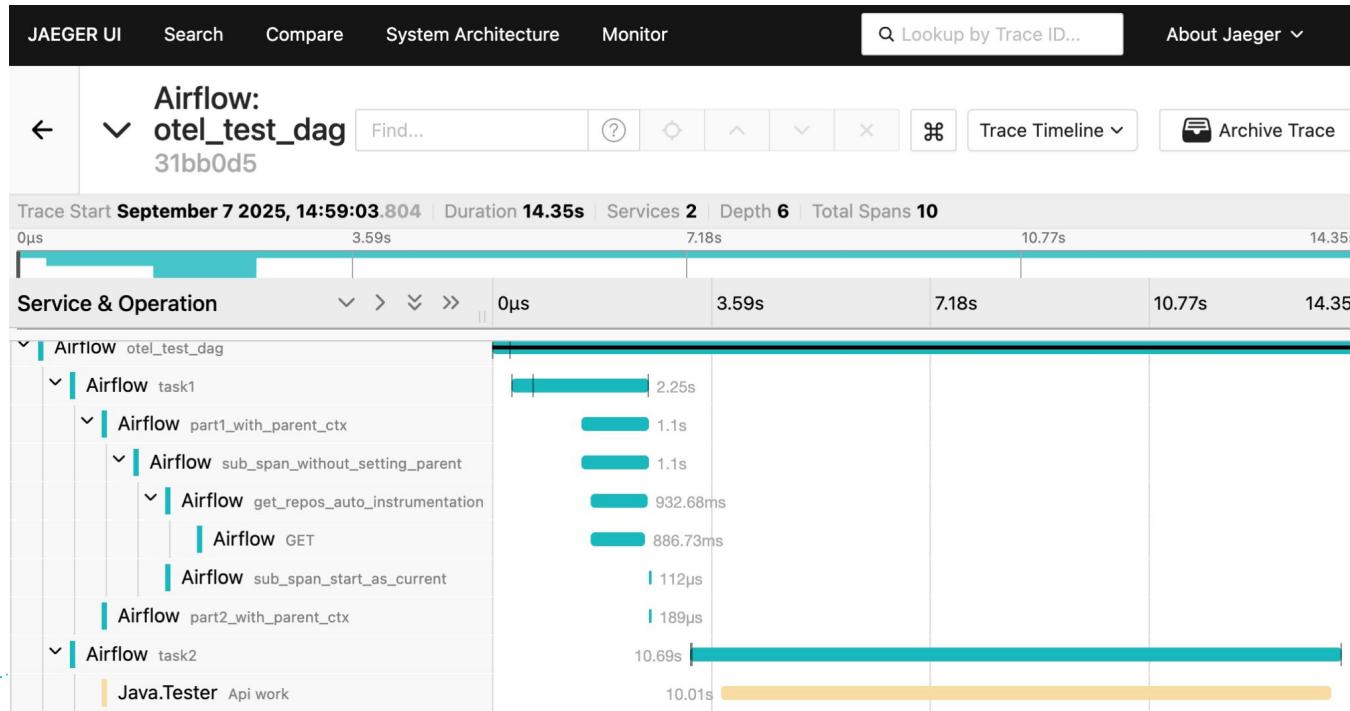
14.08
0.28

2025-09-07, 02:58:24

Run After



Jaeger UI



Next Topics to Research

- OpenTelemetry SDK initialization
- When to use
 - Simple vs Batch SpanProcessor
- Span Attributes
- Propagators
 - Context `inject()` and `extract()`
- SDK static context variable
 - Starting and ending span manually
- App or Request auto-instrumentation
- OTel Metrics and visualization with Prometheus and Grafana

Thank you! Questions?



Connect with me



Demo Code