

# Boosting dbt Core workflows performance Airflow Deferrable Capabilities

Pankaj Koti  
Senior Software Engineer  
@ Astronomer

Pankaj Singh  
Senior Software Engineer  
@ Astronomer

Tati Al-Chueyr  
Principal Software Engineer  
@ Astronomer

Airflow Summit, Seattle, USA - 8 October 2025

ASTRONOMER

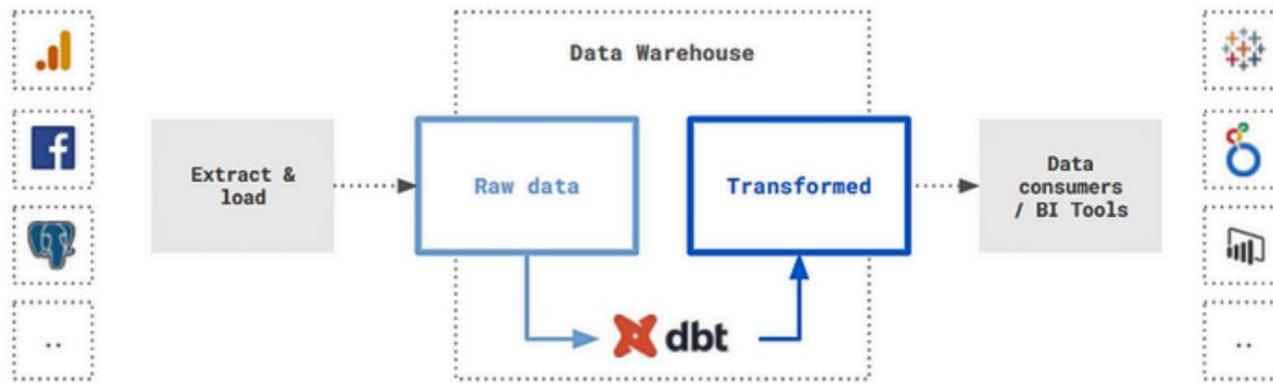
# Agenda

1. What is dbt
2. Why Cosmos
3. Performance challenge
4. Strategy 1: Airflow Deferrable
5. Strategy 2: Watcher
6. Takeaways

# 1. What is dbt

# dbt Core what is it?

dbt Core is an open-source tool that empowers **data practitioners** to transform



```
$ pip install dbt-core
```

# dbt Core project structure & syntax

```
➜ jaffle_shop git:(af-31) ✘ tree
.
├── dbt_project.yml
├── LICENSE
├── macros
│   ├── drop_table.sql
│   └── generate_alias_name.sql
└── models
    ├── customers.sql
    ├── docs.md
    ├── orders.sql
    ├── overview.md
    ├── schema.yml
    └── staging
        ├── schema.yml
        ├── stg_customers.sql
        ├── stg_orders.sql
        └── stg_payments.sql
├── packages.yml
├── profiles.yml
└── README.md
└── seeds
    ├── raw_customers.csv
    ├── raw_orders.csv
    └── raw_payments.csv

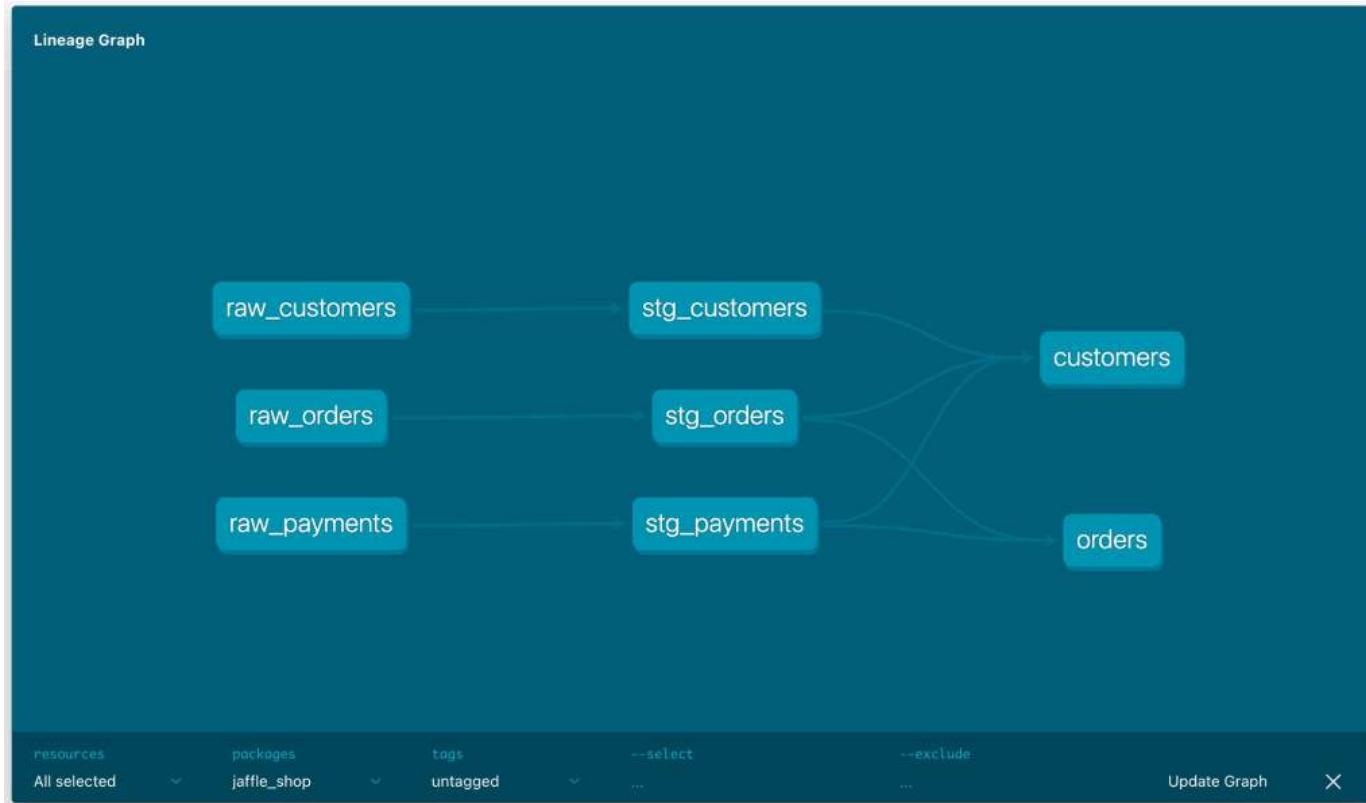
5 directories, 19 files
```

```
➜ jaffle_shop git:(af-31) ✘ cat models/staging/stg_payments.sql
with source as (
    #-
    -- Normally we would select from the table here, but we are using seeds to load
    -- our data in this project
    #
    select * from {{ ref('raw_payments') }}
),

renamed as (
    select
        id as payment_id,
        order_id,
        payment_method,
        -- `amount` is currently stored in cents, so we convert it to dollars
        amount / 100 as amount
    from source
)

select * from renamed
```

# dbt Core pipeline visualisation



# dbt Core limitations

dbt Core does not...

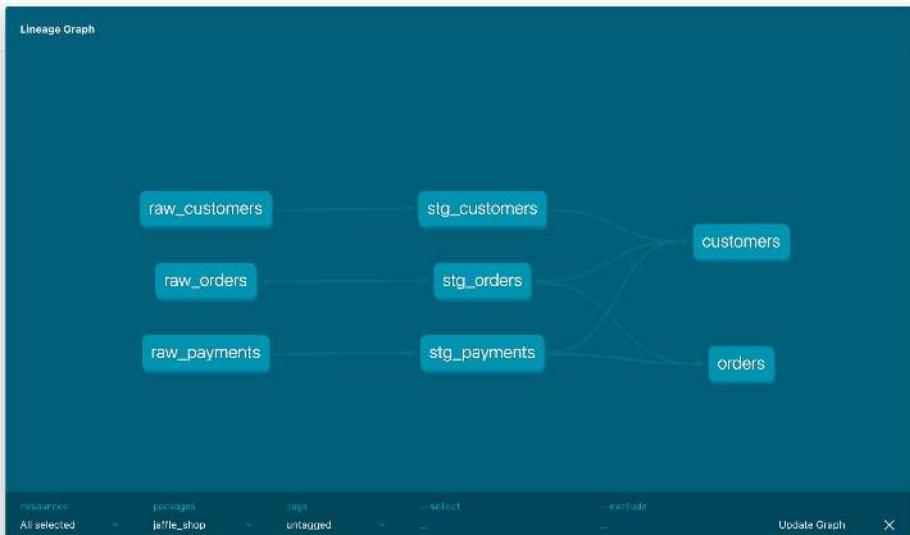
- Schedule jobs in production
- Run scalable extract and load tasks
- Have centralised logging or alerts handling
- Make the project's documentation available to other team members

dbt Labs, creator and maintainer of dbt Core, solves these issues via their proprietary and commercial platform, dbt Cloud.

## 2. Why Cosmos

# Why Cosmos?

Cosmos "magically" translates dbt pipelines in Airflow DAGs



```
$ pip install astronomer-cosmos
```

ASTRONOMER

# Why Cosmos?

```
import os
from datetime import datetime
from pathlib import Path
from cosmos import DbtDag, ProjectConfig, ProfileConfig
from cosmos.profiles import PostgresUserPasswordProfileMapping

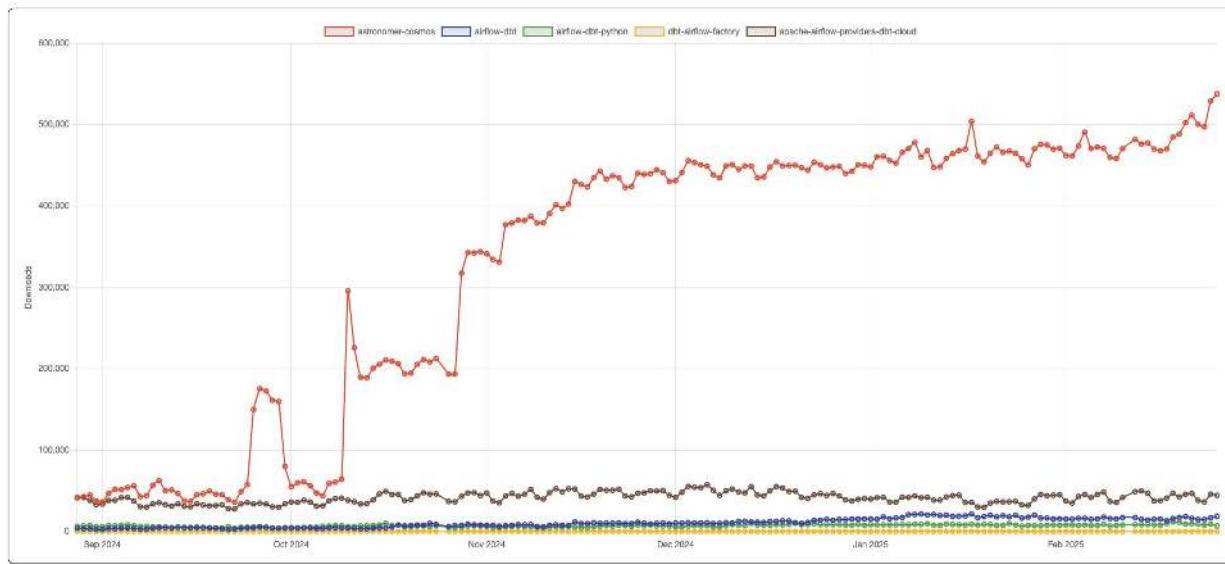
DEFAULT_DBT_ROOT_PATH = Path(__file__).parent / "dbt"
DBT_ROOT_PATH = Path(os.getenv("DBT_ROOT_PATH", DEFAULT_DBT_ROOT_PATH))

profile_config = ProfileConfig(
    profile_name="jaffle_shop",
    target_name="dev",
    profile_mapping=PostgresUserPasswordProfileMapping(
        conn_id="airflow_db",
        profile_args={"schema": "public"},
    ),
)
basic_cosmos_dag = DbtDag(
    project_config=ProjectConfig(
        DBT_ROOT_PATH / "jaffle_shop",
    ),
    profile_config=profile_config,
    schedule="@daily",
    start_date=datetime(2023, 1, 1),
    catchup=False,
    dag_id="basic_cosmos_dag",
)
```



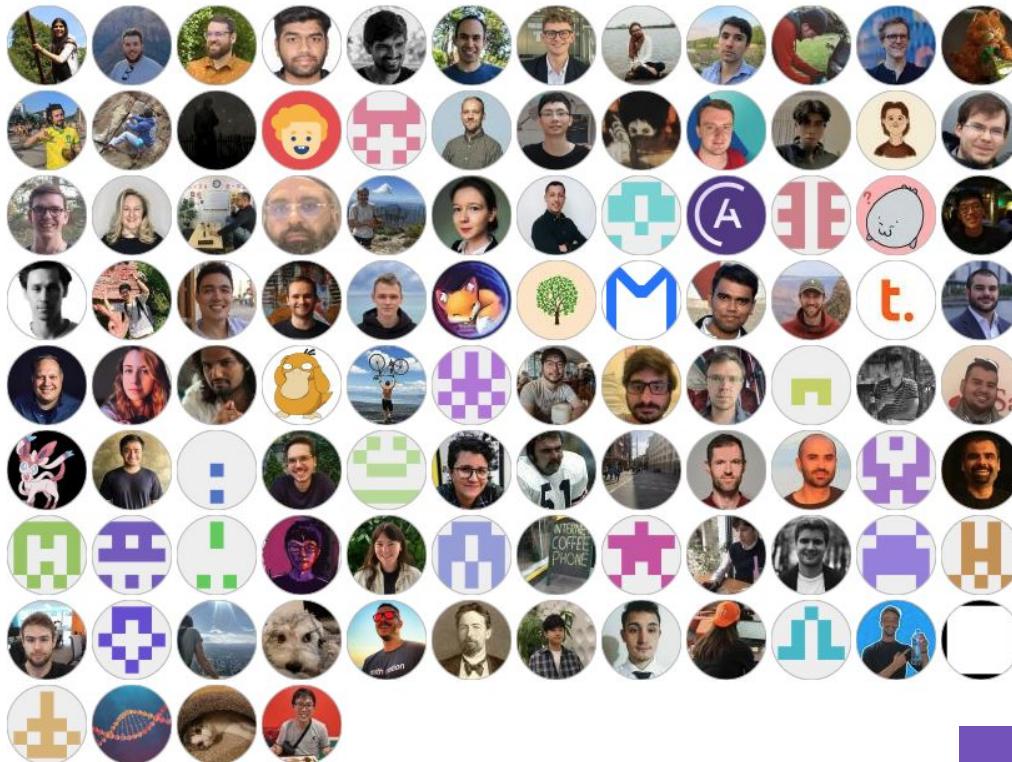
# Cosmos adoption

- 20M+ downloads in PyPI per month (August-September 2025)
- 1k stars in Github
- Millions of Cosmos tasks are run every month (>20% Astro customers)



# The driving force behind Cosmos

<https://github.com/astronomer/astronomer-cosmos/>



## 144 contributors

### **3. Performance challenge**

# Performance has been a priority

	<b>1.2.5</b>	<b>1.3 (DBT LS FILE)</b>	<b>1.4</b>	<b>1.5</b>	<b>1.6</b>
DAG Parsing time	00:00:08	00:00:02	00:00:07	00:00:02	00:00:02
Task Run time	00:00:09	00:00:08	00:00:06	00:00:05	00:00:04
Task Queue time	00:00:09	00:00:04	00:00:05	00:00:01	00:00:01
DAG Run time	00:01:29	00:00:55	00:01:18	00:00:43	00:00:42

Slides <https://airflowsummit.org/slides/2024/40-Overcoming-Performance-Hurdles-in-Integrating-dbt-with-Airflow.pdf>  
Recording: <https://www.youtube.com/watch?v=qnJPFGvqLzU>

# The cost of running dbt commands

The following approaches accomplish the same outcome.

**Which one is faster?**

# cmd: 1

```
$ dbt build
```

# cmd: 3

```
$ dbt seed  
$ dbt run  
$ dbt test
```

#cmd: 13

```
$ dbt seed --select raw_customers  
$ dbt seed --select raw_orders  
$ dbt seed --select raw_payments  
  
$ dbt run --select stg_customers  
$ dbt run --select stg_orders  
$ dbt run --select stg_payments  
$ dbt run --select customers  
$ dbt run --select orders  
  
$ dbt test --select stg_customers  
$ dbt test --select stg_orders  
$ dbt test --select stg_payments  
$ dbt test --select customers  
$ dbt test --select orders
```

# The cost of running dbt

These are the **total times** to run **the same example pipeline** (Jaffle Shop) with **dbt Core 1.10.1** using **Snowflake** in the three ways described previously (using M1 Pro)

# cmd: 1

```
$ dbt build
```

17s

# cmd: 3

```
$ dbt seed  
$ dbt run  
$ dbt test
```

34s

#cmd: 13

```
$ dbt seed --select raw_customers  
$ dbt seed --select raw_orders  
$ dbt seed --select raw_payments  
  
$ dbt run --select stg_customers  
$ dbt run --select stg_orders  
$ dbt run --select stg_payments  
$ dbt run --select customers  
$ dbt run --select orders  
  
$ dbt test --select stg_customers  
$ dbt test --select stg_orders  
$ dbt test --select stg_payments  
$ dbt test --select customers  
$ dbt test --select orders
```

61s

ASTRONOMER

# Reducing task execution time with Cosmos 1.10

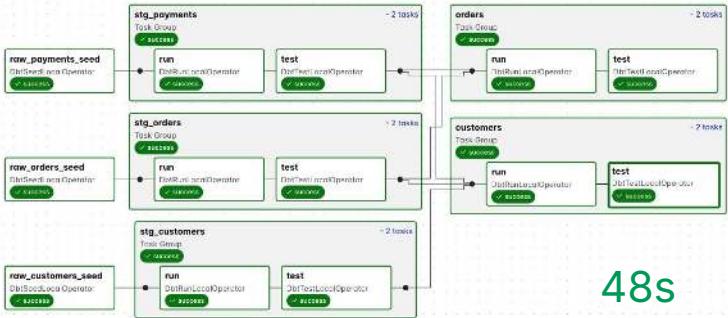
```
from datetime import datetime
from cosmos import DbtDag, ProjectConfig, ProfileConfig
from include.constants import jaffle_shop_path

project_config = ProjectConfig(
    dbt_project_path=jaffle_shop_path,
)

profile_config = ProfileConfig(
    profile_name="jaffle_shop",
    target_name="dev",
    profiles_yml_filepath=jaffle_shop_path / "profiles.yml",
)

snowflake_dag = DbtDag(
    project_config=project_config,
    profile_config=profile_config,
    start_date=datetime(2023, 1, 1),
    dag_id="snowflake_dag",
    tags=["profiles"],
)
```

Standard behaviour of running  
**one task per dbt node** in  
Cosmos



48s

# Reducing task execution time with Cosmos 1.10

```
from airflow.sdk import DAG
from cosmos import DbtRunLocalOperator, DbtSeedLocalOperator, DbtTestLocalOperator,
ProfileConfig

from include.constants import jaffle_shop_path


with DAG("snowflake_dag_per_resource") as dag:
    seed = DbtSeedLocalOperator(
        task_id="seed",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )

    run = DbtRunLocalOperator(
        task_id="run",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )

    test = DbtTestLocalOperator(
        task_id="test",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )

    seed >> run >> test
```

Running **one task per dbt node type** with Cosmos



27s

# Reducing task execution time with Cosmos 1.10

```
from airflow.sdk import DAG

from cosmos import DbtBuildLocalOperator, ProfileConfig

from include.constants import jaffle_shop_path

profile_config = ProfileConfig(
    profile_name="jaffle_shop",
    target_name="dev",
    profiles_yml_filepath=jaffle_shop_path / "profiles.yml",
)

with DAG("snowflake_dag_single_task") as dag:
    build = DbtBuildLocalOperator(
        task_id="build",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )
    build
```

Running **one task** for the whole dbt workflow with Cosmos

**build**  
DbtBuildLocalOperator  
✓ success

20s

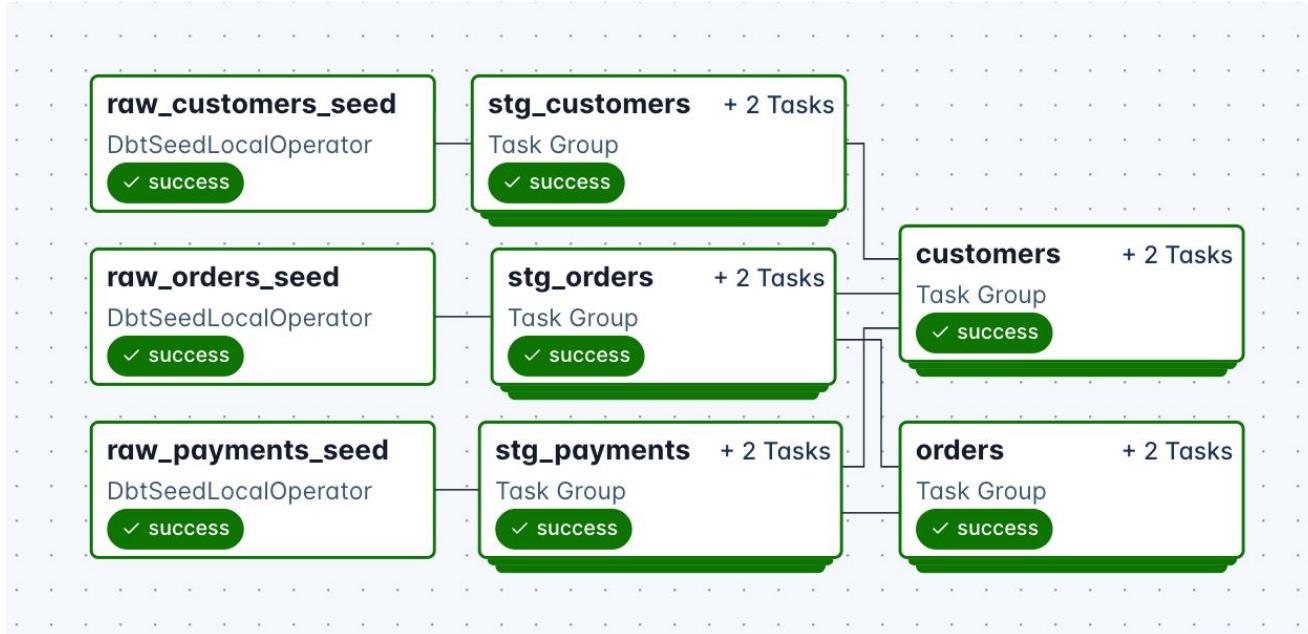
# Cosmos & dbt performance

	1 command	3 commands	13 commands
dbt Core	17s	24s	61s
dbt Fusion	15s	26s	74s
Cosmos with dbt Core (*)	20s	27s	48s

(\*) Using the `$ airflow dags test` command with Cosmos 1.11.0a1, Airflow 3.0.2, dbt Core 1.10 (installed in the same Python virtualenv as Airflow and Cosmos), against Snowflake. In practice, when using a production Airflow deployment such as Astro, the latency will be higher, due to their distributed nature.

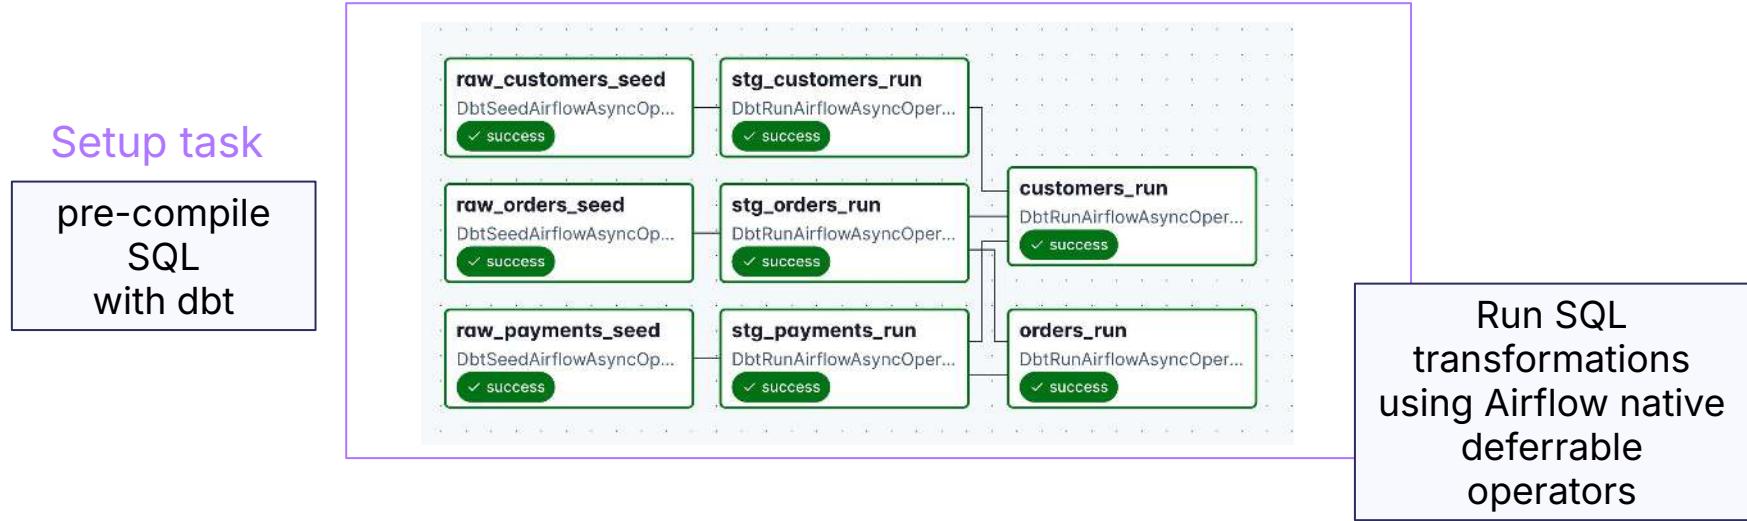
# Challenge

Is it possible to give users a **fine-grained visualisation** and **retry per model** capabilities without running a **dbt** command for every seed or model?



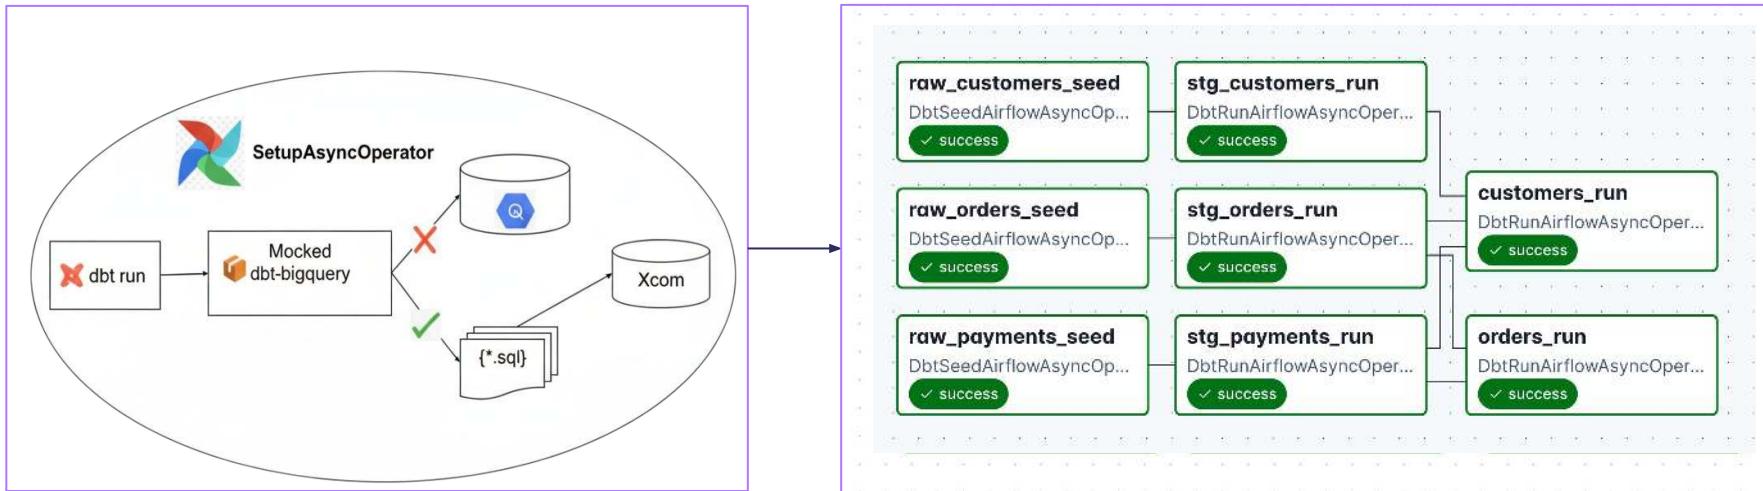
## **4. Strategy 1: Airflow Deferrable**

# Strategy 1 ExecutionMode.AIRFLOW\_ASYNC

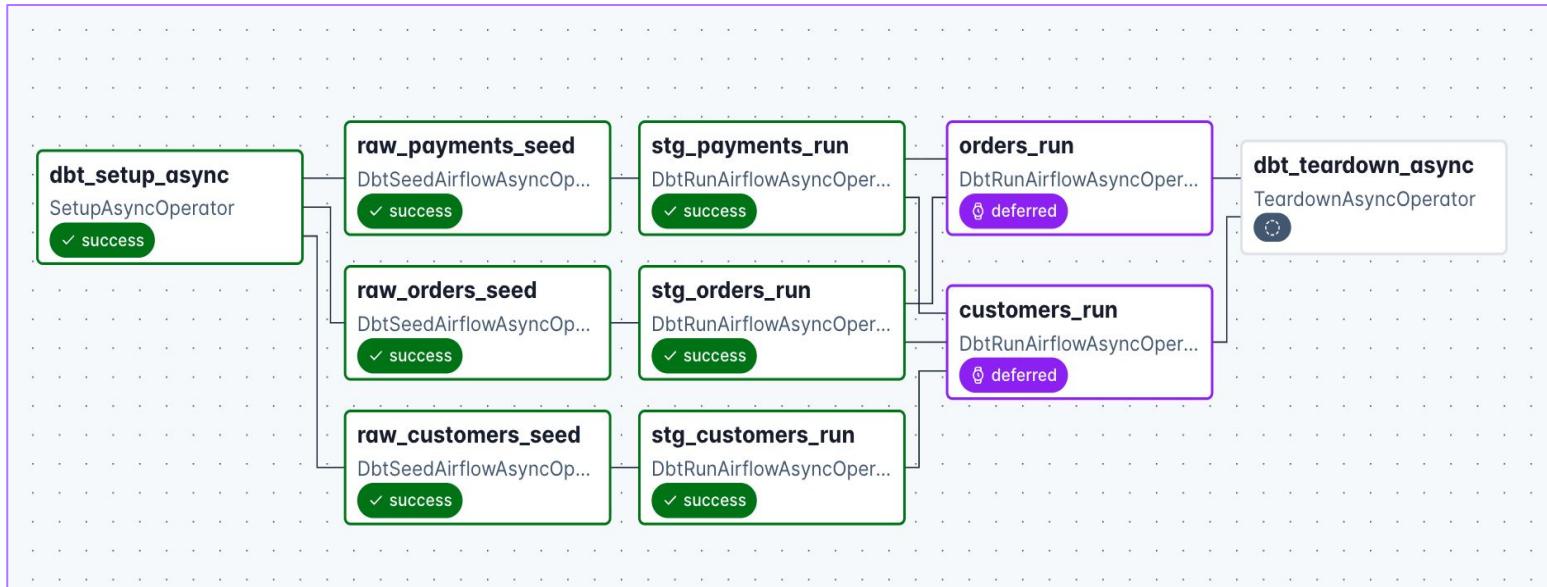


[https://astronomer.github.io/astronomer-cosmos/getting\\_started/async-execution-mode.html](https://astronomer.github.io/astronomer-cosmos/getting_started/async-execution-mode.html)

# Strategy 1 ExecutionMode.AIRFLOW\_ASYNC



# Strategy 1 ExecutionMode.AIRFLOW\_ASYNC



# Strategy 1 ExecutionMode.AIRFLOW\_ASYNC

## Performance improvements

Execution mode	Time to run project (minutes)
dbt run	13
Cosmos 1.9 (*) with <b>ExecutionMode.LOCAL</b> (Airflow with a local astro-cli setup)	11
Cosmos 1.9 (*) with <b>ExecutionMode.AIRFLOW_ASYNC</b> (Airflow with a local astro-cli setup)	11
Cosmos 1.11a6 (**) with <b>ExecutionMode.AIRFLOW_ASYNC</b> (Airflow with a local astro-cli setup)	7

(\*) Using Cosmos 1.9 with Airflow 3.0.2, dbt Core 1.10 (installed in the same Python virtualenv as Airflow and Cosmos), against Postgres. We are using a dbt project that has [129 models](#).

(\*\*) Cosmos 1.11a6 has two main improvements regarding the AIRFLOW\_ASYNC, originally released in Cosmos 1.9:

- Use of XCom instead of Remote object store to exchange compiled SQL files [#1934](#)
- Reuse Python virtualenv across tasks running in the same worker node [#1939](#)

# Strategy 1 ExecutionMode.AIRFLOW\_ASYNC

## Pros

- Reduced benchmark Airflow DAG run time by 36%
- Single dbt command invocation (less CPU/memory allocation)
- Non-blocking transformations in the data warehouse with Airflow deferrable operators

## Cons

- Currently only supports dbt models
- Currently only supports BigQuery
- Implementation specific per data warehouse
- The dbt project cannot have models with a metadata-dependency on other models
- Some users reported issues related to Airflow deferrable operators, which can be challenging to reproduce
- Only works with dbt Core (not dbt Fusion)

[https://astronomer.github.io/astronomer-cosmos/getting\\_started/async-execution-mode.html](https://astronomer.github.io/astronomer-cosmos/getting_started/async-execution-mode.html)

# Next steps

ExecutionMode.AIRFLOW\_ASYNC

1. Available since Cosmos 1.9 for BigQuery
2. Significant performance improvements in Cosmos 1.11 pre-releases
3. We need feedback!
4. Maybe add non-async version
5. We've introduced telemetry to evaluate the adoption
6. We'd love [contributions](#)

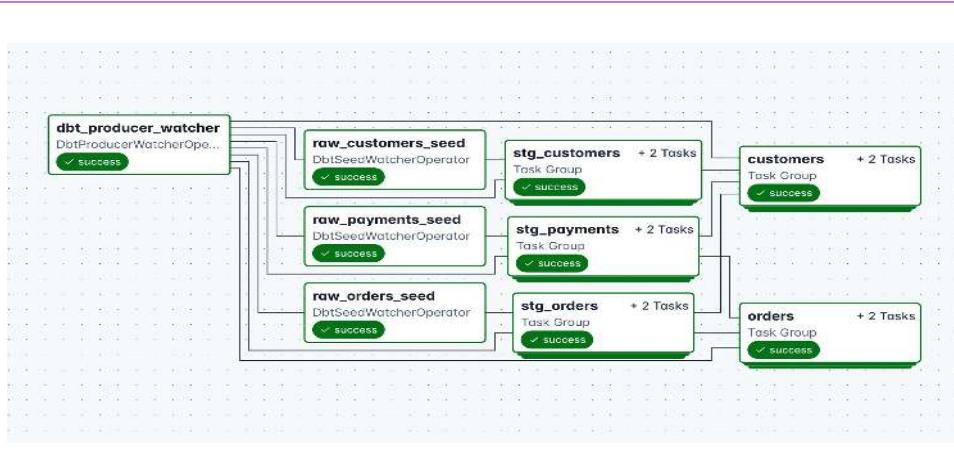
is:issue state:open label:execution:async

## 5. Strategy 2: Watcher

# Strategy 2 ExecutionMode.WATCHER

Setup task

run dbt build  
in a single  
Airflow task



all the other  
tasks are  
sensors

<https://github.com/astronomer/astronomer-cosmos/issues/1950>

# Strategy 2 ExecutionMode.WATCHER

## Registering callbacks

Register `callbacks` on dbt's `EventManager`, to access structured events and enable custom logging. The current behavior of callbacks is to block subsequent steps from proceeding; this functionality is not guaranteed in future versions.

```
from dbt.cli.main import dbtRunner
from dbt_common.events.base_types import EventMsg

def print_version_callback(event: EventMsg):
    if event.info.name == "MainReportVersion":
        print(f"We are thrilled to be running dbt{event.data.version}")

dbt = dbtRunner(callbacks=[print_version_callback])
dbt.invoke(["list"])
```

<https://docs.getdbt.com/reference/programmatic-invocations#registering-callbacks>

# Strategy 2 ExecutionMode.WATCHER

## Performance improvements

Execution mode	Number of threads	Time to run project (minutes)
dbt build in the CLI	4	6 - 7
dbt run for each model individually		30
Cosmos default <b>ExecutionMode.LOCAL</b> in Astro CLI locally		10 - 15
Cosmos proposed <b>ExecutionMode.WATCHER</b> in Astro CLI locally	1	26
	2	14
	4	7
	8	4
	16	2
The <b>ExecutionMode.WATCHER</b> in Airflow with an Astro deployment ( <b>A10</b> )	8	5

# Strategy 2 ExecutionMode.WATCHER

## Pros

- Reduced DAG run time to 1/5th of the original time
- Single dbt run
  - generates unified `run_results.json`
  - support dbt pre-hook & post-hook
- Data warehouse-agnostic implementation

## Cons/Current Limitations

- Airflow worker is blocked by transformations happening in the data warehouse
- Retries have the same performance as `ExecutionMode.LOCAL`
- Currently relies on dbt and Airflow being installed in the same Python venv (some users report conflicts between dependencies)
- Unclear how these features should work: Cosmos callback, Airflow datasets/assets and OpenLineage events

<https://github.com/astro/astro-cosmos/issues/1950>

# Next steps

ExecutionMode.WATCHER

1. We successfully ran a PoC during August 2025
2. First release estimate: end of October 2025 (available in **1.11.0a6**)
3. Work on making sensors deferrable
4. We need feedback!
5. We'd love [contributions](#)

is:issue state:open label:execution:watcher

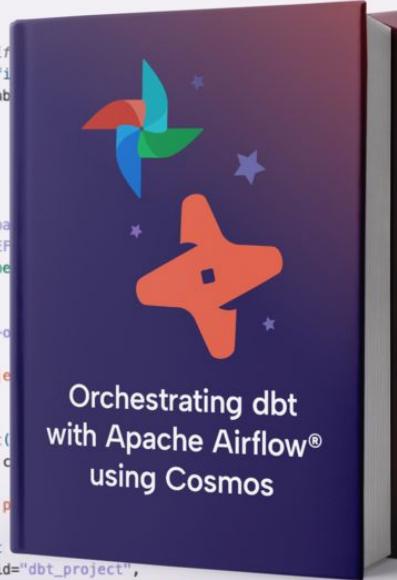
## 6. Takeaways

# Takeaways

1. To run the **same dbt pipeline with multiple dbt command** is slow
2. To use **Airflow deferrable operators** allows to not **wait for the transformation** in the **data warehouse**, which can save 36% DAG runtime
3. It is not always possible to **pre-compile a dbt project**
4. The **watcher approach** reduces the DAG runtime up to 80% and it is agnostic to **data-warehouse**
5. We need feedback and help!



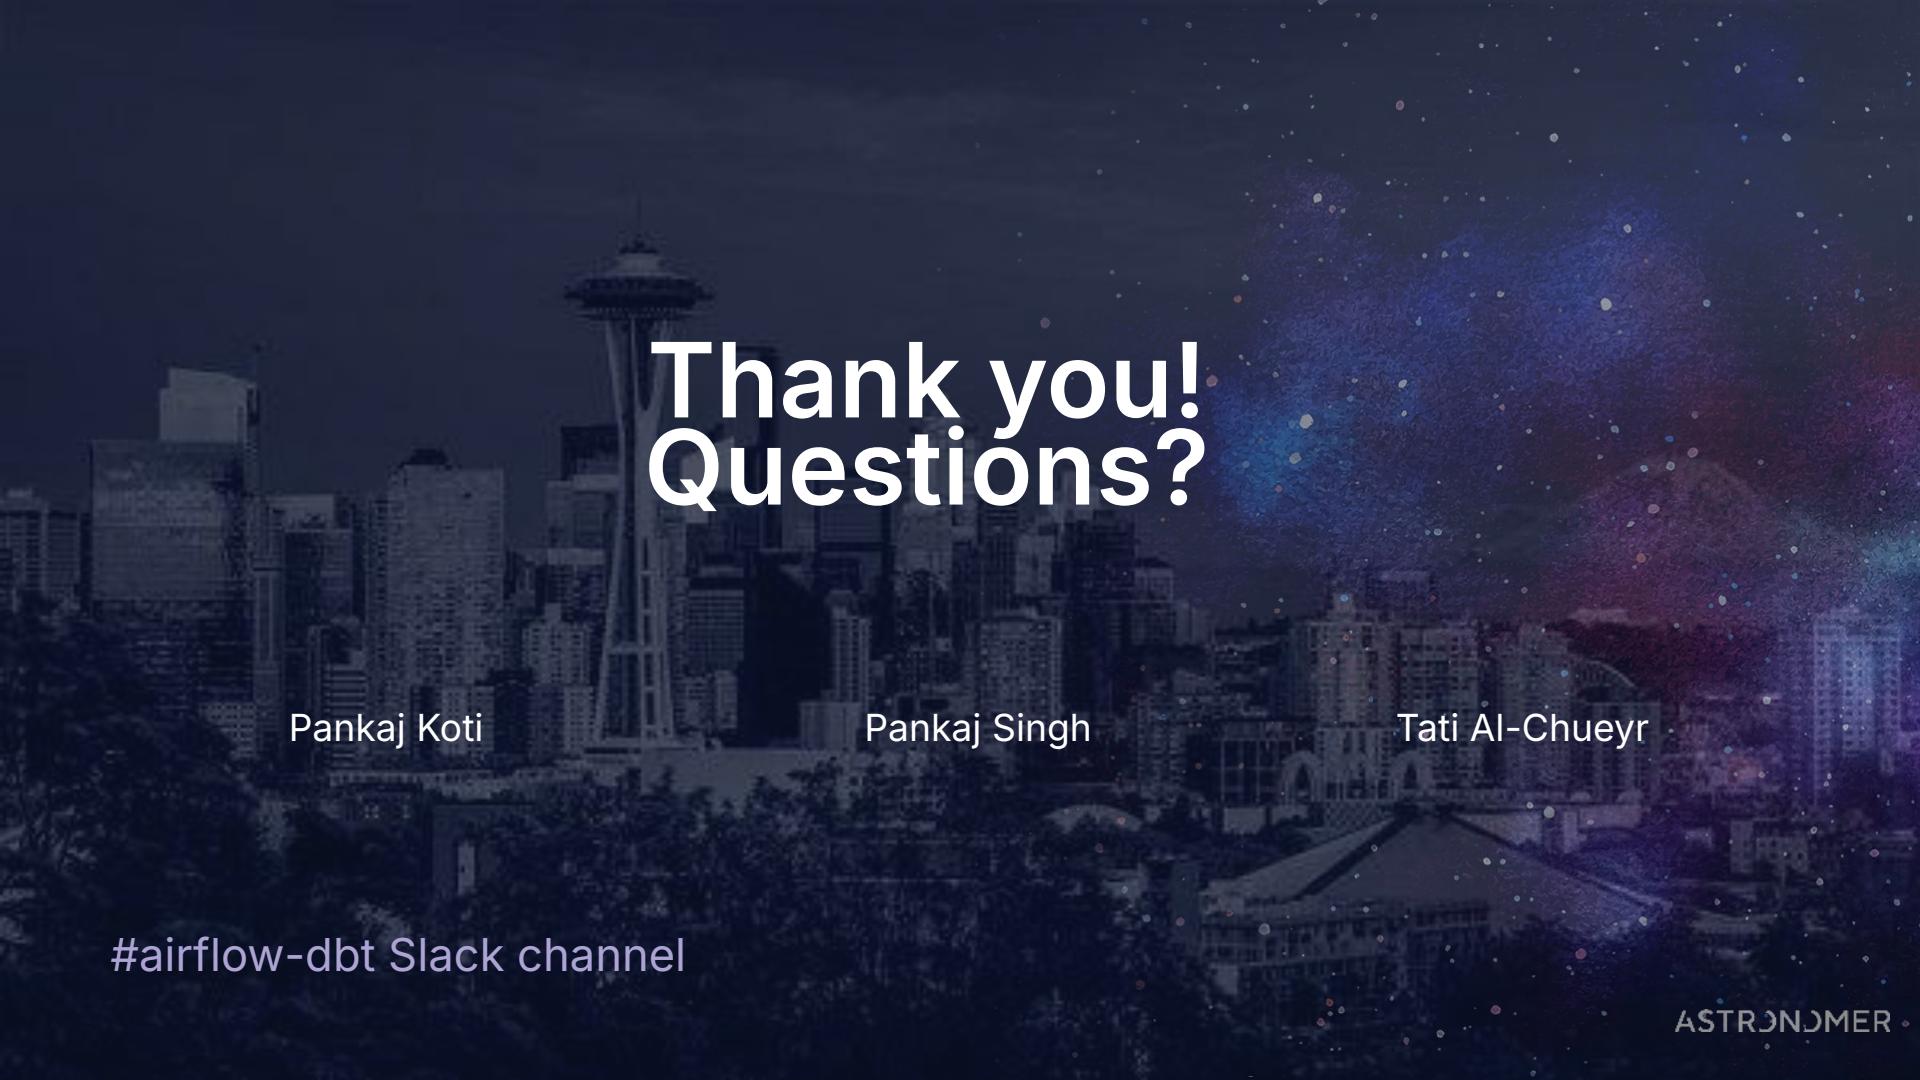
# Learn more about how to run dbt with Apache Airflow and Cosmos



# The 2025 Apache Airflow® Survey is here!

Fill it out to for a free Airflow 3  
Fundamentals or DAG Authoring in  
Airflow 3 certification code





# Thank you! Questions?

Pankaj Koti

Pankaj Singh

Tati Al-Chueyr

#airflow-dbt Slack channel

ASTRONOMER