

Boosting dbt Core workflows performance Airflow Deferrable Capabilities

Pankaj Koti
Senior Software Engineer
@ Astronomer

Pankaj Singh
Senior Software Engineer
@ Astronomer

Tati Al-Chueyr
Principal Software Engineer
@ Astronomer

Airflow Summit, Seattle, USA - 8 October 2025

ASTRONOMER

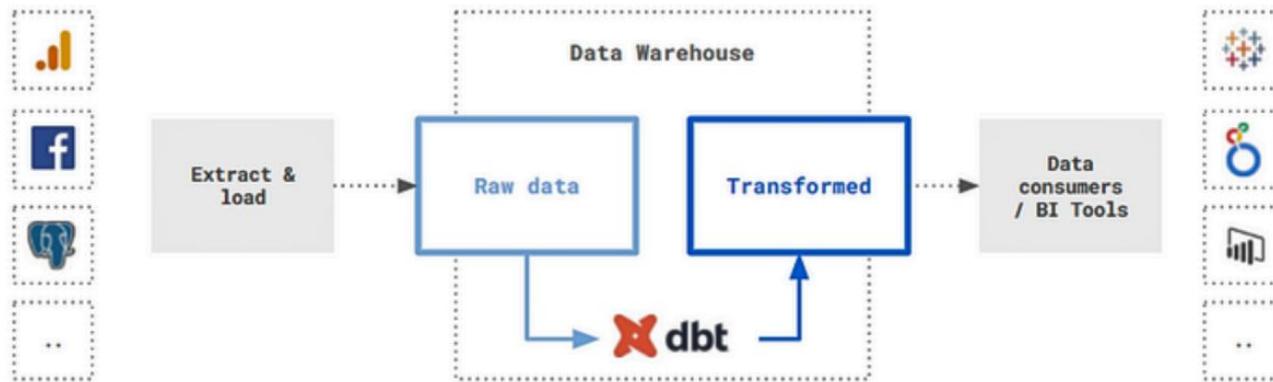
Agenda

1. What is dbt
2. Why Cosmos
3. Performance challenge
4. Strategy 1: Airflow Deferrable
5. Strategy 2: Watcher
6. Takeaways

1. What is dbt

dbt Core what is it?

dbt Core is an open-source tool that empowers **data practitioners** to transform



```
$ pip install dbt-core
```

dbt Core project structure & syntax

```
➜ jaffle_shop git:(af-31) ✘ tree
.
├── dbt_project.yml
├── LICENSE
├── macros
│   ├── drop_table.sql
│   └── generate_alias_name.sql
└── models
    ├── customers.sql
    ├── docs.md
    ├── orders.sql
    ├── overview.md
    ├── schema.yml
    └── staging
        ├── schema.yml
        ├── stg_customers.sql
        ├── stg_orders.sql
        └── stg_payments.sql
├── packages.yml
├── profiles.yml
└── README.md
└── seeds
    ├── raw_customers.csv
    ├── raw_orders.csv
    └── raw_payments.csv

5 directories, 19 files
```

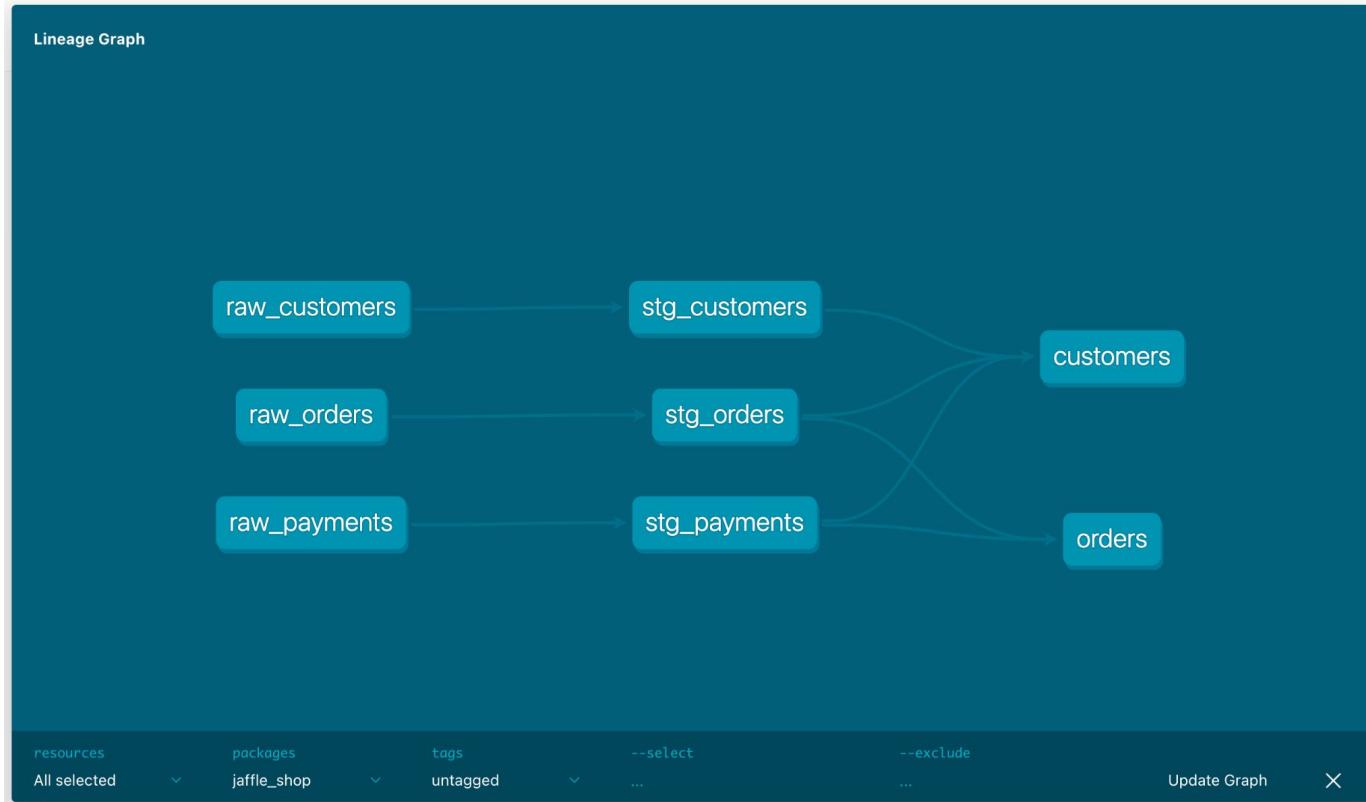
```
➜ jaffle_shop git:(af-31) ✘ cat models/staging/stg_payments.sql
with source as (
    #-
    -- Normally we would select from the table here, but we are using seeds to load
    -- our data in this project
    #
    select * from {{ ref('raw_payments') }}
),

renamed as (

    select
        id as payment_id,
        order_id,
        payment_method,
        -- `amount` is currently stored in cents, so we convert it to dollars
        amount / 100 as amount
    from source
)

select * from renamed
```

dbt Core pipeline visualisation



dbt Core limitations

dbt Core does not...

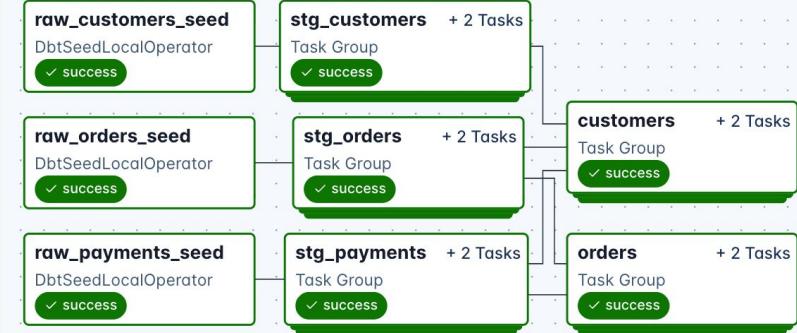
- Schedule jobs in production
- Run scalable extract and load tasks
- Have centralised logging or alerts handling
- Make the project's documentation available to other team members

dbt Labs, creator and maintainer of dbt Core, solves these issues via their proprietary and commercial platform, dbt Cloud.

2. Why Cosmos

Why Cosmos?

Cosmos "magically" translates dbt pipelines in Airflow DAGs



```
$ pip install astronomer-cosmos
```

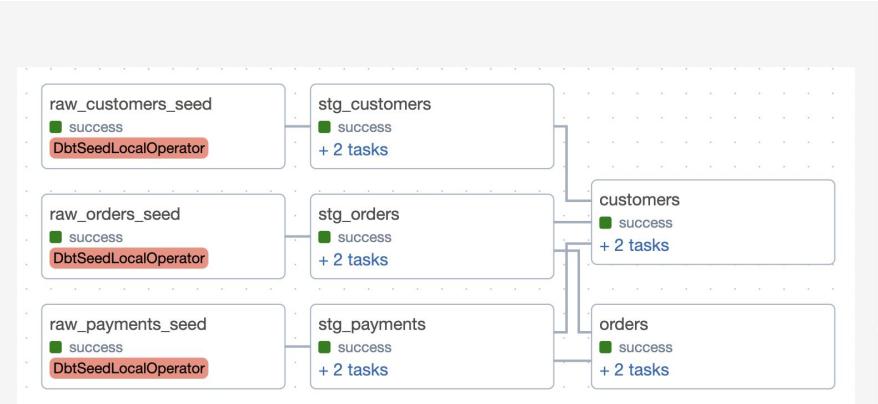
ASTRONOMER

Why Cosmos?

```
import os
from datetime import datetime
from pathlib import Path
from cosmos import DbtDag, ProjectConfig, ProfileConfig
from cosmos.profiles import PostgresUserPasswordProfileMapping

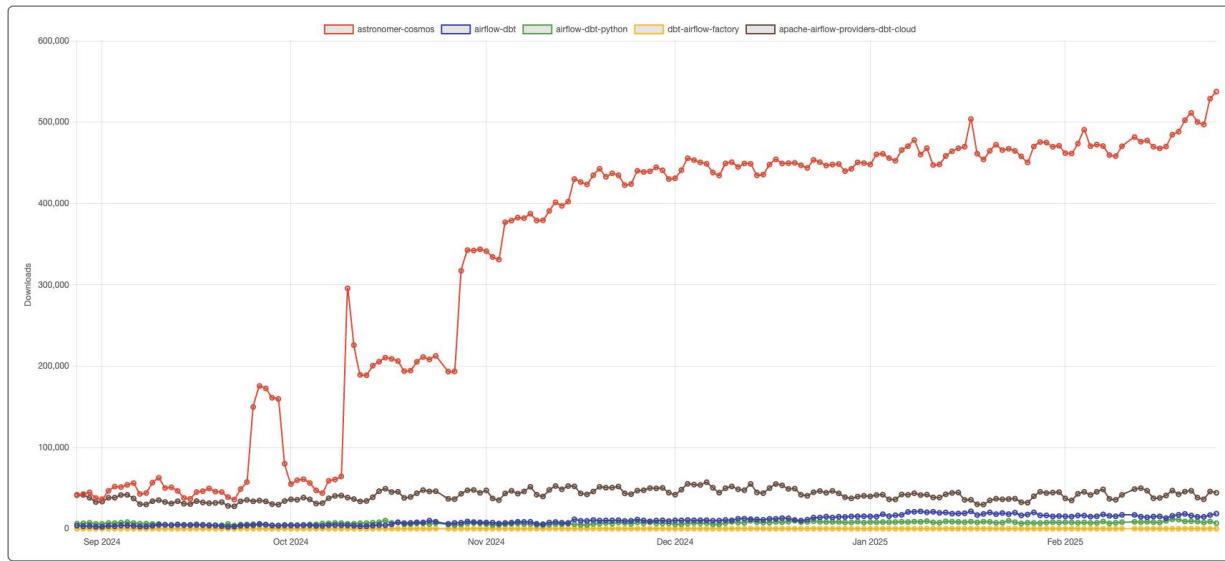
DEFAULT_DBT_ROOT_PATH = Path(__file__).parent / "dbt"
DBT_ROOT_PATH = Path(os.getenv("DBT_ROOT_PATH", DEFAULT_DBT_ROOT_PATH))

profile_config = ProfileConfig(
    profile_name="jaffle_shop",
    target_name="dev",
    profile_mapping=PostgresUserPasswordProfileMapping(
        conn_id="airflow_db",
        profile_args={"schema": "public"},
    ),
)
basic_cosmos_dag = DbtDag(
    project_config=ProjectConfig(
        DBT_ROOT_PATH / "jaffle_shop",
    ),
    profile_config=profile_config,
    schedule="@daily",
    start_date=datetime(2023, 1, 1),
    catchup=False,
    dag_id="basic_cosmos_dag",
)
```



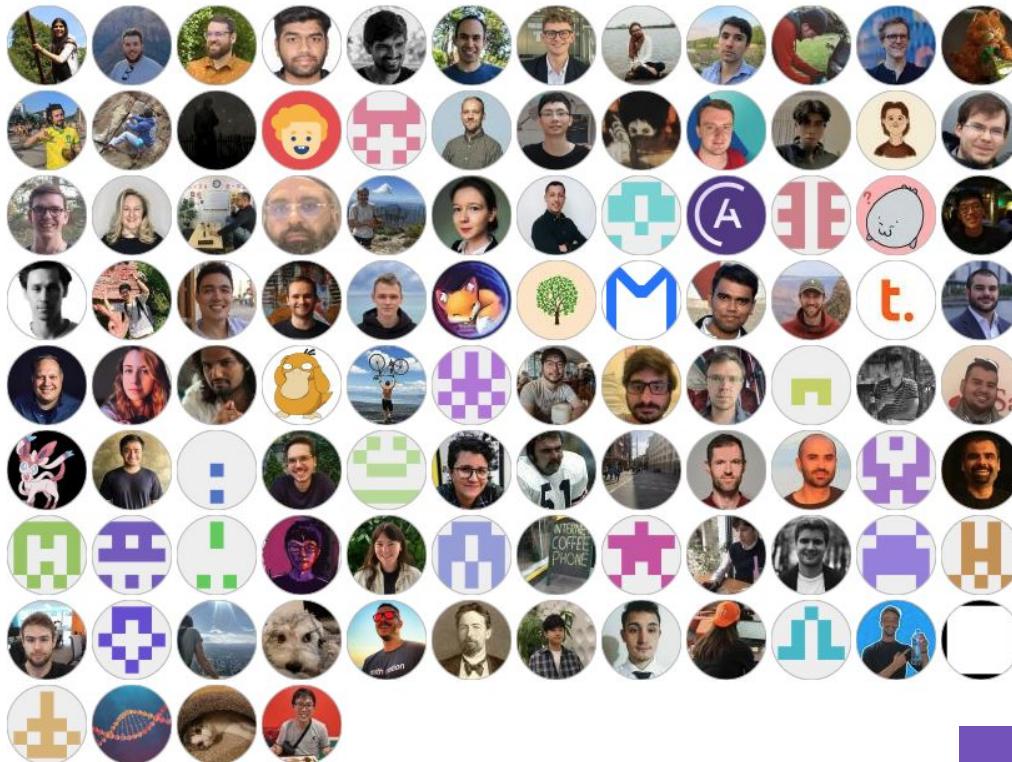
Cosmos adoption

- 20M+ downloads in PyPI per month (August-September 2025)
- 1k stars in Github
- Millions of Cosmos tasks are run every month (>20% Astro customers)



The driving force behind Cosmos

<https://github.com/astronomer/astronomer-cosmos/>



144 contributors

3. Performance challenge

Performance has been a priority

| | 1.2.5 | 1.3 (DBT LS FILE) | 1.4 | 1.5 | 1.6 |
|------------------|--------------|--------------------------|------------|------------|------------|
| DAG Parsing time | 00:00:08 | 00:00:02 | 00:00:07 | 00:00:02 | 00:00:02 |
| Task Run time | 00:00:09 | 00:00:08 | 00:00:06 | 00:00:05 | 00:00:04 |
| Task Queue time | 00:00:09 | 00:00:04 | 00:00:05 | 00:00:01 | 00:00:01 |
| DAG Run time | 00:01:29 | 00:00:55 | 00:01:18 | 00:00:43 | 00:00:42 |

Slides <https://airflowsummit.org/slides/2024/40-Overcoming-Performance-Hurdles-in-Integrating-dbt-with-Airflow.pdf>
Recording: <https://www.youtube.com/watch?v=qnJPFGvqLzU>

The cost of running dbt commands

The following approaches accomplish the same outcome.

Which one is faster?

cmd: 1

```
$ dbt build
```

cmd: 3

```
$ dbt seed  
$ dbt run  
$ dbt test
```

#cmd: 13

```
$ dbt seed --select raw_customers  
$ dbt seed --select raw_orders  
$ dbt seed --select raw_payments  
  
$ dbt run --select stg_customers  
$ dbt run --select stg_orders  
$ dbt run --select stg_payments  
$ dbt run --select customers  
$ dbt run --select orders  
  
$ dbt test --select stg_customers  
$ dbt test --select stg_orders  
$ dbt test --select stg_payments  
$ dbt test --select customers  
$ dbt test --select orders
```

The cost of running dbt

These are the **total times** to run **the same example pipeline** (Jaffle Shop) with **dbt Core 1.10.1** using **Snowflake** in the three ways described previously (using M1 Pro)

cmd: 1

```
$ dbt build
```

17s

cmd: 3

```
$ dbt seed  
$ dbt run  
$ dbt test
```

34s

#cmd: 13

```
$ dbt seed --select raw_customers  
$ dbt seed --select raw_orders  
$ dbt seed --select raw_payments  
  
$ dbt run --select stg_customers  
$ dbt run --select stg_orders  
$ dbt run --select stg_payments  
$ dbt run --select customers  
$ dbt run --select orders  
  
$ dbt test --select stg_customers  
$ dbt test --select stg_orders  
$ dbt test --select stg_payments  
$ dbt test --select customers  
$ dbt test --select orders
```

61s

ASTRONOMER

Reducing task execution time with Cosmos 1.10

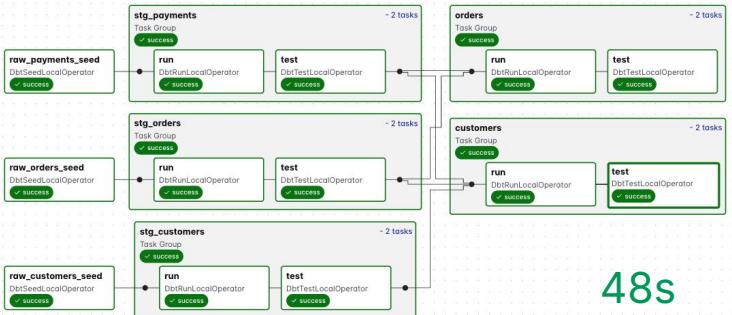
```
from datetime import datetime
from cosmos import DbtDag, ProjectConfig, ProfileConfig
from include.constants import jaffle_shop_path

project_config = ProjectConfig(
    dbt_project_path=jaffle_shop_path,
)

profile_config = ProfileConfig(
    profile_name="jaffle_shop",
    target_name="dev",
    profiles_yml_filepath=jaffle_shop_path / "profiles.yml",
)

snowflake_dag = DbtDag(
    project_config=project_config,
    profile_config=profile_config,
    start_date=datetime(2023, 1, 1),
    dag_id="snowflake_dag",
    tags=["profiles"],
)
```

Standard behaviour of running
one task per dbt node in
Cosmos



48s

Reducing task execution time with Cosmos 1.10

```
from airflow.sdk import DAG
from cosmos import DbtRunLocalOperator, DbtSeedLocalOperator, DbtTestLocalOperator,
ProfileConfig

from include.constants import jaffle_shop_path


with DAG("snowflake_dag_per_resource") as dag:
    seed = DbtSeedLocalOperator(
        task_id="seed",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )

    run = DbtRunLocalOperator(
        task_id="run",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )

    test = DbtTestLocalOperator(
        task_id="test",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )

    seed >> run >> test
```

Running **one task per dbt node type** with Cosmos



27s

Reducing task execution time with Cosmos 1.10

```
from airflow.sdk import DAG

from cosmos import DbtBuildLocalOperator, ProfileConfig

from include.constants import jaffle_shop_path

profile_config = ProfileConfig(
    profile_name="jaffle_shop",
    target_name="dev",
    profiles_yml_filepath=jaffle_shop_path / "profiles.yml",
)

with DAG("snowflake_dag_single_task") as dag:
    build = DbtBuildLocalOperator(
        task_id="build",
        profile_config=profile_config,
        project_dir=jaffle_shop_path,
    )
    build
```

Running **one task** for the whole dbt workflow with Cosmos

build
DbtBuildLocalOperator
✓ success

20s

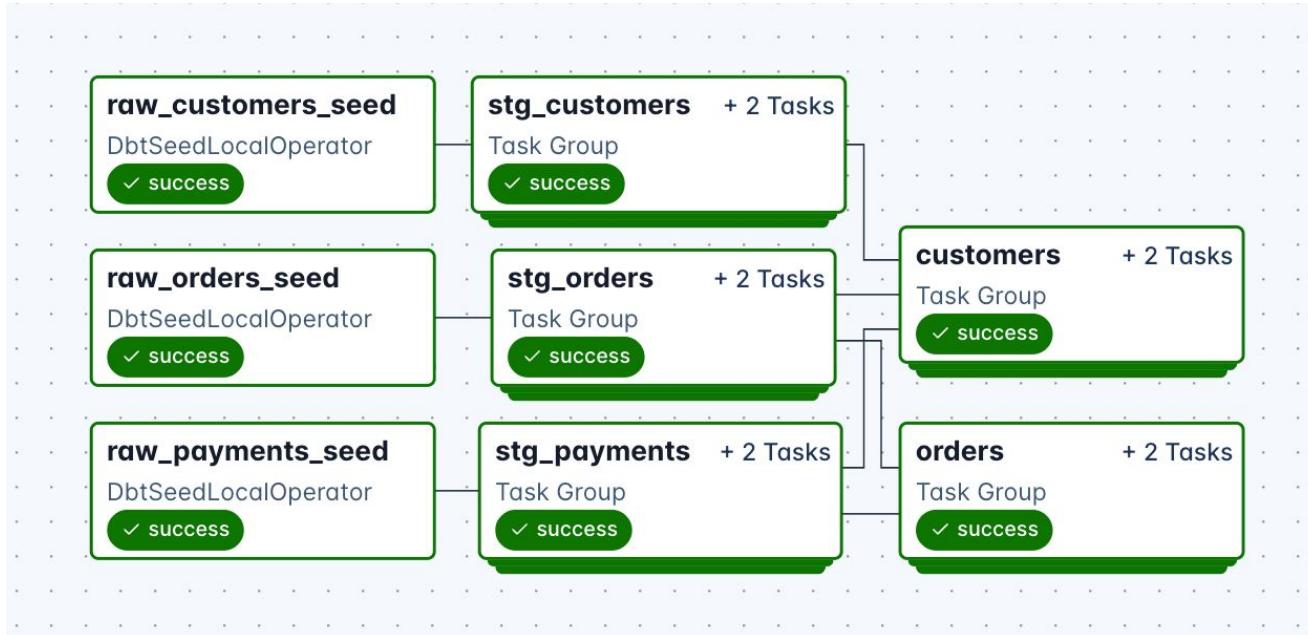
Cosmos & dbt performance

| | 1 command | 3 commands | 13 commands |
|--------------------------|-----------|------------|-------------|
| dbt Core | 17s | 24s | 61s |
| dbt Fusion | 15s | 26s | 74s |
| Cosmos with dbt Core (*) | 20s | 27s | 48s |

(*) Using the `$ airflow dags test` command with Cosmos 1.11.0a1, Airflow 3.0.2, dbt Core 1.10 (installed in the same Python virtualenv as Airflow and Cosmos), against Snowflake. In practice, when using a production Airflow deployment such as Astro, the latency will be higher, due to their distributed nature.

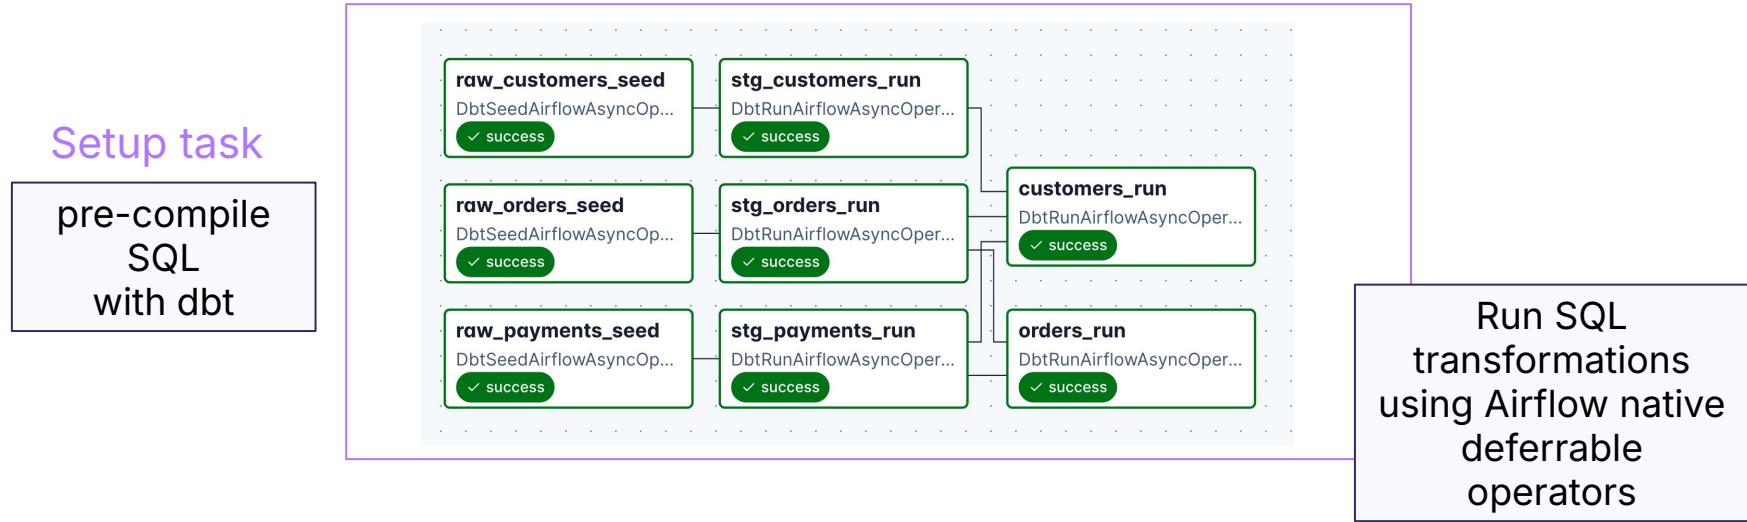
Challenge

Is it possible to give users a **fine-grained visualisation** and **retry per model** capabilities without running a **dbt** command for every seed or model?



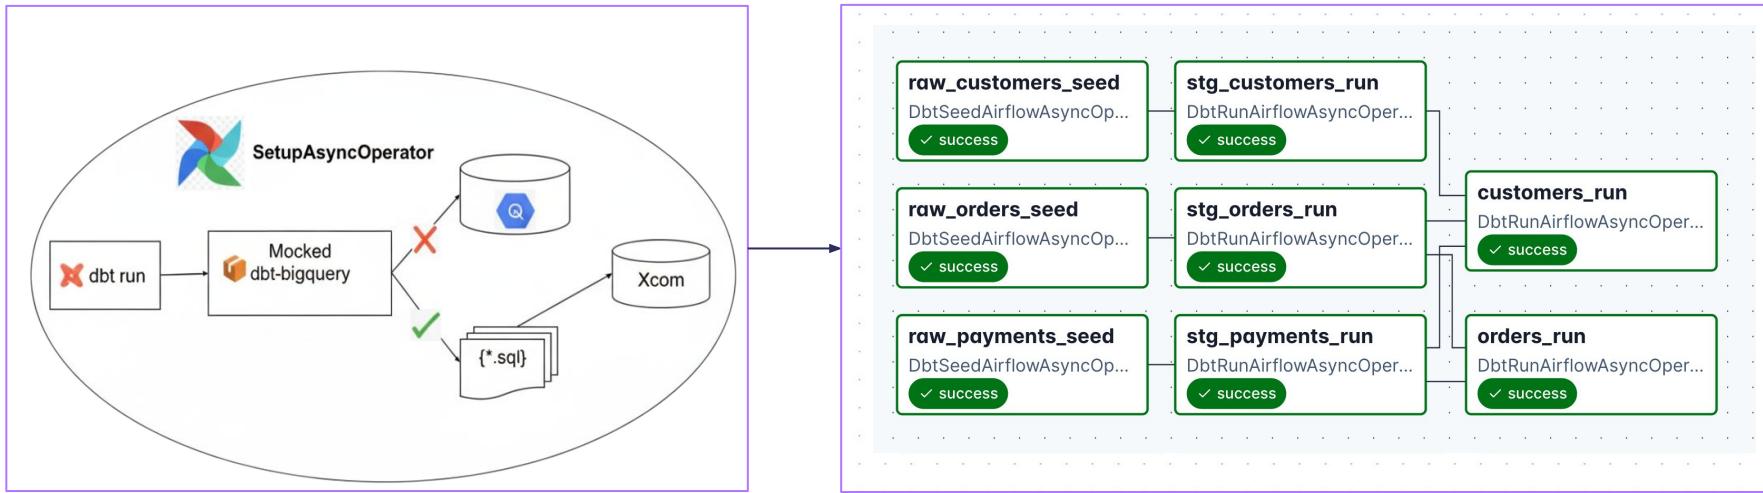
4. Strategy 1: Airflow Deferrable

Strategy 1 ExecutionMode.AIRFLOW_ASYNC

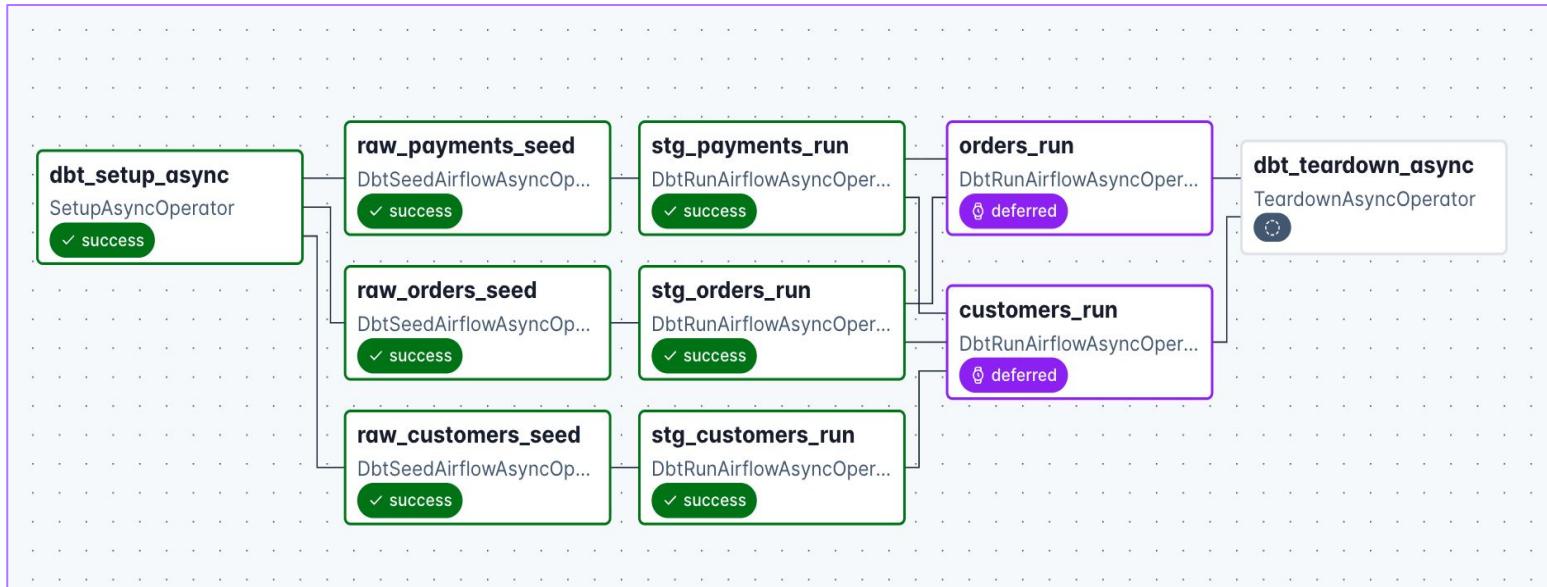


https://astronomer.github.io/astronomer-cosmos/getting_started/async-execution-mode.html

Strategy 1 ExecutionMode.AIRFLOW_ASYNC



Strategy 1 ExecutionMode.AIRFLOW_ASYNC



Strategy 1 ExecutionMode.AIRFLOW_ASYNC

Performance improvements

| Execution mode | Time to run project (minutes) |
|---------------------------------------------------------------------------------------------------|-------------------------------|
| dbt run | 13 |
| Cosmos 1.9 (*) with ExecutionMode.LOCAL (Airflow with a local astro-cli setup) | 11 |
| Cosmos 1.9 (*) with ExecutionMode.AIRFLOW_ASYNC (Airflow with a local astro-cli setup) | 11 |
| Cosmos 1.11a6 (**) with ExecutionMode.AIRFLOW_ASYNC (Airflow with a local astro-cli setup) | 7 |

(*) Using Cosmos 1.9 with Airflow 3.0.2, dbt Core 1.10 (installed in the same Python virtualenv as Airflow and Cosmos), against Postgres. We are using a dbt project that has [129 models](#).

(**) Cosmos 1.11a6 has two main improvements regarding the AIRFLOW_ASYNC, originally released in Cosmos 1.9:

- Use of XCom instead of Remote object store to exchange compiled SQL files [#1934](#)
- Reuse Python virtualenv across tasks running in the same worker node [#1939](#)

Strategy 1 ExecutionMode.AIRFLOW_ASYNC

Pros

- Reduced benchmark Airflow DAG run time by 36%
- Single dbt command invocation (less CPU/memory allocation)
- Non-blocking transformations in the data warehouse with Airflow deferrable operators

Cons

- Currently only supports dbt models
- Currently only supports BigQuery
- Implementation specific per data warehouse
- The dbt project cannot have models with a metadata-dependency on other models
- Some users reported issues related to Airflow deferrable operators, which can be challenging to reproduce
- Only works with dbt Core (not dbt Fusion)

https://astronomer.github.io/astronomer-cosmos/getting_started/async-execution-mode.html

Next steps

ExecutionMode.AIRFLOW_ASYNC

1. Available since Cosmos 1.9 for BigQuery
2. Significant performance improvements in Cosmos 1.11 pre-releases
3. We need feedback!
4. Maybe add non-async version
5. We've introduced telemetry to evaluate the adoption
6. We'd love [contributions](#)

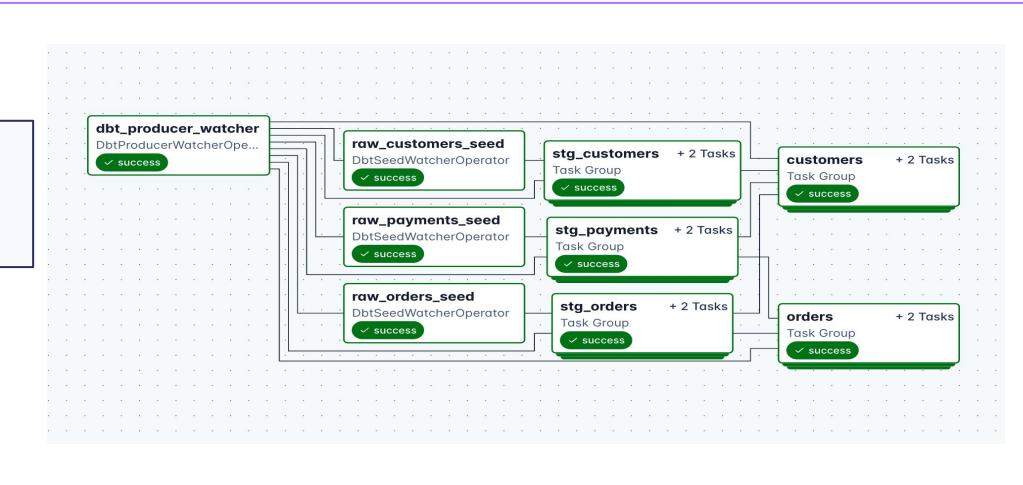
is:issue state:open label:execution:async

5. Strategy 2: Watcher

Strategy 2 ExecutionMode.WATCHER

Setup task

run dbt build
in a single
Airflow task



all the other
tasks are
sensors

<https://github.com/astronomer/astronomer-cosmos/issues/1950>

Strategy 2 ExecutionMode.WATCHER

Registering callbacks

Register `callbacks` on dbt's `EventManager`, to access structured events and enable custom logging. The current behavior of callbacks is to block subsequent steps from proceeding; this functionality is not guaranteed in future versions.

```
from dbt.cli.main import dbtRunner
from dbt_common.events.base_types import EventMsg

def print_version_callback(event: EventMsg):
    if event.info.name == "MainReportVersion":
        print(f"We are thrilled to be running dbt{event.data.version}")

dbt = dbtRunner(callbacks=[print_version_callback])
dbt.invoke(["list"])
```

<https://docs.getdbt.com/reference/programmatic-invocations#registering-callbacks>

Strategy 2 ExecutionMode.WATCHER

Performance improvements

| Execution mode | Number of threads | Time to run project (minutes) |
|-------------------------------------------------------------------------------------|-------------------|-------------------------------|
| dbt build in the CLI | 4 | 6 - 7 |
| dbt run for each model individually | | 30 |
| Cosmos default ExecutionMode.LOCAL in Astro CLI locally | | 10 - 15 |
| Cosmos proposed ExecutionMode.WATCHER in Astro CLI locally | 1 | 26 |
| | 2 | 14 |
| | 4 | 7 |
| | 8 | 4 |
| | 16 | 2 |
| The ExecutionMode.WATCHER in Airflow with an Astro deployment (A10) | 8 | 5 |

Strategy 2 ExecutionMode.WATCHER

Pros

- Reduced DAG run time to 1/5th of the original time
- Single dbt run
 - generates unified `run_results.json`
 - support dbt pre-hook & post-hook
- Data warehouse-agnostic implementation

Cons/Current Limitations

- Airflow worker is blocked by transformations happening in the data warehouse
- Retries have the same performance as `ExecutionMode.LOCAL`
- Currently relies on dbt and Airflow being installed in the same Python venv (some users report conflicts between dependencies)
- Unclear how these features should work: Cosmos callback, Airflow datasets/assets and OpenLineage events

<https://github.com/astro/astro-cosmos/issues/1950>

Next steps

ExecutionMode.WATCHER

1. We successfully ran a PoC during August 2025
2. First release estimate: end of October 2025 (available in **1.11.0a6**)
3. Work on making sensors deferrable
4. We need feedback!
5. We'd love [contributions](#)

is:issue state:open label:execution:watcher

6. Takeaways

Takeaways

1. To run the **same dbt pipeline with multiple dbt command** is slow
2. To use **Airflow deferrable operators** allows to not **wait for the transformation** in the **data warehouse**, which can save 36% DAG runtime
3. It is not always possible to **pre-compile a dbt project**
4. The **watcher approach** reduces the DAG runtime up to 80% and it is agnostic to **data-warehouse**
5. We need feedback and help!



Learn more about how to run dbt
with Apache Airflow and Cosmos

The image shows the front cover of a book titled "Orchestrating dbt with Apache Airflow® using Cosmos". The cover is dark blue with a red spine. It features a stylized orange and green graphic of a plane or rocket launching into space with stars. The title is written in white and orange text. To the left of the book, there is a vertical column of Python code. The code is a DAG (Directed Acyclic Graph) definition using the dbt Python API. It includes imports for `dbt_project`, `dbt_task`, and `dbt_node`. The DAG is named `example_injection` and defines a task `pre_dbt` that runs a command `return os.environ['AIRFLOW_HOME']` and a project `dbt_project` with various configuration parameters like `group_id`, `project_config`, and `profile_config`.

```
1  from dbt_project import dbt_project
2  from dbt_task import dbt_task
3  from dbt_node import dbt_node
4
5  # Import the dbt package
6  import dbt
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
```

```
    target_name="dev",
    profile_mapping=PostgresUserPasswordProfileMapping(
        conn_id=POSTGRES_CONN_ID,
        profile_args={"schema": SCHEMA_NAME},
    ),
)
# Only needed if you're using dbt as a dependency
_execution_config = dbt_execution_config()
dbt_executable = dbt_executable()

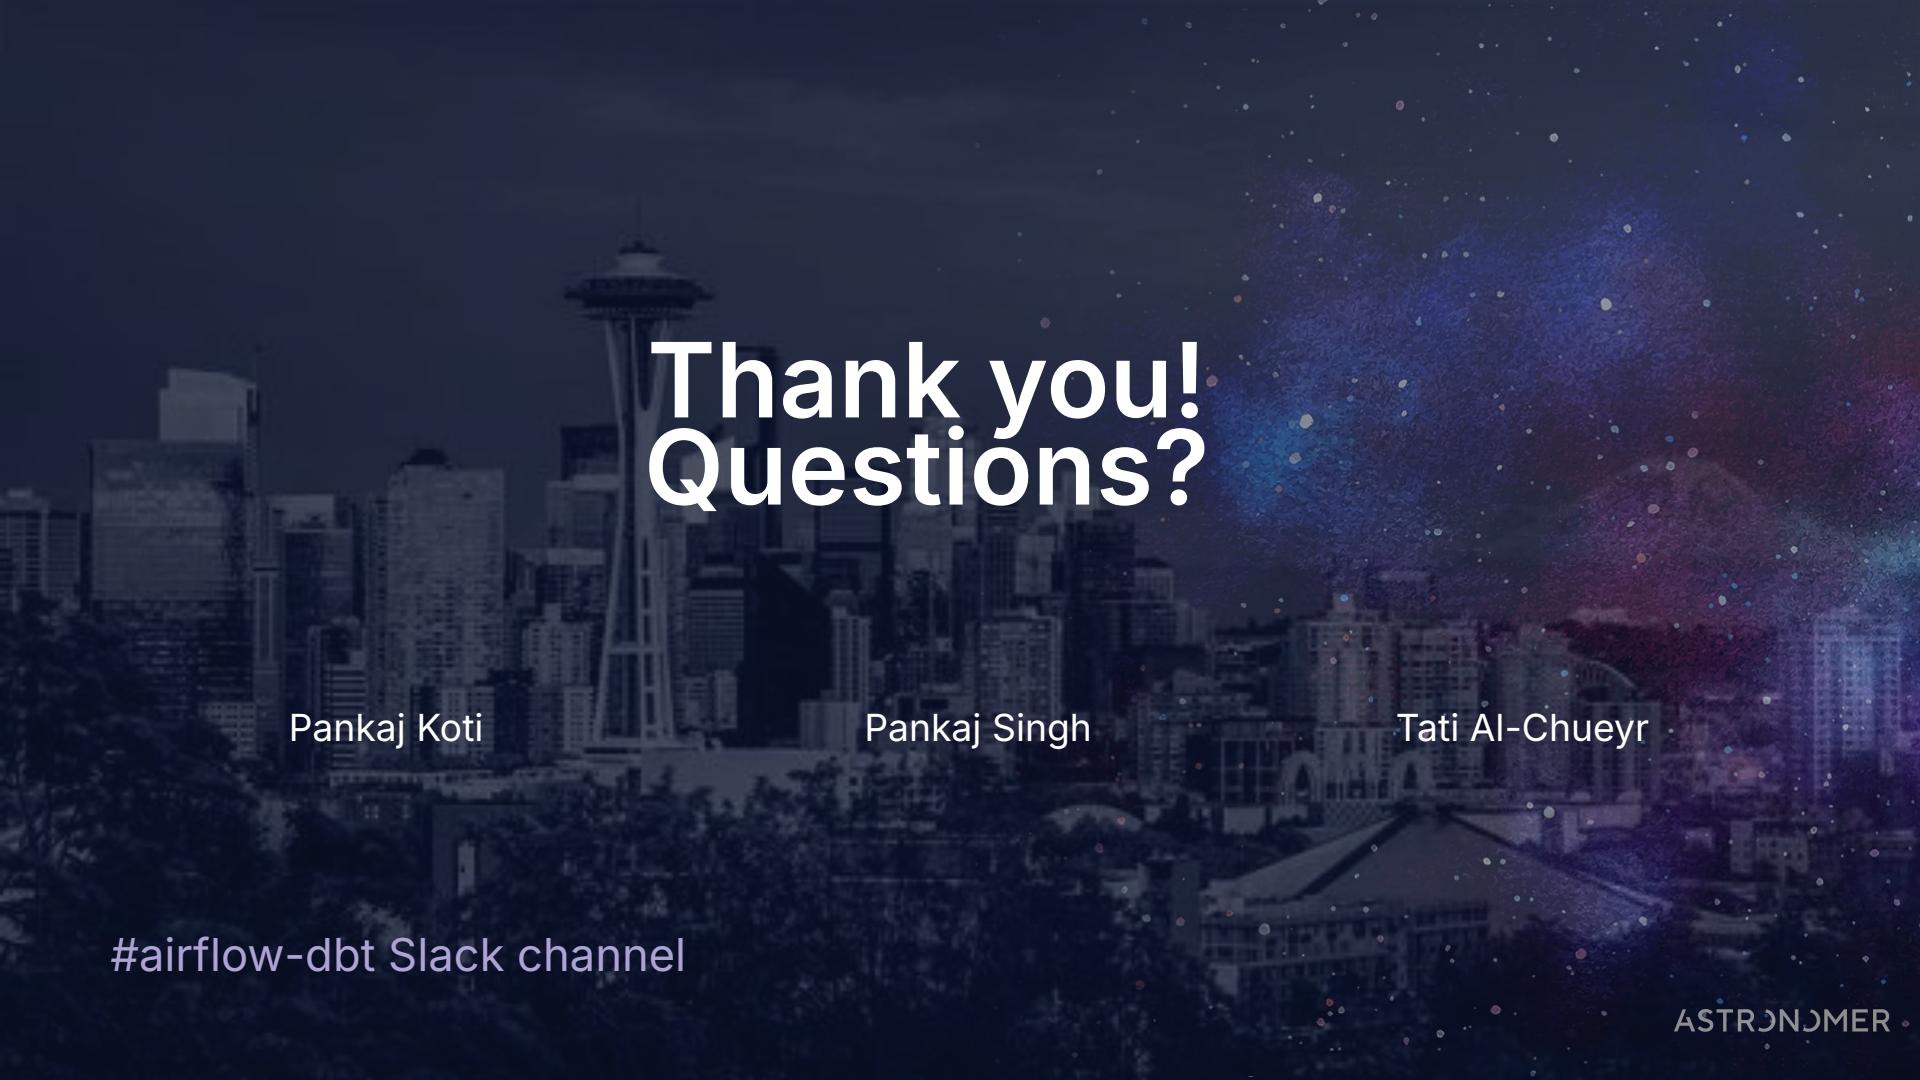
@dag(
    params={
        "my_department": "DEPT_A"
    },
    tags=["outdated"]
)
def example_injection():
    @task
    def pre_dbt():
        return os.environ['AIRFLOW_HOME']

    _pre_dbt = pre_dbt()
    dbt_project(
        group_id="dbt_project",
        project_config=_project_config,
        profile_config=_profile_config,
        execution_config=_execution_config,
        operator_args={
            "vars": {"my_department": "{{ ti.xcom_pull(task_ids='pre_dbt') }}"}}
    )
```

The 2025 Apache Airflow® Survey is here!

Fill it out to for a free Airflow 3
Fundamentals or DAG Authoring in
Airflow 3 certification code



A dark, atmospheric photograph of the Seattle skyline at night. The Space Needle is prominent in the center-left. The city lights are reflected in the water in the foreground, and a bridge is visible in the lower right. The overall mood is mysterious and professional.

Thank you! Questions?

Pankaj Koti

Pankaj Singh

Tati Al-Chueyr

#airflow-dbt Slack channel

ASTRONOMER