



Building Airflow Setups Resilient to Zonal/Regional Down Events

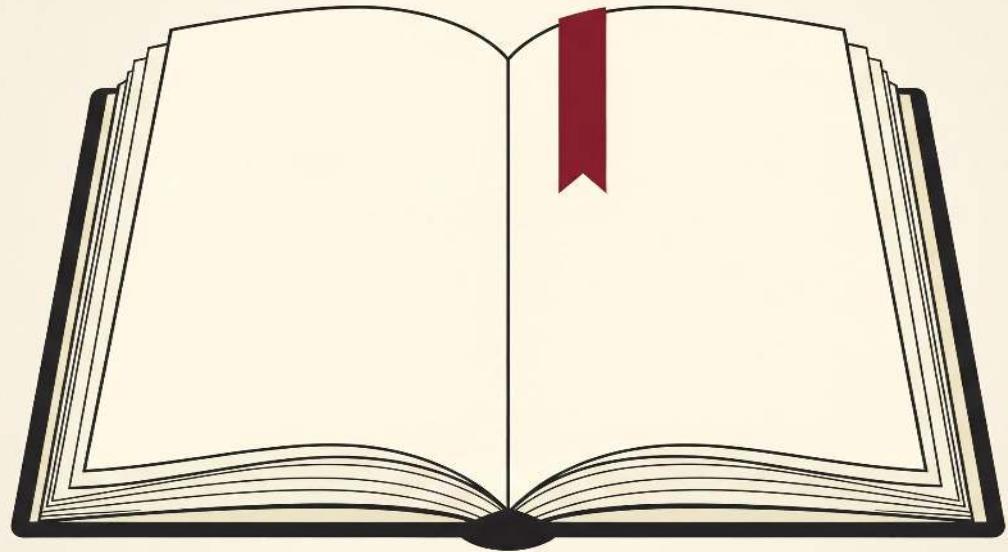
Khaled Hassan

3.0



Agenda

- The Resilience Problem
- Triggers of the problem
- Mitigation Strategies
- Conclusion
- QA

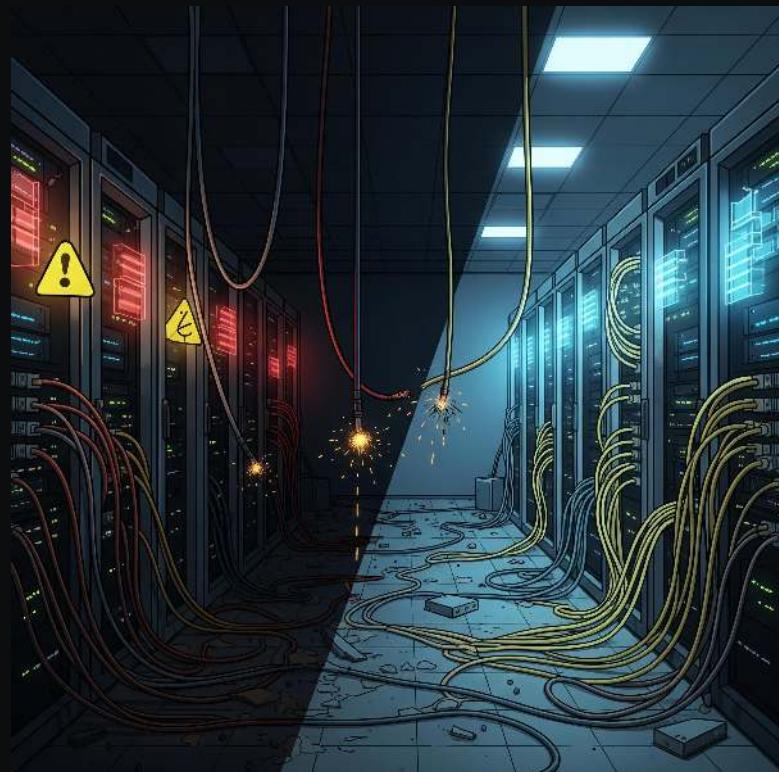


The Problem



The Problem: Vulnerability to Physical Failures

- Using Airflow as a solution for your business
- Physical failure (e.g. power/network outage)
- Running software behaviour during the outage in two aspects:
 - **Physically** - How and where are the critical infrastructure components located.
 - **Software** - How DAGs are prone to these failures

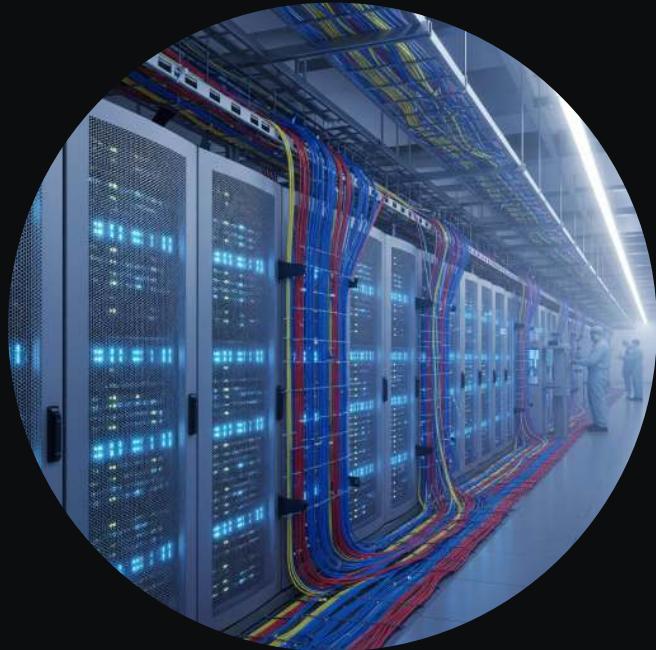


Mitigation Strategies

- **Retries should be enabled**
- **Infrastructure**
 - Deployments Replication
 - Deployments Spread
 - Data Replication
- **DAG Implementation**
 - DAGs statelessness
 - DAGs idempotency
 - DAGs dependencies replication
- **Complementary Strategies** - One on its own may not be sufficient.



Mitigation Strategy 1: Replication Of Critical Components



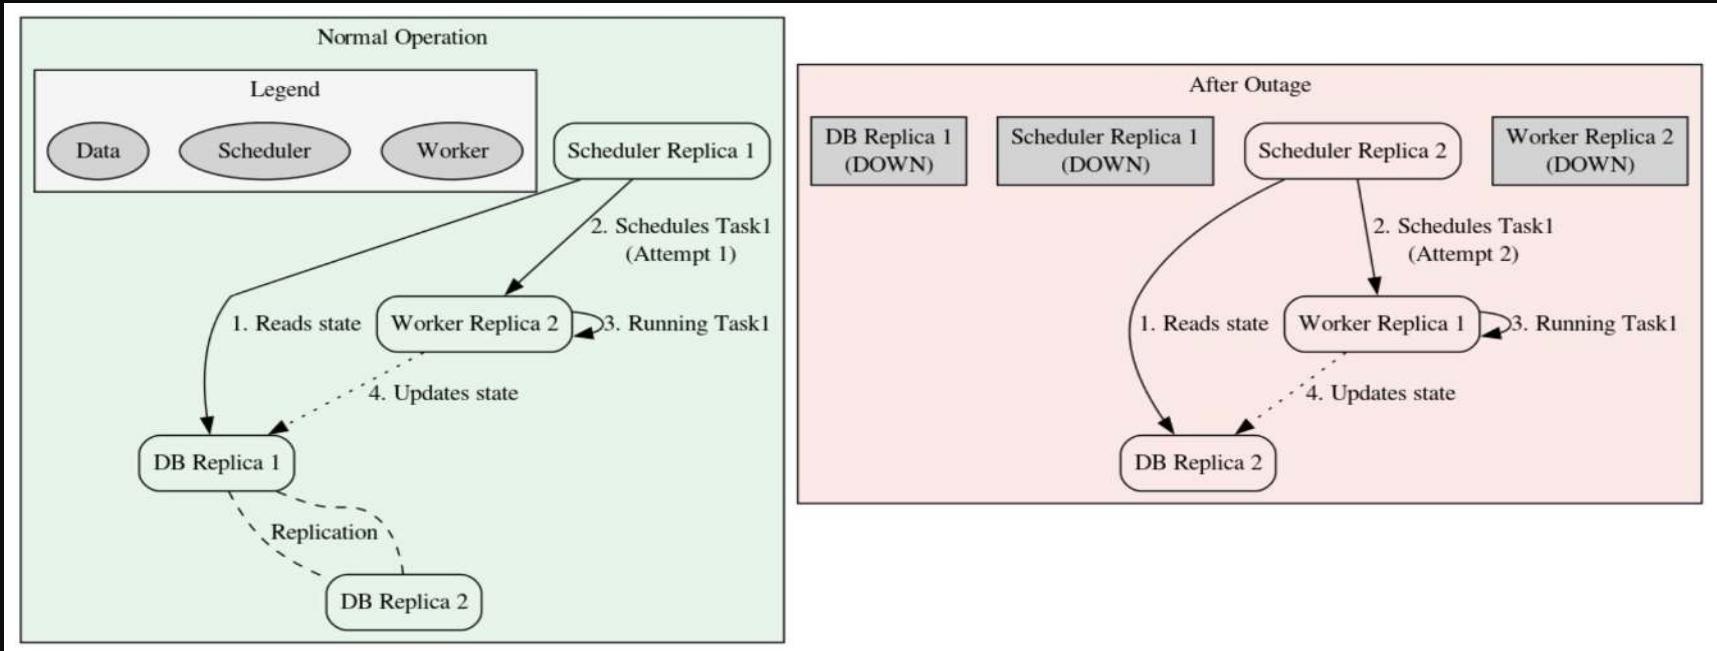
- **Deployments Replication**
 - Critical Deployments (e.g. Workers, Schedulers, Internal API in AF3)
 - Failovers to healthy replicas during failures
- **Data Replication - Airflow DB**
 - Ensure that healthy replicas can always access needed data
 - Replicating DBs might be more challenging than workloads

Real Life Scenario 1



- **Identical server rooms (2 Failure Domains):**
 - Failure Domain 1: Room A contains
 - Scheduler Replica 1
 - DB Replica 1
 - Worker Replica 2
 - AF3: Internal API Replica 2
 - Failure Domain 2: Room B contains
 - Scheduler Replica 2
 - DB Replica 2
 - Worker Replica 1
 - AF3: Internal API Replica 1
- **Only Room A Loses Power**
- **Potential Problem: Building-wide Power Outage ?**

Mitigation Strategy 1: Replication Of Critical Components



Potential Problem - What if both Rooms share the same infrastructure (e.g. power/network) ?

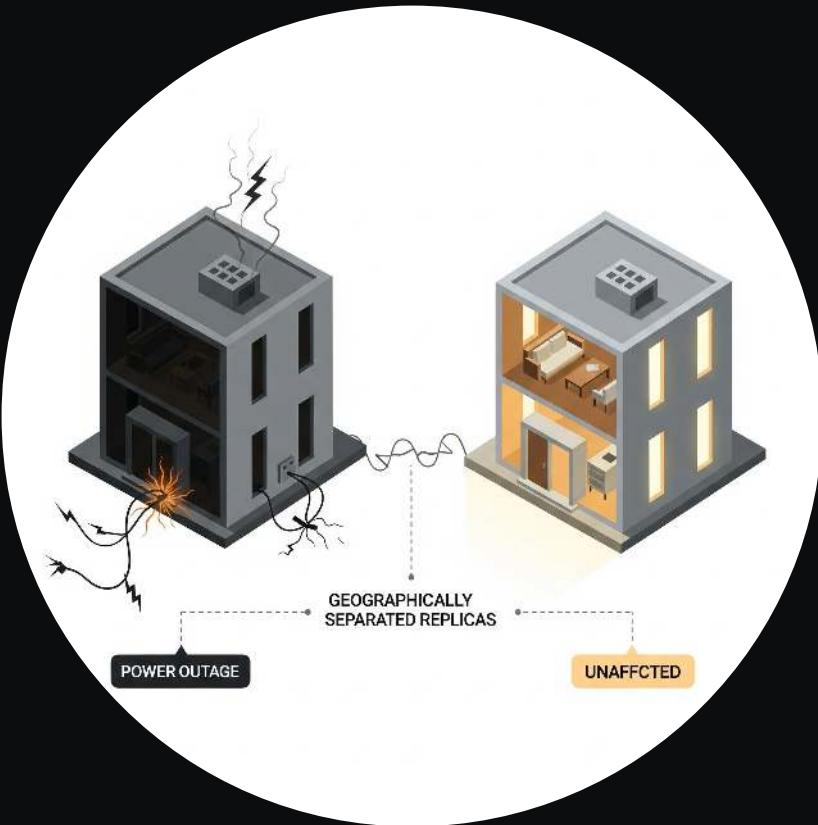
Mitigation Strategy 2: Spread of replicated components

- **Mitigation of Zonal/Regional Outages**
- **Enhanced Availability**
- **Zonal/Regional Disaster Recovery**
- **Cloud solution example:**
 - **Components are workloads in Kubernetes**
 - **Defining topologySpreadConstraints in kubernetes**

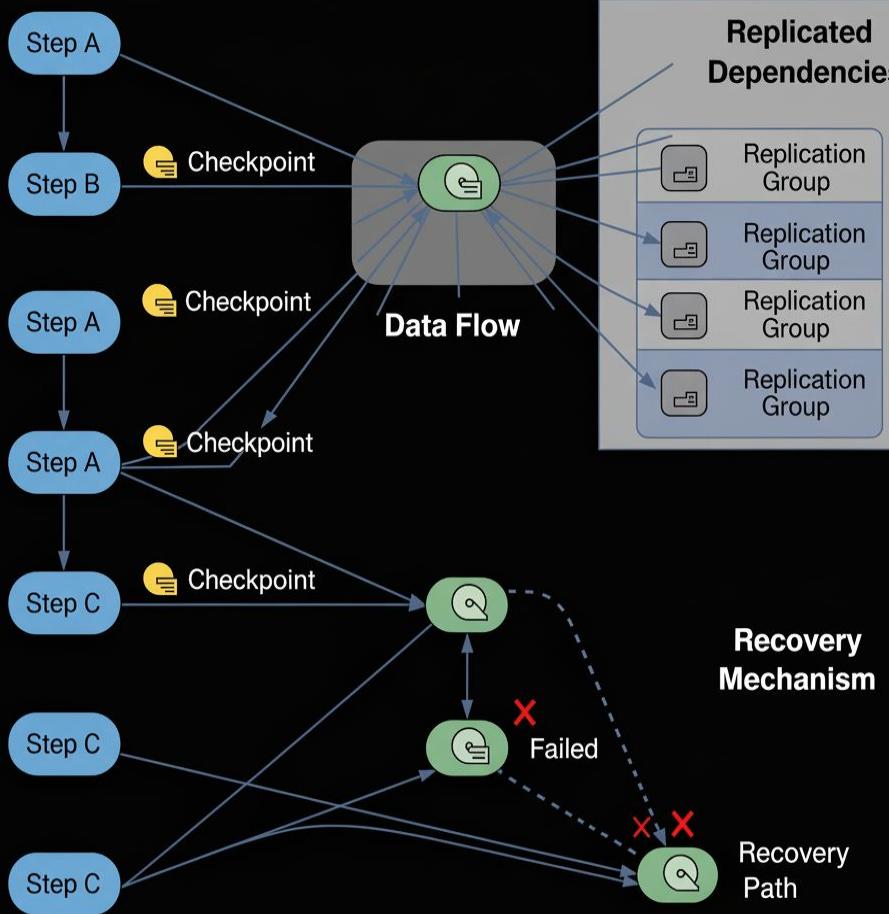
```
topologySpreadConstraints:  
  - maxSkew: 1  
    topologyKey: topology.kubernetes.io/zone
```

Real Life Scenario 2

- Room A is now in Building A and Room B is in Building B
- Each building now is a failure domain
- Whole Building A loses Power
- Continuity of processes hence continuity of business



Stateless DAG

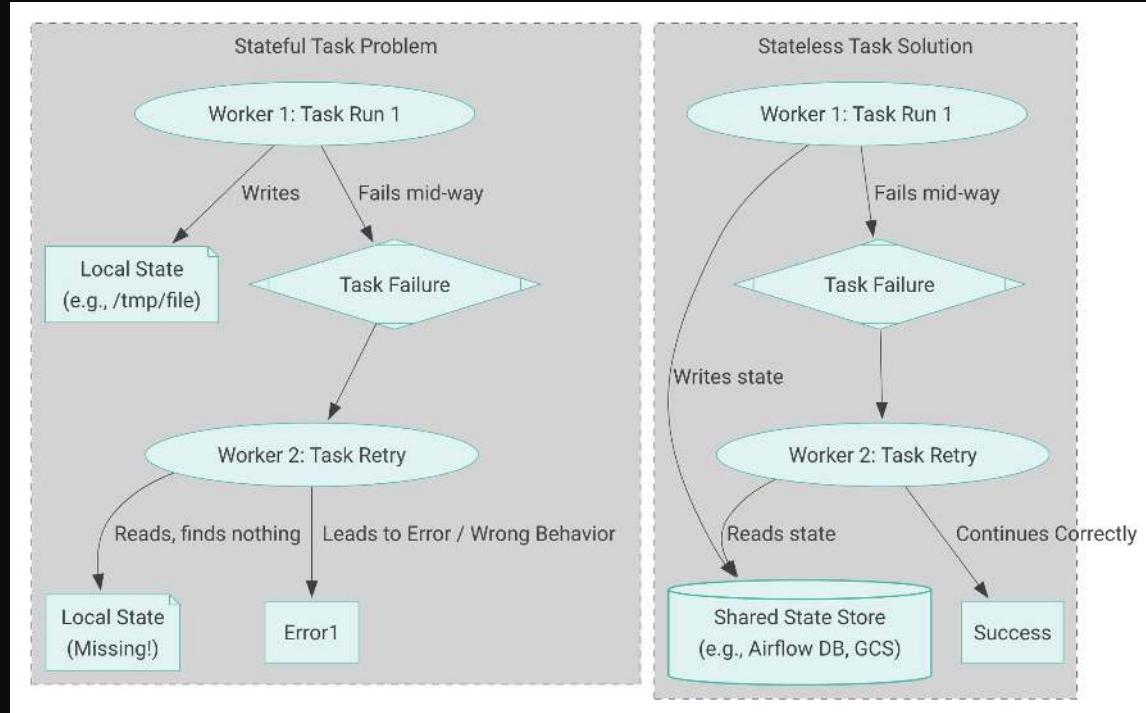


Mitigation Strategy 2: DAG Implementation

- **Stateless DAGs**
- **Idempotent Operations**
- **Replicated Dependencies**

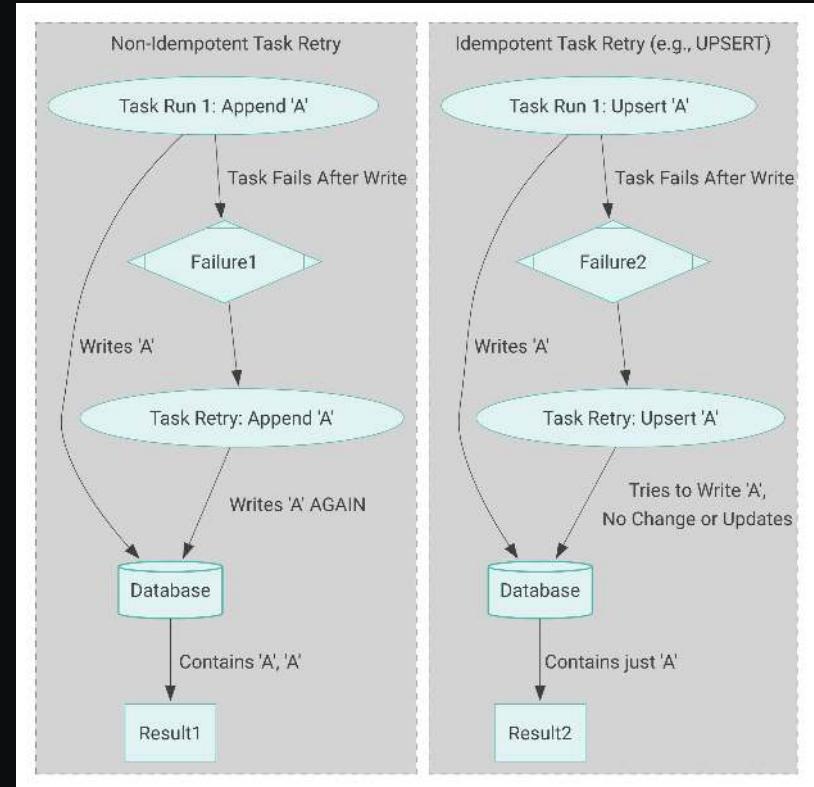
DAG Strategy 1: Stateless DAGs

- **Stateful Definition**
- **Stateless Definition**
- **Why is Statefulness a Problem?**
 - **Lost Progress**
 - **Inconsistent State on retries**
 - **Failover Issues**



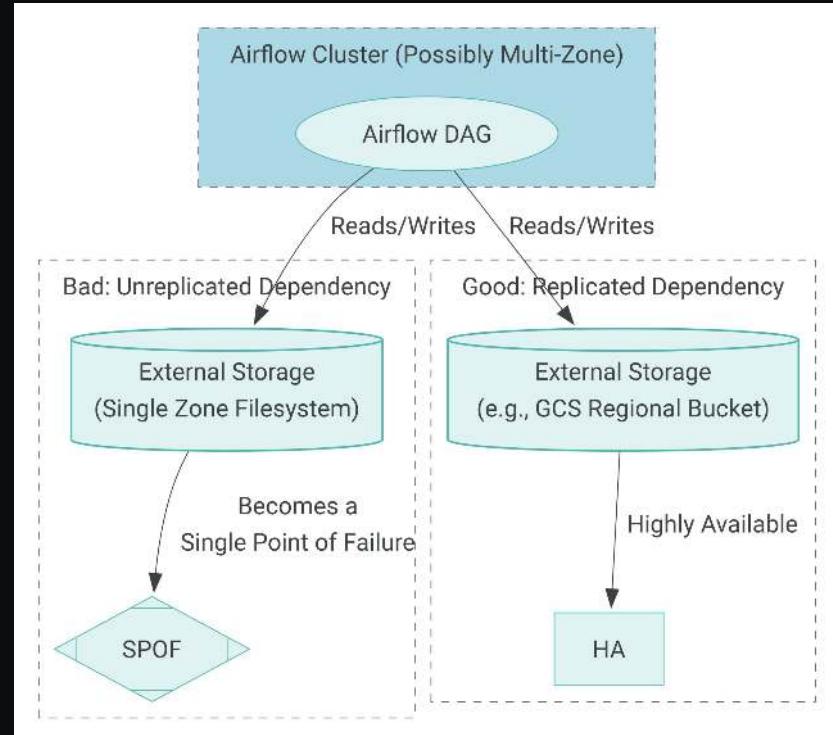
DAG Strategy 2: Idempotent Operations

- **What is Idempotency?**
- **Why Idempotency is Crucial for Reliability?**
 - Preventing Duplication
 - Simplified Recovery after failures



DAG Strategy 3: Replicating DAG Dependencies

- **External State & Single Points of Failure**
 - **DAGs often depend on external systems (storage, databases, services).**
 - **Replication of dependencies is crucial.**



Cloud Composer

- **Managed Airflow Environment**
- **High Resilience Configuration**
 - Minimum Number of Replicas
- **Underlying Infrastructure Handled automatically.**





Conclusions

- If availability is crucial for your business:
 - Replication and Spread is crucial for critical components
 - DAGs implementation analysis is crucial (Idempotency, Statelessness, etc)
- Resilience comes with a trade off :
 - Cost
 - Complexity
 - Sometimes latency
- Larger scale failures are sometimes inevitable

Questions?

hkhaled@google.com