



# DAGLint

## Elevating Airflow DAG Quality Through Automated Linting

**Snir Israeli**  
Data Engineer @Next

3.0

# Agenda

**01**

Challenges with maintaining DAGs quality at scale

**02**

Our solution - DAGLint!

**03**

Integrations

**04**

Achievements

**05**

Q&A

# DAGs at Scale Can be Messy

## Inconsistency

Teams grow, styles diverge, no standards => **inconsistent DAGs**

## Maintenance Hell

Debugging hurts, steep learning curve

## Hidden anti-patterns

Scheduler slowness, avoidable costs

## Delayed Processes

Slow Code Reviews, onboarding drags




Photo by [Nathan Cima](#) on [Unsplash](#)

## Traditional Syntax

```
from airflow import DAG
from airflow.operators.python import PythonOperator
from datetime import datetime

def say_hello():
    print("Hello World")

with DAG(
    dag_id="hello_world_traditional",
    start_date=datetime(2025, 1, 1),
    schedule_interval=None,
    catchup=False,
) as dag:
    hello = PythonOperator(
        task_id="hello_task",
        python_callable=say_hello,
    )
```

Traditional Syntax,  
logic imported

## Inconsistency

## TaskFlow API

```
from airflow.decorators import dag, task
from datetime import datetime

@dag(
    dag_id="hello_world_taskflow",
    start_date=datetime(2025, 1, 1),
    schedule_interval=None,
    catchup=False,
)
def hello_dag():
    @task
    def say_hello():
        print("Hello World")

    say_hello()

hello_dag()
```



```
from airflow import DAG
from airflow.operators.python import PythonOperator
from datetime import datetime
from hello_functions import say_hello

with DAG(
    dag_id="hello_world_external",
    start_date=datetime(2025, 1, 1),
    schedule_interval=None,
    catchup=False,
) as dag:
    hello = PythonOperator(
        task_id="hello_task",
        python_callable=say_hello,
    )
```



# DAG missing Documentation & Tags

```
# ❌ Missing doc_md - no owner, playbook, purpose
# ❌ Missing tags    - hard to navigate and filter

with DAG(
    dag_id="dag_with_no_documentation_nor_tags",
    start_date=datetime(2025,1,1),
    schedule_interval="@daily",
    catchup=False,
)
```



CodelImage

# Top-level “expansive” code

```
# ANTI-PATTERN: running expensive code at import time
import psycopg2
conn = psycopg2.connect("dbname=prod user=airflow") # X should be inside a task
cur = conn.cursor()
cur.execute("SELECT 1") # runs on every scheduler/worker import

from airflow import DAG
from airflow.operators.empty import EmptyOperator
from datetime import datetime

with DAG(
    "bad_top_level_side_effects",
    start_date=datetime(2025,1,1),
    schedule_interval=None,
    catchup=False
) as dag:

    EmptyOperator(task_id="dummy")
```



# What we tried before

## ...and why it didn't work



### Training Sessions

#### The Problem

Knowledge fades over time. People forget best practices weeks after training.

### Code Reviews

#### The Problem

Relies on reviewers remembering to check for anti-patterns. Inconsistent enforcement.

### Documentation

#### The Problem

Becomes stale quickly. Doesn't evolve with changing best practices.


### Monitor Performance

#### The Problem


Reactive approach. Issues are caught only after they've already impacted production.

# Enter, DAGLint

- ▶ Linter for Airflow DAGs
- ▶ Runs locally on terminal
- ▶ Fast and deterministic
- ▶ Identify and prevent anti-patterns
- ▶ Clear & useful output



# Architecture (High Level)



# CLI & Local Dev Workflow

# Lint single DAG


```
daglint /path/to/dags/ my_dag_name
```

# Lint all active DAGs


```
daglint /path/to/dags/ --all
```

# Run specific rules

```
daglint /path/to/dags/ my_dag --rules_to_run R01,R16
```




# Architecture (High Level)




# The Engine

- Walk DAGs directory & identify DAG files
- Lint DAG file/s
- Respect comment-based exclusions
- Scoring Mechanism



# Architecture (High Level)





# Rules Framework



File Organization Rules



DAG Structuring Rules



Code Quality Rules

Naming Conventions

File Organizations

Valid DAG ID formats



# Rules Framework



File Organization Rules



DAG Structuring Rules



Code Quality Rules

Context managers

No function definitions

No business logic



# Rules Framework



File Organization Rules



DAG Structuring Rules



Code Quality Rules

No top-level expansive calls

README.md

# How's a rule defined?

```
class CustomRule(LintRule):

    def __init__(self, **kwargs):
        super().__init__(
            name="custom_rule_name",
            description="Rule Description",
            id="R99",
            **kwargs
)

    def validate(self):
        # Custom AST analysis logic...
```



Creating rules is simple; Just inherit from *LintRule* and implement the *validate* method



New rules are automatically discovered via inheritance



CodeImage

# Exclusions (Granular & Documented)



Localized opt-outs



Force **documentation** of intent

```
# daglint-exclude: <<R04 Reason for exclusion


some_task = PythonOperator(task_id="task_id_that_violates_a_rule", ...)

# R04>>
```



CodelImage

# Architecture (High Level)



# Python AST




## What is an AST?

- **Tree Structure:** Python breaks your code into a tree of nodes, where each node represents a construct in your code (like operations, variables, functions)
- **Abstract:** It ignores unnecessary details like whitespace and focuses on the structure and meaning
- **No execution:** Code is parsed and inspected without being executed, fast and safe.
- **Used For:** Code analysis, linters, formatters, transpilers, and understanding code structure programmatically

## Your Python Code

```
x = 5 + 3
```

## How Python “Sees” It (AST)



# AST Node Visitors

## What is a Node Visitor?

A [Node Visitor](#) is a pattern that lets you "walk" through every node in the AST tree and perform actions when you encounter specific node types.

## How it works?

You create a class that inherits from `ast.NodeVisitor` and define `visit_*` methods for each node type you care about.

## Simple Example

```
import ast

class FunctionCounter(ast.NodeVisitor):
    def __init__(self):
        self.count = 0

    def visit_FunctionDef(self, node):
        self.count += 1

        self.generic_visit(node)

# Count all functions in code
counter = FunctionCounter()
counter.visit(ast.parse(code))
```



CodelImage

# Airflow Example

```
# dag_example.py
from airflow import DAG
from datetime import datetime

with DAG(
    dag_id="my_dag",
    start_date=datetime(2025, 1, 1),
    schedule_interval="@daily",
) as dag:
```



```
import ast

class DAGDocValidator(ast.NodeVisitor):
    def visit_Call(self, node):
        # Check if this is a call to DAG(...)
        if isinstance(node.func, ast.Name) and node.func.id == "DAG":
            # Collect keyword argument names
            kwarg_names = {kw.arg for kw in node.keywords if kw.arg is not None}

            if "doc_md" not in kwarg_names:
                print("✗ DAG definition is missing 'doc_md' attribute")
            else:
                print("✓ DAG definition includes 'doc_md'")


        # Keep traversing child nodes
        self.generic_visit(node)

# Load and parse the DAG file
with open("dag_example.py", "r") as f:
    tree = ast.parse(f.read())

# Run validator
validator = DAGDocValidator()
validator.visit(tree)
```



# Architecture (High Level)



# Linting Results

Clear DAG  
Score

Linting Results for [REDACTED]

Score: 73.00% 😕

Rule Name	Description	Status	Line Numbers
dag_has_no_top_level_expensive_calls	The DAG must not have top-level calls to expensive classes/services, such as AWS services directly or via bi_toolbox utility classes. i.e. SecretsManager, S3, Athena, Redshift, etc.	Passed	N/A
dag_is_defined_only_as_context_manager	The DAG object should be instantiated using a context manager	Passed	N/A
dag_has_dag_description_configured	DAGs should have a description as a README.md file, located right next to the DAG's file, configured as a doc_md keyword argument.	Failed	N/A
dag_id_does_not_match_dag_file_name	The DAG file name should match its dag_id.	Passed	N/A
dag_filename_must_be_all_lowercase_characters	The DAG file name must be all lowercase characters.	Passed	N/A
dag_has_logger_defined	The DAG should have a logger defined in the file and are configured with all relevant parameters	Passed	N/A
dag_file_enclosed_within_its_own_folder	Every DAG should be placed in its own folder, its file name is part of its folder name and the folder is under the dags folder hierarchy.	Failed	N/A
dag_uses_only_[REDACTED]	The DAG should always use [REDACTED] and not directly the DataQualityHandler class	Passed	N/A
dag_has_a_valid_team_tag	The DAG object must have a tag for the team owning the DAG in the form of a team constant. i.e. [REDACTED]	Passed	N/A
dag_dq_config_path_is_valid	The DAG should use [REDACTED] and have its corresponding config file in the config folder with the suffix '_dq_tests.json'.	Passed	N/A
[REDACTED]	[REDACTED]	Failed	26, 40, 54, 104, 131

Easily find violations

# CI/CD



**01**

## GitHub Actions

GitHub Actions runs DAGLint on every PR where a DAG file was modified



**02**

## Merge Blocks

Critical rule failures block merges with detailed output



**03**

## Uninterrupted Code Reviews

Reviewer can focus on what's important



**04**

## Clean Main Branch

Keeps main branch green & consistent

← DAGLint

## Bi 15644/di new partnership lead form experience #13619

Summary

Jobs

✓ Linting modified DAG files using D...

Run details

⌚ Usage

💾 Workflow file

### Linting modified DAG files using DAGLint

succeeded on Aug 25 in 31s

- > ✓ Set up job
- > ✓ 🛡 Check out repository code
- > ✓ Run actions/setup-python@v5
- > ✓ 📁 Get Changed Files
- > ✓ 🐍 Find Python files with DAG instantiation
- > ✓ 📦 Install DAGLint dependencies
- > ✓ 💎 Run DAGLint on modified DAG files
- > ✓ 📣 Comment PR
- > ✘ 🚫 Fail if critical rules have failed
- > ✓ Post Run actions/setup-python@v5
- > ✓ Post 🛡 Check out repository code
- > ✓ Complete job



github-actions bot commented on Aug 25 · edited by snir-israeli ▾

### DAGLint Results:

DAG File Name	Number of Violations	Score
example_dag_1	0	100.00% 😊



snir-israeli merged commit **6dddc06** into master

### 7 checks passed

- ✓ check\_user\_group
- ✓ Linting modified DAG files using DAGLIn
- ✓ lint
- ✓ pytest

# Monitoring & Analytics




Daily Jenkins pipeline runs org-wide scoring



Persist results to PostgreSQL



Tableau Dashboard



# Creative Use Case



## Gradual Migration without Regression

- Replace old custom operator usage with new version,  
**gradually**
- While transitioning safely, we didn't want new DAGs  
or updates to existing DAGs cause degradation
- Created a rule that disallow the usage of the old  
operator
- **Any code change to a DAG using the old operator  
will fail linting**


# What did we achieve?

## Achievements

- ✓ Developer Productivity
- ✓ Improved Code Reviews
- ✓ Ops & Compliance
- ✓ Improved Engineer Onboardings

Quality enforced at scale





# Takeaways

-  Violating best-practices and inconsistencies = **Quality Issues**

---
-  Airflow needs domain-aware linting

---
-  Automatic enforcement + Monitoring = **Quality & Compliance at scale!**

---
-  If you try it, focus on the developer experience!

# Medium Article



# Questions?

Thank you.

3.0

→ LinkedIn: [www.linkedin.com/in/snir-israeli](https://www.linkedin.com/in/snir-israeli)