



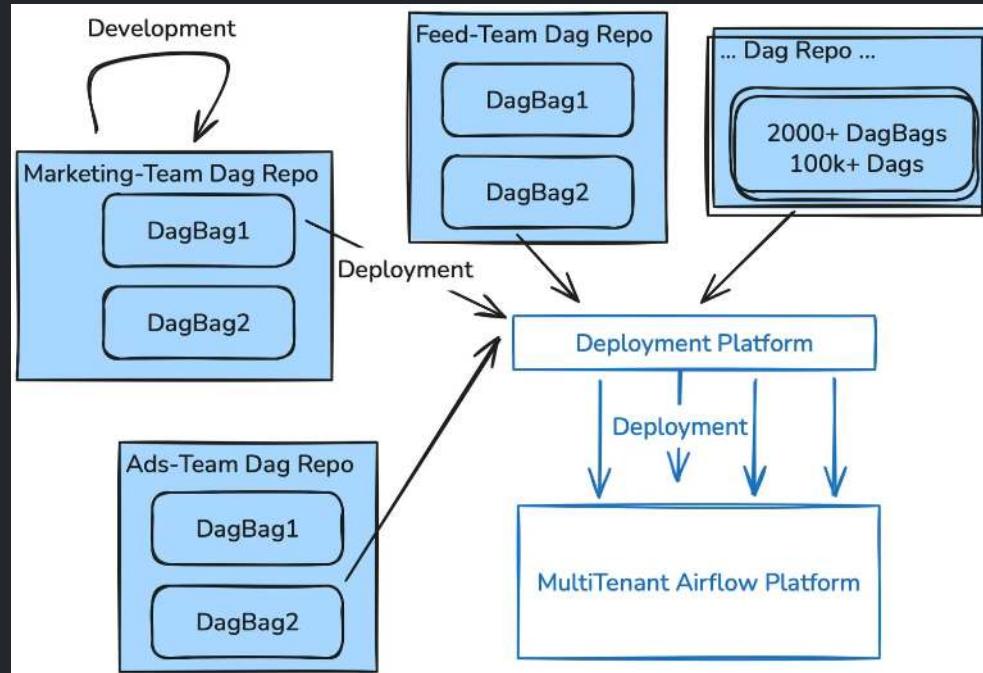
**DAGnistics: Shift-Left Airflow  
Governance With Policy  
Enforcement Framework**

**Managing 100k+ DAGs Without  
Breaking Developer Velocity**

Stefan Wang  
Senior Software Engineer  
Data Infrastructure @ LinkedIn  
Airflow Summit 2025

3.0

# The LinkedIn Scale Reality



Multi-Tenant Airflow Ecosystem

100K+

Active DAGs

2,000+

DAG Repositories

Platform-wide execution

Team-owned deployments

3,000+

DagBags

300K+

Daily Task Executions

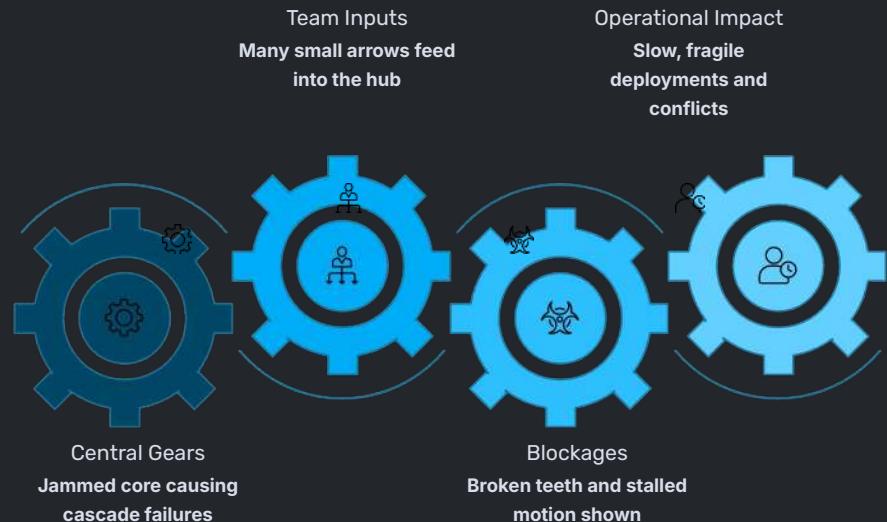
Isolated execution contexts

Critical daily volume

These complex workflows handle functional data pipelines, critical revenue generation, and compliance mandates.

# Past - The Monolithic Airflow DAG Repository

```
airflow/
├── dags/ # 2000+ teams
|   ├── marketing/
|   |   ├── weekly_email_campaign.py
|   ├── sales/
|   |   ├── monthly_sales_analysis.py
|   └── ...
├── plugins/
|   ├── shared_operators/
|   └── ...
└── requirements.txt # shared dependency across
    thousands of teams
    └── utils/ # shared utils across thousands of teams
        └── ...
```



# Asks - Autonomous DAG Repository Model



Complete Lifecycle Ownership

**Each of our 2,000+ teams controls their own DAG repository lifecycle: development, testing, and production.**



Standardized Structure

**Consistent structure mandates DAG definitions, business logic, comprehensive tests, and isolated dependencies.**



Autonomous Deployment

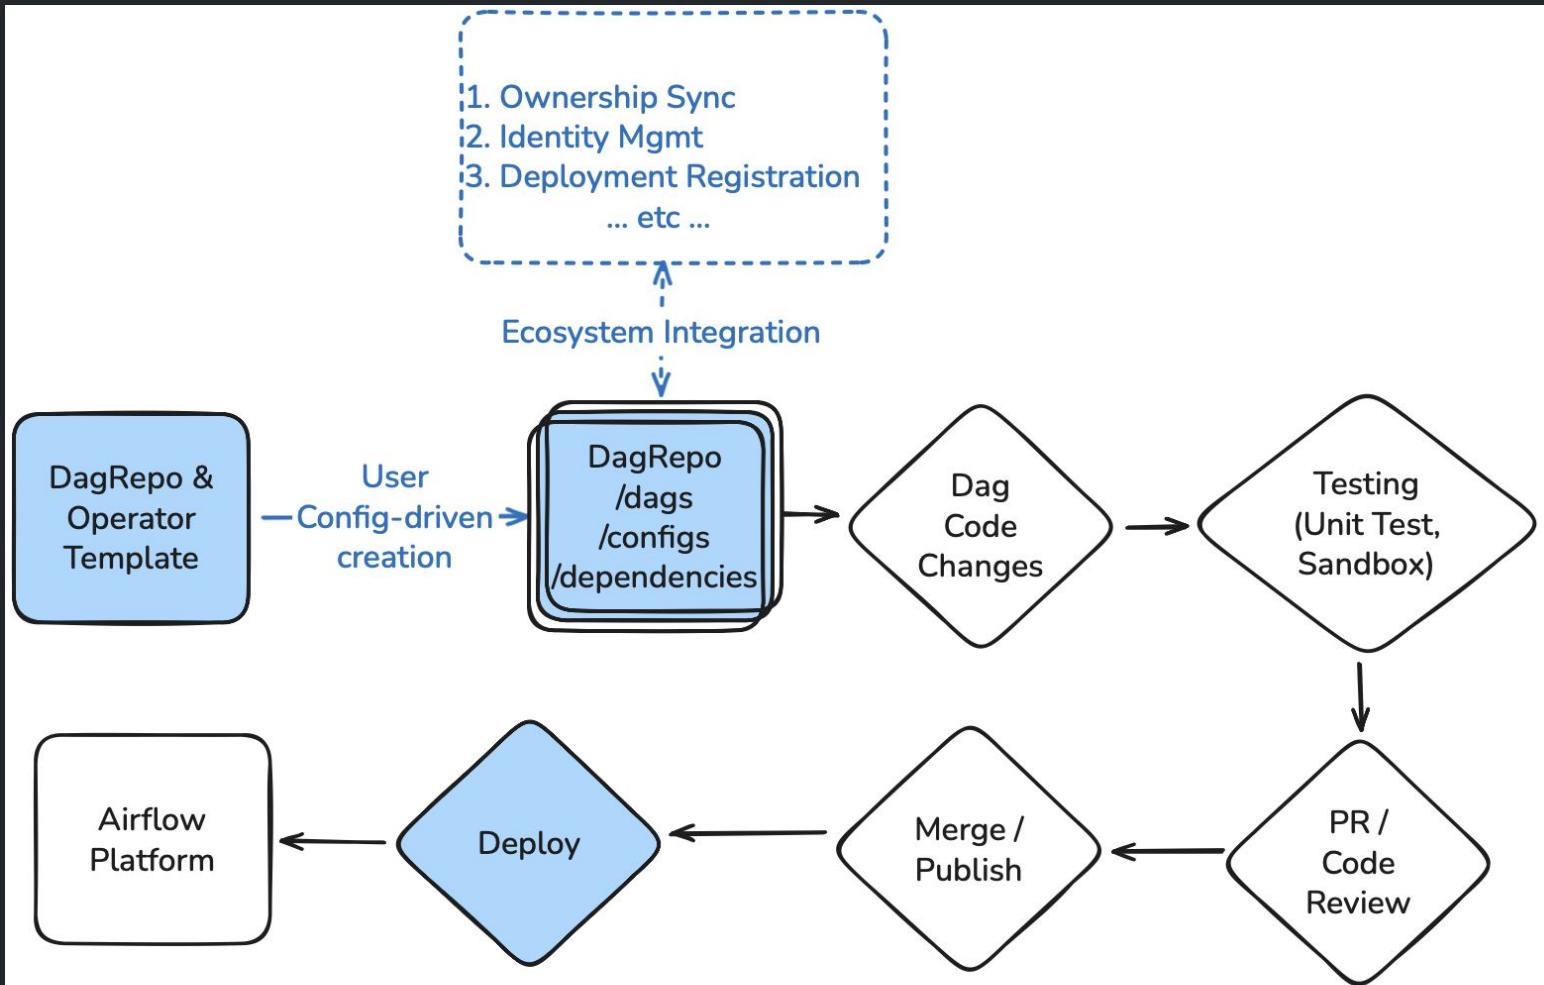
**Teams deploy independently for velocity. Sophisticated governance ensures stability.**

# Vision - Enterprise DAG CI/CD Pipeline

- Automated validation ensures rigorous policy compliance.
- Confidence built through staged verification.
- Maintains high deployment velocity.



# Future - Decentralized Dag Management



# Why Governance Matters in Multi-Tenant Airflow

## Chaos (No Governance)

- Failures only discovered in production
- Unclear ownership during critical incidents
- Frequent resource conflicts destabilize infrastructure
- Inconsistent practices across teams
- Compliance risks go undetected

## Stability (With Governance)

- Errors caught pre-deployment in CI/CD
- Clear ownership tracking and accountability
- Coordinated resource usage prevents conflicts
- Standardized monitoring and alerting
- Proactive compliance enforcement

When 2,000+ independent teams deploy on shared Airflow infrastructure, governance transforms chaos into stability.

Manual review cannot scale to hundreds of daily deployments—automated enforcement is essential.

# [Current Runtime-only] Airflow Cluster Policy

## Airflow's Native Policy System

Apache Airflow provides a cluster policy system to enforce custom rules. These policies are defined in `airflow_local_settings` and execute during DAG ~~Rollbacks~~. They can validate or mutate DAGs, tasks, task instances, and Kubernetes pods. They can reject deployments by raising exceptions.

## Two Critical Challenges

**Late Discovery:** Native policies execute within Airflow Platform Runtime. Violations are discovered after deployment in production.

**No Preflight Validation:** Developers cannot validate policies before committing code, leading to failed deployments and rollbacks.

---

So here's what we built...

# Dag Import Errors - Looks Familiar?

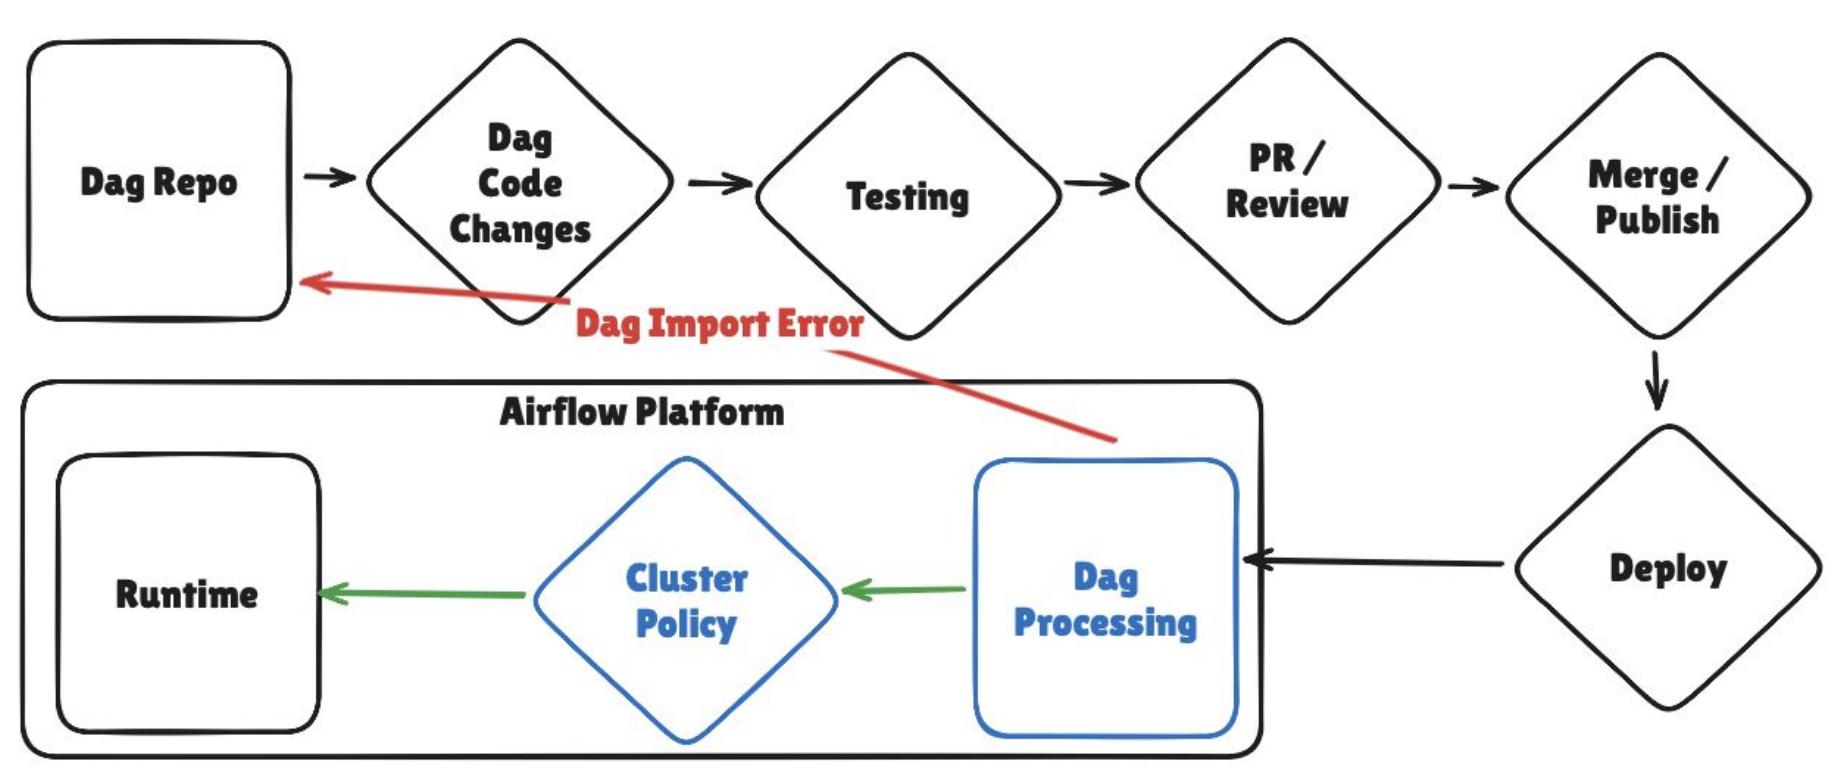


The screenshot shows the Apache Airflow web interface. At the top, there is a navigation bar with the following items: a logo, "Airflow", and menu options: "DAGs", "Datasets", "Security", "Browse", "Admin", and "Docs". The "DAGs" option is highlighted with a red border. Below the navigation bar, a red banner displays the message "DAG Import Errors (1)". Underneath the banner, a red box highlights a specific error message: "Broken DAG: [/opt/airflow/dags/11bdbf9e44aa6006233eb7f285318eb9/dags/dag\_idx\_ind.py] Traceback (most recent call last):".

! DAG Import Errors (1)

Broken DAG: [/opt/airflow/dags/11bdbf9e44aa6006233eb7f285318eb9/dags/dag\_idx\_ind.py] Traceback (most recent call last):

# Past



# Current State Challenges in Native Airflow Policies

Native Airflow policies pose significant challenges, primarily due to late-stage enforcement and rigid technical requirements.

## Late Validation & Feedback

Policies run only in production, causing feedback delays (hours/days) and increased incident risk.

## Complex Environment Coupling

Requires a full Airflow infrastructure for validation, making pre-deployment checks complex and resource-intensive.

## Developer Experience Gaps

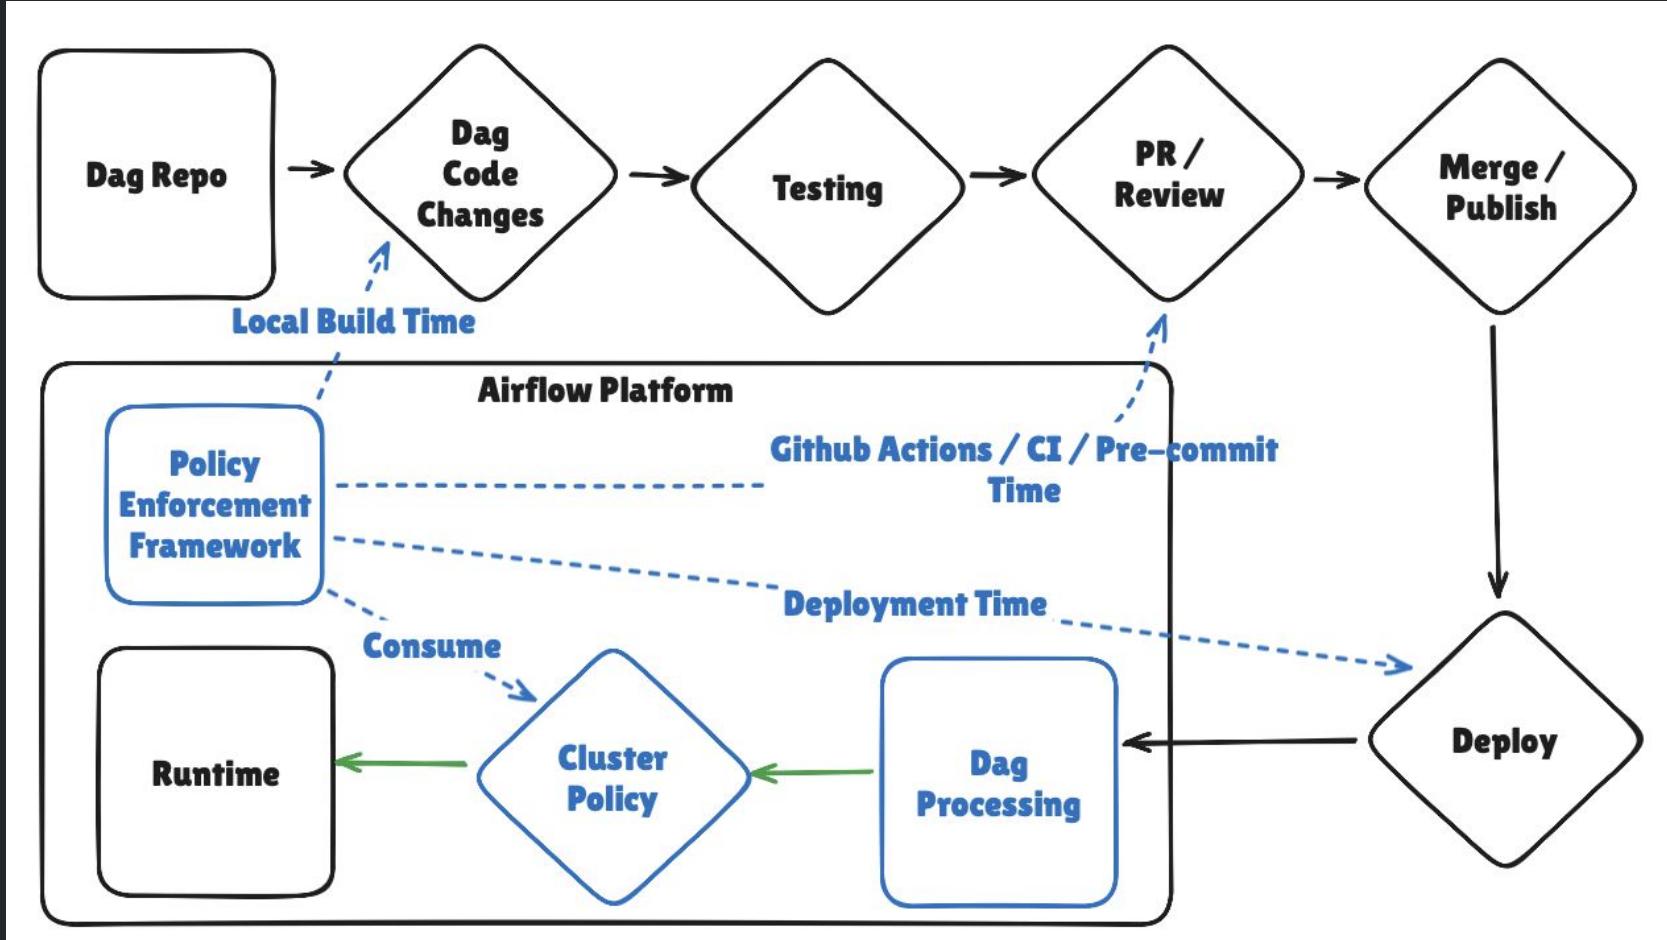
No easy local or CI/CD validation without mirroring a full Airflow environment, hindering early error detection.

## Limited Extensibility & Observability

Lacks built-in exemption systems, emergency overrides, and composable validation; failures are hard to diagnose and audit.

A full production Airflow deployment is needed for DAG validation, preventing "shift-left" policy enforcement into development and CI/CD.

# Future - No more Dag Import Errors in Production



# Technical Solution: Portable DAG Validation Framework

This lightweight framework shifts Airflow DAG policy enforcement left, enabling earlier, more flexible, and developer-friendly validation.

Decoupled Environment Interface  
**Validate DAGs anywhere via module injection, eliminating full Airflow environment requirements.**

Dynamic Dependency Resolution  
**Validate custom modules and proprietary libraries via runtime `sys.path` injection, no installation needed.**

Leveraged Native Error Detection  
**Utilize `DagBag's import_errors` for comprehensive, production-grade detection of all DAG errors.**

Resulting in fast, independent, and comprehensive validation with immediate developer feedback.

# The Policy Enforcement Engine

## Enforcement Workflow

```
class PolicyEnforcer:

    def enforce_policies(dag_repo_path, environment):

        # 1. Setup Environment
        self.setup_environment(...)

        # 2. Load DAGs (Airflow native)
        dagbag = DagBag(dag_repo_path)

        # 3. Apply Policies
        for dag in dagbag.dags.values()
            self.apply_dag_policies(dag)
```

This core loop ensures every DAG is tested against required operational standards before deployment.

## Policy Definition: Declarative Rules

```
@hookimpldef
dag_policy(dag):

    # Check ID format
    validate_dag_id_format(dag)

    # Enforce alerting
    ensure_alerting_configured(dag)

    # Verify ownership
    validate_owner_metadata(dag)

    # Enforce compliance
    enforce_compliance_rules(dag)
```

Policies are standard Python functions that access the full DAG object. They can validate, mutate, or reject deployment based on any configuration criteria.

- Composable Design: New rules can be added without modifying the core enforcement engine.

Error Type	Catch as Dag Import Error	Code Location
Missing dependency	Yes	_load_modules_from_file()
Syntax error	Yes	_load_modules_from_file()
Top-level exceptions	Yes	_load_modules_from_file()
Cycle detection	Yes	_process_modules() catch
DAG validation errors	Yes	_process_modules() catch
AirflowClusterPolicyViolation	Yes	_process_modules() catch
DAG ID collision	Yes	_process_modules() catch
Unknown executor	Yes	_process_modules() catch

# The Shift-Left Insight

## Runtime Discovery is Costly

Traditional governance enforces policies only in production. This is the worst time to find errors:

- Immediate failure and customer impact.
- Rollback is risky.
- Debugging is slow and complex.

100%

Runtime Cost

Maximum impact and risk

## The Shift-Left Advantage

Shift-Left Governance moves validation earlier in the lifecycle (Local & CI/CD). Catch issues when fixes are easiest and cheapest.

10%

CI/CD Cost

Caught before deployment

1%

Local Development Cost

Identified during coding

# Use Case #1: Missing Dag Parsing Dependency Errors

## The Problem: Environment Drift

DAGs that work locally often fail upon deployment due to critical differences in the production environment:

- Missing team-specific shared libraries
- Python version mismatches
- Conflicting package versions
- Missing system dependencies

Import failures cause DAGs to disappear from the Airflow UI, leading to broken workflows and late-stage incidents.

## The DAGnostics Solution

### CI Environment Simulation

Replicates production Python and system dependencies within the CI pipeline.

### Auto-Load Dependencies

Loads team-specific shared libraries using repository metadata.

### Validate Imports

Executes actual DAG imports using Airflow's native DagBag loader to guarantee success.

### Actionable Reporting

Surfaces import errors and full stack traces directly in the pull request.

0

Production Import Failures

Since implementing CI environment simulation

100%

Pre-Deployment Detection

All import issues caught during pull request validation

# Use Case #2: Ensuring DAG ID Uniqueness Identity & Ownership

## Challenges with 100,000+ DAGs

At massive scale (2,000+ repositories), maintaining unique identity and tracking ownership is critical. Manual coordination fails:

- ID collisions cause deployment failures.
- Incident response lacks immediate owner identification.
- Access control systems need verifiable ownership data.
- Compliance audits require clear accountability trails.

## Our Automated Enforcement

### 1 Standardized DAG ID Format

Enforced format: {dag\_name}-{repo\_name} (using globally unique repo name).

### 2 Automatic Metadata Sync

Ownership synced instantly from the central repository system.

### 3 Access Control Integration

Ownership feeds directly into permission systems for automated authorization.

### 4 Clear Audit Trail

Lineage tracked from DAG to repository to owning team for compliance.

100%

DAG Ownership Visibility

Every single DAG traceable to owning team

0

ID Collisions

Automatic format enforcement prevents conflicts

# Use Case #3: Alerting Policy Enforcement

## The Risk: Silent Failures

Critical data pipelines that fail without alerting represent one of the highest-impact operational risks.

When revenue-generating workflows break silently:

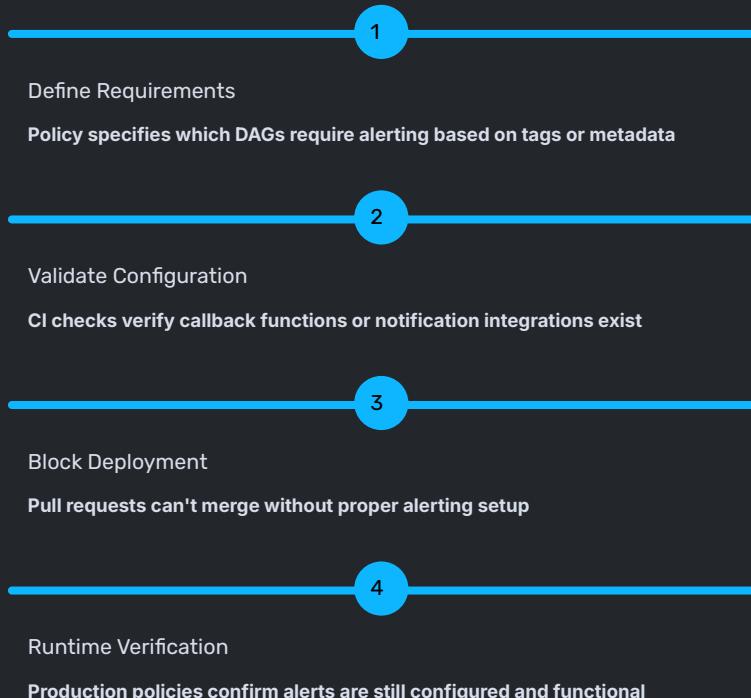
- Business metrics drift without warning
- Compliance deadlines are missed
- Customer-facing features degrade
- Problems compound before detection

Manual alerting configuration is prone to human error. Teams forget to add alerts, misconfigure integrations, or use inconsistent escalation paths.

## Our Policy Enforcement

DAGnostics enables making alerting a structural requirement, not an optional best practice:

- All production DAGs must define failure alerting
- Standardized timeout policies prevent infinite hangs
- Integration with centralized monitoring platforms
- Automatic escalation paths based on DAG criticality
- SLA monitoring for time-sensitive workflows



"Since enforcing alerting policies, we've eliminated an entire class of incidents where critical workflows failed unnoticed for hours or days."

## Questions?

Ask us About Building Airflow at Scale

Our journey scaled Airflow to 100,000+ DAGs. Shift-left governance enabled high developer velocity at enterprise scale.

---

LinkedIn

[linkedin.com/in/stefanwang](https://linkedin.com/in/stefanwang)

Let's Connect:

