



Deadline Alerts In Airflow 3.1

Dennis Ferruzzi | OSS Software Engineer @ AWS

Ramit Kataria | OSS Software Engineer @ AWS





Beyond SLA

Deadline Alerts in Airflow 3.1

3.1



Dennis Ferruzzi Ramit Kataria



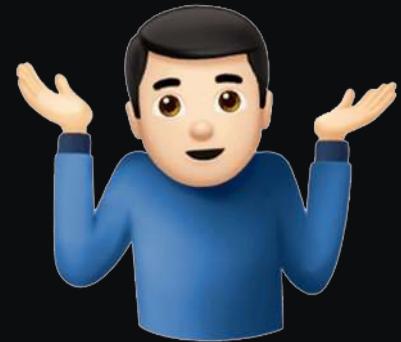
What was an SLA?

In Airflow 2, an SLA allowed users to set a duration for a Dag or Task and get a notification if they ran longer.

What is a Deadline Alert?

In Airflow 3, Deadline Alerts allow you to set time thresholds for your Dag runs, and automatically respond when those thresholds are exceeded.

That sounds awfully
similar, what's the
difference?



SLA

- `sla_miss` was evaluated **only** if the task_id was ever in SUCCESS or SKIPPED state
- SLA was defined as a timedelta relative to the `dagrun.data_interval_end`
- Callback was an attribute of the **Dag**, but the SLA was an attribute of individual **tasks**
- Callback was executed after the Dag run finished
- Dag files were parsed **each time** the callback was sent to the `DagFileProcessor`
- `DagFileProcessor` was overloaded

Deadline Alerts

- Evaluated every scheduler pass
- You can define the **Reference** (starting point) and **Interval** (timedelta)
- Callback is an attribute of the **Deadline**
- Callback is executed within seconds of the missed deadline
- Deadline is calculated at Dag run creation and stored, no need for multiple passes
- `DagFileProcessor` can focus on its job

SLA

Deadline Alerts



SLA

User defines a length of time. When (if?) the Dag run finishes: if it took longer than that amount of time, run the callback.

More flexible and timely!

Deadline Alerts

User defines a length of time and a reference point from which to start counting. When a Dag run is created, calculate that expiration time.

Each scheduler pass (5 seconds by default), if that time has passed: run the callback, even if the Dag is still currently running.

I'm Sold!

How do Deadlines Work?



How is a Deadline Alert Defined?

Deadline Alert

Reference

Interval

Callback

When to start counting

How long to wait

How to respond



What Does This Look Like?

If the DAGRUN has not finished
[Interval] 30 minutes after
[Reference] the DAGRUN_QUEUED_AT

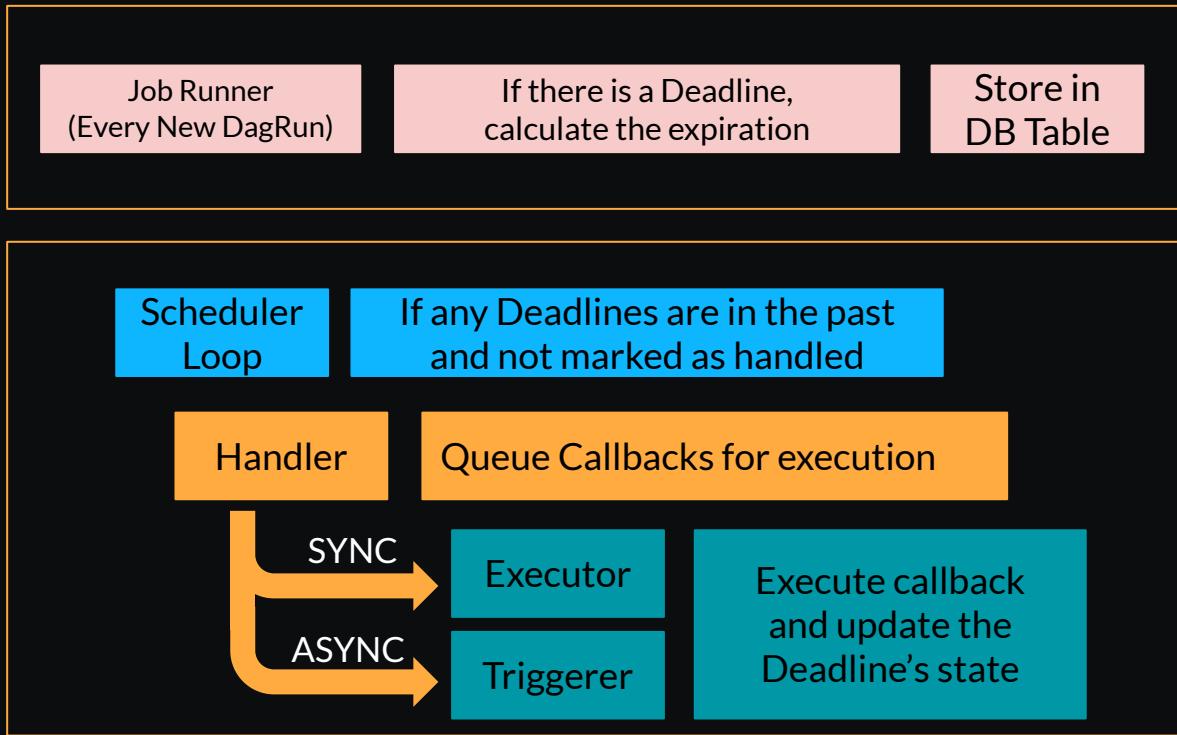
Then
[Callback] send a slack message

If the DAGRUN has still not finished
[Interval] 60 minutes after
[Reference] the DAGRUN_QUEUED_AT

Then
[Callback] send an email message

```
SLACK_TEAM = AsyncCallback(  
    SlackWebhookNotifier,  
    kwargs: {"text": "⚠ {{ dag_run.dag_id }} is running late."}  
)  
  
EMAIL_ONCALL = AsyncCallback(  
    callback_callable=SmtpNotifier,  
    kwargs={  
        "to": ONCALL_ADDRESS,  
        "subject": "⚠ {{ dag_run.dag_id }} missed deadline at {{ deadline.deadline_time }}.",  
        "html_content": "Dag Run details: {{ dag_run }}",  
    }  
)  
  
with DAG(  
    dag_id="deadline_alerts_demo",  
    deadline=[  
        DeadlineAlert(  
            reference=DeadlineReference.DAGRUN_QUEUED_AT,  
            interval=timedelta(minutes=30),  
            callback=SLACK_TEAM,  
        ),  
        DeadlineAlert(  
            reference=DeadlineReference.DAGRUN_QUEUED_AT,  
            interval=timedelta(minutes=60),  
            callback=EMAIL_ONCALL,  
        )  
    ]:  
    task1()
```

How Does It Work?



Built-in Deadline References

DAGRUN_QUEUED_AT

Measures time from when the DagRun was queued.

Useful for monitoring resource constraints.

DAGRUN_LOGICAL_DATE

Measures time from when the Dag run was scheduled to start.

Useful for ensuring scheduled Dags complete before their next scheduled run.

FIXED_DATETIME

Specifies a fixed point in time.

Useful when Dags must complete by a specific time.

AVERAGE_RUNTIME

Calculates the average historical runtime of the Dag.

Useful for detecting unusual activity in a Dag run or environment.

Callback Support

Asynchronous Callbacks
(Available in 3.1)

Synchronous Callbacks
(Coming in 3.2)

Run in the Triggerer

Run in the Executor/worker

Existing Notifiers - Async Notifiers Work

Custom Callbacks - Must be in the
Triggerer's sys.path (for example, the
Plugins folder)

Existing Notifiers - All Work

Custom Callbacks - Can be placed
anywhere in the Dag Bundle

Async vs Sync Callbacks

Asynchronous Callbacks (Available in 3.1)

Lower runtime overhead - runs on triggerer

Requires restarting the triggerer to apply changes in callback

Requires async callable

Synchronous Callbacks (Coming in 3.2)

Higher runtime overhead - runs on worker/executor

Automatically uses the callback definition from the Dag bundle

Runs any python callable

Callback Support

Asynchronous Callbacks

Existing Notifiers with async support:

Slack, Email, AWS (SES, SNS, SQS) + more
coming soon

Custom Callbacks - Must be in the
Triggerer's sys.path (for example, the
Plugins folder)

Synchronous Callbacks

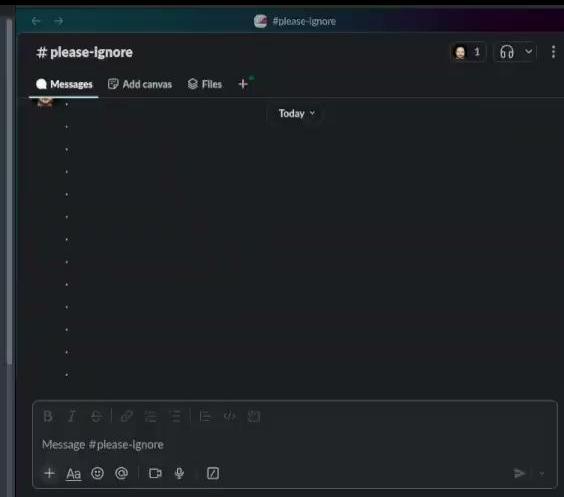
All existing Notifiers

Custom Callbacks - Can be placed
anywhere in the Dag Bundle

DEMO!

3.0

```
1 with DAG("slack_as_on_success_callback", tags=["slack"]):
2     BashOperator(
3         task_id='example_task',
4         bash_command="echo 'Hello World!'", 
5         on_success_callback=[
6             SlackWebhookNotifier(
7                 text=f"On_success callback!",
8             )
9         ],
10    )
```

A screenshot of the Airflow web interface at the URL <http://localhost:28080/dags>. The page displays a list of DAGs. The 'Dags' tab is selected. A search bar is present above the table. The table has columns: Dag ID, Schedule, Next Run, Latest Run, and Tags. Three DAGs are listed: 'slack_as_deadline_should_not_trigger', 'slack_as_deadline_should_trigger', and 'slack_as_on_success_callback'. All three DAGs are currently in the 'Success' state, indicated by green checkmark icons in the 'Required Actions' column. The 'Tags' column shows 'slack' for all three DAGs. The 'Admin' sidebar on the left is visible, containing links for Home, Dags, Assets, Browse, Admin, Docs, and User.

What's To Come?

Available in 3.1

- DAG-level Deadlines
 - Dagrun: Started
 - Dagrun: Queued
 - Fixed datetime (every day at 9AM)
 - Average runtime
- Multiple Deadlines per DAG
- Async Callbacks and Notifiers
 - Executed by the Triggerer

Future Work

- Synchronous Callbacks
 - Pick your Executor!
- Task-Level Deadlines
- Expand trigger options:
 - Dataset: Created
 - Dataset: Updated
 - Asset-Driven
 - ???
- Add more Asynchronous Notifiers

Questions?

3.0



**Get
Involved**



@Ferruzzi



@Ramit Kataria



Slide Deck