



# Enhancing Airflow REST API

## From Basic Integration to Enterprise Scale

Vishal Vijayvargiya

Sr. Software Engineer - AWS

3.0



# What is the Airflow REST API?

- RESTful interface provided by Apache Airflow
- Enables programmatic control of Airflow:
  - Trigger DAG runs
  - Pause/unpause DAGs
  - Manage connections, variables, and pools
- Why it matters:
  - Allows Airflow to integrate with external systems and automation tools
  - Forms the foundation for workflow orchestration at scale



## Example

- Trigger a DAG run

```
import requests
webserver_url = "<webserver-url>"
token = "<access-token>"
dag_id = "<dag-id>"
response = requests.post(
    url=f"https://{{webserver_url}}/api/v2/dags/{{dag_id}}/dagRuns",
    headers={
        "Authorization": f"Bearer {{token}}",
        "Content-Type": "application/json"
    },
    json={"logical_date": "2025-10-06T14:15:00Z"}
)
print(response.json())
```



## Why Vanilla REST API Isn't Enough for Enterprises

- Limited security → hard to integrate with IAM
- Exposed webserver → networking/security risks
- Hard to scale across multiple teams/orgs → inconsistent patterns
- Limited audit/compliance visibility → no centralized logs



## Amazon MWAA InvokeRestAPI

- Simplifies calling Airflow REST API endpoints securely
- Works across:
  - AWS CLI (`aws mwaa invoke-rest-api`)
  - SDKs (boto3, etc.)
  - Cloud integrations (Step Functions, Lambda, CI/CD pipelines)
  - MwaaDagRunSensor and MwaaTaskRunSensor (`apache-airflow-providers-amazon`)



## Request Syntax

```
POST /restapi/Name HTTP/1.1
Content-type: application/json

{
    "Body": JSON value,
    "Method": "string",
    "Path": "string",
    "QueryParameters": JSON value
}
```



## Response Syntax

HTTP/1.1 200

Content-type: application/json

```
{  
    "RestApiResponse": JSON value,  
    "RestApiStatusCode": number  
}
```



## Example

- Trigger a DAG run

```
import boto3
import json

client = boto3.client("mwaa")
mwaa_env = "mwaa-env-name"
dag_id = "<dag-id>"

response = client.invoke_rest_api(
    Name= mwaa_env,
    Path=f"/dags/{dag_id}/dagRuns",
    Method="POST",
    Body=json.dumps({"logical_date": "2025-10-06T10:00:00Z"})
)

print(response["RestApiResponse"])
```

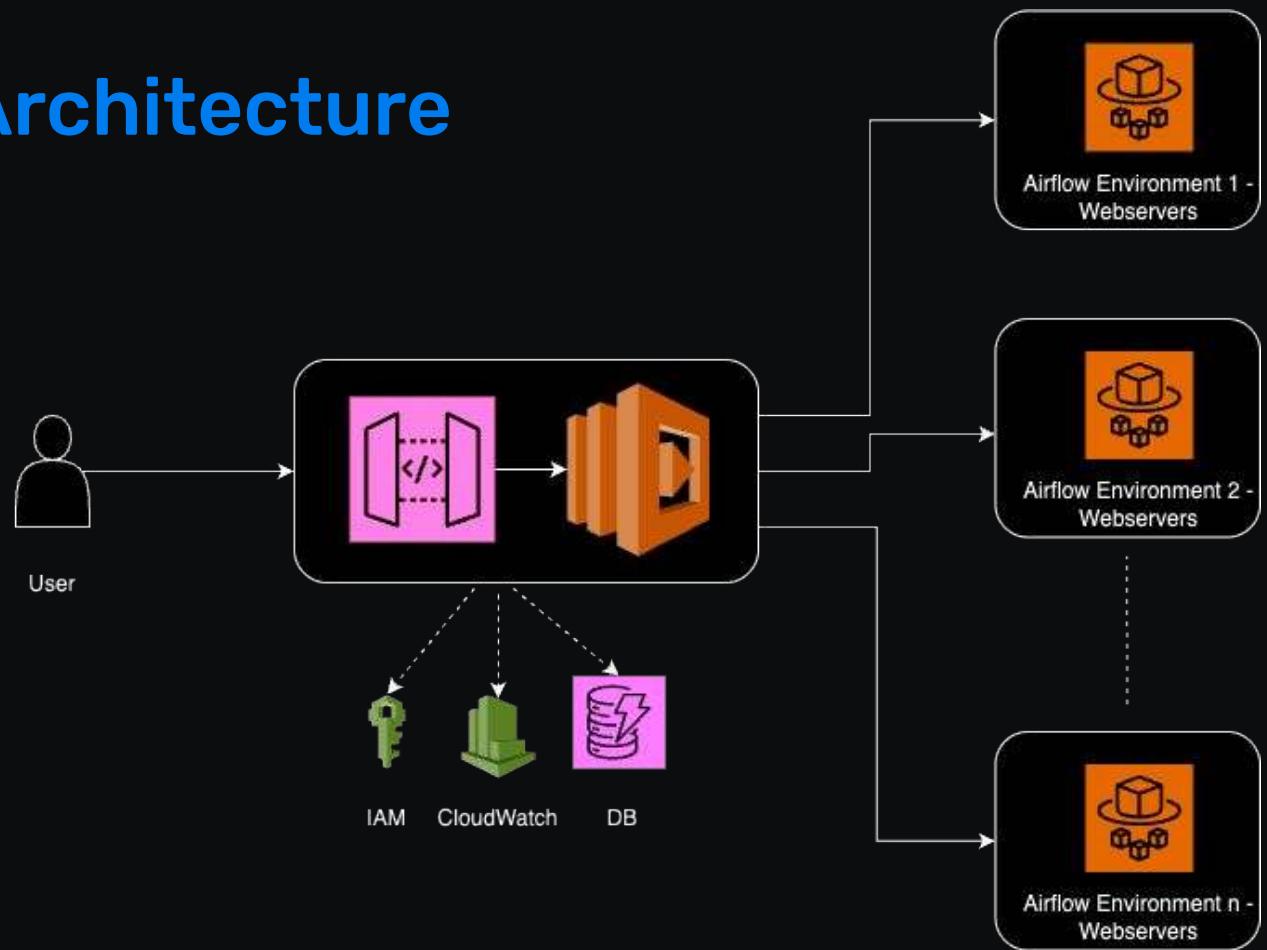


# IAM Permissions

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Sid": "AllowMwaaRestApiAccess",  
      "Effect": "Allow",  
      "Action": "airflow:InvokeRestApi",  
      "Resource": [  
        "arn:aws:airflow:us-east-1:111122223333:role/{your-environment-name}/{airflow-role}"  
      ]  
    }  
  ]  
}
```

# High Level Architecture

- Considerations:
  - Webserver
  - Autoscaling
  - CPU Utilization and Active Connection Count
  - Error Handling





## CI/CD for Data Pipelines

- Automate DAG validation & deployment via CLI/SDK

## ⚡ Event-Driven Workflows (Data Ingestion + ML Retraining)

- Trigger pipelines securely from SaaS apps
- Secure, auditable event-based triggers



## Cross-Service Orchestration

- Step Functions + Airflow
- Lambda + Airflow
- Hybrid Orchestration (On-prem + Cloud)



## Audit & Compliance

- CloudTrail logging of all API calls
- Centralized monitoring dashboards

# Example

# 3.0



## Automated Post-Deployment Health & DAG Validation

- An organization deploys Airflow DAGs frequently through automated CI/CD pipelines.
- After every DAG deployment to MWAA, teams need to ensure:
  - Airflow components are **healthy** (scheduler, triggerer, database).
  - All expected DAGs are **parsed and visible** in Airflow.
- There's **no built-in mechanism** to detect silent DAG parsing failures.
- Teams need an **automated, secure, IAM-based validation** process that runs immediately after each deployment.



## Automated Health & DAG Validation Using MWAA InvokeRestApi

- Use AWS Lambda triggered after each deployment to:
  - Call /monitor/health → ensure Airflow components are operational
  - Call /dags → verify all expected DAGs are successfully parsed
- Fails pipeline automatically if any validation step fails
- Fully IAM-authenticated, no public exposure
- CloudTrail-logged for audit and compliance

```
print(f"Checking MWAA environment: {mwaa_env}")

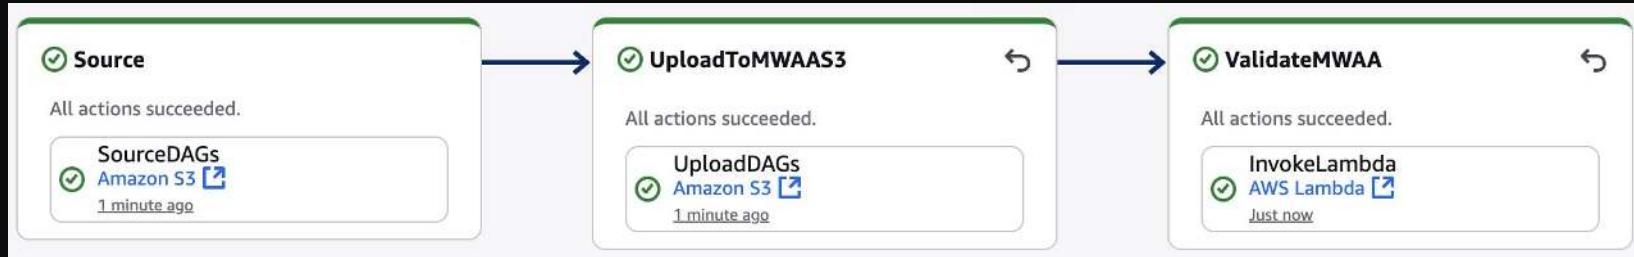
# Check Airflow Health
health_resp = client.invoke_rest_api(Name=mwaa_env, Path="/monitor/health", Method="GET")
health = health_resp["RestApiResponse"]
print("Health Response:", json.dumps(health, indent=2))

if not all(x["status"] == "healthy" for x in health.values()):
    raise Exception(f"Unhealthy Airflow components: {health}")

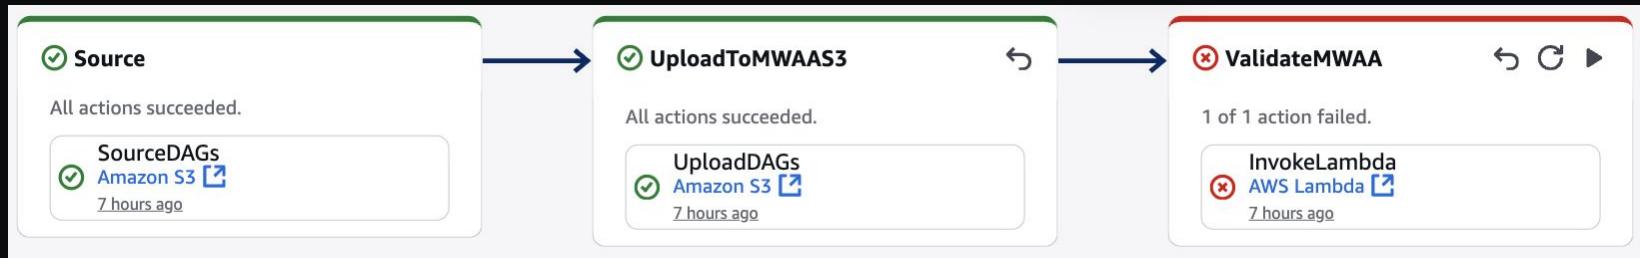
# Check DAGs
dags_resp = client.invoke_rest_api(Name=mwaa_env, Path="/dags", Method="GET")
dags_json = dags_resp["RestApiResponse"]
airflow_dags = [d["dag_id"] for d in dags_json.get("dags", [])]
print(f"Airflow DAGs found: {airflow_dags}")

missing = [d for d in expected_dags if d not in airflow_dags]
if missing:
    raise Exception(f"Missing DAGs: {missing}")

print("MWAA validation successful – environment healthy and DAGs loaded.")
```



```
2025-10-07T22:33:38.3... Checking MWAA environment: pinwheel-test-environment-demo-3-0-6
2025-10-07T22:33:38.5... Health Response: {
2025-10-07T22:33:38.5...   "metadatabase": {
2025-10-07T22:33:38.5...     "status": "healthy"
2025-10-07T22:33:38.5...   },
2025-10-07T22:33:38.5...   "scheduler": {
2025-10-07T22:33:38.5...     "status": "healthy",
2025-10-07T22:33:38.5...     "latest_scheduler_heartbeat": "2025-10-07T22:33:29.946862+00:00"
2025-10-07T22:33:38.5...   },
2025-10-07T22:33:38.5...   "triggerer": {
2025-10-07T22:33:38.5...     "status": "healthy",
2025-10-07T22:33:38.5...     "latest_triggerer_heartbeat": "2025-10-07T22:33:28.754060+00:00"
2025-10-07T22:33:38.5...   },
2025-10-07T22:33:38.5...   "dag_processor": {
2025-10-07T22:33:38.5...     "status": "healthy",
2025-10-07T22:33:38.5...     "latest_dag_processor_heartbeat": "2025-10-07T22:33:33.280598+00:00"
2025-10-07T22:33:38.5...   }
2025-10-07T22:33:38.5...
2025-10-07T22:33:38.7... Airflow DAGs found: ['basic_bash_dag', 'dynamic_task_mapping_example']
2025-10-07T22:33:38.7... MWAA validation successful - environment healthy and DAGs loaded.
```



### Execution summary

Status	Started	Completed	Duration
Failed	7 hours ago	7 hours ago	34 seconds

**Trigger**  
PollForSourceChanges - SourceDAGs

**Pipeline execution ID**  
[406432f5-94f3-4c94-87c7-35046ae6c577](#)

Latest action execution message  
Missing DAGs: ['basic\_bash\_dag']  
[Link to execution details](#)

[Diagnose with Amazon Q](#)

# CloudTrail Logs

```
▼ 1:
    eventVersion:          "1.09"
    ▼ userIdentity:
        type:                "AssumedRole"
        principalId:         "/[REDACTED]:DeploymentValidation"
        arn:                 "arn:aws:sts::[REDACTED]:assumed-role/DeploymentValidation-role-yacx15gv/DeploymentValidation"
        accountId:           "[REDACTED]"
        accessKeyId:          "[REDACTED]"
    ▼ sessionContext:
        ▼ sessionIssuer:
            type:              "Role"
            principalId:       "/[REDACTED]"
            arn:               "arn:aws:iam::[REDACTED]:role/service-role/DeploymentValidation-role-yacx15gv"
            accountId:          "[REDACTED]"
            userName:           "DeploymentValidation-role-yacx15gv"
        ▼ attributes:
            creationDate:       "2025-10-06T17:28:58Z"
            mfaAuthenticated:   "false"
            eventTime:          "2025-10-06T17:29:00Z"
            eventSource:         "airflow.amazonaws.com"
            eventName:          "InvokeRestApi"
            awsRegion:          "us-west-2"
            sourceIPAddress:    "[REDACTED]"
            userAgent:          "Boto3/1.40.4 md/Botocore#1.40.4 ua/2.1 os/linux#5.10.242-265.962.amzn2.x86_64 md/arch#x86_64 lang/python#legacy Botocore/1.40.4"
        ▼ requestParameters:
            Path:               "/monitor/health"
            Method:             "GET"
            Name:               "pinwheel-test-environment-3-0-6"
        ▼ responseElements:
            Access-Control-Expose-Headers: "x-amzn-RequestId,x-amzn-ErrorType"
            RestApiStatusCode:            200
            RestApiResponse:             "***"
```



## Closing: Airflow 3

- InvokeRestAPI usage remains fully consistent across Airflow 2 and 3 – no breaking changes
- This abstraction layer shields enterprise workflows from backend changes.

# Questions?

[linkedin.com/in/vishalvijayvargiya/](https://www.linkedin.com/in/vishalvijayvargiya/)