



From Centralization to Autonomy

Managing Airflow Pipeline through Multi-Tenancy

SILVER PANG

Software Developer @ Expedia Group

3.0

Technology landscape keep shifting



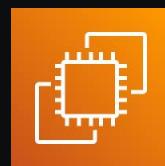
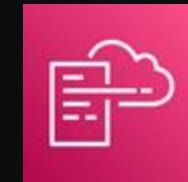
Organizational Challenges

Platform Engineering

**Design, Build, Maintain the shared
infrastructure**

Time Travel back
2018

Own its own tech stack



The Enterprise Dilemma: When Autonomy Creates Chaos.

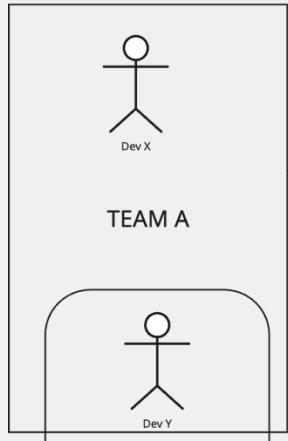
- Bottlenecks & Slowed Development
- Inconsistent Standards
- High Operational Overhead for every team



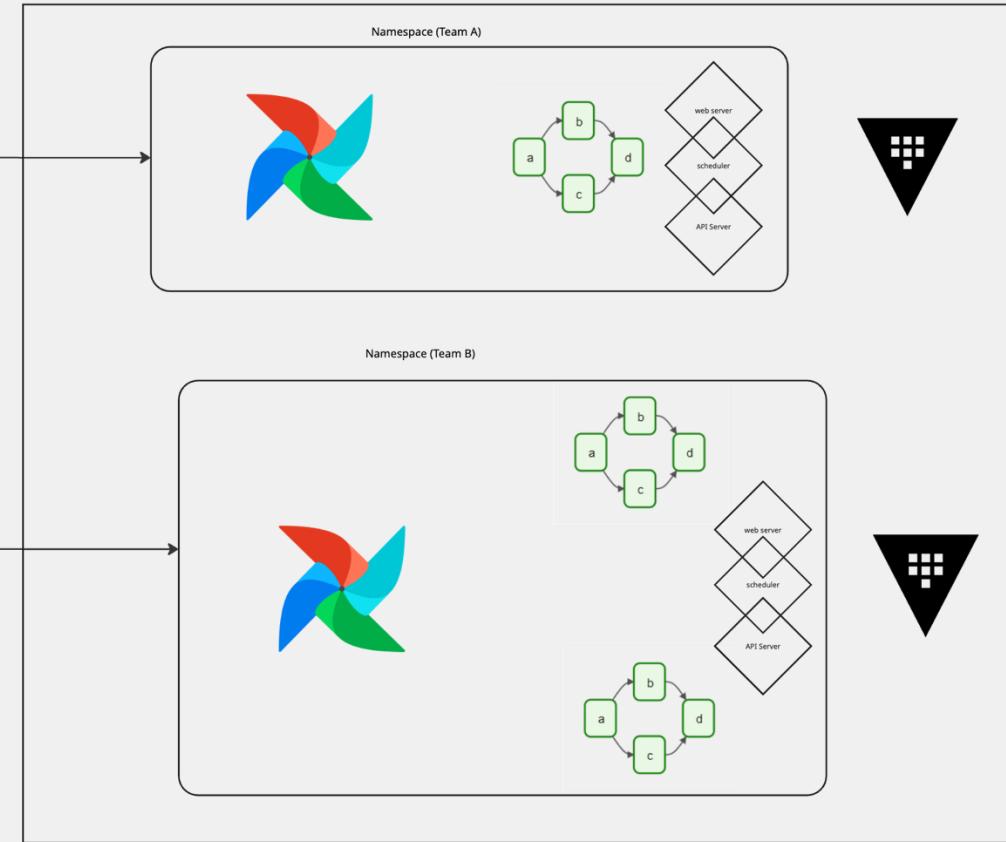
Platform Engineering stepped in

Providing toolkit
Data Processing Framework
Batch / Stream compute / Workflow Orchestration





k8s cluster



```
module "client-seattle" {
    source: "git@github.com:expedia/tf-airflow",
    client_name: "seattle-data",
    application_name: "workflow-orchestration",
    airflow_admin_security_group: "settle-admin-group",
    airflow_version: "2.10.5",
    airflow_docker_image_tag: "0.0.1",
    cluster_name: var.cluster_name,
    datadog_api_aws_secret_name:
        var.datadog_api_key_name,
    vault_secret_key: var.vault_secret_key,
    cost_center: "12345",
    environment: var.environment,
    openlineage_enabled: true
}
```



```
apache-airflow-providers-amazon==8.12.0
apache-airflow-providers-datadog==3.4.0
awscli
boto3
datadog
trino
pandas
```

The DevEx | Templates for speed When Autonomy Creates Chaos.

// TODO add vault



From zero to pipeline in “minutes”



1

Provide Application Metadata

2

Provide additional information

3

Include TechDocs

4

Review

This data is used to register your application in AMS

GitHub Instance *

Github Organization *

The GitHub Organization where the repository is located. Verify the instance above is correct for the organization chosen.

Application Name *

Name of the application. Application name should only contain characters [a - z], [0 - 9] and hyphens. Maximum number of characters allowed is 63.

Description *

Describe this service. Help others understand what this service does. What problem does it solve? What benefits does it provide? When should a person definitely use this template?

Team Name *

Name of the support team for this application service. [More details can be found here](#)

Documentation Overview Link

A link to the Overview Documentation for the Application. [More details can be found here](#)

Application Type *

Type of Application. [More details can be found here](#)

PCI Category *

PCI Category of the application. [More details can be found here](#)

Environment Scope *

Environment Scope of the application. [More details can be found here](#)

Support Slack Channel

Please include the public support Slack channel of your team for communications. Do not include # in the name, for example: my-support-channel

Funnel

The Funnel attribute is part of the PagerDuty Information for Incidents. This field is mandatory for all the Tier 0 and Tier 1 Applications. [More details can be found here](#)

Application Lifetime *

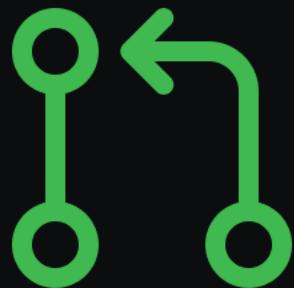
Application Lifetime specifying if the application is temporary or permanent. Choose Temporary if your application belongs to demo/hello world/bootcamp or other limited shelf-life categories. If not, choose permanent. [More details can be found here](#)

If you choose Temporary and it meets the criteria for automated decommissioning, the app will be automatically decommissioned. For more details, please refer to [FAQs on Automated App Deactivation](#).

Email *

Fully Automated CI/CD pipeline

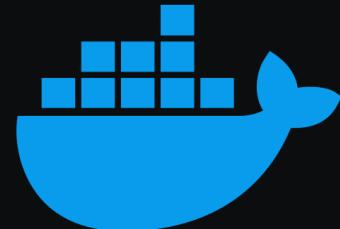
Fully Automated CI/CD pipeline



Fully Automated CI/CD pipeline



Fully Automated CI/CD pipeline



Fully Automated CI/CD pipeline



JFrog Artifactory

Fully Automated CI/CD pipeline



Fully Automated CI/CD pipeline



Fully Automated CI/CD pipeline



Fully Automated CI/CD pipeline

 All checks have passed Hide all checks

8 successful checks

	CI-branch-pr / build (pull_request)	Successful in 3m — build	Required	Details
	GHAS-CodeQL-Scan / Analyze (python) (pull_request)	Successful in 1m — Analyze...	Details	
	Sonarqube-SAST-Scan / SonarQube (pull_request)	Successful in 1m	Details	
	Code scanning results / CodeQL	Successful in 2s — No new alerts in code changed ...	Details	
	Sonarqube-SAST-Scan / JUnit Test Report (pull_request)	— 73 tests run, 0 skipped... Required	Details	
	pipeline/branch	— Pipeline 'branch' is complete	Details	

20+ New Pipelines per Month

200+ Airflow Clusters

180+ Engineering Teams

14k+ Active DAGs

1.5M+ Task Executions per Month

**Reduce Airflow cluster
deployment time**

**Reduce time for DAG
deployment**

Easier Migration

Hard Lesson Learned

Code Consistency
Compliance & Governance
The Learning Curve



05e6 in dw

```
{  
    "instance_key": "████████",  
    "uuid": "████████05e6",  
    "name": "████████-spark-job",  
    "application_name": "████████-migrator-spark",
```

[Show more](#)

To access [marketing-████████](#) Airflow instance, Ensure you are a member of the security group [-marketing-████████](#) using [Group Management](#)

Health:

- metadatabase: healthy
- scheduler: healthy
- webui: healthy

Config:

- Airflow version: 2.10.5
- Data █████ Operator version: 2.0.9

[See less](#)

Key Takeaway

Platform is the riverbed
Autonomy is flow
Adaptability is the current



From Centralization to Autonomy

Managing Airflow Pipeline through Multi-Tenancy

SILVER PANG
Software Developer @ Expedia Group

3.0