



# From Repetition to Refactor: Smarter DAG Design in Airflow 3

Steven Woods  
Director, Software Engineering  
Zelis

3.0

# Agenda: From Repetition to Refactor

## Initial DAG Design

- Common anti-patterns: hardcoded logic, duplication, rigid sequencing

## Refactor Strategy

- Applying D.R.Y. principles
- Task factories, parameterization, dynamic task mapping

## Designing for Flexibility

- Modular DAGs for batch, streaming, and ad-hoc workflows

## Pros & Cons of Refactor

- Benefits: scalability, maintainability, observability
- Trade-offs: complexity, learning curve

## Q&A

# Initial DAG Design

 **Project:** Animal image processor

## Description:

Develop an automated and scalable workflow that retrieves animal images from the internet and processes them for downstream use.

## Objective:

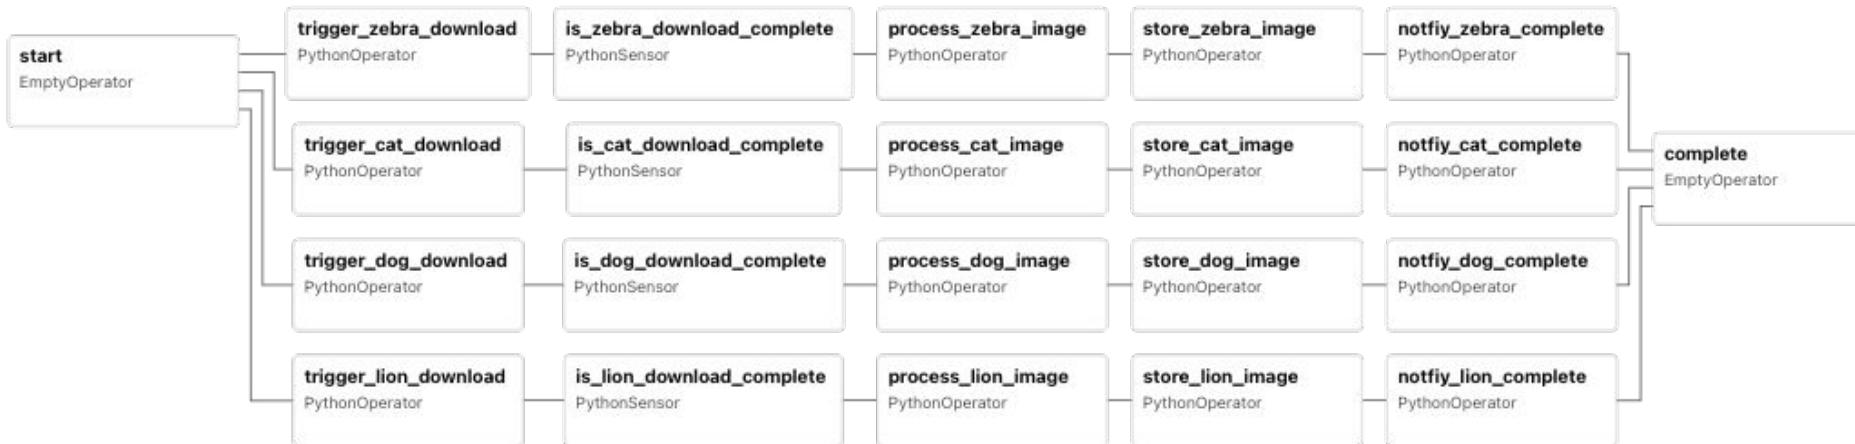
- Fetch a new animal image daily from a designated internet source
- Target animals: cat, zebra, dog, lion
- Store the image in an object storage service (e.g., Amazon S3)
- Resize to standardized dimensions & convert image



# Initial DAG Design

## Tasks

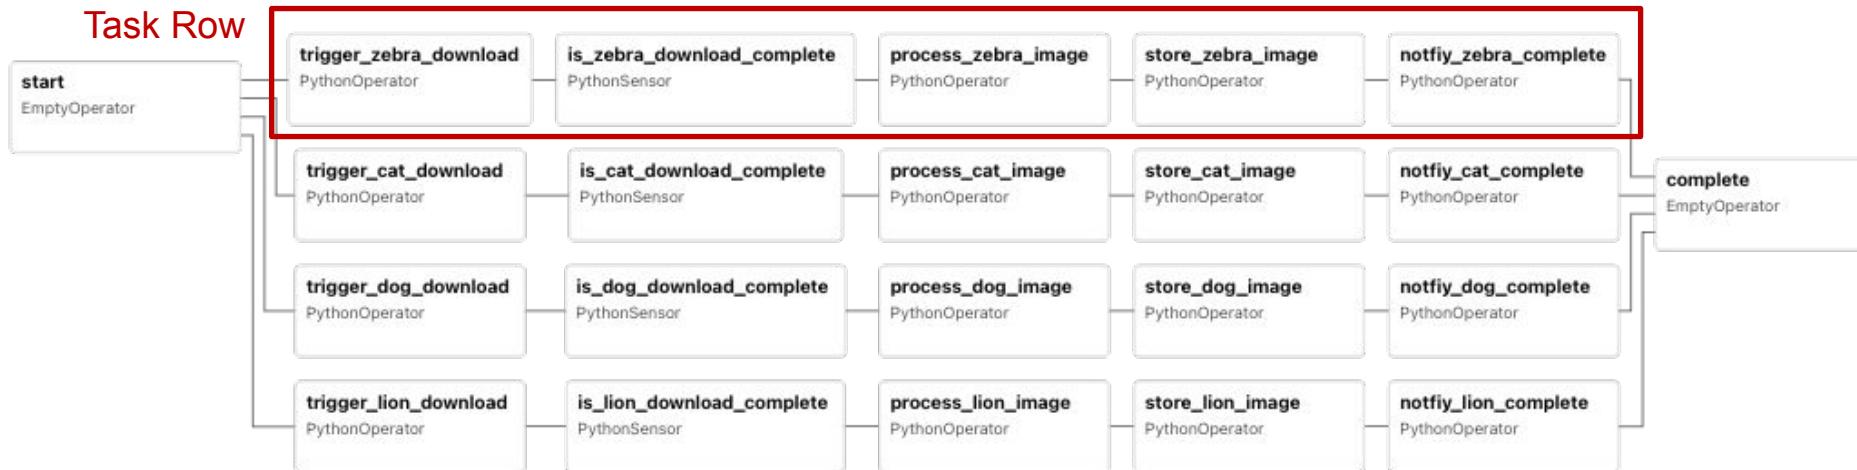
- start – Workflow entry point
- trigger\_(animal)\_download – Trigger Lambda/external process to fetch animal photos
- is\_(animal)\_download\_complete – Sensor to check if download finished (file present in S3)
- process\_(animal)\_image – Resize and convert the image
- store\_(animal)\_image – Save image metadata to the database
- notify\_(animal)\_complete – Send completion notification email
- complete – Workflow end point



# Initial DAG Design

I'll call this design “**repetitive task rows**.” It has several scalability and management issues:

- More rows = more code (PRs, deployments)
- Hard to run concurrently or isolate failures
- Doesn't scale (4 animals → 10, 20, 100+)
- Graph becomes unreadable
- Code harder to maintain
- One failing task row can break the whole DAG



# Initial DAG Design

Each new animal adds 5+ new tasks  
(download, check, process, store, notify)

Code grows linearly with every animal →  
high maintenance overhead

Hard to isolate and retry failures for a  
single animal

Increases DAG execution time and  
complexity

DAG visual quickly clutters with repetitive  
task rows



# Initial DAG Design

- Failure in cat or dog tasks → entire DAG run fails
- With DAG concurrency = 1:
  - Next DAG run blocked until current completes
  - Persistent failures = major processing slowdown



# Refactor Strategy

Recommended solution: Split into 2 DAGs:

## DAG 1: Trigger

- Runs on a schedule (e.g., @daily)
- Reads configuration from Airflow variables
- Passes parameters to the processing DAG
- Simple to enable/disable

## DAG 2: Processor

- Accepts parameters (e.g., { "animal": "zebra" })
- Has no schedule – runs only when triggered
- Can be triggered by:
  - DAG 1 (Trigger DAG)
  - Manual runs with params
  - External processes
- Supports ad-hoc requests and flexible processing

Variable:

[

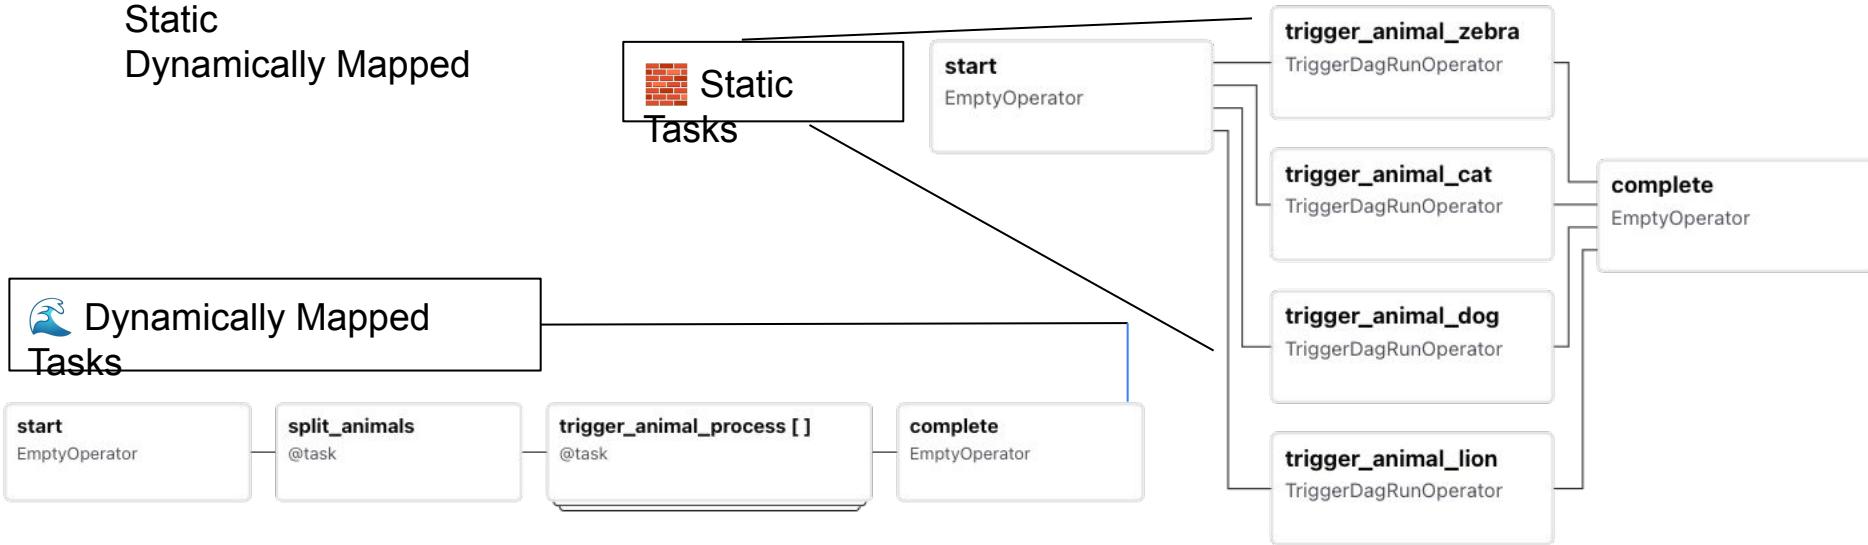
```
{"animal": "zebra"},  
 {"animal": "cat"},  
 {"animal": "dog"},  
 {"animal": "bear"}
```

]

# Refactor Strategy

## DAG 1: Trigger

- Two execution modes:  
**Group processing** → TriggerDagRunOperator(wait\_for\_completion=True)  
**Individual processing** → TriggerDagRunOperator(wait\_for\_completion=False)
- Supports running **multiple configurations** as needed  
Single  
Static  
Dynamically Mapped



# Refactor Strategy

## DAG 2: Processor

- Responsible for processing configs (accepts parameters, e.g., { "animal": "zebra" })
- No schedule – runs only when triggered
- Can be started by:
  - DAG 1 (Trigger DAG)
  - Manual runs with parameters
  - External processes
- Supports ad-hoc requests
- Easy to extend – add new configs via Airflow variables or code (no new tasks/operators required)
- Flexible & dynamic – avoids hardcoding (e.g., per-animal rules), focuses on runtime config processing



# Designing for Flexibility

The design supports a diverse range of workflows:



## Batch

- All animal downloads run on a fixed nightly schedule
- Processes the full set of animals in one batch (e.g., zebra, lion, cat, dog)
- Best for use cases where freshness isn't critical and daily updates are sufficient



## Streaming

- A data pipeline publishes animal IDs (e.g., {"animal": "zebra"}, {"animal": "cat"}) to Kafka/Kinesis. Each event tells your DAG which animal photo to download.



## Ad-hoc

- A simple website allows users to select an animal (e.g., 🦒, 🐶, 🐱)
- When triggered, the site calls Airflow's API to start the download + processing DAG
- No fixed schedule — runs only when requested

# Pros & Cons of Refactor

## Pros of Updated Design

**Scalable** – Add or remove animals by simply updating a variable (no new code required)

**Optimized execution** – Integrated use of DAG concurrency, task concurrency, and parallelism improves workload control and performance

**Separation of responsibility** – Each DAG focuses on its own role, making the codebase cleaner and easier to manage

**No code changes needed** – No PRs, reviews, or deployments for adding configs

**Simplified troubleshooting** – Issues with one animal don't block others; use Grid View to quickly spot and fix failures

**Flexible execution** – Run configurations manually, externally triggered, in bulk, or one at a time

## Cons of Updated Design

**Split codebase** – Logic is separated across two DAGs (Trigger & Processor)

**Limited validation** – New “animals” are added via variables without code review

**Fragmented visibility** – No single DAG view of the entire run; execution is split across DAGs

# Questions?

[linkedin.com/in/stevengwoods](https://linkedin.com/in/stevengwoods)