



Implementing Airflow Governance With Cluster Policies

Karen Braganza
Senior Airflow Reliability Engineer
@ Astronomer

3.0

Introduction to cluster policies

- Defined as Python functions
- Check or mutate Airflow objects.
- Enforce standards
- Useful for governance.

Types

- dag policy
- task policy
- task instance mutation hook
- pod mutation hook

Dag Policies

- Applied to a dag object
- When it is loaded in the dagbag.
- Check and mutate parameters in the [dag definition](#).

Governance Use Cases

- dag_id naming convention
- dag description, tags, owner_links
- dagrun_timeout

The screenshot shows the Airflow web interface with a dark theme. On the left, a sidebar contains navigation links: Home, Dags, Assets, Browse, and Admin. The 'Dags' link is currently selected, indicated by a green background. The main content area displays a summary of failed DAGs (0), favorite DAGs (0), and various metrics for the last 24 hours. A prominent modal dialog is centered over the page, titled "Dag Import Error". The dialog details an error from a bundle named "dags-folder" containing "dag_1.py". The timestamp of the error is 2025-10-03 21:50:50. The specific error message is "AirflowClusterPolicyViolation: If catchup is enabled, max_active_runs must be set to 1." At the top right of the modal is a search bar labeled "Search by file" with a keyboard shortcut "⌘+K". To the right of the modal, a green progress bar is partially visible, labeled "Manage Pools".

Airflow

Home

Dags

Assets

Browse

Admin

Stats

Failed

First 10 favorite Dags

No favorites yet.

Health

MetaDatabases

History

Last 24 Hours

Dag Runs

Queued

Running

0 Failed

0 First 10 favorite Dags

0 Health

0 MetaDatabases

0 Asset Events

0 Queued

0 Running

0% 0% 0%

Dag Import Error

Bundle Name: dags-folder dag_1.py

Timestamp: 2025-10-03 21:50:50

AirflowClusterPolicyViolation: If catchup is enabled, max_active_runs must be set to 1.

Manage Pools

Search by file ⌘+K



DAGs

Cluster Activity

Datasets

Security ▾

Browse ▾

Admin ▾

Docs ▾

Astronomer ▾



22:02 UTC ▾



1 DAG Import Errors (1)

Broken DAG: [/usr/local/airflow/dags/dag_1.py]

AirflowClusterPolicyViolation: If catchup is enabled, max_active_runs must be set to 1.

DAGs



All

1

Active 0

Paused

1

Running 0

Failed 0

Filter DAGs by tag

Search DAGs



DAG ▾

Owner ▾

Runs ⓘ

Schedule

Last Run ▾ ⓘ

Next Run ▾ ⓘ

Recent Tasks ⓘ

Actions

Links

map_dag

abc

airflow



@daily



2025-10-02, 00:00:00 ⓘ



Disallow catchup with max_active_runs>1

```
from airflow.exceptions import AirflowClusterPolicyViolation

def dag_policy(dag):
    allowed_dags = ["example_dag"]
    if dag.dag_id not in allowed_dags and dag.catchup and dag.max_active_runs!=1:
        raise AirflowClusterPolicyViolation("If catchup is enabled, max_active_runs must be set to 1.")
```

Require owner email in tags

```
from airflow.exceptions import AirflowClusterPolicyViolation

def dag_policy(dag):
    email_found = False
    if not dag.tags:
        raise AirflowClusterPolicyViolation("Dag tags are missing.")
    for tag in dag.tags:
        if "@astronomer.io" in tag:
            email_found = True
    if not email_found:
        raise AirflowClusterPolicyViolation("At least one owner email is required.")
```

Task Policies

- Applied to a task object (BaseOperator)
- When it is loaded in the dagbag.
- Check and mutate task parameters.
- Applies to every task instance of the task object.

Governance Use Cases

- callbacks
- worker queues
- ban operators
- executor_config

Restrict usage of a specific connection

```
from airflow.exceptions import AirflowClusterPolicyViolation

def task_policy(task: "BaseOperator") -> None:
    allowed_dags=["ex1_pgAuthorized"]
    if task.dag.dag_id not in allowed_dags:
        for attribute in dir(task):
            if attribute.endswith("conn_id"):
                val = getattr(task, attribute)
                if val=="restricted_postgres":
                    raise AirflowClusterPolicyViolation("Use of restricted_postgres connection is not authorized.")
```

Run KPO & deferrable tasks on a lightweight queue

```
def task_policy(task: "BaseOperator") -> None:  
    cls_path = task.__class__.__module__ + "." + task.__class__.__name__  
    kpo_path = "airflow.providers.cncf.kubernetes.operators.pod.KubernetesPodOperator"  
    is_deferrable = getattr(task, "deferrable", None)  
    if cls_path==kpo_path or is_deferrable:  
        task.queue="lightweight-worker-queue"
```

Task Instance Mutation Hook

- Applied to a task instance object
- When the TI is initialized on the scheduler
- Again after the TI has landed on a worker.
- Selectively applied to some task instances of a task.

Governance Use Cases

- Mutate task instances between:
 - tries
 - map indices
 - runs

Remove success callbacks for manual & backfill runs

```
def task_instance_mutation_hook(task_instance):
    if isinstance(task_instance.run_id, str):
        if "manual" in task_instance.run_id or "backfill" in task_instance.run_id:
            task_instance.task.on_success_callback=None
```

Assign map indices>5 to a separate queue

```
def task_instance_mutation_hook(task_instance):
    if task_instance.map_index>5:
        task_instance.queue="spare-high-resource-queue"
```

This does not work!

Run manual and backfill dags on a specific queue

```
def task_instance_mutation_hook(task_instance):
    if "manual" in task_instance.run_id or "backfill" in task_instance.run_id:
        task_instance.queue="manual-and-backfill-runs-queue"
```

TypeError: argument of type 'NoneType' is not iterable

Add retry callback only before the final retry

```
from airflow.providers.smtp.notifications.smtp import SmtpNotifier

def task_instance_mutation_hook(task_instance):
    if isinstance(task_instance.try_number, int) and isinstance(task_instance.max_tries, int):
        if task_instance.try_number==task_instance.max_tries:
            task_instance.task.on_retry_callback=SmtpNotifier(to="karen.braganza@astronomer.io", subject="The TI is about to run its final retry!")
```

Pod Mutation Hook

- Applied to a Kubernetes pod created using the KubernetesExecutor or KubernetesPodOperator.
- Run when building the pod.
- Examine or mutate a `kubernetes.client.models.V1Pod` object.

Governance Use Cases

- sidecar containers
- NodeSelector
- annotations
- default pod resources.
- termination grace period

Add tolerations to KE worker pod

```
def pod_mutation_hook(pod: V1Pod) -> None:  
    from kubernetes.client import V1Toleration  
  
    pod.spec.tolerations = [  
        V1Toleration(  
            key="node-group",  
            operator="Equal",  
            value="airflow-worker",  
            effect="NoSchedule",  
        )  
    ]
```

Overall Considerations

Cluster policy mutations lack user visibility.

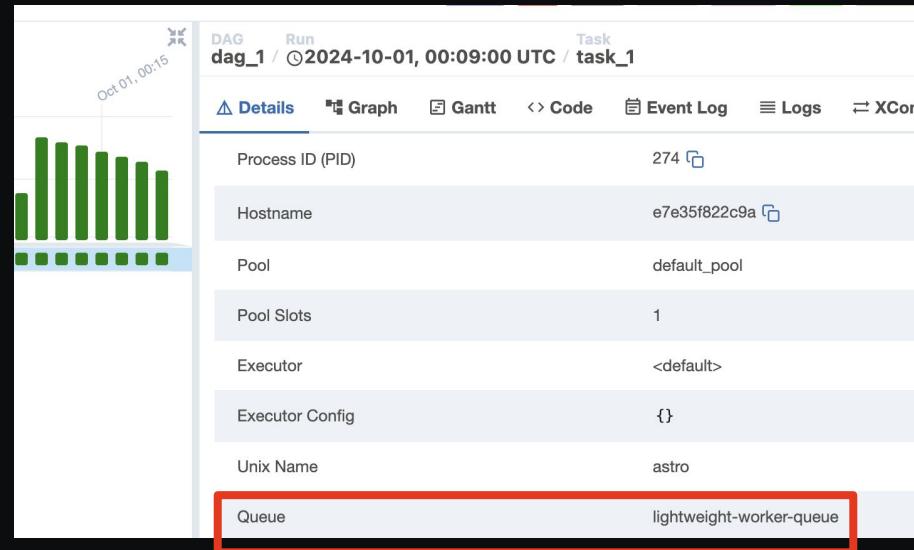
```
from datetime import datetime
from airflow.decorators import dag, task

@dag(start_date=datetime(2024, 10, 1), schedule="* * * * *")
def dag_1():

    @task(queue="default-queue")
    def task_1():
        print("This is task 1!")

    task_1()

dag_1()
```



Overall Considerations

Cluster policies can impact critical Airflow components like the scheduler!

```
import time
import logging

logger = logging.getLogger(__name__)

def task_instance_mutation_hook(task_instance):
    logger.info("Sleeping for 600 seconds...")
    time.sleep(600)
```

Overall Considerations

Cluster policies can impact critical Airflow components like the scheduler!

i	Time	Event
>	10/7/25 10:50:49.109 PM	2025-10-07T22:50:49.109233Z [info] Sleeping for 600 seconds... [policy_plugin.policy] loc=policy.py:64 host = my-splunk-connect-sck-otel-5t7dp k8s.pod.name = exact-eclipse-8805-scheduler-6968b8f78-9c89g source = kubernetes sourcetype = kube:container:scheduler
>	10/7/25 10:55:48.508 PM	2025-10-07T22:55:48.507724Z [info] Exiting gracefully upon receiving signal 15 [airflow.jobs.scheduler_job_runner.SchedulerJobRunner] loc=sched uler_job_runner.py:256 host = my-splunk-connect-sck-otel-5t7dp k8s.pod.name = exact-eclipse-8805-scheduler-6968b8f78-9c89g source = kubernetes sourcetype = kube:container:scheduler
>	10/7/25 10:55:48.509 PM	2025-10-07T22:55:48.509231Z [info] Exited execute loop [airflow.jobs.scheduler_job_runner.SchedulerJobRunner] loc=scheduler_job_runn er.py:1058 host = my-splunk-connect-sck-otel-5t7dp k8s.pod.name = exact-eclipse-8805-scheduler-6968b8f78-9c89g source = kubernetes sourcetype = kube:container:scheduler

Implementation

- airflow_local_settings.py
- custom module with pluggy (using a setuptools entrypoint)

Questions?



Karen Braganza

Senior Airflow Reliability Engineer
at Astronomer





The 2025 Apache Airflow® Survey is here!

Fill it out to for a free Airflow 3 Fundamentals or DAG Authoring in Airflow 3 certification code



References

- [OSS Airflow docs](#)
- [Astronomer docs](#)
- [OSS Airflow GitHub code](#)