



# Multi-instance Asset Synchronization

-  
**Push or Pull ?**

Sébastien Croquevieille

3.0



# Numberly

Othmane EL Metoui  
Head of Data & AI

# What we do

We generate value and growth for our clients by turning their digital marketing expenses into impactful and profitable data-driven investments.

## CRM & Loyalty Management

- 1st Party Data Collection
- Customer Journey design
- Omnichannel campaign management
- **Clean Room management**
- Incrementality & CLV

## Numberly Martech Platform

- Content builder
- Marketing Automation tool
- Messaging APIs
- Impactly

## Digital Media

- Media Strategy
- Full funnel activation
- Data Collaboration & Retail Media
- Experimentation & measurement

## Data Strategy & Insights

- Digital Consulting
- Insights & Analytics
- AI-driven solutions
- Numberly Academy

# Numberly the Marketing **Technologist**

---

- **We are around since 2000, HQ in Paris France.**



Paris



Brussels



London



Amsterdam



Dubai



Tel Aviv



New-York



Montréal

- **95% of Our Infrastructure is Self-Hosted on Bare Metal Servers**
- **We are our own Internet Service Provider.**
- **We commit to Open Source & Developer Community**
  - Numberly Github, OSS financing, Europython Sponsor 11 years straight
- **Self hosted Airflow users since 2017 - version 1.x**

# Numberly the Marketing Technologist

---

## Airflow Key figures in Production:

- **1817 DAGs on 2 instances with 341 Datasets**
- **16 500 daily DagRuns with 44 500 daily TaskInstances**

WhoAmI

# Who Am I ?

- Sébastien Crocquevieille
- Data Engineer @ Numberly
- French & Mexican
- Airflow user since 2018
- EuroPython 2023 Speaker
  - [Orchestrating Python Workflows in Apache Airflow](#)
- Pycon TW 2025 Speaker



You can call me **SEB**

# Events vs Assets Scheduling

# Disclaimer

THE CHARACTERS AND EVENTS DEPICTED IN THIS PHOTOPLAY  
ARE FICTITIOUS. ANY SIMILARITY TO ACTUAL PERSONS,  
LIVING OR DEAD, IS PURELY COINCIDENTAL.

# External Events

---

**It could be anything that happens outside Airflow**

- Any message/post
- Any change in a value
- Any detectable action or signal

# Event Driven Scheduling (Airflow 3.0)

---

## Design :

- **Execution that strictly follows external events**
- **Individual action based on payload ?**
- **1 to 1 execution**

## In practice :

- **A little complex to set up**
  - AssetWatcher: monitors external event source via triggers
- **Requires guarantees concerning:**
  - Ensuring deliverability
  - Avoiding duplicate events
  - Ordering ??

# Asset-Aware Scheduling (Airflow 2.4)

---

## Design :

- **Event that indicates a change in Data Source**
- **Data source changes are indistinguishable**

## Constraints:

- **Similar to Event Driven but with lower constraints**
  - 1 to 1 execution **not** guaranteed
  - Order of events **shouldn't** be important
- **No triggers needed, just Airflow Assets**

# Airflow (Data) Assets

---

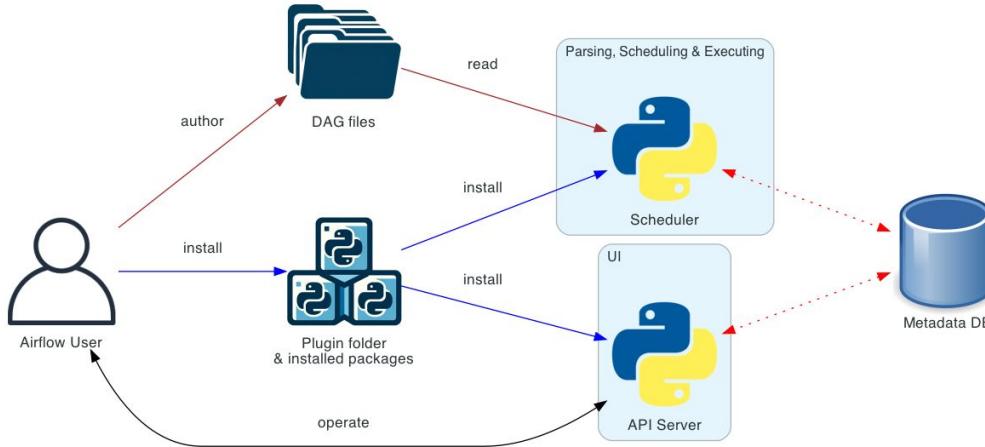
```
from airflow.sdk import Asset
```

```
example_asset = Asset("s3://asset-bucket/example.csv")
```

- It is just a string → representing a Data Source
- Stored in Airflow DB
- Maintaining link is YOUR responsibility
- Generated by DAG parsing

# Upscaling our Scheduling

# Airflow Scheduler



# Performance issues

---

- **Machine performance**
  - High CPU and RAM usage
  - Scheduler heartbeat failures
  - `dag_processing.total_parse_time` is high

# Performance issues

---

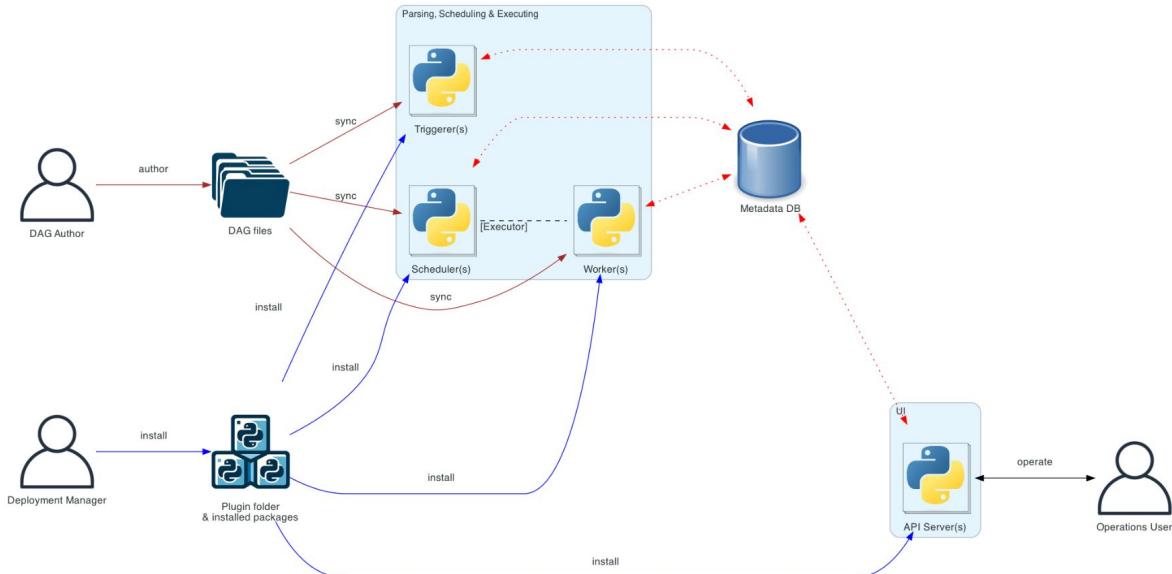
- **Machine performance**
  - High CPU and RAM usage
  - Scheduler heartbeat failures
  - `dag_processing.total_parse_time` is high
- **Slow scheduling**
  - Many tasks in “queued” or “scheduled” state
  - Delayed DAG runs

# Performance issues

---

- **Machine performance**
  - High CPU and RAM usage
  - Scheduler heartbeat failures
  - `dag_processing.total_parse_time` is high
- **Slow scheduling**
  - Many tasks in “queued” or “scheduled” state
  - Delayed DAG runs
- **Database limitations**
  - Too many connections
  - Slow queries

# Airflow Remote Executors



# Airflow Remote Executors

---

## Pros:

- Robust: decoupling workers from scheduler process
- Effective: Worker parallelization
- Available: Low latency workers always running

# Airflow Remote Executors

---

## Cons:

- Requires Infrastructure setup
- Can be Expensive: Cloud cost





# Still having issues?

# Scheduler Replication

---

Running More Than One Scheduler [\[link\]](#)

Airflow supports running more than one scheduler concurrently – both for performance reasons and for resiliency.

# Scheduler Replication

---

## Pros:

- Split scheduling load
- More resilient
- Not much work

## Cons:

- More Database operations



# Need more ?

# Separate Instances

---

Just have 2 Airflows

## Pros:

- Maximum scalability
- Separation of Concerns
- Better Access Control



# Separate Instances

---

Just have 2 Airflows ?!

## Cons:

- Even more maintenance
- Even more resources
- Redundant operations



# Assets

+

# Multiple instances

=

...

# Synchronizing Airflow Assets

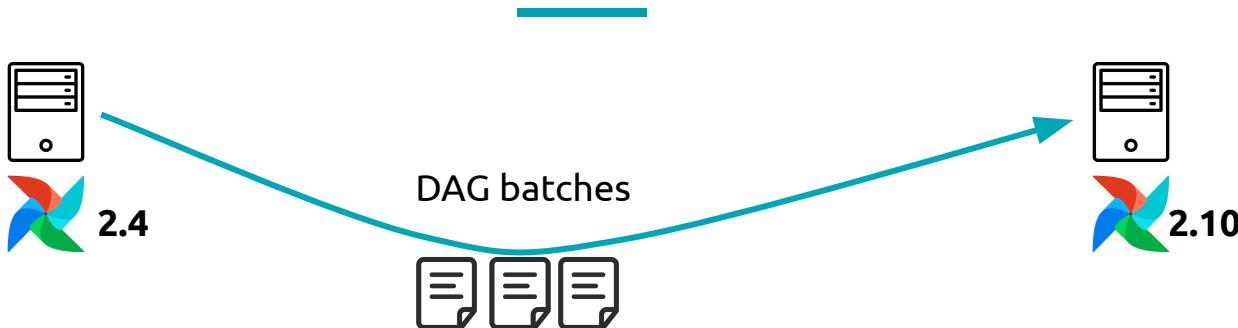
# Our use case

---

- Airflow Migration 2.4 → 2.10 (latest version)
- Complete dependency overhaul (Python & Spark)
- Big dag library (~2000 DAGs)
- From **old** instance to **new**
- Upgrading code as we migrate



# Our use case



- **Migrating DAGs by batches based on**
  - Complexity
  - Dependencies
  - Missing features
  - Urgency
- **Specific needs:**
  - Cross instance DAG Trigger
  - Cross instance Task sensor
- **What about Airflow Assets? ⇒ Synchronization !**

# Push or Pull ?

# Push

# How do we start?

---

By getting the Asset Events!

GET /datasets/events

▼

## Response samples

200

401

403

404

### Content type

application/json

[Copy](#) [Expand all](#) [Collapse all](#)

```
{  
  - "dataset_events": [  
    [- {  
      "dataset_id": 0,  
      "dataset_uri": "string",  
      "extra": { },  
      "source_dag_id": "string",  
      "source_task_id": "string",  
      "source_run_id": "string",  
      "source_map_index": 0,  
      + "created_dagruns": [ ... ],  
      "timestamp": "string"  
    }  
  ],  
  "total_entries": 0  
}
```

# Design choices

---

- **Using Airflow API to get Asset Events**
  - Filter by -timestamp & store the “offset” in Airflow Variables

## io Consumer

airflow-dataset-synchronizer.dataset-events


**Partition**
 All

 Earliest

 Latest

 Date (Paris time)

 Timestamp or oid

 Offset

**Start From**

## Control


**Msg limit**

 Keyword filter

**Filter-in**
 Headers

 Key

 Value

Date	Timestamp	Partition	Offset	Headers	Key	Value
2025-09-01 22:35:42	1756758942053	0	26492		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 98, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23520}
2025-09-01 22:07:57	1756757277595	0	26491		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 73, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23519}
2025-09-01 22:07:57	1756757277595	0	26490		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 78, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23518}
2025-09-01 22:02:29	1756756949361	0	26489		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 98, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23517}
2025-09-01 21:51:55	1756756315091	0	26488		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 98, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23516}
2025-09-01 21:51:55	1756756315091	0	26487		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 77, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23515}
2025-09-01 21:41:20	1756755680563	0	26486		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 101, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23514}
2025-09-01 21:37:49	1756755469318	0	26485		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 78, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23513}
2025-09-01 21:37:49	1756755469318	0	26484		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 72, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23512}
2025-09-01 21:32:16	1756755136838	0	26483		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 96, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23511}
2025-09-01 21:31:47	1756755107354	0	26482		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 83, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23510}
2025-09-01 21:24:43	1756754683595	0	26481		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 94, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23509}
2025-09-01 21:18:41	1756754321064	0	26480		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 115, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23508}
2025-09-01 21:07:40	1756753660001	0	26479		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 72, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23507}
2025-09-01 21:07:40	1756753660001	0	26478		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 78, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23506}
2025-09-01 21:03:34	1756753414963	0	26477		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 98, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23505}
2025-09-01 20:42:26	1756752146414	0	26476		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 230, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23504}
2025-09-01 20:37:31	1756751851238	0	26475		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 78, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23503}
2025-09-01 20:37:31	1756751851238	0	26474		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 72, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23502}
2025-09-01 20:36:23	1756751783697	0	26473		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 98, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23501}
2025-09-01 20:34:52	1756751692813	0	26472		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 91, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23500}
2025-09-01 20:34:52	1756751692813	0	26471		live.flathubprod.grp.bei.consumer.connectivity	{"\$": "followed", "id": 230, "label": "Followed", "name": "live.flathubprod.grp.bei.consumer.connectivity", "version": 1, "version_id": 23499}

# Design choices

---

- **Using Airflow API to get Asset Events**
  - Filter by -timestamp & store the “offset” in Airflow Variables
- **Push them to a Kafka topic ⇒ Single source of truth for all Asset Events**

**Almost Done !**

# Some 403 issues

---

- **No POST endpoint for Dataset Events** - Solved in 2.9
  - Dirty DB editor DAG
- **Dataset Events POST doesn't create NEW Datasets** - "Solved" in 3.0 with Materialize EP
  - Make my own materialize DAG

=====> Conflict between **parsed objects** (Dataset) and **created objects** (Dataset Events)

- **Datasets cannot be dynamically created** - Solved in 2.10 with DatasetAlias
  - Go to latest version (at the time) 2.10

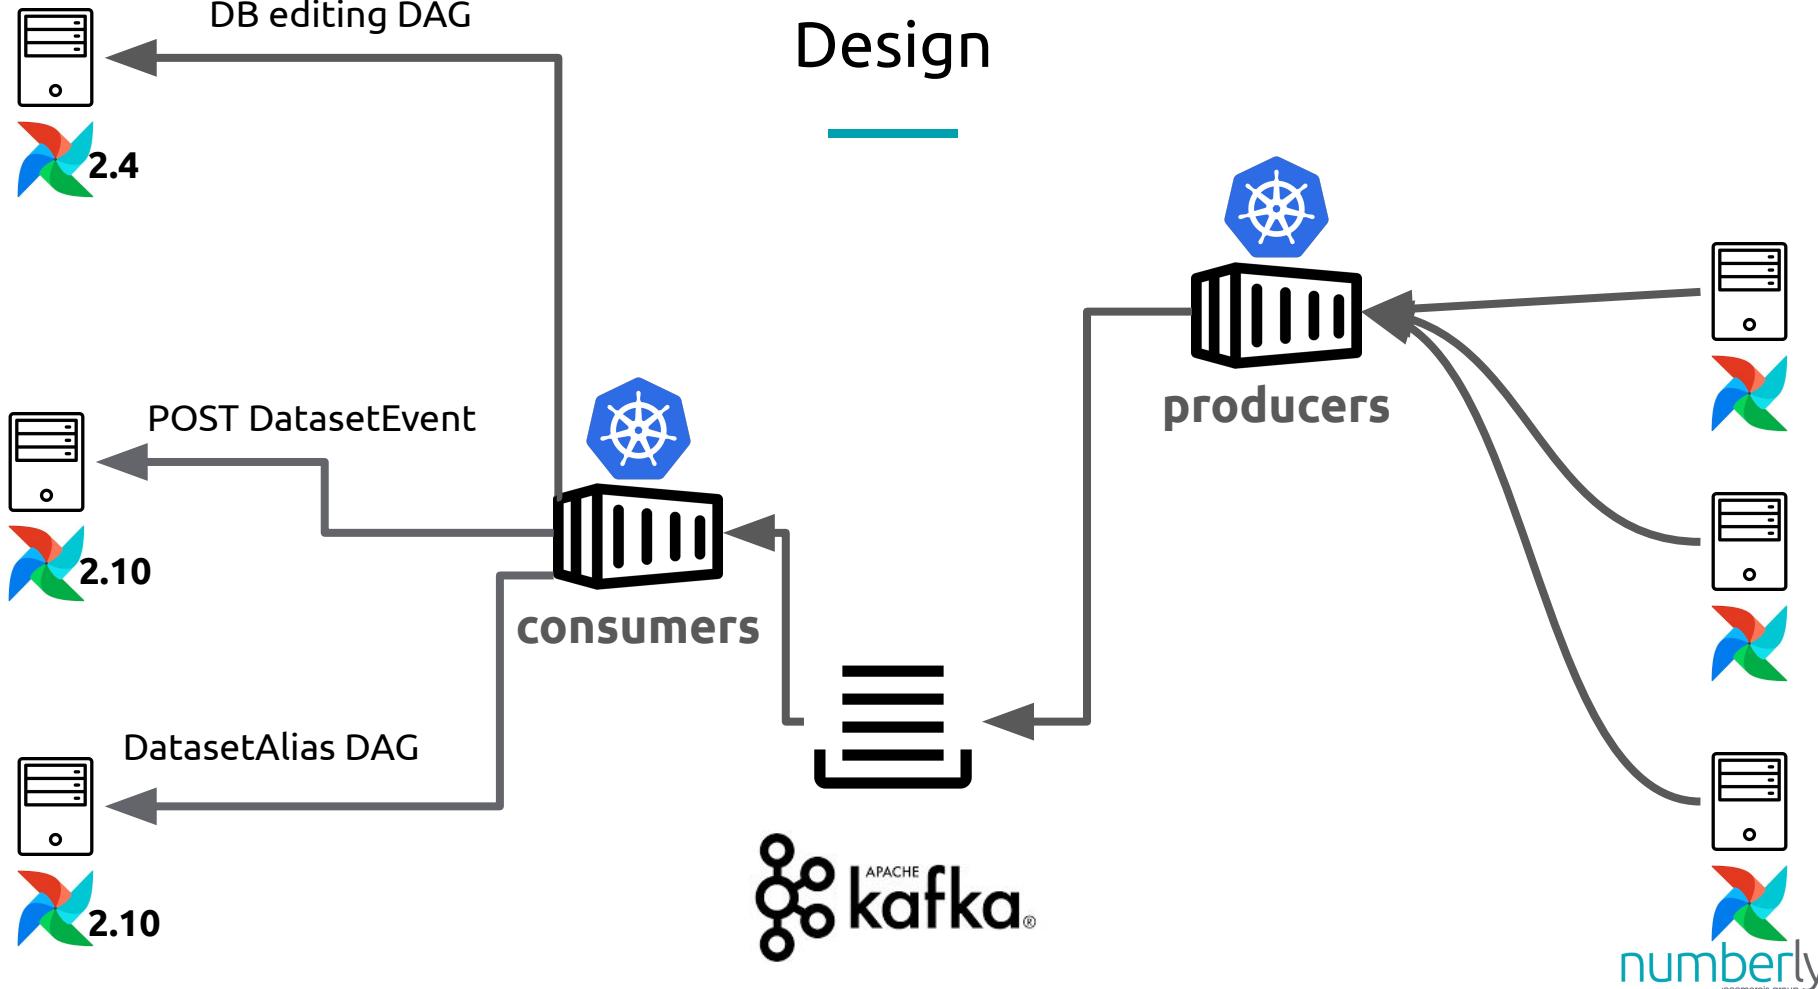
# Design choices

---

- **Using Airflow API to get Asset Events**
  - Filter by -timestamp & store the “offset” in Airflow Variables
- **Kafka topic ⇒ Single source of truth for all Asset Events**
- **Push back to Airflow through 3 different methods**

# Design

---



```
1  {
2      "dataset_id": int,
3      "dataset_name": URI,
4      "event_id": int,
5      "event_timestamp": Datetime,
6      "origin": AIRFLOW INSTANCE URL,
7      "source_dag": DAG_ID,
8      "source_task": TASK_ID
9  }
```

It works !



Or Pull ?

# Pull Based Scheduling

---

- **Airflow PULL from external instances directly**
- **Use Event based scheduling**
  - Ignore the fact that we need another migration
- **Poll Asset Events continuously to stay updated**

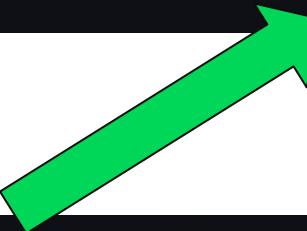
**Asset Watcher ⇒ Asset Event ⇒ DAG execution**

Let's code together

# Using an Asset Watcher

---

```
from airflow.sdk import AssetWatcher  
  
my_watcher = AssetWatcher(name="my_watcher", trigger=???)
```



```
trigger: BaseEventTrigger | dict
```

# Base Event Trigger

---

## CLASS:

- **Inherits from BaseEventTrigger**

## OBLIGATORY METHODS:

- **serialize**
  - Class path & arguments required for initialization must be serializable
- **run**
  - The actual trigger operation
- **\_\_init\_\_**

## OPTIONAL METHODS:

- **cleanup**
- **\_set\_context**

```
7  from airflow.triggers.base import BaseEventTrigger, TriggerEvent
8
9  class SimpleEventTrigger(BaseEventTrigger):
10     def __init__(self, url: str, **kwargs):
11         self.url = url
12
13     def serialize(self) → tuple[str, dict[str, Any]]:
14         return (
15             "event_scheduling.SimpleEventTrigger",
16             {"url": self.url},
17         )
18
19     async def run(self) → AsyncIterator[TriggerEvent]:
20         while True:
21             event_str = await aiohttp.get(self.url).text()
22             if not event_str:
23                 continue
24             else:
25                 yield TriggerEvent(event_str)
26
```

```
27     from airflow.sdk import Asset, AssetWatcher  
28  
29     trigger = SimpleEventTrigger(url="http://magic/airflow_endpoint")  
30     my_asset = Asset(  
31         "simple_trigger_asset", watchers=[AssetWatcher(name="my_watcher", trigger=trigger)]  
32     )
```

```
with DAG(dag_id="event_scheduled_job", schedule=[my_asset]) as dag:
```

# Triggerer Process

# Existing Event Triggers

---

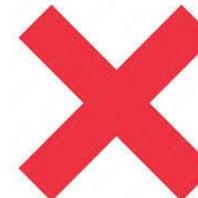
Event Triggers:

- **KAFKA:** KafkaMessageQueueTrigger
- **GENERIC MQ:** MessageQueueTrigger



Base Triggers:

- **REDIS:** AwaitMessageTrigger
- **GOOGLE:** GCSBlobTrigger
- **AWS:** EC2StateSensorTrigger, S3KeyTrigger, ...



# Pull based scheduling

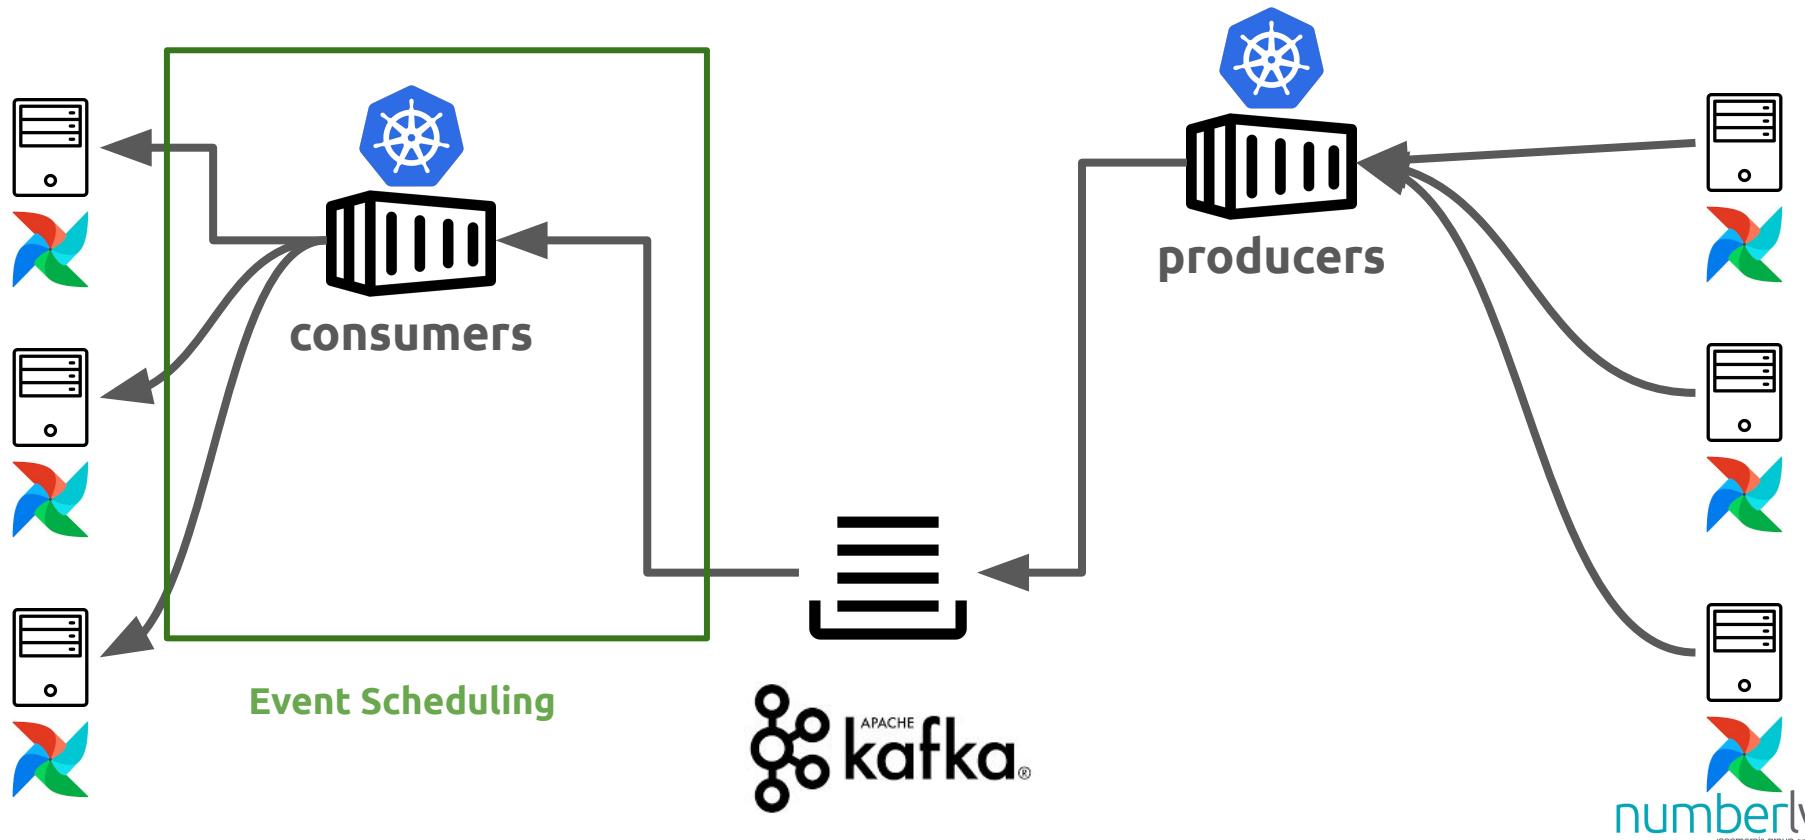
---

## Caveats:

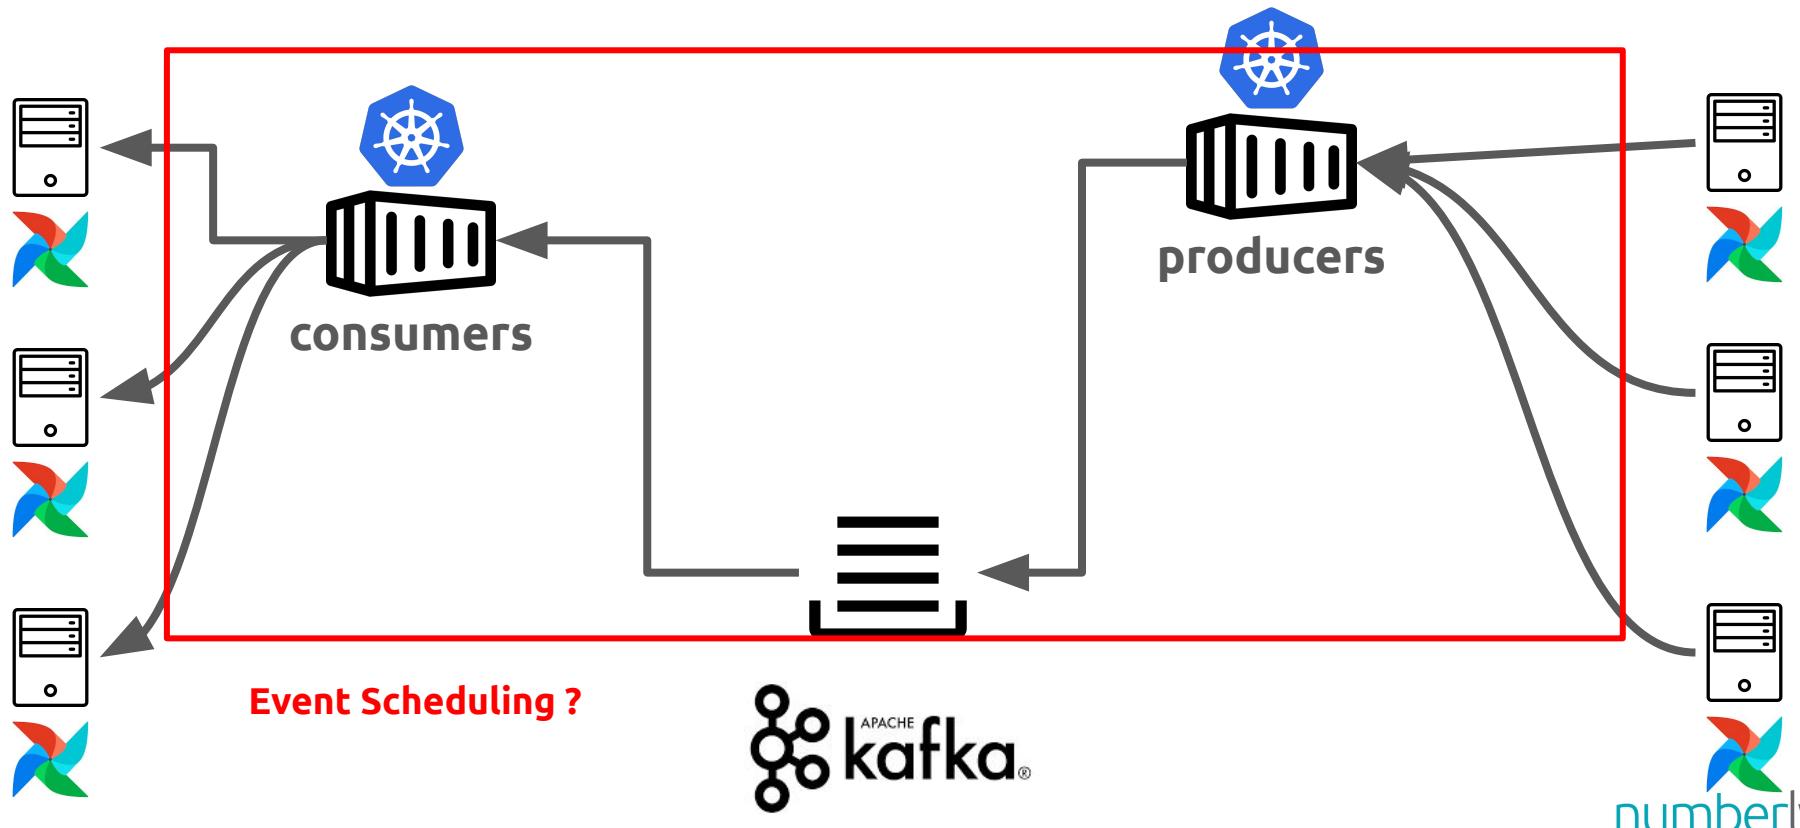
- **“run” method is king**
  - Error handling
  - Pagination
  - Authentication
  - No blocking
  - Not heavy
- **Know where your Triggerers are running**
  - Do you have one ?
  - Is it sharing scheduler resources ?
- **Use cases outside message queues are “exploratory” if possible**

So... can we pull ?

# Using Pull ?



# Using Pull ?



# Cost of pulling

---

## With the queue ?

- **Constant triggerer polling**
  - Less accurate than push
  - Or noisy
- **Still need to push to queue**
  - Keep dedicated service?
  - Use custom "high frequency" DAGs ?

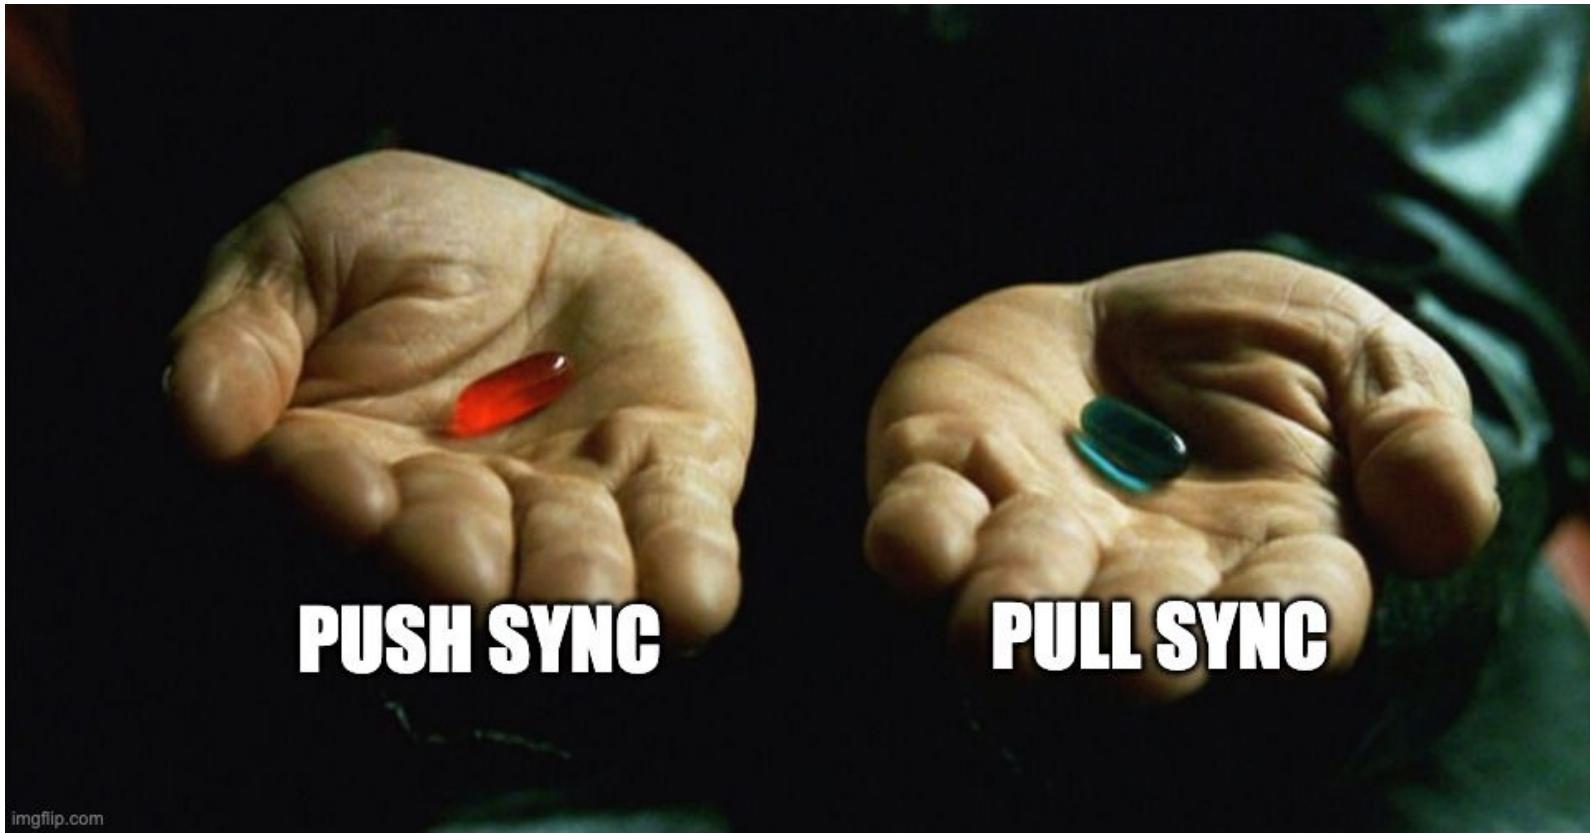


## And without the queue?

- **API request + Airflow Variable**
- **O( $n^2$ ) network calls**
- **Even More Airflow Activity**

So the choice is

I'm still on Airflow 2 !



Question time !

# Bonus

ASTRONOMER

# THE DATA FLOWCAST



Scaling On-Prem Airflow  
With 2,000 DAGs  
at Numberly  
**with Sébastien Crocquevieille**

