



# Purple is the New Green

Harnessing Deferrable Operators to  
Improve Performance & Reduce Costs

Ethan Shalev

3.0



# Blue is the New Blue

Harnessing Deferrable Operators to  
Improve Performance & Reduce Costs

Ethan Shalev

3.1

# Ethan Shalev

- Data Engineer @ Wix
- Airflow tech-lead & evangelist
- 20+ years in the field of data



[/in/eshalev](https://www.linkedin.com/in/eshalev)

# Agenda

- 1 About Wix
- 2 Motivation for using deferrable operators
- 3 Explanation of deferrable operators and their design
- 4 Implementation approach and impact
- 5 Pitfalls, lessons learned and next steps

1.

# About Wix

About Wix

Motivation

Deferral explained



Implementation

Next Steps

WIX

# About Wix

- The leading SaaS website builder platform
- Runs ~4% of all active sites on the WWW<sup>1</sup>
- Drives ~40% of traffic to sites created with simple website builders<sup>2</sup>



# Data @ Wix



## People

- 75 Data Scientists
- 120 Data Engineers
- 240 Business Analysts



## Data

- 20TB of data added and processed daily
- 1.2M SQL Queries per day



## Airflow

- 3 Production Airflow clusters
  - Migrating from 2.6.3 >> 3.1
- 7,500 DAGs
- 270,000 Daily Airflow tasks
- 10,000+ worker hours per day

2

# .Motivation


# Motivation

Reduce data-delivery  
bottlenecks

Reduce Airflow  
Worker load

Free-up  
resources

Reduce  
costs



# Exploration

Just increase resources?

Airflow Sensors?

The screenshot shows a Jira interface with a search bar at the top right. On the left, there's a sidebar with navigation links: 'For you', 'Recent', 'Starred', and 'Add'. The main content area displays a single Jira issue. At the top of the issue card, it says 'Projects / DE Frameworks / Add parent / DEFW-805'. The title of the issue is '**Investigate the feasibility of converting Platyspark operators to sensors to reduce Airflow pressure while Platyspark waits for the app to finish.**'

# Solution

# Deferrable operators!



**Eitan Shalev** May 16th, 2022 at 12:43

Deferrable Operators, new in Airflow 2.2, seem interesting, and have the potential of clearing up a lot of idle resources used by sensors and other operators that trigger external work  
<https://www.astronomer.io/guides/deferrable-operators>

# 3.

# Deferral explained

# De·fer (/də'fər/):

- To postpone or delay
- To yield to another

# Deferrable Operators


Transfer the management of a task from a worker to a triggerer service while an external process is executing

Standalone process, capable of handling 1,000+ asynchronous triggers

Frees up Airflow worker slot; increases Airflow availability

Deferral is defined at the operator level, invoked either by default or per-task.

# Typical execution vs. deferred execution



# Typical execution vs. deferred execution

No deferral:

```
class MyShinyOperator(BaseOperator):  
  
    def execute(self, context: Context):  
        result = call_external(context)  
        # connection open, waiting...  
  
        return result if self.do_xcom_push
```

```
class MyShinyOperator(BaseOperator):  
  
    def execute(self, context: Context):  
        request_id = call_external(context)  
        while True:  
            result = poll_external(request_id)  
            if result:  
                return result if self.do_xcom_push  
            else:  
                sleep(30)
```

# Typical execution vs. deferred execution

Deferral:

```
class MyShinyOperator(BaseOperator):

    def execute(self, context: Context):
        request_id = call_external(context)

        self.defer(trigger=MyShinyTrigger(request_id),
                  method_name="all_shiny",
                  kwargs=context, #optional
                  timeout=timedelta(minutes=15) #optional
                  )

    def all_shiny(context)
        return context.result if self.do_xcom_push
```

# What makes a trigger?

[GitHub](#)

```
import asyncio


from airflow.triggers.base import BaseTrigger, TriggerEvent
from airflow.utils import timezone

class DateTimeTrigger(BaseTrigger):
    def __init__(self, moment):
        super().__init__()
        self.moment = moment

    def serialize(self):
        return (
            "airflow.providers.standard.triggers.temporal.DateTimeTrigger",
            {"moment": self.moment}
        )

    async def run(self):
        while self.moment > timezone.utcnow():
            await asyncio.sleep(1)
            yield TriggerEvent(self.moment)
```

# Typical execution vs. deferred execution



# Applying deferral in your DAG

```
from airflow import DAG
from myOperators.shiny import MyShinyOperator
from datetime import datetime

with DAG(
    dag_id="deferrable_dag_demo",
    start_date=datetime(2025, 10, 8),
    schedule=None
):
    MyShinyOperator(
        task_id="so_shiny",
        shiny_params=params,
        deferrable=True
    )
```

# 4.

# Implementation

# Approach

Query Airflow  
metadata DB

Map Operator-types by  
average task duration and  
count

Identify  
quick-wins

Long average run  
time,  
Lots of them


```
SELECT
    Operator,
    Sum(duration) AS duration_seconds,
    Count(1) AS operator_count
FROM
    airflow_db..task_instance
WHERE
    state = 'success'
    AND start_date > date'2025-08-01'
    AND start_date < date'2025-09-01'
GROUP BY 2;
```

# Findings

100+ distinct types of operators

107 rows(Full Result)

operator	duration_minutes	operator_count	avg_duration_seconds
QuixflowTrinoOperator	2,298,485.419	3,184,146	43.311
PythonOperator	998,828.533	584,367	102.555
PlatypusSparkEKSRunAppOperator	966,831.021	115,104	503.978
SecuredPlatypusSparkEKSRunAppOperator	491,821.369	76,717	384.651
TrinoOperator	408,748.992	249,867	98.152
KubernetesPodCondaOperator	371,215.245	38,080	584.898
SecuredKubernetesPodCondaOperator	297,147.89	58,073	307.008
_PythonDecoratedOperator	290,920.065	764,463	22.833
PrestoBatchOperator	251,625.372	101,728	148.411
GxOperator	238,039.876	88,950	160.567
AlertingDdsSensor	156,479.901	5,888	1,594.564
QuixflowSlackActionOperator	117,760.406	86,909	81.299
BqToPrestoOperator	116,401.7	16,550	422
EmrJobFlowSensor	111,742.047	3,835	1,748.246
CustomGxOperator	100,335.102	27,079	222.316
PythonVirtualenvOperator	94,058.265	6,559	860.42
BranchPythonOperator	67,912.104	231,280	17.618
QuixflowInternalMailOperator	28,607.498	24,872	69.011
PrestoBatchSecuredOperator	19,664.64	16,145	73.08



Custom TrinoOperators  
Spark on EKS  
Custom KubernetesPod Operators  
Custom Great-Expectations (GX) Operators  
Data-transfer Operators  
(Snowflake to Iceberg to BQ...)  
Data-freshness sensors  
Slack Operator  
OpsGenie Operator  
PythonOperator(s)

# Where do you start?



trino?



Complex

Long time to delivery

Chance of failure too high

## Start where it's easy!

Create MVP

Communicate the feature

Encourage adoption

Recruit support and get others involved

Use momentum to continue

# GxOperator as a POC

Your data assets:  
database tables, flat  
files, dataframes...

Data validation with  
Great Expectations



great expectations



High quality data in  
your data products

Data documentation  
& data quality reports



Logging & alerting

[Great Expectations docs](#)

About Wix

Motivation


Deferral explained

Implementation

Next Steps

WIX

# GxOperator as a POC



# Refactoring GxOperator

- Examine Execute() method
- Identify where external polling is done
- Replace it with trigger call
- Implement asynchronous trigger

```
retry_call(self.create_expectation, fargs=(context,), tries=self.create_expectation_config.tries,
           delay=self.create_expectation_config.delay_in_seconds)


if self.deferrable:
    self.defer(trigger=GxTrigger(self.run_guid),
method_name="process_expectation_results_returned_from_trigger")
else:
    retry_call(self.get_expectation_results, tries=self.get_expectation_config.tries,
               delay=self.get_expectation_config.delay_in_seconds, exceptions=(ValueError,))
```

# GxTrigger

```
async def run(self):
    await asyncio.sleep(self.initial_wait_delay)
    session = requests.Session()
    response = session.get(self.endpoint, params={"run_guid": self.run_guid})

    while not "success" in json.loads(response.text):
        await asyncio.sleep(self.wait_delay)
        response = session.get(self.endpoint, params={"run_guid": self.run_guid})
    yield TriggerEvent({"run_guid": self.run_guid, "response": json.loads(response.text)})
```

# Demo



About Wix

Motivation

Deferral explained

**Implementation**

Next Steps

**WIX**

# 5.

# Next Steps

# Rollout

Test locally

Deploy on select DAGs

Monitor

Communicate

Track adoption

Scale triggerer

Set deferral by default



Eitan Shalev Mar 16th at 13:25

Airflow word of the day: Defer (di-'fər): to delegate to another:

Introduced in Airflow 2.2, **deferrable operators** allow tasks to "sleep" while waiting on external systems, freeing up worker slots. Instead of using up an airflow worker to run a task that polls an external service for completion, they defer the polling to an asynchronous triggerer service which can efficiently handle many hundreds of polls, till the external task is completed, at which point it returns execution to the worker.

Deferrable operators free up worker slots, resulting in better resource utilization, especially for long-running external requests (e.g., APIs, DB queries).

I encourage you to take a look at Airflow's official [documentation](#) on the topic.

I've recently modified our GX operator to support the optional deferrable mode. Feel free to take a look at the changes I made ([GitHub](#)).

Now, by adding the parameter `deferrable=True` to your GX operator, you will be allowing Airflow to free up a worker slot, so other tasks (in other DAGs) can start running, thus removing one potential bottleneck.

We're planning to gradually update DAGs to use this feature, and once we're certain it's stable, we will change the default behavior to True, so all GX operators, unless explicitly set to False, will defer.

The next steps after that are making more of our operators deferrable, so Airflow can free up workers while other services like Trino or AWS are doing their thing.

So start getting used to seeing purple tasks alongside your green ones .

Needless to say, this could not have been done without [@orene](#) & [@heorhiib](#)'s help, and the support from [@ariksa](#) and his team.

# Pitfalls

Ensure gradual rollout

Track Triggerer process performance

Scale up as needed

Track # of concurrent deferred tasks

Don't start by deferring a DAG with 1,000 tasks

All at once

On a weekend



**Elad Haziza** Jul 18th at 19:50

Hey!

Some of our deferrable tasks seem to be stuck.

I came across this error in the Airflow UI – could someone take a look?

The triggerer does not appear to be running. Last heartbeat was received ~~1000~~ seconds ago.



**Arik Sasson** 🌴 Jul 19th at 16:15

Great. ThX @Haziza & @orene  
Next deferable only with Astronomer...

# Next steps

Custom  
SparkOperators

TrinoOperators  
(SqlOperator)

Defer natively  
deferrable  
operators  
(Amazon, etc.)

Migrate custom  
PythonOperators  
to deferrables



**Almog Gelber** 17:33

The operator now has the capability to run in deferrable mode. While this functionality has been tested and confirmed to work, we are not enabling it in our environment at the moment because it introduced stability issues with our current Airflow setup.

# Expected Impact

Operator	Worker Hr/Day saved	Cost savings \$/Month	CO <sub>2</sub> Emissions reduced (Estimate)
GxOperator	450 (4.5%)	\$600	0.6 tCO <sub>2</sub> / Year
SparkOperators	2,000 (20%)	\$3,200	3 tCO <sub>2</sub> / Year
TrinoOperators	4,000 (40%)	\$6,400	6 tCO <sub>2</sub> / Year



# Conclusions

Deferrable operators  
improve performance  
& reduce costs

Easier than they seem  
at first glance

Do it!

# Thank you!

# Questions?



/in/eshalev

eitansh@wix.com

WIX