



Purple is the New Green

Harnessing Deferrable Operators to
Improve Performance & Reduce Costs

Ethan Shalev

3.0



Blue is the New Blue

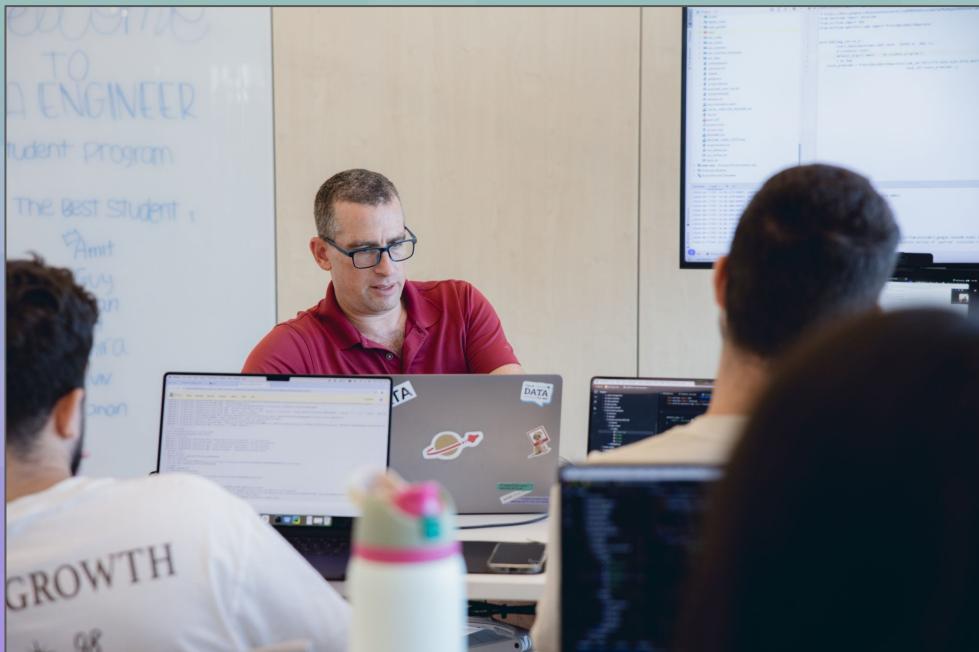
Harnessing Deferrable Operators to
Improve Performance & Reduce Costs

Ethan Shalev

3.1

Ethan Shalev

- Data Engineer @ Wix
- Airflow tech-lead & evangelist
- 20+ years in the field of data



[/in/eshalev](https://www.linkedin.com/in/eshalev)

Agenda

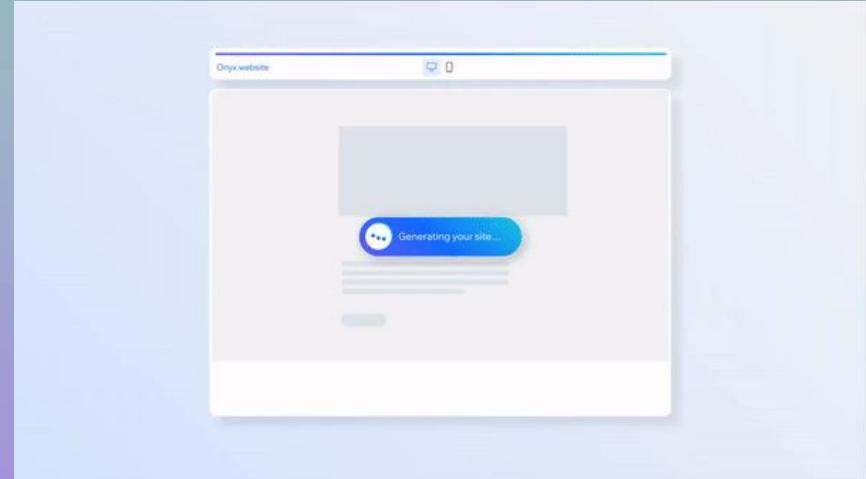
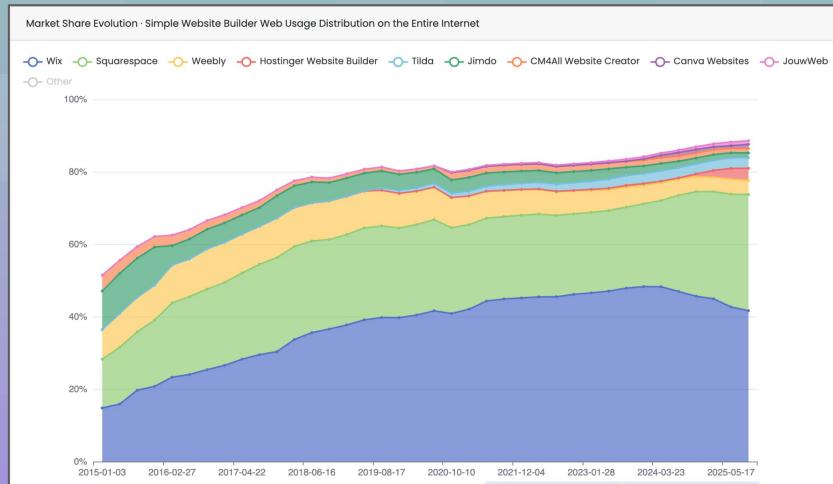
- 1** About Wix
- 2** Motivation for using deferrable operators
- 3** Explanation of deferrable operators and their design
- 4** Implementation approach and impact
- 5** Pitfalls, lessons learned and next steps

1.

About Wix

About Wix

- The leading SaaS website builder platform
- Runs ~4% of all active sites on the WWW¹
- Drives ~40% of traffic to sites created with simple website builders²



Data @ Wix



People

- 75 Data Scientists
- 120 Data Engineers
- 240 Business Analysts



Data

- 20TB of data added and processed daily
- 1.2M SQL Queries per day



Airflow

- 3 Production Airflow clusters
 - Migrating from 2.6.3 >> 3.1
- 7,500 DAGs
- 270,000 Daily Airflow tasks
- 10,000+ worker hours per day

2

.Motivation

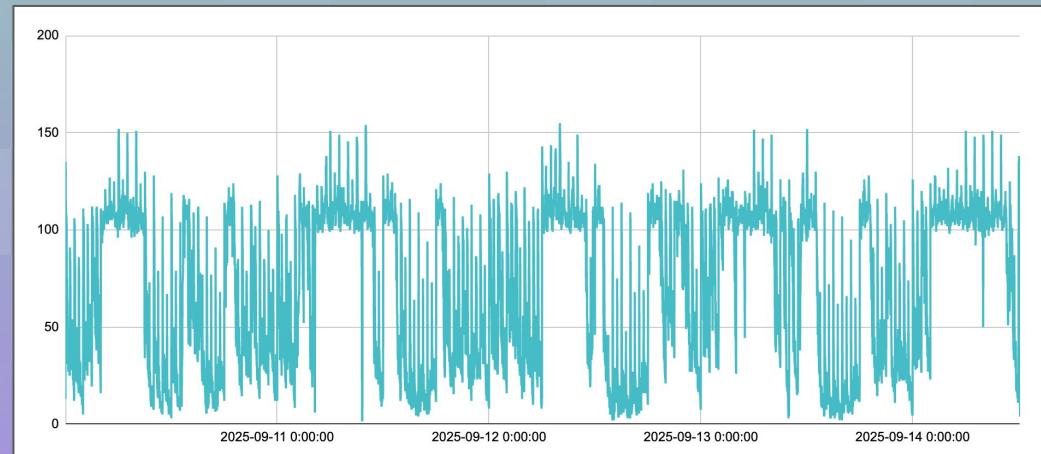
Motivation

Reduce data-delivery
bottlenecks

Reduce Airflow
Worker load

Free-up
resources

Reduce
costs



Exploration

Just increase resources?

Airflow Sensors?

The screenshot shows a Jira interface. At the top, there are icons for issues, epics, and Jira, followed by a search bar. Below the header, a sidebar on the left lists 'For you', 'Recent', and 'Starred' items. The main content area displays a single issue card. The title of the card is: **Investigate the feasibility of converting Platyspark operators to sensors to reduce Airflow pressure while Platyspark waits for the app to finish.** The card also includes navigation links for 'Projects / DE Frameworks / Add parent / DEFW-805'.

Solution

Deferrable operators!



Eitan Shalev May 16th, 2022 at 12:43

Deferrable Operators, new in Airflow 2.2, seem interesting, and have the potential of clearing up a lot of idle resources used by sensors and other operators that trigger external work
<https://www.astronomer.io/guides/deferrable-operators>

3.

Deferral explained

De·fer (/də'fər/):

- To postpone or delay
- To yield to another

Deferrable Operators

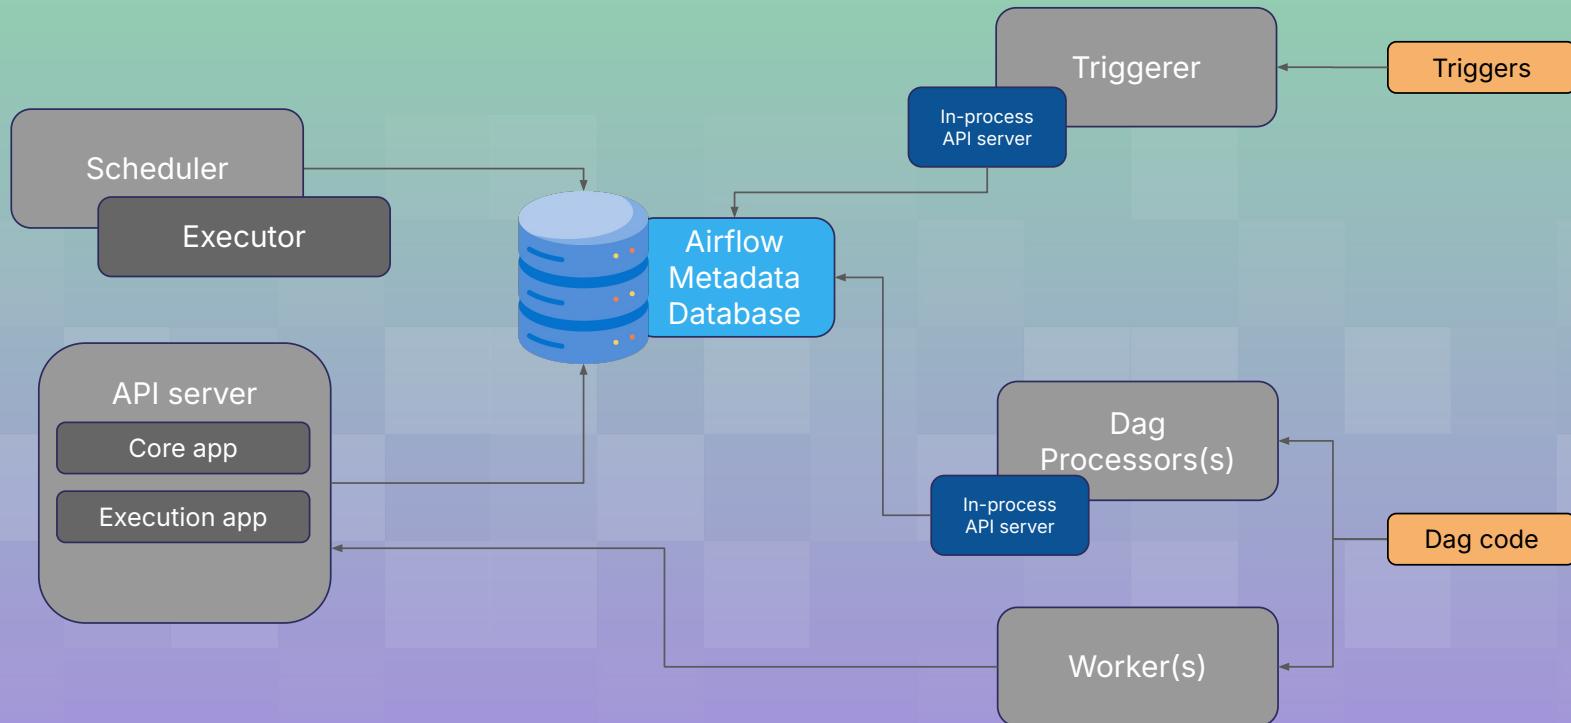
Transfer the management of a task from a worker to a triggerer service while an external process is executing

Standalone process, capable of handling 1,000+ asynchronous triggers

Frees up Airflow worker slot; increases Airflow availability

Deferral is defined at the operator level, invoked either by default or per-task.

Typical execution vs. deferred execution



Typical execution vs. deferred execution

No deferral:

```
class MyShinyOperator(BaseOperator):  
  
    def execute(self, context: Context):  
        result = call_external(context)  
        # connection open, waiting...  
  
        return result if self.do_xcom_push
```

```
class MyShinyOperator(BaseOperator):  
  
    def execute(self, context: Context):  
        request_id = call_external(context)  
        while True:  
            result = poll_external(request_id)  
            if result:  
                return result if self.do_xcom_push  
            else:  
                sleep(30)
```

Typical execution vs. deferred execution

Deferral:

```
class MyShinyOperator(BaseOperator):

    def execute(self, context: Context):
        request_id = call_external(context)

        self.defer(trigger=MyShinyTrigger(request_id),
                  method_name="all_shiny",
                  kwargs=context, #optional
                  timeout=timedelta(minutes=15) #optional
                  )

    def all_shiny(context)
        return context.result if self.do_xcom_push
```

What makes a trigger?

[GitHub](#)

```
import asyncio

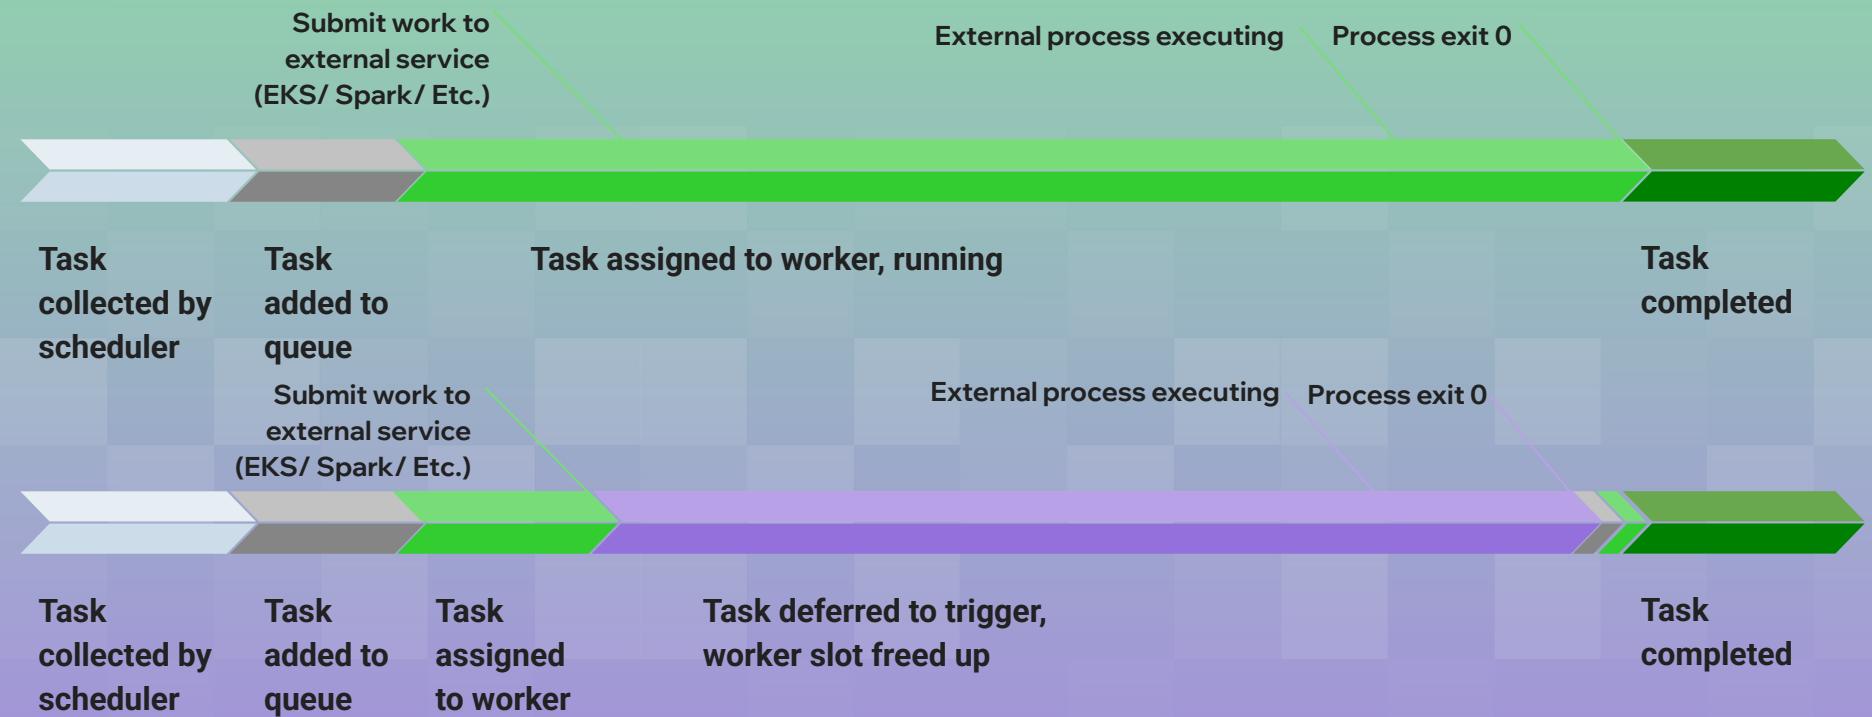
from airflow.triggers.base import BaseTrigger, TriggerEvent
from airflow.utils import timezone

class DateTimeTrigger(BaseTrigger):
    def __init__(self, moment):
        super().__init__()
        self.moment = moment

    def serialize(self):
        return (
            "airflow.providers.standard.triggers.temporal.DateTimeTrigger",
            {"moment": self.moment}
        )

    async def run(self):
        while self.moment > timezone.utcnow():
            await asyncio.sleep(1)
            yield TriggerEvent(self.moment)
```

Typical execution vs. deferred execution



Applying deferral in your DAG

```
from airflow import DAG
from myOperators.shiny import MyShinyOperator
from datetime import datetime

with DAG(
    dag_id="deferrable_dag_demo",
    start_date=datetime(2025, 10, 8),
    schedule=None
):
    MyShinyOperator(
        task_id="so_shiny",
        shiny_params=params,
        deferrable=True
    )
```

4.

Implementation

Approach

Query Airflow
metadata DB

Map Operator-types by
average task duration and
count

Identify
quick-wins

Long average run
time,
Lots of them

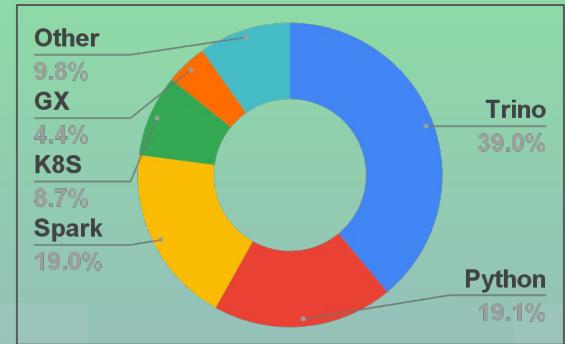
```
SELECT
    Operator,
    Sum(duration) AS duration_seconds,
    Count(1) AS operator_count
FROM
    airflow_db..task_instance
WHERE
    state = 'success'
    AND start_date > date'2025-08-01'
    AND start_date < date'2025-09-01'
GROUP BY 2;
```

Findings

100+ distinct types of operators

107 rows(Full Result)

| operator | duration_minutes | operator_count | avg_duration_seconds |
|------------------------------------|------------------|----------------|----------------------|
| QuixflowTrinoOperator | 2,298,485.419 | 3,184,146 | 43.311 |
| PythonOperator | 998,828.533 | 584,367 | 102.555 |
| PlatySparkEKSRunAppOperator | 966,831.021 | 115,104 | 503.978 |
| SecuredPlatySparkEKSRunAppOperator | 491,821.369 | 76,717 | 384.651 |
| TrinoOperator | 408,748.992 | 249,867 | 98.152 |
| KubernetesPodCondaOperator | 371,215.245 | 38,080 | 584.898 |
| SecuredKubernetesPodCondaOperator | 297,147.89 | 58,073 | 307.008 |
| _PythonDecoratedOperator | 290,920.065 | 764,463 | 22.833 |
| PrestoBatchOperator | 251,625.372 | 101,728 | 148.411 |
| GxOperator | 238,039.876 | 88,950 | 160.567 |
| AlertingDdsSensor | 156,479.901 | 5,888 | 1,594.564 |
| QuixflowSlackActionOperator | 117,760.406 | 86,909 | 81.299 |
| BqToPrestoOperator | 116,401.7 | 16,550 | 422 |
| EmrJobFlowSensor | 111,742.047 | 3,835 | 1,748.246 |
| CustomGxOperator | 100,335.102 | 27,079 | 222.316 |
| PythonVirtualenvOperator | 94,058.265 | 6,559 | 860.42 |
| BranchPythonOperator | 67,912.104 | 231,280 | 17.618 |
| QuixflowInternalMailOperator | 28,607.498 | 24,872 | 69.011 |
| PrestoBatchSecuredOperator | 19,664.64 | 16,145 | 73.08 |



Custom TrinoOperators
Spark on EKS
Custom KubernetesPod Operators
Custom Great-Expectations (GX) Operators
Data-transfer Operators
(Snowflake to Iceberg to BQ...)
Data-freshness sensors
Slack Operator
OpsGenie Operator
PythonOperator(s)

Where do you start?



trino?



Complex

Long time to delivery

Chance of failure too high

Start where it's easy!

Create MVP

Communicate the feature

Encourage adoption

Recruit support and get others involved

Use momentum to continue

GxOperator as a POC

Your data assets:
database tables, flat
files, dataframes...

Data validation with
Great Expectations



great expectations



High quality data in
your data products

Data documentation
& data quality reports



Implementation

[Great Expectations docs](#)

About Wix

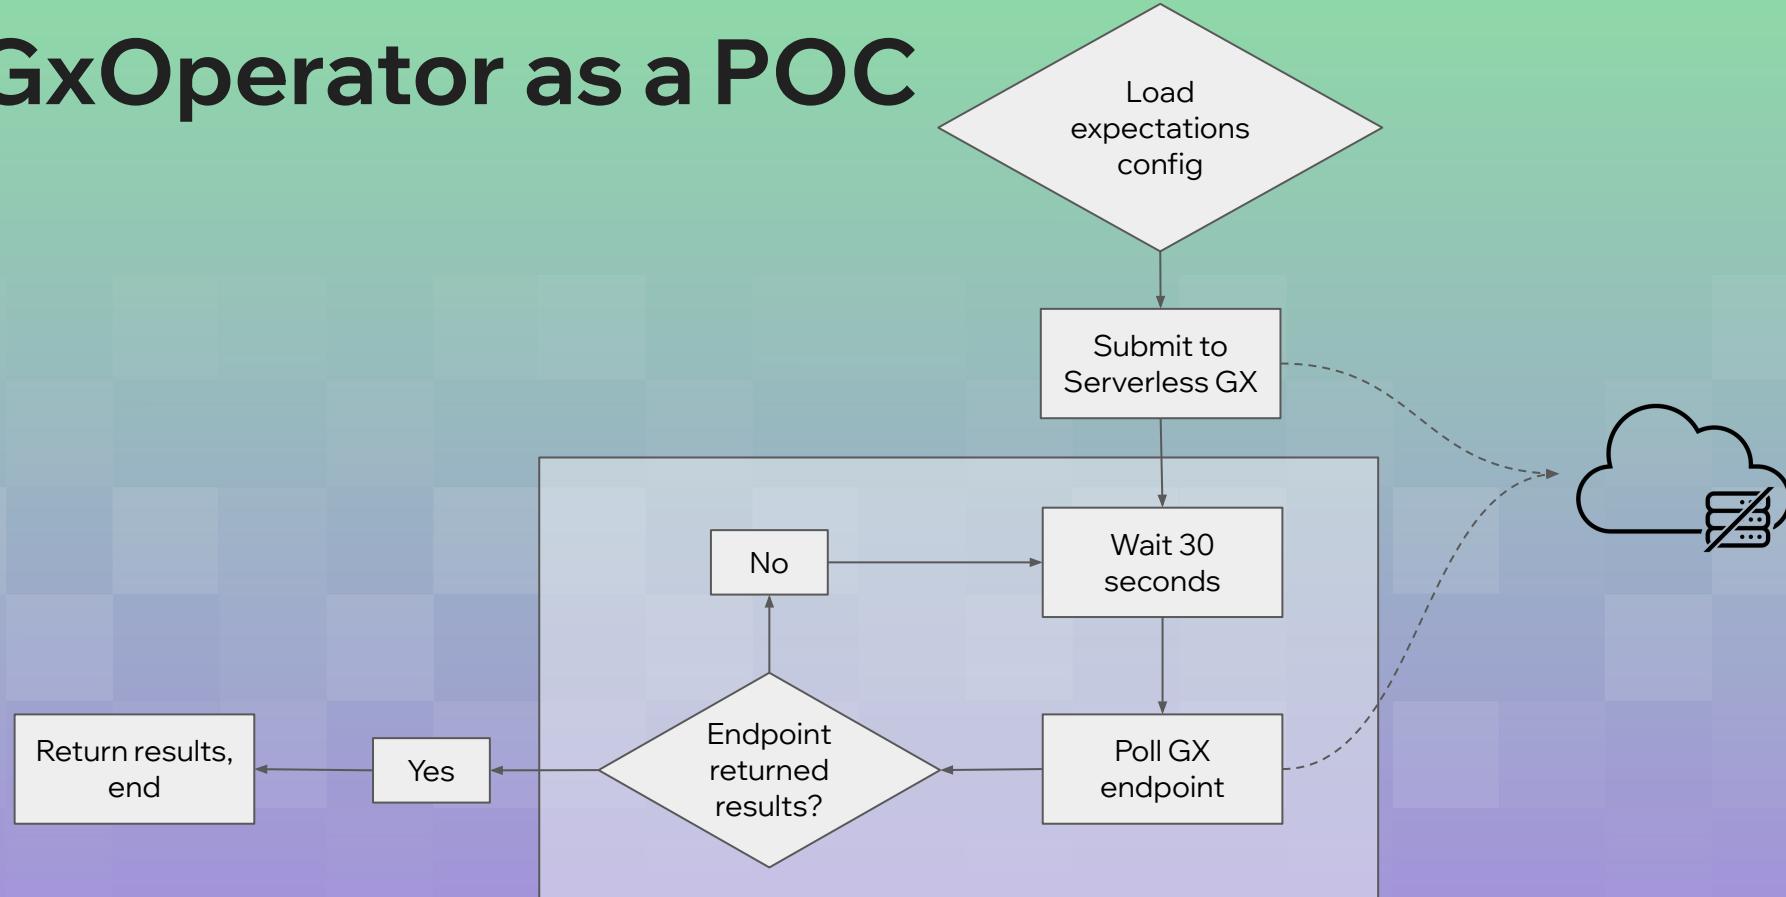
Motivation

Deferral explained

Next Steps

WIX

GxOperator as a POC



Refactoring GxOperator

- Examine Execute() method
- Identify where external polling is done
- Replace it with trigger call
- Implement asynchronous trigger

```
retry_call(self.create_expectation, fargs=(context,), tries=self.create_expectation_config.tries,
           delay=self.create_expectation_config.delay_in_seconds)

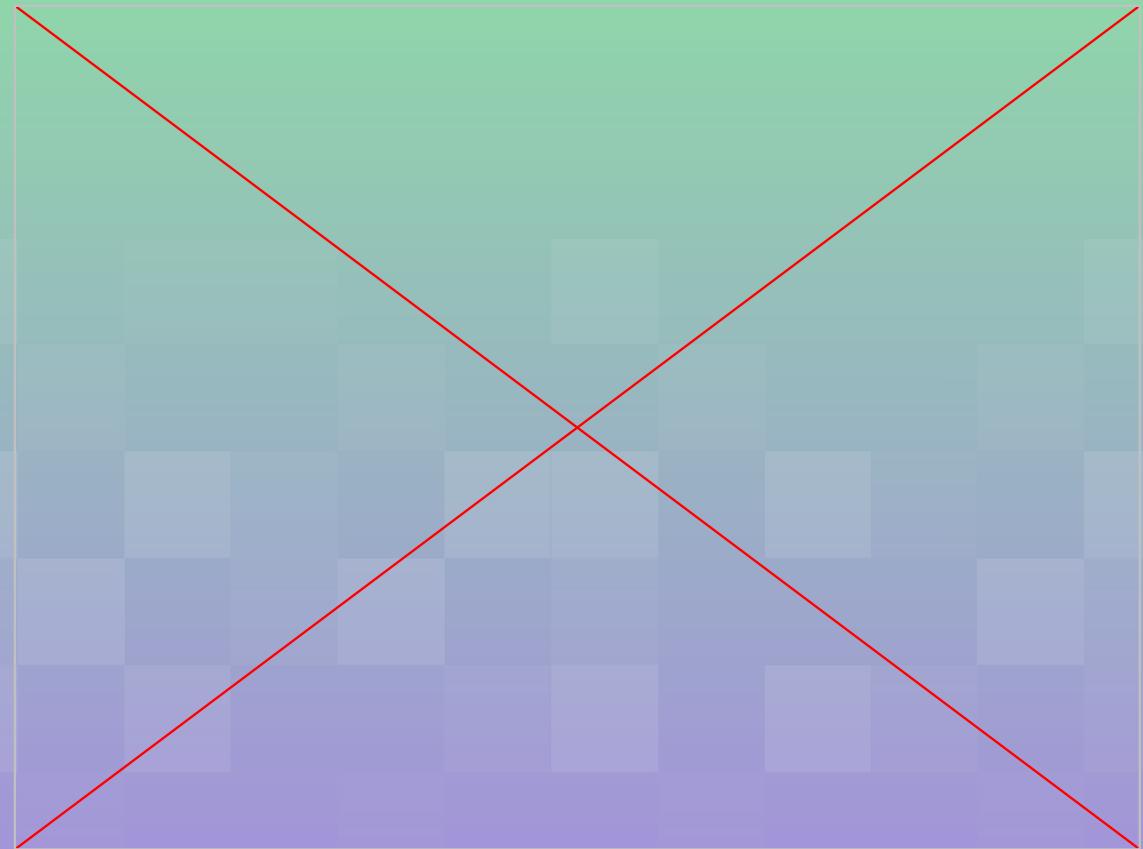
if self.deferrable:
    self.defer(trigger=GxTrigger(self.run_guid),
method_name="process_expectation_results_returned_from_trigger")
else:
    retry_call(self.get_expectation_results, tries=self.get_expectation_config.tries,
               delay=self.get_expectation_config.delay_in_seconds, exceptions=(ValueError,))
```

GxTrigger

```
async def run(self):
    await asyncio.sleep(self.initial_wait_delay)
    session = requests.Session()
    response = session.get(self.endpoint, params={"run_guid": self.run_guid})

    while not "success" in json.loads(response.text):
        await asyncio.sleep(self.wait_delay)
        response = session.get(self.endpoint, params={"run_guid": self.run_guid})
    yield TriggerEvent({"run_guid": self.run_guid, "response": json.loads(response.text)})
```

Demo



About Wix

Motivation

Deferral explained

Implementation

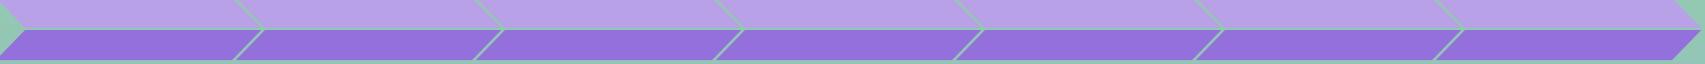
Next Steps

WIX

5.

Next Steps

Rollout



Test locally

Deploy on select DAGs

Monitor

Communicate

Track adoption

Scale triggerer

Set deferral by default



Eitan Shalev Mar 16th at 13:25

Airflow word of the day: **Defer** (di-'fər): to delegate to another:

Introduced in Airflow 2.2, **deferrable operators** allow tasks to "sleep" while waiting on external systems, freeing up worker slots. Instead of using up an airflow worker to run a task that polls an external service for completion, they defer the polling to an asynchronous triggerer service which can efficiently handle many hundreds of polls, till the external task is completed, at which point it returns execution to the worker.

Deferrable operators free up worker slots, resulting in better resource utilization, especially for long-running external requests (e.g., APIs, DB queries).

I encourage you to take a look at Airflow's official [documentation](#) on the topic.

I've recently modified our GX operator to support the optional deferrable mode. Feel free to take a look at the changes I made ([GitHub](#)).

Now, by adding the parameter `deferrable=True` to your GX operator, you will be allowing Airflow to free up a worker slot, so other tasks (in other DAGs) can start running, thus removing one potential bottleneck.

We're planning to gradually update DAGs to use this feature, and once we're certain it's stable, we will change the default behavior to True, so all GX operators, unless explicitly set to False, will defer.

The next steps after that are making more of our operators deferrable, so Airflow can free up workers while other services like Trino or AWS are doing their thing.

So start getting used to seeing purple tasks alongside your green ones .

Needless to say, this could not have been done without [@orene](#) & [@heorhiib](#)'s help, and the support from [@ariksa](#) and his team.

Pitfalls

Ensure gradual rollout

Track Triggerer process performance

Scale up as needed

Track # of concurrent deferred tasks

Don't start by deferring a DAG with 1,000 tasks

All at once

On a weekend



Elad Haziza Jul 18th at 19:50

Hey!

Some of our deferrable tasks seem to be stuck.

I came across this error in the Airflow UI – could someone take a look?



Arik Sasson Jul 19th at 16:15

The triggerer does not appear to be running. Last heartbeat was received
[redacted] seconds ago.

[redacted] and any deferred operator will remain deferred

Great. ThX @Haziza & @orene

Next deferable only with Astronomer...

Next steps

Custom SparkOperators

TrinoOperators (SqlOperator)

Defer natively
deferrable
operators
(Amazon, etc.)

Migrate custom
PythonOperators
to deferrables

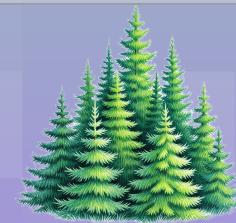


Almog Gelber 17:33

The operator now has the capability to run in deferrable mode. While this functionality has been tested and confirmed to work, we are not enabling it in our environment at the moment because it introduced stability issues with our current Airflow setup.

Expected Impact

| Operator | Worker Hr/Day saved | Cost savings \$/Month | CO ₂ Emissions reduced (Estimate) |
|----------------|---------------------|-----------------------|--|
| GxOperator | 450 (4.5%) | \$600 | 0.6 tCO ₂ / Year |
| SparkOperators | 2,000 (20%) | \$3,200 | 3 tCO ₂ / Year |
| TrinoOperators | 4,000 (40%) | \$6,400 | 6 tCO ₂ / Year |



Conclusions

Deferrable operators
improve performance
& reduce costs

Easier than they seem
at first glance

Do it!

Thank you!

Questions?



/in/eshalev

eitansh@wix.com

WIX