


# Seamless Airflow Upgrades: Migrating from 2.x to 3



**Ankit Chaurasia**

Senior Software Engineer


# Why Upgrade to Airflow 3?



Run anywhere,  
in any language



Scheduler-man  
aged backfills




Modern UI & UX



Event driven  
scheduling &  
Data assets



Task SDK



Lean core and  
Providers



DAG Versioning



Security  
upgrades



Seamless  
upgrades



Edge Executor



Inference and  
hyperparameter  
tuning




Developer  
experience


# **Seamless Airflow 3 Upgrade: Step-by-Step Checklist**

- 01 Take care of the Prerequisites**
- 02 Backup & Clean your existing Airflow instance**
- 03 Verify DAG Compatibility using Ruff**
- 04 Update Airflow Configuration**
- 05 Address known Breaking Changes**
- 06 Upgrade and test in Development**
- 07 Production Upgrade**
- 08 Post-Upgrade Validation**

# Step 1: Versions Check and Prepare for upgrade




**Airflow Version:**  
Must be 2.6.3 or higher (ideally 2.7.x) for a smooth Airflow 3 upgrade path



**Python Version:**  
Airflow 3 requires Python 3.9+; support for 3.7 and 3.8 is dropped



**Database Compatibility:**  
Metadata DB must be PostgreSQL 13+ or MySQL 8+ due to SQLAlchemy 2.x migration



**Providers & Dependencies:** Install apache-airflow-provider-standard plus any additional DAG providers needed




Astronomer users benefit from **managed runtimes**; open-source users must verify and upgrade manually

# Step 2: Clean and Back Up Your Existing Airflow Instance




## Backup Your Database

Make a backup of your Airflow metadata database before starting the migration. Shut down Airflow instances if no hot backup is available to ensure consistency.



## Avoid Migration Risks

A backup prevents issues from failed migrations or network interruptions that could leave your system in a half-migrated state.




## Clean Up Old Data

Use the airflow db clean CLI command to remove unnecessary data like old XComs to reduce database size and speed up schema changes.



## Resolve DAG Errors

Ensure no DAG processing errors remain, such as AirflowDagDuplicateIdException. Run airflow dags reserialize without errors before proceeding.



# Step 3 – Ensure DAG Code Compatibility with Ruff



Install Ruff v0.13.1 or later: `pip install --upgrade ruff`



Run lint checks for breaking changes: `'ruff check dags/ --select AIR301'`.



Preview automated corrections with `--show-fixes` to verify suggested code updates



Apply safe fixes instantly using `--fix`, updating imports and parameters without manual edits



Use `--fix --unsafe-fixes` cautiously for import path changes that may alter behavior



Re-run checks until no AIR rule violations remain, confirming code is fully upgraded



Astronomer users can execute `'astro dev upgrade-test'` to automate all steps including dependency verification

# Ruff Rules for Airflow 3 Upgrade

## AIR30 Rules (Mandatory)

1. Check for removed parameters and imports no longer available in Airflow 3
2. Mandatory changes to ensure DAGs function correctly in Airflow 3

## AIR31 Rules

1. Check for deprecated parameters and imports still available in Airflow 3
2. Recommended changes to maintain compatibility with future Airflow versions

```
$ dags/my_dag.py:19:5: AIR301 [*] `fail_stop` is removed in Airflow 3.0
>   |
> 17 |     start_date=datetime(2025, 1, 1),
> 18 |     schedule="@daily",
> 19 |     fail_stop=True
>   |     ^^^^^^^^^^ AIR301
> 20 |   ):
>   |
>     = help: Use `fail_fast` instead
>
> Found 1 error.
> [*] 1 fixable with the `--fix` option.
```

# Airflow OSS Ruff linter vs. Astro CLI

## Using Ruff Linter (OSS)

1. Install the latest Ruff linter via pip: `pip install --upgrade ruff`
2. Run Ruff on your DAG code: `ruff check --preview --select AIR30`
3. Use `--fix` flag to automatically fix some mandatory issues
4. Ruff outputs errors with suggestions, e.g., renaming deprecated DAG params like `fail\_stop` to `fail\_fast`. Fix the mandatory issues.
5. Requires manual review of breaking changes and Airflow release notes for full compatibility

## Using Astro CLI

1. Run `astro dev upgrade-test` which also includes Ruff linter as part of command and fix the ruff issues.
2. Automatic detection of Airflow 3 compatibility issues within Astro environment

# Step 4: Check and fix breaking Airflow configs



Upgrade to **Airflow 2.11.0** as **airflow config update** is available from this version.



Run **airflow config update** to detect deprecated, moved, or invalid configs in airflow.cfg



Review both **breaking** and **recommended** config changes



Use **airflow config update --fix** to auto-apply fixes with backup creation for old airflow configuration file

The following are the changes in airflow config:


- [BREAKING] Updated default value of 'core/executor' from 'SequentialExecutor' to 'LocalExecutor'.
- [BREAKING] Removed 'logging/log\_filename\_template' from configuration.
- [BREAKING] Renamed 'webserver/web\_server\_host' to 'api/host'.
- [BREAKING] Renamed 'webserver/web\_server\_port' to 'api/port'.
- [BREAKING] Renamed 'webserver/workers' to 'api/workers'.
- [BREAKING] Renamed 'webserver/web\_server\_ssl\_cert' to 'api/ssl\_cert'.
- [BREAKING] Renamed 'webserver/web\_server\_ssl\_key' to 'api/ssl\_key'.
- [BREAKING] Renamed 'webserver/access\_logfile' to 'api/access\_logfile'.
- [BREAKING] Updated default value of 'scheduler/catchup\_by\_default' from 'True' to 'False'.
- [BREAKING] Renamed 'scheduler/dag\_dir\_list\_interval' to 'dag\_processor/refresh\_interval'.
- [BREAKING] Renamed 'triggerer/default\_capacity' to 'triggerer/capacity'.

Dry-run is mode enabled. To apply above airflow.cfg run the command with '--fix'.


# Types of Airflow 3 configuration changes



Default Values Changes



Renamed Options



Removed Options



Invalidated Previously  
Valid Options

# Step 5 – Assess Breaking Changes

**Direct DB Access Removed:** Replace task code that opens DB sessions or queries metadata directly with Airflow public APIs or REST calls.

**Scheduling & Timetables Changes:** Replace `schedule\_interval` with `schedule`. New cron triggers disable automatic backfills by default; enable with `scheduler.create\_cron\_data\_intervals=True`.


**Execution Date Context:** Update DAGs to replace `execution\_date` and related variables with `logical\_date` and `data\_interval\_start/end`.

**Removed Features:** Migrate from SubDAGs to TaskGroups or Datasets; replace SLA miss alerts with external monitoring; update plugins to new system as FAB-based plugins need compatibility provider.

**XCom Serialization:** Disallow pickled XComs by default for security. Migrate to custom backends. Old pickled XComs archived during DB migration.

**Provider Changes:** Operators and hooks moved to separate provider packages requiring explicit installation.

# Step 6 – Upgrade & Test Airflow 3 Locally



Set up a local or  
staging Airflow  
environment


Upgrade Airflow  
to version 3,  
apply all config  
and DAG code  
changes, then run  
airflow db  
upgrade to  
migrate metadata  
schema

Start Airflow 3  
components: API  
server replaces  
webserver; run  
dag-processor  
alongside  
scheduler

Trigger and  
monitor all critical  
DAGs, verifying  
task execution,  
scheduling,  
XComs, and  
templated  
variables

Fix errors  
iteratively and  
re-test until all  
DAGs run cleanly

# Upgrade and Test locally (OSS method)



## Install Airflow 3

Install Airflow 3 (e.g., pip install apache-airflow==3.0.0 with constraints, or use the official Airflow 3 Docker image). Apply the same configuration changes using the updated airflow.cfg and include the DAG code fixed via Ruff.

## Run Ruff linter and airflow config update

Run Ruff's AIR rules linter to identify breaking code changes. Use airflow config update tool to migrate config files to Airflow 3 format.

## Database Migrations

Run database migrations with airflow db upgrade (equivalent to airflow db migrate) to upgrade the metadata DB schema.

## Change your startup script

Run airflow api-server and airflow dag-processor

## Run airflow and test the DAGs

Run database migrations with airflow db upgrade (equivalent to airflow db migrate) to upgrade the metadata DB schema.

# Demo Upgrade and Test locally (OSS method)



# Upgrade and test locally: Using astro cli

01

**Update Astro Project**

Change your Astro project's Dockerfile base image to the new Astro Runtime (Airflow 3) tag from astronomer.io.

02

**Run upgrade compatibility tests and fix**

Use Astro CLI command 'astro dev upgrade-test --runtime-version' to run compatibility tests including dependencies.

03

**Start Airflow 3 locally**

Execute 'astro dev start' or 'astro dev restart' to launch Airflow 3 locally with your project's DAGs and config.

04

**Verify Local Setup and Run DAGs**

Access the Airflow 3 web UI at <http://localhost:8080> to verify that components like API server, scheduler, triggerer, and local Postgres are running.

# Demo Upgrade and test locally: Using astro cli



# Choose Your Production Upgrade Strategy:

## Blue-Green vs In-Place

### Blue-Green Deployment (Recommended)

1. Set up new Airflow 3 environment alongside existing Airflow 2
2. Deploy updated DAGs and config with fresh or migrated metadata DB
3. Test fully before switching traffic to Airflow 3 to minimize downtime
4. Supports easy rollback by switching back to Airflow 2
5. Astronomer users can upgrade via UI or CLI seamlessly
6. Open-source users must provision new infrastructure and handle DB migration

### In-Place Upgrade (Advanced)

1. Upgrade existing environment by installing Airflow 3 and migrating live DB
2. Requires downtime during migration and service restarts
3. Rollback is complex
4. Must update service scripts for new API server and dag-processor

# Essential Q&A & Upgrade Resources



**Astronomer's Upgrade Guide (Airflow 2 → 3)**



**Astronomer Support & Community Slack**



**Official Apache Airflow Documentation – “Upgrading to Airflow 3”**



**Apache Airflow Slack Channel (#airflow-upgrade)**



**Astral (Ruff) Airflow Rules Documentation: Details on AIR301, AIR302 (mandatory) and AIR311, AIR312 (recommended) linter rules**

# QUESTION?

# The 2025 Apache Airflow® Survey is here!

Fill it out to get a free Airflow 3 Fundamentals or DAG Authoring in Airflow 3 certification code

