

# Seamless Integration: Building Applications That Leverage Airflow's Database Migration Framework

# Ephraim Anierobi

Senior Software Engineer @ [astronomer.io](https://astronomer.io)

Committer

PMC Member



# ASTRONOMER

ASTRONOMER

# Airflow metadata Database

Backbone:

- **SQLAlchemy (models)**
- **Alembic (migrations)**

Targets:

- **Postgres**
- **MySQL**
- (SQLite for dev/test)

# Airflow DB migrations: where • how

Location: `airflow/migrations/versions/`

State table: `alembic_version` (tracks current revision)

Use cases: first-run init, schema changes across  
upgrades/downgrades

```
# migrate to latest
airflow db migrate
# migrate to a version
airflow db migrate -n 3.1.0
# check migrations done
airflow db check-migrations
# rollback example
airflow db downgrade -n 2.7.0
# reset
airflow db reset
```

# Airflow 2.x: What **airflow db** commands Actually Migrate

## In scope

- Core Airflow metadata schema (Alembic migrations in core)

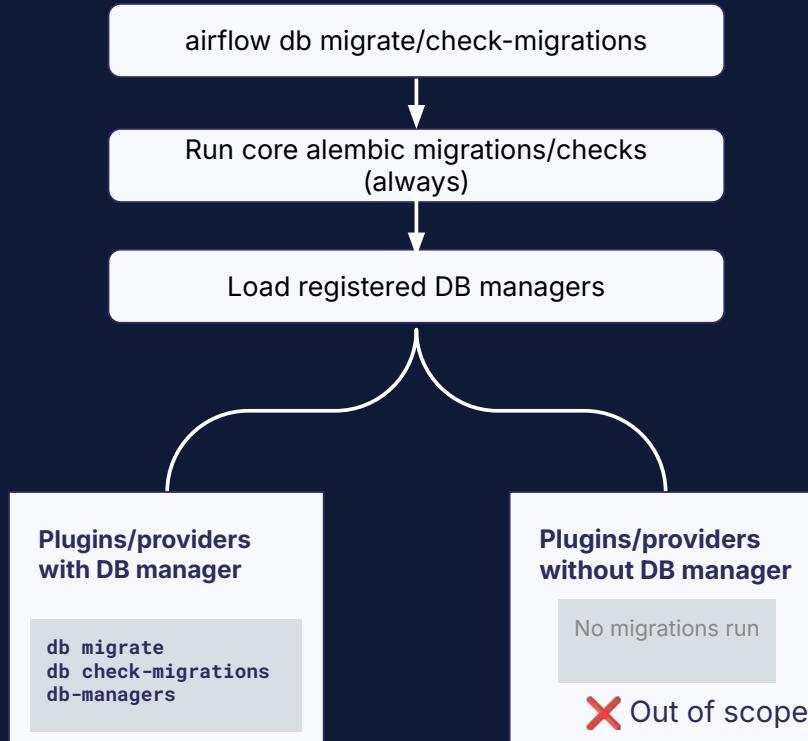
## Out of scope

- Provider/plugin-defined tables
- Non-core schemas must ship/run their own migrations (no auto-discovery)

## Implication

- **Startup wait for migrations, applies only to core.** Airflow components wait for core DB migrations, not your provider's/plugin's migrations.

# Airflow 3: When do `airflow db` commands run non-core migrations?



Scope gate; only DB managers registered are invoked; otherwise nothing happens; core only

# When to add a DB Manager (Provider/Plugin)

## Add a DB Manager if...

- Your provider/plugin owns tables (or modifies its own schema) inside the Airflow metadata DB.
- You want airflow to wait for your migrations to be done via `airflow db check-migrations` (and invoke them during `db migrate`).

## You don't need it if...

- You don't create tables (use only core models).
- Your state lives in an **external service DB** (outside Airflow's metadata DB).

## Decision rule

- If you own schema in Airflow's metadata DB → define a DB Manager. Otherwise → skip it.

**How do we integrate providers/plugins  
migrations in Airflow?**

# Demo app for this talk

A simple plugin including a listener plugin that mocks creating a ticket for every failed dag and logs these tickets with their URL in a database table (that's why we need a custom db manager!).

Github:

<https://github.com/ephraimbuddy/ticketing>

Scan the QR code or use this link to open it on any device



```
alembic init migrations
```

→ ticketing tree

```
.
```

- └── alembic.ini
- └── migrations
  - ├── env.py
  - ├── README
  - └── script.py.mako
- └── versions

3 directories, 4 files

Replace the new alembic.ini with Airflow's alembic.ini file

\*Airflow's alembic.ini file is located at [airflow-core/src/airflow/alembic.ini](#)

```
alembic revision -m "Placeholder migration"
```

*ticketing/migrations/versions/d759c6d30f5a\_placeholder\_migration.py*

```
"""placeholder migration

Revision ID: d759c6d30f5a
Revises:
Create Date: 2025-09-27 08:49:48.267152

"""

from typing import Sequence, Union

# revision identifiers, used by Alembic.
revision: str = "d759c6d30f5a"
down_revision: Union[str, None] = None
branch_labels: Union[str, Sequence[str], None] = None
depends_on: Union[str, Sequence[str], None] = None

def upgrade() -> None:
    pass

def downgrade() -> None:
    pass
```

# Use Airflow's schema and naming conventions

*ticketing/models.py*

```
...  
from airflow.configuration import conf  
from airflow.models.base import naming_convention  
...  
  
SQL_ALCHEMY_SCHEMA = conf.get("database", "SQL_ALCHEMY_SCHEMA")  
  
def _get_schema() -> str | None:  
    """Return the schema to use."""  
    if not SQL_ALCHEMY_SCHEMA or SQL_ALCHEMY_SCHEMA.isspace():  
        return None  
    return SQL_ALCHEMY_SCHEMA  
  
metadata = MetaData(schema=_get_schema(), naming_convention=naming_convention)  
  
mapper_registry = registry(metadata=metadata)  
Base = mapper_registry.generate_base()  
Base.metadata = metadata  
  
class DagRunTicket(Base):  
    ...
```

# Implement a custom DB manager

*ticketing/db\_manager.py*

```
from airflow.utils.db_manager import BaseDBManager  
...  
  
class DRTDBManager(BaseDBManager):  
    metadata = Base.metadata  
    version_table_name = "alembic_version_drt"  
    migration_dir = (PACKAGE_DIR / "migrations").as_posix()  
    alembic_file = (PACKAGE_DIR / "alembic.ini").as_posix()  
    supports_table_dropping = True  
    revision_heads_map = _REVISION_HEADS_MAP
```

```
    def upgradedb(...):  
        ...
```

```
    def downgrade(...):  
        ...
```

# Use the base metadata

*ticketing/db\_manager.py*



```
from ticketing.models import Base  
...  
  
class DRTDBManager(BaseDBManager):  
    metadata = Base.metadata  
    ...
```

# Unique version table name

*ticketing/db\_manager.py*



```
...
class DRTDBManager(BaseDBManager):
    ...
    version_table_name = "alembic_version_drt"
    ...
...
```

## *ticketing/db\_manager.py*

```
class DRTDBManager(BaseDBManager):
    ...
    supports_table_dropping = True
    ...
```

*ticketing/db\_manager.py*

```
...
_REVISION_HEADS_MAP: dict[str, str] = {
    "0.1.0": "d759c6d30f5a",
}

class DRTDBManager(BaseDBManager):
    ...
    revision_heads_map = _REVISION_HEADS_MAP
    ...

```

Linking the custom DB manager  
to the alembic migrations

*ticketing/migrations/env.py*

```
...
from ticketing.db_manager import DRTDBManager

version_table = DRTDBManager.version_table_name

...
target_metadata = DRTDBManager.metadata
...
```

*ticketing/migrations/env.py*



```
...
def include_object(object, name, type_, reflected, compare_to):
    if type_ == "table" and name not in target_metadata.tables:
        return False
    return True
...
```

*ticketing/migrations/env.py*

```
...
context.configure(
    ...
    target_metadata=target_metadata,
    ...
    version_table=version_table,
    include_object=include_object,
)
...
```

*ticketing/migrations/env.py*



```
if config.config_file_name is not None:  
    fileConfig(config.config_file_name)
```



```
if not getLogger().handlers and config.config_file_name:  
    fileConfig(config.config_file_name, disable_existing_loggers=False)
```

```
→ ticketing git:(main) tree ticketing
ticketing
├── alembic.ini
├── db_manager.py
└── listener.py
    ├── migrations
    │   ├── env.py
    │   ├── README
    │   ├── script.py.mako
    │   └── versions
    │       └── d759c6d30f5a_placeholder_migration.py
    ├── models.py
    ├── plugin.py
    └── services.py
```

3 directories, 10 files

# Linking the app DBManager to Airflow



```
[database]
external_db_managers = "ticketing.db_manager.DRTDBManager, path.to.another.dbmanager"
```

# Summary : Enabling Non-Core Migrations

- 1 Initialized Alembic for the plugin schema.
- 2 Implemented a DB Manager that points to that Alembic environment
- 3 Registered the DB Manager so Airflow can discover it.



Note: If your plugin doesn't own tables, you don't need a DB Manager.

# Possibilities

**airflow db check-migrations** will wait for your apps migration

 **airflow db migrate** will run your app's db migration

 **airflow db downgrade** will not run your app's db migration

# DB Manager Commands



```
root@4b1caeb0d3a7:/opt/airflow# airflow db-manager --help
Usage: airflow db-manager [-h] COMMAND ...
```

Manage externally connected database managers

**Positional Arguments:**

COMMAND

downgrade

Downgrade the schema of the external metadata database.

migrate

Migrates the specified external database to the latest version

reset

Burn down and rebuild the specified external database

**Options:**

-h, --help show this help message and exit

```
airflow db-manager migrate "ticketing.db_manager.DRTDBManager" --to-version 0.1.0
```

```
airflow db-manager downgrade "ticketing.db_manager.DRTDBManager" --to-version 0.0.1
```

```
airflow db-manager reset "ticketing.db_manager.DRTDBManager" --skip-init
```

FAB provider is another example implementation of this integration

# Questions?

Scan to access the demo App for the talk



# The 2025 Apache Airflow® Survey is here!

Fill it out to for a free Airflow 3  
Fundamentals or DAG Authoring in  
Airflow 3 certification code

