



Asset-Centric Federated Governance with Airflow



Akanksha Khushboo
Data & AI Architect, Google

The Illusion of Success: Tasks vs. Data Integrity



THE ORCHESTRATOR'S VIEW



SUCCESS

All DAG tasks exit with **Code 0**. The pipeline is completely green.

THE CONSUMER'S REALITY



CORRUPTED

Broken dashboards, **NULL values**, and structurally damaged metrics.

Why Green Pipelines Fail

- **Task Success $\neq$ Data Integrity:** A task exits successfully if the query executes, regardless of semantic corruption.
- **Passive Catalog Trap:** Extracting metadata post-hoc merely documents disasters rather than preventing them.

The Architecture Shift

To resolve the Green Pipeline Syndrome, engineering must transition from passive documentation to **active gatekeeping**.

- **Inline Assertion Gates:** Non-blocking and blocking quality checks directly inside the DAG runtime.
- **Ingestion Gatekeeping:** Enforcing contract and structural validation at the exact point of entry.

Balancing Autonomy and Governance at the Domain Boundary



CENTRAL MONOLITHIC

LOW AGILITY

HIGH TRUST

Centralized control preserves policy compliance but introduces structural friction.

Central team becomes a critical bottleneck for all change requests.

UNREGULATED MESH

HIGH AGILITY

LOW TRUST

Domains publish data independently, scaling consumption speed dramatically.

Creates silent data chaos with broken schemas and unannounced API changes.

ACTIVE FEDERATED MESH

HIGH AGILITY

HIGH TRUST

Decentralized execution backed by global programmatic policy gates.

Automated policy enforcement prevents bad data at the point of origin.

The Tension


- Domains require **autonomy** to move fast, build quickly, and bypass monolithic central ticketing bottlenecks.
- Downstream consumers (AI/ML, Finance) demand strict schema guarantees and highly predictable **SLA predictability**.

The Resolution: Data Contracts

Data Contracts function as **executable code**, programmatically enforcing structural and semantic validation directly at the boundary of ingestion.

This decouples organizational alignment while programmatically guaranteeing continuous, end-to-end data quality.

The Active Governance Architecture Ecosystem



Architectural Principles


- **Separation of Concerns:** Managed Airflow acts as the centralized intelligent Control Plane, leaving resource-intensive compute tasks to BQ.
- **Immutable Isolation:** Unverified raw input data remains strictly locked in a staging zone, unreachable by final downstream pipelines.

Continuous Active Enforcement

Rather than documenting failures post-hoc, Knowledge Catalog programmatically evaluates the physical schema and values against declarative Git-managed contracts during runtime.

Guaranteed Compliance: Only verified records are promoted to the target production analytics layer, completely mitigating downstream disruption.

The Gated Quarantine Pattern: Stopping Bad Data at the Border



The Logic of Programmatic Quality Boundaries

This pattern completely eliminates standard post-hoc, reactive debugging. Rather than discovering bad schemas after data has polluted dashboards, the pipeline programmatically enforces a strict operational gate.

By ensuring that **only pristine data** is promoted, Downstream SLA predictability and high-trust consumption are maintained continuously without organizational friction.

The Binary Decision Formula (fail_on_dq_failure=True)


Promotion Decision =

```
{  
  Promote to Production, if Score = 1.0  
  Halt & Quarantine, if Score < 1.0  
}
```

Operational Circuit Breakers & Frictionless Auditing




CIRCUIT BREAKER TRIPPED



Active loop diverts poisoned payloads instantly to safeguard the analytics core.

DETERMINISTIC COMPLIANCE



Every pipeline failure appends verifiable operational tags directly to line-of-custody.

Immediate Automated Isolation

Airflow monitors the Knowledge Catalog AutoDQ job status. If any contract predicate fails, the orchestrator triggers a circuit breaker, bypassing the MERGE task to keep production pure.

Frictionless DataOps Audit

Natively integrated with Google Cloud's Lineage API, leaving a deterministic audit trail showing exactly which records failed, who produced them, and which assets were protected.

Zero Pipeline Contamination

Bad schemas or corrupted values never touch production databases, maximizing uptime and sustaining reliable consumer downstream SLAs.


Automated Recovery Loop

Airflow leaves execution fingerprints, allowing Data Stewards to trace column-level validation failures back to the source GCS file effortlessly.

Garbage In, Hallucination Out: AI-Ready Data Governance



UNGOVERNED AI



Critical Business Risk

Without automated gateway controls, RAG applications treat raw, unverified data partitions as absolute truth. Downstream LLM agents consume schema drifts, duplicate keys, and stale rows, generating confident yet highly damaging hallucinations directly in production systems.

GOVERNED AI



Deterministic, Safe Generation

Embedding operational checkpoints forces strict schema compliance before publication. By matching only metadata tags marked as **'CERTIFIED'** via the Remote MCP Server, downstream AI agents dynamically isolate active queries, producing deterministic, high-trust answers.

Absolute Truth Assumption

Generative AI pipelines treat retrieved context as absolute truth. Airflow acts as the automated boundary, ensuring only pristine data is presented to the LLM.


Trusted Memory Pools for RAG

Airflow exclusively promotes clean partitions to Gold and issues the 'CERTIFIED' metadata tag.

Airflow-Gated Agentic Consumption

Downstream AI Agents use a Remote MCP Server to scan the Knowledge Catalog, programmatically refusing to answer business queries unless the underlying data holds this Airflow-issued certification.

Airflow is no longer just feeding dashboards; it is creating hallucination-safe sandboxes for autonomous AI swarms.



Federated Governance Operating Model

CENTRAL PLATFORM HUB

Infrastructure & Tooling Custody

The core platform team operates as an enabler, provisioning secure capabilities rather than policing domain workflows.

- **Airflow Maintenance:** Keeps orchestrator architecture scaled, patched, and performing globally.
- **Contract Linting:** Standard ODCS verification and automated compliance patterns.
- **Knowledge Catalog Templates:** Reusable quality assurance frameworks out-of-the-box.

DOMAIN SPOKES (E.G., CHECKOUT)

Business Logic & Ownership

Subject matter experts close to the data sources maintain local autonomy and operational lifecycle stewardship.

- **Business Rules:** Write and maintain domain-specific schemas and logic.
- **Pipeline Deployment:** Author, deploy, and monitor localized processing.
- **Incident Response:** Directly own and resolve tripped quality circuit breakers.

Scale-Free Hub-and-Spoke Topology

Traditional centralized data governance creates perpetual execution bottlenecks and structural blind spots. Federated Hub-and-Spoke models shift **accountability to the source**, allowing dozens of autonomous business teams to publish high-trust data at maximum velocity.

Shift-Left Accountability Pattern

By defining strict contracts at the producer boundary, bad data is caught and isolated long before reaching downstream consumers. This dual layout decouples global enterprise standardizations from rapid localized business changes.

Strategic Roadmap: Preparing for Native Airflow 3 Primitives

EVOLUTIONARY LEAP

Native Data-Aware Engine

The current Quarantine + Knowledge Catalog pattern successfully governs assets, but Airflow 3.0 moves these validations directly into the core execution engine with unified native primitives.

- **Physical Logical Entities:** Shift completely from scheduling pipelines on rigid cron timers to pipelines listening to physical entities (Assets).
- **Event-Driven Primitives:** Instantaneous, low-latency execution triggered natively at the system boundary the moment data stabilizes.

FUTURE HORIZONS

Domain Explorations & GCP

As native capabilities mature on GCP, enterprise domains must transition to specialized, decentralized tooling built for planet-scale data workloads.

- **Native DAG Versioning:** Track and execute parallel historical structures without code-level collisions across multi-tenant clusters.
- **Scoped Governance:** Rigorous security bounds restricting resource manipulation to specific localized domain spaces.
- **Specialized SDKs:** Rapid development templates allowing business teams to write declarative code without infrastructure overhead.

The Paradigm Shift: Primitives vs. Pipelines

Apache Airflow 3.0 completely decouples global enterprise standardizations from rapid, localized business logic modifications. By establishing native **Asset-level primitives**, modern data engineering teams escape the limitations of strict time-based orchestration. Instead of patching complex validation logic on top of legacy frameworks, the platform natively guarantees high-velocity, high-trust publishing across scale-free hub-and-spoke networks.

Securing the Future of Data Strategy

SUMMARY & CONCLUSION

Key Architectural Takeaways

As systems scale, maintaining deterministic quality requires a paradigm shift across orchestration, execution, and semantic governance.

Shift from managing "Tasks" to delivering high-trust "Data Products": Modern data mesh architectures demand that business domains own their data lifecycle entirely, exposing certified, contract-governed assets rather than raw, unmonitored tables.

Leverage Airflow as a pure control plane: Delegate intensive computations to hyper-scale cloud platforms using GCP native pushdown capabilities like Knowledge Catalog and BigQuery, combining dynamic orchestration with maximum execution efficiency.

Mandatory active governance for safe enterprise AI adoption: RAG pipelines and LLM agents demand zero-trust input boundaries. Catching semantic drifts, schema mismatches, and data corruption at the source is the only way to eliminate downstream hallucinations.



Questions?