

Democratized data workflows at scale

Emil Todorov

Mihail Petkov



Our agenda for today


- Why Airflow?
- Architecture
- Security
- Execution environment in Kubernetes



FT is a data driven organization




Time for a change





Why Airflow?



Scalable


Extendable


Dynamic

Elegant

Architecture


Architecture






Airflow will be used by multiple teams


Airflow requirements






Teams will share Airflow resources

Airflow shared components



Teams will share Kubernetes resources

Kubernetes shared components






ISSUE



SECURITY • SECURITY



How to evolve this architecture?





This way

This way

Airflow instance per team

One instance components





Instance per team problems

- Adding new team is **hard**
- Maintaining environment per team is **difficult**
- Releasing new features is **slow**
- Resources are **not fully utilised**
- Total **cost increase**






Another way?

Multitenancy

Multiple independent instances in a shared environment


Multi-tenant components




How to make AWS multi-tenant?



IAM Security



IAM Security



How to enhance Kubernetes?





System namespace

Airflow scheduler

Airflow web server

Team 1 namespace

Service Account

Resource Quota

Team 1
worker
Pod

Team 1
worker
Pod

Team 2 namespace

Service Account

Resource Quota

Team 2
worker
Pod

Team 2
worker
Pod

Team N namespace


Service Account

Resource Quota

Team 3
worker
Pod

Team 3
worker
Pod

How to improve PostgreSQL?






How to extend Airflow?



Redesign Airflow source code

Redesign Airflow source code

- Module per team



Redesign Airflow source code

- Module per team
- Connections per team

```
class ExtendedConnection(Connection):  
  
    @staticmethod  
    def get(conn_id: str) -> str:  
        team_id = DAGMetaService.get_team_id_from_dag()  
        return team_id + '_' + conn_id
```

Redesign Airflow source code

- Module per team
- Connections per team
- Extend hooks, operators and sensors

```
class ExtendedS3Hook(S3Hook):  
  
    def __init__(self, *args, **kwargs) -> None:  
        super().__init__(self, *args, **kwargs)  
        self.aws_conn_id = ExtendedConnection.get(self.aws_conn_id)
```


Redesign Airflow source code

- Module per team
- Connections per team
- Extend hooks, operators and sensors
- Use **airflow_local_settings.py**

```
def policy(task_instance: TaskInstance):
    team_id = get_team_id_from_dag_filepath(task_instance.dag_filepath)
    task_instance.executor_config['KubernetesExecutor']['labels']['team_id'] = team_id


def pod_mutation_hook(pod: Pod):
    team_id = pod.labels.get('team_id')
    pod.namespace = get_team_namespace(team_id)
```

Redesign repository structure




Execution environment in Kubernetes


ETL



Extract




Load



Transform?



Example workflow




Our goals

-  Language agnostic jobs

-  Cross task data access

KubernetesPodOperator



Our goals

 Language agnostic jobs


 Cross task data access

Unique storage pattern


- Unique team name from the multitenancy
- Unique DAG id
- Unique task id per DAG
- Unique execution date per DAG run

`/{{team}}/{{dag_id}}/{{task_id}}/{{execution_date}}`

The power of extensibility



ExecutionEnvironmentOperator



Configurable cross task data dependencies

Example input configuration


```
example_execution_environment_config = {  
    "input": {  
        "operators": [  
            {  
                "task_id": "task_1"  
            },  
            {  
                "task_id": "task_2"  
            }  
        ]  
    }  
}  
  
example_execution_environment_operator = ExecutionEnvironmentOperator(  
    task_id='task_3',  
    job_config=example_execution_environment_config,  
    image='example_docker_image',  
    tag='latest'  
)
```

Example output configuration

```
example_execution_environment_config = {  
    "output": {  
        "default": {  
            "destinations": [  
                {  
                    "type": "s3",  
                    "data": {  
                        "bucket": "<s3_bucket>",  
                        "path": "<s3_path>",  
                        "aws_conn_id": "<s3_connection_for_upload>"  
                    }  
                }  
            ]  
        }  
    }  
}  
  
example_execution_environment_operator = ExecutionEnvironmentOperator(  
    task_id='task_3',  
    job_config=example_execution_environment_config,  
    image='example_docker_image',  
    tag='latest'  
)
```


Pre-execute

- Bootstrap the environment
- Enrich the configuration
- Export the configuration to the execution environment pod




Post-execute

- Handle the execution
- Clear all bootstraps
- Deal with the output



POC with AWS S3 as intermediate storage




Is this efficient?

- ✗ Multiple downloads and uploads
- ✗ Single processing power
- ✗ Always loading the data in memory


How to evolve the execution environment?

-  Remove unnecessary data transfers
-  Parallelize the processing
-  Provide hot data access


Shared file system



Kubernetes persistent volume



Kubernetes persistent volume with EFS



So far so good

 Remove unnecessary data transfers

 Parallelize the processing

 Provide hot data access

One worker?







Benefits from Spark

- Runs perfectly in Kubernetes
- Supports many distributed storages
- Allows faster data processing
- Supports multiple languages
- Easy to use

SparkExecutionEnvironmentOperator




Spark execution environment




Our current state

- ✓ Remove unnecessary data transfers
 - ✓ Parallelize the processing
 - Provide hot data access

Hot & cold data



Alluxio



Thank you!

#apacheairflow