

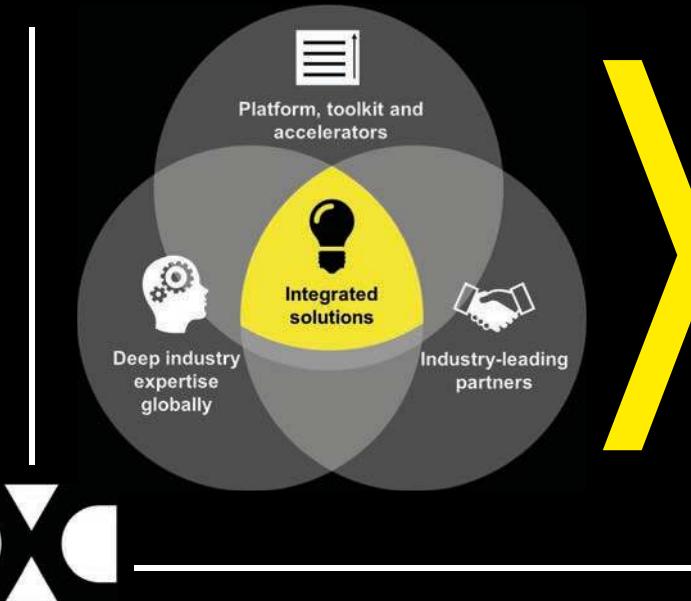
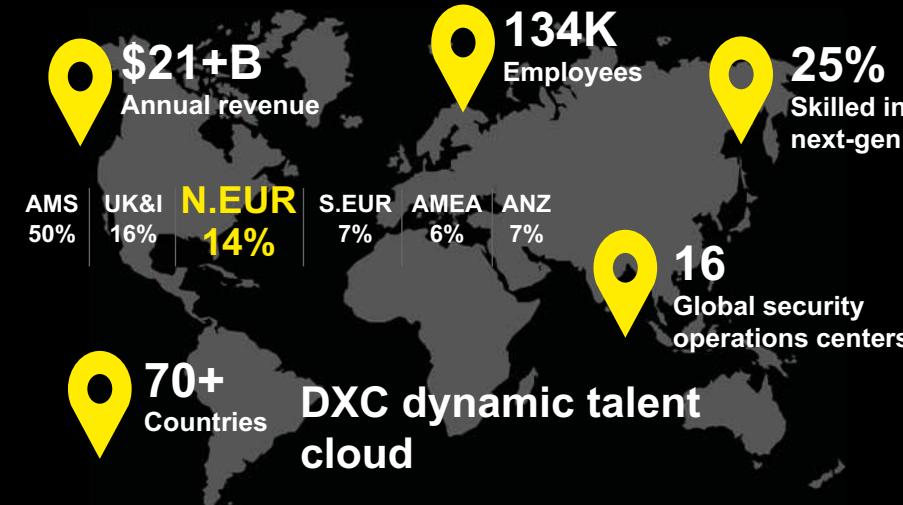
Autonomous Driving with Airflow

Where big-data meets high performance computing

Amr Noureldin – Solution Architect
Michal Dura – Big Data Engineer

The world's leading independent, end-to-end IT services company

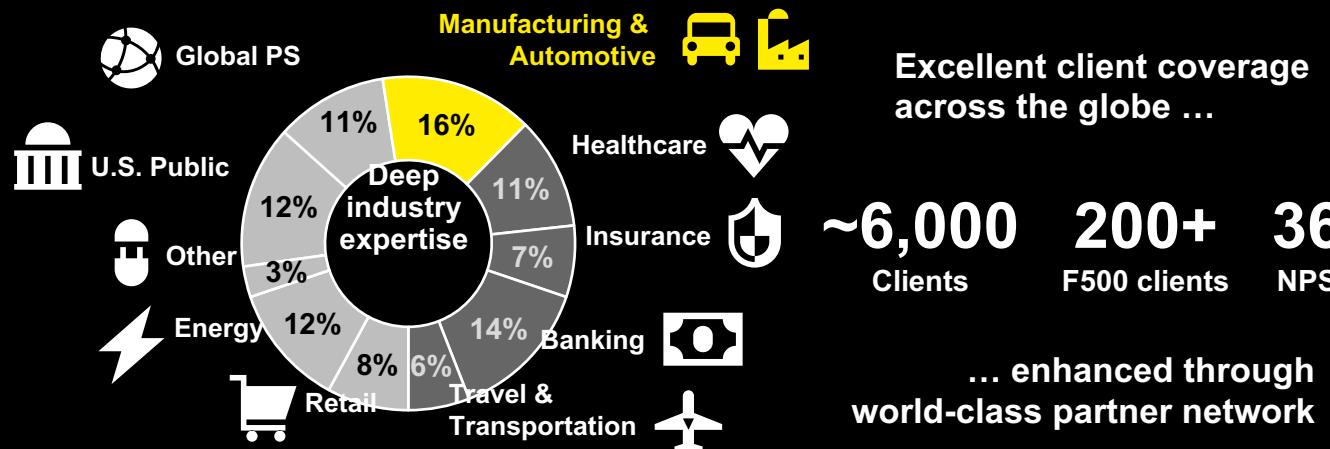
Scale & Skills



DXC Value for AD

- Accelerate time to market
- Reduce cost and risk
- Improve market leadership

Customer Intimacy

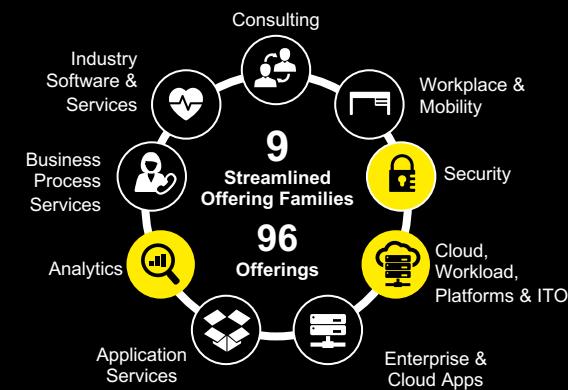


DXC.technology

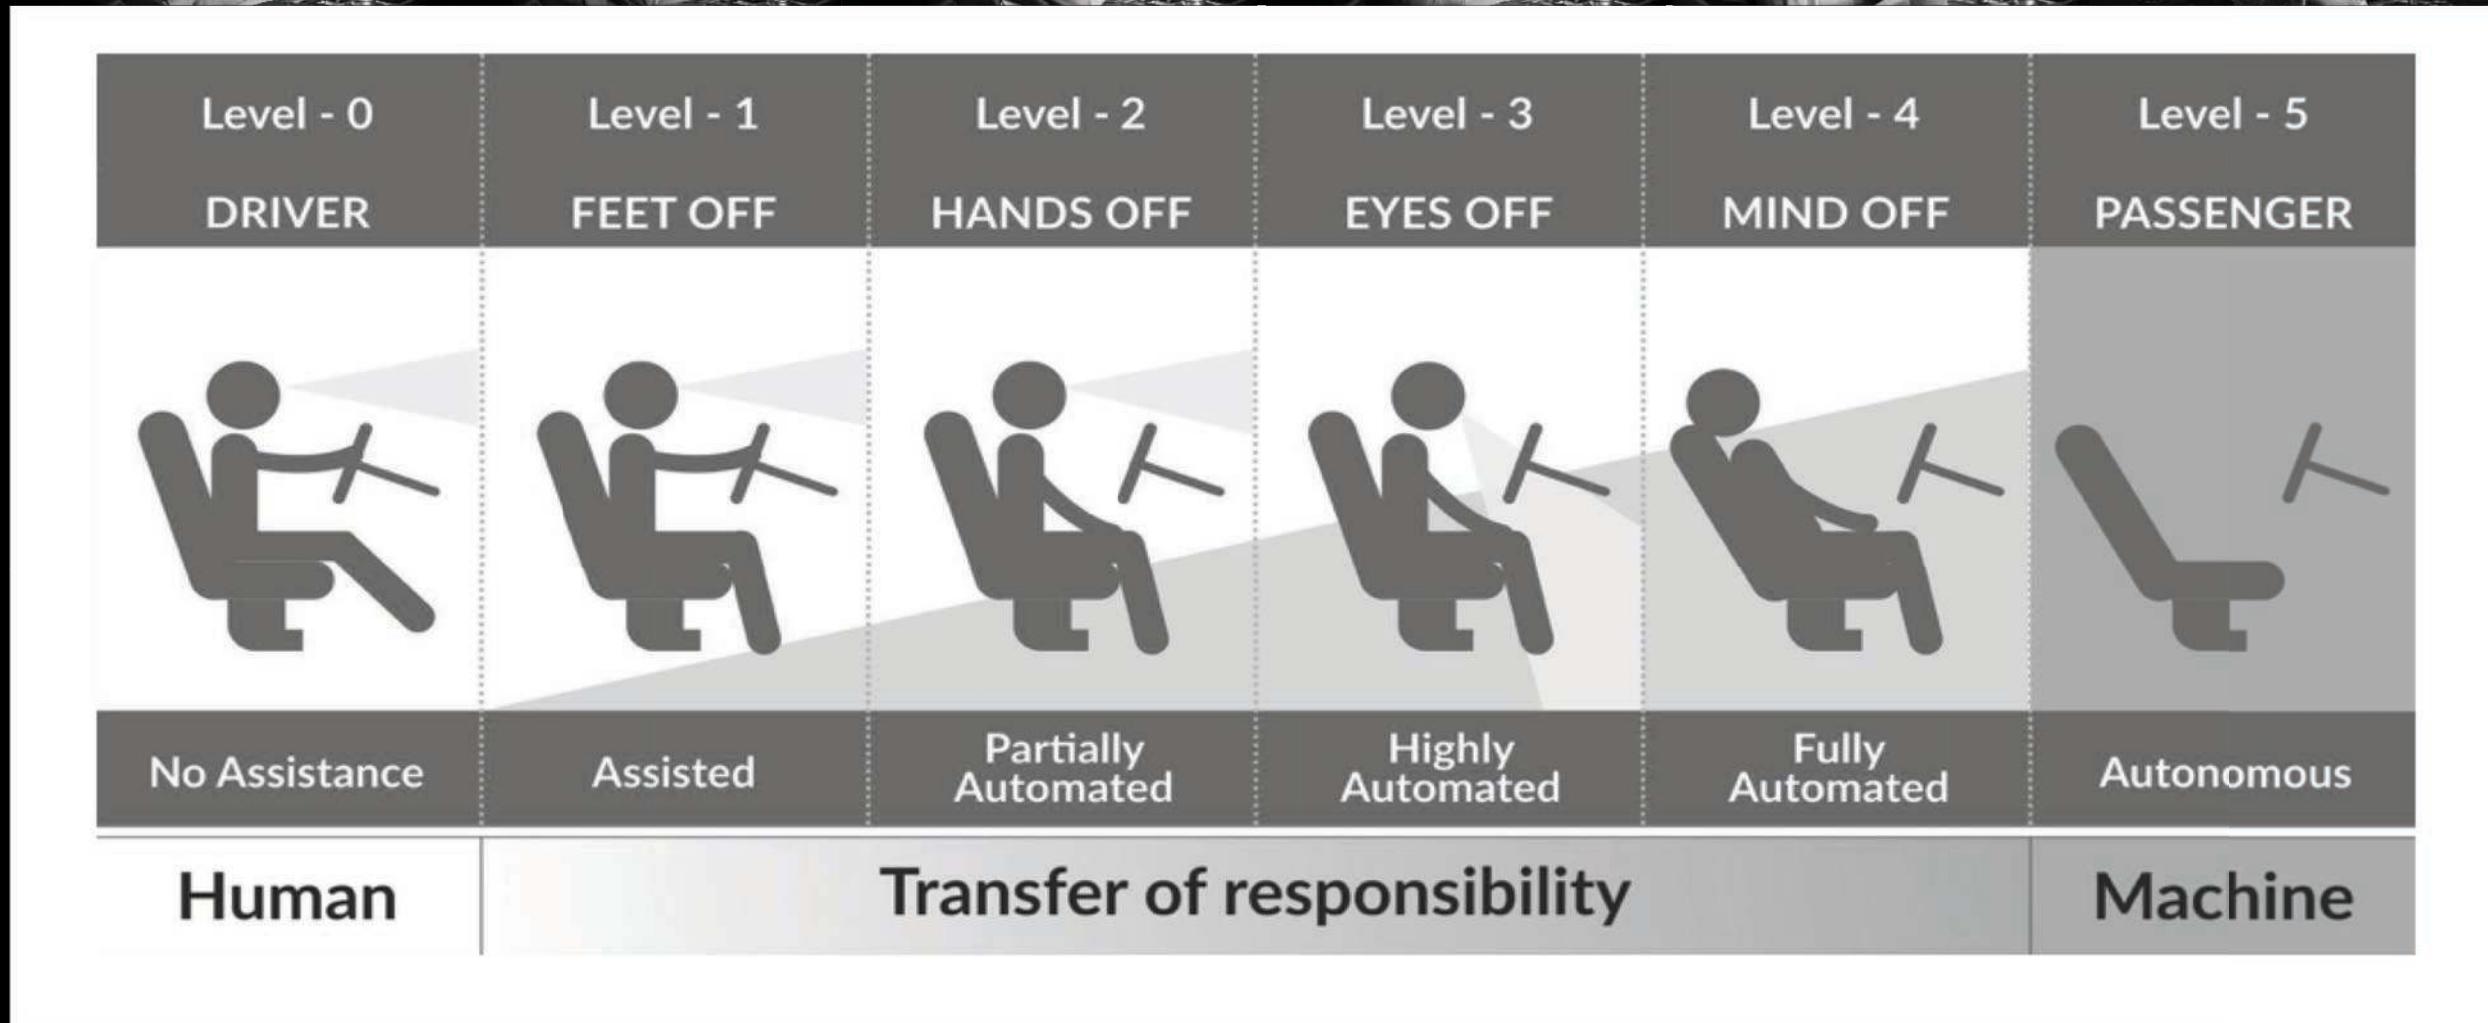
Technology-Driven Innovation

\$4B Digital revenue
250+ global partners
14 strategic co-investing partners

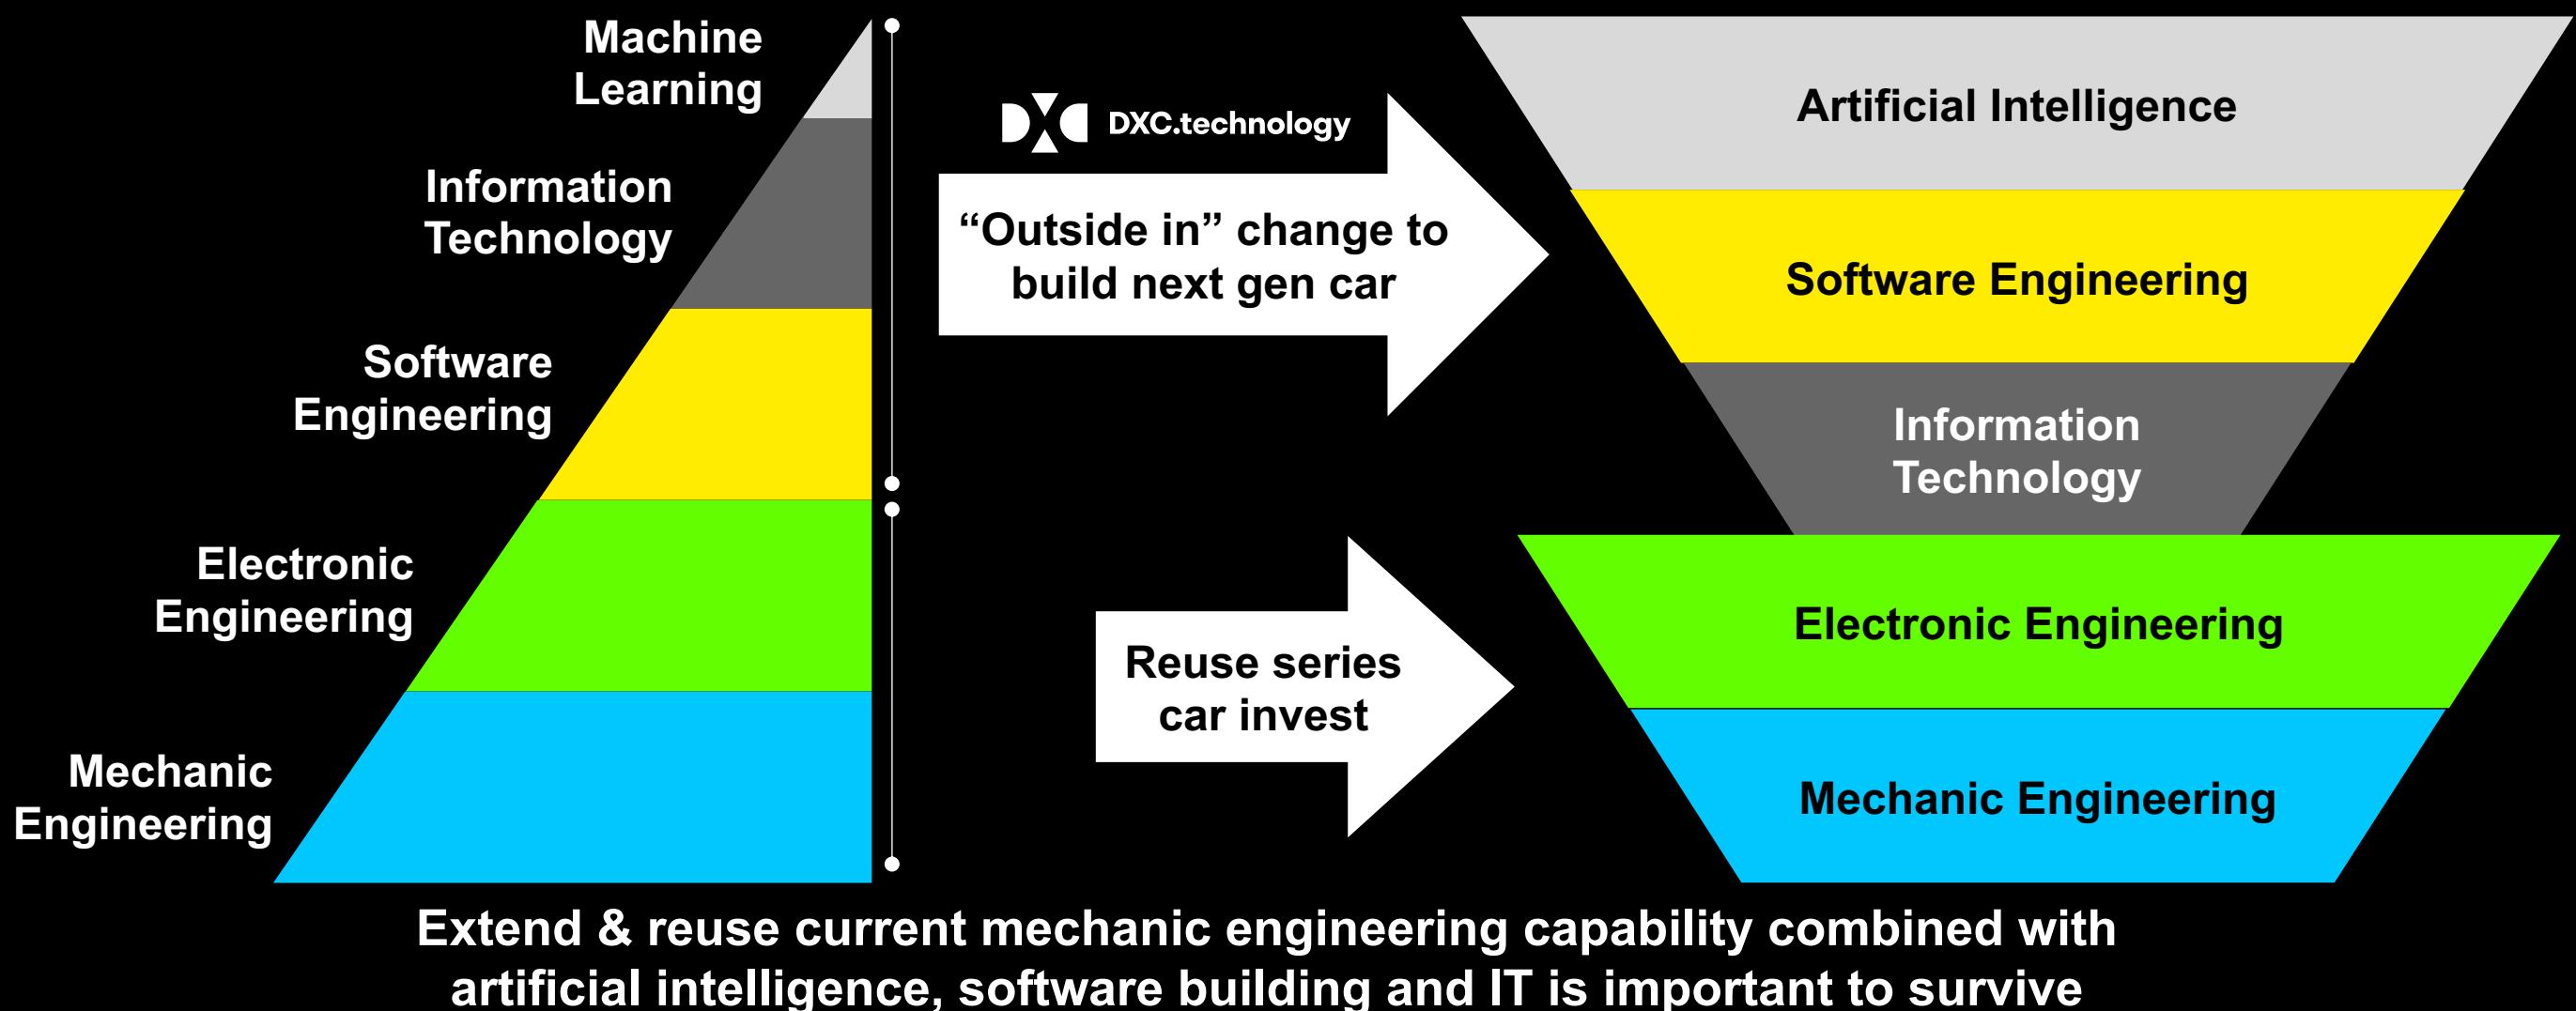
Streamlined offerings



Autonomous Driving



R&D in Automotive Industry – Capabilities are Changing

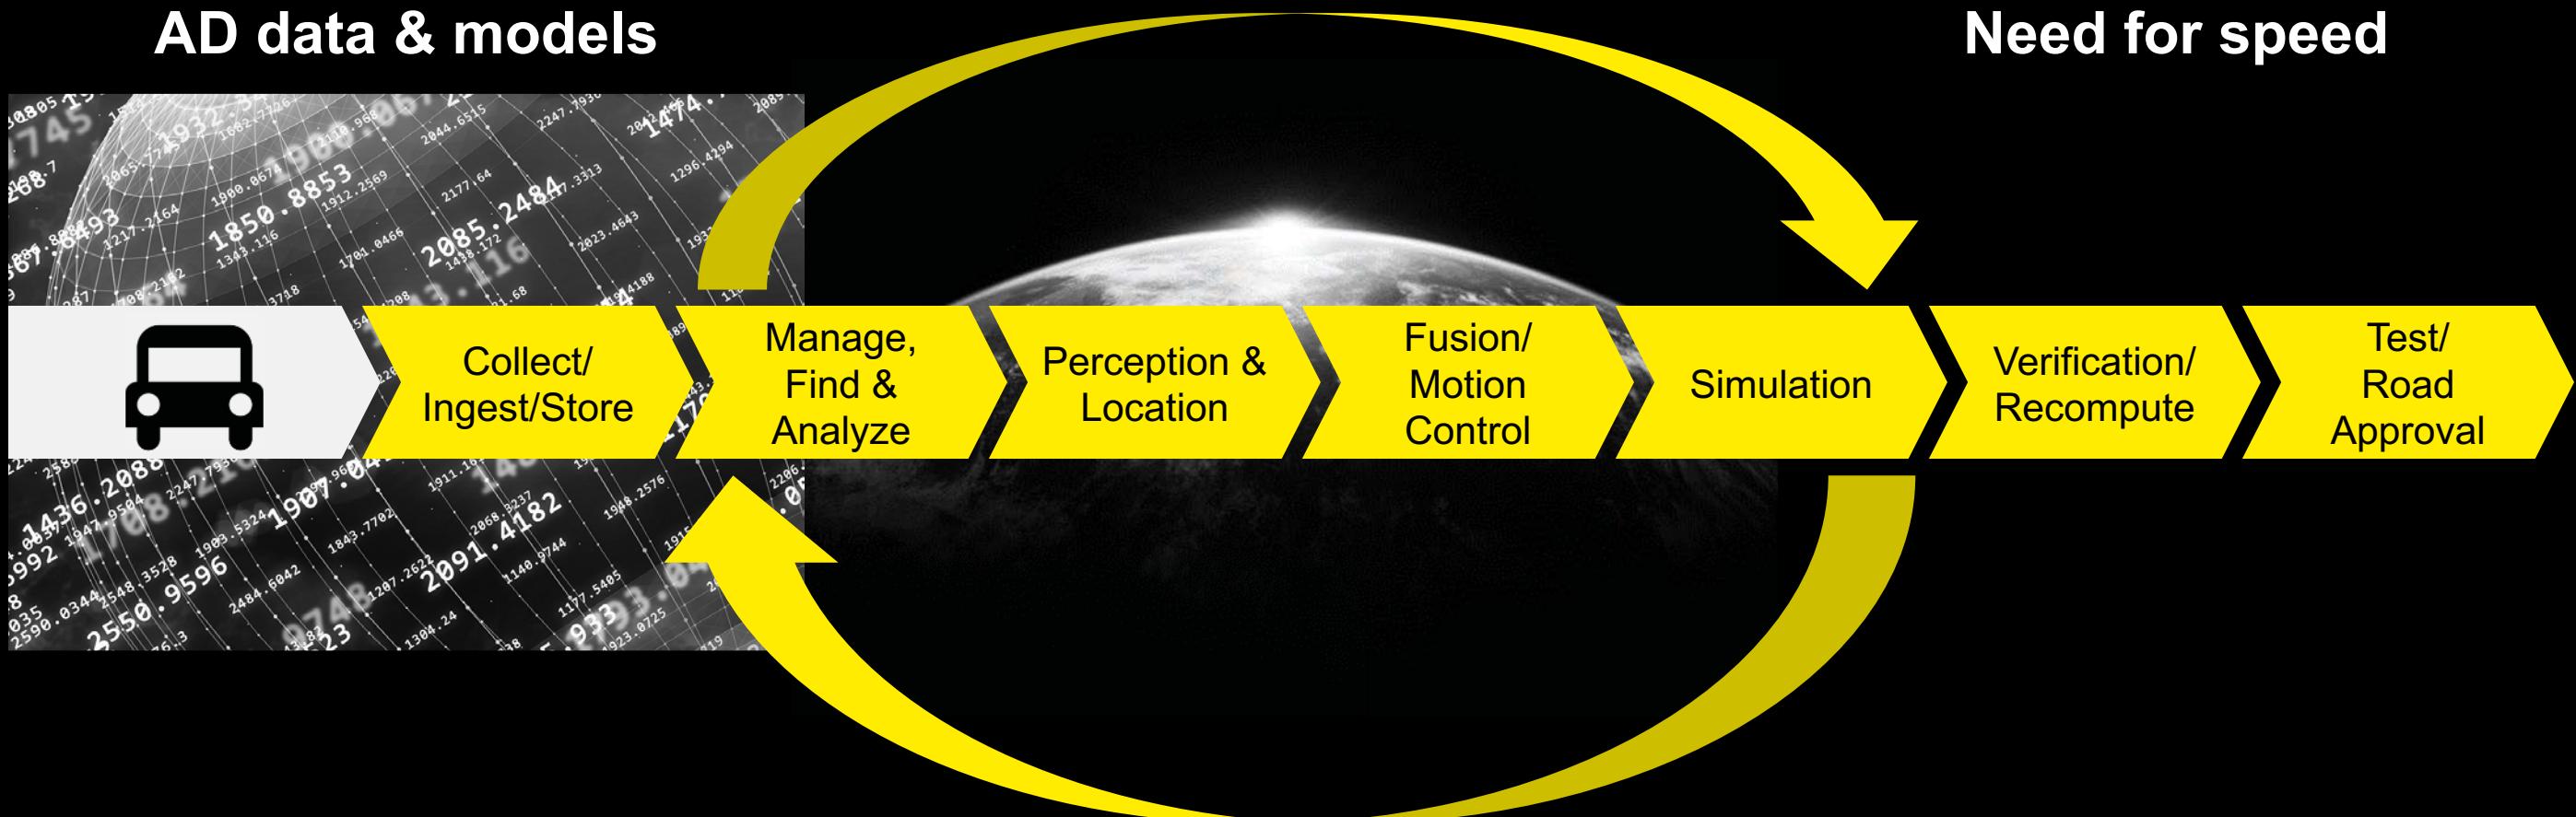


Requires an end-to-end data and AI capability ecosystem for AD development

Geographically distributed
R&D teams

AD data & models

Need for speed



HIGH PERFORMANCE DATA DRIVEN DEVELOPMENT PLATFORM

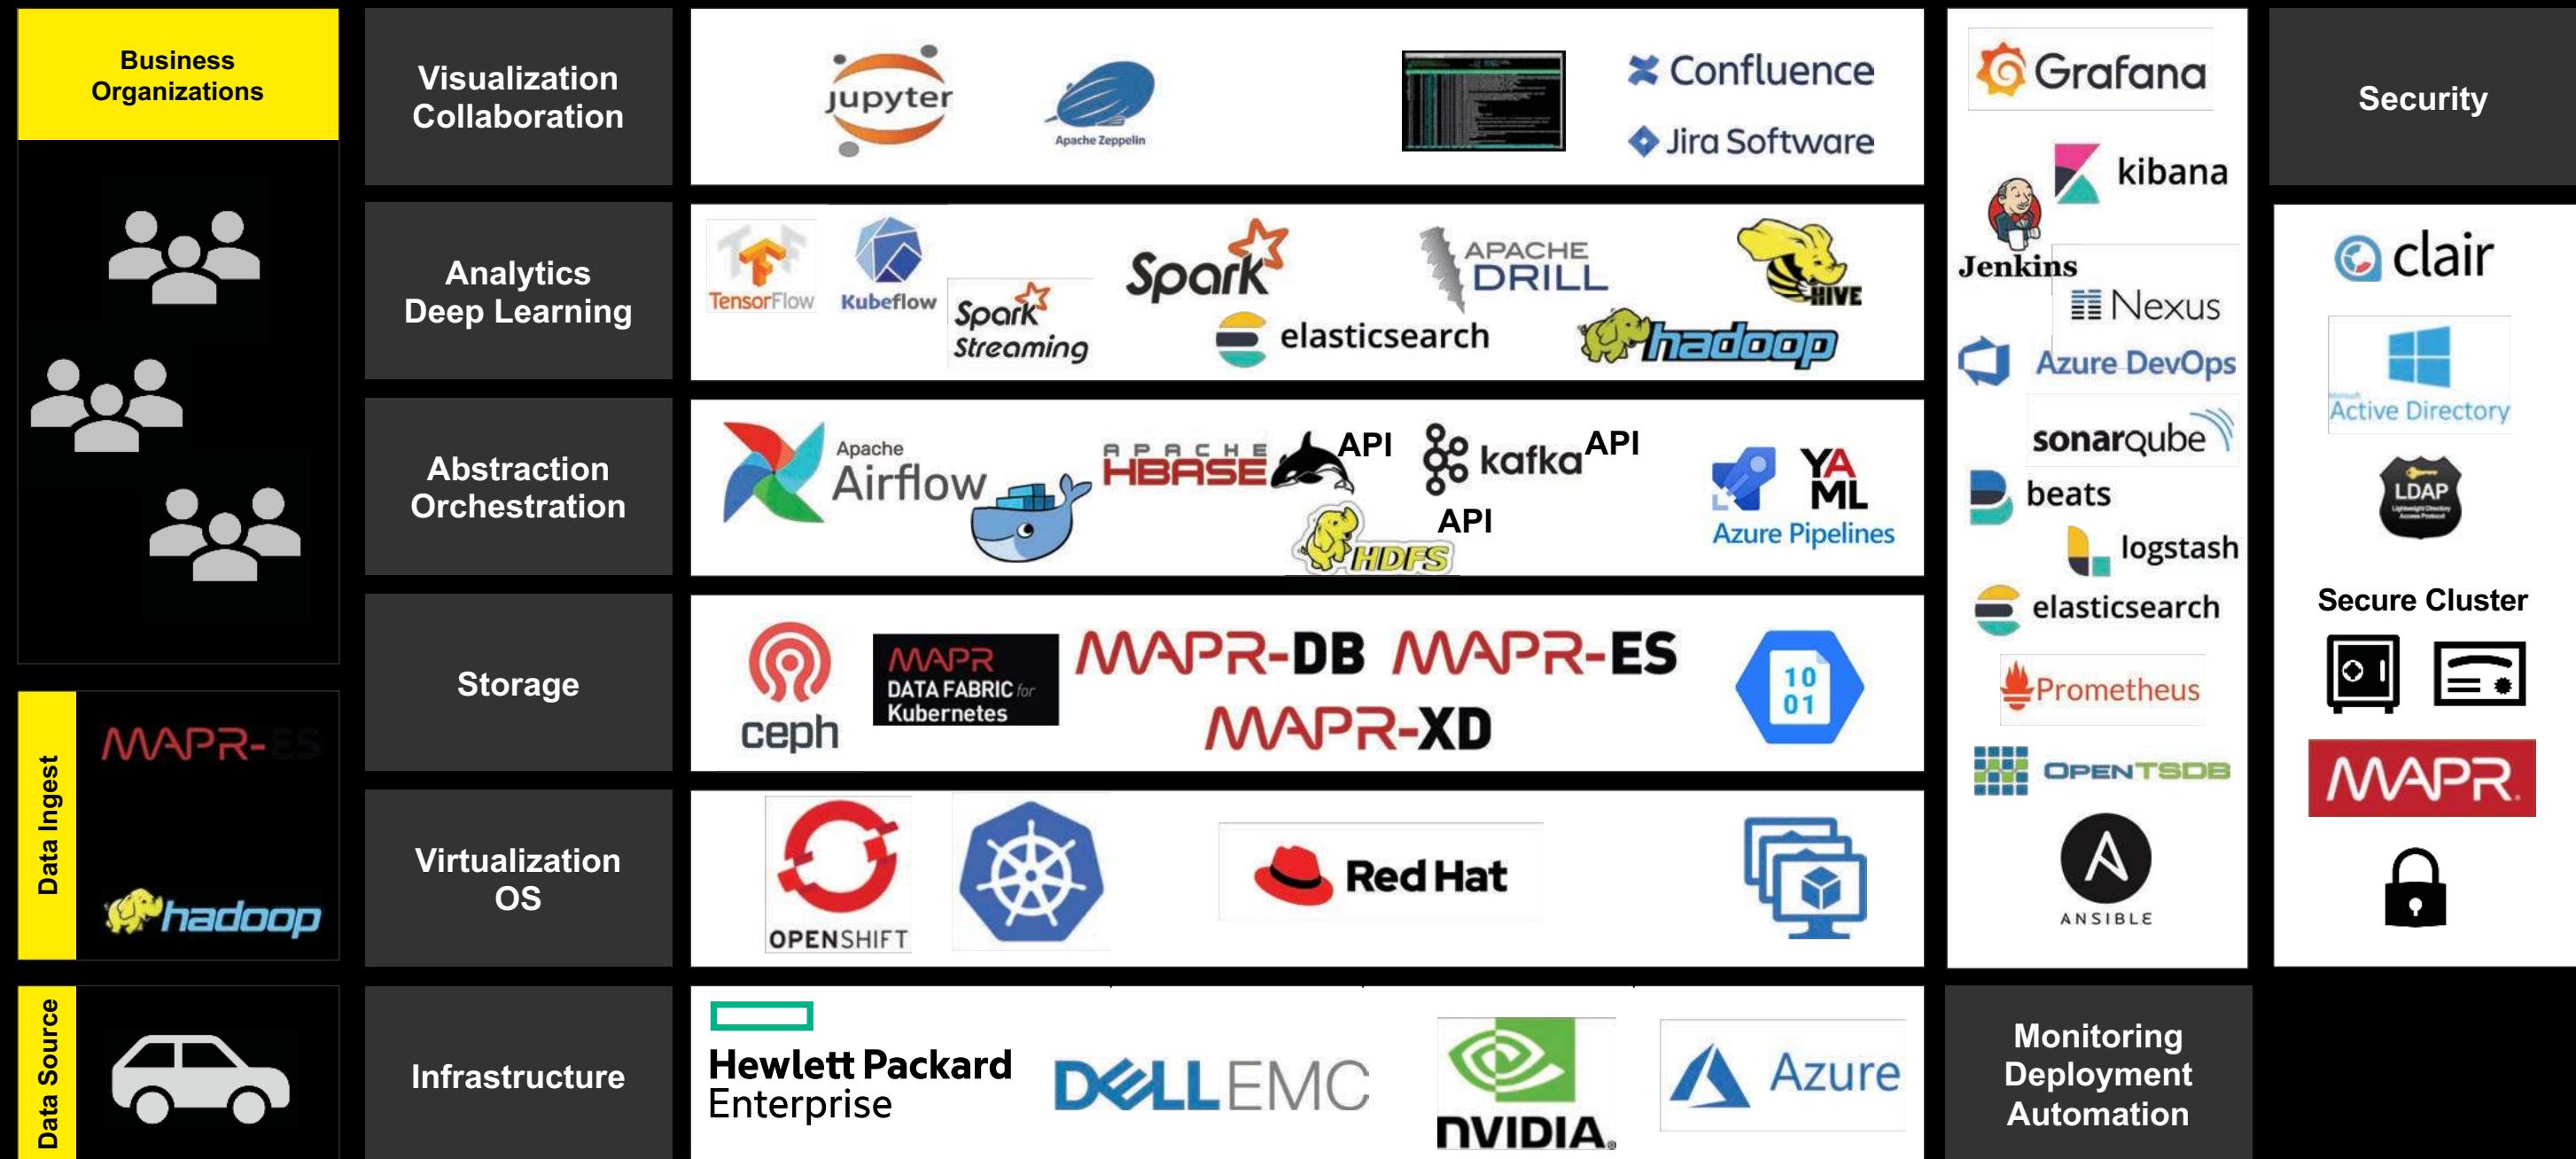
FACTS & FIGURES



DXC.technology

- >200 PB superconverged
- >100.000 processor cores
- Setup in only 3 months
- >200 GPUs
- 96 x 100 Gpbs to ADC
- >1.150 m²
>2.3 MW
- Multi-tenancy

TECHNOLOGY STACK



Autonomous Driving with Airflow – In a Nutshell (1/3)

- **Airflow on OpenShift**
 - 1 scheduler instance – big risk
 - Multiple webservers – load balanced (Rest API Calls)
 - Numerous workers – multiple queues
- **Automated deployment via Helm charts**
 - Various configurations for different instances (also different Airflow versions)
 - History tracking via version control
- **On-demand Airflow instances (ex: development purposes, isolated testing – on a service level)**

Autonomous Driving with Airflow – In a Nutshell (2/3)

- **Integration with the Data and Storage Platform – MapR**
 - Loading DAGs from different locations
 - Loading job configurations, used by the different operators
- **Integration with the Compute Platform – OpenShift**
 - KubernetesPodOperator
 - KubernetesExecutor

Autonomous Driving with Airflow – In a Nutshell (3/3)

- **Metrics Collection & Monitoring**
 - StatsD → Prometheus → Grafana
- **Log collection and aggregation: ElasticSearch + Kibana**
- **Large scale orchestration: aiming at orchestrating jobs at the scale of 100,000's / month**
 - Ingestion, simulation, reprocessing, machine learning, ...etc
 - Complex DAG dependencies

Apache Airflow - Robotic Drive orchestrator

Platform Orchestration Requirements

 Open Source

 Scalability

 Easy to adapt / extend

 Active community

What do we Orchestrate?



Data Ingestion



Machine Learning



Reprocessing



Simulation Jobs

Journey from PoC to Production

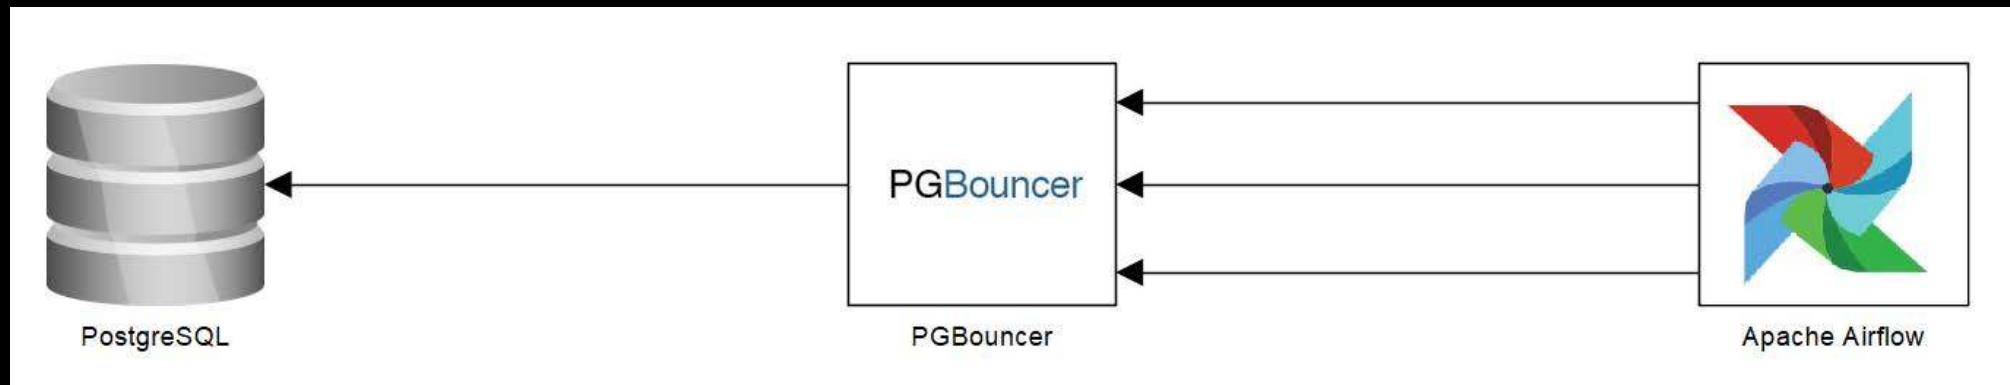
Airflow at Robotic Drive – the beginnings

- Initial work started at Q2 2019
- Airflow 1.10.2 with CeleryExecutor
- PostgreSQL 9.4
- RabbitMQ



Technical Challenges and Lessons Learned

- **Airflow stress and scalability tests**
- **Bottlenecks in the Architecture:**
 - PostgreSQL connection scalability: directly proportional relationship between number of running tasks and database connections
 - Scheduler configuration & performance



A woman with long dark hair is seated in the driver's seat of a car. She is wearing a light-colored button-down shirt and jeans. She is looking down at a tablet device she is holding in her hands. The car's interior is visible, including the steering wheel, dashboard, and side mirror. The background outside the car shows a blurred landscape, suggesting motion or a parking lot.

July 16, 2020

Tailor-made Solutions

Operators / Hooks Customizations (1/3)

Customization and standardization of SparkSubmitOperator

Included „properties_file” in operator constructor

- Spark application configuration can be provided via separate properties file
- It allows submitting jobs via Airflow, in the same fashion as submitting them standalone from the Hadoop cluster

Extend list of parameters where templating is supported

- Better DAGs reusability
- Reduced code duplication

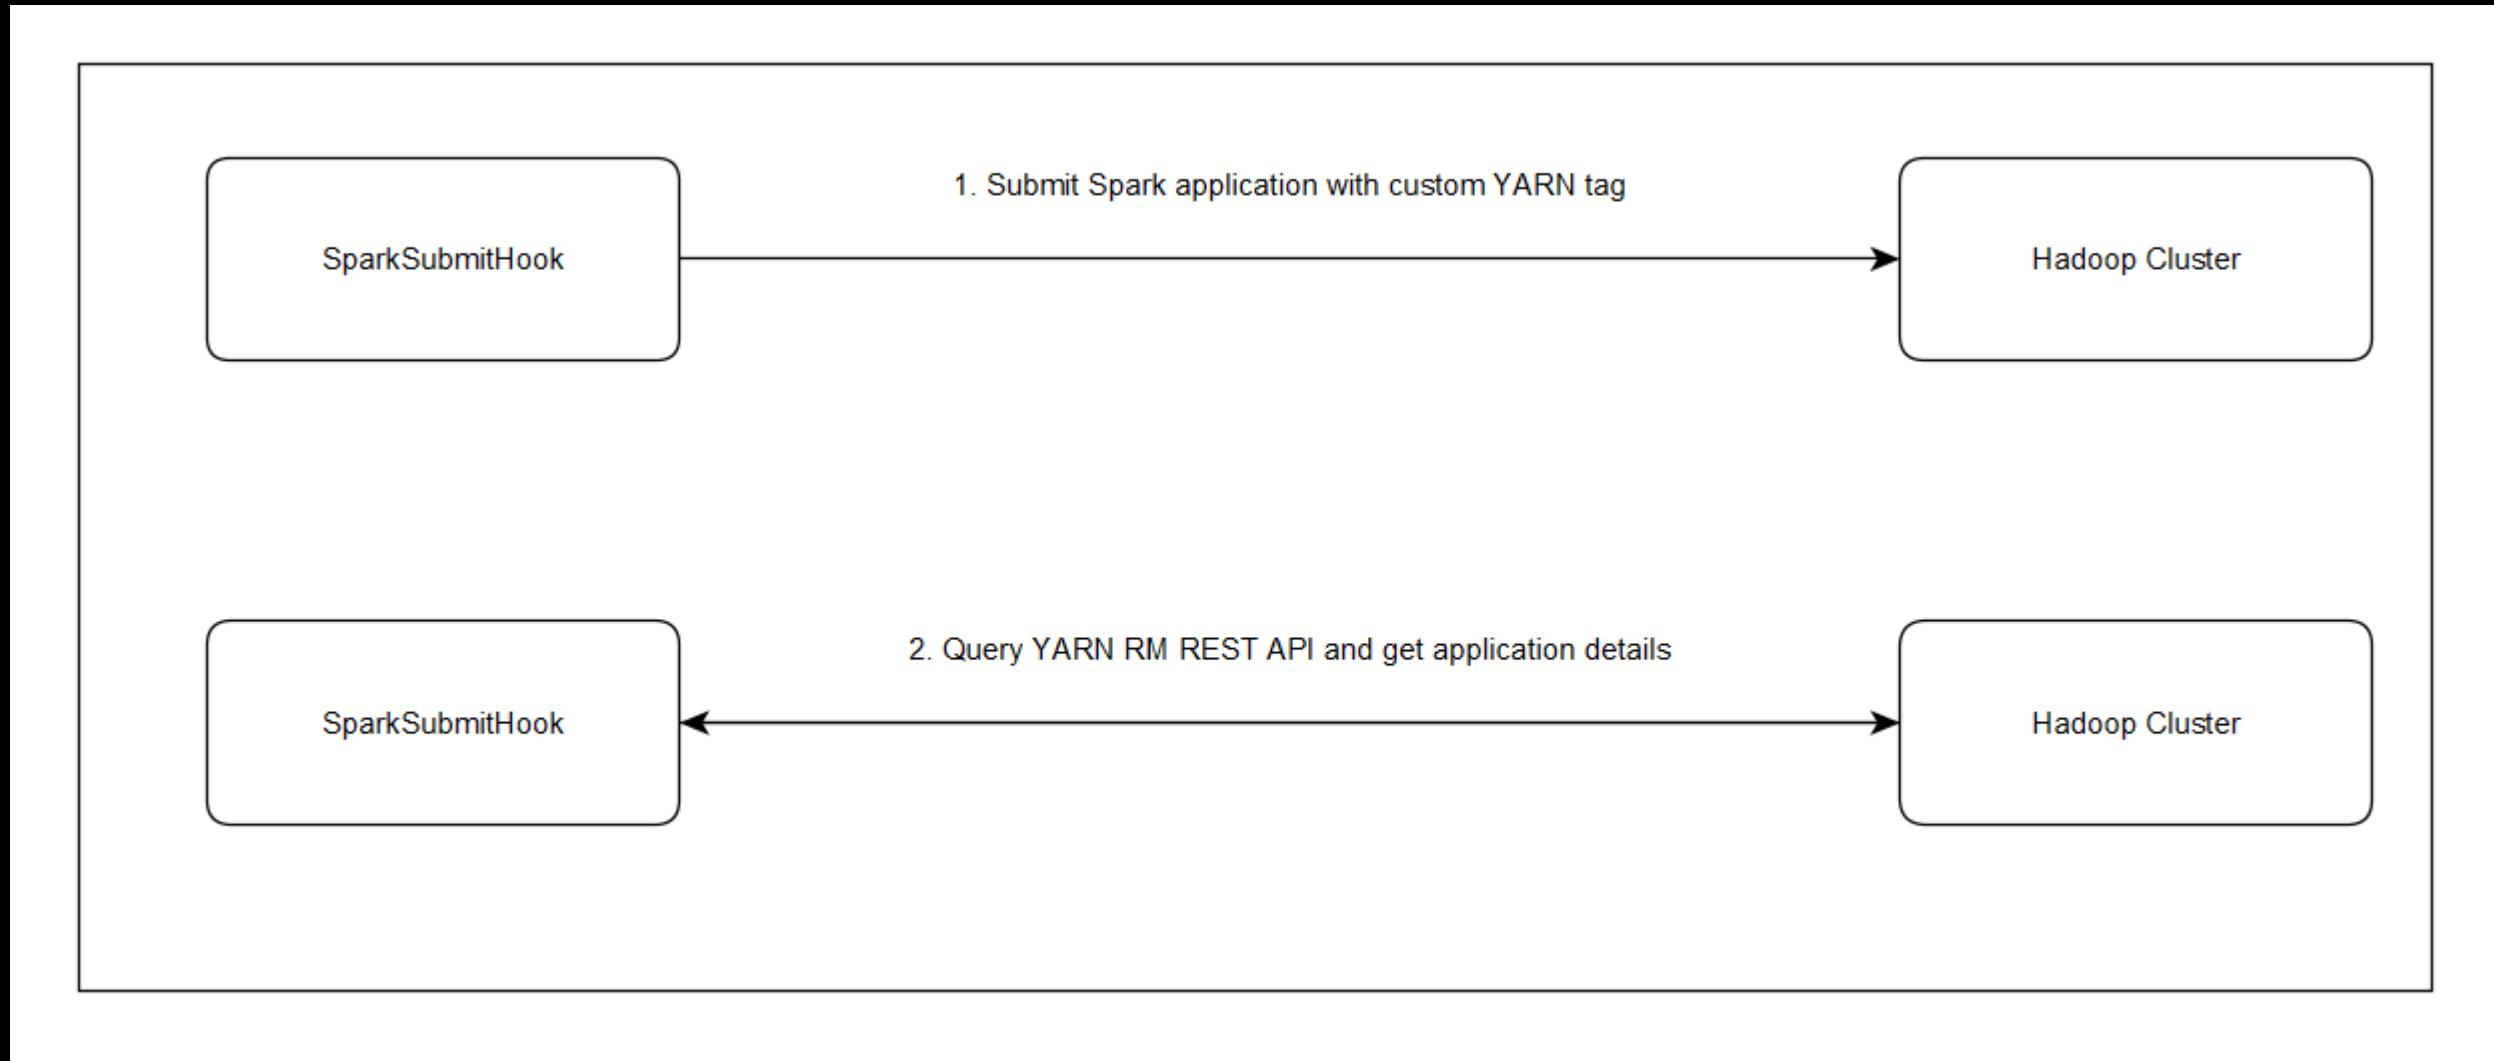
Operators / Hooks Customizations (2/3)

Enrich Airflow logs by adding YARN application details for Spark jobs

- Correlation between Spark application (triggered by Airflow) and submitted YARN job is challenging to discover when using YARN cluster mode
- Out of the box: YARN Application ID logged in the task logs only when a YARN job fails
- Extension: Extract and log the following for all Spark tasks:
 - 1) YARN Application ID
 - 2) YARN Tracking URL
 - 3) Diagnostics (Failure root cause)

Operators / Hooks Customizations (3/3)

YARN application details visible in Airflow logs for Spark jobs



Custom Authentication Methods

LDAP Secured REST API

- REST API usage allowed only for dedicated AD role
- Complete integration with LDAP
- Only one role specified for REST API – no separation between endpoints so far

A woman with long dark hair is seated in the driver's seat of a car. She is wearing a light-colored button-down shirt and jeans. She is looking down at a tablet device she is holding in her hands. The car's interior is visible, including the steering wheel, dashboard, and center console. The background shows a blurred view of the outside world through the car windows.

July 16, 2020

Airflow in Production

Production-ready Airflow instances

Robotic Drive supports by default 3 main instances used in the platform:

- **Development** (used mostly to test new Airflow deployments / features)
- **Staging** (testing new DAGs)
- **Production** (for full production usage)

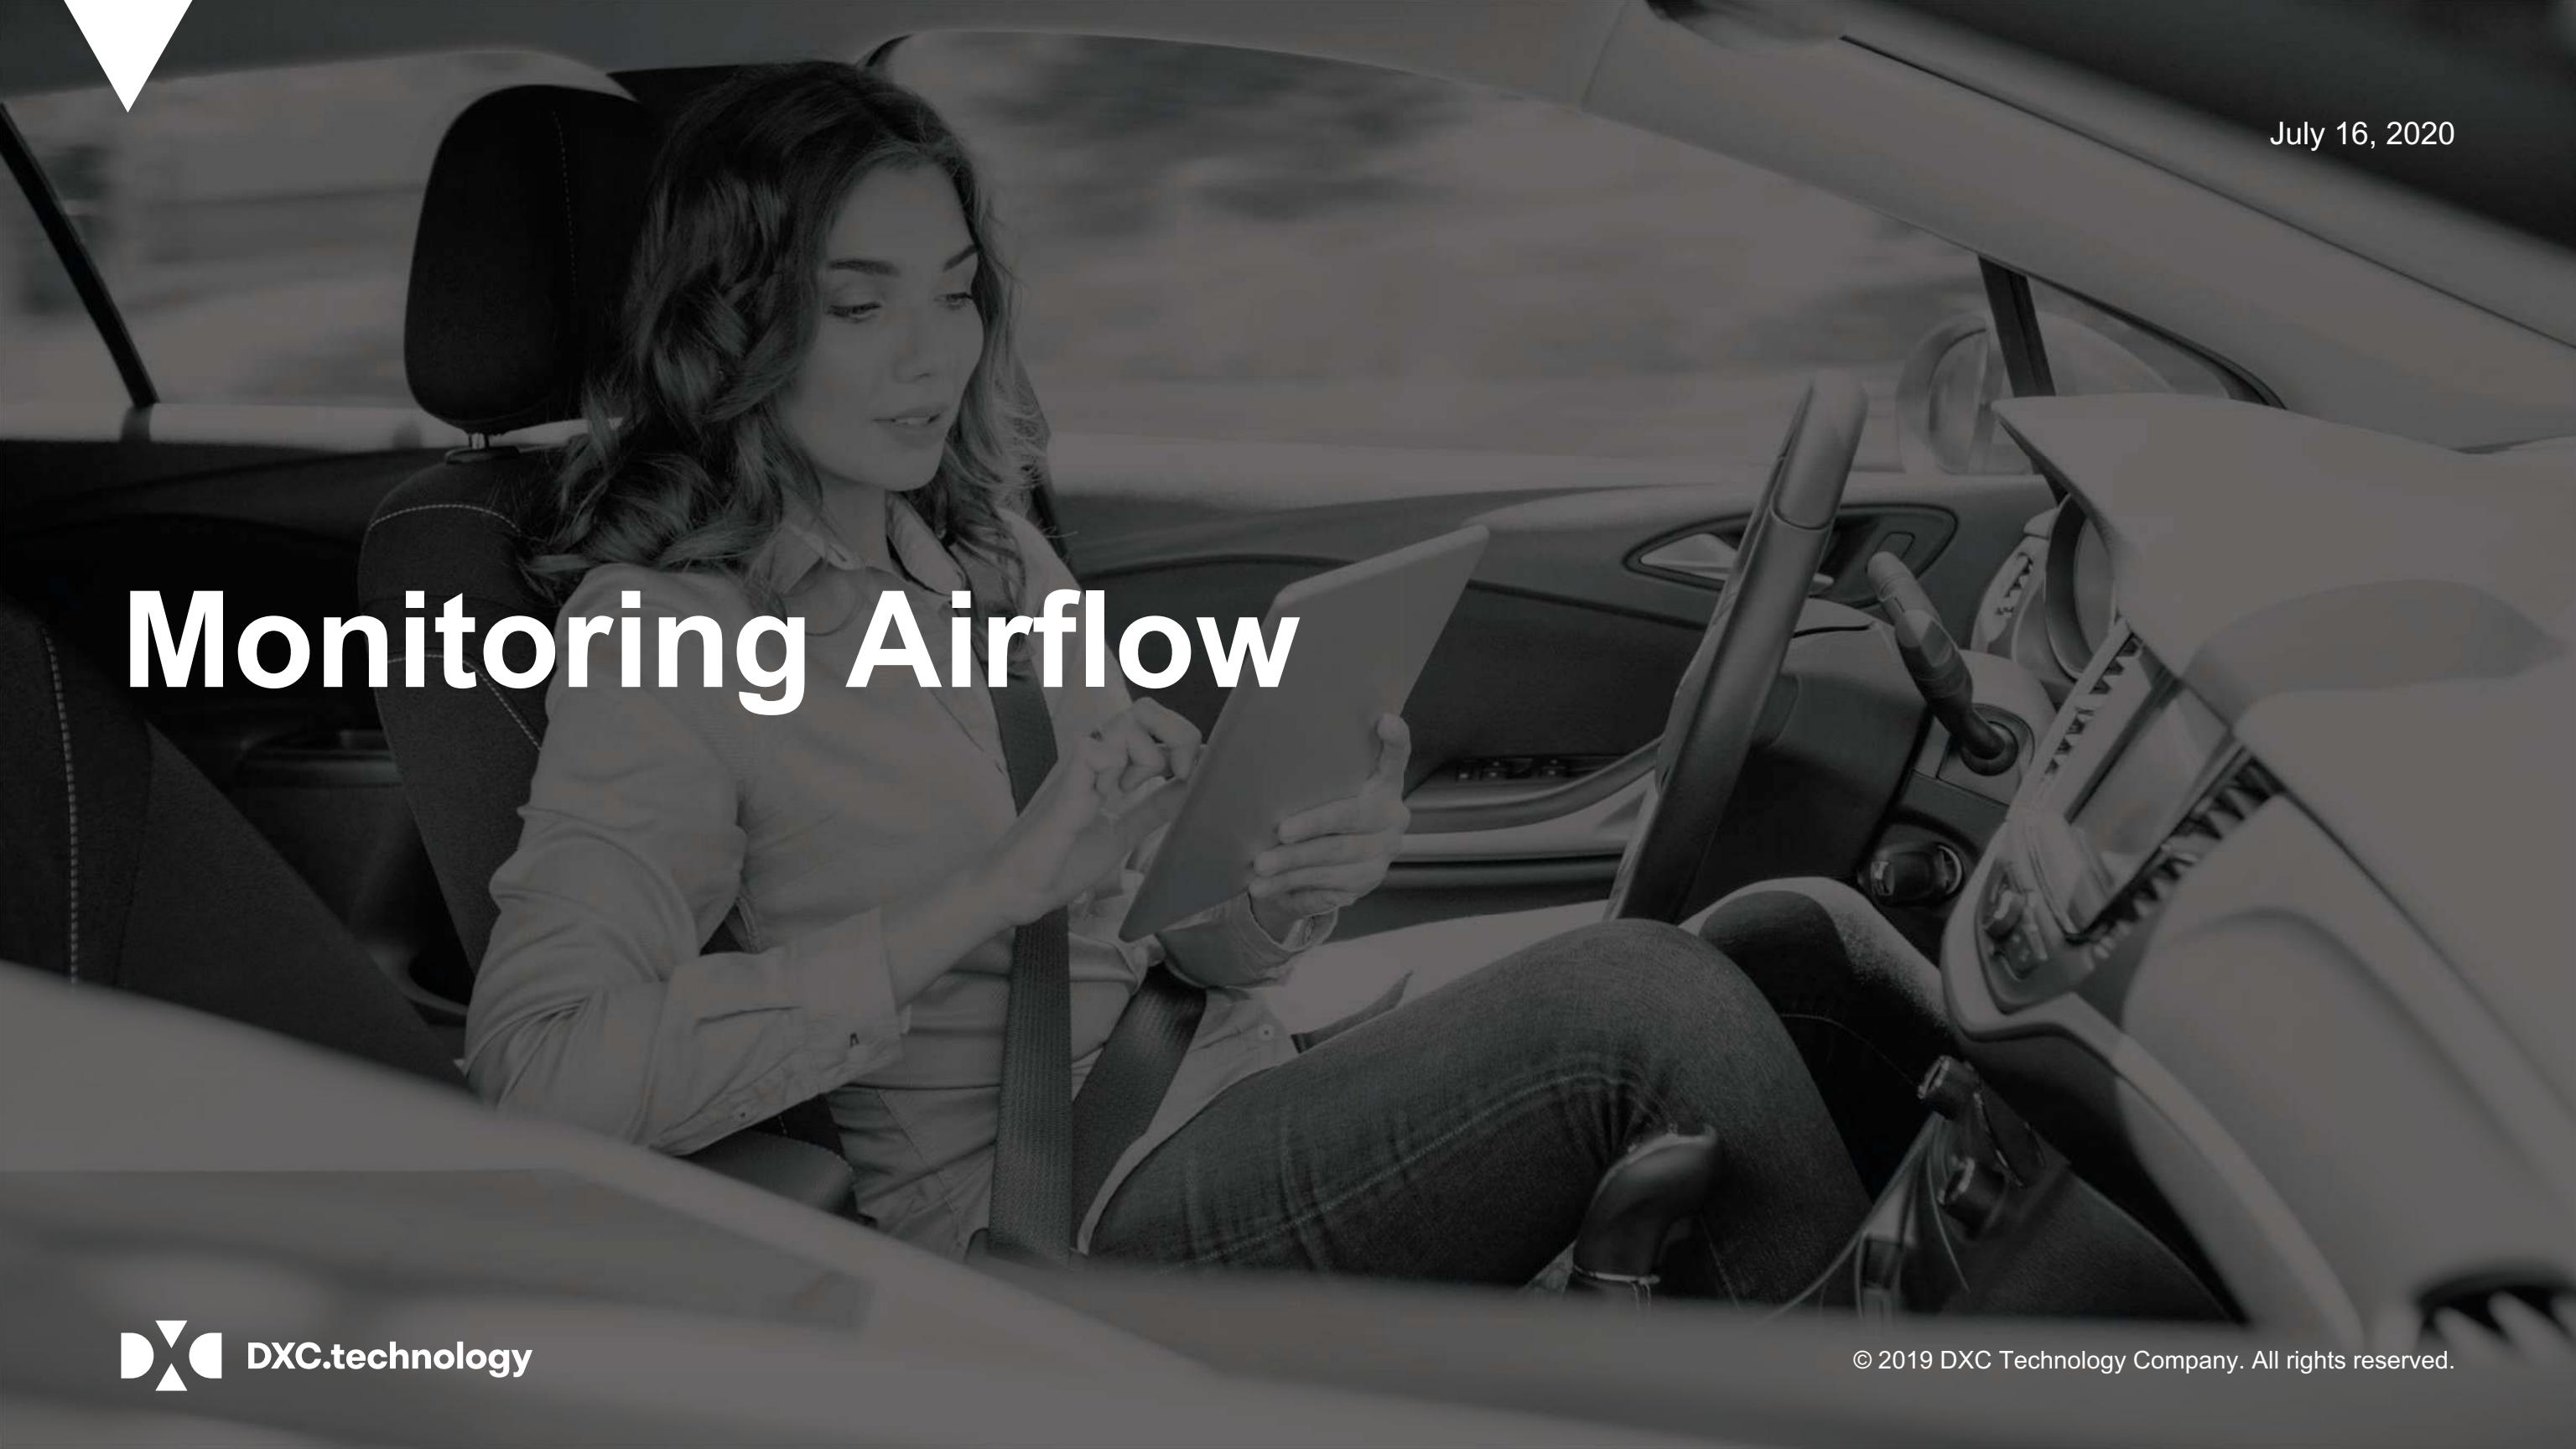
Current stable setup is created using Airflow 1.10.10 and Celery Executor

User-based Airflow instances

- Deployment automated via **Helm Charts**
- Airflow created as a Kubernetes project

These 2 points makes it possible to fully parametrize Airflow deployment and create many Airflow instances on-demand. In the Robotic Drive Platform each user is able to create their own, separate Airflow instance.

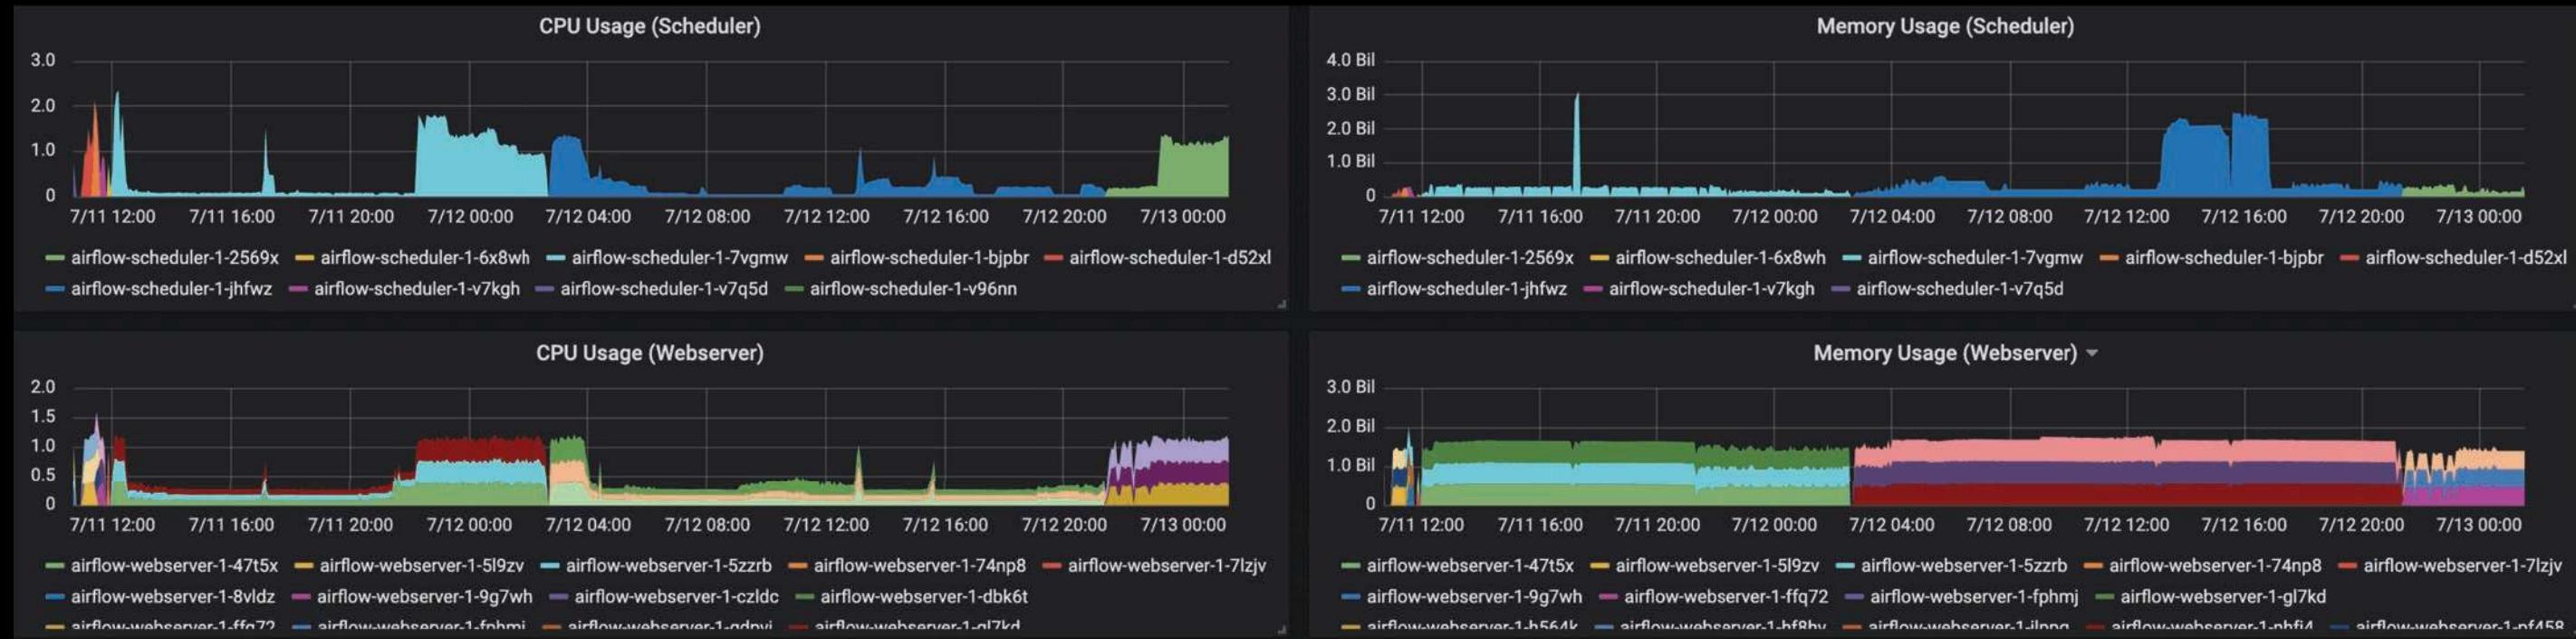
It helps to eliminate problems related to testing new DAGs, Operators and other features and changes that might impact other users, especially when multiple developers are working on the same component.



July 16, 2020

Monitoring Airflow

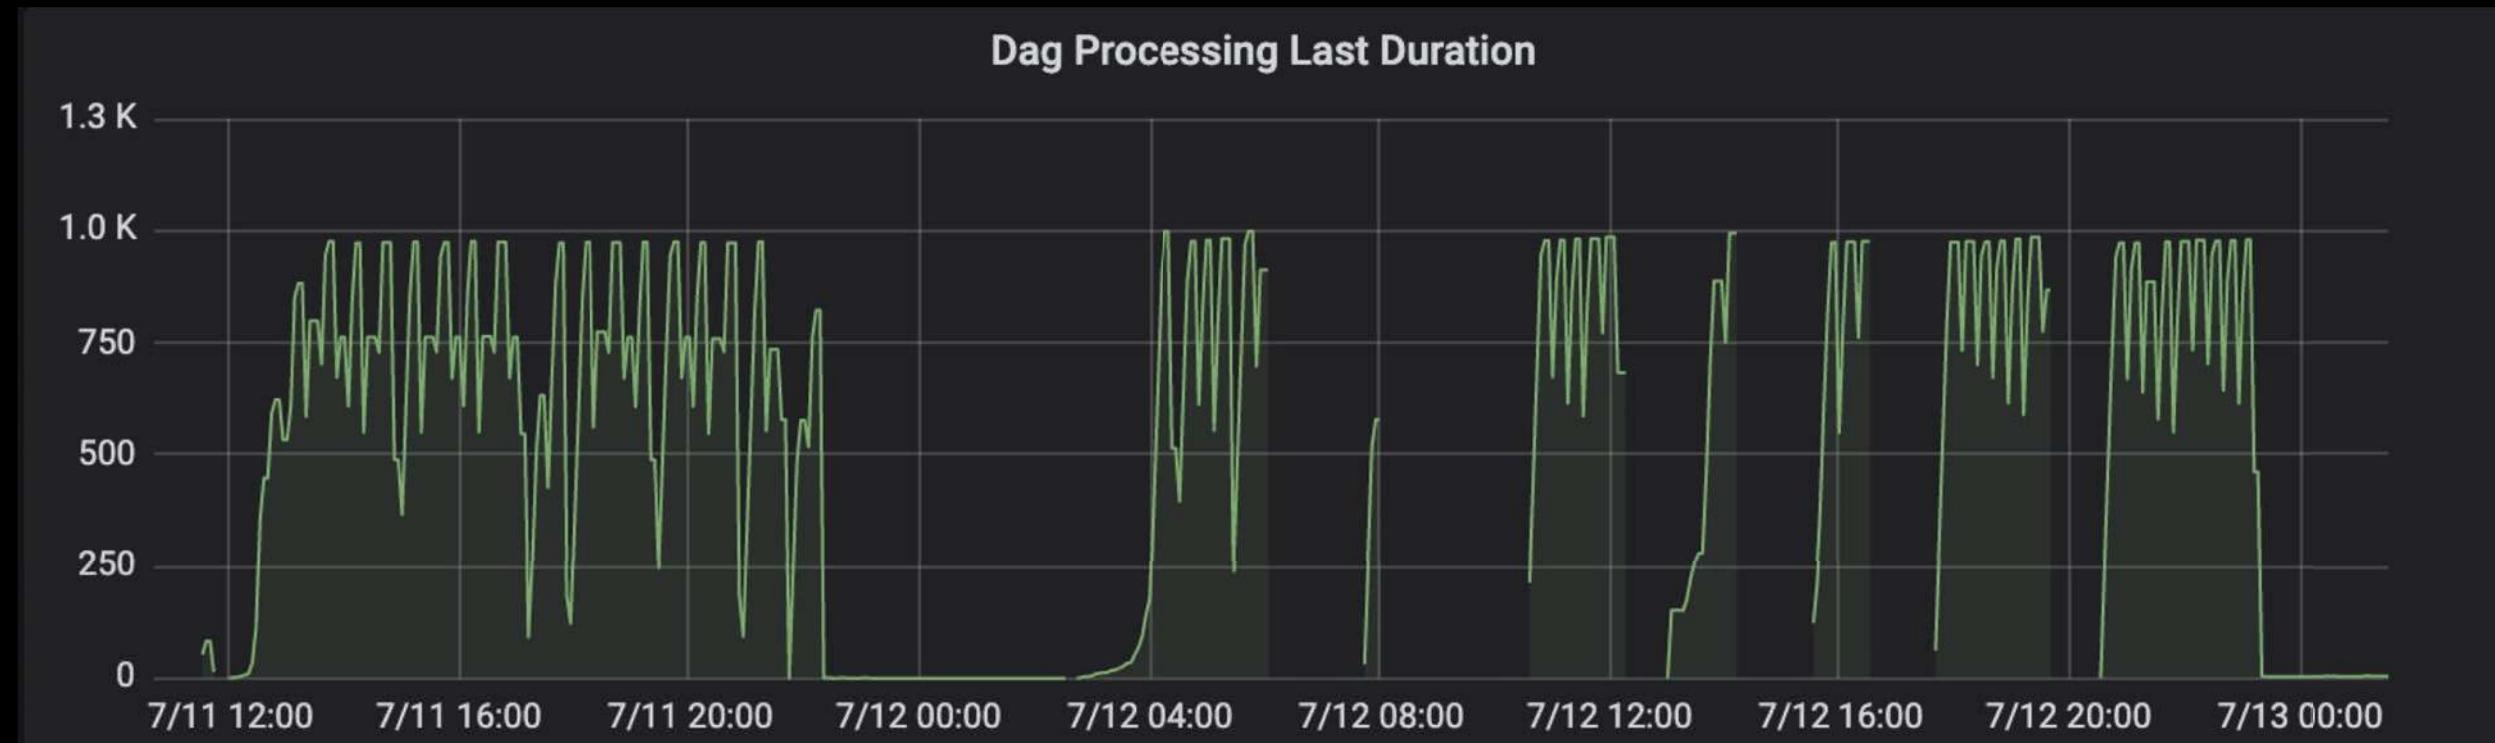
Monitoring Airflow

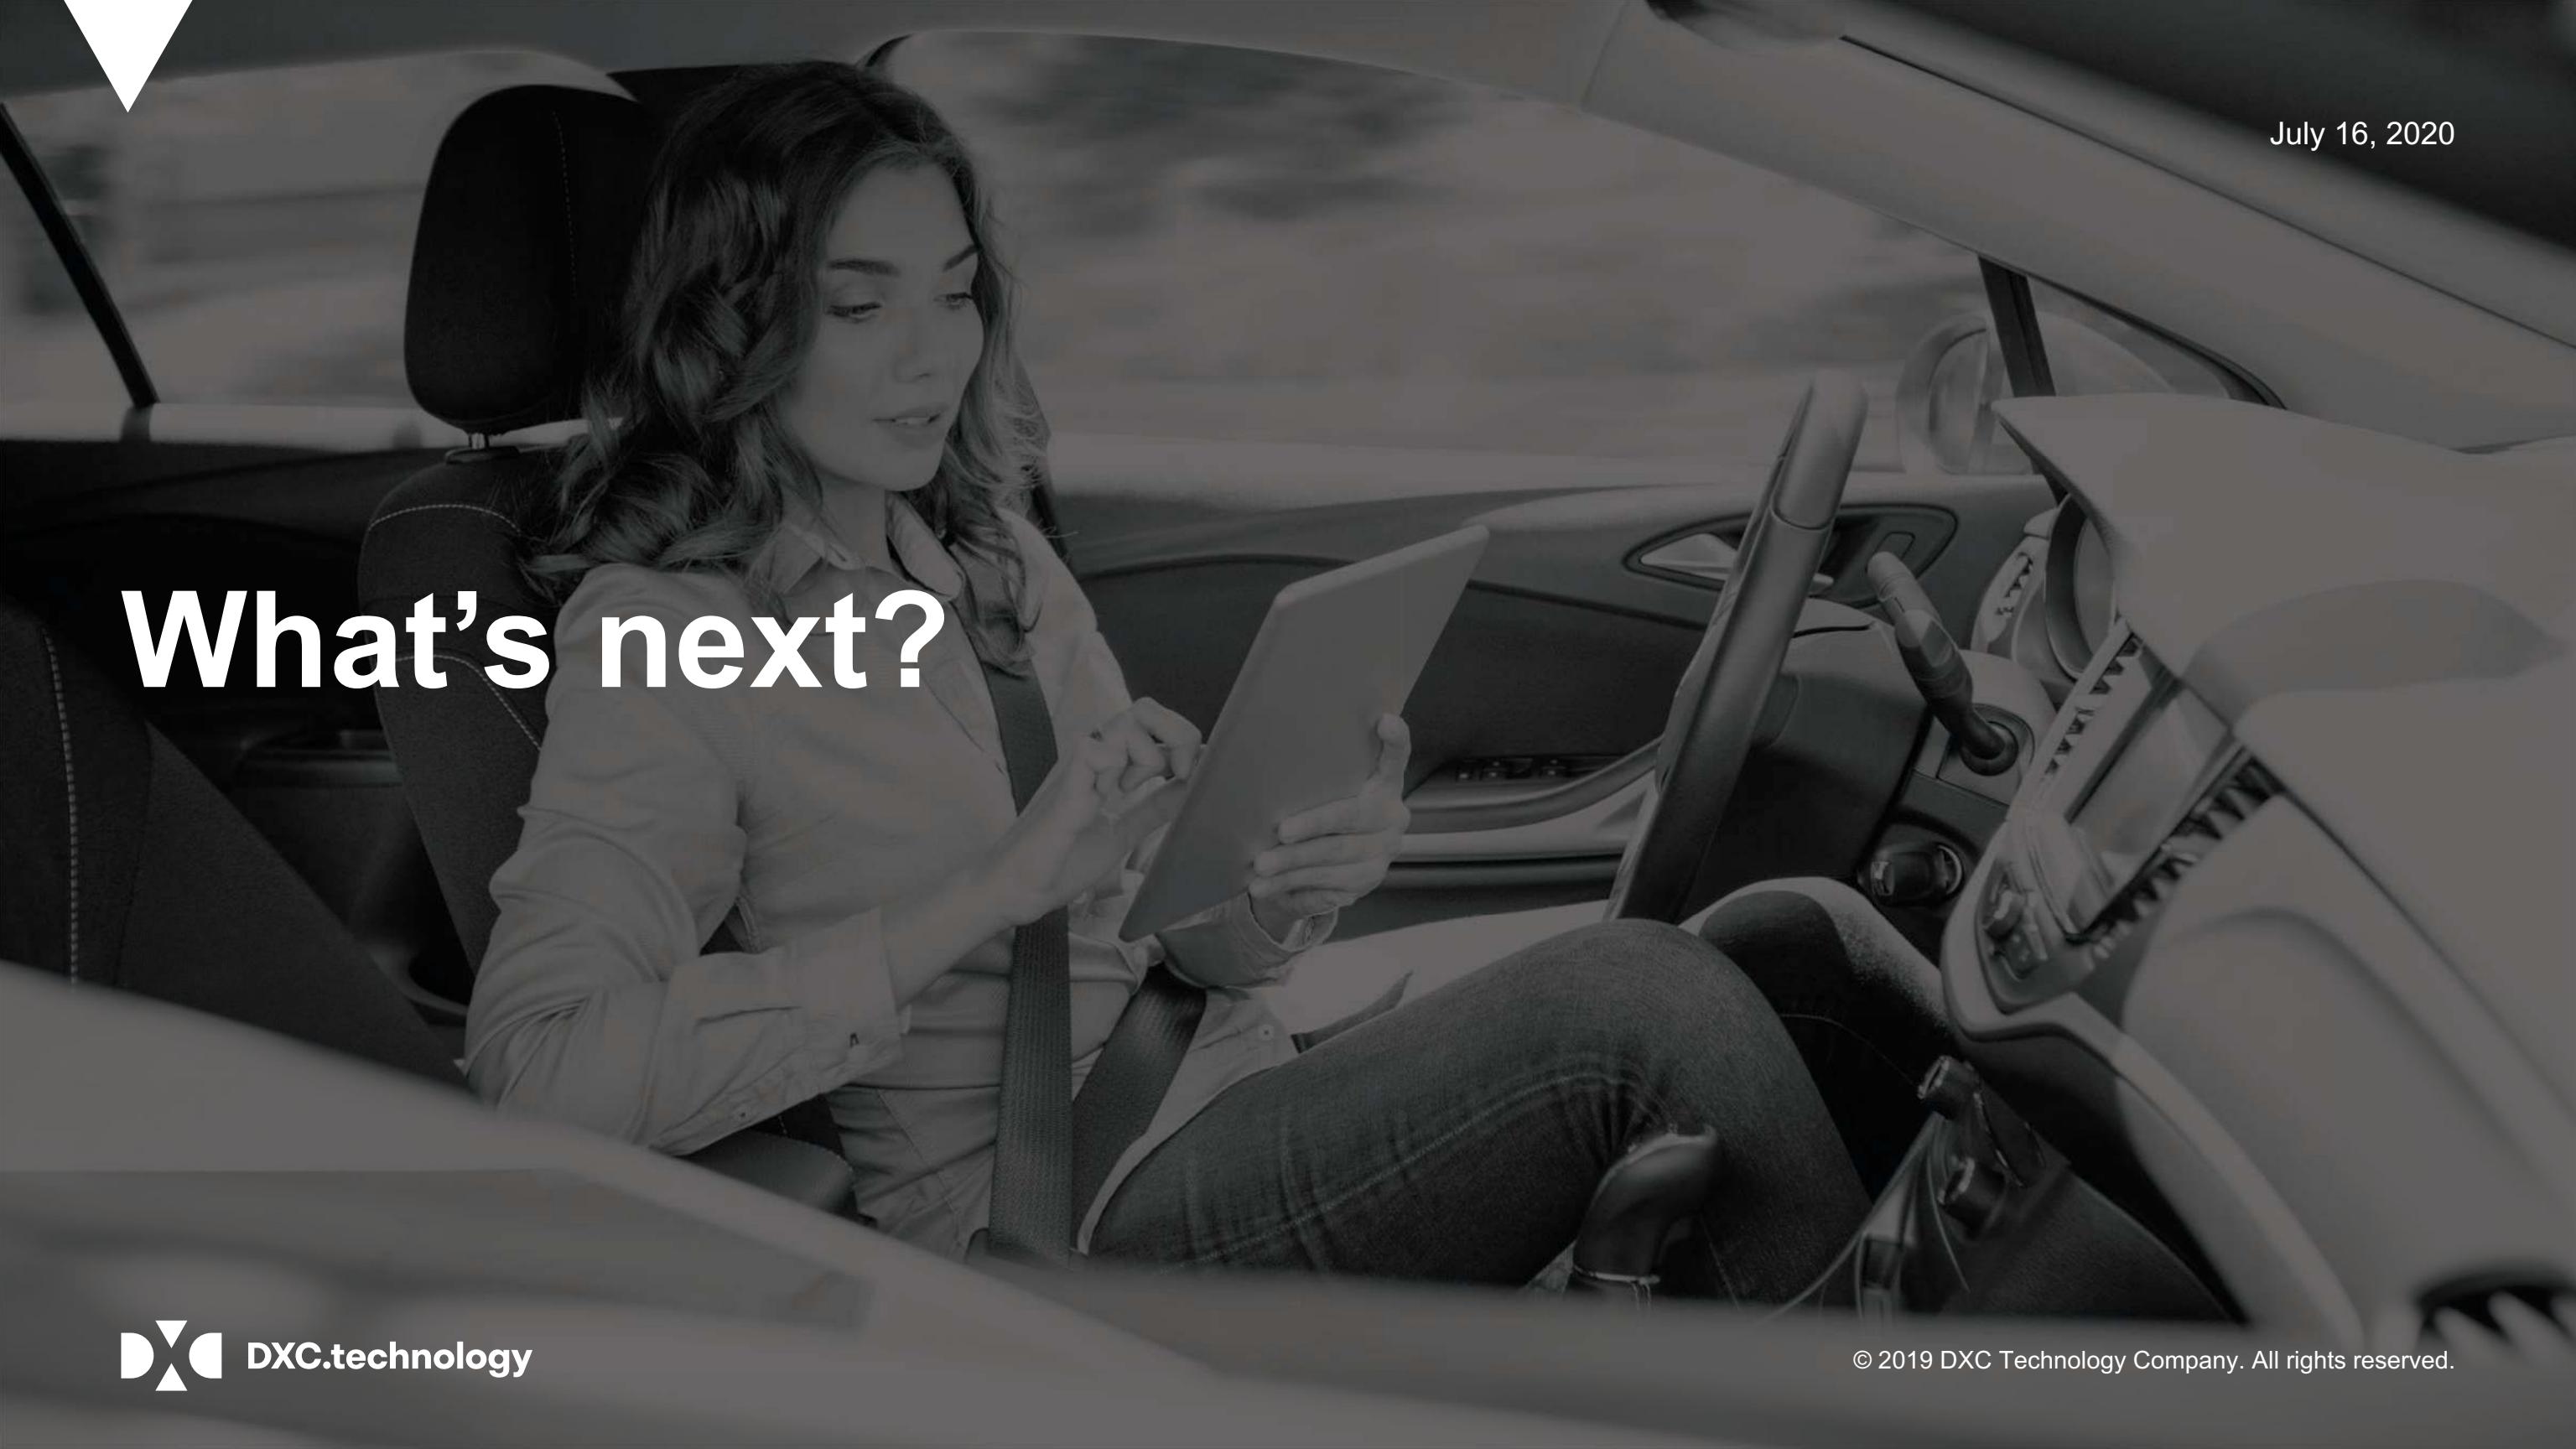


Monitoring Airflow



Monitoring Airflow



A woman with long dark hair is seated in the driver's seat of a modern car. She is wearing a light-colored button-down shirt and jeans. She is looking down at a white tablet device she is holding in her hands. The car's interior is visible, including the steering wheel, dashboard, and center console. The background outside the car shows a blurred landscape, suggesting motion.

July 16, 2020

What's next?

Looking forward to...

- **Airflow 2.0: HA Scheduler + Performance Optimizations**
- **Advanced Authentication + Authorization**
- **Extend and stabilize monitoring metrics**
- **Stable API vs Experimental API**

A woman with long dark hair, wearing a light-colored button-down shirt, is seated in the driver's seat of a modern car. She is looking down at a white tablet device held in her hands. The car's interior is visible, including the steering wheel, dashboard, and center console. The background shows a blurred view of the outside world through the car windows.

July 16, 2020

Q&A

Amr Noureldin [**<amr.noureldin@dxc.com>**](mailto:<amr.noureldin@dxc.com>)
Michal Dura [**<mdura2@dxc.com>**](mailto:<mdura2@dxc.com>)