



# From Cron to Apache Airflow

**A Startup Story**

Adam Boscarino, 2020-07-13



# Who am I?

- Data Engineer at Devoted Health
- Previously worked at DigitalOcean, Fitbit, Carbonite
- Airflow user for ~4 years
- GitHub & Twitter: ajbosco





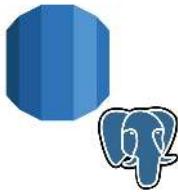
# What is Devoted Health?

- A next generation Medicare Advantage health insurer in the United States
- Founded in 2018, first enrolled members in 2019
- Offers a Clinical Services solution (Devoted Medical Group)
- Built on homegrown Devoted Tech Platform

**“ TO DRAMATICALLY IMPROVE HEALTHCARE  
FOR SENIORS IN THE UNITED STATES -- CARING  
FOR EVERYONE LIKE THEY ARE MEMBERS OF  
OUR OWN FAMILY ”**

# Devoted Health Data Platform, January 2019

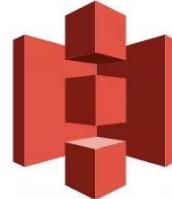
Source Data



Workflows



Storage/Data Lake



Data Warehouse



Reporting/BI



# Devoted Health Data Platform - Successes

- It did its job
- Successfully launched new health plan
- Supported key business operations and workstreams
- Powered all internal reporting





# Devoted Health Data Platform - Problems

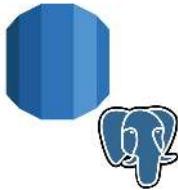
- No task dependencies
- Undetected system downtime
- Onboarding new developers
- Environment parity
- Unsure of data quality





# Devoted Health Data Platform, May 2019

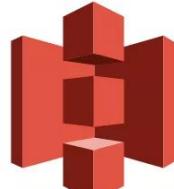
## Source Data



## Workflows



## Storage/Data Lake



Amazon S3

## Data Warehouse



## Reporting/BI



**NEW!**

**NEW!**



# Devoted Health Data Platform - Problems

- No task dependencies
- Undetected system downtime
- Onboarding new developers
- Environment parity
- Unsure of data quality

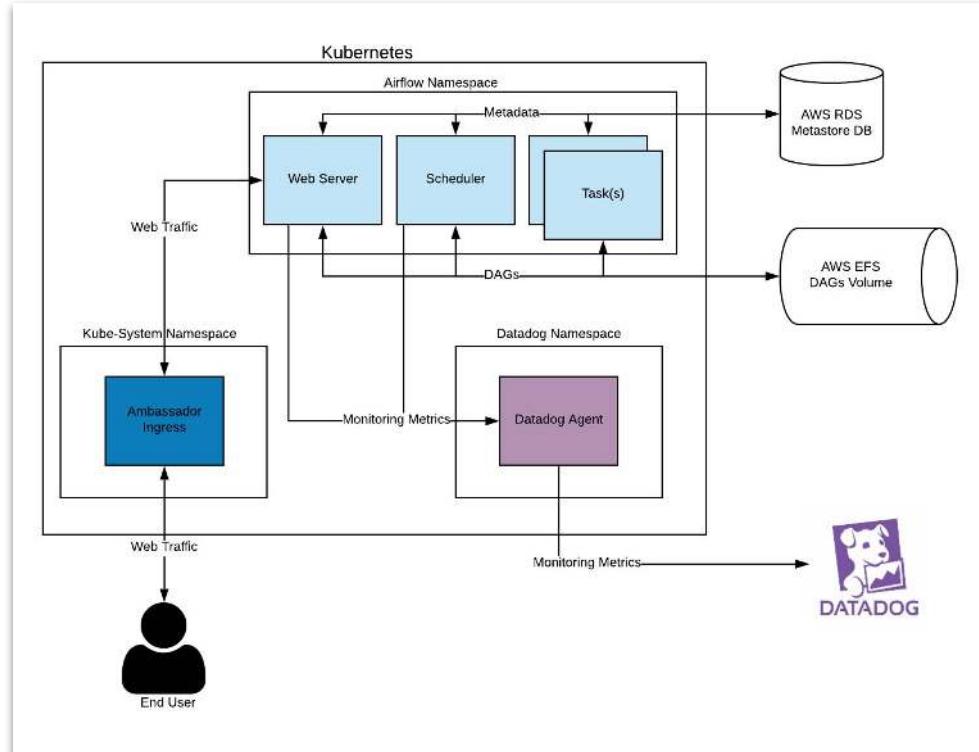


# Solution: Deploying Airflow



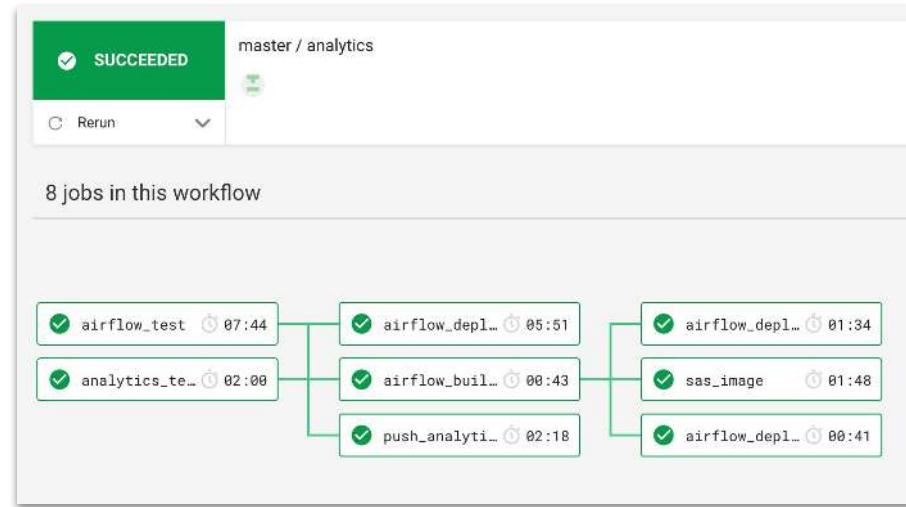
# Apache Airflow Deployment

- Kubernetes
  - Orchestrates Airflow services
  - Kubernetes Executor
- Helm
  - Kubernetes Package Manager
  - Describes Kubernetes resources
  - Official Helm Chart
- Terraform
  - Infrastructure as Code
  - Used to deploy Helm chart to Kubernetes clusters



# DAG Deployment

- DAGs are stored on AWS EFS
  - Mounted to each Airflow pod in Kubernetes
- DAGs are pushed from GitHub to AWS EFS via CircleCI
  - No manual intervention
  - Many deployments every single day





# Devoted Health Data Platform - Problems

- No task dependencies
- **Undetected system downtime**
- Onboarding new developers
- Environment parity
- Unsure of data quality



# Solution: Monitoring Airflow



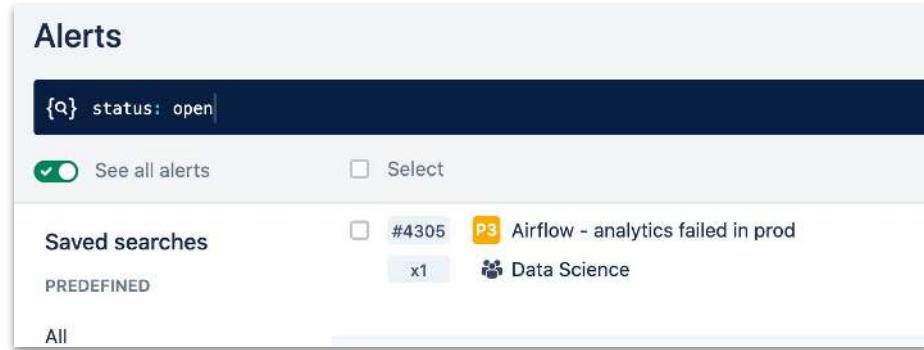
# Monitoring Airflow

- Kubernetes Liveness & Health Checks
  - Monitor /health endpoint
  - Monitor Scheduler health
  - Restart services if in bad state
- Datadog Monitors
  - Alert on-call engineer via OpsGenie and Slack
  - Airflow is not running
  - No DAGs have completed in last 2 hours
  - CPU/Memory Usage has spiked



# Monitoring DAGs

- OpsGenie alerts sent to DAG Owner (and Slack)
- DAG owners are responsible for resolving non-infrastructure failures
- Alerting is “built-in” to every DAG



The screenshot shows the Airflow UI interface for monitoring DAGs. At the top, there's a search bar with the query '{q} status: open'. Below it are two buttons: 'See all alerts' (which is checked) and 'Select'. Underneath, there are sections for 'Saved searches' and 'PREDEFINED'. A single alert is listed: '#4305 P3 Airflow - analytics failed in prod' associated with 'Data Science' (x1). There's also a 'All' link at the bottom.

# Devoted Health Data Platform - Problems

- No dependency management
- Undetected system downtime
- **Onboarding new developers**
- Environment parity
- Unsure of data quality

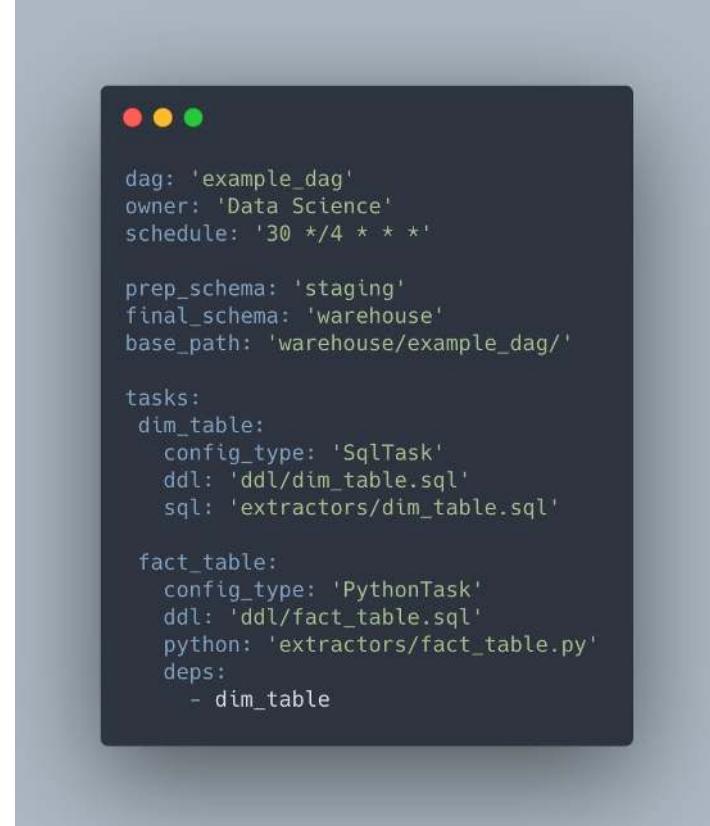


# Solution: YAML DAG Builder



# YAML DAG Builder

- Internal library to simplify and standardize DAG development.
- Abstraction on top of Airflow.
- Developers only write a DDL query and transformation in SQL or Python.
- No prior Airflow knowledge required.
- Similar structure in all DAGs makes switching teams less painful and simplifies debugging DAGs.
- Data Engineer team can bolt on additional features (alerting, monitoring, testing, etc.)



```
dag: 'example_dag'
owner: 'Data Science'
schedule: '30 */4 * * *'

prep_schema: 'staging'
final_schema: 'warehouse'
base_path: 'warehouse/example_dag/'

tasks:
  dim_table:
    config_type: 'SqlTask'
    ddl: 'ddl/dim_table.sql'
    sql: 'extractors/dim_table.sql'

  fact_table:
    config_type: 'PythonTask'
    ddl: 'ddl/fact_table.sql'
    python: 'extractors/fact_table.py'
    deps:
      - dim_table
```



# Devoted Health Data Platform - Problems

- No task dependencies
- Undetected system downtime
- Onboarding new developers
- **Environment parity**
- Unsure of data quality



# Solution: devflow



- Internal tool that wraps kubectl, Helm, and Terraform.
- Every developer gets their own Airflow deployment on Kubernetes.
- We develop on the same stack that we run production.
- Developers do not need to know anything about the infrastructure being used.



The terminal window displays the devflow command-line interface. At the top, it shows a large, stylized "DEVFLOW" logo. Below it, a message says "Run 'source with-role analyst@dev' prior to using devflow or it will not work!!!". The usage information is as follows:

```
usage:  
-h|help      display usage  
start        start Airflow instance on dev k8s cluster  
restart      restart Airflow instance on dev k8s cluster  
stop         stop Airflow instance on dev k8s cluster  
stop-all    stop all Airflow instances on dev k8s cluster that are older than 24 hours;  
            pass a different number of hours to limit to older releases, e.g. stop-all 2  
sync         sync DAGs and Plugins directories to Airflow  
status       check status of Airflow instance  
logs         view logs from Airflow web server  
logs-scheduler  view logs from Airflow scheduler  
shell        open a bash shell in the Airflow instance  
clone schema/table clone tables or schemas in * from Prod to your Dev database  
deploy udfs   deploy latest UDFs to your Dev database *  
analytics-push builds and pushes 'data-pipeline-stg' image to ECR  
push-dev-image builds and pushes 'airflow-dev' image to ECR
```



# Devoted Health Data Platform - Problems

- No task dependencies
- Undetected system downtime
- Onboarding new developers
- Environment parity
- **Unsure of data quality**

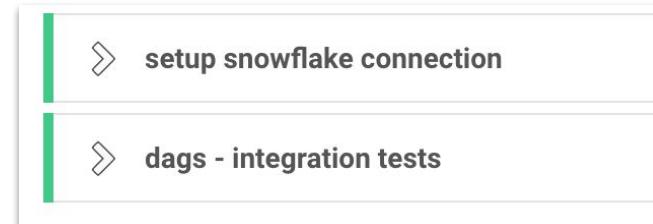
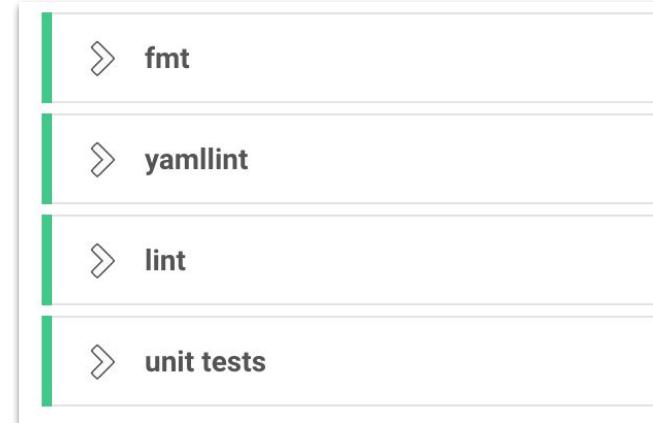


# **Solution: Testing & Validation**



# Testing DAGs

- Unit Tests
  - Used on Python transformations and core library code
- Integration Tests
  - Used for SQL tasks
  - Internal framework built on pytest
  - Executed against Snowflake using a test database
  - Mock tables are created and populated





# Data Validation Framework

- Data validation is executed at DAG run-time
- DAGs are stopped if validation fails to prevent reporting on bad data
- Started with Check Operator
- Added internal Operators
  - Runs multiple checks with one task
  - Save invalid records to table
  - Send check values to Datadog
- Checks range from primary key validation to custom business logic



```
table_a_pk_validation:
    config_type: ValidateTask
    validation_type: 'uniqueness'
    validation_table: table_a
    validation_columns:
        - id
    deps:
        - table_a_populated

table_a_count_validation:
    config_type: ValidateTask
    validation_type: 'custom'
    validation: 'validations/table_a_count_validation.sql'
    operation: '='
    pass_value: 0
    deps:
        - table_a_populated

table_c_quality_checks:
    config_type: QualityCheckTask
    description: 'Runs all data quality checks for table C.'
    quality_checks: 'validations/table_c_validation.py'
    deps:
        - table_c_populated
```

# Mission Accomplished!



# Current Issues & Future Work

- Improve SQL testing!
  - Explore tools like dbt and dataform
  - Remove need for end user to know pytest
- Improve DAG Builder
  - Make standard use cases easier
- SQL Linting/Formatting
  - Enforce best practices programmatically
- KEDA Autoscaler
  - Improve task spin-up speed



# Questions?

