

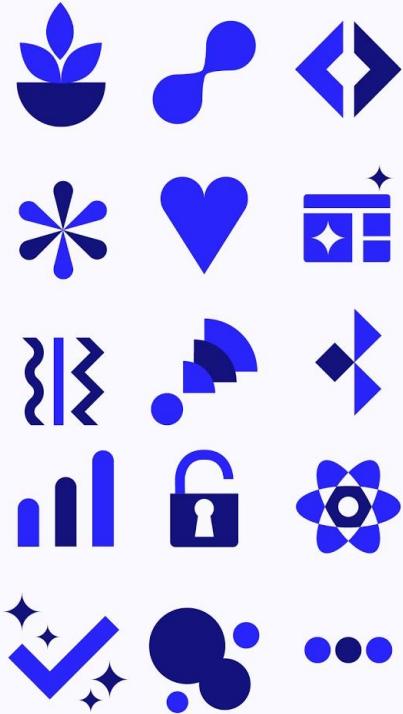
Polidea

Airflow Summit 2020 - 14.07.2020

Production Docker Image

for
**Apache
Airflow**





Polidea

Airflow Summit 2020 - 14.07.2020

Production Container Image



for
**Apache
Airflow**

Hi!



Jarek Potiuk

Apache Airflow:
PMC Member and Committer

Polidea:
Principal Software Engineer
(ex-CTO)

Airflow Summit:
Co-Organizer: Content (Lead)

@higrys



Intro



Intro

What questions will be answered?

- Context
 - What container images are and why there are important ?
- Status
 - How it looked like so far ?
 - How it is going to look like now ?
- Internals
 - What is in the image?
 - How we test the image?
- Usage
 - How to extend Airflow Image?
 - How to customize Airflow Image?
 - How you can use the Image?
- Future
 - What's next?



Intro

What this talk is NOT about?

- Basic container image knowledge
 - <https://docker-curriculum.com/>
- Details of CI container image of Airflow
 - <https://github.com/apache/airflow/blob/master/IMAGES.rst>
- Details of how Kubernetes Airflow integrate
 - “Airflow on Kubernetes” by Michael Hewitt
<https://www.crowdcast.io/e/airflowsummit/6>
- Details on deploying Airflow with the image



Intro

Who is the talk for?

- You want to deploy Airflow using container images
- You want to contribute to Airflow in Devops area
- You want to learn about best practices of using Airflow Containers
- You are a curious person that want to learn something new



Container Images

Context



Context

What is a container ?

- Standard unit of software.
 - OCI: <https://opencontainers.org/>
- Packages code and its dependencies
- Lightweight execution package of software
- Container images - binary packages



Container



Container image



Context



Container management CLI

Container ≠ Docker

- Docker is a command line tool
 - Building, Running, Sharing containers
- Docker Engine runs containers
- Alternatives: **rkt, containerd, runc, podman, lxc, ...**
- DockerHub.com is popular container registry
- Alternatives: GitHub, GCR, ECR, ACR



Container execution engine



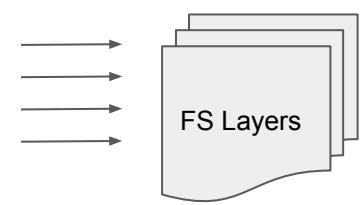
Container registry



Context: What is Container file

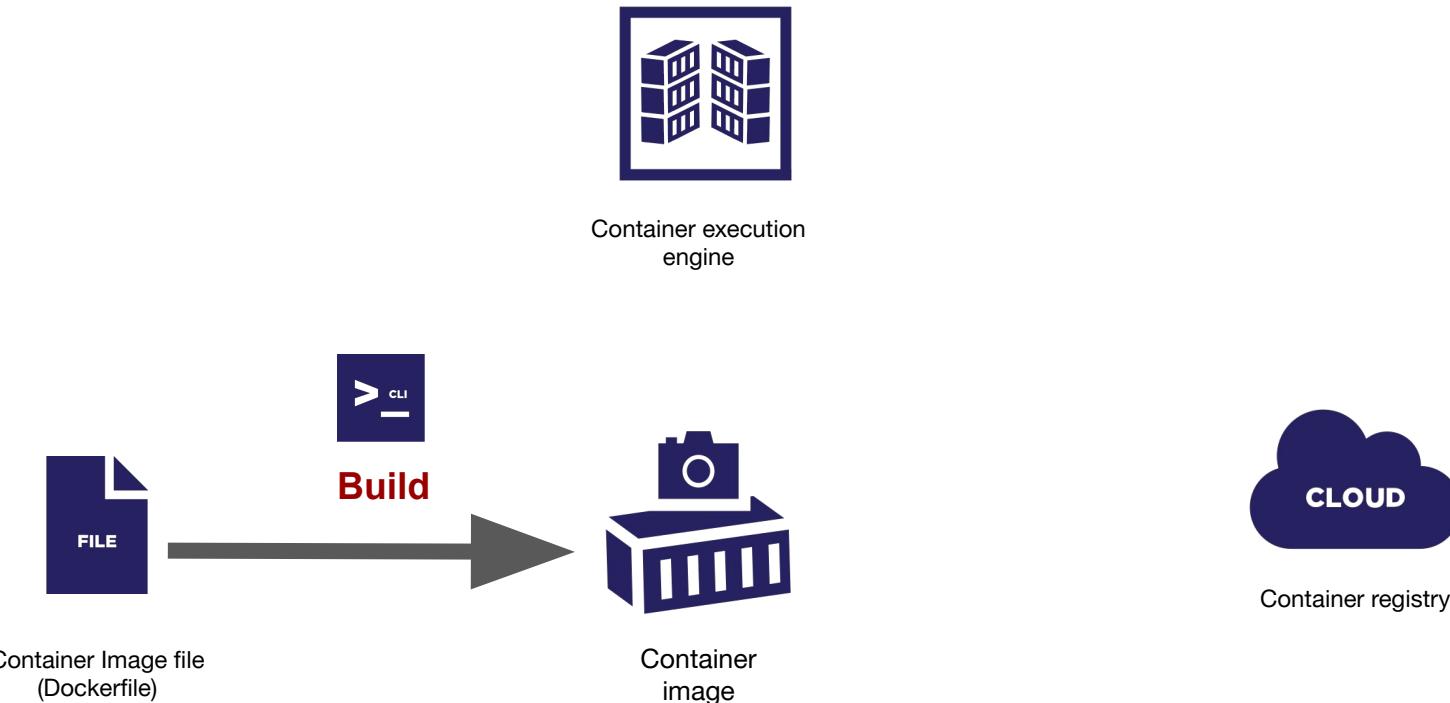
- Specify base image
- Run commands
- Copy files
- Set working directory
- Define entrypoint
- Define default command

```
FROM ubuntu:18.04
COPY . /app
RUN make /app && make install
WORKDIR /bin/project
ENTRYPOINT [ "/bin/project" ]
CMD [ "--help" ]
```





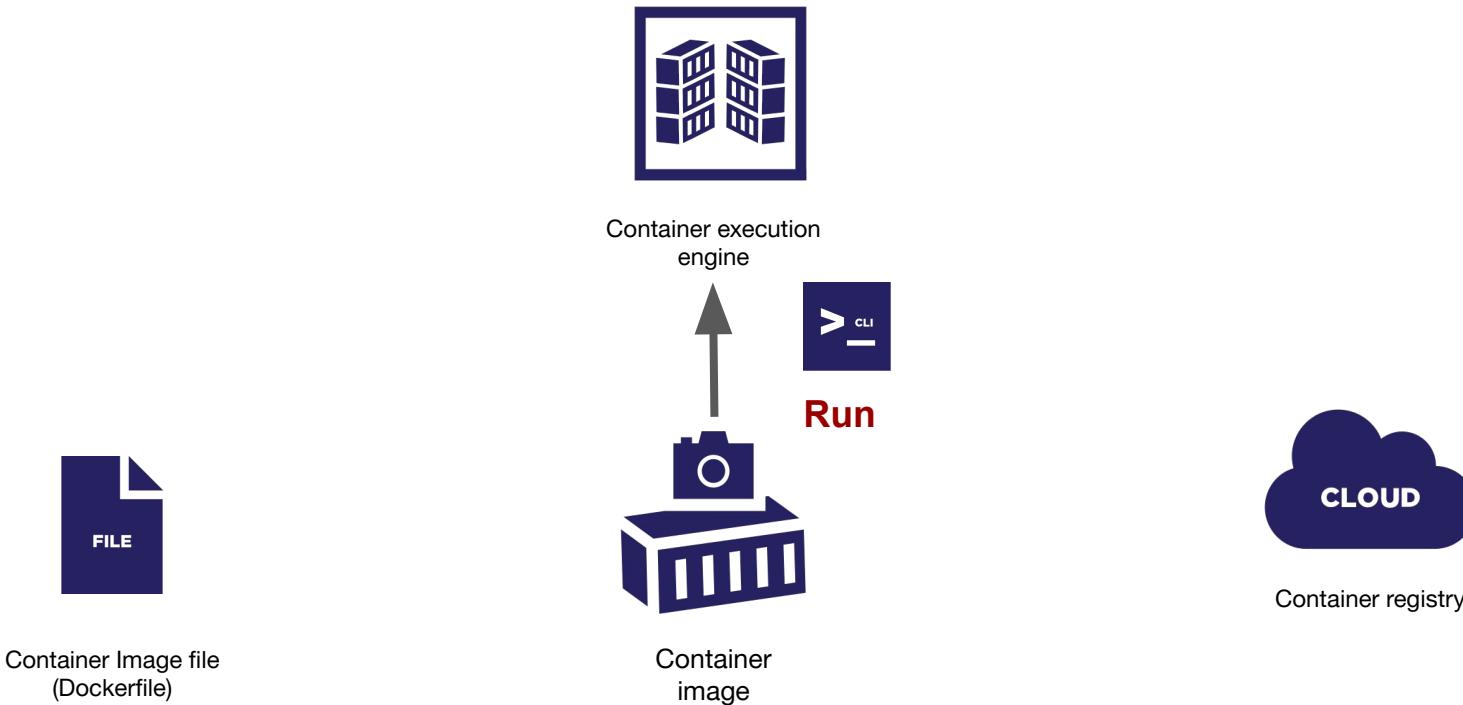
Context: Container Lifecycle: Build





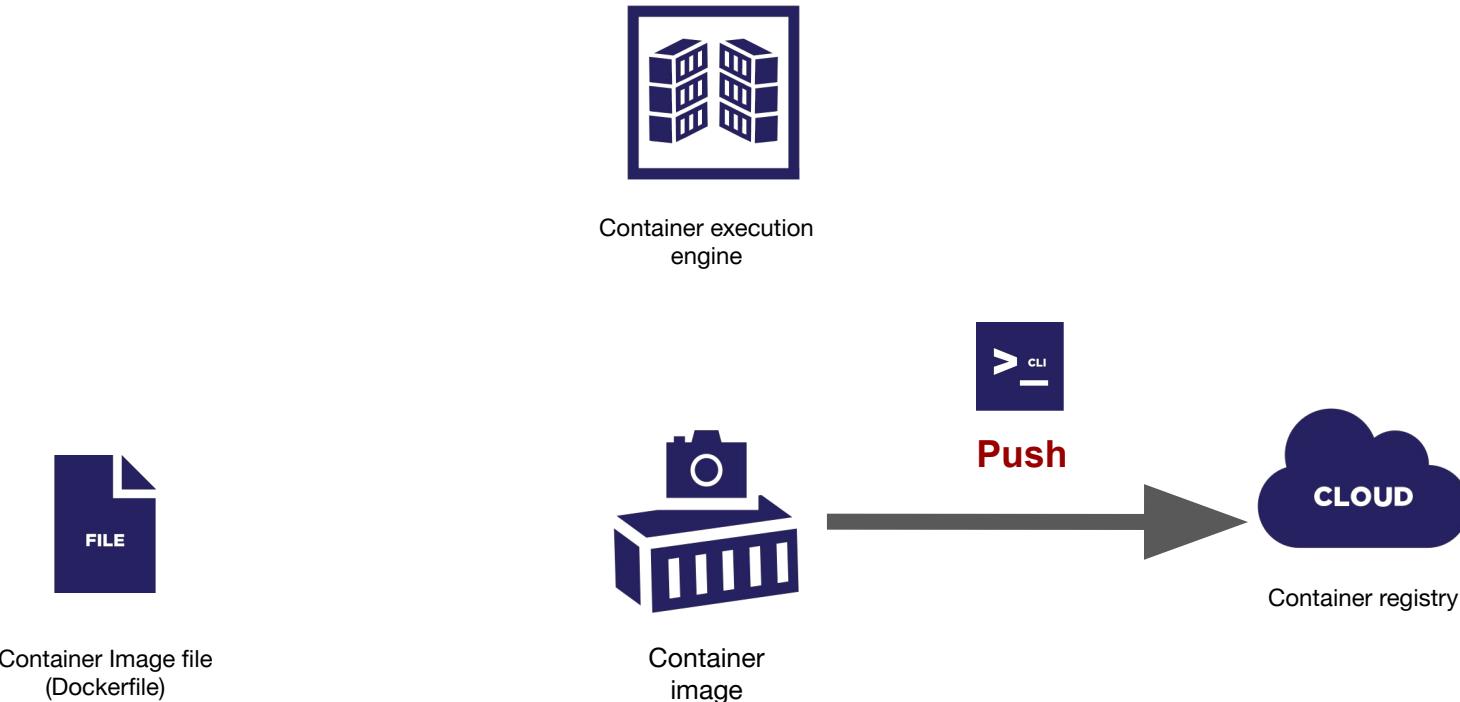
Context: Container Lifecycle: Run

—



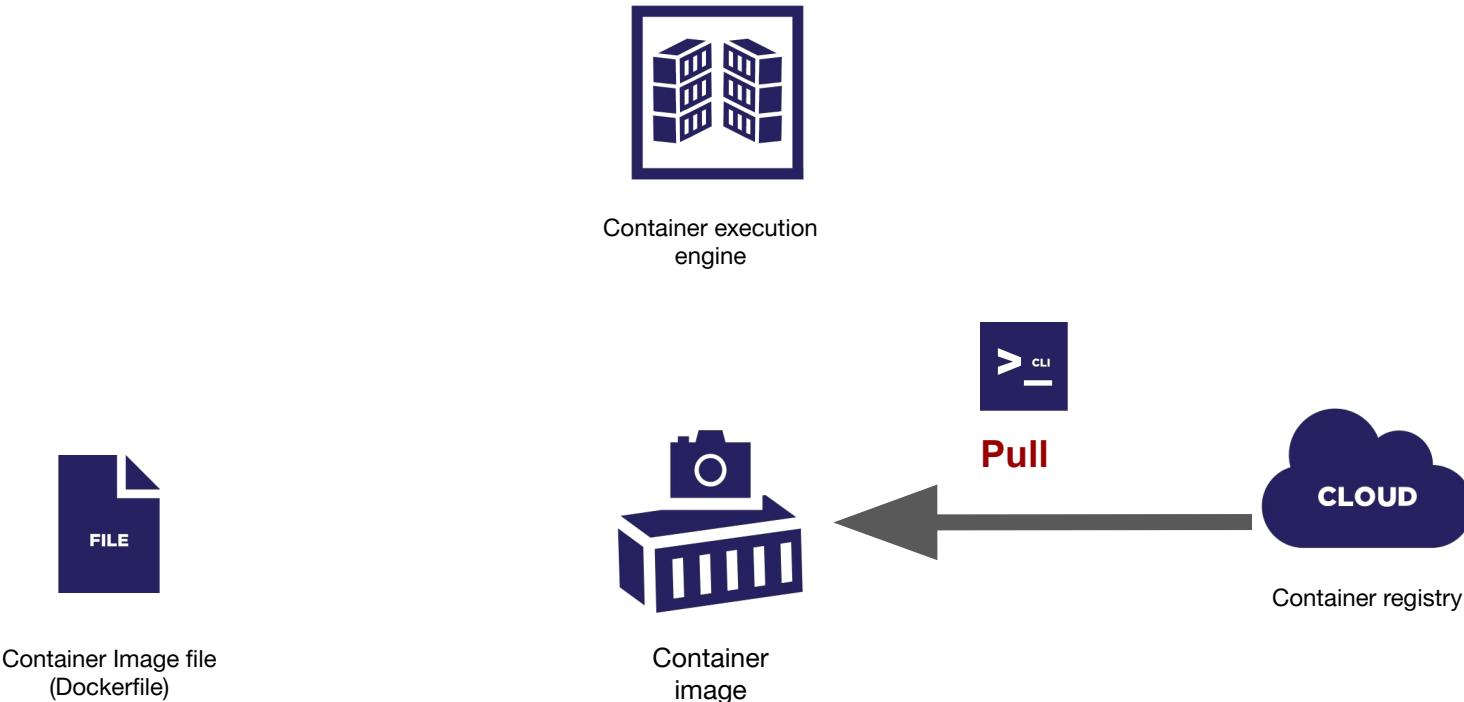


Context: Container Lifecycle: Push





Context: Container Lifecycle: Pull





Context

Why containers are important?

- Predictable, consistent development & test environment
- Predictable, consistent execution environment
- Lightweight but isolated: sandboxed view of the OS isolated from others
- Build once: run anywhere
- Kubernetes runs containers natively
- Bridge: “Development -> Operations”

```
        return None
```

```
def get(self, key):  
    if key in self.dict:  
        return self.dict[key]  
    else:  
        return None
```

```
def set(self, key, value):  
    if key in self.dict:  
        self.dict[key] = value  
    else:  
        self.dict[key] = value
```

```
def del(self, key):  
    if key in self.dict:  
        del self.dict[key]  
    else:  
        raise KeyError
```

```
def clear(self):  
    self.dict = {}
```

```
def keys(self):  
    return list(self.dict.keys)
```

```
def values(self):  
    return list(self.dict.values)
```

```
def items(self):  
    return list(self.dict.items)
```

```
def copy(self):  
    return dict(self.dict)
```

```
def update(self, dict):  
    for key, value in dict.items:  
        self.dict[key] = value
```

```
def pop(self, key):  
    value = self.dict[key]  
    del self.dict[key]  
    return value
```

```
def popitem(self):  
    key, value = next(self.dict.items)  
    del self.dict[key]  
    return key, value
```

```
def getitem(self, key):  
    value = self.dict[key]  
    return value
```

```
def setitem(self, key, value):  
    self.dict[key] = value
```

```
def delitem(self, key):  
    del self.dict[key]
```

```
def __len__(self):  
    return len(self.dict)
```

```
def __iter__(self):  
    return iter(self.dict)
```

```
def __contains__(self, key):  
    return key in self.dict
```

```
def __eq__(self, other):  
    if not isinstance(other, dict):  
        return False  
    if len(self) != len(other):  
        return False  
    for key, value in self.dict.items:  
        if key in other:  
            if value != other[key]:  
                return False  
        else:  
            return False  
    return True
```

```
def __ne__(self, other):  
    return not self == other
```

```
def __str__(self):  
    return str(self.dict)
```

```
def __repr__(self):  
    return repr(self.dict)
```



Container Images

Status



Status

History of Containers in Airflow: C

- Used for CI for > 2 years: Gerardo Curiel
- Optimized and incorporated by Breeze 1.5 years ago or so
- Docker Compose as execution engine
- Slimmed down recently (Thanks Ash!)
- Optimized for development use



Status

History of Containers in Airflow: Prod

- Puckel image created by Matthieu "Puckel_" Roisil (Thanks Matthieu!)
 - Used by many users in production
 - Used by the publicly available Helm Chart (not managed by community)
- Official Production Image (managed by community)
 - Alpha Quality community image in 1.10.10
 - Beta Quality community image in 1.10.11 (now!)



Status

State of the Official Production image

- Beta Quality - usable for production
- Most important feedback incorporated
- Already used in production
- Public Helm Chart switched to the Official Production Image
- Community Helm Chart (donated by Astronomer!) uses it for testing
- Stable version in v1-10-stable, development in master



Container Images

Internals



Internals: DockerHub releases

Released image

- ~ 210 MB compressed size
- Python: 2.7, 3.5, 3.6, 3.7, 3.8
- 1.10.11 = Python 3.6
- manually released
- using “1.10.11” tag
- latest = 1.10.11
- docker pull apache/airflow

The screenshot shows the DockerHub interface for the apache/airflow repository. The 'Tags' tab is selected, displaying several tags along with their details:

- 1.10.11**: Last updated a day ago by potiuk. Digest: 5b43a4b820eb. OS/ARCH: linux/amd64. Compressed Size: 208.86 MB.
- 1.10.11-python2.7**: Last updated a day ago by potiuk. Digest: 015eaac4f0de. OS/ARCH: linux/amd64. Compressed Size: 213.38 MB.
- 1.10.11-python3.5**: Last updated a day ago by potiuk. Digest: d30b4d7a7020. OS/ARCH: linux/amd64. Compressed Size: 213.65 MB.
- 1.10.11-python3.7**: Last updated a day ago by potiuk. Digest: a0032be900db. OS/ARCH: linux/amd64. Compressed Size: 210.01 MB.
- 1.10.11-python3.6**: Last updated a day ago by potiuk. Digest: 5b43a4b820eb. OS/ARCH: linux/amd64. Compressed Size: 208.86 MB.
- latest**: Last updated a day ago by potiuk. Digest: 5b43a4b820eb. OS/ARCH: linux/amd64. Compressed Size: 208.86 MB.

Red boxes highlight the 'IMAGE' column, the 'docker pull' command, and the 'COMRESSED SIZE' column for each tag entry. A red bracket groups the first two tag entries, and another red bracket groups the last three tag entries.



Internals: Releasing the image

Container Image or Container File ?

- Apache Software Foundation releases sources, not binaries
- Binaries can only be released for convenience of users
- Binaries must be rebuildable from released sources (PyPI, for example)
- Users should be able to build the software they need
- Should we release Container Image, Container File, or both?



Internals

Features of the production image

- Optimised for size (Compressed: ~230MB, ~800 MB on disk)
- Python 3.6, 3.7, 3.8 (2.0 and 1.10.*), 2.7, 3.5 (1.10.*)
- Extras installed:
 - async,aws,azure,celery,dask,elasticsearch,gcp,kubernetes,mysql,postgres,redis,slack,ssh,statsd,virtualenv
- OpenShift compatible (dynamic uid allocation)
- Gunicorn using shared memory (optimised parallelism)



Internals

Features of the production image file

- Builds optimised image
- Highly customizable (ARGs)
- Multi segmented (build + main)



Internals: build image

Build image

- Pass arguments
- Define variables
- Install apt dependencies (with dev ones)
- Install airflow (sources, pip, github): --user
- Include constraints
- Transpile website (yarn)
- ~700 MB compressed, ~2GB on disk
- Root user



```
ARG AIRFLOW_VERSION="2.0.0.dev0"
ARG AIRFLOW_EXTRAS="async,aws,azure,celery,dask,elasticsearch,gcp,kubernetes,mysql,postgres,redis,slack
# ...
#####
# This is the build image where we build all dependencies
#####
FROM ${PYTHON_BASE_IMAGE} as airflow-build-image
ARG PYTHON_BASE_IMAGE
ENV PYTHON_BASE_IMAGE=${PYTHON_BASE_IMAGE}
# ...
RUN apt-get update \
    && apt-get install -y --no-install-recommends \
        apt-transport-https \
        apt-utils \
        build-essential \
        freetds-bin \
        freetds-dev \
# ...
    && apt-get autoremove -yqq --purge \
    && apt-get clean \
    && rm -rf "/var/lib/apt/lists/*"
# ...
RUN pip install --user \
    https://github.com/${AIRFLOW_REPO}/archive/${AIRFLOW_BRANCH}.tar.gz#egg=apache-airflow[${AIRFLOW_EXTRAS}]
    --constraint \
    "https://raw.githubusercontent.com/${AIRFLOW_REPO}/${AIRFLOW_BRANCH}/requirements/requirements-pyton3.txt"
    && pip uninstall --yes apache-airflow;

ARG CONSTRAINT_REQUIREMENTS="requirements/requirements-pyton${PYTHON_MAJOR_MINOR_VERSION}.txt"
ENV CONSTRAINT_REQUIREMENTS=${CONSTRAINT_REQUIREMENTS}

WORKDIR /opt/airflow

# hadolint ignore=DL3020
ADD "${CONSTRAINT_REQUIREMENTS}" /requirements.txt

RUN pip install --user "${AIRFLOW_INSTALL_SOURCES}[${AIRFLOW_EXTRAS}]${AIRFLOW_INSTALL_VERSION}" \
    --constraint /requirements.txt
# ...

RUN AIRFLOW_SITE_PACKAGE="/root/.local/lib/python${PYTHON_MAJOR_MINOR_VERSION}/site-packages/airflow"; \
if [[ -f "${AIRFLOW_SITE_PACKAGE}/www_rbac/package.json" ]]; then \
    WWW_DIR="${AIRFLOW_SITE_PACKAGE}/www_rbac"; \
elif [[ -f "${AIRFLOW_SITE_PACKAGE}/www/package.json" ]]; then \
    WWW_DIR="${AIRFLOW_SITE_PACKAGE}/www"; \
fi; \
if [[ ${WWW_DIR} != "" ]]; then \
    yarn --cwd "${WWW_DIR}" install --frozen-lockfile --no-cache; \
    yarn --cwd "${WWW_DIR}" run prod; \
    rm -rf "${WWW_DIR}/node_modules"; \
fi
```

(side comment)

~ 730 modules

~ 360 MB



Internals: main image

Main image

- Pass arguments/ define variables
- Install apt dependencies (without dev!)
- Add user
- Uses root group (OpenShift)
- Copy(!) Airflow
- Copy DAGs (optionally)
- Copy entrypoint and clean-logs
- Access to /etc/passwd
- Embed dags (for tests)
- Optimized Gunicorn parallelism
- Set working dir
- Exposes port
- Set user
- Entrypoint and command
- ~230 MB compressed, ~800MB on disk

```
#####
# This is the actual Airflow image - much smaller than the build one. We copy
# installed Airflow and all it's dependencies from the build image to make it smaller.
#####

FROM ${PYTHON_BASE_IMAGE} as main
SHELL ["/bin/bash", "-o", "pipefail", "-e", "-u", "-x", "-c"]
#
ARG PYTHON_BASE_IMAGE
ENV PYTHON_BASE_IMAGE=${PYTHON_BASE_IMAGE}
#
RUN apt-get update \
    && apt-get install -y --no-install-recommends \
        apt-transport-https \
        apt-utils \
        ca-certificates \
        curl \
#
# ...
RUN addgroup --gid "${AIRFLOW_GID}" "airflow" && \
    adduser --quiet "airflow" --uid "${AIRFLOW_UID}" \
    --gid "${AIRFLOW_GID}" \
    --home "${AIRFLOW_USER_HOME_DIR}"
#
COPY --chown=airflow:root --from=airflow-build-image /root/.local "${AIRFLOW_USER_HOME_DIR}/.local"
#
COPY scripts/prod/entrypoint_prod.sh /entrypoint
COPY scripts/prod/clean-logs.sh /clean-logs
#
ARG EMBEDDED_DAGS="empty"
COPY --chown=airflow:root ${EMBEDDED_DAGS}/ ${AIRFLOW_HOME}/dags/
#
# Make /etc/passwd root-group-writeable so that user can be dynamically added by OpenShift
# See https://github.com/apache/airflow/issues/9248
RUN chmod g=u /etc/passwd
#
ENV PATH="${AIRFLOW_USER_HOME_DIR}/.local/bin:${PATH}"
ENV GUNICORN_CMD_ARGS="--worker-tmp-dir /dev/shm"
#
WORKDIR ${AIRFLOW_HOME}
#
EXPOSE 8080
#
USER ${AIRFLOW_UID}
#
ENTRYPOINT ["/usr/bin/dumb-init", "--", "/entrypoint"]
CMD ["--help"]
```

IMAGE
master-python3.6
Last updated 2 hours ago by apache
docker pull apache/airflow:master-pytho
DIGEST ba50cd9e3d4e
OS/ARCH linux/amd64
COMPRESSED SIZE 231.56 MB



Internals: entrypoint

- Creates user dynamically if missing (OpenShift)
- Fallbacks to sqlite metadata
- Waits until metadata DB is up
- Waits until broker DB is up
- If “bash” or “python” -> runs command
- Else execute airflow command

```
# In case the user is not locally created we automatically create it in /etc/passwd
# This is to handle OpenShift case where random UIDs are used
if ! whoami &> /dev/null; then
    if [[ -w /etc/passwd ]]; then
        echo "${USER_NAME:+$USER_NAME}:$(id -u):0:${USER_NAME:-default} user:${AIRFLOW_USER_HOME_DIR}::sbin/nologin" \
            >> /etc/passwd
    fi
    export HOME="${AIRFLOW_USER_HOME_DIR}"
fi

# if no DB configured - use sqlite db by default
AIRFLOW__CORE__SQLALCHEMY_CONN="${AIRFLOW__CORE__SQLALCHEMY_CONN:="sqlite:///${AIRFLOW_HOME}/airflow.db"}"

verify_db_connection "${AIRFLOW__CORE__SQLALCHEMY_CONN}"

AIRFLOW__CELERY__BROKER_URL=${AIRFLOW__CELERY__BROKER_URL:=}

if [[ -n ${AIRFLOW__CELERY__BROKER_URL} ]] && \
    [[ ${AIRFLOW_COMMAND} =~ ^(scheduler|worker|flower)$ ]]; then
    verify_db_connection "${AIRFLOW__CELERY__BROKER_URL}"
fi

if [[ ${AIRFLOW_COMMAND} == "bash" ]]; then
    shift
    exec "/bin/bash" "${@}"
elif [[ ${AIRFLOW_COMMAND} == "python" ]]; then
    shift
    exec "python" "${@}"
fi

# Run the command
exec airflow "${@}"
```



Internals: .dockerignore

- Ignores everything by default
- You must explicitly include what you want by “!”
- You can further exclude specific subdirectories/patterns
- We generate a lot of stuff in airflow sources
- Sending big context to Docker engine takes time
- You avoid accidental inclusion of unneeded artifacts

```
# NOTE! This docker ignore uses recommended technique
# Where everything is excluded by default and you deliberately
# Add only those directories/files you need. This is very useful
# To make sure that Docker context is always the same on any machine
# So that generated files are not accidentally added to the context
# This allows Docker's `COPY .` to behave in predictable way

# Ignore everything
**
```

```
# Allow only these directories
!airflow
!common
!dags
!dev
!docs
!licenses
!metastore_browser
!scripts
!tests
```

```
#....
```

```
# Now - ignore unnecessary files inside allowed directories
# This goes after the allowed directories

# Git version is dynamically generated
airflow/git_version

# Exclude static www files generated by NPM
airflow/www/static/coverage
airflow/www/static/dist
airflow/www/node_modules
# Exclude static www_rbac files generated by NPM in v1-10-test
airflow/www_rbac/static/coverage
airflow/www_rbac/static/dist
airflow/www_rbac/node_modules

# Exclude link to docs
airflow/www/static/docs

# Exclude python generated files
**/_pycache_/
**/*.py[cod]
**/*$py.class
**/.pytest_cache/
**/env/
```



Internals

How we test the image ?

- The image and chart are part of Apache Airflow monorepo
- We build the image with every PR (dependencies)
- We use it in the Kubernetes tests for master (Helm Chart integration)
- We will use released images in the Helm Chart (backward compatibility)
- We will add more tests for various Helm configurations



Container Images

Usage



Usage: Extending Airflow image - use released image



```
docker build . -t yourcompany/airflow:1.10.11-BUILD_ID
```

```
FROM apache/airflow:1.10.11

# change to root user temporarily
USER root

# Optionally install your own apt dependencies
RUN apt-get update \
    && apt-get install -y --no-install-recommends \
    emacs \
    && apt-get autoremove -yqq --purge \
    && apt-get clean \
    && rm -rf "/var/lib/apt/lists/*"

# Change back to the airflow user
USER airflow

# Add extra dependencies
RUN pip install --user numpy

# Embed DAGs (Optional) - DAGs can be baked in but also
# they can be git-synced or mounted from shared volume
COPY --chown=airflow:root dags-folder ${AIRFLOW_HOME}/dags/
```



apache/airflow:1.10.11



Container registry



Container
image

yourcompany/airflow:1.10.11-BUILD_ID



Usage

Extending image - Pros & Cons

Pros

- Use released images
- Simple build command
- Own Dockerfile
- No need for Airflow sources

Cons

- Potentially bigger size
- Predefined extras only
- Installs limited set of python dependencies



Usage: Customising Airflow image - default docker build

```
git clone git@github.com:apache/airflow.git  
cd airflow  
git checkout v1-10-stable
```



```
docker build .
```



Container image

Same as apache/airflow:1.10.11

- Python 3.6
- Default extras
- No additional dependencies



Usage: Customising Airflow image - use build args

- Installs from PyPi ==1.10.11
- Additional airflow extras, dev, runtime deps ...
- Does not use local sources (can be run from master including entrypoint!)

```
docker build . \
--build-arg PYTHON_BASE_IMAGE="python:3.7-slim-buster" \
--build-arg PYTHON_MAJOR_MINOR_VERSION=3.7 \
--build-arg AIRFLOW_INSTALL_SOURCES="apache-airflow" \
--build-arg AIRFLOW_INSTALL_VERSION=="==1.10.11" \
--build-arg CONSTRAINT_REQUIREMENTS=\
"https://raw.githubusercontent.com/apache/airflow/1.10.11/requirements/requirements-python3.7.txt" \
--build-arg AIRFLOW_SOURCES_FROM="empty" \
--build-arg AIRFLOW_SOURCES_TO="/empty" \
--build-arg ADDITIONAL_AIRFLOW_EXTRAS="jdbc"
--build-arg ADDITIONAL_DEV_DEPS="gcc g++"
--build-arg ADDITIONAL_RUNTIME_DEPS="default-jre-headless"
```



Usage: Image Customization options

- Choose Base image (python)
- Install Airflow from PyPI
- Install from GitHub branch/tag
- Install additional extras
- Install additional python deps
- Install additional apt dev deps
- Install additional apt runtime deps
- Choose different UID/GID
- Choose different AIRFLOW_HOME
- Choose different HOME dir
- Build Cassandra driver concurrently

See [IMAGES.rst](#) in the Airflow repo.

The following build arguments (--build-arg in docker build command) can be used for production images:

Build argument	Default value	Description
PYTHON_BASE_IMAGE	python:3.6-slim-buster	Base python image
PYTHON_MAJOR_MINOR_VERSION	3.6	major/minor version of Python (should match base image)
AIRFLOW_VERSION	2.0.0.dev0	version of Airflow
AIRFLOW_ORG	apache	Github organisation from which Airflow is installed (when installed from repo)
AIRFLOW_REPO	apache/airflow	the repository from which PIP dependencies are pre-installed
AIRFLOW_BRANCH	master	the branch from which PIP dependencies are pre-installed
AIRFLOW_GIT_REFERENCE	master	reference (branch or tag) from Github repository from which Airflow is installed (when installed from repo)
REQUIREMENTS_GIT_REFERENCE	master	reference (branch or tag) from Github repository from which requirements are downloaded for constraints (when installed from repo).
AIRFLOW_EXTRAS	(see Dockerfile)	Default extras with which airflow is installed
ADDITIONAL_AIRFLOW_EXTRAS		Optional additional extras with which airflow is installed
ADDITIONAL_PYTHON_DEPS		Optional python packages to extend the image with some extra dependencies
ADDITIONAL_DEV_DEPS		additional apt dev dependencies to install
ADDITIONAL_RUNTIME_DEPS		additional apt runtime dependencies to install
EMBEDDED_DAGS	empty	Folder containing dags embedded into the image in the \${AIRFLOW_HOME}/dags dir
AIRFLOW_HOME	/opt/airflow	Airflow's HOME (that's where logs and sqlite databases are stored)
AIRFLOW_UID	50000	Airflow user UID
AIRFLOW_GID	50000	Airflow group GID. Note that most files created on behalf of airflow user belong to the root group (0) to keep OpenShift Guidelines compatibility
AIRFLOW_USER_HOME_DIR	/home/airflow	Home directory of the Airflow user
PIP_VERSION	19.0.2	version of PIP to use
CASS_DRIVER_BUILD_CONCURRENCY	8	Number of processors to use for cassandra PIP install (speeds up installing in case cassandra extra is used).



Usage: It's a Breeze to build images

- Breeze - development and test environment
- Supports building production image
- Auto-complete of options
- New Breeze video showing building production images:
<https://s.apache.org/airflow-breeze>
- `./breeze build-image --help`

```
./breeze build-image --production-image --additional-extras "jira"

./breeze build-image --production-image --python 3.7 \
--additional-extras "jira"

./breeze build-image --production-image \
--additional-python-deps "torchio==0.17.10"

./breeze build-image --production-image \
--additional-dev-deps "libasound2-dev" \
--additional-runtime-deps "libasound2"

./breeze build-image --production-image \
--additional-extras "jira" --install-airflow-version="1.10.11"
```

See **BREEZE.rst** in the Airflow repo



Usage

Customising image - Pros & Cons

Pros

- Highly optimized for size
- Build image from sources
(security reviews!)
- Can add any extras
- Can add any dependency
- Breeze build commands
- Works from master and 1.10.*

Cons

- Need access to airflow sources
- Complex build command
- Need to understand internals



Why not eat and have cake ?

```
git clone git@github.com:apache/airflow.git  
cd airflow  
git checkout v1-10-stable
```



```
./breze build-image --production-image --additional-extras "jira" \  
--install-airflow-version "1.10.11"
```

When dependencies
change



Runtime
Container
image

When DAGs
change

```
FROM base-image-for-your-company:1.10.11-2020-07-14  
COPY --chown airflow:root dags-folder "${AIRFLOW_HOME}/dags"
```



base-image-for-your-company:1.10.11-2020-07-14



Base
Container
image





Usage

How to deploy the images ?

- Docker and Docker-Compose - not recommended for production
- Managed Container Services
 - Managed: Amazon ECS, Google Container on VMs, Azure Container Instances
- Kubernetes on-Prem:
 - Helm Chart
 - Airflow Operator (not recommended yet)
- Managed Kubernetes: Amazon EKS, Google GKE, Azure AKS
- OpenShift (also Kubernetes)



Container Images

Future

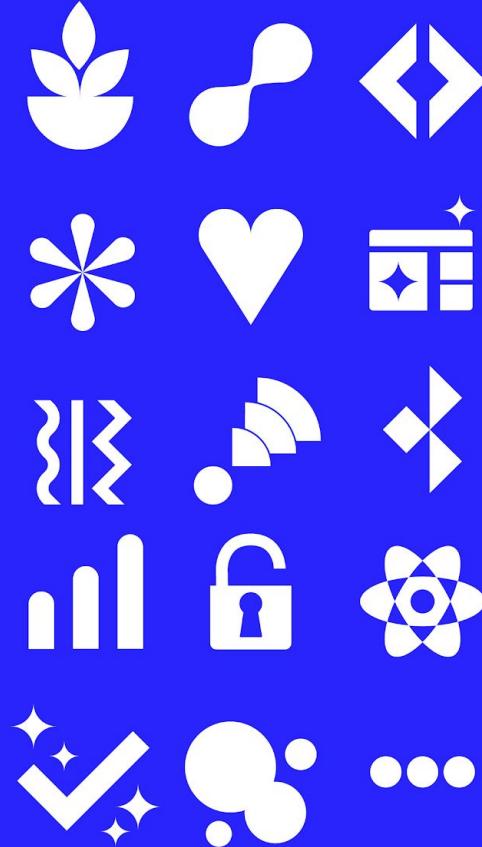


Future

What is the future for Airflow images?

- It won't change too much !
- Better automated testing via Helm Chart
- Automated releases for 2.0
- ARM support might be the big one. (Apple Mac OS)
- Official Docker Compose
- Smaller features (depends on feedback and expectations):
 - ON BUILD support ?
 - AIRFLOW__CORE__SQLALCHEMY_CONN_CMD, AIRFLOW__CELERY__BROKER_URL_CMD support ?
 - Automated user creation ?

Q&A



Thanks! Polidea.

hello@polidea.com

Bē  f  in  