

```
FROM AIRFLOW IMPORT DAG
```

# Airflow the perfect match in our Analytics Pipeline

**Sergio Camilo Fandiño Hernández**  
Senior Business Intelligence Architect @LOVOO

**Airflow**   
Summit 2020

## AGENDA

---

1. Why we met?
2. How we met?
3. The first date!
4. Fun dates!
5. Is there any dynamic in between?
6. Recap and conclusion

# About LOVOO

- LOVOO is a dating and social app and the place for chatting, live streaming, watching streams and getting to know people.
- Germany - Dresden & Berlin - 2011
- Acquired by The Meet Group (NASDAQ:MEET) in 2017
- Top 3 Dating App in Europe
- + 280 TB of Data
- ~ 6 TB Monthly Growth
- + 3 TB daily total aggregated data
- + 36 TB Swipes (162,824,303,474)

# Analytics

- 1 Head
- 6 Data Analysts
- 2 BI Architects

- 
- Product
  - Finance
  - Marketing
  - Talent Management
  - Customer Insights
  - CRM

# What can you expect?

My main purpose today is to tell you about our journey with Airflow as well as a few different use cases that could also boost the work of your Analytics/BI team on a daily basis.

- Pieces of code (examples)
- Way too many screenshots

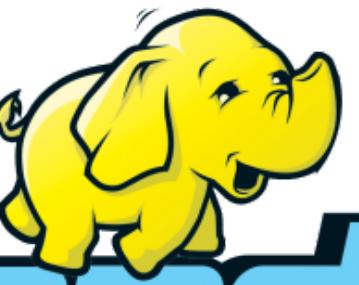
## AGENDA

---

1. Why we met?
2. How we met?
3. The first date!
4. Fun dates!
5. Is there any dynamic in between?
6. Recap and conclusion

# On-premise

cloudera

 **hadoop**

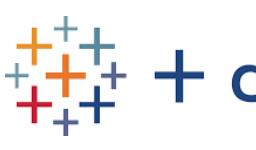


 + a b l e a u®

 oozi-e

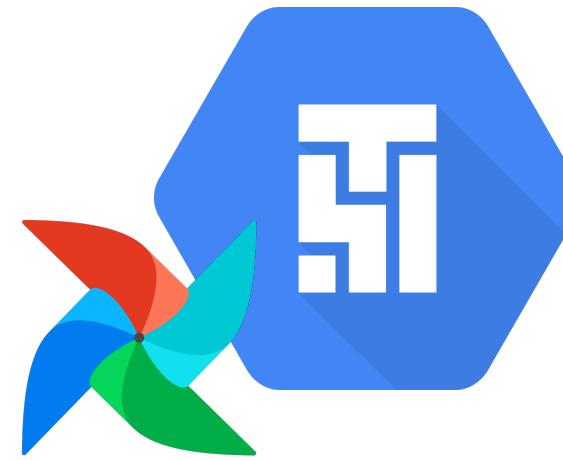
# We went Cloud



 + a b | e a u® 

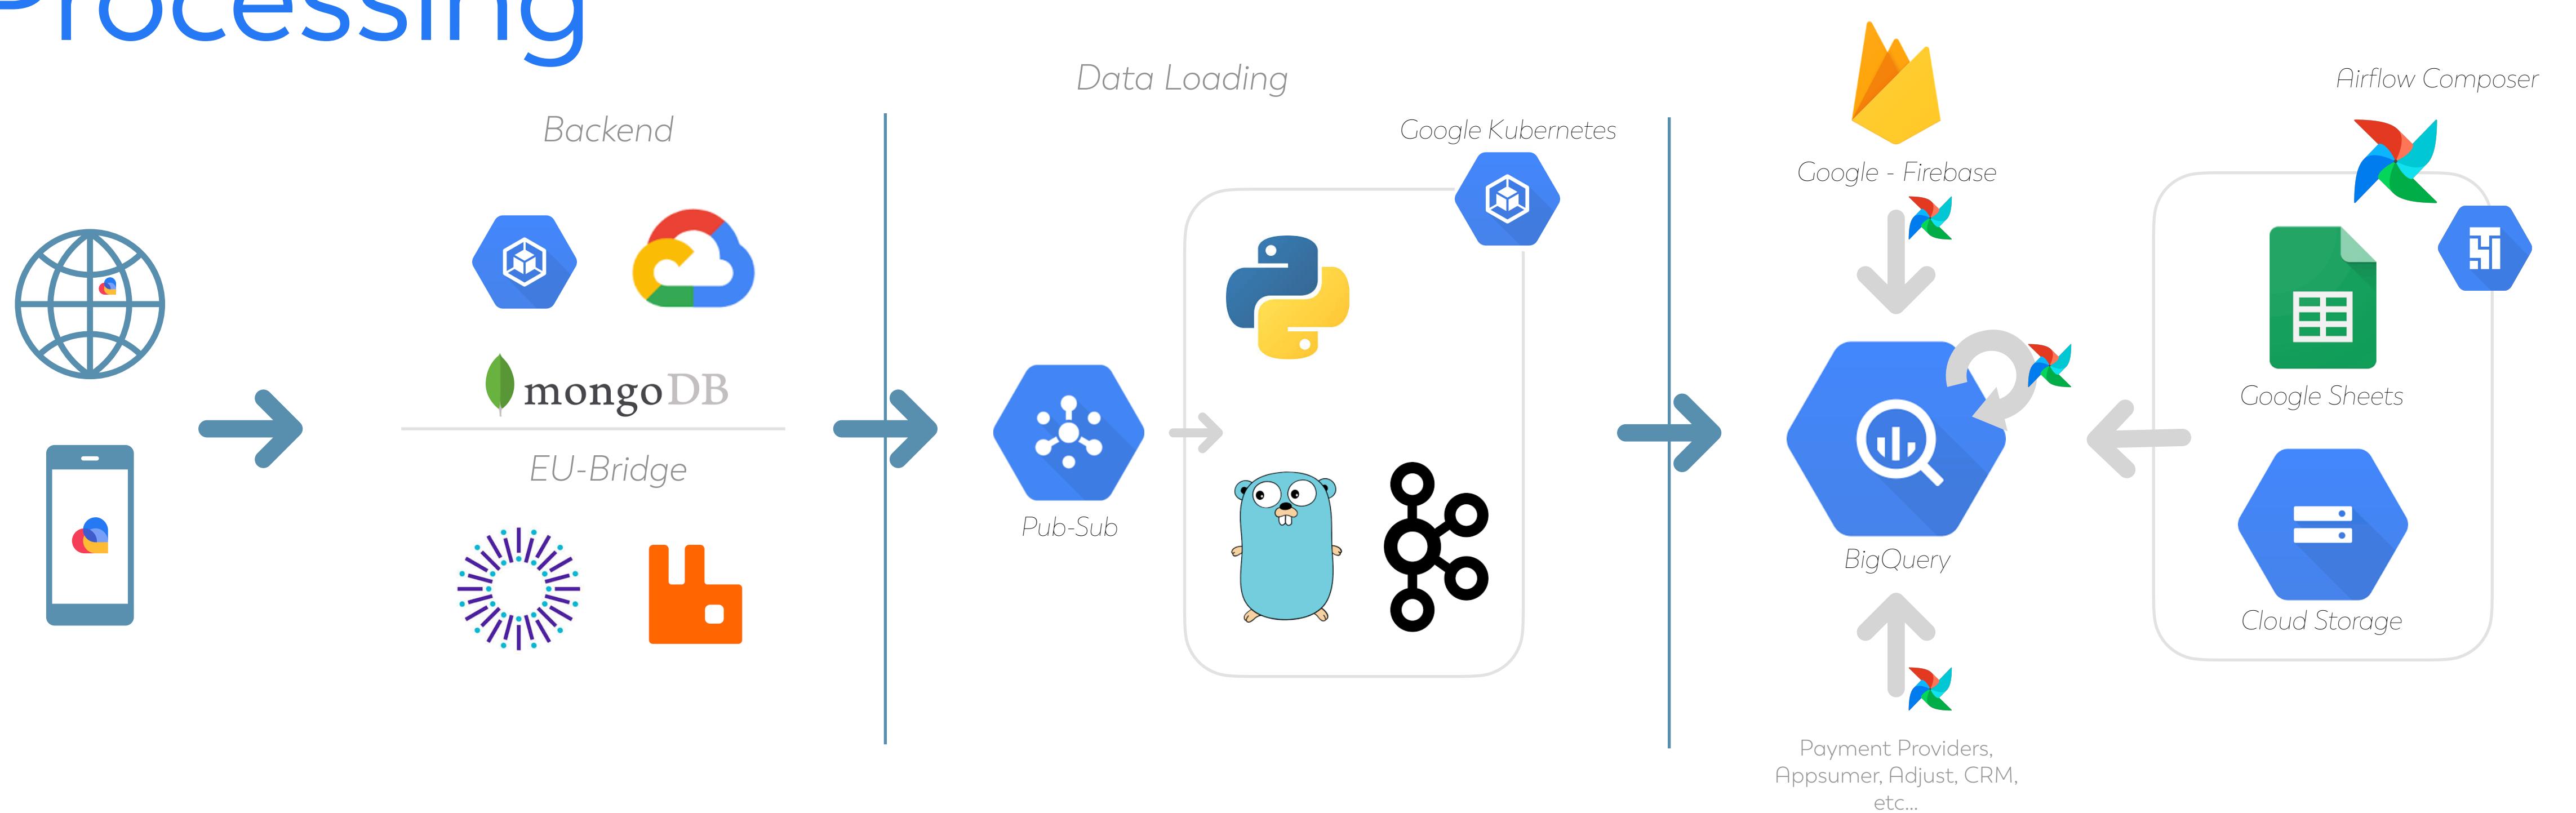


Google Cloud

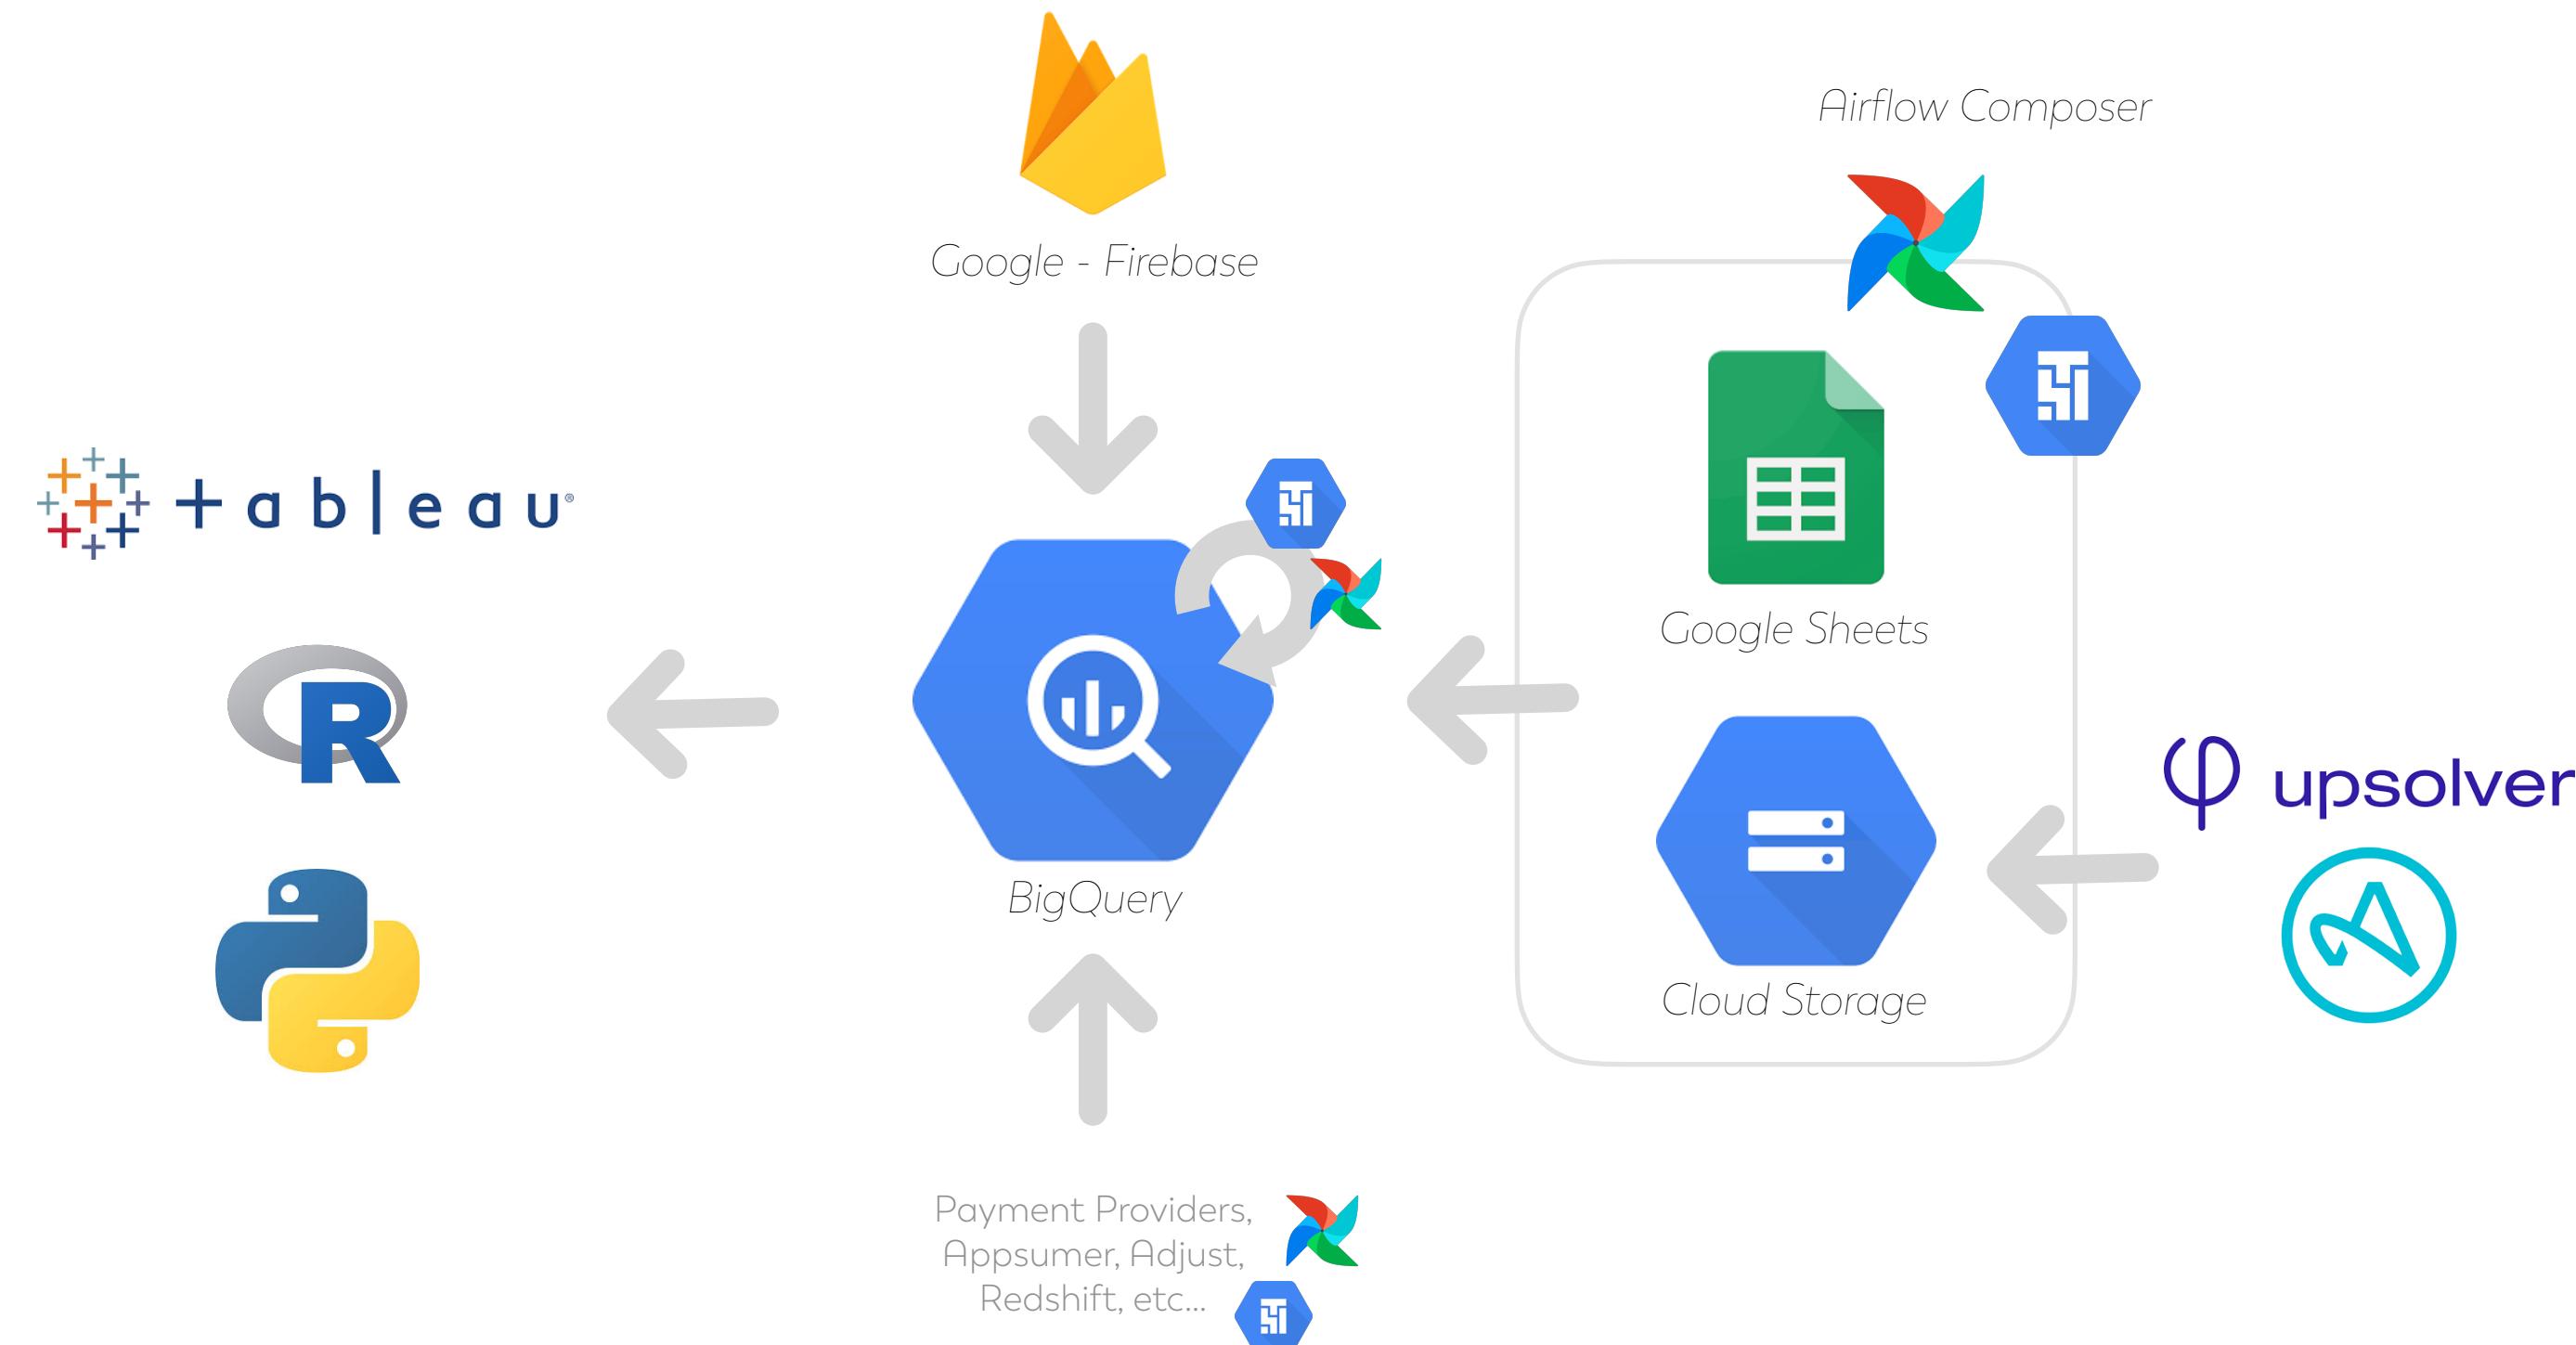


THE PROFILE DETAILS...

# Data Processing



# Analytics Data-Core



## AGENDA

---

1. Why we met?
2. How we met?
3. The first date!
4. Fun dates!
5. Is there any dynamic in between?
6. Recap and conclusion

# Orchestration Tool

- Identify what is out there
- Costs?
- Scalability?
- Data sources compatibility?
- Knowledge/Human Resources?

RIGHT SWIPE...

# Airflow



- Great community
- Game changer
- Mobile App
- Python
- BigQuery

# Google Cloud Composer



- Fully Managed Airflow
- Scalable
- IAP - Secure
- Focus on building the Analytics data pipeline
- Ease of implementation

# Google Cloud Composer



- Fully Managed Airflow



This is an alpha release of Cloud Composer. This product might be changed in backward-incompatible ways and is not recommended for production use. It is not subject to any SLA or deprecation policy. This product is not intended for real-time usage in critical applications.

- Focus on building the Analytics data pipeline
- Ease of implementation

 **Confidential Material:** This page is confidential. Do not share or discuss until authorized to do so.

## AGENDA

---

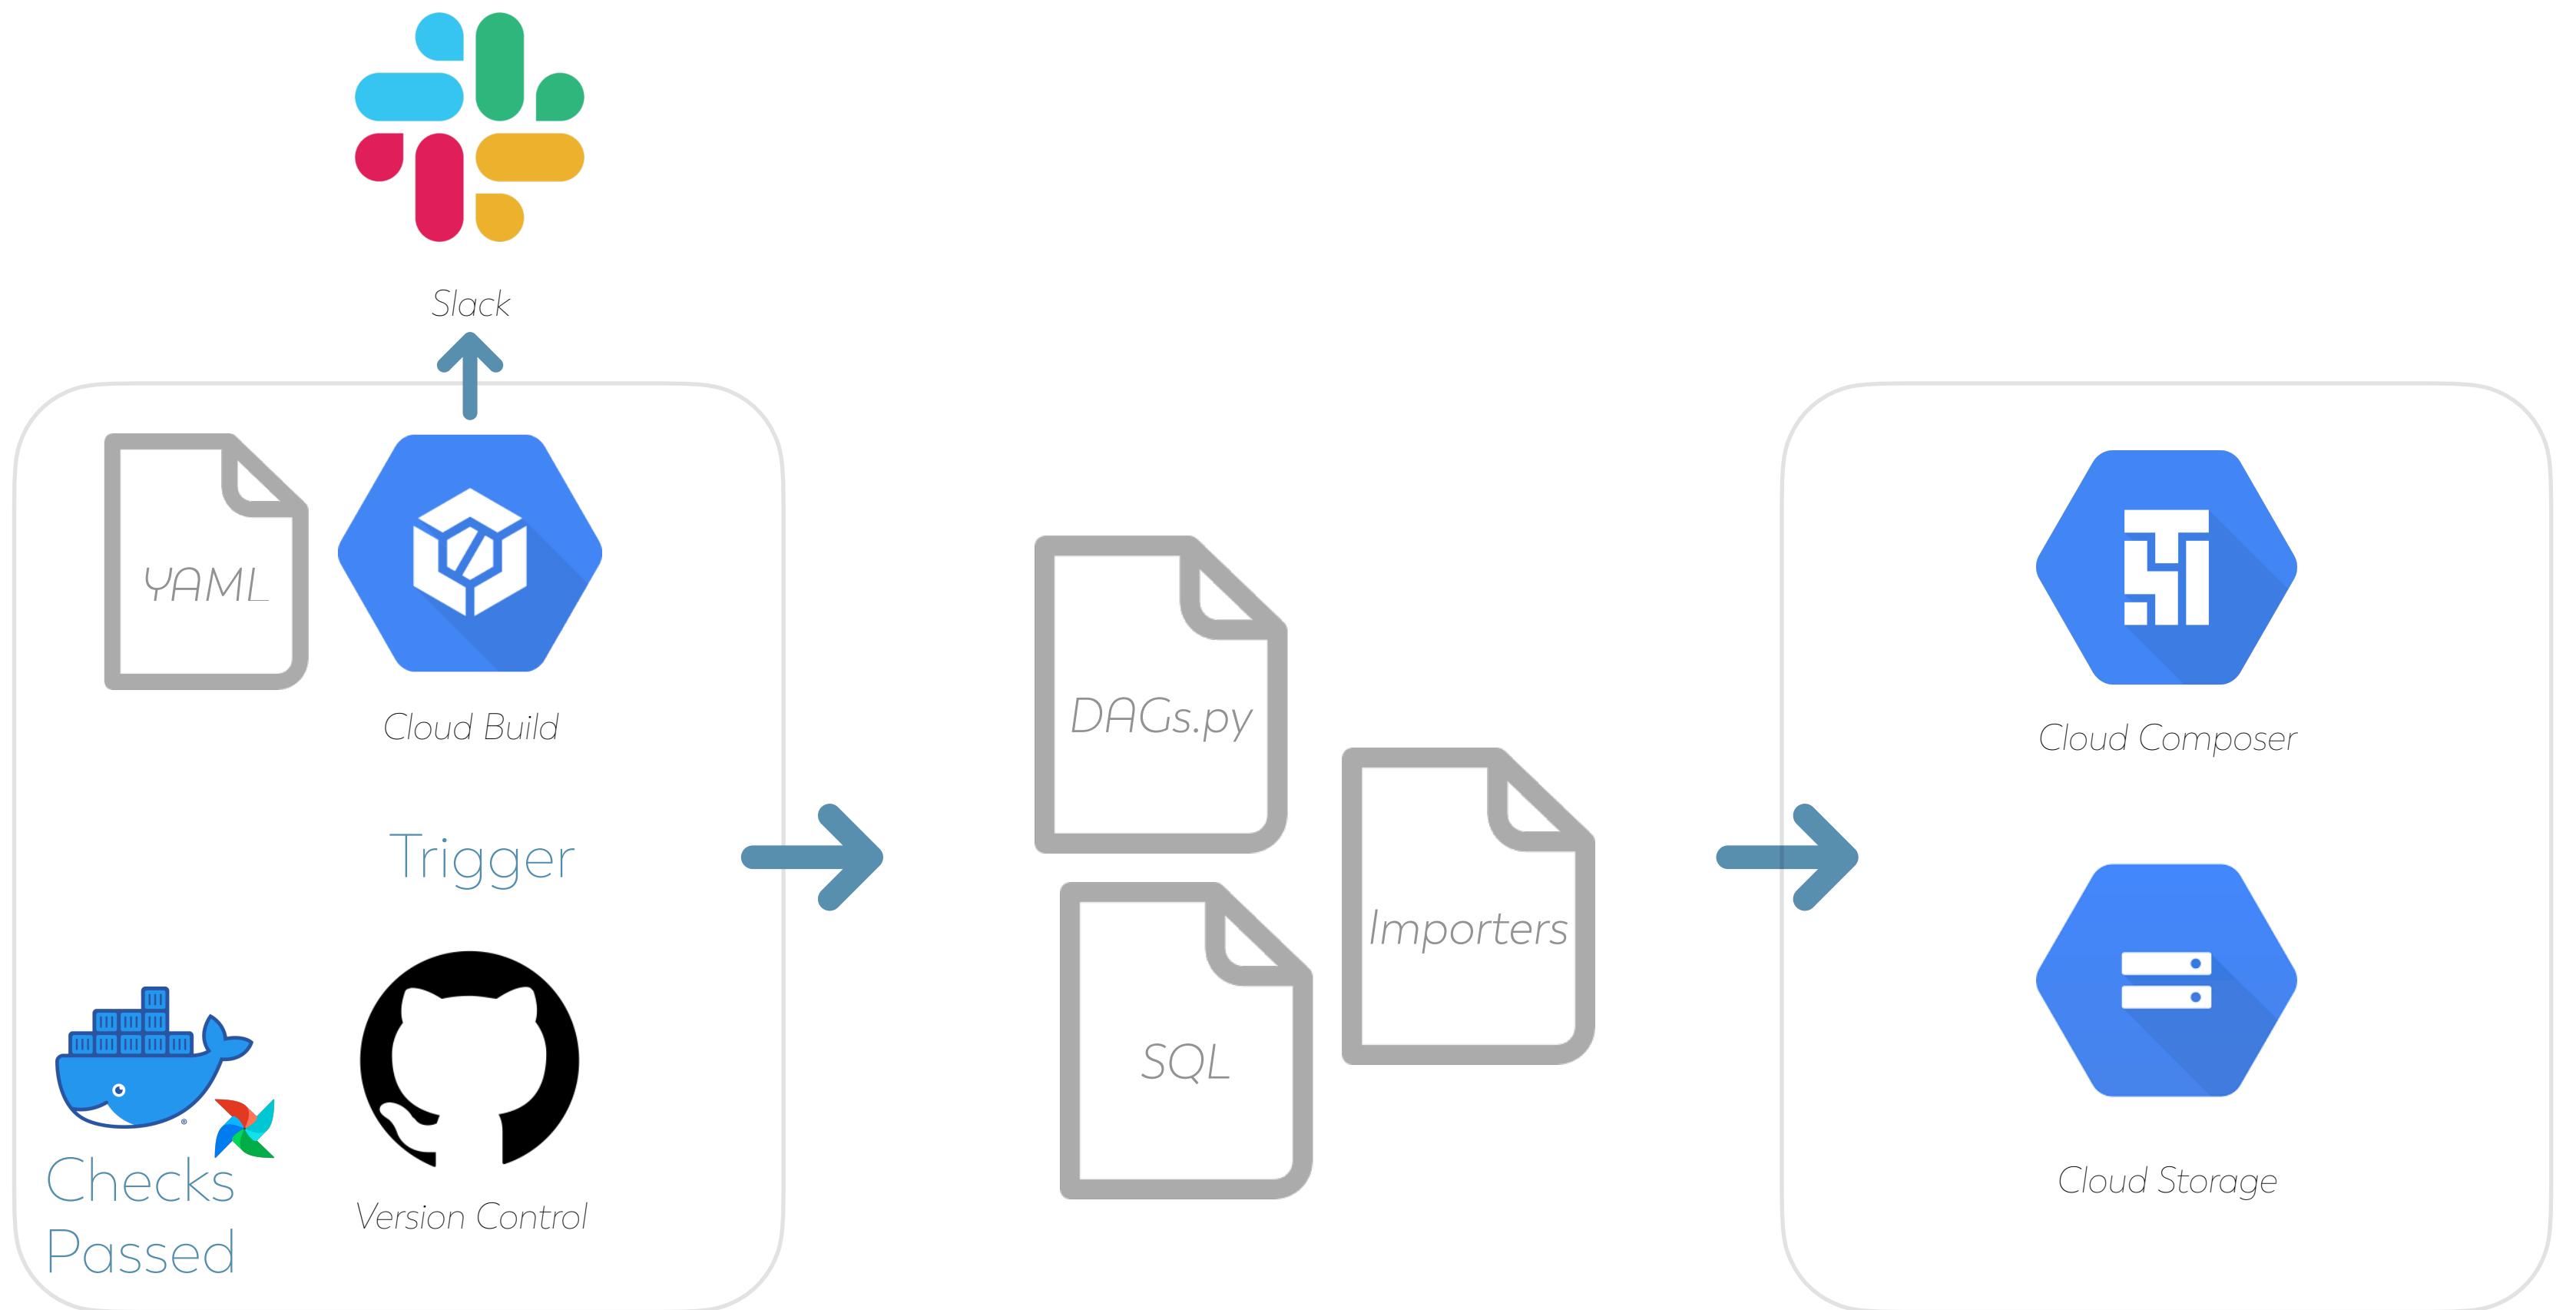
1. Why we met?
2. How we met?
- 3. The first date!**
4. Fun dates!
5. Is there any dynamic in between?
6. Recap and conclusion

# TODO List

- SQL Scripts → Data Modeling
- DAGs
- Permissions - Service Accounts
- Data Importers
- Create a Composer Environment
- How do we deploy? → CI/CD

GROWING TOGETHER!

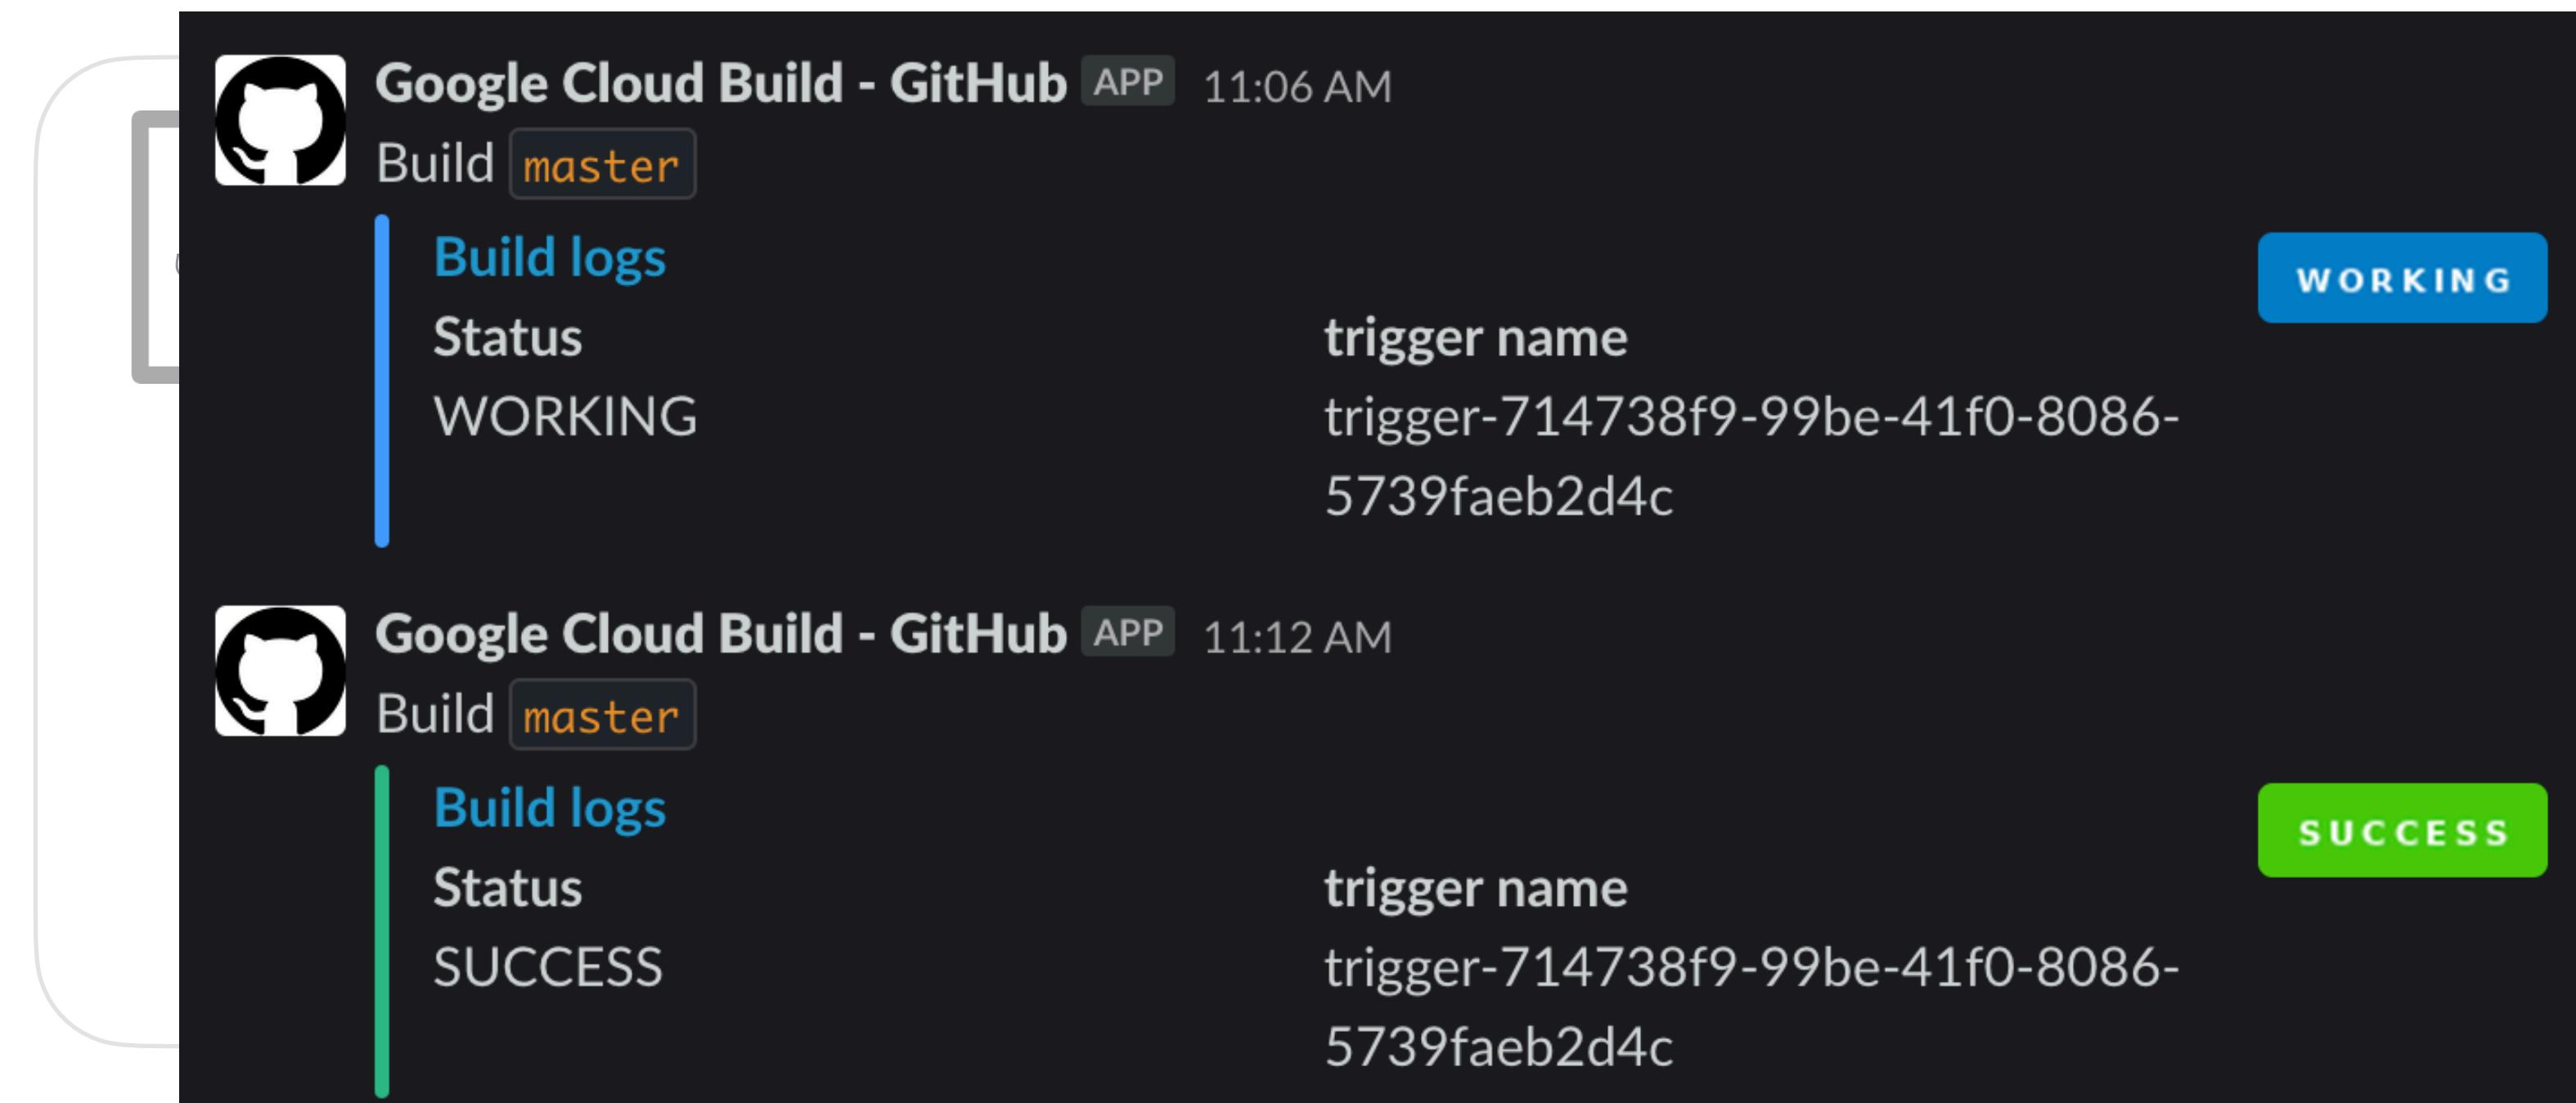
# CI/CD



# CI/CD



Slack



**Google Cloud Build - GitHub APP 11:06 AM**  
Build master

**Build logs**  
**Status**  
**WORKING**

**trigger name**  
trigger-714738f9-99be-41f0-8086-  
5739faeb2d4c

**WORKING**

**Google Cloud Build - GitHub APP 11:12 AM**  
Build master

**Build logs**  
**Status**  
**SUCCESS**

**trigger name**  
trigger-714738f9-99be-41f0-8086-  
5739faeb2d4c

**SUCCESS**

## AGENDA

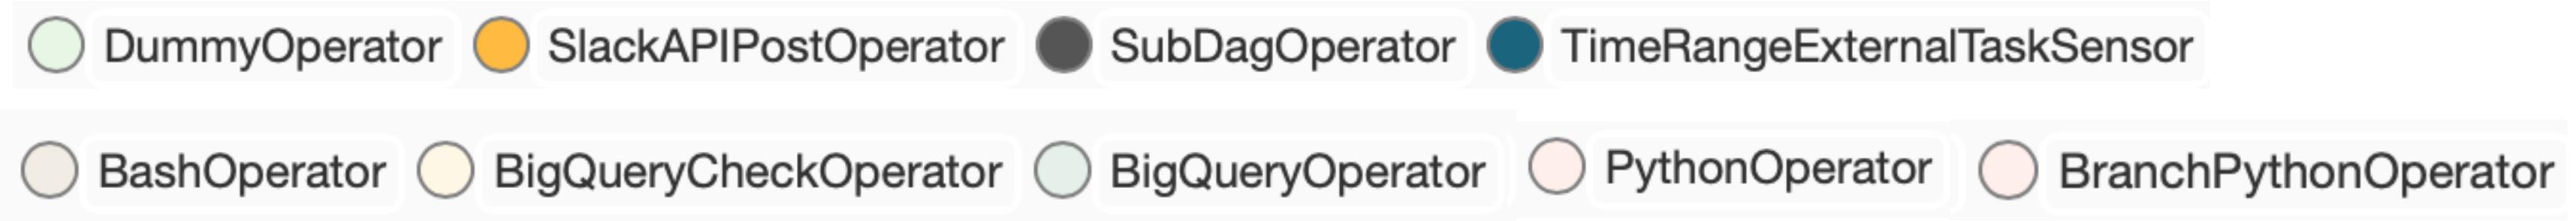
---

1. Why we met?
2. How we met?
3. The first date!
4. Fun dates!
5. Is there any dynamic in between?
6. Recap and conclusion

## HOW DOES IT LOOK LIKE?

## DAGs

## Operators



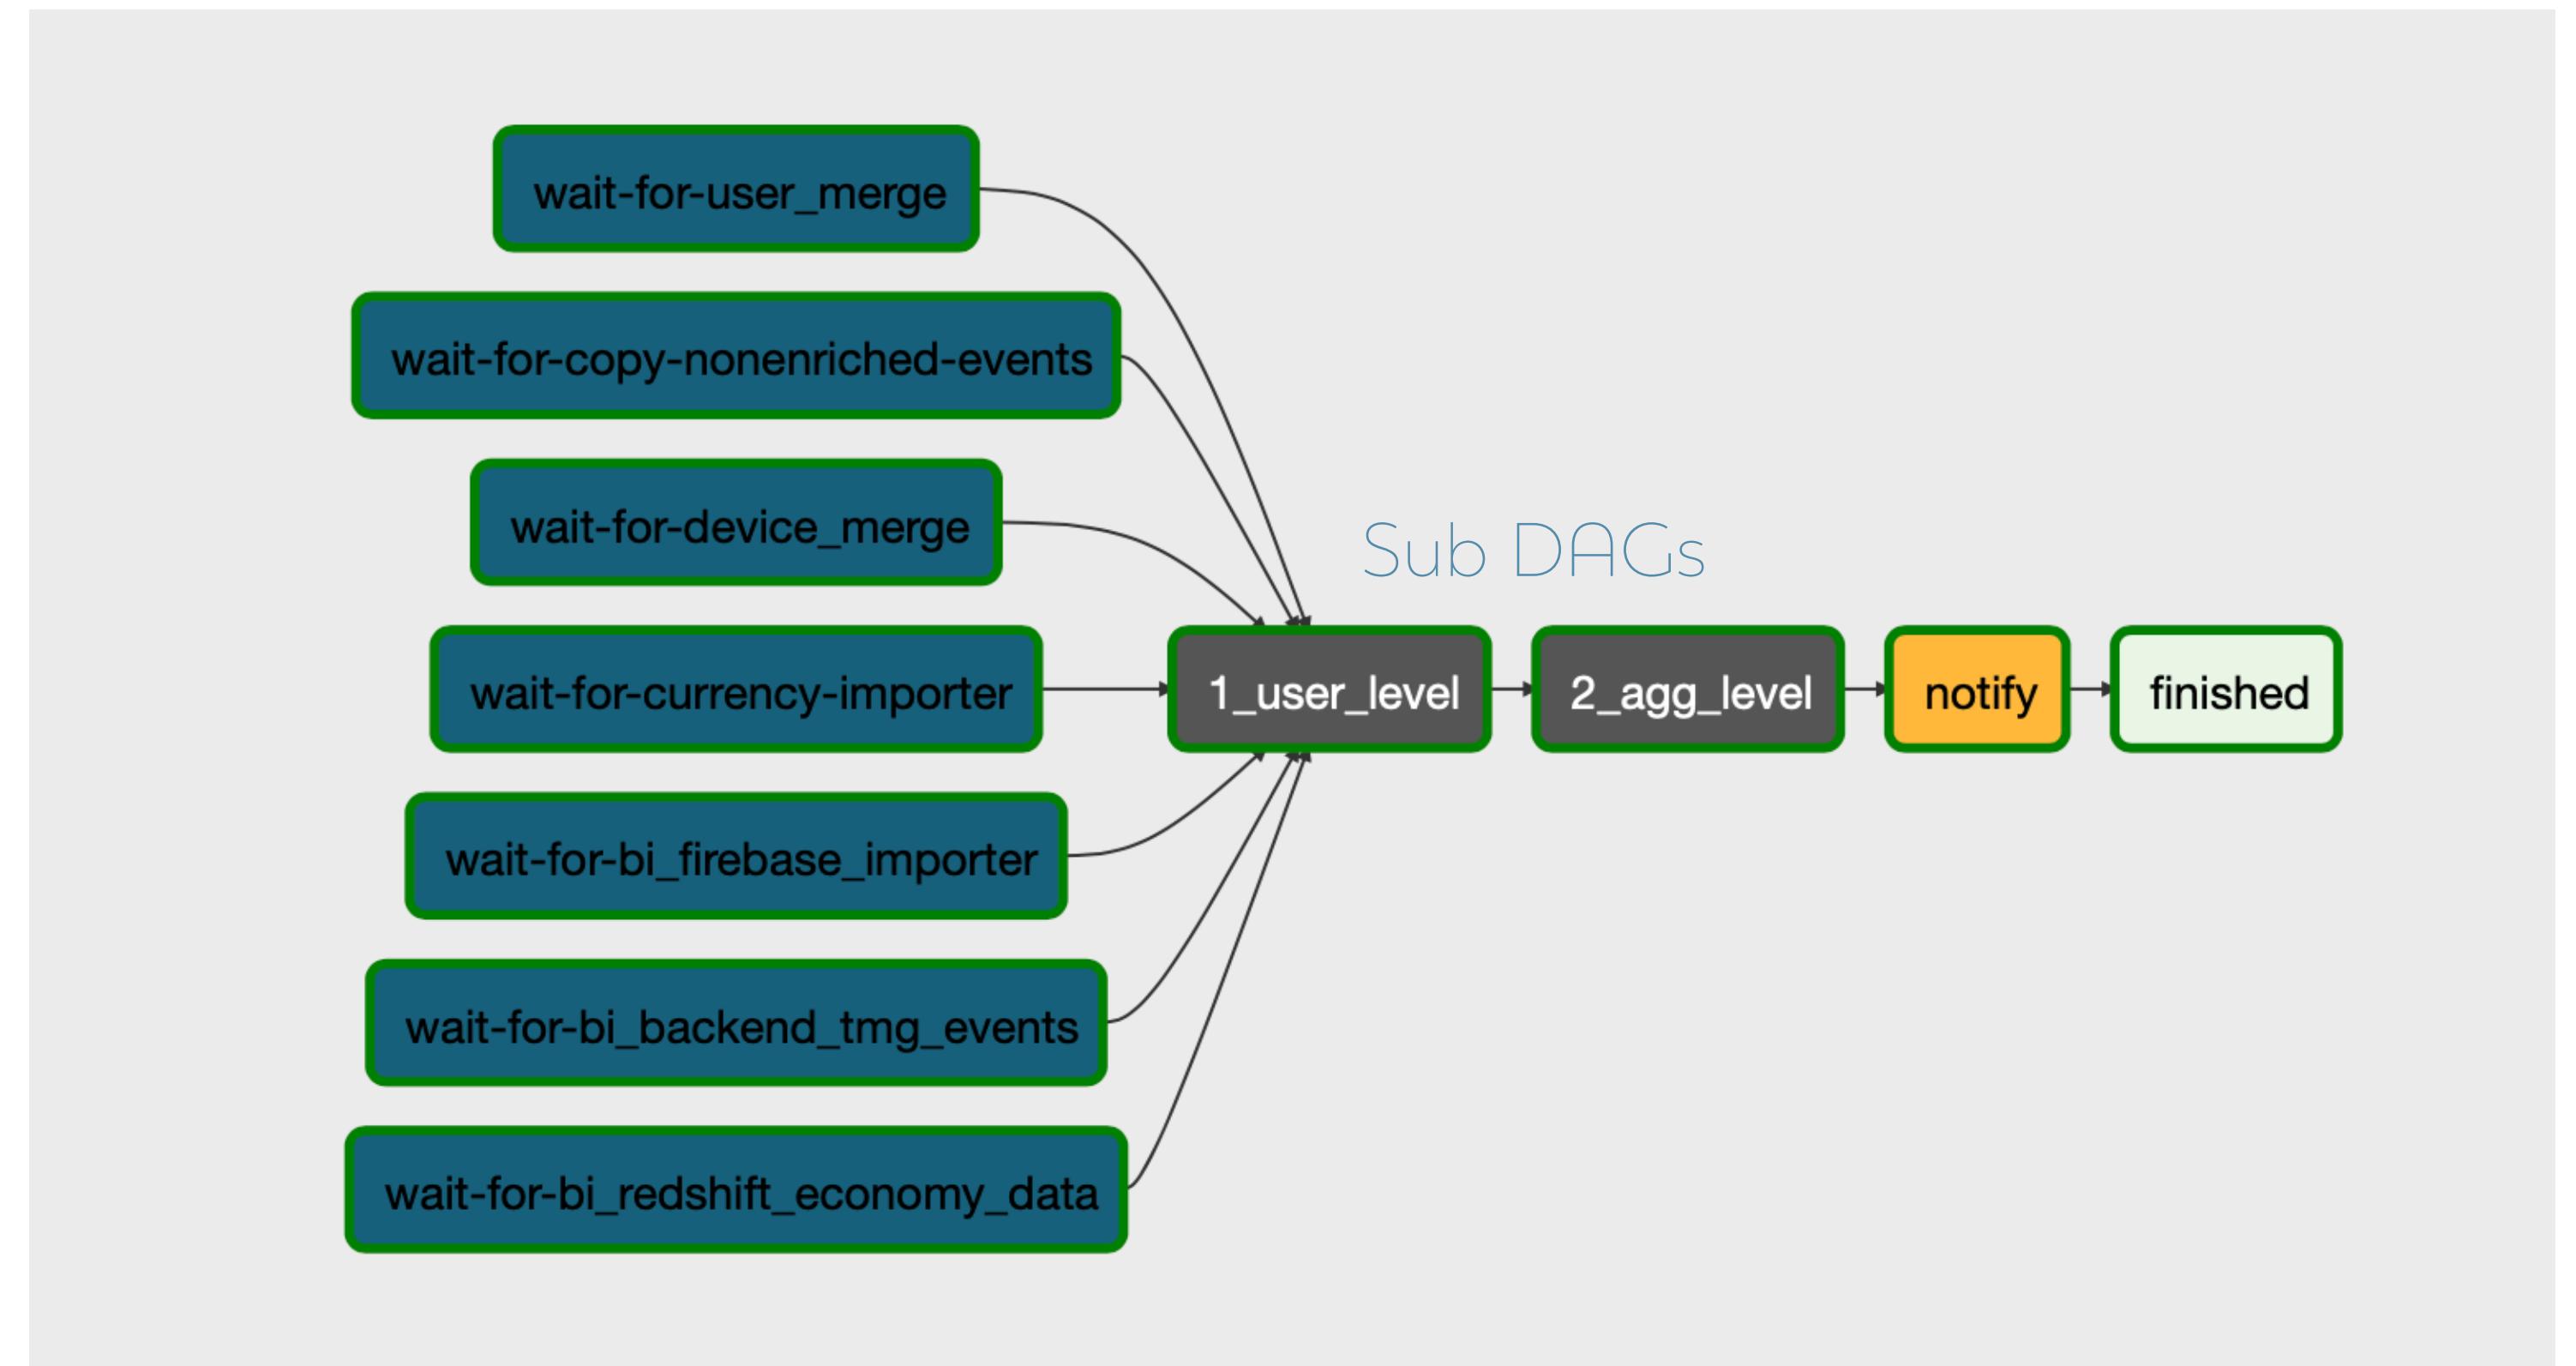
- 26 DAGs
- Sub-DAGs
- Branching
- Jinja Templating
- Hooks
- Pools
- Trigger rules

	DAG	Schedule
<input checked="" type="checkbox"/> <span>On</span>	adevents-repair	0 4 * * *
<input checked="" type="checkbox"/> <span>On</span>	airflow_monitoring	None
<input checked="" type="checkbox"/> <span>On</span>	analytics_jobs	0 5 * * *
<input checked="" type="checkbox"/> <span>On</span>	analytics_jobs_live	0 5 * * *
<input checked="" type="checkbox"/> <span>On</span>	antispam-creditfarm-detection	@daily
<input checked="" type="checkbox"/> <span>On</span>	antispam-reputation-modeltraining	@daily
<input checked="" type="checkbox"/> <span>On</span>	appsumer-importer	00 11 * * *
<input checked="" type="checkbox"/> <span>On</span>	appsumer-importer-hayi	00 12 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_backend_tmg_events	30 2 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_data_check	40 3 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_firebase_importer	40 4 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_firebase_live_events	40 4 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_marketing_events_jobs	50 6 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_marketing_jobs	50 9 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_payment_provider_apis	40 4 * * *
<input checked="" type="checkbox"/> <span>On</span>	bi_redshift_economy_data	30 4 * * *

PRETTY ON THE OUTSIDE...

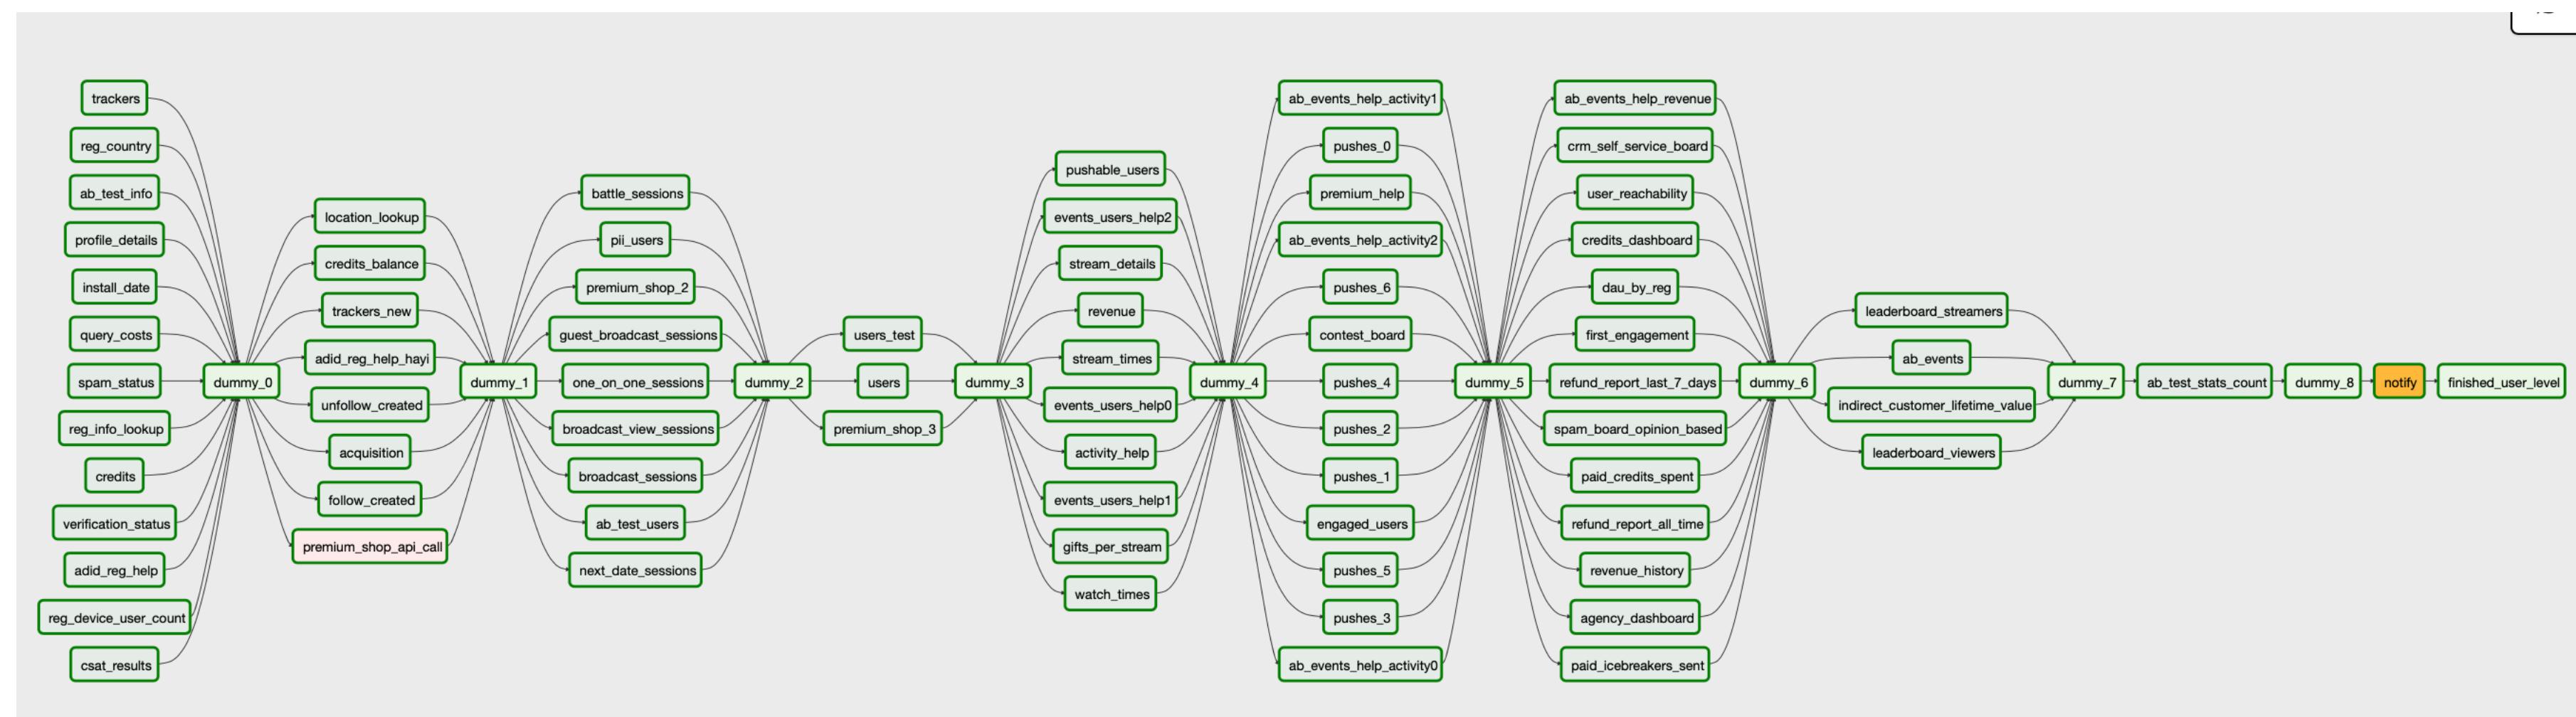
# The Core

Analytics - Workflow

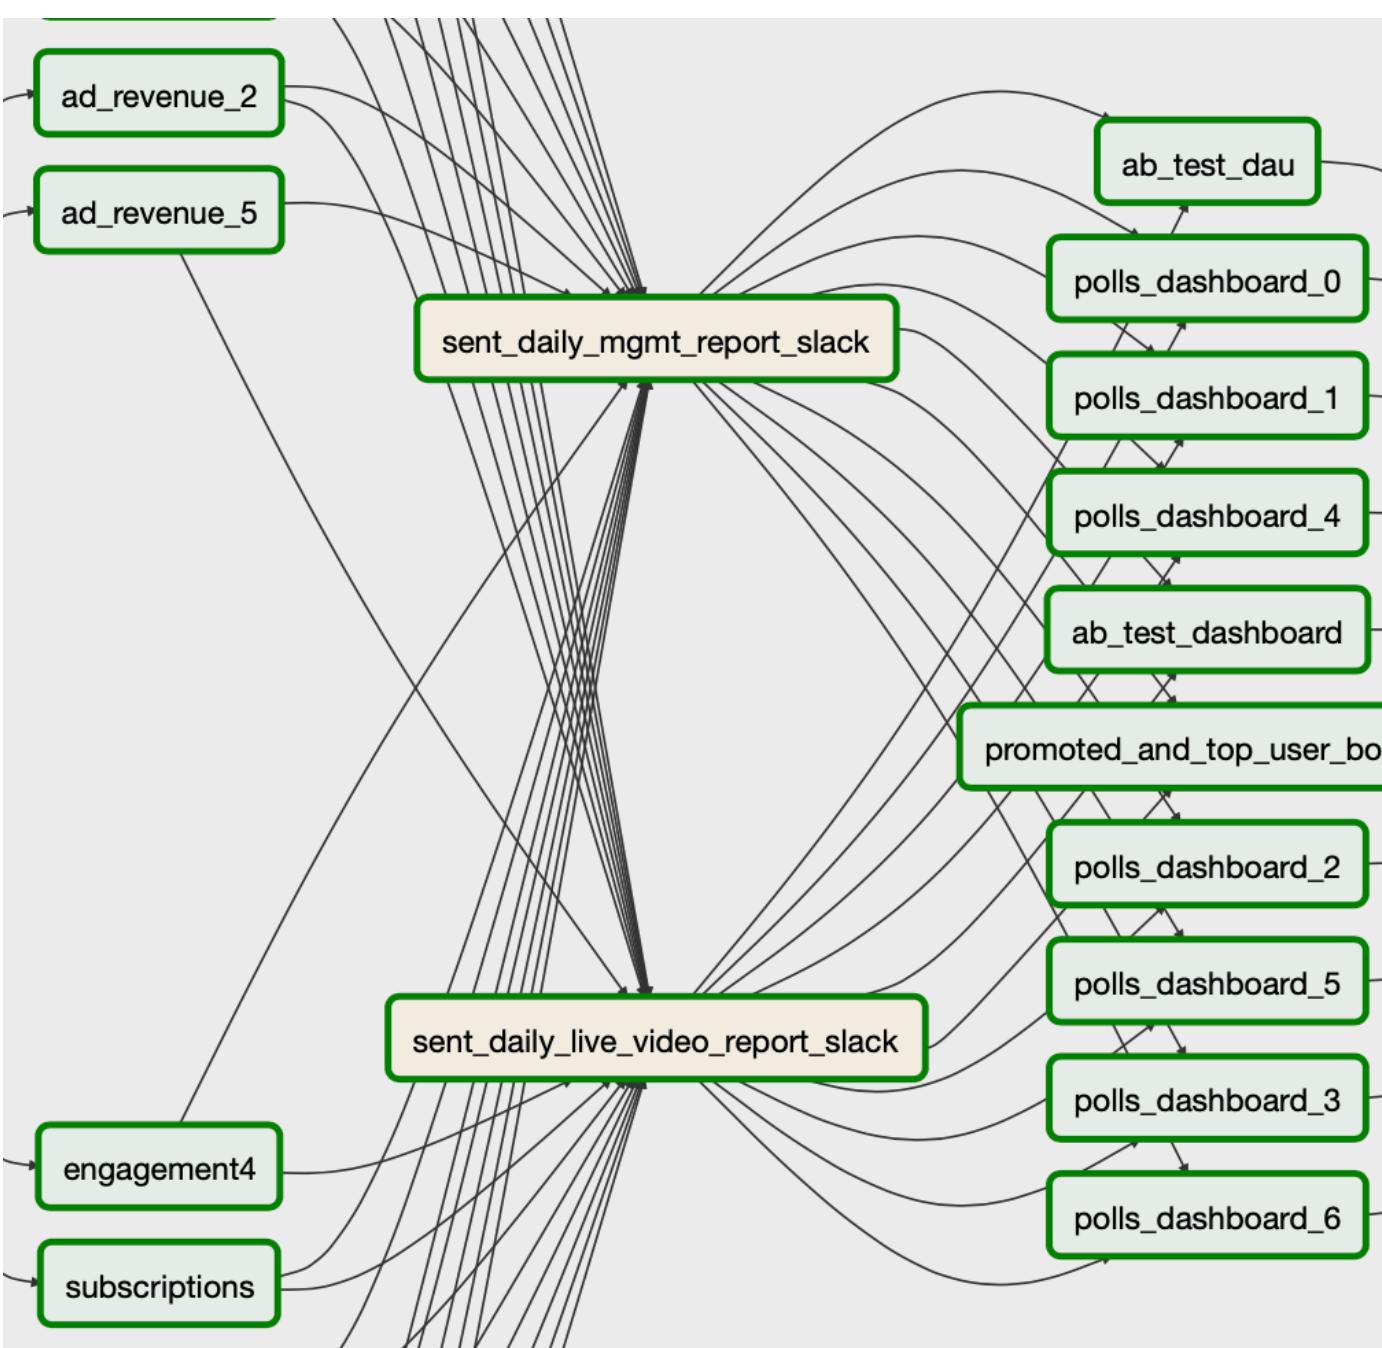


# The Core

Sub DAG



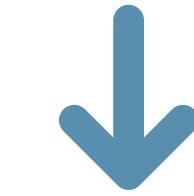
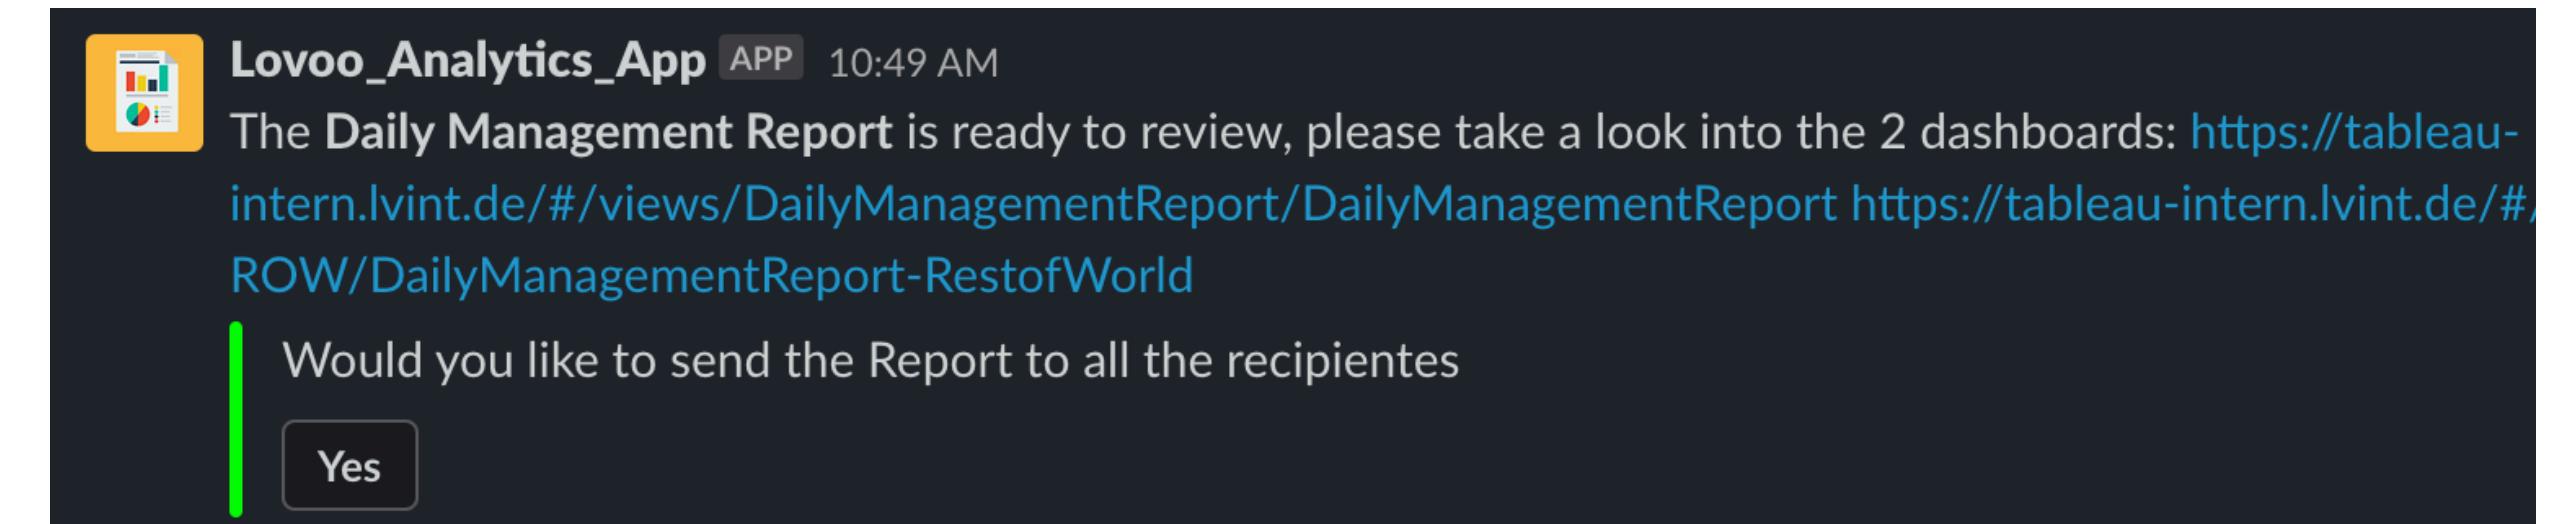
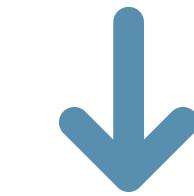
# Reports!



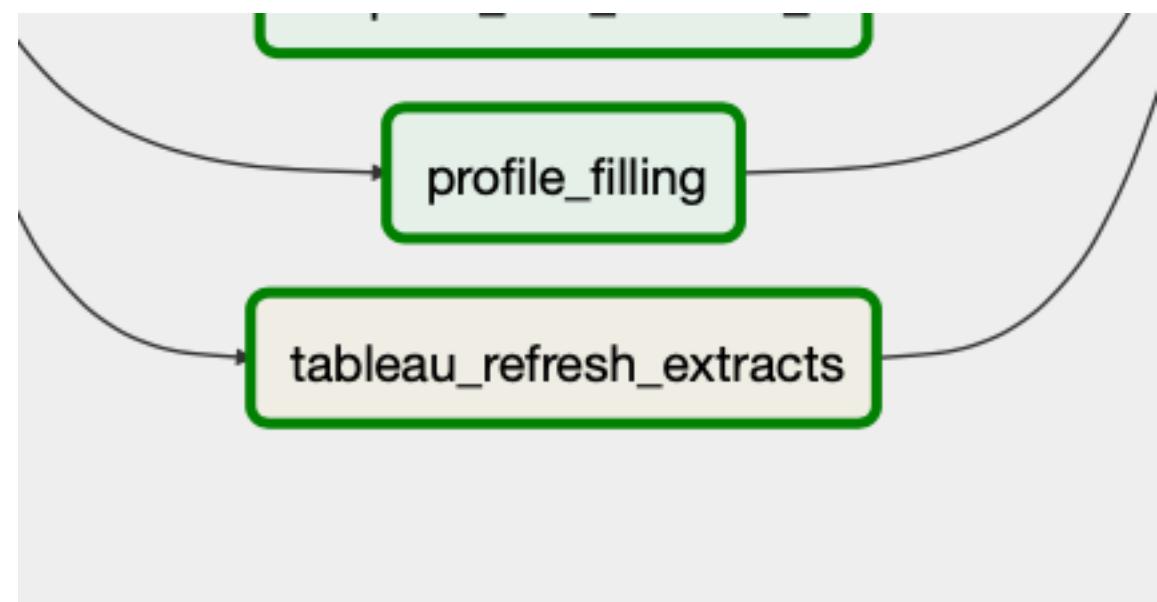
→



Slack Webhook



# Tableau Extracts



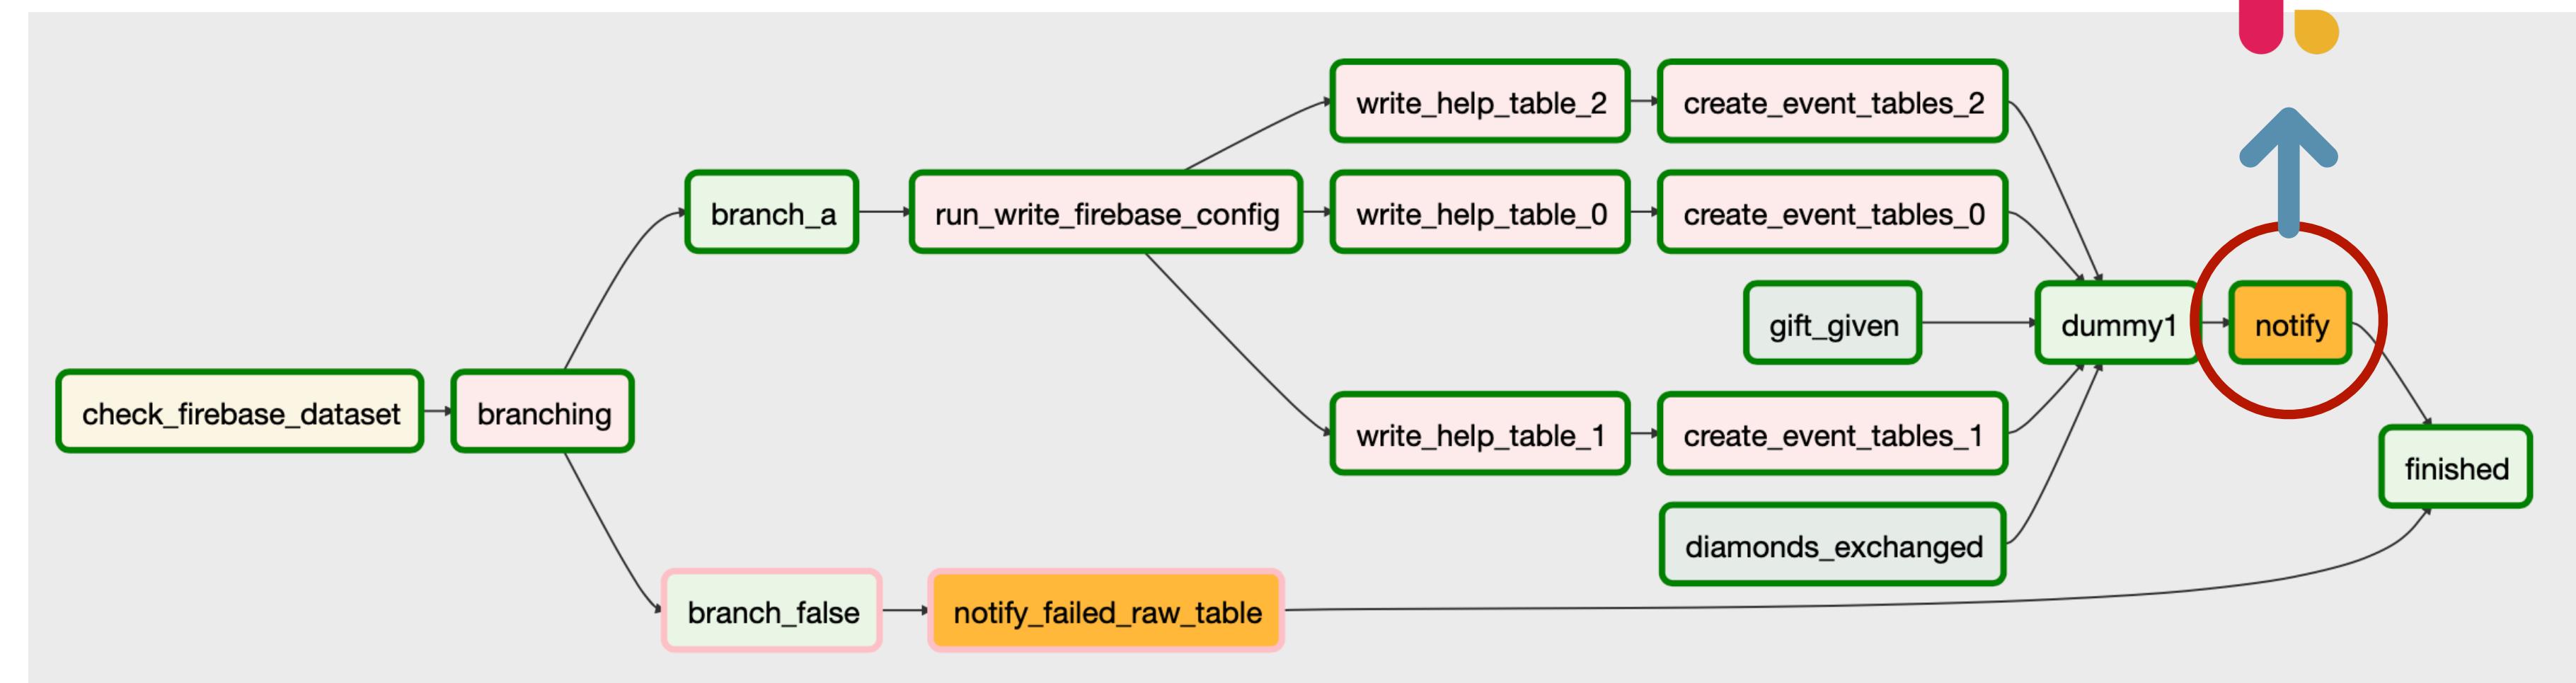
```
bash_command
1 curl -X POST http://35.205.226.12:8007
```



6  + a b | e a u°  
\$tabcmd runschedule "Daily Extract Refreshes (9:00 AM)"

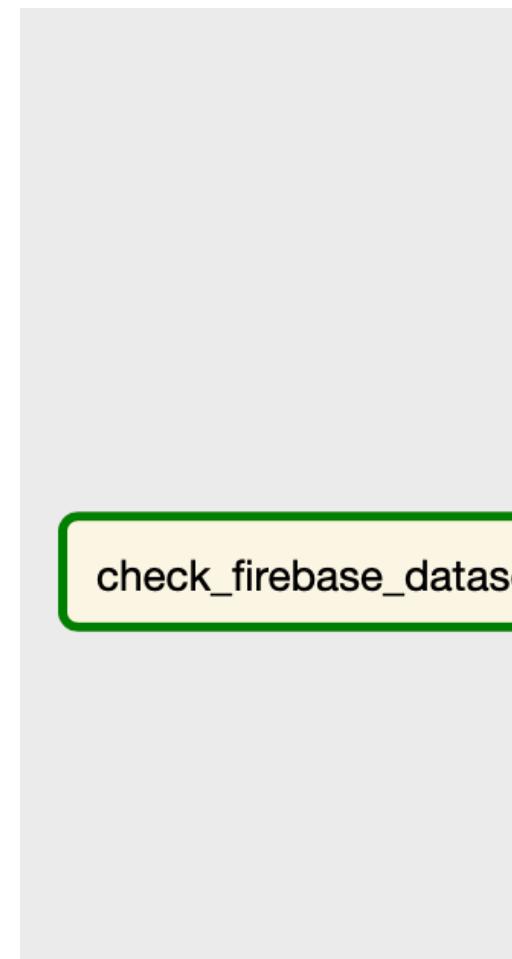
# Is Airflow finished?

by the way, this is branching...



COMMUNICATION IS VITAL...

# Is Airflow finished?



airflow-bot APP 6:47 AM bi\_firebase\_live\_events: Finished

airflow-bot APP 6:53 AM bi\_firebase\_importer: Finished

airflow-bot APP 7:04 AM 1/2 analytics\_jobs\_live-1\_user\_level\_live: Finished  
2/2 analytics\_jobs\_live-2\_agg\_level\_live: Finished  
Analytics Live Pipeline Completed analytics\_jobs\_live: Finished

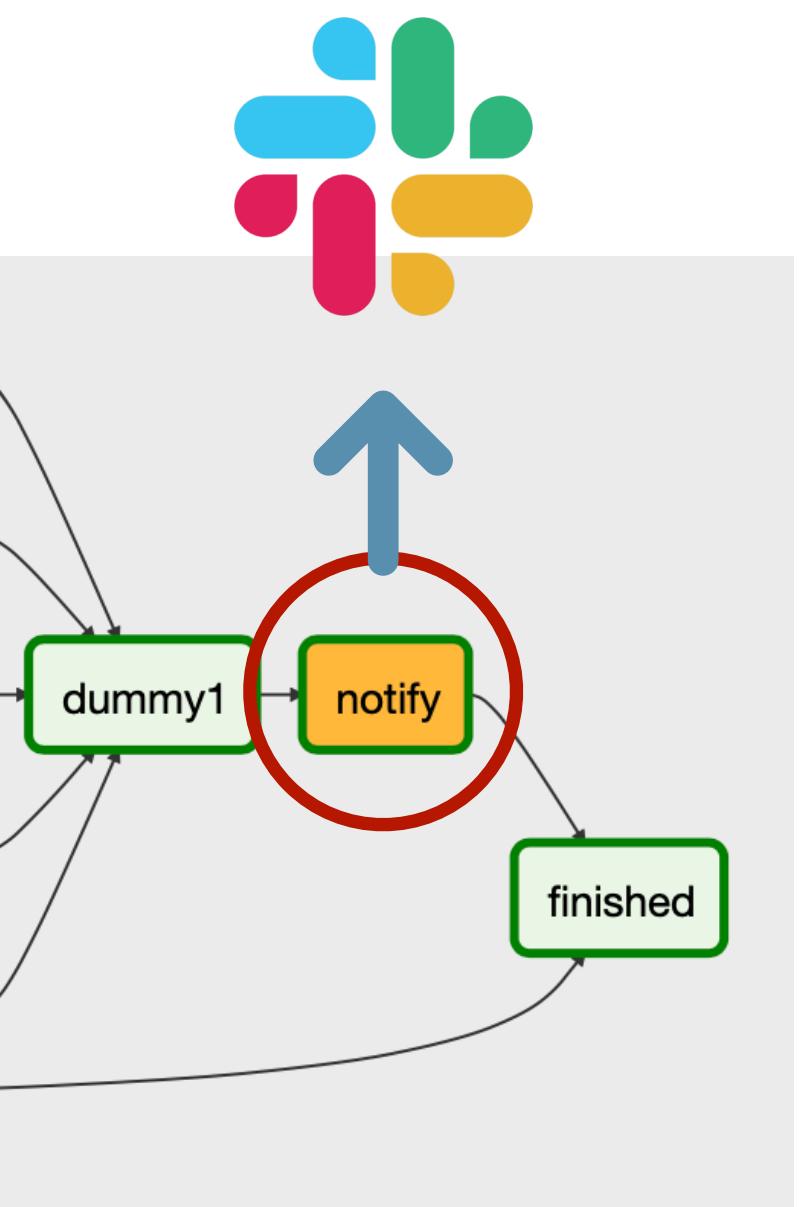
airflow-bot APP 7:21 AM 1/2 analytics\_jobs-1\_user\_level: Finished

airflow-bot APP 8:55 AM bi\_marketing\_events\_jobs: Finished

airflow-bot APP 9:21 AM 2/2 analytics\_jobs-2\_agg\_level: Finished  
Analytics Pipeline Completed analytics\_jobs: Finished

airflow-bot APP 1:26 PM AppSumer Lovoo: Finished @piotr.predkiewicz

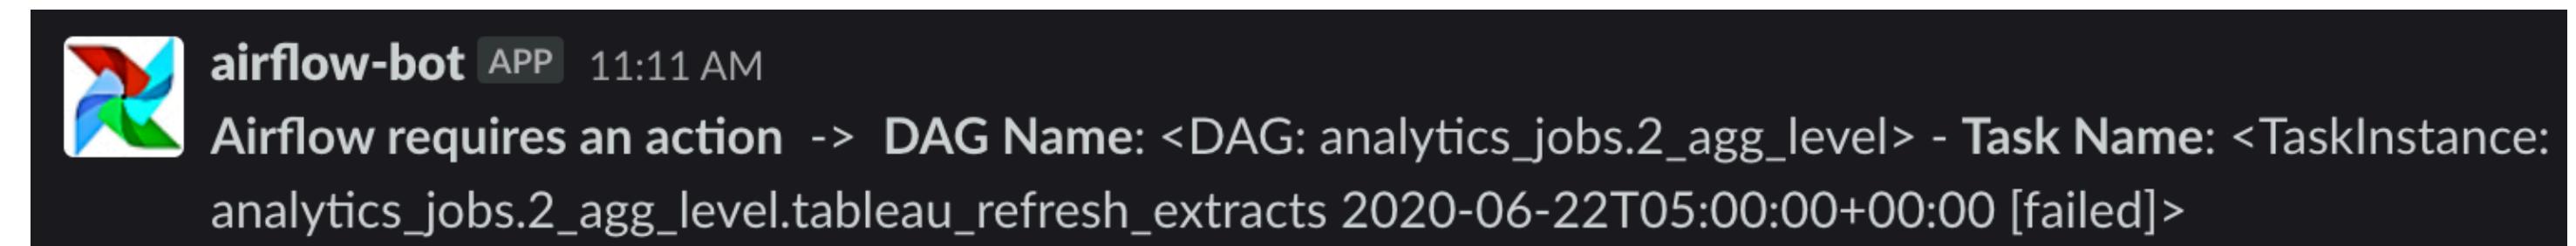
airflow-bot APP 2:35 PM AppSumer Hayi: Finished @piotr.predkiewicz



BECAUSE SH!]  
HAPPENS!

# Error Alerting

```
'on_failure_callback': on_failure_callback,
```



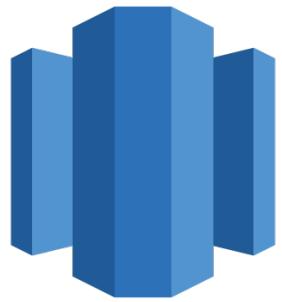
```
def on_failure_callback(context):
    operator = SlackAPIPostOperator(
        task_id='notify_fail',
        channel="#the_channel",
        token='your_Slack_bot_token',
        username='airflow-bot',
        text= str('*Airflow requires an action* {} Task: {}').format(
            str(context['dag']), str(context['task_instance'])))
    )
    return operator.execute(context=context)
```



# Integrating Data Sources

this code belongs to the DAG.py file

```
t1a = PythonOperator(  
    task_id='load_table_lovoo_transaction_groups_{}'.format(i),  
    python_callable=import_day_callable,  
    provide_context=True,  
    templates_dict={'exec_date': exec_date, 'table_name':'lovoo_transaction_groups'},  
    dag=dag)
```



# Integrating Data Sources

this code belongs to the DAG.py file

```
from BI.redshift_importer import import_datalake_redshift_data
def import_day_callable(**kwargs):
    exec_date = kwargs.get('templates_dict').get('exec_date')
    table_name = kwargs.get('templates_dict').get('table_name')
    return import_datalake_redshift_data(table_name,
                                         'load_job_dataframe_to_bq',
                                         exec_date=exec_date)
```



# Integrating Data Sources

this code belongs to the importer.py file

```
def postgreSQL_connection():
    try:
        # Using a Hook for getting the Redshift credentials from the Airflow connections -
        connection = BaseHook.get_connection("redshift_tmg")
        password = connection.password
        host = connection.host
        dbname = connection.schema
        user = connection.login
        port = connection.port

        conn = psycopg2.connect("dbname='{}' user='{}' port='{}' host='{}' password='{}'".format(dbname, user, port, host, password))
        cursor = conn.cursor()

    except Exception as e:
        print("I am unable to connect to the database: " + str(e))

    return cursor,conn
```



# Integrating Data Sources

this pseudo-code belongs to the importer.py file

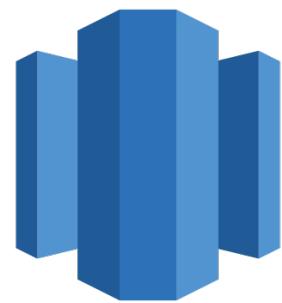
```
def import_datalake_redshift_data(table_name, method_type, exec_date, **kwargs):

    # Cursor & Connection
    cursor, conn = postgreSQL_connection()

    - Create dynamically a SQL query using the input parameters table_name and exec_date
    query = "select * from a_datalake.{} where data_updated_at::date >= '{}'".format(table_name, exec_date)

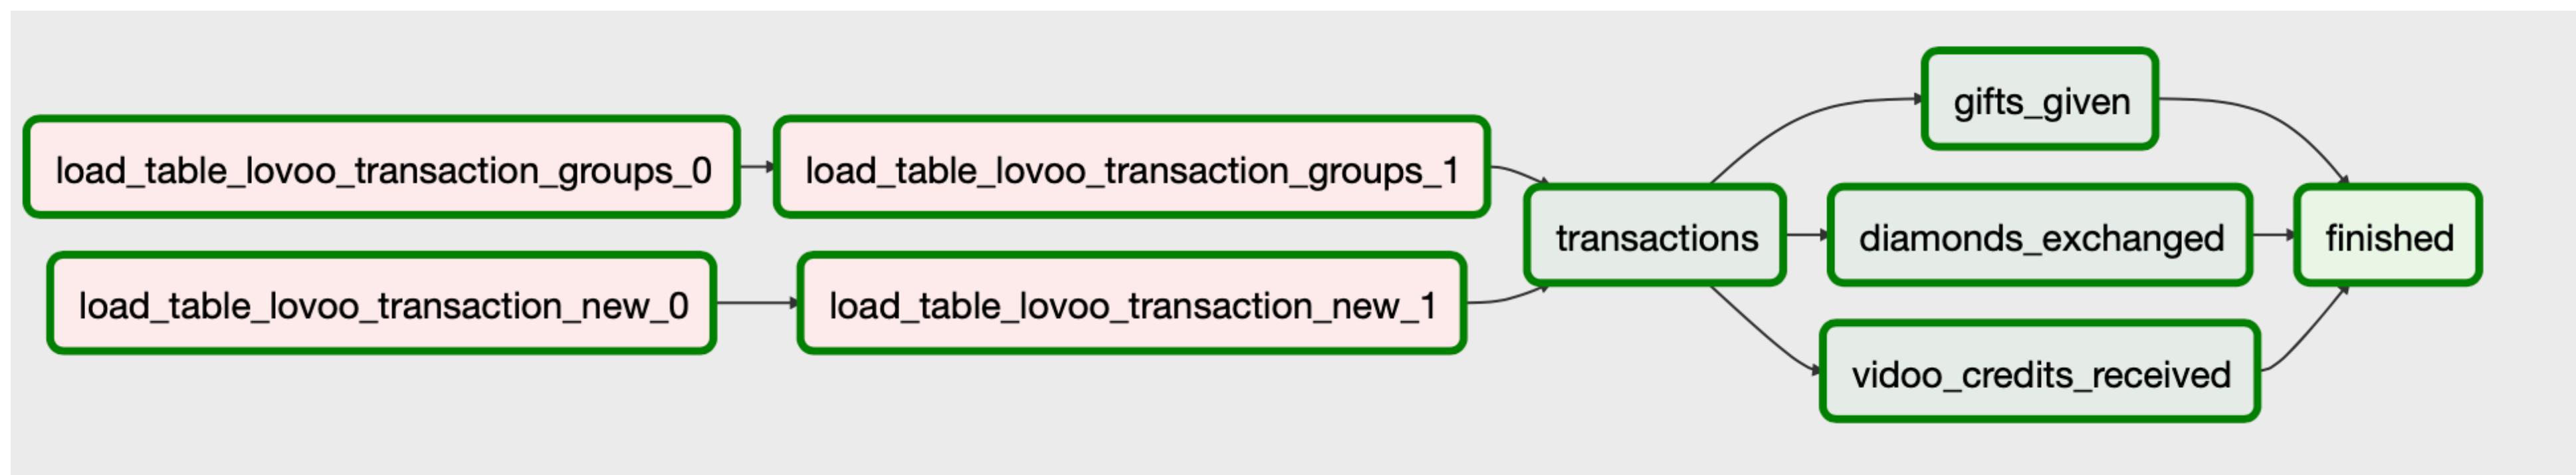
    - Use the query to request the data using the cursor
    cursor.execute(query)

    - use any method to upload the data to BigQuery
    df = cursor.fetchall()
    df = pd.DataFrame(df)
    job = client.load_table_from_dataframe(
        df, table_name, job_config=job_config
    )
    return whether it was successful or not
```



# Integrating Data Sources

2 Tables - 2 Days -> ELT in BQ



# Data Importers

- Redshift
- Firebase (very dynamic)
- Google Cloud Storage (Adjust, Merger)
- Appsumer, Shopify, Paypal, AppStore, Adyen
- S3 Storage

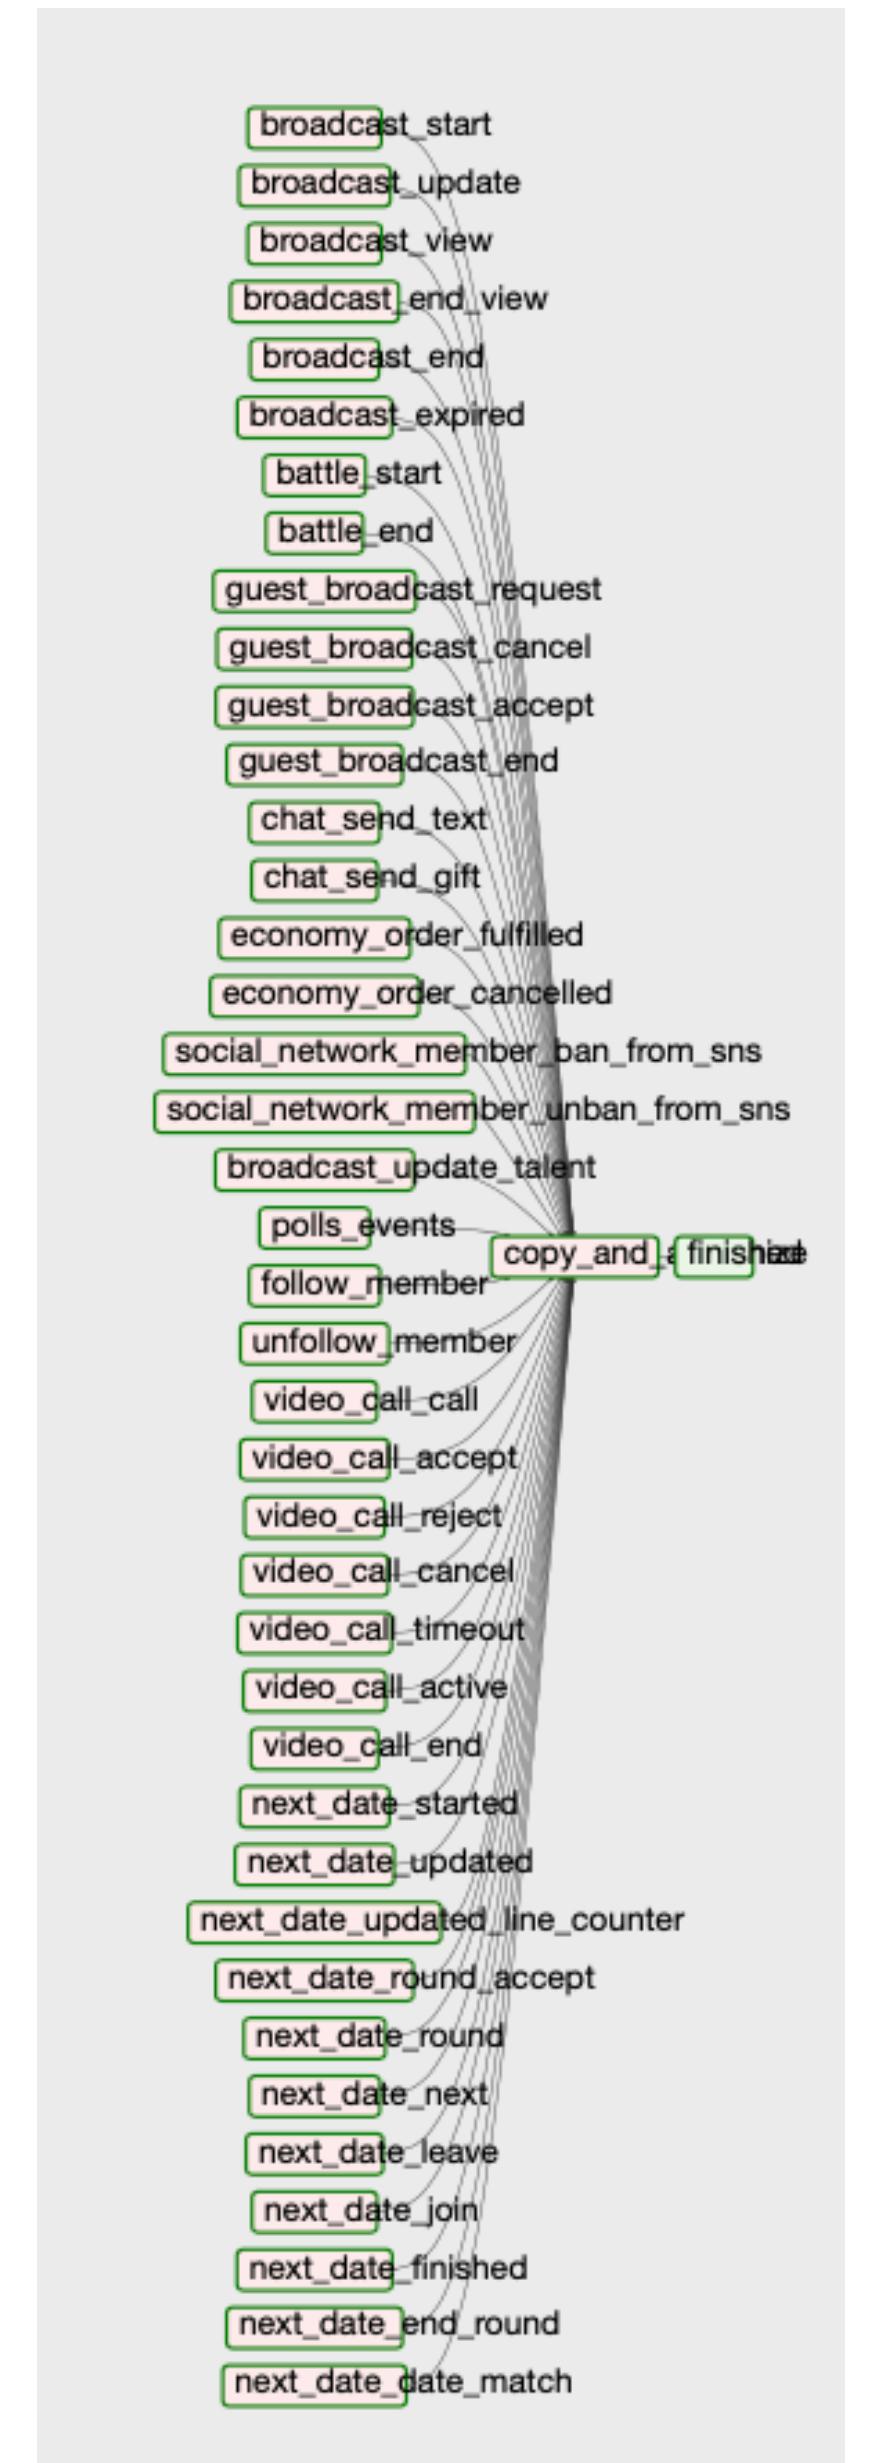
## AGENDA

---

1. Why we met?
2. How we met?
3. The first date!
4. Fun dates!
5. Is there any dynamic in between?
6. Recap and conclusion

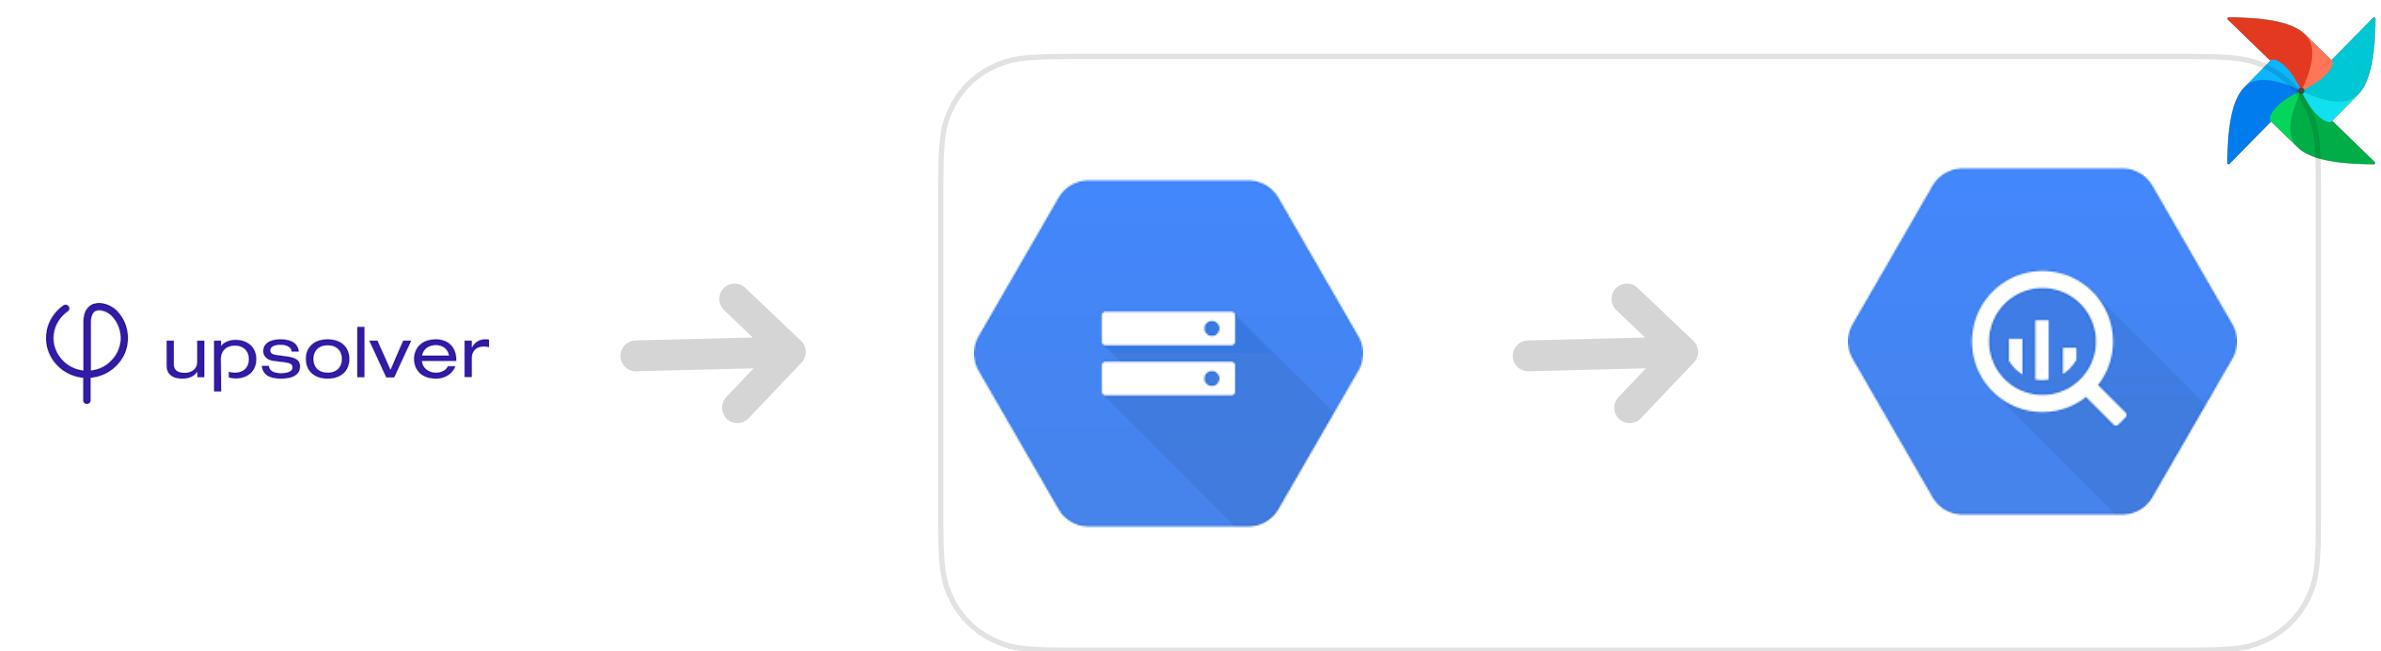
YES, VERY DYNAMIC...

# Creating Tasks Dynamically



# Creating Tasks Dynamically

1. Creating a plain text with meaningful structure
2. Create a task based on a PythonOperator
3. Define and write your Callable (your custom code)



# Creating Tasks Dynamically

JSON File

```
{  
    "broadcast_start": {  
        "id": "81ce53e3-2ca6-48a2-88f7-493f7fc2c364",  
        "format": "json"  
    },  
    "broadcast_update": {  
        "id": "de7cbe04-9b53-4b4c-bc18-4d065ed3830e",  
        "format": "json"  
    },  
    "broadcast_view": {  
        "id": "2a4a0093-baee-47a2-8817-aebdb469b1b1",  
        "format": "json"  
    },  
    "broadcast_end_view": {  
        "id": "3af65f80-94a3-42c4-8ed9-502134605d27",  
        "format": "json"  
    },  
    "broadcast_end": {  
        "id": "37fa9fdf-8384-41a7-83aa-1263814b3585",  
        "format": "json"  
    },  
    ...  
}
```

# Creating Tasks Dynamically

this code belongs to the DAG.py file

```
# Iterates over all the Mapping file and extracts the event name for generating all the task-events
for event_name in event_mapping:
    event_name_task = PythonOperator(
        task_id=str(event_name),
        provide_context=True,
        python_callable=run_import_day,
        templates_dict={'exec_date': exec_date, 'event_name': event_name,
                        'dataset': dataset, 'bucket_name': bucket_name},
        dag=dag)
```

# Creating Tasks Dynamically

this code belongs to the DAG.py file

```
# Function that will be called by the Python operator and will write a table partition in BQ
def run_import_day(**kwargs):
    dataset = 'events_input_analytics_tmgtbackend'
    bucket_name = 'lovoo-tmgt-transfer'
    import_gcs_to_bq(exec_date=kwargs.get('templates_dict').get('exec_date'),
                      event_name=kwargs.get('templates_dict').get('event_name'),
                      dataset=kwargs.get('templates_dict').get('dataset'),
                      bucket_name=kwargs.get('templates_dict').get('bucket_name'),
    )
```

# Creating Tasks Dynamically

this is your custom code (Pseudo-Code)

```
def import_gcs_to_bq(exec_date, event_name, dataset, bucket_name,**op_kwargs):  
  
    # read the structured JSON file  
    event_mapping = json.load(read_file)  
  
    # mapping the id and the event_name  
    id_event = event_mapping[event_name]['id']  
  
    # gathering the blobs inside the bucket - array of paths  
    path_array.append('gs://{}//{}//exec_date_file.json'.format(bucket_name, id_event))  
  
    # BigQuery Job to Load the JSON files to a table  
    load_job = bq_client.load_table_from_uri(  
        tuple(path_array), table_dest, job_config=job_config  
    )
```

# Creating Tasks Dynamically

```

def import_gcs_to_bq(exec_date, bucket_name, **op_kwargs):
    """This is a dynamic task creation example. It reads a JSON file containing
    event mappings and creates a BigQuery load job for each entry. The
    mapping file contains a list of events and their corresponding BigQuery
    table names. The code iterates through the list and creates a task for
    each entry. The tasks are triggered by a cron job that runs every hour.
    The tasks are triggered by a cron job that runs every hour.

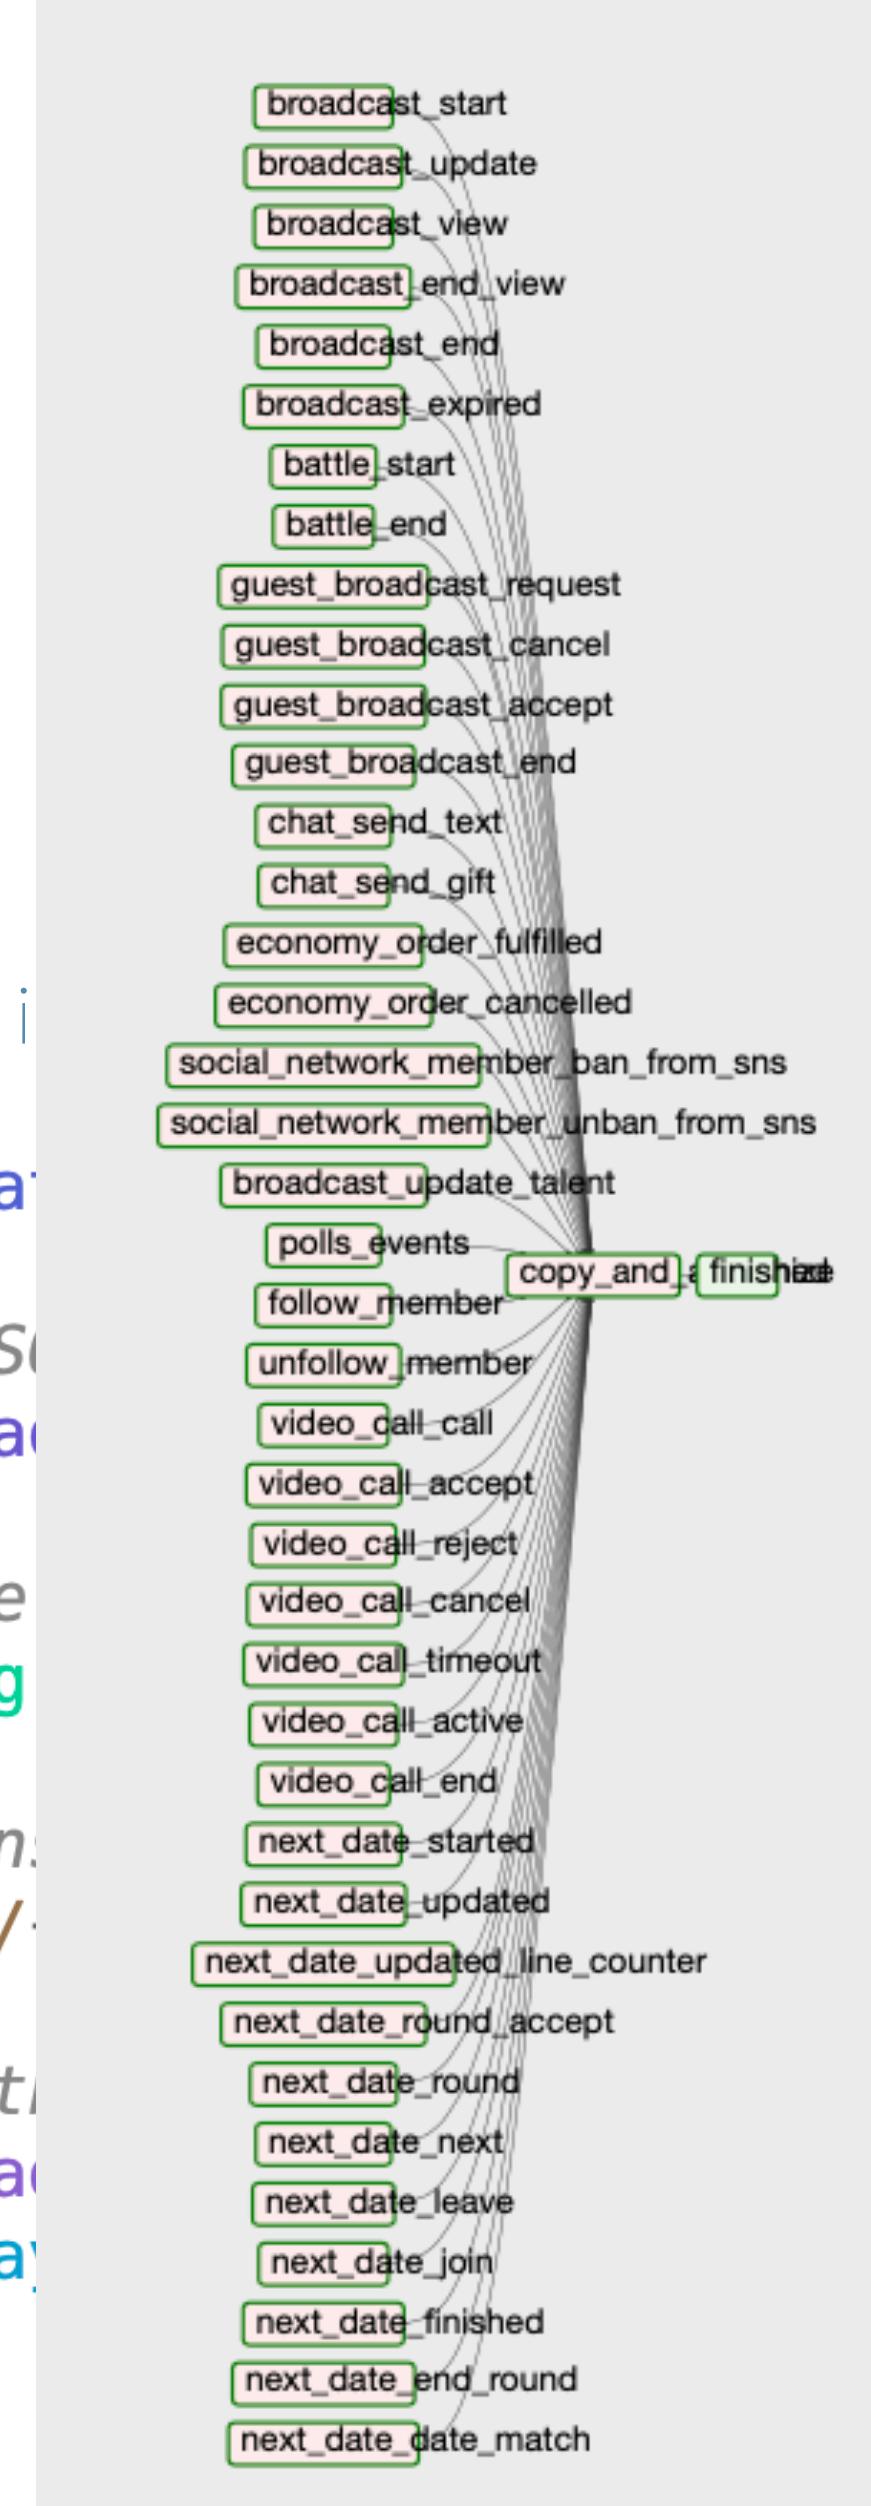
    # read the structured JSON file
    event_mapping = json.load(open('event_mapping.json'))

    # mapping the id and the event
    id_event = event_mapping['events']

    # gathering the blobs in the bucket
    path_array.append('gs://'+bucket_name)

    # BigQuery Job to Load the data
    load_job = bq_client.load_table_from_uri(
        tuple(path_array),
        'bigquery-public-data.eventify.events'
    )
    load_job.result()

```



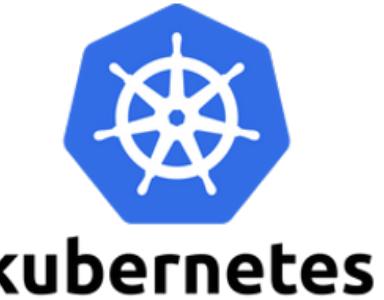
## AGENDA

---

1. Why we met?
2. How we met?
3. The first date!
4. Fun dates!
5. Is there any dynamic in between?
6. Recap and conclusion

# Recap and Conclusion

```
return KubernetesPodOperator(  
    startup_timeout_seconds=60 * 10, # we need seconds here as int, 10min now  
    task_id= 'appsumer_import_ ' + iso_date.replace('-', '_'),  
    namespace='default',  
    image=task_kwargs.get('image'),  
    cmd=task_kwargs.get('command'),  
    secrets=[appsumer_pass, service_account],  
    env_vars=env_vars,  
    name=task_kwargs.get('name'),  
    is_delete_operator_pod=True,  
    dag=dag,  
    dt=dt,  
    pool="appsumer_pool",  
    get_logs=True,  
    resources=resources,  
    affinity={  
        'nodeAffinity': {  
            # requiredDuringSchedulingIgnoredDuringExecution means in order  
            # for a pod to be scheduled on a node, the node must have the  
            # specified labels. However, if labels on a node change at  
            # runtime such that the affinity rules on a pod are no longer  
            # met, the pod will still continue to run on the node.  
        'requiredDuringSchedulingIgnoredDuringExecution': {  
            'nodeSelectorTerms': [{  
                'matchExpressions': [{  
                    'key': 'kuberunoperator',  
                    'operator': 'In',  
                    'values': [  
                        'true',  
                    ]  
                }]
```



# Recap and Conclusion

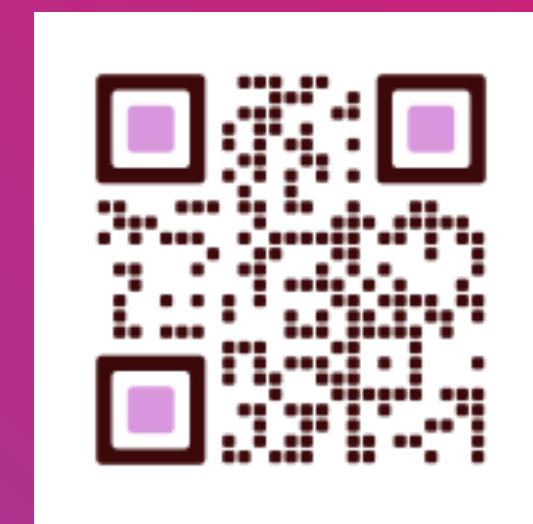
- Using an Alpha version (Google Composer) in Production was challenging!
- Focus on what's important - Google Cloud Composer
- Airflow leverages a bunch of Operators OOTB
- Always room for improvement
- No magic recipe to use - stay flexible

# Gracias.

July 16, 2020 Berlin - Germany

## Feedback and Questions

LinkedIn:



<https://www.linkedin.com/in/fandinohernandez/>

Email: [sergio.fandino@lovoo.com](mailto:sergio.fandino@lovoo.com)

